

ACT Automated Triage with Ansible Automation Platform

5W's | Information brief

ACT 0.6.18 | September 2026

BLUF. ACT (Agentic Command Tool) uses our GenAI.mil or AskSage models to do the **first pass of troubleshooting** on a Linux server automatically: it finds the **root cause** of a failure, shows the **evidence**, and proposes the **exact fix**. Run from Ansible Automation Platform (AAP), a fix is applied **only after a person approves it** (or, later, for a short list of pre-approved low-risk fixes), then verified, and every step is recorded. **Ready for a pilot** on the STIG Manager test server.

WHO

operates, approves, oversees

- **Operates:** the system administration / automation team. They own the AAP templates and the list of pre-approved fixes.
- **Approves fixes:** named approvers, given the Approve role on the AAP workflow.
- **Oversees:** the ISSO decides what data may go to the GenAI service.
- **Uses:** AAP 2.7 and GenAI.mil or AskSage, both already in place.

WHERE

what runs where

- **AAP 2.7** runs the jobs and keeps the job logs and approval records.
- **ACT runs on the server** being checked: one script copied in per run. No agent, no service, no open ports.
- **Model:** GenAI.mil (or AskSage) over HTTPS; each server must be able to reach it.
- **First target:** the STIG Manager test server. Not production, not the AAP server itself.

WHAT

an AI assistant for server triage

- Plans, runs **read-only** checks (service state, logs, disk, config), reads the real output, and reports **ROOT CAUSE, EVIDENCE, FIX, CONFIDENCE**.
- AAP runs quick health checks on every host; ACT is called **only where a check fails**.
- A workflow **pauses for approval**, applies exactly the approved fix, and checks it worked.
- Not autonomous root access, and not a replacement for monitoring.

WHY

the payoff

- Most outage time goes to the first diagnostic pass. ACT does it in **minutes, the same way every time, at any hour**.
- Faster restoration with a **person in control**: the approver sees the evidence and the exact command before anything changes.
- **Full audit trail**: every command, who approved, and the result.
- Uses tools we already have; cost follows problems, not fleet size.

WHEN

on failures, on a schedule, on demand

- When a health check fails: on a schedule (hourly or nightly), on demand, or right after patching.
- Healthy hosts: **no AI call, no cost**.
- Rollout in **4 gated phases**: pilot, scheduled reports, approval workflow, a few narrow auto-fixes.
- **Now:** Phase 1 pilot is built: stop the **nginx** or **stigman** container, ACT finds why, a person approves, it is restored.

SAFEGUARDS

how it stays safe


- **Diagnose-only by default:** unattended runs use read-only commands only.
- **Changes need approval** or an exact pre-approved pattern. Deletes, disk tools, firewall changes, reboots and file edits are **never** automatic.
- **Host names, IP addresses and accounts are pseudonymized** (replaced with placeholders) before anything goes to GenAI.
- **Every fix is verified** by a separate check. **Keys** are AAP credentials, never printed.

Decisions requested

1. Approve the **Phase 1 pilot**: diagnose and approval mode only, on the STIG Manager test server.
2. **ISSO concurrence** on sending service status and log lines to GenAI.mil (or AskSage), with passwords, tokens and keys removed and host names, IP addresses and user names pseudonymized first.
3. **Name the approvers** for the AAP workflow, and **the owner** of the pre-approved fix list.

Backup: how it works, safeguards, status

How one incident flows



If the approver denies, or nobody answers before the timeout, nothing changes. Healthy hosts stop at step 1.

Safeguard	What it means in practice
Diagnose-only default	Unattended ACT runs only proven read-only commands. Anything that would change the server is refused and reported as a proposed fix.
Human approval	The AAP workflow pauses. The approver reads ACT's root cause, evidence and the exact command; only that command can run afterwards.
Narrow pre-approval (later)	A few low-risk fixes (for example restarting one named service) may be allowed unattended, as exact patterns reviewed like code.
Never automatic	Deletes, disk and partition tools, firewall changes, reboots, file edits, chained or piped commands.
Verification	A fix is reported done only when a separate read-only check proves it. The playbook also re-checks the service itself.
Secrets and audit	The model key is an AAP credential and never printed. Every run records every command and how it was approved; AAP keeps job logs and approvals.

Topic	Status / risk and mitigation
Testing to date	Built and tested end to end against a simulated model . Phase 1 is the first run against GenAI.mil or AskSage in our environment.
Model can be wrong	Diagnoses carry evidence and a confidence level; approval and verification are the safety net, not the model.
Data sent to GenAI	The task text and command output (service status, log lines). Passwords, tokens and keys are removed. IP addresses, fully qualified names, the server's own names, names listed for it, local accounts and e-mail addresses are pseudonymized (replaced with consistent placeholders); ACT puts the real values back only on the server, to run commands. A short name ACT cannot recognize may still be sent. Released in ACT 0.6.18, which the pilot will use. Needs ISSO concurrence.
Prompt injection	A crafted log line could try to steer the model. In diagnose and approval mode it cannot change anything; pre-approved patterns stay narrow.
Network	Each server needs HTTPS access to the model endpoint.
Scope	Linux servers via AAP today. A Windows version of ACT exists; its AAP integration is not yet tested.

Rollout: each phase moves on only when its exit criteria are met

1. Pilot (now) Diagnose and approval mode on the STIG Manager test server. <i>Exit:</i> diagnoses judged useful on real incidents; data sent is acceptable.	2. Scheduled reports Diagnose-only on a wider group, hourly or nightly. <i>Exit:</i> the team uses the reports; few false alarms; usage within budget.	3. Approval workflow Diagnose, approve, apply the proposed fix. <i>Exit:</i> approved fixes verify; approvers trust the evidence shown.	4. Narrow auto-fixes A few repetitive, low-risk fixes pre-approved. <i>Ongoing:</i> every automatic fix is visible and verified; list reviewed monthly.
--	---	--	--