

# Site Automation

Operations runbooks: health checks, troubleshooting,  
ServiceNow, POA&M and patching - with AAP and, later, ACT

Setup guide for new AAP admins | site-automation 0.12.1 | October 2026

**What this is.** A step-by-step guide to setting up the site's operations runbooks in Ansible Automation Platform: health checks (system, compliance, MariaDB / MySQL, certificates, patching), troubleshooting, ServiceNow tickets, POA&M status and patching, and later ACT (GenAI). It starts with what each AAP object *is* and walks through every click, saying what each step does, why, and what you should see.

## How to read it

New to Ansible or AAP: read chapter 1 first, then do chapter 2 in order. Inventory already in AAP? Chapter 3 replaces the inventory parts of chapter 2. The later chapters are reference: open them when you build a workflow, read a finding, add ACT, change a setting, approve commands or set up the MariaDB / MySQL check.

Every setting, with its default and what it does - and where to put it - is in its own PDF, the *Settings Reference*: keep it open next to this one.

Rendered from [docs/](#) in the site-automation repository (version 0.12.1). The repository is the source of truth.

# Contents

|                                                                              |           |
|------------------------------------------------------------------------------|-----------|
| <b>1. Start here: the pieces and how they fit</b>                            | <b>7</b>  |
| What the runbooks do . . . . .                                               | 8         |
| The ideas behind Ansible, in plain words . . . . .                           | 8         |
| The AAP objects, and what each one is for . . . . .                          | 9         |
| How a health check works, start to finish . . . . .                          | 10        |
| The findings: the one idea that makes everything modular . . . . .           | 10        |
| Where things are . . . . .                                                   | 11        |
| The order to do things in . . . . .                                          | 11        |
| <b>2. Setting up AAP, step by step</b>                                       | <b>12</b> |
| Before you start . . . . .                                                   | 12        |
| Step 1. Put the repository in your work Git, with your inventory . . . . .   | 12        |
| Step 2. Credential types . . . . .                                           | 13        |
| Step 3. Credentials . . . . .                                                | 13        |
| Step 4. Project . . . . .                                                    | 14        |
| Step 5. Inventory (filled from the repository) . . . . .                     | 14        |
| Step 6. The first job template: Health check . . . . .                       | 15        |
| Step 7. Run it on one host and read the output . . . . .                     | 16        |
| Step 8. Tune the thresholds . . . . .                                        | 16        |
| Step 9. The other job templates . . . . .                                    | 17        |
| Step 10. ServiceNow . . . . .                                                | 21        |
| Step 11. Workflows and schedules . . . . .                                   | 22        |
| Step 12. Who can do what . . . . .                                           | 22        |
| If something goes wrong . . . . .                                            | 23        |
| <b>3. Using the inventory you already have in AAP</b>                        | <b>24</b> |
| The easy way: your settings files, already named after your groups . . . . . | 24        |
| How variables work in AAP (two minutes) . . . . .                            | 26        |
| Step 1. Look at what you have . . . . .                                      | 26        |
| Step 2. Create the groups the runbooks use . . . . .                         | 27        |
| Step 3. Settings for every host (inventory variables) . . . . .              | 28        |
| Step 4. Settings for particular groups . . . . .                             | 28        |
| Step 5. Settings for one host (only when one host is different) . . . . .    | 29        |
| Step 6. Use your inventory in the job templates . . . . .                    | 29        |

|                                                                                   |           |
|-----------------------------------------------------------------------------------|-----------|
| Step 7. Check that it worked . . . . .                                            | 30        |
| Which YAML files in the repository do I change? . . . . .                         | 30        |
| <b>4. ServiceNow: set up, test and verify</b>                                     | <b>31</b> |
| What you need . . . . .                                                           | 31        |
| Step 1. Ask the ServiceNow admins for an API account . . . . .                    | 31        |
| Step 2. The addresses . . . . .                                                   | 32        |
| Step 3. The network path from AAP . . . . .                                       | 33        |
| Step 4. Check from a shell first (optional, fastest) . . . . .                    | 33        |
| Step 5. The credential in AAP . . . . .                                           | 34        |
| Step 6. Your settings . . . . .                                                   | 34        |
| Step 7. The ServiceNow - test ticket template . . . . .                           | 35        |
| Step 8. Run it: first a dry run, then for real . . . . .                          | 35        |
| Step 9. Check the ticket in ServiceNow . . . . .                                  | 36        |
| When a step fails . . . . .                                                       | 37        |
| Next . . . . .                                                                    | 37        |
| <b>5. Workflows and schedules</b>                                                 | <b>39</b> |
| How to build any workflow, click by click . . . . .                               | 39        |
| Adding ACT to a workflow: what "on fail" means, and where site_act goes . . . . . | 40        |
| Linux workflows . . . . .                                                         | 41        |
| Windows workflows . . . . .                                                       | 46        |
| Linux and Windows together . . . . .                                              | 48        |
| VMware workflows . . . . .                                                        | 48        |
| Schedules . . . . .                                                               | 48        |
| Things that trip people up . . . . .                                              | 49        |
| <b>6. Dry runs: try anything, change nothing</b>                                  | <b>51</b> |
| Where to set it . . . . .                                                         | 51        |
| What each template does in a dry run . . . . .                                    | 52        |
| Examples . . . . .                                                                | 53        |
| What a dry run cannot tell you . . . . .                                          | 54        |
| <b>7. The runbooks: every check and what to do</b>                                | <b>55</b> |
| Health check (playbooks/health_check.yml) . . . . .                               | 55        |
| Troubleshoot (playbooks/troubleshoot.yml) . . . . .                               | 66        |
| Certificate report (playbooks/cert_report.yml) . . . . .                          | 67        |

|                                                                                          |           |
|------------------------------------------------------------------------------------------|-----------|
| POA&M status (playbooks/poam_status.yml) . . . . .                                       | 67        |
| ServiceNow tickets (playbooks/servicenow_tickets.yml) . . . . .                          | 68        |
| ServiceNow - test ticket (playbooks/servicenow_test_ticket.yml) . . . . .                | 68        |
| ServiceNow health (playbooks/servicenow_health.yml) . . . . .                            | 69        |
| Patch hosts (playbooks/patch_hosts.yml) . . . . .                                        | 69        |
| Apply approved ACT fix (playbooks/act_fix_approved.yml) . . . . .                        | 70        |
| VMware jobs and reports (playbooks/vm_*.yml) . . . . .                                   | 70        |
| Windows hosts . . . . .                                                                  | 70        |
| Not built yet (and why) . . . . .                                                        | 70        |
| <b>8. Adding ACT (GenAI) later</b>                                                       | <b>71</b> |
| How they fit: checks find WHAT, ACT finds WHY . . . . .                                  | 71        |
| The three levels . . . . .                                                               | 71        |
| Turn it on . . . . .                                                                     | 72        |
| Both endpoint formats and :probe (ACT 0.6.19 and newer) . . . . .                        | 75        |
| ACT for vCenter alarms (it runs on the AAP side) . . . . .                               | 78        |
| Write your own check (it works with the report, tickets and ACT automatically) . . . . . | 78        |
| Use ACT in a playbook of your own . . . . .                                              | 79        |
| What ACT is told, and what it cannot do . . . . .                                        | 80        |
| <b>9. Windows servers</b>                                                                | <b>81</b> |
| Step 1. WinRM over HTTPS on the servers . . . . .                                        | 82        |
| Step 2. The execution environment . . . . .                                              | 82        |
| Step 3. The credential . . . . .                                                         | 83        |
| Step 4. The Windows inventory . . . . .                                                  | 83        |
| Step 5. Test the connection . . . . .                                                    | 83        |
| Step 6. The job templates . . . . .                                                      | 84        |
| Step 7. Workflows and schedules . . . . .                                                | 85        |
| Start small: the pilot order . . . . .                                                   | 85        |
| Patching through Software Center . . . . .                                               | 85        |
| ACT on the servers (C:\ProgramData\act) . . . . .                                        | 87        |
| Settings: what each check looks at, and how to change it . . . . .                       | 88        |
| If something goes wrong . . . . .                                                        | 89        |
| Not in this release (next) . . . . .                                                     | 89        |
| <b>10. VMware jobs (vCenter): restart, snapshots, VLAN, reports, alarms and ACT</b>      | <b>90</b> |

|                                                                             |            |
|-----------------------------------------------------------------------------|------------|
| Before you start . . . . .                                                  | 91         |
| Which VMs a job acts on . . . . .                                           | 93         |
| The templates . . . . .                                                     | 94         |
| Snapshots: the report, the cleanup, and the approval workflow . . . . .     | 95         |
| Alarms: the report, and ACT's analysis . . . . .                            | 98         |
| Capacity planning: overall and per cluster . . . . .                        | 100        |
| Deploy a VM from a ServiceNow ticket . . . . .                              | 103        |
| ESXi security settings (every night): SSH off, timeouts, lockdown . . . . . | 110        |
| Email the report . . . . .                                                  | 112        |
| What you see . . . . .                                                      | 114        |
| Troubleshooting . . . . .                                                   | 115        |
| Limits . . . . .                                                            | 123        |
| <b>11. NetApp ONTAP health report</b>                                       | <b>124</b> |
| What the report shows . . . . .                                             | 124        |
| Before you start . . . . .                                                  | 125        |
| Set it up . . . . .                                                         | 126        |
| Settings . . . . .                                                          | 127        |
| Troubleshooting . . . . .                                                   | 128        |
| Limits . . . . .                                                            | 129        |
| <b>12. Emailed reports: the formatted email and its look</b>                | <b>130</b> |
| Which jobs email . . . . .                                                  | 131        |
| Turning it on . . . . .                                                     | 132        |
| Green or red: when a report finds problems . . . . .                        | 133        |
| What a report looks like . . . . .                                          | 133        |
| Changing the look . . . . .                                                 | 133        |
| Emailing a report from another playbook . . . . .                           | 134        |
| Troubleshooting . . . . .                                                   | 134        |
| <b>13. Secrets: where they live</b>                                         | <b>136</b> |
| Recommendation . . . . .                                                    | 136        |
| Where each secret lives . . . . .                                           | 137        |
| Rules for every playbook in this repository . . . . .                       | 138        |
| The real trust boundary . . . . .                                           | 139        |
| Keeping secrets out of Git . . . . .                                        | 139        |

|                                                                        |            |
|------------------------------------------------------------------------|------------|
| Rotation checklist . . . . .                                           | 140        |
| <b>14. Approved commands: what ACT may run by itself</b>               | <b>141</b> |
| Read this first: two safety rules that you cannot switch off . . . . . | 141        |
| How a command becomes "approved": the three ways . . . . .             | 141        |
| Add an approved command, step by step . . . . .                        | 142        |
| How to verify it worked . . . . .                                      | 144        |
| Reference: what the guard rails are . . . . .                          | 145        |
| <b>15. The MariaDB and MySQL check, step by step</b>                   | <b>147</b> |
| The short version . . . . .                                            | 147        |
| 1. Which hosts the check runs on . . . . .                             | 147        |
| 2. Tell it where MariaDB is . . . . .                                  | 148        |
| 3. Where to put the settings . . . . .                                 | 149        |
| 4. The login: the "MariaDB monitor" credential . . . . .               | 149        |
| 5. What it checks, and the limits you can change . . . . .             | 150        |
| How to verify . . . . .                                                | 151        |
| What was tested, and what was not . . . . .                            | 152        |
| If it does not work . . . . .                                          | 153        |
| <b>Appendix: credential types to paste into AAP</b>                    | <b>154</b> |
| ServiceNow API (servicenow_api.yml) . . . . .                          | 154        |
| MariaDB monitor (mariadb_monitor.yml) . . . . .                        | 154        |
| Keystore password (keystore_password.yml) . . . . .                    | 155        |
| STIG Manager API (stigman_api.yml) . . . . .                           | 155        |
| ACT model key (act_model_key.yml) . . . . .                            | 155        |
| ACT GenAI beta key (act_genai_beta_key.yml) . . . . .                  | 156        |
| STIG Manager database (stigman_database.yml) . . . . .                 | 156        |
| TLS certificate (tls_certificate.yml) . . . . .                        | 157        |
| Container registry login (hosts) (registry_login.yml) . . . . .        | 157        |
| Teams webhook (teams_webhook.yml) . . . . .                            | 157        |
| SMTP relay (smtp_relay.yml) . . . . .                                  | 158        |

# 1. Start here: the pieces and how they fit

## Confused by all the moving pieces?

Read `HOW_IT_FITS_TOGETHER.md` (or its PDF in [docs/pdf/](#)). It shows what the code, the inventory, the credentials and the survey each do, how a setting travels from a YAML file to a host, and how to update the repository safely.

This repository holds **runbooks**: jobs that check Linux servers and help troubleshoot them. You run them from **Ansible Automation Platform (AAP)**, by hand or on a schedule. This page explains the few ideas you need. `SETUP_AAP.md` then walks you through the setup, one click at a time.

## What the runbooks do

| Runbook                           | What it answers                                                                                                                                                                                                                                          | When it runs                                          |
|-----------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------|
| <b>Health check</b>               | Is anything wrong on these servers right now? Disk, mounts, services, load/memory, time, network, logging, SELinux, fapolicyd, auditd, accounts, certificates, patching, MariaDB / MySQL                                                                 | Daily (the "daily" set) and weekly (the "weekly" set) |
| <b>Database health</b>            | Are the MariaDB / MySQL databases (on the host or in podman containers) running, with connection room, no stuck replication, and disk space?                                                                                                             | Hourly or daily, on its own                           |
| <b>Troubleshoot</b>               | Someone reports a problem on a host. What does it look like?                                                                                                                                                                                             | On demand                                             |
| <b>Certificate report</b>         | Which certificates expire soon, on every server, in one list?                                                                                                                                                                                            | Weekly                                                |
| <b>POA&amp;M status</b>           | Which POA&M items are overdue or due soon? Which open CAT I/II findings have no POA&M item?                                                                                                                                                              | Weekly                                                |
| <b>ServiceNow tickets</b>         | Opens or updates a ticket for each problem a check found. Notes the ones that cleared                                                                                                                                                                    | After every check (as a workflow step)                |
| <b>ServiceNow health</b>          | Is ServiceNow answering? Are the MID Servers up?                                                                                                                                                                                                         | Hourly                                                |
| <b>Patch hosts</b>                | Installs updates, and proves every service came back. Reboots only with <code>automatic_restarts: yes</code>                                                                                                                                             | Monthly, after an approval                            |
| <b>Apply approved ACT fix</b>     | Applies exactly the fix a person approved                                                                                                                                                                                                                | Only inside a workflow, after an approval             |
| <b>NetApp ONTAP health report</b> | Is any node, disk, port, volume, aggregate, SnapMirror or certificate in trouble on the storage clusters? (NETAPP.md)                                                                                                                                    | Daily                                                 |
| <b>VMware jobs</b>                | Restart, shut down, snapshot, notes or change the VLAN of VMs through vCenter. Reports: secure boot, datastores, snapshots (with an approved cleanup), alarms and events, and ACT's analysis of the alarms: likely cause, fix and confidence (VMWARE.md) | On demand; the reports daily or weekly                |

Every check is **read-only**. It looks and reports, and changes nothing. Patching changes servers, and waits for an approval first; the VMware jobs change only what an operator asked for, and never the AAP server's VM.

## The ideas behind Ansible, in plain words

- **Host:** a server Ansible connects to, over SSH, the way you would.

- **Inventory:** the list of hosts, sorted into **groups** such as `stigman` (the STIG Manager servers) or `mariadb`. The built-in group `all` holds every host. A group can also carry **settings (variables)**, for example "on the database hosts, the MariaDB container is called servicenow-mariadb".
- **Task:** one step, such as "run `df` and read the result". A **module** is the tool a task uses (`command`, `shell`, `uri`, `set_fact`...).
- **Role:** a folder of tasks that does one job, plus its settings. `roles/check_disk/` is "check the disks". Its settings, with their defaults, are in `roles/check_disk/defaults/main.yml`.
- **Playbook:** the file you run. It says which hosts, and which roles or tasks to run on them. `playbooks/health_check.yml` says "on every host, run the checks I was given".
- **become:** run as root through sudo. The checks need it to read root-only information (the audit log, `/etc/shadow`).
- **Check mode:** a dry run. Ansible shows what it would change and changes nothing. Every runbook here works in check mode, and the checks give the same answers in both modes.
- **Idempotent:** running it twice does no harm. The second run changes nothing.

## The AAP objects, and what each one is for

AAP is a web front end that runs playbooks for you. It keeps the passwords, records every run, and runs things on a schedule. You set up a handful of objects once. Each one points at the next:

| AAP object                   | What it is                                                                                                                                       | Here                                                                                                                                                               |
|------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Project</b>               | A copy of a Git repository that AAP keeps in sync                                                                                                | <code>site-automation</code> (this repository)                                                                                                                     |
| <b>Inventory</b>             | The hosts and groups, and their settings                                                                                                         | <code>Linux servers</code> , filled from <code>inventories/site/</code> in this repository, or the inventory you already have in AAP (USING_YOUR_AAP_INVENTORY.md) |
| <b>Credential type</b>       | A form for a kind of secret, and how to hand it to the job                                                                                       | <code>ServiceNow API</code> , <code>MariaDB monitor</code> , <code>ACT model key</code> ... ( <code>aap/credential_types/</code> )                                 |
| <b>Credential</b>            | One filled-in form: the actual secret. Nobody can read it back                                                                                   | <code>Linux ssh (sudo)</code> , <code>ServiceNow API</code> ...                                                                                                    |
| <b>Execution environment</b> | The container a job runs in (Ansible, Python, tools)                                                                                             | The <b>Minimal</b> execution environment is enough. This repository needs no extra collections                                                                     |
| <b>Job template</b>          | "Run this playbook, on this inventory, with these credentials" plus the questions to ask ( <b>survey</b> ). This is what you launch              | <code>Health check</code> , <code>Troubleshoot</code> , <code>Patch hosts</code> ...                                                                               |
| <b>Workflow</b>              | Job templates chained together. The next step can depend on the last one's result (success, failure, always). It can stop for an <b>approval</b> | <code>Daily health</code> = Health check, then ServiceNow tickets                                                                                                  |
| <b>Schedule</b>              | "Launch this template or workflow at these times"                                                                                                | Daily health every day at 06:00                                                                                                                                    |

Three more words you will see:

- **Survey:** the questions a template asks when you launch it, such as "which checks?". Each answer becomes a variable for the playbook.
- **Limit:** run on only some hosts or groups, for example `snowdb01.example.mil` or `mariadb_hosts`.
- **Artifacts:** what a job hands to the next step of a workflow. The health check hands its findings to the ServiceNow step this way (`set_stats` in the playbook).

## How a health check works, start to finish

1. You click **Launch** on the *Health check* template (or its schedule fires).
2. AAP syncs the project, starts the execution environment, and connects to each host over SSH with the Machine credential.
3. On each host, each chosen check runs read-only commands and turns what it sees into **findings**. One finding is one thing that is wrong, in plain words. It includes a command you can paste to look further:

```
[CRITICAL] disk: /var is 96% full (critical at 95%), 1.2 GiB free of 40.0 GiB
look: timeout 60 du -xh --max-depth=2 /var 2>/dev/null | sort -h | tail -15
```

4. The job prints a summary per host and for the whole fleet. It then **fails the hosts that have findings**. The red status is on purpose: it is how a workflow knows it should open tickets (or ask ACT, or ask for approval).
5. In the *Daily health* workflow, the *ServiceNow tickets* step runs next. It opens a ticket for each finding, or updates the ticket that is already open.

## The findings: the one idea that makes everything modular

Every check produces the same kind of finding:

| Field                 | Example                                         | Used by                                                                    |
|-----------------------|-------------------------------------------------|----------------------------------------------------------------------------|
| <code>check</code>    | <code>disk</code>                               | reports, tickets                                                           |
| <code>id</code>       | <code>disk:/var</code>                          | ServiceNow updates the same ticket every run, instead of opening a new one |
| <code>severity</code> | <code>critical</code> or <code>warning</code>   | pass/fail, ticket urgency                                                  |
| <code>summary</code>  | <code>/var is 96% full (critical at 95%)</code> | people                                                                     |
| <code>hint</code>     | <code>du -xh --max-depth=2 /var</code><br>...   | people, and ACT, which runs it to collect evidence                         |

The report, the ServiceNow tickets and ACT only ever read findings. That is why:

- a new check you write shows up in reports and tickets with no other change;
- ACT works on top of **any** check. Add `use_act=true` and ACT reads the findings and the hint commands' output, then explains the root cause (ADDING\_ACT.md).

## Where things are

| Path                                                                                                                        | What                                                                                                                                                                                          |
|-----------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>playbooks/</code>                                                                                                     | the files AAP runs (one per job template)                                                                                                                                                     |
| <code>playbooks/group_vars/</code>                                                                                          | <b>your settings</b> , one file per AAP group: edit these (VARIABLES.md, USING_YOUR_AAP_INVENTORY.md)                                                                                         |
| <code>roles/check_*/</code>                                                                                                 | one folder per check. Its settings are in <code>defaults/main.yml</code>                                                                                                                      |
| <code>roles/site_findings/</code>                                                                                           | starts, reports and publishes findings, and decides pass/fail                                                                                                                                 |
| <code>roles/servicenow/</code> , <code>roles/troubleshoot/</code> ,<br><code>roles/patch/</code> , <code>roles/poam/</code> | the other runbooks                                                                                                                                                                            |
| <code>roles/site_act/</code>                                                                                                | the bridge to ACT (GenAI)                                                                                                                                                                     |
| <code>inventories/example/</code>                                                                                           | a sample inventory and settings. Copy it to <code>inventories/site/</code>                                                                                                                    |
| <code>aap/credential_types/</code>                                                                                          | the credential types to create in AAP                                                                                                                                                         |
| <code>poam/poam.csv</code>                                                                                                  | your POA&M list (for the POA&M runbook)                                                                                                                                                       |
| <code>docs/</code>                                                                                                          | this guide, SETUP_AAP.md, USING_YOUR_AAP_INVENTORY.md, VARIABLES.md, RUNBOOKS.md, MARIADB.md, WORKFLOWS_AND_SCHEDULES.md, ADDING_ACT.md, APPROVED_COMMANDS.md; PDFs in <code>docs/pdf/</code> |

## The order to do things in

1. **Set up and run one health check by hand** (SETUP\_AAP.md, steps 1 to 8). Your inventory is already in AAP? Use USING\_YOUR\_AAP\_INVENTORY.md for the inventory part. Read its output. Adjust thresholds in your settings files (VARIABLES.md) until the findings are the ones you care about.
2. **Add the other job templates** (step 9) and try each once.
3. **Add ServiceNow** (step 10), and build the *Daily health* workflow (WORKFLOWS\_AND\_SCHEDULES.md).
4. **Schedule** the workflows.
5. **Later: add ACT** to the runbooks you trust (ADDING\_ACT.md).

## 2. Setting up AAP, step by step

Every step says **what** you are doing, **why**, the **clicks**, and what you should **see** when it worked. Menu names are from AAP 2.7. Older versions (2.5, 2.6) may word a label or a menu differently: follow the meaning, not the exact word. If you have not yet, read `START_HERE.md` first. It explains the words used here.

### Before you start

You need:

- An AAP account that can create projects, inventories, credentials and templates in your organization (an **Organization admin** does it all).
- A **Linux account on the servers** that AAP logs in as, for example `svc_aap`. It needs an SSH key and sudo. With `NOPASSWD` sudo, add it to `check_accounts_nopasswd_allowed`, or the accounts check will report it.
- A **Git repository at work** that AAP can reach (GitLab, Bitbucket, GitHub Enterprise).
- For ServiceNow: an **API account** (step 10).

### Step 1. Put the repository in your work Git, with your inventory

#### Already have your inventory in AAP?

Then you do not need `inventories/site/`. Push the repository without it (skip items 1-3 below), skip step 5, and follow `USING_YOUR_AAP_INVENTORY.md` for the groups and settings to add.

**What.** Copy this repository into your work Git, and create your real inventory in it. (Already pushed an earlier release at work? Update it with `scripts/update-from-release.ps1` in VS Code, as in `HOW_IT_FITS_TOGETHER.md`, "Updating the repository at work". Then skip to step 2.)

**Why.** AAP reads playbooks from Git, never from your laptop. The inventory lives in the same repository, so host lists and thresholds are versioned and reviewed like code. Every change has a who, a when and a why.

1. Extract the release zip (right-click › **Extract All**). In the extracted folder, copy the folder `inventories\example` and name the copy `inventories\site`.
2. Edit `inventories/site/hosts.yml`. Replace the example host names with yours and keep the group names: the playbooks use them. A host can be in several groups.
3. Edit `inventories/site/group_vars/all.yml`. These settings apply to every host: the LDAP and SIEM servers each host must reach, the account AAP logs in as (`svc_aap`), and the thresholds. Every setting and its default is listed in `VARIABLES_REFERENCE.md`. **Two places can hold a setting:** this inventory file (it becomes AAP's *inventory* Variables box) and the project's own settings files `playbooks/group_vars/*.yml` (`all.yml` is only switched-off placeholders; a few group files, such as `stigman.yml` and `aap.yml`, ship real values). If a setting is in both, the project file wins. Use one place per setting: `VARIABLES.md` explains the order.
4. Put it in your work Git, with VS Code:
  - Create an empty repository on your Git server (for example `site-automation`) and copy its clone URL.
  - In VS Code: **Source Control › Clone Repository** › paste the URL › pick a folder, such as `C:\git\site-automation`.
  - Copy everything from the extracted release folder into that folder.

- In **Source Control**: type a message, click **Commit** (answer Yes to "stage all changes"), then **Sync Changes / Publish Branch**.

The same from a command line: `git add -A, git commit -m "site-automation <version>",` then `git push`.

5. Note the two values AAP needs in step 4: the **clone URL** (the one you pasted above) and the **branch** (shown at the bottom left of VS Code, usually `main`). AAP does not read them from your copy; you type them in once.

**You should see** the files in your Git server's web page, including `inventories/site/`.

## Step 2. Credential types

**What.** Teach AAP the kinds of secrets these runbooks use, one form per kind.

**Why.** AAP ships with types for SSH (Machine) and Git (Source Control). The ServiceNow login and the others are custom. A credential type says which fields the form has and how the secret reaches the job. Every type here passes it as an **environment variable** in the execution environment. The playbook must ask for it on purpose, and it never shows in job output or Git.

For each file in `aap/credential_types/` that you need now:

1. **Automation Execution > Infrastructure > Credential Types > Create credential type.**
2. **Name:** the name at the top of the file (for example `ServiceNow API`).
3. **Input configuration:** paste the `inputs:` part of the file, *without* the `inputs:` line.
4. **Injector configuration:** paste the `injectors:` part, *without* the `injectors:` line.
5. Click **Create credential type**.

| File                                | Name               | Needed for                                                                                          |
|-------------------------------------|--------------------|-----------------------------------------------------------------------------------------------------|
| <code>servicenow_api.yml</code>     | ServiceNow API     | tickets, ServiceNow health (step 10)                                                                |
| <code>mariadb_monitor.yml</code>    | MariaDB monitor    | the MariaDB check on a containerized database (root login refused there); optional on host installs |
| <code>keystore_password.yml</code>  | Keystore password  | optional: certs check on a PKCS12 keystore                                                          |
| <code>stigman_api.yml</code>        | STIG Manager API   | optional: POA&M cross-check with STIG Manager                                                       |
| <code>act_model_key.yml</code>      | ACT model key      | later, for ACT (ADDING_ACT.md): GenAI.mil and Ask Sage keys                                         |
| <code>act_genai_beta_key.yml</code> | ACT GenAI beta key | only for ACT with the GenAI.mil beta provider                                                       |

**You should see** each new type in the Credential Types list.

## Step 3. Credentials

**What.** Store the actual secrets, once.

**Why.** This is the only place a password or key is kept. Templates point at the credential by name. Nobody, including you, can read the secret back, and it is never in Git or job output.

**Automation Execution > Infrastructure > Credentials > Create credential:**

| Name                | Credential type | Fill in                                                                                                                                                                                                                                                                    |
|---------------------|-----------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Linux ssh (sudo)    | Machine         | <b>Username</b> <code>svc_aap</code> , <b>SSH Private Key</b> , <b>Privilege Escalation Method</b> <code>sudo</code> (plus the sudo password if sudo asks for one)                                                                                                         |
| site-automation git | Source Control  | only if the repository is private: user name and token, or an SSH key                                                                                                                                                                                                      |
| ServiceNow API      | ServiceNow API  | instance URL, API user, password (step 10)                                                                                                                                                                                                                                 |
| MariaDB monitor     | MariaDB monitor | the read-only database account, for MariaDB <b>and MySQL</b> ; needed when the database runs in a container (the ServiceNow database), usually needed for MySQL on a host, optional for MariaDB on a host. See MARIADB.md for the setup and the account to ask the DBA for |

You should see the credentials listed, with secret fields shown as **ENCRYPTED**.

## Step 4. Project

**What.** Tell AAP where the repository is.

**Why.** A project is AAP's copy of the repository. Templates pick their playbook from it.

1. **Automation Execution > Projects > Create project.**
2. **Name** `site-automation`. **Organization:** yours.
3. **Execution environment: Minimal execution environment** (`ee-minimal`). These runbooks use only modules that come with Ansible itself, so no extra collections or custom execution environment are needed (useful on a disconnected network).
4. **Source control type** `Git`. **Source control URL:** from step 1. **Source control branch/tag/commit:** `main`. Later, pin a release tag such as `v<version>` (the tags are listed on the repository's Releases page). **Source control credential:** `site-automation git` if the repository is private.
5. Tick **Update revision on launch**. Every job then syncs first, so a change you push is used right away.
6. Click **Create project**.

You should see the project's **Last job status** turn **Successful** after a few seconds. If it fails, open the sync job: it is almost always the URL, the branch name, or the credential.

## Step 5. Inventory (filled from the repository)

*(Skip this step if your inventory is already in AAP: see `USING_YOUR_AAP_INVENTORY.md`.)*

**What.** Create the `Linux servers` inventory and fill it from `inventories/site/hosts.yml`.

**Why.** Sourcing the inventory from the project brings in hosts, groups and the settings in `group_vars/`, all at once, and keeps them in step with Git. You do not type hosts into AAP by hand.

1. **Automation Execution > Infrastructure > Inventories > Create inventory > Create inventory.** Name `Linux servers`. Click **Create inventory**.
2. Open it > **Sources** tab > **Create source:**
  - **Name** `site inventory from Git`
  - **Source** `Sourced from a Project`

- **Project** `site-automation`
- **Inventory file** `inventories/site/hosts.yml`
- Tick **Overwrite**, **Overwrite variables** and **Update on launch**.

3. Click **Create source**, then **Launch inventory update**. The rocket icon also starts it.

**You should see**, under **Groups**, `stigman`, `mariadb_hosts`, `patch_hosts`... Under **Hosts** you should see your servers. Open a group: its **Variables** are the ones from `group_vars/<group>.yml`.

## Step 6. The first job template: Health check

**What.** A template that runs `playbooks/health_check.yml`.

**Why.** A job template is the thing you launch. It fixes which playbook, which inventory and which credentials. Its survey asks which checks to run.

1. **Automation Execution > Templates > Create template > Create job template.**
2. Fill in:

| Field                 | Value                                                                         | Why                                              |
|-----------------------|-------------------------------------------------------------------------------|--------------------------------------------------|
| Name                  | <code>Health check</code>                                                     |                                                  |
| Job type              | Run                                                                           | (Check also works: the checks read the same way) |
| Inventory             | <code>Linux servers</code>                                                    |                                                  |
| Project               | <code>site-automation</code>                                                  |                                                  |
| Playbook              | <code>playbooks/health_check.yml</code>                                       |                                                  |
| Execution environment | Minimal execution environment                                                 |                                                  |
| Credentials           | <code>Linux ssh (sudo)</code> (+ <code>MariaDB monitor</code> if you made it) | how AAP logs in                                  |
| Limit                 | leave empty, tick <b>Prompt on launch</b>                                     | so you can run it on one host                    |
| Privilege escalation  | ticked                                                                        | the checks read root-only files                  |

3. Click **Create job template**.
4. **Survey tab > Create survey question:**
  - **Question** `Which checks?`
  - **Answer variable name** `health_checks`
  - **Answer type** *Multiple Choice (multiple select)*
  - **Choices**, one per line: `daily`, `weekly`, `all`, `disk`, `mounts`, `services`, `performance`, `time`, `network`, `logging`, `selinux`, `fapolicyd`, `auditd`, `accounts`, `certs`, `patching`, `mariadb`, `mysql`, `database`, `containers`
  - **Default answer** `daily`. Tick **Required**. Click **Create question**.

**Made this survey with an earlier release?** AAP never updates a survey for you: open the **Health check** template > **Survey tab** > click the `health_checks` question > **Edit** > add `containers` on a new line under **Choices** > **Save**. (`daily` and `all` already include the containers check without that; the choice is for running

it on its own.)

5. Turn the survey **on** (the switch at the top of the Survey tab).

**You should see** the template in the Templates list with a rocket icon.

## Step 7. Run it on one host and read the output

**What.** Launch the health check on a single server.

**Why.** Start small. You will see what the checks find on a real host, and whether the thresholds suit you.

1. Click the rocket on **Health check**. **Limit:** one host name. Survey: **daily**. Click **Next**, then **Launch**.
2. The job output scrolls by. The parts to read:
  - **Check disk | run, Check services | run...**: each check running (green **ok** lines are normal).
  - **Report | result for this host**: the verdict in plain words. Either **HEALTHY. Checks run: ...** or a list of **[CRITICAL] / [WARNING]** findings. Each finding has a **look:** line: a command you can paste on the host to see more.
  - **Report | fleet summary**: one line for all hosts, for example **3 healthy, 1 with findings**.
  - A red **FAIL this host because it has findings** at the end. This is **on purpose**: a job with findings fails, so a workflow can react (open tickets). It does not mean the playbook broke.
3. The **Details** tab shows who launched it, when, and the exact revision of the repository.

**You should see** a finished job, green if the host is healthy, red with findings otherwise.

## Step 8. Tune the thresholds

**What.** Make the findings match what you care about.

**Why.** Defaults are a starting point. If a warning is expected on some host, change the setting for that host or group, not the code.

1. Change the setting in the project's settings file `playbooks/group_vars/<group>.yaml` (every host: `all.yaml`; one host: `playbooks/host_vars/<host>.yaml`, create it). This is the place `VARIABLES.md` recommends, and it wins over the same setting in `inventories/site/group_vars/` or in an AAP group or inventory Variables box. The file name must be the exact group or host name. Examples:

YAML

```
check_disk_overrides:
  /var/log/audit: {warn: 70, crit: 85}      # this filesystem: warn earlier
check_services_ignore_failed: ['^dnf-makecache']
site_fail_on: [critical]                   # warnings are reported but do not fail the job
```

(With `[critical]`, a check that could not run, `<check>:check-error`, still fails the host: a check that did not run is not a healthy result.)

2. Commit and push. `playbooks/group_vars/` is read straight from the project, which syncs on launch: nothing else to do. Only if you changed `inventories/site/group_vars/` instead, run the inventory source again in AAP (**Inventories > Linux servers > Sources > Launch inventory update**): those files reach AAP only through the inventory source.
3. Launch the health check again.

**You should see** the finding gone, or at its new threshold.

## Step 9. The other job templates

Create each the same way as step 6: **Inventory** `Linux servers`, **Project** `site-automation`, **Execution environment** Minimal. For the ones that log in to hosts, tick **Privilege escalation** and leave **Limit** empty with **Prompt on launch** ticked. Only the differences are listed. A survey question is *Multiple Choice (single select)* unless it says otherwise.

### Three templates run on the controller, not on hosts:

POA&M status, ServiceNow tickets and ServiceNow health. For these, do **not** tick Prompt on launch for Limit. A workflow's Limit (for example `patch_hosts`) is passed to every step that prompts for one. These playbooks run on `localhost`, so the Limit would match nothing and the step would silently do nothing, and still show green.

### Troubleshoot

On demand, when someone reports a problem. Always launch it with a **Limit**: the host with the problem. Findings do not fail this job.

| Field                           | Value                                                                                                                                                                                                                                                                                                                                       |
|---------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Playbook                        | <code>playbooks/troubleshoot.yml</code>                                                                                                                                                                                                                                                                                                     |
| Credentials                     | <code>Linux ssh (sudo)</code>                                                                                                                                                                                                                                                                                                               |
| Survey: <code>ts_area</code>    | "What kind of problem?" Choices: <code>overview</code> , <code>disk</code> , <code>performance</code> , <code>service</code> , <code>network</code> , <code>selinux</code> , <code>fail2ban</code> , <code>login</code> , <code>time</code> , <code>logs</code> , <code>mariadb</code> , <code>mysql</code> . Default <code>overview</code> |
| Survey: <code>ts_service</code> | <i>Text</i> , not required. "Service name (for area = service), e.g. nginx"                                                                                                                                                                                                                                                                 |
| Survey: <code>ts_target</code>  | <i>Text</i> , not required. "host:port it cannot reach (for area = network)"                                                                                                                                                                                                                                                                |
| Survey: <code>ts_since</code>   | <i>Text</i> . "How far back to read logs". Default <code>-2h</code>                                                                                                                                                                                                                                                                         |

### Database health - MariaDB / MySQL

The database check on its own, for **MariaDB and MySQL**, installed on the host or in **podman containers** (one, several, found automatically, rootless). Give it to the database team. It only reads: no data is read, nothing is changed. Set the databases up first (MARIADB.md): the container name in a settings file, and the account for the credential.

| Field                | Value                                                                                           |
|----------------------|-------------------------------------------------------------------------------------------------|
| Name                 | <code>Database health - MariaDB / MySQL</code>                                                  |
| Playbook             | <code>playbooks/database_health.yml</code>                                                      |
| Credentials          | <code>Linux ssh (sudo)</code> and <code>MariaDB monitor</code> (the read-only database account) |
| Privilege escalation | ticked                                                                                          |
| Job type             | Run, with <b>Prompt on launch</b> ticked, so you can pick <i>Check</i> first                    |

| Field  | Value                                                                                                                                                                                                                                                                                                                                                                                          |
|--------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Limit  | empty, <b>Prompt on launch</b> ticked. The job runs on the groups <code>mariadb_hosts</code> , <code>mysql_hosts</code> , <code>database_hosts</code> , <code>mariadb</code> and <code>mysql</code> ; a Limit narrows that to one host or group. A Limit cannot add a server outside those groups: for that, put <code>target: &lt;group or host&gt;</code> in the template's <b>Variables</b> |
| Survey | none. (Optional: a <i>Multiple Choice</i> <code>use_act</code> , choices <code>no</code> and <code>yes</code> , to add ACT's explanation; then also attach the <i>ACT model key</i> credential, <code>ADDING_ACT.md</code> )                                                                                                                                                                   |

Steps: create the template as in step 6, then launch it twice.

1. Launch with **Job type** `Check`. You should see, for every database host, a `MariaDB | status` line such as `MySQL 8.4.11 - container snow-mysql (podman, image ...): up 312.4 h, 23/151 connections ...`, then `HEALTHY` (or a list of findings, each with a command to look further), and at the end `N healthy, M with findings`. *Check* is safe here: the job only reads either way.
2. Launch with **Job type** `Run`. The result is the same. A host with findings shows as failed, so a workflow can open a ticket (workflow 8).

If a host prints `mariadb:connect` with `Access denied`, the credential is missing or the account lacks a privilege: `MARIADB.md`.

## Service watch

Two templates and a workflow: the steps are in `SERVICE_WATCH_DEMO.md` (Parts 5 and 6). The one field to decide here is which hosts it watches.

| Field       | Value                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
|-------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Playbooks   | <code>playbooks/service_watch.yml</code> ( <i>Service watch - check</i> ) and <code>playbooks/service_fix_approved.yml</code> ( <i>Service watch - apply approved fix</i> )                                                                                                                                                                                                                                                                                                                                                                                                                                             |
| Credentials | <code>Linux ssh (sudo)</code> , plus <code>ACT model key</code> on the check template                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| Variables   | optional: <code>target: &lt;group or host&gt;</code> . Without it the jobs run on the group <code>stigman</code> - a <b>group inside the template's inventory</b> , not an inventory called <code>stigman</code> (if there is no such group, the job stops and lists the groups the inventory has). To watch the containers of other servers, put that group (or <code>all</code> ) here, on <b>both</b> templates. Which containers are watched on each host is found by itself: every container that should be running, root's and each user's ( <code>SERVICE_WATCH_DEMO.md</code> , "Which containers are watched") |
| Survey      | on the workflow: <code>watch_mode</code> ( <code>approval</code> / <code>self-heal</code> )                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |

## Certificate report

Weekly: one table of every certificate on every host, soonest expiry first.

| Field                                      | Value                                                                          |
|--------------------------------------------|--------------------------------------------------------------------------------|
| Playbook                                   | <code>playbooks/cert_report.yml</code>                                         |
| Credentials                                | <code>Linux ssh (sudo)</code> (+ <code>Keystore password</code> if you use it) |
| Survey: <code>check_certs_warn_days</code> | <i>Integer</i> . "Warn this many days before expiry". Default <code>30</code>  |

## POA&M status

Weekly. Runs on the controller and reads `poam/poam.csv` from the project.

| Field                                 | Value                                                                             |
|---------------------------------------|-----------------------------------------------------------------------------------|
| Playbook                              | <code>playbooks/poam_status.yml</code>                                            |
| Credentials                           | none (+ <b>STIG Manager API</b> for the cross-check)                              |
| Privilege escalation                  | not needed                                                                        |
| Limit                                 | empty, and <b>no</b> Prompt on launch (it runs on the controller)                 |
| Survey: <code>poam_stigman_api</code> | <i>Text</i> , not required. "STIG Manager API URL (blank = skip the cross-check)" |

## ServiceNow - test ticket

Run once before the tickets step, and again after a password or certificate change. It opens a test incident, reads it back, adds a work note and resolves it, with PASS / WARN / FAIL for each step (SERVICENOW\_SETUP.md). Runs on the controller.

| Field                | Value                                                                                                         |
|----------------------|---------------------------------------------------------------------------------------------------------------|
| Playbook             | <code>playbooks/servicenow_test_ticket.yml</code>                                                             |
| Credentials          | <b>ServiceNow API</b>                                                                                         |
| Job type             | Run, with <b>Prompt on launch</b> ticked, so you can pick <i>Check</i> (a dry run: it only logs in and reads) |
| Privilege escalation | not needed                                                                                                    |
| Limit                | empty, and <b>no</b> Prompt on launch (it runs on the controller)                                             |
| Survey               | none                                                                                                          |

## ServiceNow tickets

Only useful as a workflow step after a check. Runs on the controller.

| Field                | Value                                                             |
|----------------------|-------------------------------------------------------------------|
| Playbook             | <code>playbooks/servicenow_tickets.yml</code>                     |
| Credentials          | <b>ServiceNow API</b>                                             |
| Privilege escalation | not needed                                                        |
| Limit                | empty, and <b>no</b> Prompt on launch (it runs on the controller) |
| Survey               | none                                                              |

## ServiceNow health

Hourly. Runs on the controller.

| Field                | Value                                                             |
|----------------------|-------------------------------------------------------------------|
| Playbook             | <code>playbooks/servicenow_health.yml</code>                      |
| Credentials          | <code>ServiceNow API</code>                                       |
| Privilege escalation | not needed                                                        |
| Limit                | empty, and <b>no</b> Prompt on launch (it runs on the controller) |
| Survey               | none                                                              |

## Patch hosts

Monthly, inside the *Patch with checks* workflow. Run it as *Check* first: it lists what would be updated and changes nothing.

| Field                                    | Value                                                                                                                                                                                                                                                                                                                                                                                                                                |
|------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Playbook                                 | <code>playbooks/patch_hosts.yml</code>                                                                                                                                                                                                                                                                                                                                                                                               |
| Credentials                              | <code>Linux ssh (sudo)</code>                                                                                                                                                                                                                                                                                                                                                                                                        |
| Job type                                 | Run, with <b>Prompt on launch</b> ticked, so you can pick <i>Check</i> (a dry run)                                                                                                                                                                                                                                                                                                                                                   |
| Variables                                | <b>required:</b> <code>target: &lt;group&gt;</code> , the name of a <b>group (or host) inside the inventory</b> - not the inventory's own name. For example <code>target: stigman</code> , or <code>target: all</code> when the template's inventory holds only the servers to patch. Without it the job looks for a group called <code>patch_hosts</code> , and stops with a message that lists the groups your inventory does have |
| Survey: <code>patch_security_only</code> | "Security updates only?" Choices <code>false</code> , <code>true</code> . Default <code>false</code>                                                                                                                                                                                                                                                                                                                                 |
| Survey: <code>automatic_restarts</code>  | "Reboot hosts automatically when the updates need it?" Choices <code>no</code> , <code>yes</code> . Default <code>no</code> (no automatic restarts: the job lists the hosts that need a reboot)                                                                                                                                                                                                                                      |

## Apply approved ACT fix

Only inside the approval workflow (ADDING\_ACT.md). Create it when you add ACT.

| Field       | Value                                                                                                                            |
|-------------|----------------------------------------------------------------------------------------------------------------------------------|
| Playbook    | <code>playbooks/act_fix_approved.yml</code>                                                                                      |
| Credentials | <code>Linux ssh (sudo)</code> (no ACT key: it runs the approved commands itself)                                                 |
| Job type    | Run, with <b>Prompt on launch</b> ticked, so a workflow can run it as <i>Check</i> (a dry run that only shows what it would run) |
| Variables   | leave <b>Prompt on launch</b> unticked                                                                                           |
| Survey      | none                                                                                                                             |

## VMware jobs and reports

They run against vCenter, not your hosts, and have their own setup: the VMware execution environment, a vCenter credential, one template per job. Follow VMWARE.md; the emails are in EMAIL\_REPORTS.md.

## NetApp ONTAP health report

It reads the storage clusters' REST API, not your hosts: a read-only ONTAP account, a NetApp ONTAP credential, one template. Follow NETAPP.md.

## Try each one

With a Limit of one host:

- **Troubleshoot:** `ts_area = disk`. Each command's output appears under a `### what it shows` line.
- **Patch hosts:** Job type **Check**. It lists what would be updated and changes nothing.
- **Certificate report:** read the table under `All certificates on all hosts, soonest expiry first`.

## Step 10. ServiceNow

**What.** An API account, its credential, and a test ticket that proves it works.

**Why.** The tickets step and the ServiceNow health check use ServiceNow's REST Table API with a user name and password. No plugin or collection is installed.

**The whole setup, step by step, is in `SERVICENOW_SETUP.md`:** the account to ask for (with a request you can copy), the instance URL and every API call, the network path (firewall, proxy, the DoD CA), the credential, the settings, the test ticket, and how to check the ticket in ServiceNow. In short:

1. Ask the ServiceNow admins for an **integration account** (web service access only) with:
  - `itil` (or a role that can create and update incidents);
  - `mid_server` (read the MID Server list), if you want the MID Server check.
2. Create the **ServiceNow API** credential (step 3) with the instance URL, for example `https://servicenow.example.mil`, and the account.
3. In your settings file (`playbooks/group_vars/all.yml`; the example inventory's `inventories/site/group_vars/all.yml` already has these two lines, so change the value in one place only), set who gets the tickets:

YAML

```
servicenow_assignment_group: Linux Operations    # a group that exists in ServiceNow
servicenow_min_severity: warning                # or critical: tickets for critical findings only
```

4. Create **ServiceNow - test ticket** (step 9) and launch it, first as *Check*, then as *Run*. **You should see** six steps that say PASS, and a link to the test incident. It is resolved by the same job.
5. Launch **ServiceNow health** once: `API answers in N s`, then the MID Servers and their status.
6. Build the *Daily health* workflow (`WORKFLOWS_AND_SCHEDULES.md`). Its tickets step then runs after every check.

Try it safely first: launch the workflow **without** the ServiceNow API credential on the tickets template. The step then only prints `NOT CONNECTED - what would be opened` and sends nothing.

## Step 11. Workflows and schedules

See WORKFLOWS\_AND\_SCHEDULES.md: *Daily health, Weekly compliance, Patch with checks, Fix with approval (ACT)*, and when to schedule each.

### Going further, once the basics work:

- ACT (the model): ADDING\_ACT.md. If a model answers `HTTP 400`, see its section Both endpoint formats and `:probe`.
- Letting ACT run fixes you trust without asking: APPROVED\_COMMANDS.md.
- Every setting you can change, and where to put it: VARIABLES.md and VARIABLES\_REFERENCE.md (in print: the *Settings Reference* PDF).
- MariaDB and MySQL checks (also in podman containers): MARIADB.md.

## Step 12. Who can do what

Open a template or workflow › **User Access** (or **Team Access**) › **Add roles**:

| Role    | Can                                            |
|---------|------------------------------------------------|
| Execute | launch it (and answer its survey)              |
| Approve | approve or deny its approval steps (workflows) |
| Read    | see it and its job output                      |
| Admin   | change it                                      |

For example, give the help desk **Execute** on *Troubleshoot*, and only the server team **Execute** on *Patch hosts* and **Approve** on the patch workflow.

## If something goes wrong

| You see                                                                                | It means                                                                                | Do                                                                                                                                                        |
|----------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------|
| Project sync fails                                                                     | wrong URL, branch or Git credential                                                     | open the sync job and read the last lines                                                                                                                 |
| UNREACHABLE! ... Permission denied (publickey)                                         | SSH key or user wrong for that host                                                     | check the Machine credential; test <code>ssh svc_aap@host</code> from a workstation                                                                       |
| Missing sudo password / a password is required                                         | sudo wants a password                                                                   | put it in the Machine credential (Privilege Escalation Password)                                                                                          |
| Unknown check(s): ...                                                                  | a survey choice that is not a check name                                                | fix the survey choices (step 6)                                                                                                                           |
| the <name> check could not run: ...                                                    | that check hit an error. The others still ran                                           | the message says why, often a missing command such as <code>semanage</code>                                                                               |
| skipped: the MariaDB/MySQL check is off for <host>                                     | the database check only runs on database hosts                                          | add the host to <code>mariadb_hosts</code> , <code>mysql_hosts</code> or <code>database_hosts</code> in the inventory, or name its container (MARIADB.md) |
| Findings you do not care about                                                         | defaults do not fit that host                                                           | step 8: change the setting for its group or host                                                                                                          |
| NOT CONNECTED - what would be opened                                                   | the tickets template has no ServiceNow API credential                                   | attach it (step 10)                                                                                                                                       |
| Tickets step fails: No check results reached this job                                  | the check step before it failed before checking any host, or it was launched on its own | read the first red job in the workflow; run the tickets step only inside a workflow                                                                       |
| A controller step (tickets, POA&M, ServiceNow health) shows skipping: no hosts matched | its template prompts for a Limit and the workflow passed one                            | untick Prompt on launch for Limit on that template (step 9)                                                                                               |
| A VMware job or report fails                                                           | see the troubleshooting tables                                                          | VMWARE.md, "Troubleshooting"                                                                                                                              |
| No email, or an email error                                                            | see the troubleshooting table                                                           | EMAIL_REPORTS.md, "Troubleshooting"                                                                                                                       |

### 3. Using the inventory you already have in AAP

You imported your inventory into AAP. You do not need `inventories/site/` in Git, and you can skip the Git inventory parts of `SETUP_AAP.md` (the `cp inventories/example ...` in step 1, and all of step 5). Your hosts and groups stay in AAP. Your **settings** can go in either of two places, and you can use both as long as **one setting lives in one place**:

- **this project's settings files** in `playbooks/group_vars/` (edited in VS Code, reviewed in Git, never overwritten by an update); the next section; or
- **AAP's Variables boxes** (inventory, group, host): the steps after it. What you paste there is YAML, and the blocks to paste are below.

Which place wins when a setting is in both, and the full list of places, is in `VARIABLES.md`. Every setting and its default: `VARIABLES_REFERENCE.md`.

The runbooks need two things from your inventory:

1. **Groups**: which hosts are which (steps 1-2). The safety lists are built in and already cover the usual group names (`aap`, `aap_hosts`, `sat`, `idm`, `netapp`, `netapp_console_hosts`...).
2. **Settings (variables)**: your LDAP server, your SIEM, the account AAP logs in as, the ServiceNow assignment group... (steps 3-5).

#### The easy way: your settings files, already named after your groups

The repository has one settings file per AAP group, in `playbooks/group_vars/`. **The file name is the group name**: `mariadb.yml` applies to the hosts in AAP group `mariadb`, and `all.yml` to every host. Ansible reads these files by itself on every run, and the hosts and groups stay in AAP. **If your group has another name, copy the file to your group's name** (for example `mariadb.yml` to `mariadb_hosts.yml`, `aap.yml` to `aap_hosts.yml`, `netapp.yml` to `netapp_console_hosts.yml`, which are the names the example inventory uses). A file for a group your inventory does not have is never used, without any error. The files are:

| File                                                                                           | For AAP group                                | What is in it                                                                                                                                               |
|------------------------------------------------------------------------------------------------|----------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>all.yml</code>                                                                           | every host (and the ServiceNow / POA&M jobs) | placeholders, all switched off: the safety lists (already built in), ACT provider and models, LDAP/log server, service account, ServiceNow assignment group |
| <code>stigman.yml</code>                                                                       | <code>stigman</code>                         | the containers service watch watches ( <code>stigman</code> , <code>nginx</code> ); the web certificate                                                     |
| <code>mariadb.yml</code> , <code>mysql.yml</code>                                              | <code>mariadb</code> , <code>mysql</code>    | placeholders: turn the database check on, and the container's name (MariaDB or MySQL; <code>MARIADB.md</code> )                                             |
| <code>aap.yml</code> , <code>sat.yml</code> , <code>idm.yml</code> , <code>logstash.yml</code> | the same names                               | disk limits, certificates, services that must run                                                                                                           |
| <code>netapp.yml</code>                                                                        | <code>netapp</code>                          | nothing yet (the group is protected by the built-in lists)                                                                                                  |
| <code>rhel8_all.yml</code>                                                                     | <code>rhel8_all</code>                       | a placeholder for Python 3.9 (only with a newer AAP)                                                                                                        |

## To change a setting:

1. Open the file in VS Code (for example `playbooks\group_vars\mariadb.yml`).
2. A line that starts with `#` is **off**. To turn a placeholder on, delete the `#` in front of it (and in front of the lines that belong to it), and replace `CHANGE-ME` with your value. Keep the indentation exactly as shown: spaces, never tabs.
3. **Commit** and **Sync Changes**. AAP uses it on the next run: the project syncs itself when **Update revision on launch** is ticked.

Example (the full guide, with several containers and rootless ones, is MARIADB.md): MariaDB runs in a container called `snow-mariadb`. In `mariadb.yml`, change

```
check_mariadb_enabled: true
# check_mariadb_container: CHANGE-ME
```

into

```
YAML
check_mariadb_enabled: true
check_mariadb_container: snow-mariadb
```

**A group that has no file yet** (`sn_prod` here is only an example name; use one of yours): create `playbooks/group_vars/sn_prod.yml` with the settings for it. The same name as the AAP group, `.yml` at the end. For one host only: `playbooks/host_vars/<host name>.yml` (create the folder).

**One thing the files cannot do: choose which hosts a template runs on.** That is decided before any setting is read. Each playbook therefore has a built-in choice, and the job template's own **Variables** box (the *Extra variables* on its Details page, not the inventory's) can change it with one `target:` line:

| Job templates                                                            | Variables box empty: runs on                                                                                                               | To run on something else                                                                                                                                                                                                        |
|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Health check, Troubleshoot, Certificate report, Apply approved ACT fix   | <b>every host</b> in the template's inventory (the built-in group <code>all</code> )                                                       | one run: the <b>Limit</b> at launch (tick <b>Prompt on launch</b> next to Limit). Always: <code>target: "rhel8_all:rhel9_all"</code>                                                                                            |
| Database health - MariaDB / MySQL                                        | the groups <code>mariadb_hosts</code> , <code>mysql_hosts</code> , <code>database_hosts</code> , <code>mariadb</code> , <code>mysql</code> | <code>target: &lt;group or host&gt;</code> (a Limit only narrows those groups; it cannot add a server outside them)                                                                                                             |
| Service watch, Service watch - apply approved fix, STIG Manager - deploy | the group <code>stigman</code>                                                                                                             | <code>target: &lt;group&gt;</code>                                                                                                                                                                                              |
| Patch hosts                                                              | the group <code>patch_hosts</code> only (no such group: the job stops and lists your groups)                                               | <b>always set it</b> , for example <code>target: "rhel8_all:rhel9_all"</code> , and pick the hosts with the Limit at launch. <code>target</code> is a group or host <b>inside</b> the inventory, never the inventory's own name |

`target` and the Limit take group names, host names, or several joined with `:`. `stigman:mariadb` means either group, `rhel9_all:&sn_prod` only hosts in both, and `rhel9_all:!sn_dev` all of `rhel9_all` except `sn_dev`. Put a space after `target:`, quote a value that has a `:` in it, and use the group name exactly as the inventory's **Groups** tab shows it (`rhel9_all`, not `rhel9`).

**Is every host in this inventory a Linux server?** If some are not (firewalls, switches, appliances), an empty Variables box tries them too: they show as failed, and in a workflow that opens ServiceNow tickets each of them gets a ticket (a host that never reports is ticketed), on every run. Then put `target: "rhel8_all:rhel9_all"` on those templates **before** you schedule anything.

**These files are yours.** The update script adds a settings file you do not have yet, and never changes one you already have. A new release can therefore never overwrite your settings.

**Or put the same lines in AAP instead.** Everything below explains the AAP Variables boxes, which take exactly the same lines. **Keep each setting in one place.** If the same setting is in both a file here and an AAP *group* or *inventory* box, **the file here wins** (every job prints a note when that happens). The exception is one host: a *host's* Variables box in AAP beats a group file, and only a `playbooks/host_vars/` file beats the host box. Full order: VARIABLES.md.

## How variables work in AAP (two minutes)

A variable is a setting with a name, such as `check_disk_warn_pct: 85`. Where you type it decides which hosts get it:

| Where in AAP                                                                          | Applies to                                        | The Git-file equivalent                   |
|---------------------------------------------------------------------------------------|---------------------------------------------------|-------------------------------------------|
| <b>Inventories</b> › <b>your inventory</b> › <b>Edit inventory</b> › <b>Variables</b> | every host in the inventory                       | <code>group_vars/all.yml</code>           |
| <b>...</b> › <b>Groups</b> › <b>a group</b> › <b>Edit group</b> › <b>Variables</b>    | the hosts in that group (and in its child groups) | <code>group_vars/&lt;group&gt;.yml</code> |
| <b>...</b> › <b>Hosts</b> › <b>a host</b> › <b>Edit host</b> › <b>Variables</b>       | that one host                                     | <code>host_vars/&lt;host&gt;.yml</code>   |

- **Inside AAP the more specific one wins:** host beats group, a child group beats its parent group, and a group beats the inventory. A survey answer, or the template's own **Variables**, beats all of them. **Files in this project's** `playbooks/group_vars/` **beat the AAP group and inventory boxes** (see the previous section).
- **Anything you do not set uses the default** in `roles/<role>/defaults/main.yml`. Every setting is explained in VARIABLES\_REFERENCE.md. Do not edit the role files: set the value in a settings file or in AAP instead. A new release would overwrite edits in the roles, but never your settings files or AAP variables.
- The boxes take YAML. Keep the indentation exactly as shown (spaces, not tabs). AAP refuses to save YAML it cannot read.

## Step 1. Look at what you have

**Inventories** › **your inventory**: open the **Groups** tab, then the **Hosts** tab. Write down which of your groups (or hosts) are:

- all your Linux servers;
- the AAP server(s);
- the NetApp Console host;
- the STIG Manager host(s);
- the ServiceNow database host(s);

- the ServiceNow MID Server host(s);
- the servers monthly patching may touch.

On the old ansible-core machine, `ansible-inventory -i <your inventory> --graph` prints the same tree.

## Step 2. Create the groups the runbooks use

A host can be in many groups. Creating these groups leaves your existing groups as they are.

| Create group                                                                              | Put in it                                                                                                                                        | Why                                                                                            | Needed?                      |
|-------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------|------------------------------|
| <code>aap_hosts</code>                                                                    | the AAP server(s)                                                                                                                                | <b>safety:</b> ACT never fixes anything here by itself, and patching never reboots it          | <b>yes</b>                   |
| <code>netapp_console_hosts</code>                                                         | the NetApp Console host                                                                                                                          | <b>safety:</b> ACT only diagnoses here (vendor-managed)                                        | <b>yes</b> , if you have one |
| <code>no_patch</code>                                                                     | the groups <code>aap_hosts</code> and <code>netapp_console_hosts</code> (as children)                                                            | <b>safety:</b> <i>Patch hosts</i> skips these even when a job targets them                     | <b>yes</b>                   |
| <code>patch_hosts</code>                                                                  | the servers monthly patching may touch                                                                                                           | <i>Patch hosts</i> only patches these                                                          | only for patching            |
| <code>mariadb_hosts</code> ,<br><code>mysql_hosts</code> ,<br><code>database_hosts</code> | the database host(s), MariaDB or MySQL (any of the three, or the groups <code>mariadb</code> / <code>mysql</code> ; the check treats them alike) | the database check runs here, and the <b>Database health</b> template targets these by default | only for the database check  |
| <code>stigman</code>                                                                      | the STIG Manager host(s)                                                                                                                         | service watch and STIG Manager deploy run here                                                 | only for those               |

### The clicks:

1. **Inventories > your inventory > Groups > Create group.** Type the **Name** (exactly as in the table), then **Create group**.
2. Add hosts: open the new group > **Hosts** tab > **Add existing host** > tick the hosts > **Add**.
3. Or add one of your groups inside it: open the new group > **Related Groups** tab > **Add existing group** > tick your group > **Add**. Its hosts are now in the new group too.

**Keep your own names instead?** If you already have, say, an `aap` group and do not want a second group, tell the runbooks your names. Put these in the inventory **Variables** (step 3). A list you write **replaces** the shipped one, so repeat every name you want kept; `aap` and `aap_hosts` are always protected, whatever you write:

YAML

```
site_act_diagnose_only_groups: [aap, netapp]    # ACT never fixes these
patch_never_reboot_groups: [aap]              # patching never reboots these
patch_never_patch_groups: [aap, netapp]       # patching skips these
```

**Updated from 0.4.1 or older?** Your `playbooks/group_vars/all.yml` then still sets these three lists (and `site_act_provider`, `site_act_models`) as real lines, and that file silently wins over the inventory Variables box. Open it and put `#` in front of those lines (or delete them), or make your change there instead. The job's **NOTE - settings in this project's files win ...** line lists what the file still sets.

The checks run on every host of the inventory. To keep them to your Linux servers, put **target: linux\_servers** (your all-Linux group) in the **Variables** of the Health check, Troubleshoot, Certificate report and Apply approved ACT fix templates.

### Step 3. Settings for every host (inventory variables)

**Inventories › your inventory › Edit inventory › Variables.** If something is already there, add below it; do not delete what is there. Paste, then change every line marked **CHANGE**:

```

--- what every host must be able to reach, from itself ---

check_network_dns_names: [ldap.yoursite.mil]           # CHANGE: names that must resolve
check_network_tcp:                                     # CHANGE: services every host needs
  - {name: LDAP, host: ldap.yoursite.mil, port: 636}
check_logging_remote:                                  # CHANGE: your SIEM / log collector
  - {name: SIEM, host: siem.yoursite.mil, port: 6514}

# ---- the account AAP logs in as ----
check_accounts_nopasswd_allowed: [svc_aap]             # CHANGE: your AAP service account
check_accounts_expiry_watch: [svc_aap]                 # CHANGE: the same account

# ---- STIG ----
check_fapolicyd_required: true                         # false if some hosts do not run
  fapolicyd

# ---- where tickets go ----
servicenow_assignment_group: Linux Operations          # CHANGE: a group that exists in
  ServiceNow                                           # or critical: tickets for critical
servicenow_min_severity: warning                      only

```

The ACT provider and models are **not** set here: they are chosen in `playbooks/group_vars/all.yml` (lines ready as comments) or per template (ADDING\_ACT.md, step 2). Do not also set them here: if the file ever sets them, the file wins and the value here is ignored.

Not sure about a value yet? Leave the line out. The check then uses its default, or checks nothing for it (no `check_network_tcp` means no TCP checks).

### Step 4. Settings for particular groups

Open the group › **Edit group › Variables**, paste, change the **CHANGE** lines, **Save**. Put these on your own groups or on the groups from step 2. Both work.

**The AAP server(s)**, for example on `aap_hosts`:

```

YAML

check_disk_overrides:
  /var: {warn: 80, crit: 90}          # AAP keeps job output and container images under /var

```

**The STIG Manager host(s)**, on `stigman`:

YAML

```

check_certs_endpoints:          # the certificate browsers actually see
- {name: STIG Manager, host: 127.0.0.1, port: 443, servername: stigman.yoursite.mil} # CHANGE
check_certs_files: [/etc/stigman/*.crt, /etc/stigman/*.pem] # CHANGE:
  where your cert files are
watch_containers:              # optional since 0.6.0: service watch finds the containers
- name: stigman                 # that should run by itself; this list is watched first,
- name: nginx                   # in this order (dependencies first), and a missing one is
  reported

```

If STIG Manager runs as systemd units (`systemctl list-units | grep -i stig` shows them), also add them, so the services check watches them: `check_services_required_extra: [<unit>, <unit>]`.

**The ServiceNow database host(s)** (MariaDB or MySQL): also add this group to `mariadb_hosts` or `mysql_hosts` (step 2).

YAML

```

check_mariadb_container: servicenow-mariadb # CHANGE: the container name from `sudo podman ps
-a`

```

Then attach the **MariaDB monitor** credential to the **Database health**, Health check and Troubleshoot templates (it serves MySQL too). Container images give root a password, so the check logs in with a monitoring account instead. The whole setup (several containers, discovery, rootless containers, the account to ask the DBA for, and how to verify) is in MARIADB.md.

**The ServiceNow MID Server host(s):**

YAML

```

check_services_required_extra: [mid] # CHANGE: its unit name (systemctl list-units |
  grep -i mid)
check_certs_keystores:
- path: /opt/servicenow/mid/agent/security/agent_keystore # CHANGE: your MID Server's path

```

**The NetApp Console host:** nothing is needed. It is checked like every host; being in `netapp_console_hosts` and `no_patch` keeps automation from changing it.

## Step 5. Settings for one host (only when one host is different)

**Hosts › the host › Edit host › Variables.** For example, a host whose `/data` is always nearly full on purpose:

YAML

```

check_disk_overrides:
  /data: {warn: 95, crit: 98}

```

Or a host with a known-broken vendor unit you want ignored:

YAML

```

check_services_ignore_failed: ['^vendor-agent']

```

## Step 6. Use your inventory in the job templates

On every job template (and workflow), set **Inventory** to your inventory. That is the only link between them. Everything else (project, playbook, credentials, survey) is as in SETUP\_AAP.md, steps 6 and 9.

## Step 7. Check that it worked

1. Launch **Health check** with **Limit** = one server and the survey set to **daily**.
2. In **Report | result for this host**, look for your own values. A network finding names *your* LDAP server; an accounts finding names *your* service account. If a value is not there, the variable did not reach the host: check its spelling, and which group or host it is on.
3. **skipping: no hosts matched** (a green job that did nothing) means the Limit or **target** names a group or host this inventory does not have. Check the spelling against the inventory's **Groups** tab, and that the template uses the right **Inventory**.
4. **skipped: the MariaDB/MySQL check is off for <host>** means the host is in none of **mariadb\_hosts**, **mysql\_hosts**, **database\_hosts**, **mariadb** or **mysql** and has no container name set (MARIADB.md).
5. Findings you do not care about: change the setting for that group or host (steps 4-5), not the code.

## Which YAML files in the repository do I change?

| File                                                       | Change it?                                                                                                   |
|------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------|
| playbooks/group_vars/*.yaml,<br>playbooks/host_vars/*.yaml | <b>Yes, these are your settings</b> (the update script never overwrites them)                                |
| inventories/example/...                                    | <b>No.</b> Only a reference to copy blocks from; AAP does not read it                                        |
| roles/*/defaults/main.yaml                                 | <b>No.</b> Read VARIABLES_REFERENCE.md to see every setting; set values in a settings file or in AAP instead |
| playbooks/*.yaml                                           | <b>No</b>                                                                                                    |
| aap/credential_types/*.yaml                                | <b>No.</b> Paste their contents into AAP (credential types)                                                  |
| ansible.cfg                                                | <b>No.</b> AAP ignores its <b>inventory =</b> line and uses the template's inventory                         |
| poam/poam.csv                                              | <b>Yes.</b> Your POA&M export (see <b>poam/README.md</b> )                                                   |

## 4. ServiceNow: set up, test and verify

The runbooks open ServiceNow incidents for the problems the checks find (the **ServiceNow tickets** step), and check that ServiceNow itself is healthy (**ServiceNow health**). This page takes you from nothing to a **verified test ticket**, one step at a time:

1. an API account from the ServiceNow admins;
2. the addresses: the instance URL and the API calls;
3. the network path from AAP to ServiceNow;
4. a quick check from a shell (optional);
5. the credential in AAP;
6. your settings: which assignment group gets the tickets;
7. the **ServiceNow - test ticket** template;
8. a dry run, then a real run;
9. checking the ticket in ServiceNow.

**Sign-in method.** This integration uses **basic authentication only** (an account name and password over HTTPS). OAuth, CAC/PKI login and SSO are **not supported**. If your instance requires one of them for the API, ask the ServiceNow admins for an integration account that is allowed to use basic authentication for the Table API, and exempt from SSO.

Plan about 30 minutes, plus the time the ServiceNow admins need for step 1.

### What you need

| Item                                    | Example                                                                                                 | Who gives it to you                                   |
|-----------------------------------------|---------------------------------------------------------------------------------------------------------|-------------------------------------------------------|
| The instance URL                        | <code>https://yourinstance.servicenowservices.com</code> or <code>https://servicenow.example.mil</code> | the ServiceNow admins (or your browser's address bar) |
| An API account and its password         | <code>svc_aap_sn</code>                                                                                 | the ServiceNow admins                                 |
| Roles on that account                   | <code>itil</code> (and <code>mid_server</code> for the MID Server check)                                | the ServiceNow admins                                 |
| The assignment group for the tickets    | <code>Linux Operations</code>                                                                           | the ServiceNow admins or your team lead               |
| A network path from AAP to the instance | TCP 443 from the AAP execution nodes                                                                    | the network or firewall team                          |
| Trust for the instance's certificate    | the DoD root CA, if the instance uses a DoD certificate                                                 | you (step 3)                                          |

### Step 1. Ask the ServiceNow admins for an API account

**What.** An account that only the automation uses, allowed to read, open and update incidents.

**Why.** The playbooks log in to ServiceNow's REST API with a user name and password (basic authentication). A dedicated account keeps the tickets traceable ("opened by `svc_aap_sn`"), and its access can be limited to what the automation needs.

Send them this (change the names in brackets):

```
Request: a ServiceNow integration account for Ansible Automation Platform (AAP)

- User ID: svc_aap_sn (or your naming standard)
- "Web service access only": ticked (API use only, no interactive login)
- "Internal Integration User": ticked, if our release has it
- Roles: itil (read, create and update incidents; read groups and users)
        mid_server (read the MID Server list; only if we want the MID Server check)
  No admin role.
- Authentication: basic authentication over HTTPS for the Table API. If MFA or single
  sign-on is enforced, this account must be exempt, as web-service accounts usually are.
- If IP address access control is on: allow the AAP execution nodes' addresses: [addresses]
- Password: set, with "Password needs reset" unticked. Tell me the expiry or rotation
  rule, so the AAP credential can be updated in time.
- Tickets go to the assignment group: [Linux Operations]. Please confirm its exact name.
- If our instance has custom incident states or close codes: the value for "Resolved"
  and a valid close code.
```

**You should get back** the user ID, the password (through your normal secure channel), the exact assignment group name, and, if they are custom, the resolved state and close code.

## Step 2. The addresses

**The instance URL** is the start of the address you see in your browser on ServiceNow: the `https://` and the host name, **nothing after it**.

| In the browser                                                             | Instance URL for AAP                                     |
|----------------------------------------------------------------------------|----------------------------------------------------------|
| <code>https://yourinstance.servicenowservices.com/now/nav/ui/home</code>   | <code>https://yourinstance.servicenowservices.com</code> |
| <code>https://servicenow.example.mil/nav_to.do?uri=incident_list.do</code> | <code>https://servicenow.example.mil</code>              |

**The API calls.** Everything goes to that host over HTTPS (TCP 443), with basic authentication and JSON. These are the only calls the playbooks make:

| Method and path                                                                 | What for                                   | Used by              |
|---------------------------------------------------------------------------------|--------------------------------------------|----------------------|
| <code>GET /api/now/table/incident?sysparm_limit=1</code>                        | log in and read (the connection test)      | test ticket, health  |
| <code>GET /api/now/table/sys_user_group?sysparm_query=name=&lt;group&gt;</code> | does the assignment group exist?           | test ticket          |
| <code>GET /api/now/table/sys_user?sysparm_query=user_name=&lt;caller&gt;</code> | does the caller exist?                     | test ticket          |
| <code>GET /api/now/table/incident?sysparm_query=correlation_id=...</code>       | is a ticket for this problem already open? | tickets              |
| <code>POST /api/now/table/incident</code>                                       | open an incident                           | test ticket, tickets |

| Method and path                        | What for                       | Used by              |
|----------------------------------------|--------------------------------|----------------------|
| GET /api/now/table/incident/<sys_id>   | read an incident back          | test ticket          |
| PATCH /api/now/table/incident/<sys_id> | add a work note;<br>resolve    | test ticket, tickets |
| GET /api/now/table/ecc_agent           | MID Servers: up and validated? | health               |

### Step 3. The network path from AAP

**Which machine talks to ServiceNow.** The ServiceNow jobs run on the **AAP execution node**, inside the execution environment container, not on your managed hosts. To see the nodes, open **Automation Execution > Infrastructure > Instances** (or **Instance Groups**). Those are the addresses the firewall (and ServiceNow's IP access control, if any) must allow.

**Firewall.** Allow the execution nodes > the instance, TCP 443.

**A proxy.** If the execution nodes reach the outside only through a proxy, give the jobs the proxy settings. Go to **Settings > Automation Execution > Job > Edit > Extra Environment Variables**, and set (add to what is already there):

```
{"https_proxy": "http://proxy.example.mil:8080", "no_proxy": "localhost,127.0.0.1,.example.mil"}
```

This applies to every job. Leave hosts that must not go through the proxy in **no\_proxy**.

**The certificate.** A ServiceNow instance with a DoD certificate is only trusted if the execution environment trusts the DoD root CA. Choose one:

| Option                                        | How                                                                                                                                                                                                                                                                                   | When                                                                                          |
|-----------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------|
| <b>A. The CA in the execution environment</b> | have the AAP admins build the execution environment with the DoD root CAs in its trust store                                                                                                                                                                                          | best: every job trusts it                                                                     |
| <b>B. A CA file in the repository</b>         | put the CA certificate (PEM) in the project, for example (an example path: the folder does not exist until you create it)<br><code>playbooks/files/ca/servicenow-ca.pem</code> , and set<br><code>servicenow_ca_path: "{{ playbook_dir }}/files/ca/servicenow-ca.pem"</code> (step 6) | when you cannot change the execution environment. CA certificates are public, so this is fine |
| C. Switch the check off                       | <code>servicenow_validate_certs: false</code>                                                                                                                                                                                                                                         | only for one test, never for real use                                                         |

### Step 4. Check from a shell first (optional, fastest)

On the AAP execution node, or any Linux machine on the same network:

```
shell
curl -sS -u svc_aap_sn -H 'Accept: application/json' \
  'https://yourinstance.servicenowservices.com/api/now/table/incident?sysparm_limit=1&sysparm_fields=number'
```

`curl` asks for the password, so it is not left in your shell history.

**You should see** `{"result":[{"number":"INC0012345"}]}` (any number). Anything else:

| What you get                                                                                 | What it means                                                                                                  |
|----------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------|
| SSL certificate problem: unable to get local issuer certificate                              | the CA is not trusted (step 3, the certificate)                                                                |
| Could not resolve host                                                                       | DNS: check the host name                                                                                       |
| nothing, then Connection timed out, or Connection refused                                    | firewall or proxy (step 3)                                                                                     |
| <code>{"error":{"message":"User Not Authenticated"...}}</code> (HTTP 401)                    | wrong user or password, the account is locked, or it is not allowed to use basic authentication                |
| <code>{"error":{"message":"Operation Failed","detail":"ACL Exception..."}}</code> (HTTP 403) | the account lacks the <code>itil</code> role                                                                   |
| an HTML page                                                                                 | the URL has something after the host name, or a login page (SSO) answered: the account must be exempt from SSO |

## Step 5. The credential in AAP

- If you have not yet, create the **ServiceNow API** credential type (SETUP\_AAP.md, step 2). It uses the file `aap/credential_types/servicenow_api.yml`.
- Automation Execution > Infrastructure > Credentials > Create credential:**
  - Name** `ServiceNow API`;
  - Credential type** `ServiceNow API`;
  - Instance URL**: from step 2 (only `https://` and the host name);
  - API user name** and **API password**: from step 1.
- Create credential.**

**You should see** the credential listed, with the password shown as `ENCRYPTED`. The job gets the three values as `SN_HOST`, `SN_USERNAME` and `SN_PASSWORD`. They never appear in job output, and never reach your managed hosts.

## Step 6. Your settings

In VS Code, open `playbooks/group_vars/all.yml`. (If your inventory comes from Git, the same lines in `inventories/site/group_vars/all.yml` end up in AAP's inventory Variables box, and the project's file wins over them: keep each setting in **one** place. See VARIABLES.md.) In the ServiceNow block, delete the `#` in front of the lines and fill in your values:

```

---- ServiceNow tickets -----
servicenow_assignment_group: Linux Operations    # the exact name from step 1
servicenow_min_severity: warning                 # or critical: tickets for critical findings only
# servicenow_caller: svc_aap_sn                  # optional; empty = the API account is the caller
# servicenow_ca_path: "{{ playbook_dir }}/files/ca/servicenow-ca.pem" # step 3, option B

```

**Commit and Sync Changes.** Every other ServiceNow setting (urgency, impact, resolve, close code) is described in `roles/servicenow/defaults/main.yml`.

**All settings**, with their defaults: `servicenow`.

## Step 7. The ServiceNow - test ticket template

**Automation Execution › Templates › Create template › Create job template:**

| Field                | Value                                                                                     |
|----------------------|-------------------------------------------------------------------------------------------|
| Name                 | ServiceNow - test ticket                                                                  |
| Inventory            | any (for example <code>rhel-inventory</code> ); it runs on AAP itself, not on these hosts |
| Project              | site-automation                                                                           |
| Playbook             | playbooks/servicenow_test_ticket.yml                                                      |
| Credentials          | ServiceNow API                                                                            |
| Job type             | Run, with <b>Prompt on launch</b> ticked (so you can pick <i>Check</i> , a dry run)       |
| Limit                | empty, and <b>no</b> Prompt on launch                                                     |
| Privilege escalation | not needed                                                                                |

**Create job template.**

## Step 8. Run it: first a dry run, then for real

**The dry run. Launch › Job type Check › Launch.** Steps 1 and 2 run for real, because they only read. Steps 3 to 6 are only printed:

```

PASS 1 connection: the instance answers and svc_aap_sn logged in and read the incident table (0 s)
PASS 2 settings: assignment group "Linux Operations" exists; caller not set: the tickets show the
  API account as the caller
DRY RUN (Check mode): steps 1-2 ran (they only read). Steps 3-6 would:
3 open: POST https://yourinstance.servicenowservices.com/api/now/table/incident {...}

```

Nothing was created in ServiceNow.

**The real run.** Low priority does not mean silent: ServiceNow may still send its usual emails, to the assignment group for the new incident and to the caller when it is resolved. Warn the group first, or point the test at a test group (`servicenow_assignment_group: <test group>` in the template's **Variables**), or at a test instance if you have one.

**Launch › Job type Run.** At the end it prints a summary like this:

```

ServiceNow test at https://yourinstance.servicenowservices.com as svc_aap_sn:
PASS | 1 connection | the instance answers and svc_aap_sn logged in and read the incident table (0
s)
PASS | 2 settings | assignment group "Linux Operations" exists; caller not set: ...
PASS | 3 open | opened INC0012346 (sys_id 5f1c...)
PASS | 4 read back | INC0012346: state New, priority 5 - Planning, assignment group Linux
Operations, ...
PASS | 5 work note | added a work note to INC0012346 (see its Activity stream)
PASS | 6 resolve | resolved INC0012346 (state 6)

```

Step 4 also prints a link: **Open it in ServiceNow:** [https://.../nav\\_to.do?uri=incident.do?sys\\_id=...](https://.../nav_to.do?uri=incident.do?sys_id=...)

| Result | Meaning                                                                                                                               |
|--------|---------------------------------------------------------------------------------------------------------------------------------------|
| PASS   | that step works                                                                                                                       |
| WARN   | it works, but look at it (for example, no assignment group is set, or the instance would not resolve the ticket). The job stays green |
| FAIL   | that step does not work. The job is red, and the line says why and what usually causes it                                             |

To leave the test ticket open and look at it first, add `servicenow_test_resolve: false` to the template's **Variables**.

## Step 9. Check the ticket in ServiceNow

- Click the link from step 4 of the job output, or go to **Incident > All** and search for the number (INC0012346).
- Check:
  - Short description:** `[site-automation test] test ticket from AAP job <number> - safe to close;`
  - Assignment group:** yours (step 6);
  - Caller:** your `servicenow_caller`, or the API account;
  - Urgency** 3 and **Impact** 3, so normally **Priority** 5 - Planning, the lowest (your instance's priority table decides). The real tickets use `servicenow_urgency` and `servicenow_impact`;
  - Activity:** the work note `Test work note from AAP job ...;`
  - State:** Resolved, with the close notes `Test ticket from AAP job ...: resolved automatically by the same job` (unless you set `servicenow_test_resolve: false`).
- To find every test ticket: filter **Incident > All** on **Correlation display** = `site-health-test`.

**The test passes when** all six steps say PASS (or WARN for step 6, if you do not resolve tickets automatically), and the ticket in ServiceNow looks like the list above.

## When a step fails

| Step and message                                      | Cause                                                                                                        | Fix                                                                                                                                                                                                            |
|-------------------------------------------------------|--------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1: Attach the "ServiceNow API" credential             | the template has no ServiceNow credential                                                                    | step 7: add it under Credentials                                                                                                                                                                               |
| 1: HTTP -1 ... certificate verify failed              | the execution environment does not trust the instance's CA                                                   | step 3, the certificate                                                                                                                                                                                        |
| 1: HTTP -1 ... Name or service not known              | DNS                                                                                                          | the Instance URL's host name; DNS from the execution node                                                                                                                                                      |
| 1: HTTP -1 ... timed out or Connection refused        | firewall or proxy                                                                                            | step 3                                                                                                                                                                                                         |
| 1: HTTP 401                                           | wrong user or password; account locked or inactive; "Password needs reset"; basic authentication not allowed | step 1 with the admins, then update the credential                                                                                                                                                             |
| 1: HTTP 403                                           | the account may not read incidents                                                                           | the <code>itil</code> role                                                                                                                                                                                     |
| 1: HTTP 404                                           | the Instance URL has a path after the host name                                                              | only <code>https://</code> and the host name                                                                                                                                                                   |
| 2: assignment group "... does not exist               | the name differs from ServiceNow's, even in capitals or spaces                                               | copy the exact name from ServiceNow into <code>servicenow_assignment_group</code>                                                                                                                              |
| 2: could not look up ... (WARN)                       | the account may not read groups or users                                                                     | fine if step 4 shows the group; otherwise ask for read access to <code>sys_user_group</code>                                                                                                                   |
| 3: HTTP 403 ... ACL Exception Insert Failed           | the account may read but not create incidents                                                                | the <code>itil</code> role                                                                                                                                                                                     |
| 3: HTTP 400                                           | ServiceNow refused a value, and says which                                                                   | usually <code>servicenow_category</code> or a custom field                                                                                                                                                     |
| 4: did not stick (WARN)                               | the ticket was opened without the group                                                                      | the group name (step 6)                                                                                                                                                                                        |
| 5: HTTP 403                                           | the account may create but not update                                                                        | the <code>itil</code> role, or an ACL on work notes                                                                                                                                                            |
| 6: could not resolve ... Data Policy Exception (WARN) | the instance needs other values to resolve                                                                   | ask the admins for the resolved state and close code; set <code>servicenow_resolved_state</code> and <code>servicenow_close_code</code> . Only needed if you use <code>servicenow_resolve_cleared: true</code> |

## Next

- **ServiceNow health** (hourly): the instance's API and the MID Servers (RUNBOOKS.md).
- **The tickets step** in the *Daily health* workflow (WORKFLOWS\_AND\_SCHEDULES.md): one incident per problem, updated on later runs, noted or resolved when it clears.

- Keep **ServiceNow - test ticket**. Run it again after a password change, a new certificate, or a ServiceNow upgrade.

## 5. Workflows and schedules

A **workflow** chains job templates. Each arrow says when the next step runs:

- **Run on success:** only if the step before it succeeded.
- **Run on fail:** only if it failed.
- **Run always:** either way.

The checks *fail on purpose* when they find something. So "Run on fail" means "a problem was found", and "Run always" means "whatever happened". A step also receives what the steps before it published (**artifacts**). That is how the tickets step learns what the health check found.

This chapter first shows how to build any workflow, then a catalog of ready-made workflows for Linux, for Windows, and for both, and how to schedule them.

### How to build any workflow, click by click

1. **Automation Execution › Templates › Create template › Create workflow job template.**
2. **Name** (e.g. `Daily health`), **Organization**.
  - **Inventory: leave it empty.** Each step then uses its own template's inventory, so one workflow can hold Linux and Windows steps. A workflow inventory only reaches steps whose template prompts for the inventory.
  - Tick **Prompt on launch** next to **Limit** if you want to pick hosts at launch.
  - Leave **Variables** empty: workflow variables override every step's own answers.
3. **Create workflow job template.** The **workflow visualizer** opens.
4. **The first step: Add step › Node type:**
  - *Job Template:* one of the templates (the usual case);
  - *Approval:* a person must click Approve (give it a **Timeout**);
  - *Workflow Job Template:* another workflow, as one step;
  - *Project Sync / Inventory Source Sync:* refresh the code or the inventory first.

Pick the template, then fill in **this step's** answers:

- **Survey answers**, e.g. `health_checks = daily`;
- **Job type** `Check` for a dry run (`DRY_RUNS.md`);
- a **Limit**.

The wizard only offers what the template asks in its survey or has **Prompt on launch** ticked for. › **Finish**.

5. **Next steps:** hover over a box › **Add step and link** (the plus sign) › the node, and the **link type** (success / fail / always).
  - **Two things at once:** add a second first step with **Add step**. Both start together.
  - **Two arrows into one step:** open the step › **Convergence:** *Any* runs it when one parent finishes as linked, *All* waits for all of them.
6. **Save** (top right). Closing without saving throws the graph away.
7. **Try it:** launch it once with the risky steps set to `Check` (step 4), then set them back.

**Who may approve:** open the workflow › **User Access** (or **Team Access**) › **Add roles** › **Approve**. Pending approvals are in **Automation Execution › Administration › Workflow Approvals**, and in the workflow

job's view.

### Which templates fit where:

| Template                                        | Fails when                                                     | Publishes (for the next step)             | Typical next step                                                    |
|-------------------------------------------------|----------------------------------------------------------------|-------------------------------------------|----------------------------------------------------------------------|
| Health check / Windows health check             | a finding ( <code>site_fail_on</code> ), or ACT proposes a fix | the findings ( <code>site_health</code> ) | ServiceNow tickets ( <b>always</b> ); an approval ( <b>on fail</b> ) |
| Certificate report / Windows certificate report | a certificate expires soon                                     | the findings                              | ServiceNow tickets ( <b>always</b> )                                 |
| POA&M status                                    | an item is overdue / due soon                                  | the findings                              | ServiceNow tickets ( <b>always</b> )                                 |
| Troubleshoot / Windows troubleshoot             | never for findings (a look, not a verdict)                     | the findings                              | a health check, tickets if you want them                             |
| Patch hosts / Windows patch                     | a host fails, or its services do not come back                 | -                                         | a health check ( <b>always</b> )                                     |
| Apply approved ACT fix                          | the problem is still there                                     | the findings of the re-check              | ServiceNow tickets ( <b>always</b> )                                 |
| Service watch - check                           | a container is down (NEEDS APPROVAL)                           | the fix to approve                        | an approval ( <b>on fail</b> )                                       |
| ServiceNow tickets                              | problems exist (it opened or updated tickets)                  | -                                         | - (the end)                                                          |

## Adding ACT to a workflow: what "on fail" means, and where `site_act` goes

**You do not add `site_act` to the workflow.** There is no ACT template and no ACT box. ACT runs **inside the check job**. The check playbooks (Health check, Troubleshoot, Service watch, the Windows ones) already include `roles/site_act`, and it runs when that job's answers say `use_act = yes`, at the level in `site_act_level`. In the workflow you only give the check box those two answers.

**"On fail" is the check job's own result.** Nothing extra is checked. The check playbook decides by itself, against its thresholds (the role defaults and your `group_vars`: disk full %, certificate days, ...). The check box ends **red** when:

- it found a problem at or above `site_fail_on` (default: `critical` and `warning`), or
- ACT (`diagnose`) proposed a fix that waits for a person: `NEEDS APPROVAL ... ACT proposes: <commands>`, or
- a host could not be checked (unreachable, login or sudo failed).

A **Run on fail** link from the check box therefore leads to the approval. **Run on success** means the check box ended green; **Run always**, either way.

| Box                                                                                     | Link into it                | Runs when               |
|-----------------------------------------------------------------------------------------|-----------------------------|-------------------------|
| 1. Health check ( <code>use_act = yes</code> , <code>site_act_level = diagnose</code> ) | (start)                     | always                  |
| 2. Approval                                                                             | <b>Run on fail</b> , from 1 | the check job ended red |

| Box                       | Link into it                          | Runs when                          |
|---------------------------|---------------------------------------|------------------------------------|
| 3. Apply approved ACT fix | <b>Run on success</b> , from 2        | a person clicked Approve           |
| 4. ServiceNow tickets     | <b>Run always</b> , from 1 and from 3 | after the check, and after the fix |

The steps in words:

1. **Health check** box: answers `use_act = yes`, `site_act_level = diagnose`. This needs those two survey questions on the template (ADDING\_ACT.md, step 3).
2. **Approval** box, **Run on fail** from 1.
3. **Apply approved ACT fix** box (`playbooks/act_fix_approved.yml`), **Run on success** from 2. It needs no answers: AAP hands it ACT's proposed commands from the health check job, and it runs exactly those.
4. **ServiceNow tickets** boxes, **Run always**, from 1 and from 3.

The approver opens the health check job and reads `ACT | what ACT says` and the red `NEEDS APPROVAL` line. If there is no such line (the job is red because of findings ACT had no fix for, or an unreachable host), there is nothing to approve: **Deny**. Approving does no harm either: the apply step skips every host with nothing approved.

**Without approvals:** `site_act_level = self-heal` with the fixes you allow in `site_act_allow`. ACT applies those itself, inside the check job, and the checks run again. Then the workflow is simply `[Health check] --Run always--> [ServiceNow tickets]`.

**A playbook of your own** gets ACT in the **playbook**, not in the workflow: it includes `site_findings` (start, your check), then `site_act`, then `site_findings` report. See ADDING\_ACT.md, "Use ACT in a playbook of your own". Its job template then gets the same two survey questions, and it fits the same workflow.

## Linux workflows

### 1. Daily health

```
[Health check] --Run always--> [ServiceNow tickets]
health_checks = daily
```

| Step | Template           | Answers                            | Link                            |
|------|--------------------|------------------------------------|---------------------------------|
| 1    | Health check       | <code>health_checks = daily</code> | (start)                         |
| 2    | ServiceNow tickets | none                               | <b>Run always</b> , from step 1 |

**Why "always".** The tickets step opens tickets for what failed, and also notes the tickets of problems that cleared. That needs the healthy runs too. It fails when problems exist, so the workflow shows red until they are fixed.

**Schedule:** every day at 06:00.

### 2. Weekly compliance

Three independent chains in one workflow. Add each first box with **Add step** (not "Add step and link"), so they all start together:

```
[Health check]      --Run always--> [ServiceNow tickets]
health_checks = weekly
[Certificate report] --Run always--> [ServiceNow tickets]
[POA&M status]     --Run always--> [ServiceNow tickets]
```

Give each chain **its own** tickets box. A box with two parents gets only one parent's findings: the artifacts overwrite each other.

**Schedule:** Mondays at 07:00.

### 3. Patch with checks

```
[Health check: pre] --success--> [Approval: patch?] --success--> [Patch hosts]
disk, mounts, services
|
Run on fail
v
[ServiceNow tickets]
[ServiceNow tickets]

[Health check: post] --always-->
health_checks = daily
```

| Step | Template           | Answers                                                                                      | Link                                                               |
|------|--------------------|----------------------------------------------------------------------------------------------|--------------------------------------------------------------------|
| 1    | Health check       | <code>health_checks = disk, mounts, services</code>                                          | (start)                                                            |
| 2    | <i>Approval</i>    | Name <code>Patch these hosts now?</code> ,<br><b>Timeout</b> 8 hours                         | <b>Run on success</b> , from 1                                     |
| 3    | Patch hosts        | <code>patch_security_only</code> as you like;<br><code>automatic_restarts</code> (see below) | <b>Run on success</b> , from 2                                     |
| 4    | Health check       | <code>health_checks = daily</code>                                                           | <b>Run always</b> , from 3                                         |
| 5    | ServiceNow tickets |                                                                                              | <b>Run always</b> , from 4                                         |
| 6    | ServiceNow tickets |                                                                                              | <b>Run on fail</b> , from 1 (tickets for what the pre-check found) |

#### Why these links:

- **Healthy first.** Patching a host that is already failing makes the next problem impossible to diagnose. The pre-check must pass before anyone is asked. **If any one host fails it, the whole run stops before the approval and no host is patched.** Fix that host (the ticket says what is wrong), or launch again with a Limit that leaves it out. Warnings count too while `site_fail_on` includes `warning`. A disk at 86% blocks patching, so clear it before the maintenance window.
- **A person approves.** The approval step waits (up to its **Timeout**). Deny, or no answer, stops the workflow and nothing is changed.
- **Always check after.** The post-check runs whatever happened in patching, so a half-patched host is still reported.

**No automatic restarts** unless you ask for them. Patching installs the updates; a host that then needs a reboot is **not** rebooted, and the job says so. The post-check's `patching` finding ("reboot needed") keeps it visible until someone reboots it. To let the job reboot hosts that need it, give step 3 `automatic_restarts = yes`. Add a survey question for it on the Patch hosts template (*Multiple Choice*, `no / yes`, default `no`), or put `automatic_restarts: true` in the step's variables. Hosts in `patch_never_reboot_groups` are never rebooted

either way.

**Launch** with **Limit** `patch_hosts`, or one group at a time (`stigman` this week, `servicenow_mid_hosts` next week). **Schedule:** monthly, in your maintenance window, for example the second Wednesday at 20:00. The approval then waits for the on-call person.

#### 4. Patch with a dry run first

The same as 3, but the approver sees exactly what will be installed before approving:

```
[Health check: pre]      health_checks = disk, mounts, services
| Run on success
[Patch hosts]           Job type = Check: lists what it would install
| Run on success
[Approval]              "Install what the dry run listed?"
| Run on success
[Patch hosts]           Job type = Run
| Run always
[Health check: post]    health_checks = daily
| Run always
[ServiceNow tickets]
```

| Step | Template           | Answers                                                                                                                                                | Link                           |
|------|--------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------|
| 1    | Health check       | <code>health_checks = disk, mounts, services</code>                                                                                                    | (start)                        |
| 2    | Patch hosts        | <b>Job type</b> <code>Check</code>                                                                                                                     | <b>Run on success</b> , from 1 |
| 3    | <i>Approval</i>    | Name <code>Install what the dry run listed?</code> .<br>Description: <code>Open the "Patch hosts" dry-run job: it lists every package per host.</code> | <b>Run on success</b> , from 2 |
| 4    | Patch hosts        | <b>Job type</b> <code>Run; automatic_restarts</code> as you decide                                                                                     | <b>Run on success</b> , from 3 |
| 5    | Health check       | <code>health_checks = daily</code>                                                                                                                     | <b>Run always</b> , from 4     |
| 6    | ServiceNow tickets |                                                                                                                                                        | <b>Run always</b> , from 5     |

Needs **Prompt on launch** for **Job type** on the Patch hosts template (steps 2 and 4 set it).

#### 5. Fix with approval (ACT)

For when you use ACT. See `ADDING_ACT.md` for the three levels.

```
[Health check] --Run on fail--> [Approval: apply ACT's fix?] --success--> [Apply approved ACT
fix]
use_act = yes
site_act_level = diagnose
|
Run always
v
[ServiceNow tickets]

|
Run always
v
[ServiceNow tickets]
```

| Step | Template                                     | Answers                                                                                                                                                                              | Link                                                                         |
|------|----------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------|
| 1    | Health check (with the ACT survey questions) | <code>use_act = yes, site_act_level = diagnose</code>                                                                                                                                | (start)                                                                      |
| 2    | ServiceNow tickets                           |                                                                                                                                                                                      | <b>Run always</b> , from 1                                                   |
| 3    | <i>Approval</i>                              | Name <code>Apply the fix ACT proposed?</code> , <b>Timeout</b> 30 minutes.<br>Description: <code>Read ACT's report and the NEEDS APPROVAL line in the health check job first.</code> | <b>Run on fail</b> , from 1                                                  |
| 4    | Apply approved ACT fix                       |                                                                                                                                                                                      | <b>Run on success</b> , from 3                                               |
| 5    | ServiceNow tickets                           |                                                                                                                                                                                      | <b>Run always</b> , from 4 (notes the tickets whose problem the fix cleared) |

**Caveat: not yet proven in a live AAP.** The whole approve-and-fix path depends on the check job's result (the proposed commands) reaching the apply job through the Approval node. The playbooks were tested on their own, but that hand-over inside AAP itself has not been run end to end. Rehearse it first with the dry run below, and read `SERVICE_WATCH_DEMO.md` (Part 9) if the apply job says `Nothing to apply`.

The approver reads, in the health check job, `ACT | what ACT says` (root cause, evidence, the fix) and the red `NEEDS APPROVAL ... ACT proposes: <command>` line. **That exact command** is all the apply step can run: it runs it itself (no model, no key), stops at the first command that fails, and runs the same checks again. The tickets carry ACT's analysis too.

**Rehearse it: see ACT's fixes, change nothing.**

- **One job, no workflow.** Launch **Health check** (or **Troubleshoot**) with `use_act = yes` and `site_act_level = diagnose`. ACT investigates with read-only commands and proposes fixes; every change is refused. The job ends red with `NEEDS APPROVAL ... ACT proposes: <commands>`: that list is what would run. Nothing on the host changed.
- **The whole workflow, as a dry run.** On the **Apply approved ACT fix** template (and *Service watch - apply approved fix*) tick **Prompt on launch** next to **Job type**. In the workflow visualizer click the apply box › **Edit** (pencil) › **Job type Check** › **Save**. Now run the workflow and approve: the apply job prints `DRY RUN on <host>: approved, and would run: <commands>`. `Nothing was changed (Check mode)`. and ends green. When you trust it, set that node's **Job type** back to **Run**.

## 6. Security posture (weekly)

The STIG-related checks, and the POA&M list, with tickets:

```
[Health check] --Run always--> [ServiceNow tickets]
selinux, fapolicyd, auditd, accounts
[POA&M status] --Run always--> [ServiceNow tickets]
```

| Step | Template     | Answers                                                           | Link    |
|------|--------------|-------------------------------------------------------------------|---------|
| 1    | Health check | <code>health_checks = selinux, fapolicyd, auditd, accounts</code> | (start) |

| Step | Template           | Answers | Link                           |
|------|--------------------|---------|--------------------------------|
| 2    | ServiceNow tickets |         | <b>Run always</b> , from 1     |
| 3    | POA&M status       |         | (start, with <b>Add step</b> ) |
| 4    | ServiceNow tickets |         | <b>Run always</b> , from 3     |

**Schedule:** Wednesdays 07:00. Use it instead of the **weekly** set in *Weekly compliance* if your security team wants these on their own day.

## 7. Certificate watch (daily)

Certificates expire on a day, not in a week:

```
[Certificate report] --Run always--> [ServiceNow tickets]
```

**Schedule:** every day 06:30. Set **check\_certs\_warn\_days** (survey) to 30 so tickets start a month ahead. The de-duplication keeps it to one ticket per certificate.

## 8. Database watch: MariaDB and MySQL (hourly)

```
[Database health - MariaDB / MySQL] --Run always--> [ServiceNow tickets]
```

| Step | Template                          | Answers                                                                                                                     | Link                       |
|------|-----------------------------------|-----------------------------------------------------------------------------------------------------------------------------|----------------------------|
| 1    | Database health - MariaDB / MySQL | none; it targets <b>mariadb_hosts</b> , <b>mysql_hosts</b> and <b>database_hosts</b> by default (a <b>Limit</b> narrows it) | (start)                    |
| 2    | ServiceNow tickets                |                                                                                                                             | <b>Run always</b> , from 1 |

The database template runs only the MariaDB / MySQL check (host installs and podman containers), so it is quick and the database team can own it and its tickets (MARIADB.md). Prefer to keep it in the general Health check? Use the *Health check* as step 1 with **health\_checks** = **disk**, **services**, **mysql** (**mariadb** and **database** are the same check) and a **Limit** of your database group. The template *Database health* also supports **use\_act=true**.

**Schedule:** hourly. It is quick, and a ticket comes within the hour of a problem.

## 9. Investigate a host (on demand)

One launch gives the troubleshooting commands, ACT's explanation and every check:

```
[Troubleshoot] --Run always--> [Health check]
ts_area, use_act = yes,          health_checks = all
site_act_level = explain
```

| Step | Template     | Answers                                                                                                                     | Link                       |
|------|--------------|-----------------------------------------------------------------------------------------------------------------------------|----------------------------|
| 1    | Troubleshoot | <b>ts_area</b> = <b>overview</b> (or asked at launch), <b>use_act</b> = <b>yes</b> , <b>site_act_level</b> = <b>explain</b> | (start)                    |
| 2    | Health check | <b>health_checks</b> = <b>all</b>                                                                                           | <b>Run always</b> , from 1 |

Give the workflow **Prompt on launch** for **Limit** and launch it with one host. No tickets: you are already looking. Add `[ServiceNow tickets]` after step 2 (**Run always**) if you want them.

## 10. STIG Manager deploy with checks

```
[STIG Manager - deploy]      Job type = Check (+ Show changes): what would change
| Run on success
[Approval]                  "Apply the changes the dry run listed?"
| Run on success
[STIG Manager - deploy]      Job type = Run
| Run always
[Service watch - check]
| Run always
[Health check]              health_checks = certs, services
```

| Step | Template              | Answers                                                 | Link                           |
|------|-----------------------|---------------------------------------------------------|--------------------------------|
| 1    | STIG Manager - deploy | <b>Job type</b> <b>Check</b> (and <b>Show changes</b> ) | (start)                        |
| 2    | <i>Approval</i>       | <b>Apply the changes the dry run listed?</b>            | <b>Run on success</b> , from 1 |
| 3    | STIG Manager - deploy | <b>Job type</b> <b>Run</b>                              | <b>Run on success</b> , from 2 |
| 4    | Service watch - check |                                                         | <b>Run always</b> , from 3     |
| 5    | Health check          | <b>health_checks = certs, services</b>                  | <b>Run always</b> , from 4     |

The dry run shows which files and units would change, and how. The approver approves only changes they expected. (This dry run has not been tried on a real server yet: run step 1 on its own once, on the test server, before you build the workflow.)

## 11. Service watch

The approval or self-heal workflow for containers: SERVICE\_WATCH\_DEMO.md, Part 6.

If you schedule it every 15 minutes, turn **off** "Enable concurrent jobs" on the workflow (edit the workflow), because its approval waits up to 30 minutes: without that, a new run starts every 15 minutes while the last one still waits for a person.

## Windows workflows

Build them the same way, with the Windows templates (WINDOWS.md). The ServiceNow tickets step is the same template as for Linux: it runs on AAP itself.

## 12. Windows daily health

```
[Windows health check]  --Run always--> [ServiceNow tickets]
health_checks = daily
```

**Schedule:** every day 06:15.

## 13. Windows weekly compliance

```
[Windows health check]      --Run always--> [ServiceNow tickets]
health_checks = weekly
[Windows certificate report] --Run always--> [ServiceNow tickets]
```

**Schedule:** Mondays 07:15.

## 14. Windows patch with a dry run first

```
[Windows health check: pre]    health_checks = disk, services
| Run on success
[Windows patch]               Job type = Check: lists the updates Software Center has
| Run on success
[Approval]                    "Install what the dry run listed?"
| Run on success
[Windows patch]               Job type = Run (automatic_restarts?)
| Run always
[Windows health check: post]   health_checks = daily
| Run always
[ServiceNow tickets]
```

| Step | Template             | Answers                                                                              | Link                             |
|------|----------------------|--------------------------------------------------------------------------------------|----------------------------------|
| 1    | Windows health check | <b>health_checks = disk, services</b>                                                | (start)                          |
| 2    | Windows patch        | <b>Job type Check</b>                                                                | <b>Run on success,</b><br>from 1 |
| 3    | <i>Approval</i>      | <b>Install what Software Center offers (listed in the dry run)?, Timeout 8 hours</b> | <b>Run on success,</b><br>from 2 |
| 4    | Windows patch        | <b>Job type Run; automatic_restarts</b> as you decide                                | <b>Run on success,</b><br>from 3 |
| 5    | Windows health check | <b>health_checks = daily</b>                                                         | <b>Run always,</b> from 4        |
| 6    | ServiceNow tickets   |                                                                                      | <b>Run always,</b> from 5        |

**No automatic restarts** unless step 4 gets **automatic\_restarts = yes**: the servers that need a restart are listed, and the post-check's **patching** finding ("a restart is pending") keeps them visible. **Schedule:** inside the servers' ConfigMgr maintenance window. Outside it, ConfigMgr holds the updates and the patch step stops and says so.

## 15. ACT for Windows rollout

```
[ACT for Windows - install]    Job type = Check: would install / update / already current
| Run on success
[Approval]
| Run on success
[ACT for Windows - install]    Job type = Run
```

The dry run lists, per server, whether ACT would be installed, updated, or is already current. Launch it with a Limit (a group) to roll out group by group. Run it again after every update of this repository.

## 16. Windows: investigate a server (on demand)

```
[Windows troubleshoot] --Run always--> [Windows health check]
ts_area                      health_checks = all
```

Launch with one server in the Limit.

## Linux and Windows together

### 17. Everything, every morning

One workflow, two independent chains (each first box added with **Add step**). Each step uses its own template's inventory, so leave the workflow's Inventory empty:

```
[Health check]          --Run always--> [ServiceNow tickets]
health_checks = daily
[Windows health check]  --Run always--> [ServiceNow tickets]
health_checks = daily
[Certificate report]    --Run always--> [ServiceNow tickets]
```

**Schedule:** every day 06:00. It replaces workflows 1, 7 and 12.

## VMware workflows

The VMware templates run on the AAP side against vCenter: no Limit, and each uses its own template's credentials (VMWARE.md).

### 18. VMware alarms: the report, then ACT's analysis (daily)

```
[VM - alarms report]  --Run always--> [VM - alarms ACT analysis]
```

**Always**, not On success: the report is red when `vm_alarm_fail: true` is set, and the analysis should run anyway. Each step sends its own email (`report_email_to` on each template, or a workflow survey question `report_email_to` for both). **Schedule:** every day 06:00. Only the analysis is wanted? Schedule `VM - alarms ACT analysis` on its own: its email lists every problem.

### 19. Snapshot cleanup with approval

```
[VM - snapshot report] --On success--> [Approval] --On success--> [VM - snapshot cleanup]
```

Click by click, with the survey and the dry run: VMWARE.md, "Snapshots".

### 20. ESXi security settings (nightly)

No workflow needed: a schedule on the `ESXi - security settings` template, every day at 00:00 in **your** time zone (VMWARE.md, "ESXi security settings").

### 22. vSphere capacity planning (weekly)

No workflow needed: a schedule on `VM - capacity planning`, e.g. every Monday at 07:00. One template per cluster (`vm_capacity_clusters`) gives each cluster's owners their own report (VMWARE.md, "Capacity planning").

### 21. NetApp ONTAP health (daily)

No workflow needed: a schedule on the `NetApp - ONTAP health report` template, e.g. every day at 06:00 in your time zone (NETAPP.md). In a workflow with `ontap_report_fail: true`, its "On failure" path can open a ticket.

## Schedules

**What.** A schedule launches a template or workflow at set times.

**How.** Open the workflow (or template) > **Schedules** tab > **Create schedule:**

1. **Name, Start date/time, Time zone** (use your site's, not UTC, so "06:00" means 06:00).
2. **Repeat frequency** and its details: *Day* (every 1 day), *Week* (on Monday), *Month* (on the second Wednesday)...
3. **Prompts**: the survey answers, the Limit and the Job type **for this schedule**. Without them, the schedule uses the template's **defaults**. A schedule can only change what the template **prompts on launch** or asks in its survey.
4. **Finish**. The schedule's page shows the next run times. Check them once.

| What                                                | When                                                                               | Launches                                                                         |
|-----------------------------------------------------|------------------------------------------------------------------------------------|----------------------------------------------------------------------------------|
| Daily health (or <i>Everything, every morning</i> ) | every day 06:00                                                                    | workflow 1 (or 17)                                                               |
| Certificate watch                                   | every day 06:30                                                                    | workflow 7                                                                       |
| Database watch                                      | hourly                                                                             | workflow 8                                                                       |
| ServiceNow health                                   | hourly                                                                             | template <i>ServiceNow health</i> , with an email <b>notification on failure</b> |
| Weekly compliance                                   | Monday 07:00                                                                       | workflow 2                                                                       |
| Security posture                                    | Wednesday 07:00                                                                    | workflow 6                                                                       |
| Patch preview                                       | the Monday before patch night                                                      | template <i>Patch hosts</i> , <b>Prompts › Job type Check</b>                    |
| Patch with checks                                   | monthly, in your window                                                            | workflow 3 or 4                                                                  |
| Service watch                                       | every 15 minutes (with <b>Enable concurrent jobs</b> switched off on the workflow) | workflow <i>Service watch - approve or self-heal</i>                             |
| Windows daily health / weekly compliance            | 06:15 daily / Monday 07:15                                                         | workflows 12 and 13                                                              |
| Windows patch                                       | monthly, in the ConfigMgr maintenance window                                       | workflow 14                                                                      |

**Notifications** (optional). **Automation Execution › Administration › Notifiers › Create notifier** (type *Email*, your mail relay). Then on a template or workflow › **Notifications** tab › switch on **Failure**. Use it where a ticket cannot help. ServiceNow health is the main case: if ServiceNow is down, it cannot take a ticket about itself.

## Things that trip people up

- **Red is not broken**. A check job with findings is red by design. Read **Report | result for this host**: findings mean the playbook worked.
- **Save the visualizer**. Closing without **Save** loses the graph.
- **One tickets box per check box**. Artifacts from two parents overwrite each other.
- **A step can only be set to Check** (or given a Limit, or survey answers) if its template has **Prompt on launch** for it, or a survey.
- **Schedules use survey defaults** unless the schedule sets its own answers (Prompts step).

- **Approvals need a timeout.** Otherwise a forgotten approval keeps the workflow open forever. After the timeout it counts as denied.
- **Workflow variables override node answers.** Leave the workflow's **Variables** empty, and set answers on each step.
- **Leave the workflow's Inventory empty** unless every step should use it: it only reaches steps whose template prompts for the inventory, and a Linux inventory on a Windows step matches no host.
- **Limit on the workflow** applies to every step whose template prompts for a Limit. The controller-side templates (ServiceNow tickets, ServiceNow health, POA&M status) must **not** prompt for one, or a workflow Limit makes them match no host and silently do nothing.
- **A tickets step fails when nothing reached it.** If the check step died before checking any host (credential, project sync, execution environment), the tickets step fails too, so a broken schedule shows red instead of green. Read the first red job.
- **No automatic restarts** unless a patch step gets `automatic_restarts = yes`.

## 6. Dry runs: try anything, change nothing

A **dry run** runs a job without changing anything. Ansible calls it *check mode*; AAP calls it **Job type: Check**. Every step that would change something only says what it *would* change. The steps that only read (the health checks, looking at logs, asking ServiceNow what is open) still run, so the answers are real.

Use a dry run the first time you use any template, on one host, and before every change you are not sure about: patching, applying an approved fix, installing ACT.

### Where to set it

There are five places. Pick the one that fits:

| Where                     | How                                                                                                                                         | Good for                                                                                                               |
|---------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------|
| <b>1. The template</b>    | Template › <b>Edit</b> › <b>Job type</b> <b>Check</b>                                                                                       | a template that is <i>always</i> a dry run, e.g. a copy named <b>Patch hosts (dry run)</b> that more people may launch |
| <b>2. At launch</b>       | Template › <b>Edit</b> › tick <b>Prompt on launch</b> next to <b>Job type</b> . Then at launch, pick <b>Check</b> or <b>Run</b>             | the usual way: the same template, you choose each time                                                                 |
| <b>3. A workflow step</b> | (needs 2 on the template) Workflow › <b>Visualizer</b> › click the step › <b>Edit</b> (pencil) › <b>Job type</b> <b>Check</b> › <b>Save</b> | rehearsing a whole workflow, or a "dry run first, then real" workflow                                                  |
| <b>4. A schedule</b>      | (needs 2) Template or workflow › <b>Schedules</b> › the schedule › <b>Prompts</b> › <b>Job type</b> <b>Check</b>                            | a regular preview, e.g. "what would patch night install" a week before                                                 |
| <b>5. Command line</b>    | <b>ansible-playbook</b> ... <b>--check</b> (add <b>--diff</b> to see file changes)                                                          | testing from a shell                                                                                                   |

**Show changes.** Tick **Show changes** on the template (or **--diff**) and a dry run also shows the line-by-line changes it would make to files, for example in STIG Manager's configuration.

**How to tell.** The job's **Details** tab says **Job type: Check**. Our playbooks also print **DRY RUN (Check mode): would ...** where it matters.

## What each template does in a dry run

| Template                                                                        | In a dry run it...                                                                                                                                                                                                                                                                              |
|---------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Health check, Database health, Troubleshoot, Certificate report                 | runs <b>exactly as usual</b> : they only read, so you get the same findings. Nothing is written: no syslog lines, no disk-forecast sample, and on Windows no event log entries                                                                                                                  |
| POA&M status                                                                    | as usual: reads the CSV and asks STIG Manager                                                                                                                                                                                                                                                   |
| ServiceNow tickets                                                              | asks ServiceNow what is already open, then prints <b>would open: ..., would add a note to: ..., would resolve: .... Nothing is sent</b>                                                                                                                                                         |
| ServiceNow health                                                               | as usual: only reads                                                                                                                                                                                                                                                                            |
| ServiceNow - test ticket                                                        | logs in and checks the assignment group and caller (it only reads), then prints the incident it would open, note and resolve. <b>Nothing is created</b>                                                                                                                                         |
| Patch hosts (Linux)                                                             | lists the packages dnf would update. Installs nothing, restarts nothing                                                                                                                                                                                                                         |
| Apply approved ACT fix                                                          | <b>DRY RUN on &lt;host&gt;: approved, and would run: &lt;commands&gt;</b> . Runs nothing. A command the apply-time guard would refuse (reboot, mkfs, ...) shows <b>Refused on &lt;host&gt;: ...</b> and that host goes red, as in the real run                                                  |
| Service watch - check                                                           | checks the containers and shows what it would propose. Writes no incident, does not call ACT                                                                                                                                                                                                    |
| Service watch - apply approved fix                                              | <b>DRY RUN on &lt;host&gt;: approved, and would run: &lt;commands&gt;</b> (or <b>Refused on &lt;host&gt;: ...</b> , as above)                                                                                                                                                                   |
| STIG Manager - deploy                                                           | says which files and units would change (with Show changes: how), and verifies nothing (nothing was started). Changes nothing. <i>Not yet tried on a real server: try it once on the test server</i>                                                                                            |
| Windows health check, troubleshoot, certificate report, connection test         | as usual: they only read                                                                                                                                                                                                                                                                        |
| Windows patch                                                                   | asks Software Center to rescan (without starting anything) and prints <b>DRY RUN (Check mode): would install &lt;updates&gt;</b>                                                                                                                                                                |
| ACT for Windows - install                                                       | <b>DRY RUN: would install ACT in C:\ProgramData\act because ...</b> , or that it is already current                                                                                                                                                                                             |
| VMware jobs (restart, shut down, snapshot, delete snapshot, notes, change VLAN) | finds the VMs and runs every check (the AAP VM is refused as in a real run), then prints <b>Would restart ..., Would take snapshot ..., would delete ...</b> , the notes <b>AFTER, would move .... vCenter is not changed</b> , and no report email is sent ( <b>DRY RUN: would email ...</b> ) |

| Template                                                                           | In a dry run it...                                                                                                                                                                                                |
|------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| VM - secure boot report, VM datastore report, VM snapshot report, VM alarms report | as usual: only read. No report email is sent                                                                                                                                                                      |
| VM - alarms ACT analysis                                                           | reads vCenter and lists the problems it would give ACT ( <b>A dry run (Check): ACT was not called</b> ). ACT is not called, no email is sent                                                                      |
| NetApp - ONTAP health report                                                       | as usual: it only reads (GET requests), and prints the same report. No email is sent                                                                                                                              |
| VM - capacity planning                                                             | as usual: it only reads. With <b>vm_capacity_act</b> , ACT is not called ( <b>ACT did not run: A dry run</b> ). No email                                                                                          |
| ESXi - security settings                                                           | reads every host's SSH service, timeouts and lockdown mode and lists what it would change ( <b>ESXi security settings (dry run): 3 of 40 host(s) would change</b> ). <b>No host is changed</b> , no email is sent |
| VM - snapshot cleanup                                                              | lists what it would delete ( <b>WOULD DELETE ...</b> ), deletes nothing, sends no email. In the approval workflow: set the cleanup step's Job type to Check (docs/VMWARE.md, "Dry run")                           |

**ACT does not run in a dry run.** It would call the model, and ACT could change things. A dry run with **use\_act = yes** reports the checks and says ACT gave no result. To see ACT's ideas without changes, do a *real* run with **site\_act\_level = diagnose**: ACT then only reads and proposes (ADDING\_ACT.md).

## Examples

**Before patching one server.** Launch **Patch hosts** › **Job type** **Check**, **Limit** **web01**. The output lists every package that would be updated. Nothing is installed.

```
TASK [patch : Patch | what was updated (dry run - what WOULD be updated)]
ok: [web01] => {"msg": "DRY RUN (Check mode), nothing was installed - would make 4 change(s):
  Installed: expat-2.5.0-6.el9_8.5.x86_64; Installed: libxml2-2.9.13-14.el9_8.5.x86_64;
  Removed: libxml2-2.9.13-14.el9_8.4.x86_64; Removed: expat-2.5.0-6.el9_8.3.x86_64"}
```

(**Removed** is the old version that the new one replaces.)

**Before patch night, every month.** On **Patch hosts** add a schedule for the Monday before patch night, with **Prompts** › **Job type** **Check**. The job output shows what patch night will install on every host, in time to exclude something.

**A dry-run copy for the team.** Template **Patch hosts** › ... › **Copy**. Rename the copy **Patch hosts (dry run)**, set **Job type** **Check**, untick **Prompt on launch** for it, and give your team **Execute** on the copy only. They can preview; only you can patch.

**Which tickets would a run open?** In the *Daily health* workflow click the **ServiceNow tickets** step › **Edit** › **Job type** **Check**. Run the workflow: the tickets job prints what it would open, note and resolve. Set it back to **Run** when the tickets look right.

**A Windows server, before its first real patch.** Launch **Windows patch** › **Job type** **Check**, **Limit** one server:

```
"DRY RUN (Check mode): would install 2026-09 Cumulative Update (KB5040001); Malicious Software Removal Tool (KB890830)"
```

**The approve-and-fix path, end to end.** See WORKFLOWS\_AND\_SCHEDULES.md, "Fix with approval (ACT)": set the apply step to **Check**, run the workflow, approve. The apply job shows what it would run.

**From a shell** (on a machine with Ansible and this repository):

```
shell
ansible-playbook -i <inventory> playbooks/patch_hosts.yml --check --diff -l web01
ansible-playbook -i <windows inventory> playbooks/win_patch.yml --check -l web01
```

## What a dry run cannot tell you

- Whether the updates will **install cleanly**, or how long they take: it only lists them.
- Whether a **restart** would succeed: nothing restarts in a dry run.
- Anything that **depends on an earlier change**. In a dry-run workflow the post-check after "patching" still sees the unpatched host, because nothing was patched.
- What **ACT** would say (see above).

## 7. The runbooks: every check and what to do

For each runbook: what it answers, its settings (defaults in `roles/<role>/defaults/main.yml`; change them in your settings files or AAP's Variables boxes, as `VARIABLES.md` shows), and each finding with what usually fixes it. The `look:` command printed with a finding is always the first thing to run.

### Health check (playbooks/health\_check.yml)

Runs the checks you pick on each host: a list of names, or the shortcuts:

- `daily` = disk, mounts, services, performance, time, network, logging, mariadb (the database check: MariaDB or MySQL; `mysql` and `database` also pick it), containers (podman containers, rootful and rootless)
- `weekly` = selinux, fapolicyd, auditd, accounts, certs, patching
- `all` = both

A host with a finding at a severity in `site_fail_on` (default: critical and warning) fails. That is what a workflow reacts to. A template that only emails its report to people: `site_fail_on: []` keeps the job green (`EMAIL_REPORTS.md`). A check that itself breaks becomes a warning finding (the X check could not run: `..., id <check>:check-error`) and the other checks still run. That finding fails the host even when `site_fail_on` is `[critical]` (a check that did not run is not a healthy result); only `site_fail_on: []` never fails.

**All settings** of this job, with their defaults: `site_findings`.

#### disk

| Finding                                          | Means                                                                                                      | Usually                                                                                                                                                        |
|--------------------------------------------------|------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>/var is 96% full</code>                    | local filesystem space<br>( <code>check_disk_warn_pct 85</code> ,<br><code>check_disk_crit_pct 95</code> ) | find what grew with the <code>look:</code> command:<br>old logs, core dumps, a runaway log file.<br><code>journalctl --vacuum-size=500M</code> for the journal |
| <code>... has used 97% of its inodes</code>      | too many small files (space may still be free)                                                             | usually a directory with millions of small files (session/cache/spool); <code>look:</code> finds it                                                            |
| <code>/var will be full in about 3.2 days</code> | growth-rate forecast from earlier runs (needs a day of history)                                            | same as above, before it becomes an outage                                                                                                                     |
| <code>could not read filesystem usage</code>     | <code>df</code> failed or timed out                                                                        | a hung mount; see the mounts check                                                                                                                             |

Per-mount limits: `check_disk_overrides: {/var/log/audit: {warn: 70, crit: 85}}`. The forecast keeps a small history in `/var/lib/site-health/disk-history.json` (the only file the checks write; a Check-mode run does not write it).

**All settings** of this check, with their defaults: `check_disk`.

## mounts

| Finding                                                         | Means                                                          | Usually                                                                                |
|-----------------------------------------------------------------|----------------------------------------------------------------|----------------------------------------------------------------------------------------|
| <code>/data (xfs) is in /etc/fstab but is not mounted</code>    | it did not mount at boot, or was unmounted                     | <code>journalctl -b</code> for the mount error; a missing disk or LUN; a typo in fstab |
| <code>/var is mounted READ-ONLY</code>                          | the kernel remounted it read-only after disk/filesystem errors | <b>urgent:</b> <code>journalctl -k</code> for I/O errors; storage team; plan an fsck   |
| <code>network mount /mnt/share does not answer within 5s</code> | NFS/CIFS server down, network, or a stale handle               | check the server and the network; diagnose only (NetApp is vendor-managed)             |
| <code>network mount ... is 96% full</code>                      | the share is full                                              | the storage owner                                                                      |

**All settings** of this check, with their defaults: `check_mounts`.

## services

| Finding                                          | Means                                                                                | Usually                                                                                                               |
|--------------------------------------------------|--------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------|
| <code>systemd unit X has FAILED</code>           | a unit crashed or exited with an error                                               | <code>systemctl status X</code> , <code>journalctl -u X</code> ; fix the cause, then <code>systemctl restart X</code> |
| <code>required service X is inactive</code>      | a service that must run ( <code>check_services_required + _extra</code> ) is stopped | same; find out why it stopped before starting it                                                                      |
| <code>required service X is not installed</code> | the unit does not exist here                                                         | remove it from the list for this group, or install it                                                                 |
| <code>required service X is disabled</code>      | it will not start after a reboot                                                     | <code>systemctl enable X</code>                                                                                       |
| <code>X has restarted 5 times</code>             | a restart loop: it keeps crashing                                                    | its log shows the crash                                                                                               |

Per group: `check_services_required_extra`: `[mariadb]`. To ignore a failed unit: `check_services_ignore_failed`: `['^dnf-makecache']`.

**All settings** of this check, with their defaults: `check_services`.

## performance

| Finding                                         | Means                                         | Usually                                                                |
|-------------------------------------------------|-----------------------------------------------|------------------------------------------------------------------------|
| 15-minute load is 14.2 on 4 CPUs (3.55 per CPU) | more work than CPUs, for 15 minutes           | the top CPU processes are in the finding; a runaway job, or undersized |
| only 4% of memory is available                  | memory is nearly exhausted                    | the top memory processes are in the finding; a leak, or undersized     |
| 63% of swap is in use                           | the host is short of memory                   | as above                                                               |
| CPUs spend 35% of their time waiting for disk   | slow or overloaded storage                    | <code>iostat -x</code> shows which device; storage latency             |
| the hypervisor takes 15% of CPU time (steal)    | the VM host is overcommitted                  | the virtualization team                                                |
| the kernel killed processes for lack of memory  | OOM kills in the last 24 h (with the victims) | the victim's memory limit, or the host's memory                        |
| 27 zombie processes                             | a parent process is not reaping its children  | restart the parent ( <code>look</code> : shows it)                     |

All settings of this check, with their defaults: `check_performance`.

## time

| Finding                               | Means                                            | Usually                                                                             |
|---------------------------------------|--------------------------------------------------|-------------------------------------------------------------------------------------|
| the clock is NOT synchronized         | chrony uses no time source                       | <code>chronyc sources -v</code> : unreachable servers, firewall (UDP 123)           |
| chrony is installed but not answering | chronyd is not running                           | <code>systemctl start chronyd</code>                                                |
| the clock is 2.4s ahead of NTP time   | offset above <code>check_time_offset_warn</code> | chrony corrects it slowly; very large offsets break Kerberos/AD logins at 5 minutes |

All settings of this check, with their defaults: `check_time`.

## network

| Finding                                                              | Means                                                                             | Usually                                                |
|----------------------------------------------------------------------|-----------------------------------------------------------------------------------|--------------------------------------------------------|
| there is no default route                                            | the host cannot reach other networks                                              | NetworkManager connection, gateway                     |
| this host's own name ... does not resolve                            | its FQDN is not in DNS or <code>/etc/hosts</code>                                 | DNS record                                             |
| <code>ldap.example.mil</code> does not resolve (DNS)                 | a name in <code>check_network_dns_names</code> fails                              | <code>/etc/resolv.conf</code> , DNS server             |
| LDAP ( <code>ldap.example.mil:636</code> ) does not answer within 5s | a TCP port in <code>check_network_tcp</code> is unreachable <b>from this host</b> | firewall (local or network), the service down, routing |

All settings of this check, with their defaults: `check_network`.

## logging

| Finding                                           | Means                                                        | Usually                                                                                                   |
|---------------------------------------------------|--------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------|
| the journal is kept in memory only                | <code>/var/log/journal</code> missing: logs vanish at reboot | <code>mkdir /var/log/journal &amp;&amp; systemctl restart systemd-journald (or Storage=persistent)</code> |
| log server SIEM (...) does not answer             | forwarding target unreachable                                | network/firewall; the SIEM collector                                                                      |
| rsyslog has 850 MB of undelivered messages queued | the log server has been unreachable for a while              | as above; the queue drains once it answers                                                                |
| 1203 error messages in the journal                | something complains a lot (top senders listed)               | look at the top sender first                                                                              |

**All settings** of this check, with their defaults: `check_logging`.

## selinux

| Finding                                                                                               | Means                                                                                             | Usually                                                                                                                                                               |
|-------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| SELinux is Permissive, it must be enforcing                                                           | not enforcing (STIG)                                                                              | <code>setenforce 1</code> after checking denials; <code>SELINUX=enforcing</code> in <code>/etc/selinux/config</code>                                                  |
| <code>/etc/selinux/config</code> says <code>SELINUX=permissive</code>                                 | a reboot would change the mode                                                                    | fix the config file                                                                                                                                                   |
| locally added permissive domain(s)                                                                    | SELinux only logs for those domains                                                               | <code>semanage permissive -d &lt;domain&gt;</code> once the policy is right                                                                                           |
| 12 SELinux denial(s) since yesterday - top: nginx (httpd_t) denied read on a file labeled user_home_t | SELinux blocked something                                                                         | a wrong label (fix: <code>restorecon -Rv &lt;path&gt;</code> ), a boolean ( <code>setsebool -P ...</code> ), or a real policy need. <code>audit2why</code> says which |
| SELinux boolean X is off, expected on                                                                 | a boolean in <code>check_selinux_booleans</code> differs                                          | <code>setsebool -P X on</code>                                                                                                                                        |
| N file(s) under <code>/etc/nginx</code> have the wrong SELinux label                                  | a dry-run <code>restorecon</code> would relabel them ( <code>check_selinux_relabel_paths</code> ) | <code>restorecon -Rv &lt;path&gt;</code>                                                                                                                              |

**All settings** of this check, with their defaults: `check_selinux`.

## fapolicyd

| Finding                                                                     | Means                            | Usually                                                                                                       |
|-----------------------------------------------------------------------------|----------------------------------|---------------------------------------------------------------------------------------------------------------|
| fapolicyd is inactive                                                       | application allow-listing is off | <code>systemctl start fapolicyd</code> ; its log says why it stopped                                          |
| fapolicyd is in permissive mode                                             | it logs but blocks nothing       | <code>permissive = 0</code> in <code>/etc/fapolicyd/fapolicyd.conf</code>                                     |
| N file(s) no longer match the fapolicyd trust database                      | files changed outside RPM        | <code>fapolicyd-cli --update</code> , or re-add the file's trust entry                                        |
| N fapolicyd denial(s) since yesterday - top: /usr/bin/bash -> /opt/app/tool | it blocked programs              | a legitimate program needs a trust entry ( <code>fapolicyd-cli --file add &lt;path&gt;</code> ), after review |

`check_fapolicyd_required: true` (the example inventory sets it) makes a missing fapolicyd a finding.

**All settings** of this check, with their defaults: `check_fapolicyd`.

## auditd

| Finding                                                                                                                                                                                             | Means                                                                                                                                        | Usually                                                                                                                                                                            |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| auditd is failed                                                                                                                                                                                    | nothing is being audited                                                                                                                     | <code>systemctl status auditd</code> ; often a full audit log partition                                                                                                            |
| auditing is DISABLED in the kernel                                                                                                                                                                  | <code>auditctl -s</code> says enabled 0                                                                                                      | <code>audit=1</code> on the kernel command line                                                                                                                                    |
| only 1 audit rule(s) are loaded                                                                                                                                                                     | the STIG rules are missing or failed to load                                                                                                 | <code>augenrules --check</code> , <code>augenrules --load</code> ; a syntax error in <code>/etc/audit/rules.d/</code>                                                              |
| N audit event(s) were LOST                                                                                                                                                                          | the kernel backlog overflowed                                                                                                                | raise <code>-b</code> (backlog) in the rules                                                                                                                                       |
| the audit log filesystem is 92% full                                                                                                                                                                | auditd may stop, or halt the host, when it fills                                                                                             | rotate/offload; check <code>space_left_action</code>                                                                                                                               |
| audit records are not forwarded                                                                                                                                                                     | au-remote is not active ( <code>check_auditd_remote_required</code> )                                                                        | configure audisp-remote                                                                                                                                                            |
| audit volume is 850 MB a day (... events in 24 h, at most N in one hour) - mostly account <code>svc_x</code> (60%), rule key <code>delete</code> (55%), program <code>/usr/bin/python3</code> (40%) | more audit records than <code>check_auditd_volume_warn_mb_day</code> (200), excluded accounts not counted                                    | what the named account/program does that the named STIG rule records: a job deleting temp files in a loop, an app in a rootless container (its processes keep the user's login ID) |
| the audit logs on disk reach back only 9.5 h (5 files of 8 MB, then the oldest is deleted)                                                                                                          | with <code>max_log_file_action = ROTATE</code> older records are gone ( <code>check_auditd_retention_warn_hours</code> 24)                   | bigger <code>max_log_file</code> / more <code>num_logs</code> in <code>/etc/audit/auditd.conf</code> , or fewer records (the finding names who makes them)                         |
| auditd dropped records N time(s) ... on the way to its plugins                                                                                                                                      | the queue to au-remote / the SIEM forwarder was full: the SIEM copy is missing records                                                       | fewer records; a larger <code>q_depth</code> in <code>auditd.conf</code>                                                                                                           |
| audit rule(s) stop recording <code>audit=1234</code> (someone) altogether                                                                                                                           | a <code>-a never</code> (or <code>-a always,exclude</code> ) rule hides an account that is not in <code>check_auditd_exclude_accounts</code> | remove the rule, or approve the account (below)                                                                                                                                    |
| too much to break down by account within 120 s                                                                                                                                                      | very large logs                                                                                                                              | <code>aureport</code> with the <code>look:</code> command, or a larger <code>check_auditd_volume_timeout</code>                                                                    |

**How much is audited, and the vulnerability scanner.** The volume, retention and rule findings come from the audit logs themselves. The breakdown (which account, which STIG rule key, which program) is read only when the volume or the retention is over its limit, at low priority, and is always printed in the job output (`Auditd | how much is audited`), with excluded accounts shown separately.

A credentialed vulnerability scan logs in to every host and runs thousands of commands with `sudo`: during a scan its account usually makes most of the audit records. List it in `check_auditd_exclude_accounts` (in `playbooks/group_vars/all.yml`, so it applies to every host) and it no longer counts toward the volume warning. It still counts for retention and disk space, which are physical: if the scan rotates a week of records

away in an hour, that is still reported, with the scanner named as the cause. The job output shows how much of the scanner's volume comes from syscall rules and how much from sudo/PAM records, which tells you what an audit rule exclusion would remove.

To stop auditd recording that account at all (a host change, not part of the health check):

- It is a deliberate exception to the STIG audit requirements: get it approved and recorded first.
- The rules `-a never,exit -F arch=b64 -S all -F auid=<UID>` and the same with `arch=b32` (the form of audit's own sample rules), placed before the STIG rules, remove the syscall records: a file like `/etc/audit/rules.d/05-exclude-scanner.rules` (`augenrules` merges the files in name order). The sudo/PAM records (`USER_CMD`, `USER_START`, `CRED_*`) come from user space and need `-a never,user -F auid=<UID>`. Leaving that one out keeps the sudo records, which still show what the account ran.
- Use the numeric UID. A name is looked up when the rules load at boot, before IdM/AD (sssd) can answer, and a rule that fails to load can stop the rest of the rules from loading.
- STIG hosts make the rules immutable (`-e 2`): a new rule takes effect only at the next reboot.
- The auditd check then reports rules like this only for accounts that are NOT on `check_auditd_exclude_accounts`.

**All settings** of this check, with their defaults: `check_auditd`.

### accounts

| Finding                                                    | Means                                                               | Usually                                                                                                            |
|------------------------------------------------------------|---------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------|
| <code>account(s) other than root have UID 0</code>         | full root power under another name                                  | <b>investigate first</b> ; remove or renumber the account                                                          |
| <code>account(s) with an EMPTY password</code>             | log in without a password                                           | <code>passwd -l &lt;user&gt;</code>                                                                                |
| <code>3 account(s) not used for 35+ days</code>            | inactive and not locked (STIG)                                      | lock them ( <code>usermod -L</code> ), or list break-glass accounts in <code>check_accounts_ignore_inactive</code> |
| <code>sudo without a password (NOPASSWD) for: alice</code> | a sudoers entry skips the password                                  | remove it, or allow it on purpose in <code>check_accounts_nopasswd_allowed</code> (e.g. <code>svc_aap</code> )     |
| <code>the password of svc_aap expires in 6 days</code>     | an account in <code>check_accounts_expiry_watch</code> expires soon | <b>change it before it breaks automation</b> (and update the Machine credential)                                   |

**All settings** of this check, with their defaults: `check_accounts`.

## certs

| Finding                                                          | Means                                                                                                              | Usually                                                                                                                                |
|------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------|
| certificate CN=... in /etc/stigman/stigman.crt expires in 9 days | a certificate file expires within <code>check_certs_crit_days</code> (critical) / <code>warn_days</code> (warning) | renew it. For STIG Manager: update the <i>TLS certificate</i> credential and run <i>STIG Manager - deploy</i> ; it restarts only nginx |
| ... in STIG Manager (https://127.0.0.1:443) expires in ...       | the certificate a TLS endpoint serves                                                                              | the service's certificate, even if the file on disk was already replaced (restart the service)                                         |
| ... in /opt/.../agent_keystore (alias mid) expires in ...        | a Java keystore entry                                                                                              | renew in the keystore (MID Server: ServiceNow's procedure)                                                                             |
| could not check the certificate in X                             | unreadable file, endpoint not answering, or a PKCS12 keystore without its password                                 | fix the path, or add the Keystore password credential                                                                                  |

What it looks at: `check_certs_files` (globs), the usual directories (`check_certs_discover_dirs`, skipping CA bundles), `check_certs_endpoints`, `check_certs_keystores`. For a PKCS12 keystore, name the variable that holds the password (`password_env: KEYSTORE_PASSWORD`), never the password.

**All settings** of this check, with their defaults: `check_certs`.

## patching

| Finding                                                  | Means                                                 | Usually                                         |
|----------------------------------------------------------|-------------------------------------------------------|-------------------------------------------------|
| 12 security advisories pending, 2 of them rated Critical | updates that fix security issues are waiting          | patch (the <i>Patch with checks</i> workflow)   |
| a reboot is needed to finish earlier updates             | new kernel or core libraries are not in use yet       | reboot in a window                              |
| no package has been updated for 60 days                  | older than <code>check_patching_max_age_days</code>   | patch                                           |
| cannot check for updates                                 | dnf cannot reach its repositories (Satellite, mirror) | repository configuration, network, subscription |

**All settings** of this check, with their defaults: `check_patching`.

## mariadb and mysql (health only): the database check

Checks **MariaDB and MySQL** (5.7, 8.0, 8.4; MariaDB 10.x and 11.x; Percona). Pick it in the Health check as `mariadb`, `mysql` or `database` (the same check), or run it alone with the **Database health - MariaDB / MySQL** template (`playbooks/database_health.yml`: only this check, on every database server it is pointed at; it fails a host that has findings, so a workflow can open a ticket). In the Health check it runs on hosts in `mariadb_hosts`, `mysql_hosts`, `database_hosts`, `mariadb` or `mysql`, and on any host that has a container name set. It runs status queries only: no data is read and nothing is changed. It handles a database installed on the host and one in **podman containers** (one, several, found automatically, or rootless): the client inside each container is used. **How to set it up, step by step, and how to verify it: MARIADB.md.**

| Finding                                                                     | Means                                                                                                     | Usually                                                                                                            |
|-----------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------|
| MariaDB is not running                                                      | the service or container is down                                                                          | its log; disk full; a crash                                                                                        |
| MariaDB is running but the status query failed                              | cannot log in, or it hangs                                                                                | the login (see below); <code>max_connections</code> reached                                                        |
| MariaDB restarted 12 minutes ago                                            | a restart nobody may have planned                                                                         | the error log: crash, OOM kill                                                                                     |
| 143 of 151 connections in use (95%)                                         | clients are about to be refused                                                                           | the <code>look:</code> query shows which user and client hold them; a connection leak in an app                    |
| N client connection(s) were REFUSED                                         | <code>max_connections</code> was reached since the last start                                             | as above; raise <code>max_connections</code> only if the load is real                                              |
| 40 queries are executing at this moment                                     | a pile-up: locking or one slow query                                                                      | <code>look:</code> lists the oldest                                                                                |
| 2 query/queries running longer than 300s                                    | long-running queries (user / db / seconds / state; never the query text)                                  | the owner decides; <code>KILL &lt;id&gt;</code> only after that                                                    |
| replication (default) is broken: IO thread No                               | a replica stopped copying                                                                                 | <code>Last_IO_Error</code> / <code>Last_SQL_Error</code> in the finding                                            |
| replica ... is 1900 seconds behind                                          | replication lag                                                                                           | load on the replica; long transactions on the primary                                                              |
| Galera node is not healthy / cluster has 2 node(s), expected 3              | Galera cluster trouble                                                                                    | the node's error log; network between nodes                                                                        |
| MySQL Group Replication is not healthy / has 2 ONLINE member(s), expected 3 | a Group Replication member is not ONLINE, or the group is short ( <code>check_mariadb_group_size</code> ) | <code>SELECT ... FROM performance_schema.replication_group_members</code> (in the finding); the member's error log |
| could not read the replication status / ... Group Replication members       | the monitoring account lacks a privilege                                                                  | the grant named in the finding (MARIADB.md)                                                                        |
| only 91% of reads come from memory                                          | <code>innodb_buffer_pool_size</code> may be too small                                                     | the DBA                                                                                                            |
| 3 [ERROR] line(s) in the MariaDB log                                        | errors in its log in the last day (latest shown)                                                          | read them                                                                                                          |
| the MariaDB data directory is on a filesystem 93% full                      | MariaDB stops writing when it fills                                                                       | disk space                                                                                                         |

**Logging in** (the full steps are in `MARIADB.md`). Without a credential, the check logs in as root through the local socket. That is MariaDB's default for a host install on RHEL (no password needed with sudo); MySQL's root usually has a password. **A database in a container usually gives root a password** (the official images do), and the check then reports `Access denied ... attach the "MariaDB monitor" credential`. For containers, or whenever you prefer not to use root, ask the DBA to create a monitoring account with **no data access**, inside that database with host `localhost`, and attach the *MariaDB monitor* credential to the Health check template:

SQL

```
CREATE USER 'aap_monitor'@'localhost' IDENTIFIED BY '<password>';
GRANT PROCESS, REPLICATION CLIENT ON *.* TO 'aap_monitor'@'localhost';    -- MySQL; MariaDB before
10.5.9
GRANT PROCESS, REPLICATION MONITOR ON *.* TO 'aap_monitor'@'localhost';    -- MariaDB 10.5.9 and
newer instead
GRANT SELECT ON performance_schema.* TO 'aap_monitor'@'localhost';    -- MySQL Group
Replication only
```

`PROCESS` shows other sessions (for long queries). Replication status needs `REPLICATION CLIENT` (MySQL) or, on MariaDB 10.5.9 and newer, `REPLICATION MONITOR: REPLICATION CLIENT` **alone is not enough there**. The password reaches the client only through its environment, never its command line.

**All settings** of this check, with their defaults: `check_mariadb`.

### containers: podman containers, root's and every user's

Finds the podman containers on the host by itself - you list nothing - and checks each one. Hosts without podman are skipped (`skipped: podman is not installed`, no finding). It is read-only: it runs `podman ps`, one `podman container inspect` and `systemctl show`, and nothing that starts, stops or creates anything.

**Which containers.** Root's (rootful podman) and each user's (rootless podman: only that user can see them). A user is looked at when a podman process runs as them, when they have *linger* on (a file in `/var/lib/systemd/linger/`), or when they are in `/etc/subuid` and have container storage in their home (`~/ .local/share/containers/storage`). A user's containers are read **as that user** (`runuser -u USER -- env XDG_RUNTIME_DIR=/run/user/UID podman ...`), and only when `/run/user/UID` exists: otherwise podman would create it, and a check must not change anything. Such a user is listed as `not checked` instead (and reported, below).

**Which should be running.** A container should be running when a systemd unit runs it and starts it at boot (a Quadlet `.container` file with an `[Install]` section, or an enabled unit from `podman generate systemd` or written by hand), or when its restart policy is `always` or `unless-stopped`. A container that has neither and stopped with exit code 0 was stopped on purpose: no finding. A Quadlet or `--new` unit removes its container while the unit is stopped, so the check also reads the Quadlet files and unit files: a container that should run but does not exist at all is found too (state `missing`).

The output has a table of what was found:

| OWNER<br>RUN                                                                              | NAME       | STATE      | UNIT                 | SHOULD |
|-------------------------------------------------------------------------------------------|------------|------------|----------------------|--------|
| root                                                                                      | systemd-db | running    | db.service (active)  | yes    |
| root                                                                                      | web        | running    | -                    | yes    |
| alice                                                                                     | web        | missing    | web.service (failed) | yes    |
| bob                                                                                       | web        | exited (3) | -                    | no     |
| not checked: carol - no /run/user/1003 (the user is not logged in and linger is off): ... |            |            |                      |        |

Every finding names the owner, and its **look:** command runs as that owner (paste it as root):

| Finding                                                                                                                                   | Means                                                                                                                                                                                                                                                                                                | Usually                                                                                                                                                                                                                                                                                                                                           |
|-------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| container X (owner 0) should be running (...) but it does not exist / it exited with code N (critical)                                    | a container that should run is down, or its systemd unit failed                                                                                                                                                                                                                                      | read the <b>look:</b> output (the unit's status, the container's log, the journal). Start it the right way: <b>systemctl start UNIT</b> for a unit (as a user: <b>runuser -u 0 --env XDG_RUNTIME_DIR=/run/user/UID systemctl --user start UNIT</b> ), <b>podman start X</b> for a plain container. Service watch does this for you, with approval |
| required container X ... does not exist / is not running (critical)                                                                       | a container in <b>check_containers_required</b> is missing or stopped                                                                                                                                                                                                                                | create or start it; fix its name in the setting                                                                                                                                                                                                                                                                                                   |
| container X (owner 0) is running but its healthcheck reports unhealthy (critical)                                                         | its own healthcheck fails                                                                                                                                                                                                                                                                            | <b>podman inspect X</b> shows the last healthcheck output ( <b>State.Health.Log</b> ); the container's log                                                                                                                                                                                                                                        |
| container X (owner 0) stopped with an error: it exited with code N, 2.5 hours ago (warning)                                               | a container that is not meant to run all the time stopped with an error in the last <b>check_containers_recent_hours</b> (24) hours                                                                                                                                                                  | its log. Older crashes are not reported                                                                                                                                                                                                                                                                                                           |
| container X (owner 0) has restarted N times - a restart loop (warning)                                                                    | more restarts than <b>check_containers_restart_warn</b> (3), counted by podman or by its systemd unit                                                                                                                                                                                                | its log shows why it keeps stopping                                                                                                                                                                                                                                                                                                               |
| container X (owner 0) was killed by the out-of-memory killer (warning)                                                                    | the kernel killed it for memory                                                                                                                                                                                                                                                                      | more memory on the host, or a higher memory limit on the container                                                                                                                                                                                                                                                                                |
| user U (uid N) has containers that should run (...) but linger is off (warning)                                                           | the user's containers stop when they log out and do not start at boot                                                                                                                                                                                                                                | <b>loginctl enable-linger U</b> (as root)                                                                                                                                                                                                                                                                                                         |
| user U (uid N) has podman containers ... and N unit(s) that should start at boot, but they could not be checked: no /run/user/N (warning) | the user has units that should start at boot (a Quadlet with <b>[Install]</b> , an enabled unit) but is not logged in and has no linger: none of their containers run now. A user with only container storage (an admin who once ran podman) is listed under <b>not checked:</b> and gets no finding | if they should run: <b>loginctl enable-linger U</b> . If not, leave them out: <b>podman_discover_ignore: ['U/. *']</b>                                                                                                                                                                                                                            |
| the containers of 0 could not be read: podman ps failed ... (warning)                                                                     | podman (or systemctl) failed for that owner: a blind spot                                                                                                                                                                                                                                            | run the <b>look:</b> command to see the error (storage problems, <b>podman system migrate</b> needed after an upgrade ...)                                                                                                                                                                                                                        |

**Settings** (roles `podman_discover` and `check_containers`; `VARIABLES_REFERENCE.md`):

| Setting                                    | Default           | What it does                                                                                                                                                        |
|--------------------------------------------|-------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>podman_discover_enabled</code>       | <code>true</code> | <code>false</code> = do not look for containers (the check says <code>skipped</code> )                                                                              |
| <code>podman_discover_rootless</code>      | <code>true</code> | also look at users' containers                                                                                                                                      |
| <code>podman_discover_users</code>         | <code>[]</code>   | more users to look at, e.g. <code>[appsvc]</code> when the rules above miss one                                                                                     |
| <code>podman_discover_ignore</code>        | <code>[]</code>   | containers to leave out: regular expressions that must match the <b>whole</b> name or <b>OWNER/NAME</b> , e.g. <code>['test-.*', 'alice/.*', 'root/scratch']</code> |
| <code>check_containers_required</code>     | <code>[]</code>   | containers that must exist and run: a name (any owner) or <b>OWNER/NAME</b> , e.g. <code>[stigman-api, alice/web]</code>                                            |
| <code>check_containers_recent_hours</code> | <code>24</code>   | how recent a crash must be to be reported                                                                                                                           |
| <code>check_containers_restart_warn</code> | <code>3</code>    | restarts above this = a restart loop                                                                                                                                |

**To verify it** on one host: run the Health check with `health_checks: containers` and a Limit of that host; compare the table with `sudo podman ps -a` and, for a user, `sudo runuser -u alice -- env XDG_RUNTIME_DIR=/run/user/$(id -u alice) podman ps -a`. Service watch uses the same way of finding containers (`SERVICE_WATCH_DEMO.md`).

**All settings** of this check, with their defaults: `check_containers`, `podman_discover`.

## Troubleshoot (playbooks/troubleshoot.yml)

For "something is wrong on this host": pick the area and the job runs the commands an experienced admin runs first. Each is printed under a line saying what it shows, and the matching health checks run too. Read-only. Findings do not fail this job: you asked for a look, not a verdict.

| Area                                        | Commands (what they show)                                                                                               | Checks                                             |
|---------------------------------------------|-------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------|
| overview                                    | uptime, failed units, disk, memory, recent errors, recent reboots                                                       | disk, mounts, services, performance, time, logging |
| disk                                        | space and inodes, biggest directories on the fullest filesystem, deleted-but-open files, files over 1 GiB, journal size | disk, mounts                                       |
| performance                                 | load, top, vmstat, memory users, disk latency (iostat), load history (sar), OOM kills                                   | performance                                        |
| service ( <code>ts_service</code> )         | its status, why it stopped (result, exit code, restarts), its log, dependencies, unit file, listening ports             | services (plus that unit)                          |
| network ( <code>ts_target</code> host:port) | addresses, routes, DNS servers, name lookup, TCP connect, route taken, firewall, connection counts                      | network (plus that target)                         |

| Area               | Commands (what they show)                                                                | Checks    |
|--------------------|------------------------------------------------------------------------------------------|-----------|
| selinux            | mode, today's denials, audit2why, changed booleans, local file labels                    | selinux   |
| fapolicyd          | status, today's denials, rules, enforcing?                                               | fapolicyd |
| login              | sshd log, faillock lockouts, failed logins, SSSD domain status, authselect, clock offset | accounts  |
| time               | chrony tracking and sources, timedatectl, configured servers                             | time      |
| logs               | errors, top error senders, rsyslog and its queue, journal size                           | logging   |
| mariadb (or mysql) | (the MariaDB / MySQL health check)                                                       | mariadb   |

Add `use_act=true` and ACT reads everything collected and explains the root cause (ADDING\_ACT.md).

**All settings** of this job, with their defaults: troubleshoot.

## Certificate report (playbooks/cert\_report.yml)

The certs check on every host, and **one table**, soonest expiry first: days left, date, host, subject, and where the certificate is. Hand it to whoever renews certificates. It fails the hosts with a certificate inside `check_certs_warn_days`, so a workflow can open tickets (`site_fail_on: []` keeps it green: a report only).

**Emailed** with `report_email_to` (Linux and Windows: `win_cert_report.yml`). In the email the table is coloured: **red** = expired, **amber** = expires within `check_certs_warn_days` (30), **blue** = within `check_certs_notice_days` (60), the rest plain with a green "ok". Windows: `win_check_certs_warn_days` / `win_check_certs_notice_days`.

**Found by itself** (the report only; `cert_report_discover: false` turns it off):

- **TLS ports:** every listening TCP port is asked once for its certificate, on the address it listens on, with a 3-second limit (`check_certs_probe_timeout`), at most 25 ports (`check_certs_discover_ports_max`). Ports that are plain text are never touched (`check_certs_discover_ports_skip: 22, 25, 53, 80, 389, 3306, 5432 ...`), nor the ones you listed in `check_certs_endpoints`. Shown as TLS port 8443, java (`https://127.0.0.1:8443`).
- **Java keystores:** `*.jks`, `*.keystore`, `*.p12`, `*.pfx` in `check_certs_discover_keystore_dirs` (`/etc`, `/opt`, `/srv`, Tomcat's folders), and the ones running Java programs use (`-Djavax.net.ssl.keyStore=...`). A JKS is read without its password; a PKCS12 that needs one is listed under **Found but not checked** - no password is ever guessed. To check it, add it to `check_certs_keystores` with `password_env` and a *Keystore password* credential. The JDK's own trust stores (`cacerts`) are skipped.

The weekly health check (`certs` in `health_checks`) does not search by itself, so it opens no tickets for certificates nobody listed; set `check_certs_discover_ports` / `_keystores: true` on it when you want that.

**All settings** of this job, with their defaults: `check_certs`.

## POA&M status (playbooks/poam\_status.yml)

Reads `poam/poam.csv` (see `poam/README.md` for the columns and how to export from eMASS):

| Finding                                                                                    | Means                                                                        |
|--------------------------------------------------------------------------------------------|------------------------------------------------------------------------------|
| POA&M 1001 [CAT II]: ... - OVERDUE by 12 days                                              | open item past its scheduled completion date (critical)                      |
| POA&M 1002 ...: due in 20 days                                                             | due within <code>poam_due_soon_days</code> (30)                              |
| ... scheduled completion date "TBD" is missing or not a date                               | fix the date in the list                                                     |
| open CAT I finding V-123456 (...) on 2 asset(s) in collection Production has no POA&M item | STIG Manager cross-check: an open CAT I/II finding that no POA&M item covers |

It also lists "maybe ready to close": items whose STIG findings are no longer open in STIG Manager.

**The STIG Manager cross-check** is optional. It needs a Keycloak client that may use the client-credentials grant, with the scope `stig-manager:collection:read`, and a grant in the collections it reads. **Your realm owner creates that client**; nothing here changes the realm. Then create the *STIG Manager API* credential (token URL `https://<keycloak>/realms/<realm>/protocol/openid-connect/token`, client ID, secret), and set `poam_stigman_api`: `https://<stigman>/api`.

**All settings** of this job, with their defaults: `poam`.

## ServiceNow tickets (playbooks/servicenow\_tickets.yml)

A workflow step after a check. For each finding at or above `servicenow_min_severity`:

- **no open ticket** for that problem on that host › it opens an incident: short description [`site-health`] `host: summary`; description with the host, check, finding, the `look: command`, and ACT's analysis if ACT ran; urgency/impact from the severity; assignment group, category and configuration item from the settings.
- **an open ticket exists** › it adds a work note `Still failing at ... (servicenow_note_repeats)`.
- **a host that was targeted but never reported** › a ticket too (unreachable, login failed).
- **a problem that cleared** (its check ran again and passed) › a work note on its ticket, or it resolves it with `servicenow_resolve_cleared: true`.

Tickets are matched by a correlation ID (`site-health:<host>:<check>:<finding>`, hashed), never by text, so the same problem always updates the same ticket. Without the ServiceNow API credential it only prints what it would open. If no check results reach it at all (the check step died before checking any host), it fails, so a broken schedule shows red instead of green.

**All settings** of this job, with their defaults: `servicenow`.

## ServiceNow - test ticket (playbooks/servicenow\_test\_ticket.yml)

Proves the ServiceNow setup end to end, with the same calls and settings as the tickets step. Six steps, each PASS, WARN or FAIL with what usually causes a failure:

| Step         | What it does                   | FAIL usually means                                                                                                                                           |
|--------------|--------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1 connection | logs in and reads one incident | certificate (CA), DNS, firewall or proxy; HTTP 401 = user or password; 403 = no <code>itil</code> role; 404 = a path after the host name in the Instance URL |

| Step        | What it does                                                                            | FAIL usually means                                                                                                                                  |
|-------------|-----------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------|
| 2 settings  | <code>servicenow_assignment_group</code> and <code>servicenow_caller</code> exist       | the name differs from ServiceNow's                                                                                                                  |
| 3 open      | opens a low-priority incident:<br><code>[site-automation test] ... safe to close</code> | the account may not create incidents ( <code>itil</code> ), or ServiceNow refused a value (it says which)                                           |
| 4 read back | reads it back: number, state, group, caller; prints a link                              | the account may not read it; WARN if the group did not stick                                                                                        |
| 5 work note | adds a work note                                                                        | the account may not update incidents                                                                                                                |
| 6 resolve   | resolves it ( <code>servicenow_test_resolve: false</code> leaves it open)               | WARN only: your instance needs another resolved state or close code ( <code>servicenow_resolved_state</code> , <code>servicenow_close_code</code> ) |

As *Check* (a dry run) steps 1 and 2 run, and 3 to 6 are only printed. Set up, run and verify it: `SERVICENOW_SETUP.md`.

## ServiceNow health (playbooks/servicenow\_health.yml)

From the controller: does the instance answer its REST API, and how fast (`servicenow_slow_seconds`)? Is every MID Server **Up** and **validated**? Hourly is a good schedule. Alert with an AAP **notification** (email) on failure, not with a ticket: a broken ServiceNow cannot take tickets. The MID Server **hosts** (the Linux service, its keystore) are covered by the health check: see `inventories/example/group_vars/servicenow_mid_hosts.yml`.

**All settings** of this job, with their defaults: `servicenow`.

## Patch hosts (playbooks/patch\_hosts.yml)

Installs updates with `dnf` (`patch_security_only` for security fixes only; `patch_exclude` for packages never updated here). Then:

1. **One host at a time** (`patch_serial`). The rollout **stops at the first host that fails**.
2. Refuses with less than `patch_min_free_mb` free in `/var`.
3. **No automatic restarts** unless `automatic_restarts: true` (the survey or the template's Variables). Without it, a host that needs a reboot is listed ("reboot it by hand, or run again with `automatic_restarts: true`"). With it, the job reboots only when `needs-restarting -r` says so (`patch_reboot: when_needed`), and **never** hosts in `patch_never_reboot_groups` (`aap_hosts`).
4. After the update (and reboot), **every service that was running before must be running again**. If one is not, the host fails and the rollout stops.
5. Never patches hosts in `no_patch` (vendor appliances), `aap`, `aap_hosts`, `sat`, `netapp` or `netapp_console_hosts`. The AAP host is **always** protected: even a list you write yourself (`patch_never_patch_groups`) cannot remove `aap` or `aap_hosts`. It also refuses a host that is the machine running the job, but only when the connection is `local` or the name is `localhost`; a normal SSH host that happens to be the controller is stopped by the group rule above, not by this one.

Run it as **Job type: Check** first: it lists what would be updated and changes nothing (`DRY_RUNS.md`).

**All settings** of this job, with their defaults: `patch`.

## Apply approved ACT fix (playbooks/act\_fix\_approved.yml)

The step after an approval in the *Fix with approval (ACT)* workflow. It runs **exactly** the commands ACT proposed and a person approved, then re-runs the checks that found the problem. It refuses to run on its own, and skips diagnose-only hosts (`aap_hosts`, `netapp_console_hosts`): there a person applies the fix by hand. See `ADDING_ACT.md`.

**All settings** of this job, with their defaults: `site_act`.

## VMware jobs and reports (playbooks/vm\_\*.yml)

Restart, shut down, snapshot, notes, change VLAN; the secure boot, datastore, snapshot and alarms reports; the snapshot cleanup; ACT's analysis of the alarms. They work through vCenter: setup, every report's sections and troubleshooting are in `VMWARE.md`.

**All settings** of this job, with their defaults: `vmware_vm`.

## Windows hosts

Windows has its own playbooks, checks and troubleshooting areas (for example `updates` and `security` in *Windows troubleshoot*). They are described, with their settings, in `WINDOWS.md`. Settings for them are read the same way as any other: `VARIABLES.md`.

## Not built yet (and why)

| Idea                                                                    | Why not yet                                                                                                                                                              |
|-------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| SCAP scan › import into STIG Manager                                    | Needs the same Keycloak client as the POA&M cross-check, plus a decision on SCC vs OpenSCAP content. STIG Manager's own importer (stigman-watcher) is the supported path |
| Certificate renewal workflow                                            | Depends on how your PKI issues certificates. For STIG Manager, renewal already works: update the TLS credential, run <i>STIG Manager - deploy</i>                        |
| AAP self-health (job failure rate, execution environment and token age) | Needs an AAP API credential; AAP's own dashboards cover most of it today                                                                                                 |
| Container runbooks (image updates with rollback, drift)                 | Deferred on purpose; <i>Service watch</i> covers container outages                                                                                                       |
| MariaDB backups, security baseline, password rotation                   | Out of scope for now: health only, as asked                                                                                                                              |

## 8. Adding ACT (GenAI) later

The runbooks work without ACT. Add it when the checks are running well and you want the next step done for you: **why** it broke and **what** fixes it. This page shows how the runbooks and ACT fit together, how to turn it on per template, and how to write your own check so it works with ACT automatically.

### How they fit: checks find WHAT, ACT finds WHY

| check roles (deterministic, cheap)                                            | site_act (the bridge)                                                                                   | ACT (GenAI)                                                            |
|-------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------|
| -----                                                                         | -----                                                                                                   | -----                                                                  |
| disk, services, selinux, mariadb ... --> each finding: summary + hint command | findings as one line each --> + the hint commands' output + your level (explain / diagnose / self-heal) | root cause, evidence, the exact fix (runs, or proposes, per the level) |

- The **checks** decide *whether* something is wrong. They are plain Ansible: the same answer every time, and they are free. The model is only called for hosts with findings.
- Every finding carries a **hint**: the read-only command a person would run next. ACT gets those commands' output as evidence. A new check you write is therefore usable by ACT with no extra work.
- `roles/site_act` is the only place that talks to ACT. Every runbook uses it the same way. Turning ACT on is a variable, not a code change.
- Host names, IP addresses and user names are **pseudonymized** before anything reaches the model (ACT 0.6.18 and newer). Its replies are translated back on the host.

### The three levels

| Level            | ACT does                                                                                                                                                                                                                             | Changes to hosts      | Approval step?                                                                                |
|------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------|-----------------------------------------------------------------------------------------------|
| <b>explain</b>   | reads the evidence (the hint commands' output, collected by Ansible) and writes the root cause and a suggested fix. <b>It runs no commands itself</b>                                                                                | none                  | no                                                                                            |
| <b>diagnose</b>  | investigates by running <b>read-only</b> commands itself, then proposes exact fix commands. Every change is refused and waits for a person                                                                                           | only after approval   | yes: the job fails with <b>NEEDS APPROVAL</b> › approval step › <i>Apply approved ACT fix</i> |
| <b>self-heal</b> | like <b>diagnose</b> , but may <b>run fixes that match site_act_allow</b> by itself (for example <code>systemctl restart chronyd</code> ). Anything else still waits for approval. The checks then run again to prove the fix worked | only fixes you listed | only for the other fixes                                                                      |

Start with **explain**. Move a runbook to **diagnose** when ACT's explanations have been right for a while. Use **self-heal** only for fixes you would let a junior admin do unasked: start a stopped service, restart a hung one.

**Never fixed automatically, at any level:** hosts in `site_act_diagnose_only_groups` (default: `aap`, `aap_hosts`, `sat`, `idm`, `netapp`, `netapp_console_hosts`; `aap` and `aap_hosts` are always kept, whatever list you

write). There `self-heal` becomes `diagnose`, and the apply step skips them: it prints the approved command for a person to run by hand.

## Turn it on

### 1. The ACT model key credential

Create the credential type `ACT model key` (`aap/credential_types/act_model_key.yml`) and one credential with your key: the GenAI key for GenAI.mil, or the AskSage key for Ask Sage (SETUP\_AAP.md, steps 2-3). For the GenAI beta proxy, also create `ACT GenAI beta key` (`act_genai_beta_key.yml`). Attach the credential(s) to the **Health check** and **Troubleshoot** templates (and *Service watch - check*), next to `Linux ssh (sudo)`. This is the only place a key lives. The apply steps do not call the model: they run the approved commands themselves, so they need no key.

### 2. Pick the model provider

ACT can talk to three providers. Pick one with `site_act_provider`:

| site_act_provider            | Provider                              | Key (credential)                                                                       | Also set                                                                                                                                                   |
|------------------------------|---------------------------------------|----------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>genai</code> (default) | GenAI.mil                             | <i>ACT model key</i> , <b>GenAI API key</b> field                                      | nothing (optional: <code>site_act_model</code> )                                                                                                           |
| <code>asksage</code>         | Ask Sage                              | <i>ACT model key</i> , <b>AskSage API key</b> field                                    | <code>site_act_url</code> : your organization's Ask Sage API URL. The default is <code>https://api.genai.army.mil/server/openai/v1/chat/completions</code> |
| <code>genai-beta</code>      | GenAI.mil beta proxy (preview models) | <i>ACT GenAI beta key</i> ( <code>aap/credential_types/act_genai_beta_key.yml</code> ) | nothing                                                                                                                                                    |

The settings:

| Setting                                                 | What                                                                                                                                                                          | Default                        |
|---------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------|
| <code>site_act_provider</code>                          | <code>genai</code> , <code>asksage</code> or <code>genai-beta</code>                                                                                                          | <code>genai</code>             |
| <code>site_act_models</code>                            | the model <b>for each provider</b> , so the model follows the provider you pick, e.g. <code>{genai: gemini-3.8-flash, asksage: gpt-5.6-sol-gov}</code>                        | ACT's default for the provider |
| <code>site_act_urls</code>                              | the API URL for each provider, e.g. <code>{asksage: https://&lt;host&gt;/server/openai/v1/chat/completions}</code>                                                            | ACT's default for the provider |
| <code>site_act_model</code> , <code>site_act_url</code> | one model / URL for this run, whatever the provider (a survey answer)                                                                                                         | empty: use the two above       |
| <code>site_act_ca</code>                                | a CA bundle file <b>on the hosts</b> , if the provider's certificate comes from a CA the hosts do not trust yet (for example a DoD CA not in <code>/etc/pki/ca-trust</code> ) | the host's normal trust store  |

### Where to set them:

- **For everyone:** `playbooks/group_vars/all.yml` in this project (the file has these lines ready as comments). A setting there beats the *inventory's* Variables box in AAP, so put it in one of the two places, not both. Which place wins in every case: `VARIABLES.md`.

YAML

```

site_act_provider: genai          # the default provider
site_act_models:                 # the model follows the provider you pick
  genai: gemini-3.8-flash
  asksage: gpt-5.6-sol-gov
site_act_urls:                   # only if your Ask Sage is not ACT's default URL
  asksage: https://<your ask sage host>/server/openai/v1/chat/completions

```

To make Ask Sage the default, change the first line to `site_act_provider: asksage`.

- **For one group:** the group's settings file (`playbooks/group_vars/<group>.yaml`), or the group's **Variables** box in AAP. For example `site_act_provider: asksage` there, and only that group's hosts use Ask Sage (with its model from `site_act_models`).
- **For one template:** its **Variables** box (*Extra variables*), for example `site_act_provider: asksage`. It beats everything above, for every run of that template.
- **Per launch:** the survey questions in step 3 (a survey answer beats everything too).

The API **key** is never a variable: it comes from the *ACT model key* credential on the template.

These settings apply to every runbook that uses ACT, including *Service watch* and *Apply approved ACT fix*. If the key for the chosen provider is missing, ACT does not run: the job says `ACT did not run: this job has no API key for provider asksage. Attach ...`, and the findings are still reported.

**Network:** ACT runs **on each host**, not in AAP. So each host must be able to reach the provider's URL over HTTPS. If they go through a proxy, add it: `site_act_env: {HTTPS_PROXY: "http://proxy.yoursite.mil:8080"}`.

**Models served only on** `/v1/messages` (ACT 0.6.19 and newer). Some gateways offer a model only through the Anthropic endpoint and answer HTTP 400 for it on `.../v1/chat/completions`. ACT tries the other endpoint by itself, with nothing to set. To test your models and set the switches on purpose, follow Both endpoint formats and `:probe` below.

**How long ACT may take.** `site_act_model_timeout` (default **600** seconds) is the time the model may take for one answer, retries after a provider error included (ACT's `GENAI_TIMEOUT`; ACT's own default, 90, is too short for a long analysis or a busy provider). `site_act_timeout` (default **1200**) is the whole ACT run, several answers. They apply to every ACT job: the runbooks with `use_act`, the VMware alarms and capacity analyses, *Service watch*. A `model turn exceeded the 600s total timeout after a transient provider failure` means the provider was failing and retrying for that long: run it again; if it repeats, raise both in `all.yaml` (`site_act_model_timeout: 900, site_act_timeout: 1800`).

**Streaming.** Every ACT job has the model stream its answer (`ACT_STREAM=1`): data keeps flowing while it writes, so a firewall or proxy that cuts quiet connections (`network error reaching API: connection reset by peer`, often on the longer answers only) no longer does. Where a gateway cannot stream, ACT notices and sends normal requests by itself. To turn it off: `site_act_env: {ACT_STREAM: "0"}`.

**"ACT gave no analysis: the model answered as a chat".** Now and then a model loses the thread and answers like a chat window ("I am ready. What task or command would you like me to assist you with?") instead of analysing the findings; whatever it tried to run then is no fix. The report says so (ACT status `no_analysis`) and nothing is passed on for approval. Run the job again; if it repeats with one model, switch that template to another (`site_act_models`).

Anything else ACT should know about your site goes in the same places (for example `playbooks/group_vars/all.yaml`):

YAML

```
site_act_extra_instructions: >-
  RHEL 9 servers under DISA STIG. fapolicyd and SELinux are enforcing. Never suggest disabling
  them.
```

All settings, with their defaults: `site_act`.

### 3. Survey questions

On **Health check** and **Troubleshoot** › **Survey** › **Create survey question**:

| Question                                                                                 | Variable                       | Type                            | Choices                                   | Default                  |
|------------------------------------------------------------------------------------------|--------------------------------|---------------------------------|-------------------------------------------|--------------------------|
| Ask ACT (GenAI) about the findings?                                                      | <code>use_act</code>           | Multiple Choice (single select) | <code>no, yes</code>                      | <code>no</code>          |
| What may ACT do?                                                                         | <code>site_act_level</code>    | Multiple Choice (single select) | <code>explain, diagnose, self-heal</code> | <code>explain</code>     |
| Which model provider? (optional)                                                         | <code>site_act_provider</code> | Multiple Choice (single select) | <code>genai, asksage, genai-beta</code>   | the one in your settings |
| Which model? (optional, blank = the provider's model from <code>site_act_models</code> ) | <code>site_act_model</code>    | Text, not required              |                                           | blank                    |

Leave out the last two if everyone uses the same provider. Your setting (section 2) then applies.

### 4. Try explain by hand

Launch **Troubleshoot** on a host with a known problem: `ts_area = service`, `ts_service = the unit`, `use_act = yes`, `site_act_level = explain`. In the output, read `ACT | what ACT says`. With ACT on, the host's ServiceNow ticket includes this analysis too.

### 5. diagnose with the approval workflow

Build *Fix with approval (ACT)* (WORKFLOWS\_AND\_SCHEDULES.md). What happens:

1. The health check finds, for example, `required service chronyd.service is inactive`.
2. ACT (`diagnose`) looks at the unit, its log and the config, and proposes `systemctl start chronyd`. The command is refused, so the job fails: `NEEDS APPROVAL - ... ACT proposes: systemctl start chronyd`.
3. The workflow stops at the approval step. The approver reads ACT's report and approves.
4. **Apply approved ACT fix** runs **exactly** `systemctl start chronyd` itself: no model and no key, so nothing can be reworded between the approval and the fix. A last check refuses a few commands even after approval (reboot, `mkfs`, `usermod ...`; see APPROVED\_COMMANDS.md).
5. The same checks run again, without ACT, and the job is green only if the problem is gone. The tickets step notes the ticket as cleared.

### 6. self-heal for the fixes you trust

```
playbooks/group_vars/all.yml (or the settings file of one group)

site_act_allow:
  - 'systemctl (start|restart) (chronyd|rsyslog|crond)'
```

Each entry is a regular expression that must match the **whole** command. Keep them narrow: name the services, never `.*`. Then pick `site_act_level = self-heal`. After ACT applies an allowed fix, the checks with findings run again. The report shows `cleared: ...` or `still there: ...` and reflects the state **after** the fix.

## Both endpoint formats and :probe (ACT 0.6.19 and newer)

**What this is for.** A model gateway can offer a model on two web addresses: the *OpenAI* one (the provider URL, ending in `/v1/chat/completions`) and the *Anthropic* one (ending in `/v1/messages`). Some models answer on only one of them. If ACT prints `HTTP 400 ... bad request` for a model that the gateway lists, this is the usual cause. ACT can now speak both, and `:probe` tests which one each model needs.

### Two facts that shape the steps:

- `:probe` is a command you type at the ACT prompt. **It cannot run inside an AAP job** (jobs are not interactive). So you run it **once, by hand, on one lab host**, and then put the result into settings the jobs read.
- In a job, ACT already works out the format by itself (one refused request, then it switches), so the jobs may need no change at all. The settings below only save that one wasted request per run, or point at an Anthropic address that is not the default.

**Not verified against a real gateway.** These steps were written from ACT's own documentation and tests. Nobody has run them against your gateway yet; treat the first run as the test.

### Step 1: check which ACT the project carries

`vendor/act/VERSION` must say `0.6.19` or newer. This release carries **0.6.23**. It sends the key only to an `https://` provider URL (an `http://` URL is refused unless `site_act_env: {ACT_ALLOW_HTTP: "1"}`), and more dangerous commands always need a person (see `APPROVED_COMMANDS.md`). New in 0.6.22 and 0.6.23, all automatic:

- **A bigger output limit for thinking models (0.6.23).** Gemini 2.5 and later, GPT-5 and the o-series think before they answer, and the thinking counts against the output limit. At ACT's old limit (4096 tokens) the thinking could use it all and the answer came back **empty** - on GenAI.mil every Gemini model did that in `:probe`. These models now get 16384 (others keep 4096). It is a ceiling, not a charge: you pay for the tokens used. To force one value: `site_act_env: {ACT_MAX_TOKENS: "8000"}`.
- **Temperature per model.** Gemini 3 models (and GPT-5 / o-series reasoning models) get **no** temperature, so they use their own default (1.0): Google's Gemini 3 developer guide says to keep it at 1.0 and that lower values can make the model loop. Other models keep 0.2. To force a value for every model: `site_act_env: {ACT_TEMPERATURE: "0.2"}`.
- **Fewer empty answers.** A reply cut off at the output limit (the model's thinking used it all) is asked again once with a higher limit; an empty reply is asked again once with a strict JSON schema of ACT's actions; a reply the gateway's content filter blocked is reported as that.
- **Structured output.** When a model refuses ACT's function-calling format, ACT asks for a strict JSON schema instead of "any JSON object", so the gateway itself guarantees a valid answer.
- **HTTP 429 (too many requests):** ACT waits as long as the gateway asks (`Retry-After`). A spent credit quota still stops the run (it would not help to wait).
- **Clearer errors.** A model alias the gateway retired (GenAI.mil retires aliases 60 days after deprecation) is reported as that, not as a missing endpoint or a permission problem; a locked or wrong key says to enter a new one.

- **Interactive only:** replies stream in (a counter shows progress) and **Esc** cancels a reply that is still coming. Jobs do not stream (nothing to watch, and some proxies break streaming).
- **Tool-result turns** (command results sent back the way function-calling models expect): only for models `:probe` confirmed - see step 5. If it is older, update it first: README section

"Updating ACT" (`scripts/update-act.sh`).

## Step 2: put ACT on one lab Linux host

Any host that can reach the gateway will do; it does not have to be one of your managed hosts. Copy the file `vendor/act/act` there (WinSCP, or `scp` in PowerShell), then on the host:

```
shell
chmod 700 act
```

## Step 3: give it your key without writing the key anywhere

Type this on the host. It asks for the key without showing it, and keeps it only in that terminal session (use `GENAI_BETA_KEY` or `ASKSAGE_KEY` for the other providers):

```
shell
read -rs GENAI_KEY; export GENAI_KEY
```

Paste the key, press Enter, and do not `echo` it. If the provider URL is not ACT's default, also `export GENAI_URL='https://<your gateway>/v1/chat/completions'`.

## Step 4: run `:probe`

```
shell
./act
```

At the ACT prompt:

```
:probe all
```

`:probe all` tries every model the key lists, on both endpoints, with a small request and then ACT's full request. For each it shows what worked and, for a refusal, **the gateway's own reason**. `:probe <model name>` tests just one; `:probe model1 model2` (spaces or commas) tests several. Write down, for each model, which endpoint worked (`openai` or `anthropic`). Leave ACT with an empty line (or Ctrl-D). From 0.6.23 `full OK` means the model gave a usable answer; `full empty (...)` or `full no action (...)` says what came back instead (for example `finish_reason=length`: the model ran out of output room), and a test that needed more room says `OK (needed a higher output limit: N)` - ACT remembers that limit for the model.

From 0.6.22 each model also gets these lines (under its `basic` / `full` lines):

| Line                                                                               | Meaning                                                       | What to do                                |
|------------------------------------------------------------------------------------|---------------------------------------------------------------|-------------------------------------------|
| <code>stream OK / stream not supported (...)</code>                                | whether replies can stream in (interactive use only)          | nothing                                   |
| <code>structured output OK (strict) / OK (non-strict) / not supported (...)</code> | whether the gateway enforces ACT's JSON schema for this model | nothing: ACT uses the best one that works |

| Line                                                                 | Meaning                                                             | What to do                                                 |
|----------------------------------------------------------------------|---------------------------------------------------------------------|------------------------------------------------------------|
| <code>tool results OK / tool results not supported (...)</code>      | whether this model accepts command results as tool turns            | <b>write it down:</b> step 5                               |
| <code>temperature: model default / temperature: 0.2</code>           | what ACT sends                                                      | nothing (to force a value: <code>ACT_TEMPERATURE</code> )  |
| <code>output limit 16384 (thinking model) / output limit 4096</code> | the most the model may write per answer, thinking included (0.6.23) | nothing                                                    |
| <code>full empty (finish_reason=length ...)</code>                   | even the higher limit was not enough for that test                  | tell the maintainer; or <code>ACT_MAX_TOKENS</code> higher |

The result is remembered for the session. If a setup file already exists for your user on that host (`~/.config/act/config.json`), ACT also saves the format there (only the format table, never the key).

**Windows.** ACT for Windows (`vendor/act-windows/act.ps1`) has the same `:probe` command. Run it the same way on a lab Windows host, in PowerShell, and enter the key with `:setup`. Its result goes into the Windows jobs' settings as above. (Not run by hand for this page.)

### Step 5: tell the jobs

A job does not read the lab host's file (it runs as another account, on another host), so put the result in settings, in the same places as the other ACT settings (section 2):

| Your probe result                                                                                                                           | Setting in playbooks/group_vars/all.yml (or a group file)                                                                                                                                                                                    |
|---------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| All the models you use work on the OpenAI endpoint                                                                                          | nothing to set                                                                                                                                                                                                                               |
| Some models work only on the Anthropic endpoint                                                                                             | nothing to set (ACT learns it during the run); optional, to skip the one refused request per run when <b>every</b> model you use is Anthropic-only: <code>site_act_env: {ACT_API_FORMAT: anthropic}</code>                                   |
| The Anthropic endpoint is at a different address than the OpenAI one with <code>/messages</code> in place of <code>/chat/completions</code> | <code>site_act_env: {GENAI_ANTHROPIC_URL: "https://&lt;host&gt;/.../v1/messages"}</code><br>( <code>ASKSAGE_ANTHROPIC_URL</code> or <code>GENAI_BETA_ANTHROPIC_URL</code> for the other providers)                                           |
| A model you want is listed for the key but works on neither                                                                                 | it is not available to your key; ask the gateway's owner. Choose another model in <code>site_act_models</code>                                                                                                                               |
| <code>tool results OK</code> for the model your jobs use (0.6.22)                                                                           | optional: <code>site_act_env: {ACT_TOOL_RESULTS: tool}</code> . Jobs then send command results as tool turns, which function-calling models handle best on long runs. If a model refuses them anyway, ACT switches that model back by itself |
| <code>tool results not supported</code>                                                                                                     | nothing to set (the default sends results as ordinary messages, as before)                                                                                                                                                                   |

`ACT_API_FORMAT: anthropic` applies to **every** model ACT uses in the job. If you use both kinds, leave it out.

## Many hosts at once: `site_act_concurrency`

Every host in one job uses the same key, and GenAI.mil's default quota is **60 requests and 200,000 tokens a minute** per key (and a daily credit budget). A Health check with `use_act` on 20 hosts with findings would otherwise run ACT on as many hosts at once as the job's forks allow, and the gateway answers HTTP 429. `site_act_concurrency` (default **3**) caps how many hosts run ACT at the same time; the rest wait their turn. **0** = no cap. ACT itself also waits when the gateway sends a 429 with `Retry-After`.

### How to verify

1. **The probe result:** in step 4 every model you plan to use shows one endpoint that worked.
2. **A job:** launch **Troubleshoot** with `use_act = yes`, `site_act_level = explain`, and the model in `site_act_model` (or set through `site_act_models`) on a host with a known finding. In the output, task `ACT | what ACT says` holds an analysis, not an `HTTP 400` message.
3. **The old symptom is gone:** search the job output for `400`. A line saying ACT switched to the other endpoint is fine; a final `HTTP 400` for the model is not (then the probe said neither endpoint works: step 5, last row).

If a step fails, the message after `HTTP 400` is the gateway's reason: copy it exactly when you ask its owner for help, but remove the key and host names first.

## ACT for vCenter alarms (it runs on the AAP side)

**VM - alarms ACT analysis** (VMWARE.md, "Alarms") uses ACT differently from the health checks: ACT runs **inside the job's execution environment on the AAP node**, not on a managed host, and only reads the evidence the job collected from vCenter (it runs no command). So the ESXi hosts need nothing; the **AAP node** needs HTTPS to the model's URL, and the job checks that first. It uses the same settings as everything else here: `site_act_provider`, `site_act_models`, `site_act_url`, `site_act_env` (a proxy), `site_act_ca` (a CA file **in the execution environment**, e.g. one from the project: `{{ playbook_dir }}/files/ca/model-ca.pem`), and the **ACT model key** credential on its template.

## Write your own check (it works with the report, tickets and ACT automatically)

A check is a role named `check_<name>` that **adds findings** to `site_findings`. That is the whole contract. As an example, here is a check that `/tmp` is not world-writable without the sticky bit.

1. Create `roles/check_tmp/defaults/main.yml`, with its settings and a comment saying what it checks:

```
YAML
---
# TMP: /tmp and /var/tmp must have the sticky bit (mode 1777).
check_tmp_paths: [/tmp, /var/tmp]
```

2. Create `roles/check_tmp/tasks/main.yml`. Collect with read-only commands, and always use `changed_when: false` and `check_mode: false`: read commands then run in a dry run too. Turn what you saw into findings, and add them:

YAML

```

---
- name: tmp | modes
  ansible.builtin.command:
    argv: "{{ ['stat', '-c', '%a %n'] + check_tmp_paths }}"
  register: _tmp
  changed_when: false
  failed_when: false
  check_mode: false

- name: tmp | findings
  ansible.builtin.set_fact:
    _tmp_findings: >-
      {%- set f = [] -%}
      {%- for line in _tmp.stdout_lines if line.split()[0] != '1777' -%}
      {%-   set path = line.split()[1] -%}
      {%-   set _ = f.append({'check': 'tmp', 'id': 'tmp:' ~ path, 'severity': 'warning',
                            'summary': path ~ ' has mode ' ~ line.split()[0] ~ ' (expected
                            1777)',
                            'hint': 'stat ' ~ path ~ '; findmnt ' ~ path}) -%}

      {%- endfor -%}
      {{ f }}

- name: tmp | add them to this host's findings
  ansible.builtin.set_fact:
    site_findings: "{{ site_findings | default([]) + _tmp_findings }}"

```

- Let the health check know it: in `playbooks/health_check.yml`, add `tmp` to `_health_sets` (to `weekly` and `all`), and add `tmp` to the survey choices of the Health check template.
- Test it on one host: `ansible-playbook playbooks/health_check.yml -l <host> -e health_checks=tmp`. Then break the condition on purpose (`chmod 777 /tmp` on a test host) and check that the finding appears. A check that never fires proves nothing.

That is all. The report prints it, `set_stats` publishes it, ServiceNow tickets it (the `id` keeps it on one ticket), and with `use_act=yes` ACT receives the summary and runs the `hint` as evidence.

### Rules for a good finding:

- summary:** one line a person understands, with the number and the limit (`/var is 96% full (critical at 95%)`).
- id:** stable across runs, with no numbers that change (`disk:/var`, not `disk:/var:96`). Otherwise every run opens a new ticket.
- hint:** **read-only**, runs in under a minute (`timeout 60 ...`), and shows the cause, not just the symptom. It runs as root when ACT is on (Ansible runs it in `explain` mode and hands ACT the output), so never put a command that changes something there.
- severity:** `critical` = act now (outage, security); `warning` = act soon.

## Use ACT in a playbook of your own

Any playbook that fills `site_findings` can hand them to ACT the same way:

YAML

```
- name: My check
  hosts: my_hosts
  become: true
  gather_facts: false
  tasks:
    - ansible.builtin.include_role: {name: site_findings, tasks_from: start.yml}
    - ansible.builtin.include_role: {name: site_findings, tasks_from: run_check.yml}
      vars: {_site_check: tmp}                                # runs roles/check_tmp
    - ansible.builtin.include_role: {name: site_act}
      when: use_act | default(false) | bool and site_findings | length > 0
    - ansible.builtin.include_role: {name: site_findings, tasks_from: report.yml}
```

`report.yml` goes last. It prints, publishes (for tickets and the approval workflow), and fails the host when there are findings or a fix waits for approval.

## What ACT is told, and what it cannot do

- Its task: the findings (one line each), "find WHY, confirm it with evidence, give the exact fix; several findings may share one cause", plus your `site_act_extra_instructions`.
- `explain`: it runs nothing. Ansible ran the hint commands.
- `diagnose`: only commands ACT can prove are read-only run on their own. Everything else is refused and becomes a proposal.
- `self-heal`: in addition, only commands matching `site_act_allow`.
- Apply step: no model at all. The playbook runs only the approved commands, character for character, in order, and stops at the first one that fails. Run as *Check*, it only shows what it would run (a dry run: see WORKFLOWS\_AND\_SCHEDULES.md, "Rehearse it").
- Every run is recorded in the AAP job output. The ticket carries ACT's analysis. For the service-watch style incident log, see SERVICE\_WATCH\_DEMO.md.

## 9. Windows servers

The Windows runbooks work the way the Linux ones do. Checks find what is wrong and say it in plain words, with a command to look further. The same report, the same ServiceNow tickets and POA&M, and the same workflows and schedules follow. Only the checks themselves are different: each one is a PowerShell script that runs on the server over WinRM.

**Nothing extra to install on AAP or on the servers.** The playbooks use only Ansible's built-in `script` module, so no Ansible collection is needed. The execution environment needs the Python library for WinRM (`pywinrm`); step 2 shows how to check. The servers need WinRM over HTTPS, which you already use. The check scripts run in Windows PowerShell 5.1 and PowerShell 7, and also in Constrained Language Mode (see step 5).

| Template                         | Playbook                                       | What it does                                                                                                                       |
|----------------------------------|------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------|
| Windows connection test          | <code>playbooks/win_connection_test.yml</code> | run it first: can AAP reach the servers, with the right account?                                                                   |
| Windows health check             | <code>playbooks/win_health_check.yml</code>    | the checks below; daily and weekly sets                                                                                            |
| Windows troubleshoot             | <code>playbooks/win_troubleshoot.yml</code>    | on demand: an admin's first PowerShell commands for one kind of problem                                                            |
| Windows certificate report       | <code>playbooks/win_cert_report.yml</code>     | every certificate on every server, soonest expiry first                                                                            |
| Windows patch                    | <code>playbooks/win_patch.yml</code>           | install what Software Center offers (ConfigMgr); restart only with <code>automatic_restarts: yes</code> ; prove the server is back |
| ACT for Windows - install        | <code>playbooks/win_act_install.yml</code>     | put ACT on the servers in <code>C:\ProgramData\act</code> , locked down and unblocked; keep it current                             |
| ServiceNow tickets, POA&M status | the ones you have                              | unchanged: they take the Windows results the same way                                                                              |

**The checks** (`roles/win_check_<name>`):

| Set   | Check                    | Finds                                                                                                     |
|-------|--------------------------|-----------------------------------------------------------------------------------------------------------|
| daily | <code>disk</code>        | volumes too full                                                                                          |
| daily | <code>services</code>    | required services not running; automatic services that stopped with an error; services that keep crashing |
| daily | <code>performance</code> | CPU, memory or page file too busy; names the busiest processes                                            |
| daily | <code>time</code>        | Windows Time not running, not synchronized, synced too long ago, clock off                                |
| daily | <code>network</code>     | no gateway or DNS; names that do not resolve; ports that do not answer; a broken domain trust             |

| Set    | Check                 | Finds                                                                                                                                   |
|--------|-----------------------|-----------------------------------------------------------------------------------------------------------------------------------------|
| daily  | <code>eventlog</code> | unexpected shutdowns, crashes, disk errors, a cleared Security log, many errors, full logs                                              |
| weekly | <code>security</code> | Defender real-time protection off or old signatures; endpoint services down; firewall off; SMBv1; RDP without NLA; UAC off; TLS 1.0/1.1 |
| weekly | <code>audit</code>    | audit policy settings the STIG asks for; event log sizes; the log forwarding agent                                                      |
| weekly | <code>accounts</code> | Guest enabled; local passwords old or never expiring; inactive local accounts; unexpected local Administrators                          |
| weekly | <code>certs</code>    | certificates that expire or expired, including the WinRM listener AAP connects through                                                  |
| weekly | <code>patching</code> | a restart pending for installed updates; days since the last update                                                                     |

Everything is read-only. The one write is one Application event log entry per finding (source `site-health`), so a SIEM that collects event logs sees the findings too.

## Step 1. WinRM over HTTPS on the servers

You already use it. To see it on a server:

```
winrm enumerate winrm/config/listener
```

It shows a listener with `Transport = HTTPS`, `Port = 5986` and a `CertificateThumbprint`. The firewall must let port 5986 in from the AAP execution nodes. The `certs` check warns before that listener's certificate expires: when it expires, AAP can no longer connect.

## Step 2. The execution environment

The job runs inside an execution environment (EE), and that EE needs these Python libraries:

- `pywinrm`, always;
- `requests-ntlm`, for NTLM (the usual choice);
- or `requests-kerberos` plus a Kerberos client (`kinit`), for Kerberos.

AAP's *Default execution environment* has them. Your own EE (for example `platforms-ee`) may not. **The Windows connection test tells you:** its first part runs inside the EE and prints `YES` or `no` for each library. If `pywinrm` is missing, you have two options:

- use AAP's Default execution environment for the Windows templates; or
- add the libraries to your EE's definition and rebuild it:

```
YAML
dependencies:
  python:
    - pywinrm
    - requests-ntlm
```

## Step 3. The credential

### Automation Execution > Infrastructure > Credentials > Create credential:

- **Name:** `Windows (WinRM)`; **Credential type:** `Machine`.
- **Username:** `YOURDOMAIN\svc_aap` for NTLM. For Kerberos, use `svc_aap@YOURDOMAIN.MIL`, with the domain in capitals.
- **Password:** the account's password.
- **Privilege escalation:** leave it empty. An administrator's WinRM session is already elevated.

The account must be in the **local Administrators group** on every server, usually through a domain group added by GPO. The checks need that to read the Security log, the audit policy, local accounts and certificate stores.

## Step 4. The Windows inventory

Keep the Windows servers in their **own inventory**. Go to **Inventories > Create inventory**, name it `Windows servers`, and paste this in its **Variables** box:

```
YAML
ansible_connection: winrm
ansible_port: 5986
ansible_winrm_scheme: https
ansible_winrm_transport: ntlm # or kerberos
ansible_winrm_server_cert_validation: validate # needs your CA bundle in the EE, see below
ansible_winrm_ca_trust_path: /etc/pki/ca-trust/source/anchors/your-ca.pem # path inside the EE
```

- **Certificate validation.** `validate` means the EE checks the server's WinRM certificate against the CAs it trusts. EEs usually do not include your organisation's root CAs, so `validate` fails with `CERTIFICATE_VERIFY_FAILED` until you add your CA bundle to the EE and point `ansible_winrm_ca_trust_path` at it (the path inside the EE, as shown above). **Use this setting for real servers.** `ignore` keeps the traffic encrypted but does not check who the server is, so anyone who can intercept the connection could pose as it: use it **only** for a short test on a lab network, then switch back.
- **Hosts.** Add them by full name (`web01.yoursite.mil`); Kerberos needs the full name.
- **Groups.** Use whatever helps: by site, by role (`win_dc`, `win_web...`). The settings below can be set per group. `inventories/example-windows/hosts.yml` shows an example.

Linux and Windows do not mix: the Linux templates skip Windows servers, and the Windows templates skip Linux servers. So a template pointed at the wrong inventory changes nothing and opens no tickets.

## Step 5. Test the connection

Create a job template, then launch it with **Limit** set to one server:

| Field     | Value                                          |
|-----------|------------------------------------------------|
| Name      | <code>Windows connection test</code>           |
| Inventory | <code>Windows servers</code>                   |
| Project   | <code>site-automation</code>                   |
| Playbook  | <code>playbooks/win_connection_test.yml</code> |

| Field                 | Value                                  |
|-----------------------|----------------------------------------|
| Execution environment | one with <code>pywinrm</code> (step 2) |
| Credentials           | <code>Windows (WinRM)</code>           |
| Limit                 | empty, <b>Prompt on launch</b> ticked  |

A good result ends with:

```
Connected: WEB01 (Microsoft Windows Server 2022 Standard, build 20348), domain yoursite.mil
PowerShell 5.1.20348.2849, FullLanguage
Logged on as yoursite\svc_aap; local administrator: yes
OK: the Windows runbooks can run here.
```

`ConstrainedLanguage` instead of `FullLanguage` means WDAC or AppLocker restricts PowerShell on that server. The check scripts are written and tested for that mode; only two optional parts are skipped in it (TLS endpoints and the search for missing updates). Whether Ansible itself can run on such a server depends on your ansible-core version and the policy: if this test got through and printed its result, it can.

## Step 6. The job templates

For all of these, use: **Inventory** `Windows servers`, **Project** `site-automation`, the EE from step 2, **Credentials** `Windows (WinRM)`, **Limit** empty with **Prompt on launch** ticked. Leave **Variables** empty: with no `target`, they run on every server in the inventory.

### Windows health check

| Field                              | Value                                                                                                                                                                                                    |
|------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Playbook                           | <code>playbooks/win_health_check.yml</code>                                                                                                                                                              |
| Survey: <code>health_checks</code> | <i>Multiple Choice (multiple select):</i> <code>daily, weekly, all, disk, services, performance, time, network, eventlog, security, audit, accounts, certs, patching</code> . Default <code>daily</code> |

### Windows troubleshoot

| Field                           | Value                                                                                                                            |
|---------------------------------|----------------------------------------------------------------------------------------------------------------------------------|
| Playbook                        | <code>playbooks/win_troubleshoot.yml</code>                                                                                      |
| Survey: <code>ts_area</code>    | <code>overview, disk, performance, service, network, login, time, logs, updates, security</code> . Default <code>overview</code> |
| Survey: <code>ts_service</code> | <i>Text</i> , not required. "Service name (for area = service), e.g. W3SVC"                                                      |
| Survey: <code>ts_target</code>  | <i>Text</i> , not required. "host:port it cannot reach (for area = network)"                                                     |
| Survey: <code>ts_hours</code>   | <i>Integer</i> , default <code>2</code> . "How many hours of event logs to read"                                                 |

Each command's output appears under a `### what it shows` line, followed by the matching health checks. Findings do not fail this job.

**All settings** of this job, with their defaults: `win_troubleshoot`.

## Windows certificate report

| Field    | Value                                      |
|----------|--------------------------------------------|
| Playbook | <code>playbooks/win_cert_report.yml</code> |
| Survey   | none                                       |

## Step 7. Workflows and schedules

These are the same as for Linux, with the Windows templates. The full list, with every step and link, is in `WORKFLOWS_AND_SCHEDULES.md` (workflows 12 to 17):

| Workflow                           | Steps                                                                                                                                                              | Schedule                                                  |
|------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------|
| Windows daily health               | Windows health check ( <b>daily</b> ) › ServiceNow tickets ( <b>Run always</b> )                                                                                   | every day 06:15                                           |
| Windows weekly compliance          | Windows health check ( <b>weekly</b> ) › tickets; Windows certificate report › its own tickets box (both <b>Run always</b> )                                       | Monday 07:15                                              |
| Windows patch with a dry run first | Windows health check ( <b>disk, services</b> ) › Windows patch as <b>Check</b> › <i>Approval</i> › Windows patch › Windows health check ( <b>daily</b> ) › tickets | monthly, inside the servers' ConfigMgr maintenance window |

The ServiceNow tickets template runs on AAP itself, so it stays as it is.

## Start small: the pilot order

Everything Windows here is new, so try it on **one** server that does not matter much, in this order, and widen only when each step looks right:

1. **Windows connection test**: the account, WinRM and PowerShell work.
2. **Windows health check**: the findings make sense for that server.
3. **Windows patch as Check** (the dry run): it lists what Software Center would install.
4. **Windows patch for real**, inside that server's ConfigMgr maintenance window.
5. **ACT for Windows - install** on the same server (tell your security team first, see below).

Every step can be a dry run first: `DRY_RUNS.md`.

## Patching through Software Center

`playbooks/win_patch.yml` does, on each server, what an admin does in Software Center: install every software update your ConfigMgr admins deployed to it ("Install all"), required and available alike. AAP decides nothing about *which* updates: ConfigMgr does, as it does today. Applications listed in Software Center are not installed; only software updates are. Then:

1. **No automatic restarts** unless you ask for them with `automatic_restarts: yes` (the survey, or the template's Variables). Without it, the job installs the updates and lists the servers that need a restart: **A RESTART IS NEEDED to finish the updates ... automatic restarts are off**. A person restarts them (the post-check's `patching` finding keeps them visible). With it, the job restarts a server only when the updates or the ConfigMgr client say so, and never one in `win_patch_never_reboot_groups`.

2. **Prove the server is back** (after a restart, and also when none was needed). A new boot time, and every automatic service that ran before runs again. Delayed services get `win_patch_services_grace_min` (5 minutes) to start.
3. **Look again.** After a restart, servicing-stack and cumulative updates often reveal more updates. There is a second round (`win_patch_max_rounds: 2`), only when it restarted.

### Safety.

- One server at a time, and the rollout stops at the first server that fails.
- By default it only patches the group `win_patch_hosts`. Set `target:` in the template's Variables (for example `target: all`) and pick the servers with the Limit at launch.
- Servers in `win_patch_never_patch_groups` (default `no_patch`) are skipped.
- **Maintenance windows still apply.** Outside the server's ConfigMgr maintenance window, ConfigMgr holds the updates. The job then stops with `ConfigMgr will not install these now ... (waiting for a maintenance window)` instead of waiting forever. Run it inside the window, or ask the ConfigMgr admin for a window that covers the patch time.
- **A dry run first:** run the template as *Check*. It rescans and lists what it would install, and changes nothing (DRY\_RUNS.md). The *Windows patch with a dry run first* workflow puts that list in front of the approver.

**Needs:** the ConfigMgr client (Software Center) on the servers, and full PowerShell (not Constrained Language Mode) for the install call. Restarts use `shutdown.exe`, and AAP watches the WinRM port go down and come back. No Ansible collection is needed.

**Scheduling.** With `target: all` and no Limit, a scheduled run patches every server in the Windows inventory, one after another. Choose that on purpose. Otherwise, make one schedule per group (for example the test servers the week before the others), each with its own Limit.

### The job template:

| Field                                   | Value                                                                                                                             |
|-----------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------|
| Name                                    | <code>Windows patch</code>                                                                                                        |
| Playbook                                | <code>playbooks/win_patch.yml</code>                                                                                              |
| Credentials                             | <code>Windows (WinRM)</code>                                                                                                      |
| Job type                                | Run, with <b>Prompt on launch</b> ticked (Check = the dry run)                                                                    |
| Variables                               | <code>target: all</code> (or a group), and pick servers with the Limit                                                            |
| Survey: <code>automatic_restarts</code> | "Restart servers automatically when the updates need it?"<br>Choices <code>no</code> , <code>yes</code> . Default <code>no</code> |

If a server fails:

| What you see                                  | What to do                                                 |
|-----------------------------------------------|------------------------------------------------------------|
| <code>has no working ConfigMgr client</code>  | the SMS Agent Host service (CcmExec) is missing or stopped |
| <code>waiting for a maintenance window</code> | run inside the window, or ask the ConfigMgr admin          |

| What you see                               | What to do                                                                        |
|--------------------------------------------|-----------------------------------------------------------------------------------|
| these updates failed: ... (error 0x...)    | C:\Windows\CCM\Logs\UpdatesDeployment.log and WUAHandler.log on the server        |
| still not done after 120 minutes           | Software Center on the server; C:\Windows\CCM\Logs\UpdatesHandler.log             |
| did not accept a logon after the restart   | look at the server's console                                                      |
| ... ran before patching and do not run now | the service did not come back: <code>Get-Service &lt;name&gt;</code> , System log |

**All settings** of this job, with their defaults: win\_patch.

## ACT on the servers (C:\ProgramData\act)

playbooks/win\_act\_install.yml puts ACT for Windows (vendor/act-windows/act.ps1) on the servers in C:\ProgramData\act (setting win\_act\_dir). Run it again after you update this repository, and every server gets the new version.

- **Locked down.** Ordinary users may create folders and files in C:\ProgramData (and in C:\temp). Someone could then place, or swap, act.ps1 before an administrator runs it, and their code would run with administrator rights. So the job:
  - takes ownership of C:\ProgramData\act;
  - lets only Administrators and SYSTEM in (nothing inherited from C:\ProgramData);
  - checks the file's SHA256 against the one in the repository, every run, and replaces it if it differs.
- **Unblocked.** The file is written on the server itself, so Windows never marks it as "downloaded from the internet", and the job runs Unblock-File on it anyway. Every run also checks it is not blocked, and fixes it if it is.
- **Using it on a server:** an administrator opens PowerShell and runs powershell -ExecutionPolicy Bypass -File C:\ProgramData\act\act.ps1. The first time, :setup asks for their own GenAI or Ask Sage key and keeps it in their profile. No key is put on the server for everyone.
- **Dry run:** run as Check and it says what it would install or fix. To remove it: win\_act\_state: absent.
- **Tell your security team before the first run.** The installer is a large PowerShell script that writes a .ps1 file, and some endpoint-protection products (Defender for Endpoint, Trellix) flag that pattern. Let them know it is coming from AAP, and try one server first.
- **Versions and :probe.** ACT for Windows has the same endpoint formats and :probe command as the Linux one (see ADDING\_ACT.md); the copy in vendor/act-windows/ is refreshed when the repository's vendored ACT is. On Windows, an unattended run never approves the danger tier of commands: only a person at the prompt can. See APPROVED\_COMMANDS.md.
- ACT itself needs full PowerShell. Where WDAC or AppLocker enforce Constrained Language Mode, ACT's own notes (sign it with your code-signing certificate) apply.

**All settings**, with their defaults: win\_act.

## Settings: what each check looks at, and how to change it

Every setting is in `roles/win_check_<name>/defaults/main.yml`, with a comment saying what it does. To change one:

- **for all Windows servers**, use the `Windows servers` inventory's **Variables** box;
- **for a group**, use the group's **Variables** box, or a settings file named after the group (for example `playbooks/group_vars/win_web.yml` for a group called `win_web`);
- **for one server**, use the host's **Variables** box.

The settings you will most likely want:

```
network: what every Windows server must resolve and reach
win_check_network_dns_names: [dc01.yoursite.mil]
win_check_network_tcp:
  - {name: LDAPS, host: dc01.yoursite.mil, port: 636}
  - {name: SIEM, host: siem.yoursite.mil, port: 9997}

# security: your endpoint protection's services (for example Trellix / McAfee)
win_check_security_av_services: [masvc, mfemms]

# audit: your log forwarding agent
win_check_audit_forwarder_services: [SplunkForwarder]

# accounts: who may be a local administrator (wildcards allowed)
win_check_accounts_admins_allowed: ['*\Administrator', 'YOURDOMAIN\Domain Admins',
  'YOURDOMAIN\svc_aap']

# services: more services that must run, on a group (for example IIS servers)
win_check_services_required_extra: [W3SVC, WAS]
```

A few defaults worth knowing:

- **disk**: warning at 85% used, critical at 95%. Per volume: `win_check_disk_overrides: {'D:': {warn: 92, crit: 98}}`.
- **services**: `EventLog`, `RpcSs`, `Dnscache`, `W32Time`, `WinRM` and `mpssvc` (the firewall) must run. Automatic services that stop cleanly by design are not reported.
- **security**: Defender is checked only when it is the active antivirus (`win_check_security_av: auto`). With another product, list its services instead.
- **certs**: warning 30 days before expiry, critical 7 days before. Old expired certificates that nothing uses are listed in the report, not raised as findings.
- **patching**: warning when the last update is 35 days old, critical at 60 days.

**All settings**, with their defaults: `win_check_accounts`, `win_check_audit`, `win_check_certs`, `win_check_disk`, `win_check_eventlog`, `win_check_network`, `win_check_patching`, `win_check_performance`, `win_check_security`, `win_check_services`, `win_check_time`.

## If something goes wrong

| What you see                                                                           | What to do                                                                                                                                                                                                                                                    |
|----------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>winrm or requests is not installed / connection test: no winrm</code>            | the EE has no pywinrm: step 2                                                                                                                                                                                                                                 |
| <code>CERTIFICATE_VERIFY_FAILED</code>                                                 | the CA bundle (step 4);<br><code>ansible_winrm_server_cert_validation: ignore</code> only as a short lab test                                                                                                                                                 |
| <code>the specified credentials were rejected by the server</code>                     | username format ( <code>DOMAIN\user</code> for NTLM), password, or the account is locked                                                                                                                                                                      |
| <code>Connection refused / timed out on 5986</code>                                    | the firewall, or no HTTPS listener (step 1)                                                                                                                                                                                                                   |
| <code>kerberos: ... kinit or Server not found in Kerberos database</code>              | no Kerberos in the EE, or a short host name: use NTLM, or full names                                                                                                                                                                                          |
| <code>local administrator: NO</code>                                                   | add the account to the local Administrators group                                                                                                                                                                                                             |
| <code>the &lt;name&gt; check could not run: ...</code>                                 | the text after it is PowerShell's error. The other checks still ran                                                                                                                                                                                           |
| <code>skipping: no hosts matched</code>                                                | the template's inventory is not <code>Windows servers</code> , or the Limit matches no server there                                                                                                                                                           |
| <code>... is not digitally signed or running scripts is disabled on this system</code> | a GPO sets the PowerShell execution policy for the machine (for example <code>AllSigned</code> ), which overrides what Ansible asks for. Set it to <code>RemoteSigned</code> by GPO for these servers, or sign the scripts with your code-signing certificate |

## Not in this release (next)

- ACT in the Windows jobs: `use_act` on the Windows health check and troubleshoot (explain, propose fixes for approval), like on Linux. ACT is already on the servers (above).
- Service watch for Windows services and IIS application pools.
- Checks for domain controllers, IIS and SQL Server.

## 10. VMware jobs (vCenter): restart, snapshots, VLAN, reports, alarms and ACT

Small jobs for virtual machines, done through vCenter (they replace the vCenter runbooks of System Center Orchestrator):

| Job template            | Playbook                                         | What it does                                                                                                                                                                                |
|-------------------------|--------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| VM - restart            | <code>playbooks/vm_reboot.yml</code>             | restarts VMs. <b>guest</b> (default): the operating system restarts itself through VMware Tools. <b>hard</b> : a reset, like the reset button                                               |
| VM - shut down          | <code>playbooks/vm_shutdown.yml</code>           | shuts VMs down and waits until they are off. <b>guest</b> (default) through VMware Tools; <b>hard</b> : power off at once                                                                   |
| VM - snapshot           | <code>playbooks/vm_snapshot.yml</code>           | takes a snapshot <b>without memory</b> (the disks only). The AAP job number and user go into its description                                                                                |
| VM - delete snapshot    | <code>playbooks/vm_snapshot_delete.yml</code>    | deletes one snapshot by its exact name, or every snapshot of the VM. Lists what the VM has first                                                                                            |
| VM - notes              | <code>playbooks/vm_notes.yml</code>              | edits the VM's Notes in vCenter: append a line (with date and user), replace, or clear. Prints the old notes                                                                                |
| VM - change VLAN        | <code>playbooks/vm_vlan.yml</code>               | moves one network adapter to another port group (VLAN). The VM's other adapters are not touched                                                                                             |
| VM - secure boot report | <code>playbooks/vm_secure_boot_report.yml</code> | read-only: every VM's firmware and secure boot. Lists BIOS VMs and EFI VMs with secure boot off separately                                                                                  |
| VM - datastore report   | <code>playbooks/vm_datastore_report.yml</code>   | read-only: every datastore's capacity, free space, used and provisioned %; the ones at 85 / 90 % used listed apart                                                                          |
| VM - snapshot report    | <code>playbooks/vm_snapshot_report.yml</code>    | read-only: every snapshot of every VM - name, when taken and by whom, age, size, folder - oldest first; the ones 3 days or older listed apart                                               |
| VM - snapshot cleanup   | <code>playbooks/vm_snapshot_cleanup.yml</code>   | deletes the snapshots 3 days or older - never "keep" ones, excluded VMs, the AAP server or ones over 1 TB - then reports what it deleted. Best run in the workflow below, after an approval |

| Job template             | Playbook                                         | What it does                                                                                                                                                                                                                                               |
|--------------------------|--------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| VM - alarms report       | <code>playbooks/vm_alarm_report.yml</code>       | read-only: the triggered alarms on vCenter, datacenters, clusters and hosts; hosts not connected; the last 24 hours' failed logins, VM shutdowns / HA restarts and other errors and warnings; configuration issues - coloured by severity                  |
| VM - alarms ACT analysis | <code>playbooks/vm_alarm_act_analysis.yml</code> | read-only: the same problems, with ACT's likely cause, fix and confidence for each. ACT runs on the AAP side, never on the hosts                                                                                                                           |
| VM - capacity planning   | <code>playbooks/vm_capacity_report.yml</code>    | read-only: overall and per cluster - CPU, memory (also with a host down), vCPUs per core, VMs that still fit, hosts recommended, datastores - with growth from vCenter's own history and the days until each limit; optionally GenAI's estimate next to it |
| ESXi - security settings | <code>playbooks/esxi_security.yml</code>         | on every ESXi host: SSH off, the shell timeouts at 600 s, lockdown mode normal. Changes only what differs, so it can run every night and puts back whatever drifted                                                                                        |

They run on the AAP side and talk to vCenter. They never log in to the VMs.

**Never the AAP server.** Restart, shut down and VLAN change refuse the AAP server's VM. A job would cut off the system running it, and AAP must never be restarted from AAP. The VM is recognised by its name, its guest host name or its IP address, matched against the hosts in your inventory groups `aap` and `aap_hosts`. If the template's inventory has no such group and you have not named the AAP VM (below), these three jobs refuse to run at all. Snapshots, notes and the report are allowed on every VM, since they do not stop it.

## Before you start

### 1. The execution environment must have the VMware collection

The jobs use Red Hat's certified collection `vmware.vmware` (2.11 or later) and its Python libraries (pyVmomi and the vSphere Automation SDK). Find out whether your execution environment (EE) has them. On the AAP server, as the user AAP runs as:

```
podman images | grep -i ee-
podman run --rm <the EE image your templates use> ansible-galaxy collection list vmware.vmware
podman run --rm <the EE image your templates use> python3 -c "import pyVmomi, vmware.vapi;
    print('ok')"
```

- **Both answer:** nothing to do.
- **Not there:** build an EE that has it, with `ansible-builder`. Start from your AAP version's `ee-minimal` image, and pull the collection from your Private Automation Hub when the network is air-gapped. `ansible-builder` reads the collection's Python requirements by itself:

```

execution-environment.yml

version: 3
images:
  base_image:
    name: registry.redhat.io/ansible-automation-platform-27/ee-minimal-rhel9:latest
dependencies:
  galaxy:
    collections:
      - name: vmware.vmware

```

Then in AAP: **Execution Environments > Add** (the image you pushed). Pick it on the VM templates.

Do **not** add a `collections/requirements.yml` to this repository for it. AAP would then try to download from Galaxy on every project sync, and on an air-gapped network every sync would fail.

vCenter 7 or 8 is needed: the jobs also use vCenter's REST API.

## 2. A vCenter account for AAP

A service account with a vCenter role that has only what the jobs need:

| Job                                  | vCenter privileges                                                                                                                                                                                                                                                           |
|--------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| restart                              | Virtual machine > Interaction > Reset                                                                                                                                                                                                                                        |
| shut down                            | Virtual machine > Interaction > Power off                                                                                                                                                                                                                                    |
| snapshot / delete snapshot           | Virtual machine > Snapshot management > Create snapshot / Remove snapshot                                                                                                                                                                                                    |
| notes                                | Virtual machine > Change configuration > Set annotation                                                                                                                                                                                                                      |
| change VLAN                          | Virtual machine > Change configuration > Modify device settings, and Network > Assign network on the port groups                                                                                                                                                             |
| ESXi security settings               | Host > Configuration > Security profile and firewall (the SSH service), Advanced settings (the timeouts), Change settings (lockdown mode and its exception users). If one is missing, the job's error names it (e.g. <code>lacks the privilege Host.Config.Settings</code> ) |
| all (to read), and the alarms report | Read-only, which every vCenter role includes                                                                                                                                                                                                                                 |

Give it the role on the folders or clusters of the VMs these jobs may touch, and nowhere else.

## 3. The credential

AAP: **Credentials > Add**, type **VMware vCenter**: vCenter Host (its name, e.g. `vcsa01.yoursite.mil`), Username, Password. It hands the jobs `VMWARE_HOST`, `VMWARE_USER` and `VMWARE_PASSWORD`. Nothing about vCenter goes in Git.

**vCenter's certificate**: the jobs check it. If the EE does not trust your CA, a job fails with `certificate verify failed`. The right fix is to add your CA to the EE. Until then, `vmware_validate_certs: false` turns the check off, in the template's Variables (not recommended).

**Several vCenters**: the reports (alarms, datastores, snapshots and their cleanup, capacity planning) and the ESXi security settings read **every vCenter in one job, into one report**: keep one VMware vCenter credential and list the vCenters in `vmware_vcenters` (see "Several vCenters, one report" below). The jobs that act on VMs by name (restart, shut down, snapshot, notes, VLAN, secure boot) still use the credential's

vCenter only: for a VM in another vCenter, use a second credential with that vCenter's address, or tick **Prompt on launch** for Credentials and pick it at launch.

## 5. Several vCenters, one report

With two or more vCenters (for example in Enhanced Linked Mode, as one tree in the vSphere Client), list them all in `playbooks/group_vars/all.yml` (for every job) or in a template's Variables:

YAML

```
vmware_vcenters: [vc01.yoursite.mil, vc02.yoursite.mil]
```

- **One credential** is enough: the jobs log in to every vCenter with its user name and password. In Enhanced Linked Mode the vCenters share single sign-on, so the same account works on all of them. Give its role as a **Global Permission** (vSphere Client: **Administration > Access Control > Global Permissions**), or on each vCenter's root with **Propagate to children**.
- The credential's own vCenter is always read. Listing it again (or by another name, its IP) is fine: each vCenter is read once.
- A vCenter on another port: `vc03.yoursite.mil:8443`.
- **The report:** one email for all. The **Datacenter** column says which vCenter (`vc02 / DC1`); the footer lists the vCenters read. A cluster or datastore name that exists in two places is shown with its vCenter and datacenter, e.g. `PROD (vc02 / DC1)`, and never added together.
- **A vCenter that cannot be read** (down, a firewall, a refused login) does not stop the others: the report covers the others, lists it in red on top ("vCenters not read") and its title ends with "- 1 vCenter(s) not read". The job stays **green**: the report completed (`vmware_vcenter_fail_unread: true` makes it fail, e.g. for a workflow's failure path).
- **Snapshot cleanup:** each snapshot is deleted on its own vCenter. A VM's ID repeats in every vCenter (both have a `vm-42`), so the report hands the cleanup the vCenter with each snapshot, and the cleanup deletes only that one.
- The ESXi security settings set the hosts of every vCenter listed.

## 4. Name the AAP VM (and vCenter's)

In your `playbooks/group_vars/all.yml`:

YAML

```
vmware_protected_vms: [AAP01, VCSA01]      # the exact VM names in vCenter
```

This covers the case where the AAP VM's name differs from its host name and VMware Tools is not running, so neither the guest host name nor the IP is known. More groups to refuse: `vmware_protected_groups: [aap, aap_hosts, my_group]` (aap and aap\_hosts are always refused anyway).

## Which VMs a job acts on

You type them. Every job (except the report) asks for **VM names**, one per line: one VM, or several. These are the names as vCenter shows them in its inventory. They are exact, upper/lower case included, with no wildcards.

The job runs on the AAP server and asks vCenter for each name. It does **not** use the template's inventory hosts or a Limit to choose VMs; it uses the inventory only to recognise the AAP server. So a VM does not have to be in any AAP inventory to be restarted or snapshotted.

- Names that do not exist, or exist twice in vCenter, stop the job before anything is done.
- With several VMs, every check is done for all of them first. If one is refused (the AAP server, no VMware Tools for a clean restart), nothing is done to any of them.
- Several VMs are handled in the order given. Restarts are requested one right after the other, without waiting for each to come back up, so restart the members of a cluster (two database nodes, for example) in separate runs. A shutdown waits for each VM to be off before the next.

## The templates

Make one job template per job:

- Inventory:** your Linux inventory, the one with the group **aap**. The jobs need it for the AAP check; they run nothing on its hosts.
- Credentials:** the VMware vCenter credential.
- Playbook:** one of the files in the table at the top.
- Execution environment:** the one with **vmware.vmware**.
- Limit:** leave it **empty** and do **not** prompt for it. These jobs run on the AAP side (**localhost**); a Limit, or a workflow's Limit, would make them skip without doing anything.
- Job type:** tick **Prompt on launch** if you want dry runs. As **Check**, every job says what it would do and changes nothing.

Surveys (**Survey > Add**, then switch the survey **on**):

| Template           | Question (variable)                                    | Type               | Default / choices                                                     |
|--------------------|--------------------------------------------------------|--------------------|-----------------------------------------------------------------------|
| all but the report | VM names, one per line ( <b>vm_names</b> )             | Textarea, required |                                                                       |
| restart, shut down | How ( <b>vm_power_mode</b> )                           | Multiple choice    | <b>guest</b> , <b>hard</b> ; default <b>guest</b>                     |
| snapshot           | Snapshot name ( <b>vm_snapshot_name</b> )              | Text, optional     | empty = <b>aap-&lt;date&gt;-&lt;time&gt;</b>                          |
| snapshot           | Why / change number ( <b>vm_snapshot_description</b> ) | Text               |                                                                       |
| delete snapshot    | Snapshot name ( <b>vm_snapshot_name</b> )              | Text               |                                                                       |
| delete snapshot    | Delete ALL snapshots ( <b>vm_snapshot_delete_all</b> ) | Multiple choice    | <b>false</b> , <b>true</b> ; default <b>false</b>                     |
| notes              | Text ( <b>vm_notes</b> )                               | Textarea           |                                                                       |
| notes              | How ( <b>vm_notes_mode</b> )                           | Multiple choice    | <b>append</b> , <b>replace</b> , <b>clear</b> ; default <b>append</b> |
| change VLAN        | Network adapter number ( <b>vm_nic</b> )               | Integer            | <b>1</b>                                                              |
| change VLAN        | Port group ( <b>vm_portgroup</b> )                     | Text, required     | its name as vCenter shows it                                          |
| any (optional)     | Email the result to ( <b>report_email_to</b> )         | Text, not required | empty = no email ("Email the report" below)                           |

The secure boot and datastore reports need no survey. Schedule them weekly with **report\_email\_to** in their Variables. They stay green when they find VMs or datastores to fix - the report and email say which;

`vm_secure_boot_fail: true` / `vm_datastore_fail: true` mark the job failed instead, for a workflow that opens a ticket. `vm_secure_boot_folder` limits the secure boot report to one vCenter folder; `vmware_datacenter` limits either report to one datacenter; `vm_datastore_warn_pct` / `vm_datastore_crit_pct` (85 / 90) set the datastore thresholds.

## Snapshots: the report, the cleanup, and the approval workflow

### The idea, in one picture

|                     |                                                                                                     |
|---------------------|-----------------------------------------------------------------------------------------------------|
| 1. SNAPSHOT REPORT  | looks at every VM and makes a list of its snapshots.<br>It changes nothing. It emails you the list. |
|                     |                                                                                                     |
| v                   |                                                                                                     |
| 2. APPROVAL         | AAP stops and waits. A person reads the list and<br>clicks Approve (go ahead) or Deny (stop).       |
|                     |                                                                                                     |
| v (only on Approve) |                                                                                                     |
| 3. SNAPSHOT CLEANUP | deletes the old snapshots that were ON THAT LIST -<br>nothing else - and emails what it deleted.    |

Think of it like cleaning out a fridge: first someone writes down everything old in it (the report), then a grown-up says "yes, throw those out" (the approval), then the helper throws out **only what is on the paper** (the cleanup). Anything someone put a "keep" label on stays.

### What the report shows

For every snapshot of every VM: the VM, the snapshot's name, when it was taken, **who took it**, how many days old it is, how big it is (MB, GB or TB), the VM's folder and the description. Oldest first. **Every snapshot is in that list, whatever its age** - even one taken an hour ago. Then two short lists on top:

- **To delete:** snapshots older than 3 days (`vm_snapshot_max_age_days`).
- **Held back:** old snapshots the cleanup will never delete, each with the reason.

The email's heading and subject always say the limit, e.g. `VMware snapshot report (older than 2 days): 4 to delete (310 GB); 23 snapshot(s) on 17 VM(s)`.

### To use 1 or 2 days instead of 3:

| Where                          | Change                                                                                                                                                                                                                        |
|--------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| In the workflow                | the survey answer when you launch it - or for good: the workflow > <b>Survey</b> tab > the question > <b>Default 2</b> . The report and the cleanup both use it: after the approval, snapshots older than 2 days are deleted. |
| The report template on its own | Templates > <code>VM - snapshot report</code> > Edit > <b>Variables:</b><br><code>vm_snapshot_max_age_days: 2</code>                                                                                                          |

In the workflow, the survey answer wins over the template's Variables.

"Who took it" comes from vCenter's history. vCenter forgets old history (often after 30 days), so for an old snapshot it may say **unknown**. Snapshots taken by the "VM - snapshot" job also have the AAP user in their description.

## What the cleanup never deletes

| Never deleted                                                                 | Why                                                                                    | Change it with (in playbooks/group_vars/all.yml)                          |
|-------------------------------------------------------------------------------|----------------------------------------------------------------------------------------|---------------------------------------------------------------------------|
| A snapshot whose name or description says <b>keep</b> or <b>do not delete</b> | someone wants it kept                                                                  | <code>vm_snapshot_keep_regex</code>                                       |
| Any snapshot of the VMs on your do-not-touch list                             | domain controllers, vCenter, ...                                                       | <code>vm_snapshot_cleanup_exclude_vms: [DC*, VCSA*] (* = anything)</code> |
| Any snapshot of the VMs in <code>vmware_protected_vms</code>                  | the same protected VMs as for restart                                                  | <code>vmware_protected_vms</code>                                         |
| Any snapshot of the AAP server                                                | AAP must never be put at risk by AAP                                                   | (always)                                                                  |
| A snapshot bigger than <b>1 TB</b> , or whose size is unknown                 | merging a huge snapshot can take hours and slow the VM: do it by hand, at a quiet time | <code>vm_snapshot_cleanup_max_size_gb: 1024</code>                        |
| More than 50 snapshots in one run                                             | a typo in the days would otherwise delete everything                                   | <code>vm_snapshot_cleanup_max: 50</code>                                  |
| Anything that was not on the approved list                                    | the approver only said yes to that list                                                | (always, in the workflow)                                                 |

## Set it up (once)

**Step 0 - what you need first** (see "Before you start" above): the VMware vCenter credential, the execution environment with `vmware.vmware`, and the mail relay in `all.yml` ("Email the report").

**Step 1 - your do-not-touch list.** In VS Code, in `playbooks/group_vars/all.yml`, add (with your VM names; \* matches anything) - then Commit, Sync Changes, and sync the project in AAP:

```
YAML
vm_snapshot_cleanup_exclude_vms: [DC*, VCSA*, AAP*]
```

**Step 2 - the report template.** Templates > Create template > Create job template:

- Name: `VM - snapshot report`
- Inventory: your inventory with the `aap` group
- Project: site-automation. Playbook: `playbooks/vm_snapshot_report.yml`
- Execution environment: the one with `vmware.vmware`
- Credentials: your **VMware vCenter** credential
- Leave Limit empty. No survey. Save.

**Step 3 - the cleanup template.** The same again:

- Name: `VM - snapshot cleanup`
- Playbook: `playbooks/vm_snapshot_cleanup.yml`
- Tick **Prompt on launch** next to **Job type** (this lets you do a dry run - below).
- No survey. Save.

**Step 4 - the workflow.** Templates > Create template > **Create workflow job template:**

- Name: **VM snapshot cleanup (with approval)**, your organization. Save.
- **Survey** tab > Create survey question - two questions, then switch the survey on:

| Question                                   | Answer variable name                  | Answer type                  | Default             |
|--------------------------------------------|---------------------------------------|------------------------------|---------------------|
| Delete snapshots older than how many days? | <code>vm_snapshot_max_age_days</code> | Integer, required, minimum 1 | <b>3</b>            |
| Email the reports to                       | <code>report_email_to</code>          | Text                         | your team's address |

The workflow hands these answers to both jobs, so the report and the cleanup use the same days.

- **Visualizer** (the boxes):
1. Click **Add step**. Node type: **Job template**. Pick **VM - snapshot report**. Save.
  2. Hover over that box > **Add step and link** (the plus sign). Node type: **Approval**. Name: **Delete the snapshots listed in the report?**. Timeout: e.g. **1 day** (no answer by then = no). Run type: **Run on success**. Save.
  3. Hover over the approval box > **Add step and link**. Node type: **Job template**. Pick **VM - snapshot cleanup**. Run type: **Run on success** (after an approval, that means: approved). Save.
  4. Click **Save** at the top of the Visualizer.

**Step 5 - who may approve.** On the workflow: **User Access** (or Team Access) > Add > pick the people > role **Approve**. They need no other rights. Optional: **Notifications** tab > turn on the approval notification so they get an email when something waits for them.

**Step 6 - when it runs.** On the workflow: **Schedules** > Create schedule (e.g. every Monday 7:00). Or launch it by hand.

**Run it**

1. Launch the workflow (or the schedule does). Answer the survey (or keep 3 days).
2. The report runs and **emails the list**: what will be deleted, what is held back, and why.
3. AAP waits at the approval. Open **Automation Execution > Administration > Workflow Approvals** (or the workflow job's view) > the approval > **Approve** or **Deny**.
4. On **Approve**: the cleanup deletes the snapshots from that list that are still there, and still that old, and **emails what it deleted** (and how much space came back). On **Deny** or timeout: nothing is deleted.

The cleanup's email can be sent only when something was actually deleted:

`vm_snapshot_cleanup_email_only_if_deleted`: **true** (in the cleanup template's Variables).

**Dry run: try it, delete nothing**

- **The safest first try: say no.** Launch the workflow, read the report email, then **Deny** the approval. The report only reads, so nothing changed anywhere.
- **See exactly what the cleanup would delete, without deleting:** in the workflow's Visualizer, click the **VM - snapshot cleanup** box > **Edit** (pencil) > Job type **Check** > Save, and save the Visualizer. Run the workflow and **Approve**: the cleanup job's output lists **WOULD DELETE ...** for each snapshot and deletes nothing (and sends no email). Set the box back to **Run** afterwards. (This needs the Step 3 tick "Prompt on launch" for Job type.)

- **Without the workflow:** launch `VM - snapshot cleanup` on its own. With no approved list and no `vm_snapshot_cleanup_confirm: true` it only lists what it would delete - it never deletes.

### Without the workflow

To clean up by hand, make a separate template (not the one in the workflow) on `playbooks/vm_snapshot_cleanup.yml` with a survey: `vm_snapshot_max_age_days` (Integer, 3) and `vm_snapshot_cleanup_confirm` ("Delete them?", Multiple choice `false / true`, default `false`). With `false` it only lists; with `true` it deletes - with every rule above still in place.

## Alarms: the report, and ACT's analysis

Two read-only jobs. The **report** shows what is wrong in vCenter now and what happened in the last 24 hours. The **ACT analysis** has ACT (the AI assistant) work out the likely cause and the fix of each of those problems, and say how sure it is. Neither changes anything.

### What the report shows

| Section                   | What is in it                                                                                                                                  | Colour                                                                           |
|---------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------|
| Triggered alarms          | the alarms vCenter shows now on vCenter itself, the datacenters, clusters and hosts: severity, object, alarm, since when, acknowledged by whom | red = critical, amber = warning                                                  |
| Host connection           | hosts not connected to vCenter now, and hosts that lost the connection in the last 24 h and came back                                          | red / amber                                                                      |
| Failed logins             | wrong user name or password: one row per user, source address and server, with the count and the first and last time                           | amber; red from 10 failures                                                      |
| Virtual machine events    | guest shutdowns and restarts, power-offs, resets, HA restarts and failures, and who did it                                                     | blue = someone did it, amber = HA or nobody, red = a failover or power-on failed |
| Other errors and warnings | every other error and warning event; the same event on the same object counted once                                                            | red = error, amber = warning                                                     |
| Configuration issues      | what vCenter shows as configuration issues (SSH left on, HA problems, ...)                                                                     | amber                                                                            |
| All hosts                 | every host: connection, maintenance mode, ESXi version and build, hardware, its alarms and its errors / warnings                               | red / amber / blue row                                                           |

On top: the counts in coloured boxes. The subject says the most important ones, e.g. `[AAP] VMware alarms report: 1 critical, 2 warning alarm(s), 47 failed login(s) (events: last 24 h)`.

### Set up the report

#### 1. Templates > Create template > Create job template:

- Name: `VM - alarms report`
- Inventory: any (e.g. your VMware one: the job runs on the AAP side)

- Project: site-automation. Playbook: `playbooks/vm_alarm_report.yml`
- Execution environment: your VMware one (this job needs only pyVmomi, which it has)
- Credentials: your **VMware vCenter** credential. A read-only vCenter account is enough.
- Variables: `report_email_to: vmteam@yoursite.mil`
- Save.

2. **Schedules > Create schedule:** e.g. every day at 06:00.

| Setting (template Variables)         | Default                                           | What it does                                                                                                |
|--------------------------------------|---------------------------------------------------|-------------------------------------------------------------------------------------------------------------|
| <code>vm_alarm_hours</code>          | 24                                                | the events of the last this-many hours (168 for a weekly report; 0 = alarms and hosts only)                 |
| <code>vm_alarm_types</code>          | <code>[vcenter, datacenter, cluster, host]</code> | whose alarms; add <code>vm</code> , <code>datastore</code> or <code>network</code> for theirs too           |
| <code>vm_alarm_login_critical</code> | 10                                                | this many failed logins of one user from one place, or more, is red                                         |
| <code>vm_alarm_fail</code>           | false                                             | <code>true</code> = the job shows <b>failed</b> on a critical alarm or a host not connected, for a workflow |
| <code>vmware_datacenter</code>       | (all)                                             | one datacenter only                                                                                         |

## The ACT analysis

**What you get:** an email "VMware alarms ACT analysis". First "What to do first", then one row per problem (P1, P2, ...): severity, object, the problem, the **likely cause**, the evidence, the **fix** and a **confidence**. The confidence is ACT's own estimate, from 0 to 100: green 80 and up (the evidence shows it), amber 50-79 (likely), red under 50 (a guess: check first). These are likely causes, not certain ones: check before you change anything. ACT runs no command and changes nothing.

**Where ACT runs, and what about the hosts?** ACT runs inside this job's execution environment, on the AAP server (or the execution node that runs the job). Not on the ESXi hosts, not on the VMs. So:

- The **ESXi hosts need nothing**: no Python, no network path to ACT.
- **Python** is in every execution environment (Ansible itself runs on it; ACT needs 3.8 or later).
- The **AAP node must reach the model's address over HTTPS** (port 443): `https://api.genai.mil` for GenAI.mil, or your `site_act_url`. The job checks this first and says plainly when it cannot:
- *"cannot reach the model ... Open HTTPS ... or set a proxy"*: ask for a firewall rule from the AAP node to that address, or set your proxy in `playbooks/group_vars/all.yml`: `site_act_env: {HTTPS_PROXY: "http://proxy.yoursite.mil:8080", NO_PROXY: "localhost,.yoursite.mil"}`
- *"Its certificate is not trusted"*: the execution environment does not trust the model's certificate authority. Put that CA certificate (public, `.pem`) at `playbooks/files/ca/model-ca.pem`, add `site_act_ca: "{{ playbook_dir }}/files/ca/model-ca.pem"` to `all.yml`, and add the line `playbooks/files/ca/` to `.site-local` (so an update keeps it).

**Privacy:** before anything leaves, ACT replaces host, cluster, datacenter and VM names, IP addresses and user names with placeholders, and puts the real names back into its answer.

**Set it up:**

1. **The key.** If your other ACT jobs have an **ACT model key** credential, use it. If not: create the credential type from `aap/credential_types/act_model_key.yml`, then a credential of that type with your GenAI.mil (or Ask Sage) key (SECRETS.md).
2. **Templates > Create template > Create job template:**
  - Name: `VM - alarms ACT analysis`
  - Inventory, project, execution environment: as for the report. Playbook: `playbooks/vm_alarm_act_analysis.yml`
  - Credentials: **VMware vCenter** and **ACT model key** (both)
  - Tick **Prompt on launch** next to **Job type** (for the dry run)
  - Variables: `report_email_to: vmteam@yoursite.mil`
  - Save. The provider and model are your usual ACT settings (`site_act_provider`, `site_act_models` in `all.yml`); the default is GenAI.mil.
3. **First run, as a dry run:** Launch, Job type **Check**. It reads vCenter, lists the problems it would give ACT, and calls nothing. Then launch it normally.

**Both in one go:** a workflow with `VM - alarms report`, then `VM - alarms ACT analysis` linked with **Always**. Each sends its own email. Or schedule just the ACT analysis: its email lists every problem too.

**Green or red:** nothing wrong = ACT is not called and no email is sent (`vm_alarm_act_email_if_none: true` sends one anyway). The job is **red** when ACT could not run (no key, no network, an ACT error): the email still lists the problems, and says why.

| Setting                                 | Default                                   | What it does                                                                                               |
|-----------------------------------------|-------------------------------------------|------------------------------------------------------------------------------------------------------------|
| <code>vm_alarm_act_max_items</code>     | <code>40</code>                           | give ACT at most this many problems, the most severe first (the rest are listed as not analyzed)           |
| <code>vm_alarm_act_config_issues</code> | <code>true</code>                         | configuration issues too                                                                                   |
| <code>vm_alarm_act_check_url</code>     | <code>true</code>                         | check that the AAP node can reach the model first                                                          |
| <code>vm_alarm_act_model_timeout</code> | <code>site_act_model_timeout</code> (600) | seconds the model may take for its answer (a <code>GENAI_TIMEOUT</code> in <code>site_act_env</code> wins) |
| <code>site_act_timeout</code>           | <code>1200</code>                         | seconds the whole ACT run may take (keep it above the model's)                                             |

## Capacity planning: overall and per cluster

`playbooks/vm_capacity_report.yml` answers "how much room is left, and for how long?" - for all clusters together and for each cluster. It only reads. The history comes from **vCenter's own daily statistics** (it keeps a year of them at the default statistics level), so it works from the first run: nothing has to be collected first.

What the report shows

| Section                                                  | What it shows                                                                                                                                                                                                                                                                                                                                                                                                       |
|----------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Overall                                                  | all clusters together: CPU and memory capacity, used, used with a host down in each cluster, growth per month and the days until the target; hosts usable and <b>hosts recommended</b> (now, and within 12 months); vCPUs per core and room for more VMs; shared storage with its growth and runway                                                                                                                 |
| Clusters: CPU and memory                                 | per cluster: hosts usable, CPU and memory used, <b>memory used if the biggest host fails</b> (N+1: the cluster must still run everything), vCPUs per core, VMs and their average size, memory growth, <b>memory runway</b> (days and date until the target with a host down), <b>room for more VMs</b> of the average size (and what limits it), <b>hosts recommended</b> , VMs added / removed in the last 30 days |
| Clusters: storage                                        | per cluster, its shared datastores together: capacity, free, used, provisioned (thin disks at full size), growth, days until 85 % and until full                                                                                                                                                                                                                                                                    |
| Datastores: runway                                       | each shared datastore, the one that fills first on top: growth, days (and date) until 85 %, 90 % and full, and how well a straight line fits its history                                                                                                                                                                                                                                                            |
| Math vs GenAI (with <code>vm_capacity_act: true</code> ) | for every item, the job's runway (the math) next to ACT's own estimate (GenAI), whether they agree, ACT's recommendation and reasoning, and its confidence                                                                                                                                                                                                                                                          |
| Hosts per cluster                                        | every host: state, cores, CPU and memory used, VMs - an unbalanced cluster shows here                                                                                                                                                                                                                                                                                                                               |

**Hosts recommended** is a count of hosts of the cluster's average size (cores, memory) that keeps memory and CPU at their targets (80 %) with one host down - for example **1 now (memory 59% on average then); 2 within 12 months**. "Within 12 months" applies the current growth. It is a sizing guide in hosts, not a list of models.

**How the runway is worked out:** a straight line through the last 30 days of vCenter's daily values (memory consumed, CPU used, datastore space used), continued until it reaches the limit. "Not growing" means the line is flat or falling. The **trend fit** (R<sup>2</sup>) says how straight the history is: poor means jumps (VMs moved, snapshots deleted, a datastore grown), and the runway is a rough guide then.

Math vs GenAI

With `vm_capacity_act: true`, ACT (the AI assistant, GenAI.mil or Ask Sage) reads the same numbers and makes **its own estimate** for the overall picture, each cluster and each datastore: when it needs action (in days), what to do, why, and how sure it is. The report puts the two side by side:

| Math (straight line) | GenAI (ACT) | Agreement            |
|----------------------|-------------|----------------------|
| 102 days (90 % full) | 55 days     | GenAI 47 days sooner |
| 0 days (memory)      | 0 days      | agree                |

They **agree** when they are within 30 days or 25 % of each other. Where they differ, read ACT's reasoning (a recent jump, a limit that comes first, a poor fit), and plan for the sooner one. ACT states needs in hosts, cores, GB and TB, and timing in days or months. It runs on the AAP side and runs no command (as in "The

ACT analysis" above: the same key, URL, proxy and CA settings).

## Set it up

### 1. **Templates > Create template > Create job template:**

- Name: `VM - capacity planning`
- Inventory: any (e.g. your VMware one). Project: site-automation. Playbook: `playbooks/vm_capacity_report.yml`
- Execution environment: your VMware one. Credentials: **VMware vCenter** (read-only is enough)
- Tick **Prompt on launch** next to **Variables** (to try one cluster)
- Variables: `report_email_to: vmteam@yoursite.mil`
- Save. Launch it once and read the report.

### 2. **Per cluster:** the same report for one or a few clusters - Variables `vm_capacity_clusters: [PROD-CL01]` (\* works: `[PROD-*]`). A second template with that variable gives each cluster its own report.

### 3. **GenAI's estimate (optional):** add the **ACT model key** credential to the template, and `vm_capacity_act: true` to its Variables. Launch.

### 4. **Schedule: Schedules > Create schedule**, e.g. every Monday at 07:00 (your time zone).

| Setting                                                        | Default               | What it does                                                                     |
|----------------------------------------------------------------|-----------------------|----------------------------------------------------------------------------------|
| <code>vm_capacity_clusters</code>                              | <code>[]</code> (all) | only these clusters                                                              |
| <code>vm_capacity_failover_hosts</code>                        | <code>1</code>        | hosts down to plan for (N+1); a larger HA admission control failover level wins  |
| <code>vm_capacity_mem_target_pct / _cpu_target_pct</code>      | <code>80 / 80</code>  | memory / CPU used at most this, with those hosts down                            |
| <code>vm_capacity_vcpu_per_core_max</code>                     | <code>4</code>        | vCPUs per physical core (amber from here)                                        |
| <code>vm_capacity_storage_warn_pct / _crit_pct</code>          | <code>85 / 90</code>  | datastore used % (the storage runway is to 90 %)                                 |
| <code>vm_capacity_runway_warn_days / _crit_days</code>         | <code>90 / 30</code>  | a runway shorter than this is amber / red                                        |
| <code>vm_capacity_horizon_months</code>                        | <code>12</code>       | "hosts recommended within N months"                                              |
| <code>vm_capacity_trend_days / vm_capacity_history_days</code> | <code>30 / 90</code>  | the days the trend is fitted on / read                                           |
| <code>vm_capacity_include_local_datastores</code>              | <code>false</code>    | count a host's own datastores too                                                |
| <code>vm_capacity_act</code>                                   | <code>false</code>    | GenAI's estimate next to the math                                                |
| <code>vm_capacity_fail</code>                                  | <code>false</code>    | <code>true</code> = the job shows failed when a cluster or datastore is critical |

## Deploy a VM from a ServiceNow ticket

`playbooks/vm_deploy_from_ticket.yml` builds one VM from a vCenter template for a ServiceNow ticket (an incident or a task assigned to your team). You launch it with the ticket number; it reads the ticket, checks the request against the rules you set in AAP, writes on the ticket what it builds, builds it, and writes the result - with the VM's IP address - on the ticket.

### How it flows, step by step

The operator launches **one workflow**, `VM - deploy from ServiceNow ticket (workflow)`, with the ticket number. It runs the deploy job twice - as its *propose* step and its *build* step - with an approval between them:

1. Requester — writes the ticket in ServiceNow (plain words, or key: value lines)

2. Operator — launches the workflow in AAP with the ticket number (ONE launch)

— workflow step 1: PROPOSE (the deploy job, vm\_deploy\_phase: propose) —————

3. AAP → ServiceNow: reads the ticket

4. AAP → GenAI: "what does this ticket ask for?" (plain words only; key: value

5. GenAI → AAP: name, template, CPU, memory, disks, site lines skip 4-6)

- each with the ticket's own words

6. AAP checks every value really is in the ticket

questions? → AAP → ServiceNow: questions to the requester (a comment); the workflow STOPS

7. AAP checks the rules (templates, site, caps, a free name, the group) and finds the cluster and datastore (read-only)

8. AAP → ServiceNow: "[AAP] GenAI proposal" + "[AAP] Checked - ready to build: ... it would go to ..."

AAP → email: the same, to the approver. NOTHING IS BUILT

— workflow step 2: APPROVAL —————

9. Approver — reads the email (or the ticket), clicks Approve in AAP (Deny = nothing built)

— workflow step 3: BUILD (the deploy job, vm\_deploy\_phase: build) —————

10. AAP → ServiceNow: reads the ticket and the posted proposal (GenAI is NOT asked again; a ticket

changed since the proposal = nothing built)

11. AAP checks the rules again; AAP → ServiceNow: "[AAP] Deploying this VM" (refused = nothing built)

12. AAP → vCenter: clone the template, CPU / memory / disks, customize (or network boot), power on

13. AAP → ServiceNow: "[AAP] VM ... deployed" - IP (or MAC for MECM), where it went

14. AAP → email: the result

For a demo with no person in between: `vm_deploy_genai_approval: false` on the job template and launch the job itself - one run reads, checks and builds (steps 3-8, then 11-14), when every value is traced to the ticket and within the rules; questions or a refusal still stop it.

### Who does what:

| Who            | Does                                                                                             | Never does                                                                          |
|----------------|--------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|
| ServiceNow     | holds the request; shows the job's notes and questions                                           | -                                                                                   |
| AAP (this job) | reads the ticket, checks everything, writes the notes, <b>builds the VM in vCenter</b> , reports | builds anything the rules do not allow, or without a "Deploying" note on the ticket |

| Who                               | Does                                                                                                                                | Never does                                                                                                                                          |
|-----------------------------------|-------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>GenAI</b> (ACT, evidence mode) | <b>reads the ticket text and turns it into values</b> (name, template, CPU, memory, disks, site), quoting the words each comes from | builds or changes anything, runs a command, sees a password or a credential, sees vCenter, chooses a cluster or datastore, decides whether to build |
| <b>vCenter</b>                    | does what AAP asks: clone, size, customize, power on                                                                                | -                                                                                                                                                   |
| <b>The operator / approver</b>    | launches the workflow; reads the proposal and approves (or denies) in AAP                                                           | -                                                                                                                                                   |

**GenAI only interprets text into values.** It answers once, in run 1, with values and quotes; AAP checks each quote is really in the ticket and each value against the rules, and writes the result on the ticket for a person to see. Every action - the notes on the ticket and everything in vCenter - is done by AAP, with the VMware vCenter and ServiceNow credentials, which GenAI never sees. Run 2 builds from what is written on the ticket, not from a new GenAI answer.

### What the requester writes in the ticket

**key:** **value** lines, in the ticket's **description** or in a **comment / work note**. A later comment overrides an earlier value (to correct a request, add a comment with the new line):

```
vm_name: app01
template: rhe19-gold
cpu: 4
memory_gb: 16
```

| Key                               | Meaning                                                                                                                                     | Not given                                                            |
|-----------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------|
| <b>vm_name</b>                    | the VM's name in vCenter and its host name                                                                                                  | <b>required</b>                                                      |
| <b>template</b>                   | the template, by its name in vCenter                                                                                                        | <b>required</b>                                                      |
| <b>cpu, memory_gb</b>             | vCPUs, memory in GB                                                                                                                         | the template's                                                       |
| <b>disks_gb</b>                   | disk sizes in GB, system disk first:<br><b>disks_gb: 120, 200</b> grows disk 1 to 120 GB (never shrinks it) and adds one 200 GB disk (thin) | the template's disks                                                 |
| <b>site</b>                       | which site: its cluster, datastore, folder (and vCenter) come from<br><b>vm_deploy_sites</b>                                                | the ticket's <b>Location</b> field, when it is one of the sites      |
| <b>cluster, datastore, folder</b> | where it goes ( <b>folder</b> as a path: <b>Linux/App</b> ; <b>datastore</b> may be a datastore cluster) - overrides the site's             | the site's, else <b>vm_deploy_defaults</b> , else the template's own |
| <b>domain</b>                     | the DNS domain set in the guest (Linux)                                                                                                     | <b>vm_deploy_domain</b>                                              |
| <b>customization_spec</b>         | a customization specification from vCenter ( <b>needed for Windows</b> )                                                                    | Linux: none needed                                                   |

| Key                                            | Meaning                                        | Not given        |
|------------------------------------------------|------------------------------------------------|------------------|
| <code>vcenter</code> , <code>datacenter</code> | which one, when the template exists in several | found by the job |
| <code>notes</code>                             | added to the VM's notes                        | -                |

Other lines (a greeting, `Name: Jane Doe`) are ignored. Upper and lower case do not matter (`Memory: 16 GB` works too).

### What the job does, in order

1. Reads the ticket (`vm_deploy_ticket: INC...`, `TASK...`, `SCTASK...`, `RITM...`, `CHG...`).
2. Checks: the ticket is open and assigned to `vm_deploy_assignment_group`; the template is in `vm_deploy_templates`; the cluster, datastore and folder are allowed; CPU and memory are within the caps; the name is valid and **not used in any vCenter** (`vmware_vcenters`). Anything wrong: a work note on the ticket says what, and nothing is built.
3. Writes "[AAP] Deploying this VM" with the exact request on the ticket - and builds nothing if ServiceNow refuses that note.
4. Clones the template (powered off), sets CPU, memory and the VM's notes (the ticket number), applies **guest customization** (Linux: host name = `vm_name`, every network adapter on DHCP; Windows: the named customization specification, its computer name set to `vm_name`), powers it on.
5. Waits for vCenter's customization result (`vm_deploy_wait_customization`, 900 s) and for an IPv4 address from VMware Tools (`vm_deploy_wait_ip`, 600 s).
6. Writes "[AAP] VM app01 deployed" on the ticket: IP address, customization result, where it went, size. On a failure: what failed, and whether a (partly set up) VM exists.

The network is the template's: the template's adapter is on your staging VLAN, and so is the new VM.

### Sites: site A goes to this cluster, site B to that one

```
YAML

vm_deploy_sites:
  SiteA: {cluster: A-PROD-CL01, datastore: A-DSC01, folder: Servers/New}
  SiteB: {cluster: B-PROD-CL01, datastore: B-DSC01, folder: Servers/New, vcenter:
    vc02.yoursite.mil}
```

The ticket says `site: SiteA` (upper / lower case does not matter) - or nothing, when the ticket's **Location** field is one of the site names. The site's values fill in what the ticket does not say (over `vm_deploy_defaults`); a `cluster`: the ticket names itself still wins, within the allow-lists. With sites set up, a ticket that names no site is refused (`vm_deploy_site_required: false` uses `vm_deploy_defaults` instead). Any setting can go in a site: `cluster`, `datastore`, `folder`, `vcenter`, `datacenter`, `domain`, `cpu`, `memory_gb`.

### Automatic placement: the job picks the cluster and datastore

When neither the ticket, its site nor `vm_deploy_defaults` names a cluster (or a datastore), the job chooses one by rules - never from all of vCenter, only among explicit candidates:

- **Candidates:** a site's `clusters` / `datastores` lists, else `vm_deploy_clusters` / `vm_deploy_datastores`. With neither, the VM goes where the template is.
- **Cluster:** the one with the most memory left after this VM, planning one host down (`vm_deploy_failover_hosts`) - it must stay under `vm_deploy_mem_target_pct` (80 %).

- **Datastore:** shared, accessible, not in maintenance, mounted by that cluster, and under `vm_deploy_datastore_max_used_pct` (85 %) with the VM at its **full** size (thin disks counted full).
- The ticket note and the report give the ranking: `cluster chosen: PROD-A2: memory 61.3% used with this VM and 1 host down (target 80%)`.

YAML

```
vm_deploy_sites:
  SiteA: {clusters: [A-PROD-*], datastores: [A-DS-*], folder: Servers/New}
  SiteB: {clusters: [B-PROD-*], datastores: [B-DS-*], folder: Servers/New, vcenter:
    vc02.yoursite.mil}
```

## Requests in plain words: GenAI reads the ticket (proof of concept)

With `vm_deploy_genai: true`, a ticket does not need the `key: value` lines. A requester can write *"We need a Windows Server 2022 box called winapp05 for SiteA, 8 GB RAM, 2 CPUs, a 120 GB system disk and a 200 GB data disk"*, and ACT (GenAI.mil, the same ACT model key credential) reads it.

**Propose, approve, build.** In the workflow (below), the propose step has ACT read the ticket (in evidence mode: it runs no command, sees no credential and no inventory). The job posts what it read on the ticket - `[AAP] GenAI proposal (ticket text sha256:...)`, each value with the words of the ticket it came from - then checks it and finds where it would go, and emails that to the approver. **Nothing is built** until someone approves; the build step then builds **the posted proposal** - GenAI is not asked again - with every check of a normal build. If something is missing or unclear, the questions go to the requester as a **comment** (they see comments, not work notes) and the workflow stops. A value the ticket gives itself (`cpu: 4` in a comment) always wins.

## The guardrails:

| Guardrail                             | What it does                                                                                                                                                                                                                                                      |
|---------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Off by default                        | <code>vm_deploy_genai: false</code> ; the <code>key: value</code> lines work without GenAI                                                                                                                                                                        |
| A person approves                     | the workflow's approval sits between the proposal and the build; the build step builds the proposal posted on the ticket - never a fresh answer, never unseen ( <code>vm_deploy_genai_approval: false</code> drops the person, for a demo)                        |
| The ticket must not change in between | the proposal carries a hash of the ticket text; editing the description or adding a comment voids it: the build step builds nothing, and the workflow is run again from the start                                                                                 |
| Every value traceable to the ticket   | each value comes with the ticket's words it is from; the job checks those words are in the ticket and contain the value (the name word for word, the numbers as written) - anything else is dropped and becomes a question. GenAI's own confidence is not trusted |
| Templates from a catalog              | GenAI may only pick a key of <code>vm_deploy_template_catalog</code> , and it must be in <code>vm_deploy_templates</code>                                                                                                                                         |
| Same rules as any request             | allow-lists, CPU / memory / disk caps, name rules, the group, a free name: what the ticket text asks for ("use 64 CPUs") is refused like a typed request                                                                                                          |
| No placement by GenAI                 | cluster and datastore come from the site, the defaults or the rules above                                                                                                                                                                                         |
| Instructions in the ticket are data   | the prompt tells the model to ignore them; the two checks above make them harmless anyway                                                                                                                                                                         |

| Guardrail                | What it does                                                                                                                 |
|--------------------------|------------------------------------------------------------------------------------------------------------------------------|
| Everything on the record | the proposal, the questions, the "Deploying" note and the result are on the ticket; the job's artifacts keep what GenAI read |

### Set it up:

1. The deploy job template (Set it up, above) gets the **ACT model key** credential too, and **Prompt on launch** ticked next to **Variables** (the workflow hands each step its `vm_deploy_phase`). Its Variables gain:

YAML

```
vm_deploy_genai: true
vm_deploy_template_catalog:
  win2022-pxe: "Windows Server 2022 (installed over the network by MECM)"
  rhel9-gold: "Red Hat Enterprise Linux 9"
```

2. **Templates > Create template > Create workflow job template:** name `VM - deploy from ServiceNow ticket (workflow)`; a **survey** with one question, *Ticket number*, variable `vm_deploy_ticket` (the workflow passes it to every step). Save, then open the **Workflow Visualizer**:
  - **Start > node 1:** job template `VM - deploy from ServiceNow ticket`; its prompt, Variables: `vm_deploy_phase: propose`
  - **node 1 > (On success) > node 2: Approval**, name `Build the VM on the ticket?`, a timeout of, e.g., 8 hours
  - **node 2 > (On success = approved) > node 3:** job template `VM - deploy from ServiceNow ticket`; Variables: `vm_deploy_phase: build`
  - Nothing on *On failure*: a question for the requester, a refusal or a denied approval ends it.
3. Give the approvers the workflow's **Approve** permission (Access > the workflow > Roles), and `report_email_to` their address on the job template: the propose step's email is what they read.
4. Launch the workflow with a test ticket. The proposal and "ready to build" are on the ticket and in the email; approve in **Automation Execution > Workflow Approvals**; the VM is built.

The same workflow works for tickets with `key: value` lines: step 1 then only checks and places.

**Without the workflow** (a demo): launch the job template itself with `vm_deploy_genai_approval: false` - one run reads, checks and builds. With `vm_deploy_genai_approval: true` and no workflow, the job proposes and stops; launching it again builds the posted proposal.

The ticket text goes to the model (GenAI.mil): only use it for tickets whose text may go there.

### Windows from a network-boot (PXE / MECM) template

A template with **no operating system** that boots from the network into an MECM (SCCM) task sequence: list it in `vm_deploy_pxe_templates` (e.g. `[win2022-pxe]`). For those, the job:

- does **no guest customization** (there is no OS yet) and does not wait for an IP address;
- refuses a `vm_name` longer than 15 characters (the Windows / MECM computer name limit);
- powers the VM on - it boots from the network to MECM - and writes on the ticket its **MAC address** and the computer name to use.

On the MECM side, one of the two usual ways:

1. **Unknown computers:** the task sequence is deployed to *All Unknown Computers* and asks for (or sets) the computer name - use the VM's name from the ticket.

2. **Known computers:** an MECM admin imports the computer (**Assets and Compliance > Devices > Import Computer Information**: name = `vm_name`, MAC = the one on the ticket) into the collection the task sequence is deployed to. (This step can be automated later through MECM's AdminService.)

Check that the template boots from the network first (VM Options > Boot Options: EFI, network first or a boot delay), that its adapter is on the staging VLAN, and that the VLAN forwards PXE to your MECM distribution point (DHCP / IP helper).

**A dry run (Check)** reads the ticket, prints the request with where each value came from (`cpu: 4 (from comment 2026-10-08 10:05 - Jane Doe ...)`), checks everything and says where the VM would go. It builds nothing and writes nothing on the ticket. Run it first on every new kind of request.

## Set it up

1. **The template:** VMware Tools (open-vm-tools on Linux) installed and running; for Linux with cloud-init, VMware guest customization allowed (`disable_vmware_customization: false`); its network adapter set to **Connect at power on**.
2. **The vCenter account** needs, on the template, the target cluster, datastore(s), folder(s) and the staging port group: *Virtual machine > Provisioning > Deploy template, Customize guest, Read customization specifications; Virtual machine > Inventory > Create from existing; Virtual machine > Change configuration > Change CPU count, Change memory, Set annotation; Virtual machine > Interaction > Power on; Resource > Assign virtual machine to resource pool; Datastore > Allocate space; Network > Assign network*. A missing one is named in the error (`lacks the privilege ...`).
3. **The ServiceNow API credential** (the one the ticket jobs use): its account reads the ticket's table (incident, task ...) and adds work notes to it.
4. **Templates > Create template > Create job template:**
  - Name: `VM - deploy from ServiceNow ticket`
  - Inventory: any. Project: site-automation. Playbook: `playbooks/vm_deploy_from_ticket.yml`
  - Execution environment: your VMware one. Credentials: **VMware vCenter** and **ServiceNow API**
  - Tick **Prompt on launch** next to **Job type** (for the dry run)
  - **Survey:** one question, *Ticket number*, variable `vm_deploy_ticket`, text, required
  - Variables (the rules - a ticket cannot go beyond them):

YAML

```
vm_deploy_templates: [rhel9-gold, rhel8-gold]      # REQUIRED: only these are deployed (* allowed)
vm_deploy_assignment_group: Linux Operations      # only tickets assigned to this group
vm_deploy_clusters: [PROD-*]                    # [] = any
vm_deploy_defaults: {cluster: PROD-CL01, datastore: DSC-PROD, folder: Linux/New}
vm_deploy_max_cpu: 16
vm_deploy_max_memory_gb: 128
vm_deploy_domain: yoursite.mil
report_email_to: vmteam@yoursite.mil
```

5. **First run:** a test ticket assigned to your group with the four lines above; launch as **Check**, read the request and where it would go; then launch it for real.

| Setting                          | Default         | What it does                                                                          |
|----------------------------------|-----------------|---------------------------------------------------------------------------------------|
| <code>vm_deploy_templates</code> | <code>[]</code> | the templates a ticket may ask for ( <b>required</b> : nothing is deployed until set) |

| Setting                                                                                | Default                | What it does                                                                                                                                                    |
|----------------------------------------------------------------------------------------|------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>vm_deploy_clusters / _datastores / _folders</code>                               | <code>[]</code> (any)  | where a ticket may put the VM                                                                                                                                   |
| <code>vm_deploy_defaults</code>                                                        | <code>{}</code>        | what a ticket does not say (cluster, datastore, folder, cpu, memory_gb, domain)                                                                                 |
| <code>vm_deploy_max_cpu / vm_deploy_max_memory_gb</code>                               | 16 / 128               | caps                                                                                                                                                            |
| <code>vm_deploy_name_regex</code>                                                      | letters, digits, -     | the names allowed ('^[a-z0-9-]{1,15}\$' for Windows-safe names)                                                                                                 |
| <code>vm_deploy_assignment_group</code>                                                | <code>" "</code> (any) | only tickets assigned to this group                                                                                                                             |
| <code>vm_deploy_domain</code>                                                          | localdomain            | the DNS domain Linux customization sets                                                                                                                         |
| <code>vm_deploy_customize / vm_deploy_power_on</code>                                  | true / true            | guest customization; power on                                                                                                                                   |
| <code>vm_deploy_wait_customization / vm_deploy_wait_ip</code>                          | 900 / 600              | seconds to wait for each                                                                                                                                        |
| <code>vm_deploy_done_fields</code>                                                     | <code>{}</code>        | fields set on the ticket when the VM is up, e.g. <code>{state: "6", close_code: "Solution provided", close_notes: "VM deployed"}</code>                         |
| <code>vm_deploy_ticket_table</code>                                                    | from the prefix        | the table, for a ticket number with another prefix                                                                                                              |
| <code>vm_deploy_sites</code>                                                           | <code>{}</code>        | per site: <code>cluster</code> , <code>datastore</code> , <code>folder</code> , <code>vcenter</code> ... (a ticket's <code>site</code> : or Location picks one) |
| <code>vm_deploy_site_required</code>                                                   | true                   | with sites set up, a ticket must name its site                                                                                                                  |
| <code>vm_deploy_pxe_templates</code>                                                   | <code>[]</code>        | templates that boot from the network (PXE / MECM): no customization, MAC on the ticket                                                                          |
| <code>vm_deploy_max_disks / vm_deploy_max_disk_gb / vm_deploy_max_total_disk_gb</code> | 4 / 2048 / 4096        | disk caps: how many, the biggest one, all together                                                                                                              |
| <code>vm_deploy_mem_target_pct / vm_deploy_failover_hosts</code>                       | 80 / 1                 | automatic placement: a cluster fits when memory with this VM and that many hosts down stays under the %                                                         |

| Setting                                       | Default                                   | What it does                                                                                                                                  |
|-----------------------------------------------|-------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------|
| <code>vm_deploy_datastore_max_used_pct</code> | 85                                        | automatic placement: a datastore fits when it stays under this % with the VM at its full size                                                 |
| <code>vm_deploy_genai</code>                  | false                                     | GenAI reads requests written in plain words (propose, approve, build)                                                                         |
| <code>vm_deploy_genai_approval</code>         | true                                      | a GenAI proposal waits for a person (the workflow's approval); false = one run reads, checks and builds                                       |
| <code>vm_deploy_phase</code>                  | " "                                       | the workflow's steps: propose (read, check, place - build nothing), build (build what the ticket says - never ask GenAI); " " = the job alone |
| <code>vm_deploy_template_catalog</code>       | {}                                        | the templates GenAI may pick, each with a plain description                                                                                   |
| <code>vm_deploy_genai_model_timeout</code>    | <code>site_act_model_timeout</code> (600) | seconds the model may take                                                                                                                    |

**All settings** of the VMware jobs, with defaults: `VARIABLES_REFERENCE.md`.

**Who can change the request:** anyone who can comment on the ticket can change its lines, and an incident or a task has no approval of its own: the person who launches the job is the control, and the "Deploying" note records exactly what was built. When the job later picks up tickets by itself (a schedule), it will need a stronger gate - an approved change, or an approval step in a workflow.

## ESXi security settings (every night): SSH off, timeouts, lockdown

One job, `playbooks/esxi_security.yml`, sets these on every ESXi host, through vCenter:

| What                                                                                                                                                                                                                                                                          | Setting                             | Default                                                                                                     |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------|-------------------------------------------------------------------------------------------------------------|
| <b>SSH</b> off: the service stopped, and set to start manually                                                                                                                                                                                                                | <code>esxi_security_ssh</code>      | disabled (enabled, or " " = leave it)                                                                       |
| <b>ESXi Shell</b> off (the host's local shell, a separate service from SSH): the same                                                                                                                                                                                         | <code>esxi_security_shell</code>    | disabled (enabled, or " " = leave it)                                                                       |
| <b>Timeouts</b> at 600 s (10 minutes): SSH and the ESXi Shell stop by themselves 600 s after someone starts them ( <code>UserVars.ESXiShellTimeOut</code> ), and an idle SSH or shell session is logged out after 600 s ( <code>UserVars.ESXiShellInteractiveTimeOut</code> ) | <code>esxi_security_settings</code> | those two, at 600                                                                                           |
| <b>Lockdown mode normal:</b> the host is managed only through vCenter; the DCUI (the host's console) still works                                                                                                                                                              | <code>esxi_security_lockdown</code> | normal (disabled, or " " = leave it). <b>strict is refused</b> , and a host already in strict is left alone |

It changes **only what differs**: a host that already matches is not touched. So it is safe to run every night: it puts back whatever someone changed (SSH turned on, a timeout changed, lockdown turned off) and the email says which host and what. A dry run (Check) changes nothing and lists what it would change. Hosts not connected to vCenter are skipped and listed.

**Another timeout?** `esxi_security_settings` is a list of advanced settings, name: value. Add or remove lines, e.g. the DCUI and Host Client idle timeouts:

YAML

```
esxi_security_settings:
  UserVars.ESXiShellTimeout: 600
  UserVars.ESXiShellInteractiveTimeout: 600
  UserVars.DcuiTimeout: 600
  UserVars.HostClientSessionTimeout: 600
```

### Before the first real run: who logs in to the hosts directly?

Lockdown mode **blocks direct logins** to the hosts: SSH, the Host Client (<https://<host>/ui>), and any tool that connects to a host instead of vCenter. Most often that is the **vulnerability scanner** with ESXi credentials, and sometimes a monitoring or backup tool. Ask those teams first. For each such account, add it to the host's **exception users**: the job adds them **before** it turns lockdown on, and never removes one.

YAML

```
esxi_security_lockdown_exception_users: [svc_scanner, svc_monitor]
```

Exception users must be **local accounts on the host**, or **Active Directory accounts with their own permission on the host**. A `vsphere.local` account or an AD group does not work. Your vCenter service account for AAP is not affected: lockdown keeps vCenter's access.

### Set it up

1. **The vCenter account** needs the privileges in "Before you start", step 2 (the ESXi security settings row), on the hosts or their cluster.
2. **Templates > Create template > Create job template:**
  - Name: `ESXi - security settings`
  - Inventory: any (e.g. your VMware one). Project: site-automation. Playbook: `playbooks/esxi_security.yml`
  - Execution environment: your VMware one. Credentials: your **VMware vCenter** credential
  - Tick **Prompt on launch** next to **Job type** (for the dry run) and next to **Variables** (to try it on one host)
  - Variables: `report_email_to: vmteam@yoursite.mil` and, if you have them, `esxi_security_lockdown_exception_users: [...]`
  - Save.
3. **Dry run:** Launch, Job type **Check**. Read the report: every host and what it would change.
4. **One host for real:** Launch, Job type **Run**, Variables `esxi_security_hosts: [esx01.yoursite.mil]`. Check the host in the vSphere Client: **Configure > System > Services** (SSH and ESXi Shell stopped, policy Start and stop manually), **Advanced System Settings** (the two timeouts), **Security Profile > Lockdown Mode** (Normal). Have the scanner team run a scan of it.
5. **All hosts:** Launch again without `esxi_security_hosts`.
6. **Every night at midnight:** the template's **Schedules** tab > **Create schedule:**

- Name: **Nightly 00:00**
- Start date: tomorrow, start time **00:00**
- **Time zone: yours** (e.g. **America/New\_York**). AAP does not assume your local time: with the wrong zone, "midnight" runs at another hour.
- Repeat: **Daily**, every 1 day. No end. Save.

The nightly email comes only when a host was changed or failed (an empty night sends nothing); `esxi_security_email_only_if_changed: false` emails every night.

### A host that needs SSH (a support case)

The nightly run would turn SSH off again. Exempt the host while the case is open, in the template's Variables:

```
YAML
esxi_security_exclude_hosts: [esx07.yoursite.mil]
```

Then turn SSH on in the vSphere Client (and lockdown off if the vendor needs the Host Client). Remove the line when the case is closed: the next night sets the host back. To **undo** on purpose for every host: `esxi_security_lockdown: disabled`, and `esxi_security_ssh: ""` to leave SSH as it is.

| Setting                                  | Default               | What it does                                                                              |
|------------------------------------------|-----------------------|-------------------------------------------------------------------------------------------|
| <code>esxi_security_hosts</code>         | <code>[]</code> (all) | only these hosts (names, <code>*</code> and <code>?</code> allowed; the short name works) |
| <code>esxi_security_clusters</code>      | <code>[]</code> (all) | only the hosts of these clusters                                                          |
| <code>esxi_security_exclude_hosts</code> | <code>[]</code>       | never these hosts                                                                         |
| <code>vmware_datacenter</code>           | (all)                 | one datacenter only                                                                       |

## Email the report

Every VMware job can email its result: the secure boot report its whole list, the other jobs what they did. Email is off until there is an address to send to.

**Step 1 - ask your mail admins** for:

- the mail relay's name and port (e.g. `smtp.yoursite.mil`, port 25);
- whether it uses encryption: STARTTLS (usual on 25 and 587), SSL (port 465), or none;
- whether it needs a login, or instead accepts mail from allowed servers. Then the **AAP server's IP address** must be on its allowed list: the email leaves from the AAP server (with separate AAP execution nodes, from the node that runs the job - list those);
- which sender address to use (e.g. `aap-noreply@yoursite.mil`).

**Step 2 - the relay, in your** `playbooks/group_vars/all.yml` (VS Code, then Commit and Sync Changes; AAP picks it up at the next project sync):

```
YAML
report_email_smtp_host: smtp.yoursite.mil      # the relay from step 1
report_email_from: aap-noreply@yoursite.mil   # the sender from step 1
```

Add these lines only when step 1 says so:

| The relay                                                                                         | Add to all.yml                                                                                                                                                                                                                                                                           |
|---------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| port 25 or 587 with STARTTLS                                                                      | nothing: that is the default                                                                                                                                                                                                                                                             |
| a port other than 25                                                                              | <code>report_email_smtp_port: 587</code> (its port)                                                                                                                                                                                                                                      |
| SSL on port 465                                                                                   | <code>report_email_smtp_port: 465</code> and<br><code>report_email_security: ssl</code>                                                                                                                                                                                                  |
| no encryption at all                                                                              | <code>report_email_security: none</code> (then no login can be used)                                                                                                                                                                                                                     |
| its certificate is from your own CA (DoD) and the job says <code>certificate verify failed</code> | put the CA certificate (.pem, public) at <code>playbooks/files/ca/smtp-ca.pem</code> in the repository, <code>report_email_ca_path: "{{ playbook_dir }}/files/ca/smtp-ca.pem"</code> , and the line <code>playbooks/files/ca/</code> in <code>.site-local</code> (so an update keeps it) |

**Step 3 - only if the relay needs a login:** a credential, so the password stays out of Git.

1. AAP: **Automation Execution > Infrastructure > Credential Types > Create credential type**. Name it **SMTP relay**. Paste the `inputs:` part of `aap/credential_types/smtp_relay.yml` into **Input configuration** and the `injectors:` part into **Injector configuration**. Save.
2. **Credentials > Create credential**: type **SMTP relay**, the user name and password. Save.
3. On each VMware template: **Edit > Credentials**, add it next to the vCenter credential.

**Step 4 - who gets it.** One of:

- every VMware job, always: `report_email_to: [vmteam@yoursite.mil]` in `all.yml`;
- one template: its **Variables** box, `report_email_to: vmteam@yoursite.mil, ops@yoursite.mil`;
- whoever launches it: a survey question, variable `report_email_to`, type **Text**, not required. Left empty = no email.

Several addresses: separate them with commas. Copies: `report_email_cc`, the same way.

**Step 5 - test it** with the secure boot report: it only reads, so it is safe to run any time. Launch it with `report_email_to` set. In the job output, the task **Email | result** says `emailed "[AAP] VMware secure boot report: ..." to vmteam@yoursite.mil`, and the email arrives.

The other reports (health checks, certificate reports, ...) email the same way: EMAIL\_REPORTS.md lists them, and explains when a report job is green or red.

**What arrives:** a formatted email, subject `[AAP] VMware secure boot report: 29 VM(s) to fix` (or `[AAP] VM restart: web01`, `[AAP] VM snapshot before-patch: web01`, ...). The secure boot report has number boxes (VMs checked, secure boot on / off, BIOS) and a table per problem: each VM with its folder, ESXi host, guest host name, IP, power state and guest OS - so VMs with the same name can be told apart. At the end: the AAP job, who started it, when, and which vCenter. The email also carries a plain-text copy. A dry run (Check) sends nothing; it says `DRY RUN: would email ....` To change the look of the email: EMAIL\_REPORTS.md.

The same report is in the job's **Details > Artifacts** (`vm_report`, and for the secure boot report `vm_secure_boot` with every VM's details), which is also what the next step of a workflow receives.

**If it fails**, the job fails (so nobody misses it) and says why:

| The job says                                                                   | Fix                                                                                 |
|--------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|
| set report_email_smtp_host ... and report_email_from                           | step 2 is missing (or not synced to AAP yet)                                        |
| could not send the email through ...: [Errno 111] Connection refused/timed out | wrong name or port, or a firewall between the AAP server and the relay              |
| does not offer STARTTLS                                                        | the relay uses SSL (port 465) or no encryption: see step 2                          |
| TLS with the mail relay ... failed: ... certificate verify failed              | your own CA: the last row of step 2                                                 |
| refused the login                                                              | the user name or password in the SMTP relay credential                              |
| could not send ...: 530 ... authentication required                            | the relay needs a login (step 3), or the AAP server's IP on its allowed list        |
| refused every recipient/(refused: x@...)                                       | the relay does not accept that address (wrong address, or outside mail not allowed) |
| the password would cross the network unencrypted                               | a login with report_email_security: none: use starttls or ssl                       |

AAP's own **Notifications** (the template's Notifications tab) only say that a job succeeded or failed. This sends the report itself.

## What you see

- **restart:** Restart requested for (operating system restart through VMware Tools): web01. The job does not wait for the server to come back up; check its services afterwards, or run the health check.
- **shut down:** Shut down (...): web01., once the VM is off (up to vm\_power\_timeout, 600 s). VMs that were already off are listed as such.
- **snapshot:** Took snapshot "before-patch" (no memory) of: web01.
- **delete snapshot:** web01: deleted before-patch. If the name is not there: no snapshot named "x" (it has: ...).
- **notes:** the notes BEFORE and AFTER, for each VM.
- **change VLAN:** web01: Network adapter 2 moved VLAN200, or already on VLAN200.
- **secure boot report:** a count, then one line per VM to fix, with what to do.
- **datastore report, snapshot report, alarms report:** the report as text (the same tables as the email), with the counts in the task name, e.g. VMware alarms report: 1 critical, 2 warning alarm(s), 47 failed login(s) (events: last 24 h).
- **alarms ACT analysis:** VMware alarms ACT analysis: likely causes and fixes for 12 problem(s), then one row per problem with ACT's cause, fix and confidence.

### Refusals, and what to do:

- **VMware Tools is not running:** a guest restart or shutdown cannot work. Fix Tools, or use **hard** if losing unsaved data is acceptable.
- **No VM is named ...:** names are exact (case included), with no wildcards.
- **2 VMs are named ...:** rename one, or set vmware\_datacenter if they are in different datacenters.
- **Refused: ... is the AAP server or a protected VM:** by design. Do it in vCenter yourself.

## Troubleshooting

Find the message in the job's output (the red task near the end), or in the email, then do what the last column says.

### Every VMware job

| You see                                                                                                                                                                               | It means                                                                                                               | Do                                                                                                                                                                                                                   |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| No vCenter to talk to: attach a credential of type "VMware vCenter"                                                                                                                   | the template has no vCenter credential                                                                                 | Template > Edit > Credentials > add your <b>VMware vCenter</b> credential                                                                                                                                            |
| vCenter ... refused the login                                                                                                                                                         | wrong user name or password, or the account is locked                                                                  | fix the credential; log in to the vSphere Client with the same account to test it                                                                                                                                    |
| Could not connect to vCenter ...: ...<br>CERTIFICATE_VERIFY_FAILED                                                                                                                    | the execution environment does not trust vCenter's certificate                                                         | add vCenter's CA to the execution environment; to test only, <b>vmware_validate_certs: false</b> in the template's Variables                                                                                         |
| Could not connect to vCenter ...: timed out or Connection refused                                                                                                                     | the AAP node cannot reach vCenter on port 443                                                                          | a firewall rule from the AAP node (or execution node) to vCenter, port 443                                                                                                                                           |
| pyVmomi is not installed in the execution environment, or couldn't resolve module/action 'vmware.vmware...'                                                                           | the template uses an execution environment without the VMware parts                                                    | pick your VMware execution environment on the template ("Before you start", step 1)                                                                                                                                  |
| A VM, host or datastore you see in vSphere is missing from a report                                                                                                                   | the vCenter account cannot see it                                                                                      | give its role on that folder or cluster, with <b>Propagate to children</b> ("Before you start", step 2)                                                                                                              |
| A whole vCenter's datacenters and clusters are missing (the vSphere Client shows two or more vCenters)                                                                                | the jobs read the credential's vCenter, plus only those in <b>vmware_vcenters</b>                                      | <b>vmware_vcenters: [vc01..., vc02...]</b> in <b>all.yml</b> ("Before you start", step 5)                                                                                                                            |
| Red <b>vCenters not read</b> on top, - 1 <b>vCenter(s) not read</b> in the title (the job failed at <b>every vCenter was read</b> only with <b>vmware_vcenter_fail_unread: true</b> ) | that vCenter could not be reached, or refused the login; the rest of the report covers the others                      | the reason is in the row: <b>refused the login</b> = the account is not known there (in linked mode: give it a Global Permission); <b>timed out / Connection refused</b> = a firewall rule to that vCenter, port 443 |
| Could not find imported module support code for ...<br>site_vmware_... Looked for ... a nsible.module_utils.site_vcenter s                                                            | the folder <b>roles/vmware_vm/module_utils/</b> is missing: an update copied the changed files but not that new folder | copy <b>roles/vmware_vm/module_utils/</b> from the release into your repository, commit, sync the project                                                                                                            |
| the same vCenter as ...: read once in the artifacts                                                                                                                                   | a vCenter listed twice (a short name and its FQDN, or its IP)                                                          | nothing: it is read once                                                                                                                                                                                             |
| Refused: ... is the AAP server or a protected VM                                                                                                                                      | by design                                                                                                              | do it in vCenter yourself                                                                                                                                                                                            |

## Datastore and snapshot reports, snapshot cleanup

| You see                                                                                  | It means                                                                              | Do                                                                                           |
|------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------|
| The datastore heading still says 80%                                                     | <code>vm_datastore_warn_pct: 80</code> is set somewhere (it wins over the default 85) | remove it from <code>all.yml</code> and the template's Variables, or set the number you want |
| <code>vm_snapshot_max_age_days</code> must be a number of days                           | a survey or variable answer like <code>three</code>                                   | a number: <code>1, 2, 3</code>                                                               |
| The snapshot report says 3 days although you set 2                                       | in the workflow, the <b>survey answer</b> wins over the template's Variables          | change the survey's answer or its Default (section "Snapshots")                              |
| Refused: N snapshots are older than ... - more than <code>vm_snapshot_cleanup_max</code> | more than 50 at once (often a wrong number of days)                                   | check the days; for a big cleanup on purpose, raise <code>vm_snapshot_cleanup_max</code>     |
| NOT DELETED ...: no longer there, or no longer old enough / now kept                     | the snapshot changed between the report and the approval                              | nothing: it is safe. Run the workflow again                                                  |
| Taken by says unknown                                                                    | vCenter no longer keeps the event (often after 30 days)                               | nothing: the snapshot's description may name the AAP user                                    |

## ESXi security settings

| You see                                                                            | It means                                                                           | Do                                                                                          |
|------------------------------------------------------------------------------------|------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------|
| ... the vCenter account lacks the privilege Host.Config.... on this host           | the vCenter role misses that privilege                                             | add it to the role ("Before you start", step 2), on the hosts or their cluster              |
| lockdown strict is refused                                                         | esxi_security_lockdown: strict                                                     | use normal                                                                                  |
| No host matches (hosts: ...)                                                       | a typo in esxi_security_hosts / _clusters, or vmware_datacenter                    | the names as the vSphere Client shows them                                                  |
| X: the host has no such setting                                                    | that advanced setting does not exist on this ESXi version                          | check the name (case matters), or remove it from esxi_security_settings                     |
| lockdown exception users: ... - lockdown left disabled, so they are not locked out | the exception users could not be added, so lockdown was not turned on (on purpose) | fix the cause (often a privilege, or an account the host does not know) and run again       |
| lockdown is strict: left as it is                                                  | someone set strict on that host                                                    | nothing, or set it to normal in vCenter if strict was a mistake                             |
| A host is listed as <b>skipped</b>                                                 | it is not connected to vCenter                                                     | reconnect it; the next run sets it                                                          |
| The scanner (or another tool) cannot log in to a host any more                     | lockdown blocks direct logins                                                      | add its account to esxi_security_lockdown_exception_users and run again, or exempt the host |
| SSH is off again on a host you opened for support                                  | the nightly run did its job                                                        | esxi_security_exclude_hosts while the case is open                                          |
| The ESXi Shell stays on, and the report does not mention it                        | esxi_security_shell: "" is set somewhere (leave it as it is)                       | remove it: the default is disabled                                                          |
| Every night's email lists the same host as changed                                 | something turns the setting back every day (a host profile, a script, a person)    | find what it is: two tools fighting over one setting                                        |

## Deploy from a ServiceNow ticket

| You see                                                          | It means                                                                                                       | Do                                                                                                                                                                                |
|------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| No ServiceNow connection: attach the "ServiceNow API" credential | the template has no ServiceNow credential                                                                      | add it (Set it up, step 4)                                                                                                                                                        |
| No incident ticket INC... in ServiceNow                          | a wrong number, or the account cannot see it                                                                   | check the number; the account needs read access to that table                                                                                                                     |
| ServiceNow refused the read (HTTP 403)                           | the account may not read that table                                                                            | give it read access (e.g. the itil role)                                                                                                                                          |
| REFUSED: template '...' is not in vm_deploy_templates            | a template you have not allowed                                                                                | add it to <code>vm_deploy_templates</code> , or correct the ticket                                                                                                                |
| REFUSED: the ticket is assigned to ...                           | not your team's ticket                                                                                         | reassign it, or <code>vm_deploy_assignment_group</code>                                                                                                                           |
| the ticket does not say vm_name                                  | no <code>vm_name</code> : line                                                                                 | the requester adds a comment <code>vm_name: ...</code>                                                                                                                            |
| Nothing was built: ServiceNow refused the work note              | the account cannot update that table                                                                           | give it write access to work notes; nothing was built, by design                                                                                                                  |
| a VM named ... already exists in vCenter ...                     | the name is taken (or the job ran twice)                                                                       | another name in a comment, or nothing to do                                                                                                                                       |
| template ... is in vc01 / DC1, vc02 / DC1                        | the same template name in two places                                                                           | <code>vcenter</code> : (or <code>datacenter</code> :) in the ticket                                                                                                               |
| template ... is Windows: name a customization specification      | Windows needs Sysprep settings                                                                                 | <code>customization_spec</code> : <name> in the ticket (vCenter: Policies and Profiles) - or, for a network-boot (MECM) template, list it in <code>vm_deploy_pxe_templates</code> |
| the ticket does not say which site                               | sites are set up, the ticket names none (and its Location is not a site)                                       | <code>site</code> : ... in a comment; or <code>vm_deploy_site_required</code> : <code>false</code>                                                                                |
| site '...' is not one of vm_deploy_sites                         | a site name you have not set up                                                                                | correct the ticket, or add the site                                                                                                                                               |
| ... ready to build in ... - waiting for the approval             | the propose step (or the job alone with <code>vm_deploy_genai_approval</code> : <code>true</code> ): by design | approve the workflow (or launch the job again)                                                                                                                                    |
| has no GenAI proposal for its text as it is now                  | the build step, but the ticket changed after the proposal (or there is none)                                   | run the workflow from the start                                                                                                                                                   |
| questions for the requester (nothing built)                      | GenAI could not trace a value to the ticket, or something is missing                                           | the requester answers in a comment ( <code>memory_gb</code> : 16); run the job again                                                                                              |
| DROPPED (not in the ticket): vm_name ...                         | GenAI gave a value the ticket does not contain (made up, or guessed)                                           | nothing: it is a question now                                                                                                                                                     |

| You see                                                 | It means                                                                              | Do                                                                                                                        |
|---------------------------------------------------------|---------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------|
| GenAI did not run: ...                                  | no ACT model key, the model not reachable, or a timeout                               | the ACT model key credential on the template; see docs/ADDING_ACT.md; or <b>key: value</b> lines                          |
| no cluster among ... has room                           | automatic placement: every candidate would go over the memory target with a host down | another site / cluster, less memory, or a higher <b>vm_deploy_mem_target_pct</b>                                          |
| no datastore among ... has room                         | every candidate datastore would go over 85 % with the VM                              | another datastore list, smaller disks                                                                                     |
| disk 1 cannot shrink                                    | <b>disks_gb</b> asks for a smaller first disk than the template's                     | the first size at least the template's                                                                                    |
| A network-boot VM sits at "PXE-E..." / "No boot device" | the VLAN does not reach MECM's PXE, or MECM does not answer this computer             | the DHCP / IP helper on the staging VLAN; the computer imported in MECM, or unknown-computer support on the task sequence |
| lacks the privilege VirtualMachine.Provisioning...      | the vCenter account's role                                                            | grant it (Set it up, step 2)                                                                                              |
| Customization no result within 900 s                    | VMware Tools does not run in the template, or cloud-init blocks VMware customization  | fix the template; the VM exists - set its host name by hand, or delete it and run again                                   |
| no IPv4 address within 600 s                            | no DHCP lease on the staging VLAN, or VMware Tools does not run                       | check DHCP; the IP shows in vCenter when it comes                                                                         |
| the template's ... is not set to connect at power on    | the template's adapter                                                                | tick <b>Connect at power on</b> on the template                                                                           |

## Capacity planning

| You see                                                               | It means                                                                      | Do                                                                                                         |
|-----------------------------------------------------------------------|-------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------|
| not enough history / No trend: vCenter has fewer than 7 daily samples | a new vCenter, or its statistics level or daily interval was changed          | wait a week; check vCenter > Configure > General > Statistics: the 1-day interval enabled, level 1 or more |
| vCenter's daily statistics interval is turned off: no history         | the 1-day interval is disabled                                                | enable it (same place); usage now is still shown                                                           |
| A runway that jumps from run to run                                   | the trend fit is poor: VMs were moved, snapshots deleted, a datastore grown   | read the Trend fit column; use <code>vm_capacity_trend_days: 60</code> for a longer, steadier line         |
| cannot lose a host                                                    | the cluster has as many hosts down (or in maintenance) as it plans to survive | bring hosts back; or set <code>vm_capacity_failover_hosts: 0</code> for a cluster that needs no failover   |
| No cluster matches ...                                                | a typo in <code>vm_capacity_clusters</code>                                   | the cluster's name as the vSphere Client shows it                                                          |
| ACT did not run: ...                                                  | as for the alarms analysis (key, network, certificate)                        | the ACT rows above                                                                                         |
| Math and GenAI differ a lot                                           | ACT weighed something the line does not (a jump, a limit that comes first)    | read ACT's reasoning; plan for the sooner of the two                                                       |

## Alarms report

| You see                                                                                   | It means                                                                                                                                      | Do                                                                                                                                                              |
|-------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>types: unknown object type(s)</code><br>...                                         | a typo in <code>vm_alarm_types</code>                                                                                                         | use <code>vcenter</code> , <code>datacenter</code> , <code>cluster</code> , <code>host</code> , <code>datastore</code> , <code>vm</code> , <code>network</code> |
| <code>vCenter ... has no datacenter named ...</code>                                      | a typo in <code>vmware_datacenter</code>                                                                                                      | the exact name, as the vSphere Client shows it                                                                                                                  |
| An alarm the vSphere Client shows is missing                                              | it is on a VM, datastore or network (not read by default), or in another datacenter                                                           | add its type to <code>vm_alarm_types</code> (e.g. <code>[vcenter, datacenter, cluster, host, vm]</code> ); check <code>vmware_datacenter</code>                 |
| No events at all, or fewer than in the vSphere Client                                     | <code>vm_alarm_hours: 0</code> , or the account sees only part of the inventory (vCenter shows an account only the events of what it can see) | give the role at the top of the vCenter inventory, with Propagate to children                                                                                   |
| Only the newest 5000 events were read                                                     | a busy vCenter                                                                                                                                | lower <code>vm_alarm_hours</code> , or raise <code>vm_alarm_max_events</code>                                                                                   |
| Only the first 200 are shown                                                              | a long table                                                                                                                                  | the job's <b>Artifacts</b> ( <code>vm_alarms</code> ) have all of them; or raise <code>vm_alarm_max_rows</code>                                                 |
| The job shows <b>failed</b> with <code>critical alarm(s) ... host(s) not connected</code> | <code>vm_alarm_fail: true</code> is set                                                                                                       | expected (for a workflow); remove it for a report only                                                                                                          |
| Reading alarms and events from vCenter ... <code>failed: ...</code>                       | something this job did not expect from your vCenter                                                                                           | send that line to the people who maintain this repository                                                                                                       |

## Alarms ACT analysis

| You see                                                                                                          | It means                                                                                                                                     | Do                                                                                                                                                                                                                                        |
|------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| nothing to analyze, and no email                                                                                 | nothing is wrong: no alarm, no problem in the window                                                                                         | nothing. <code>vm_alarm_act_email_if_none: true</code> sends an email anyway                                                                                                                                                              |
| A dry run (Check): ACT was not called                                                                            | the job ran as Check                                                                                                                         | launch it with Job type Run                                                                                                                                                                                                               |
| This job has no API key for provider genai: attach ...                                                           | no <b>ACT model key</b> credential on the template, or its GenAI field is empty                                                              | add the credential (with the key for your provider) to the template                                                                                                                                                                       |
| This AAP node cannot reach the model at ... (... timed out / Connection refused / Name or service not known ...) | a firewall, DNS or a missing proxy between the AAP node and the model                                                                        | ask for HTTPS (443) from the AAP node to that address, or set your proxy: <code>site_act_env: {HTTPS_PROXY: "http://proxy.yoursite.mil:8080"}</code> in <code>all.yml</code>                                                              |
| ... Its certificate is not trusted: set <code>site_act_ca</code> ...                                             | the execution environment does not trust the model's certificate authority                                                                   | put that CA certificate at <code>playbooks/files/ca/model-ca.pem</code> , <code>site_act_ca: "{{ playbook_dir }}/files/ca/model-ca.pem"</code> in <code>all.yml</code> , and <code>playbooks/files/ca/</code> in <code>.site-local</code> |
| ACT failed: ... with 401 or Unauthorized                                                                         | the key is wrong or expired                                                                                                                  | a new key in the credential                                                                                                                                                                                                               |
| ACT failed: ... with 429 or Too Many Requests                                                                    | the key's quota per minute                                                                                                                   | run it less often, or lower <code>vm_alarm_act_max_items</code>                                                                                                                                                                           |
| model turn exceeded the 600s total timeout (older releases: 90s, 300s), often after a transient provider failure | the model (or the provider, retrying after an error) took longer than ACT allows for one answer: <code>site_act_model_timeout</code> , 600 s | run it again (a provider hiccup passes); <code>site_act_model_timeout: 900</code> in <code>all.yml</code> (keep <code>site_act_timeout</code> above it); fewer items ( <code>vm_alarm_act_max_items: 20</code> ); or a faster model       |
| ACT produced no result (it timed out after 1200s ...)                                                            | a slow model, or a lot of evidence                                                                                                           | <code>site_act_timeout: 1800</code> , or <code>vm_alarm_act_max_items: 20</code>                                                                                                                                                          |
| ACT's answer is below as text                                                                                    | the model did not use the table format                                                                                                       | the answer is still in the email. Run it again, or try another model ( <code>site_act_models</code> )                                                                                                                                     |
| A row says ACT gave no answer for this one                                                                       | the model skipped that problem                                                                                                               | run it again; fewer problems ( <code>vm_alarm_act_max_items</code> ) help a small model                                                                                                                                                   |
| The job is red, but the email came                                                                               | ACT could not run; the email lists the problems and says why                                                                                 | fix what the email says, then run it again                                                                                                                                                                                                |

**Test the way to the model by hand**, from the AAP server (the execution environment uses the same network):

```
shell
curl -sS -o /dev/null -w '%{http_code}\n' https://api.genai.mil/v1/chat/completions
```

Any number (401, 404, 405) means the way is open. `Could not resolve host`, a timeout, or `SSL certificate problem` is the problem the job reports.

## Limits

- **Change VLAN between a standard and a distributed switch.** Moving an adapter from a standard switch port group to a distributed one, or back, fails inside the `vmware.vmware` collection (2.11); the job says so and changes nothing. Do that one in vCenter. Moves between port groups of the same kind of switch work.
- The guest keeps its IP settings after a VLAN change. A new VLAN usually needs new ones.
- `vm_names` are exact VM names. To act on many VMs, list them.
- Tested against vcsim, VMware's vCenter simulator (`tests/vmware/run_vmware_test.sh`), not against a real vCenter. Run each job once as **Check**, then on a test VM, before you use it on real servers. The simulator raises no alarms: the alarms report's alarm handling is tested with recorded data (`tests/test_filters.py`, `tests/test_vmware_modules.py`). The ACT analysis is tested with the real ACT against a scripted model (`tests/test_alarm_act_contract.py`), not a real one.

# 11. NetApp ONTAP health report

One job, `playbooks/ontap_health_report.yml`, reads every NetApp ONTAP cluster you list and emails one report: what is broken first, then capacity, protection and services. Each row is coloured: red = critical, amber = warning, green = OK, blue = for your information.

**It only reads.** It sends GET requests to each cluster's REST API (the same API System Manager uses) and nothing else: the code refuses any other kind of request. It runs on the AAP side, like the VMware jobs: no inventory of the clusters is needed, no SSH, no Limit.

## What the report shows

At the top: the title with the number of critical and warning findings, then boxes with the main counts (nodes down, health alerts, hardware, full volumes and aggregates, SnapMirror, network, certificates). Then, in this order:

| Section                                  | What it shows                                                                                                                                                                           | Colours                              |
|------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------|
| <b>Clusters</b>                          | each cluster: ONTAP version, system health, nodes up, space used on its aggregates, NTP servers. A cluster that could not be read is listed in red with the reason                      | health, nodes, space                 |
| <b>Nodes</b>                             | state, uptime (days hours:minutes), model, version, whether the partner could take over (storage failover), service processor, NVRAM battery, the node's clock against the AAP server's | green UP / red DOWN, and each status |
| <b>System health</b>                     | ONTAP's health monitors ( <code>system health subsystem show</code> ) that are not OK                                                                                                   | amber degraded                       |
| <b>Health alerts</b>                     | what the health monitors raised ( <code>system health alert show</code> ): node, severity, alert, resource, probable cause, corrective action                                           | red critical / major, amber minor    |
| <b>Error events (last 24 h)</b>          | ONTAP's event log (EMS): emergency, alert and error events, the same event on the same node counted once                                                                                | red alert / emergency, amber error   |
| <b>Hardware: controllers and shelves</b> | each controller and disk shelf: power supplies, fans, temperature, failed parts (FRUs)                                                                                                  | red failed                           |
| <b>Disks</b>                             | broken disks; disks rebuilding, copying or in maintenance; unassigned disks                                                                                                             | red / amber / blue                   |
| <b>Spare disks</b>                       | spare disks per node (a failed disk is rebuilt onto a spare)                                                                                                                            | red none, amber few                  |
| <b>Network interfaces (LIFs)</b>         | LIFs down, not on their home port, or disabled                                                                                                                                          | red / amber / blue                   |
| <b>Ports down</b>                        | enabled Ethernet and FC ports that are not up                                                                                                                                           | red in use, blue no cable / not used |

| Section                                  | What it shows                                                                                                                                                                                            | Colours                                                     |
|------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------|
| <b>Port errors (CRC)</b>                 | ports with receive errors or CRC errors since the node booted (CRC counters: ONTAP 9.11+)                                                                                                                | amber from 1, red from 10,000                               |
| <b>Aggregates</b>                        | every data aggregate, fullest first: node, home node, size, used, available, used %                                                                                                                      | amber 85 %, red 90 %                                        |
| <b>Aggregates not on their home node</b> | after a takeover, aggregates still on the partner                                                                                                                                                        | amber                                                       |
| <b>Volumes</b>                           | volumes 85 % full or more, and volumes offline or restricted, with their aggregate, state, type, size, available, used %                                                                                 | amber 85 %, red 90 %, red offline                           |
| <b>Inodes (files)</b>                    | volumes using 85 % or more of their inodes (no new file can be created when they run out)                                                                                                                | amber 85 %, red 90 %                                        |
| <b>Snapshot space</b>                    | the 25 volumes whose snapshots hold the most: type, volume size, snapshot count, reserve, snapshot used, % of reserve, % of volume                                                                       | amber from 100 % of the reserve (spilling into the volume)  |
| <b>SnapMirror</b>                        | relationships that are unhealthy, lagging, broken off, paused or uninitialized, and transfers running now: mirror state, status, healthy, lag, last transfer, progress, progress updated, policy, reason | red unhealthy / 48 h lag, amber 24 h lag, blue transferring |
| <b>Cluster and SVM peers</b>             | cluster peers not available, SVM peers not peered                                                                                                                                                        | red / amber                                                 |
| <b>SVMs not running</b>                  | stopped storage VMs (a stopped DR destination is normal: blue)                                                                                                                                           | red                                                         |
| <b>CIFS (SMB) servers</b>                | each CIFS server: enabled or not, and whether its domain controllers answer                                                                                                                              | red stopped / no DC, amber a DC down                        |
| <b>LUNs</b>                              | LUNs not mapped to any host (forgotten?) and LUNs not online                                                                                                                                             | amber / red                                                 |
| <b>Certificates expiring</b>             | certificates expiring within 60 days or expired                                                                                                                                                          | amber 60 days, red 14 days                                  |

A section with nothing to report says so in one line ("Every LIF is up and at home."). Long tables show the problems first and are cut at 50 rows (`ontap_report_max_rows`); the job's artifacts (`ontap_health`) have every finding.

## Before you start

### 1. A read-only account on every cluster

The report needs an ONTAP account with the built-in **readonly** role, for the **http** application (the REST API). On each cluster, as an administrator (`cluster1` is the cluster's admin SVM, its name):

```
security login create -vserver cluster1 -user-or-group-name svc_aap_ro -application http
-authentication-method password -role readonly
```

Use the **same user name and password on every cluster**: the job has one credential for all. (An Active Directory account works too, with `-authentication-method domain`, if the cluster is set up for AD logins.)

## 2. The way from AAP to the clusters

The AAP server (or the execution node that runs the job) must reach each cluster's **management address** on HTTPS, port 443. Test from the AAP server (it asks for the password):

```
curl -sk -u svc_aap_ro "https://cluster1.yoursite.mil/api/cluster?fields=version"
```

A short JSON answer with the version means the account, the network and the API are fine.

## 3. The cluster's certificate

ONTAP's management certificate is often self-signed, or from your own CA. The job checks it by default. Either:

- put the CA certificate (public, `.pem`) at `playbooks/files/ca/ontap-ca.pem`, add `ontap_ca_path: "{{ playbook_dir }}/files/ca/ontap-ca.pem"` to `playbooks/group_vars/all.yml`, and the line `playbooks/files/ca/` to `.site-local` (so an update keeps it); or
- set `ontap_validate_certs: false` (the password still travels encrypted; the job just does not check who answers).

## 4. The credential

1. **Automation Execution > Infrastructure > Credential Types > Create credential type**: name `NetApp ONTAP`; paste the `inputs` block of `aap/credential_types/netapp_ontap.yml` into **Input configuration** and the `injectors` block into **Injector configuration**. Save.
2. **Credentials > Create credential**: type `NetApp ONTAP`, name e.g. `NetApp ONTAP read-only`, the user name and password from step 1. Save.

## Set it up

1. **Templates > Create template > Create job template**:

- Name: `NetApp - ONTAP health report`
- Inventory: any (e.g. your VMware one: the job runs on the AAP side)
- Project: `site-automation`. Playbook: `playbooks/ontap_health_report.yml`
- Execution environment: any with Python 3 (no NetApp collection is needed)
- Credentials: **NetApp ONTAP read-only**
- Tick **Prompt on launch** next to **Variables** (to try one cluster first)
- Variables:

YAML

```
ontap_clusters: [cluster1.yoursite.mil, cluster2.yoursite.mil]
report_email_to: storage-team@yoursite.mil
```

- Leave Limit empty. Save.

2. **First run: one cluster.** Launch with Variables `ontap_clusters: [cluster1.yoursite.mil]`. Read the report in the job's output. Lines that start with `cluster1: could not read ...` or `... reduced detail ...` under a section say what this cluster's ONTAP version or the account's role did not allow: see Troubleshooting.
3. **All clusters**, then a schedule: **Schedules > Create schedule**, e.g. every day at 06:00 (set **your** time zone).

A dry run (Check) reads and prints the same report, and sends no email.

**SnapMirror** relationships are read on their **destination** cluster: list the destination clusters too, or the relationships into them are not checked.

## Settings

Set them in the template's Variables, or in `playbooks/group_vars/all.yml` for every run.

| Setting                                                                 | Default                             | What it does                                                                                                 |
|-------------------------------------------------------------------------|-------------------------------------|--------------------------------------------------------------------------------------------------------------|
| <code>ontap_clusters</code>                                             | <code>[]</code>                     | the clusters' management addresses                                                                           |
| <code>ontap_validate_certs</code> / <code>ontap_ca_path</code>          | <code>true</code> / (none)          | check the clusters' certificates; the CA file to check them with                                             |
| <code>ontap_volume_warn_pct</code> / <code>_crit_pct</code>             | <code>85</code> / <code>90</code>   | volume used % for amber / red                                                                                |
| <code>ontap_aggr_warn_pct</code> / <code>_crit_pct</code>               | <code>85</code> / <code>90</code>   | aggregate used %                                                                                             |
| <code>ontap_inode_warn_pct</code> / <code>_crit_pct</code>              | <code>85</code> / <code>90</code>   | inodes used %                                                                                                |
| <code>ontap_snapshot_reserve_warn_pct</code> / <code>_crit_pct</code>   | <code>100</code> / <code>150</code> | snapshot space as % of the snapshot reserve                                                                  |
| <code>ontap_snapmirror_lag_warn_hours</code> / <code>_crit_hours</code> | <code>24</code> / <code>48</code>   | time since the last good SnapMirror transfer                                                                 |
| <code>ontap_cert_warn_days</code> / <code>_crit_days</code>             | <code>60</code> / <code>14</code>   | days before a certificate expires                                                                            |
| <code>ontap_time_drift_warn_seconds</code> / <code>_crit_seconds</code> | <code>60</code> / <code>300</code>  | a node's clock against the AAP server's                                                                      |
| <code>ontap_port_errors_warn</code> / <code>_crit</code>                | <code>1</code> / <code>10000</code> | receive or CRC errors on a port, since boot                                                                  |
| <code>ontap_min_spares</code>                                           | <code>1</code>                      | fewer spare disks than this on a node is amber (0 is red)                                                    |
| <code>ontap_ems_hours</code> / <code>ontap_ems_max</code>               | <code>24</code> / <code>500</code>  | the event log window, and the most events read per cluster                                                   |
| <code>ontap_report_max_rows</code>                                      | <code>50</code>                     | rows per table                                                                                               |
| <code>ontap_snapshot_top</code>                                         | <code>25</code>                     | volumes in the snapshot space table                                                                          |
| <code>ontap_report_fail</code>                                          | <code>false</code>                  | <code>true</code> = the job shows <b>failed</b> on a critical finding or a cluster not read (for a workflow) |
| <code>ontap_timeout</code>                                              | <code>30</code>                     | seconds for one request                                                                                      |

**All settings** of this job, with their defaults: `ontap_report`.

## Troubleshooting

| You see                                                                              | It means                                                                                                                         | Do                                                                                                                        |
|--------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------|
| No ONTAP account: attach a credential of type "NetApp ONTAP"                         | the template has no NetApp ONTAP credential                                                                                      | add it (Before you start, step 4)                                                                                         |
| ontap_clusters is empty                                                              | no cluster listed                                                                                                                | ontap_clusters: [...] in the template's Variables                                                                         |
| A cluster <b>UNREACHABLE</b> with <b>timed out</b> or <b>Connection refused</b>      | the AAP node cannot reach the cluster on port 443                                                                                | a firewall rule from the AAP node to the cluster's management address; test with the <b>curl</b> above                    |
| <b>UNREACHABLE</b> with <b>Name or service not known</b>                             | the name does not resolve on the AAP node                                                                                        | use the name DNS knows, or the IP address                                                                                 |
| <b>UNREACHABLE</b> with <b>CERTIFICATE_VERIFY_FAILED</b>                             | the cluster's certificate is not trusted                                                                                         | Before you start, step 3                                                                                                  |
| the ONTAP account was refused (HTTP 401)                                             | wrong user or password, the account does not exist on <b>that</b> cluster, it is locked, or it lacks the <b>http</b> application | create it the same on every cluster (step 1); unlock it ( <b>security login unlock</b> )                                  |
| cluster1: could not read alerts - HTTP 403 ... the account's role may not allow this | the role does not allow that command                                                                                             | use the built-in <b>readonly</b> role (step 1)                                                                            |
| ... reduced detail - ONTAP refused a field ...                                       | this ONTAP version does not know a field; the job asked again with fewer                                                         | nothing: the section is still shown, a column may say <b>?</b> . Send the line to the people who maintain this repository |
| CRC counters need ONTAP 9.11 or later                                                | older ONTAP has no CRC counter in its REST API                                                                                   | nothing: receive errors are still shown                                                                                   |
| SnapMirror is empty, but you have relationships                                      | they are read on the destination cluster                                                                                         | add the destination cluster to <b>ontap_clusters</b>                                                                      |
| An aggregate you know is missing                                                     | root aggregates are not returned by ONTAP's REST API                                                                             | by design: root aggregates are not listed                                                                                 |
| A volume shows <b>?</b> for size                                                     | it is offline: ONTAP reports no space for it                                                                                     | it is listed in red as offline                                                                                            |
| The email is too long                                                                | many findings                                                                                                                    | lower <b>ontap_report_max_rows</b> / <b>ontap_snapshot_top</b> ; the artifacts keep everything                            |
| could not read ems - ... only the first 500 records were read                        | many error events                                                                                                                | raise <b>ontap_ems_max</b> , or lower <b>ontap_ems_hours</b>                                                              |

## Limits

- The fields the job reads were checked against NetApp's ONTAP REST API reference for ONTAP 9.10 to 9.16. It was tested against a fake ONTAP API (`tests/netapp/`), not a real cluster: run it on one cluster first, and send any `could not read` or `reduced detail` line you see.
- The CLI commands it reads through REST (system health, storage failover, SnapMirror progress) are ONTAP `show` commands at admin level, which the `readonly` role may run.
- It reports; it changes nothing. Fixes are done in System Manager or the ONTAP CLI.

## 12. Emailed reports: the formatted email and its look

Jobs can email what they found or did, as a **formatted (HTML) report** with a plain-text copy for mail clients that do not show HTML. One email per job run, whatever the number of hosts.

## Which jobs email

| Template                             | Playbook                                                             | The email                                                                                                                                                                                   |
|--------------------------------------|----------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Health check (Linux)                 | <code>health_check.yml</code>                                        | findings of every host (critical first), hosts that did not report, healthy hosts                                                                                                           |
| Windows health check                 | <code>win_health_check.yml</code>                                    | the same, for Windows servers                                                                                                                                                               |
| Database health                      | <code>database_health.yml</code>                                     | the same, for the database check                                                                                                                                                            |
| Certificate report (Linux / Windows) | <code>cert_report.yml</code> ,<br><code>win_cert_report.yml</code>   | findings, plus every certificate, soonest expiry first: red expired, amber within 30 days, blue within 60                                                                                   |
| POA&M status                         | <code>poam_status.yml</code>                                         | overdue, due-soon and missing items                                                                                                                                                         |
| ServiceNow health                    | <code>servicenow_health.yml</code>                                   | the instance and the MID Servers                                                                                                                                                            |
| Troubleshoot (Linux / Windows)       | <code>troubleshoot.yml</code> ,<br><code>win_troubleshoot.yml</code> | what the checks found                                                                                                                                                                       |
| Apply approved ACT fix               | <code>act_fix_approved.yml</code>                                    | the checks after the fix                                                                                                                                                                    |
| VM secure boot report                | <code>vm_secure_boot_report.yml</code>                               | VMs with secure boot off, BIOS VMs (folder, host, guest, IP)                                                                                                                                |
| VM datastore report                  | <code>vm_datastore_report.yml</code>                                 | datastores 85 / 90 % used or more, then every datastore, fullest first                                                                                                                      |
| VM snapshot report                   | <code>vm_snapshot_report.yml</code>                                  | snapshots older than N days, kept ones, then every snapshot with size and age                                                                                                               |
| VM snapshot cleanup                  | <code>vm_snapshot_cleanup.yml</code>                                 | what was deleted (and the space freed), what could not be                                                                                                                                   |
| VM alarms report                     | <code>vm_alarm_report.yml</code>                                     | triggered alarms, host connection, failed logins, VM events, other errors and warnings, configuration issues, every host - each row coloured by severity                                    |
| VM alarms ACT analysis               | <code>vm_alarm_act_analysis.yml</code>                               | per problem: ACT's likely cause, evidence, fix and confidence (coloured); what to do first                                                                                                  |
| NetApp ONTAP health report           | <code>ontap_health_report.yml</code>                                 | clusters, nodes, health alerts, events, hardware, disks, network, aggregates, volumes, inodes, snapshots, SnapMirror, SVMs, CIFS, LUNs, certificates - problems first, coloured (NETAPP.md) |

| Template                                                             | Playbook                            | The email                                                                                                                                                                  |
|----------------------------------------------------------------------|-------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| VM capacity planning                                                 | <code>vm_capacity_report.yml</code> | overall and per cluster: CPU, memory with a host down, vCPUs per core, VMs that fit, hosts recommended, datastores with runways; with GenAI, its estimate next to the math |
| ESXi security settings                                               | <code>esxi_security.yml</code>      | the hosts it changed (what), the ones that failed (why), every host's settings before the run. By default only when a host was changed or failed                           |
| VM restart, shut down, snapshot, delete snapshot, notes, change VLAN | <code>vm_*.yml</code>               | what was done to which VM                                                                                                                                                  |

## Turning it on

1. **The mail relay, once**, in your `playbooks/group_vars/all.yml` (then Commit, Sync Changes, and sync the project in AAP). A relay that accepts the AAP server by a whitelist needs nothing more:

```
YAML
report_email_smtp_host: smtp.yoursite.mil
report_email_from: aap-noreply@yoursite.mil
report_email_security: none           # only if the relay does not offer STARTTLS
```

Encryption, your own CA and a relay login: VMWARE.md, "Email the report".

2. **Who gets each report: on each template.** Templates > the template > Edit > **Variables**:

```
YAML
report_email_to: linux-team@yoursite.mil, ops@yoursite.mil
```

Or a survey question with the variable `report_email_to` (Text, not required). Putting `report_email_to` in `all.yml` instead makes **every** job in the table above email those people, each time it runs - usually more than you want.

Not in an AAP inventory **group's** Variables: the email is sent by the AAP side of the job, which never gets a group's variables. The template's Variables, a survey, the inventory's own Variables and `all.yml` all work.

| Setting                                                     | Default                     | What it does                                                         |
|-------------------------------------------------------------|-----------------------------|----------------------------------------------------------------------|
| <code>report_email_to</code> / <code>report_email_cc</code> | (none)                      | who gets it; a list, or text with commas                             |
| <code>report_email_html</code>                              | <code>true</code>           | send the formatted version too; <code>false</code> = plain text only |
| <code>report_email_html_template</code>                     | <code>report.html.j2</code> | the look of the formatted version (below)                            |
| <code>report_email_subject_prefix</code>                    | <code>[AAP]</code>          | put before every subject                                             |

The subject is the report's title with its counts, e.g. `[AAP] Health check: 3 finding(s) on 2 of 48 host(s)`. Every email ends with the AAP job number, the template, who started it and when. A dry run

(Check) sends nothing. A job without `report_email_to` says `No email: report_email_to is not set`.

**All settings**, with their defaults: `site_email`.

## Green or red: when a report finds problems

AAP shows a job as **successful (green)** or **failed (red)**, nothing in between. What a report job does when it finds problems:

| Job                                                                           | Default                                                                                           | Change it with (template Variables)                                                    |
|-------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------|
| VM secure boot report, VM datastore report, VM snapshot report                | <b>green</b> : the report and email say what to fix                                               | <code>vm_secure_boot_fail: true / vm_datastore_fail: true</code> = red                 |
| VM snapshot cleanup                                                           | green; <b>red</b> when a deletion failed                                                          | -                                                                                      |
| VM alarms report                                                              | <b>green</b> : the report and email say what is wrong                                             | <code>vm_alarm_fail: true</code> = red on a critical alarm or a host not connected     |
| VM alarms ACT analysis                                                        | green; <b>red</b> when ACT could not run (no key, no network): the email still lists the problems | -                                                                                      |
| VM capacity planning                                                          | <b>green</b> ; red when GenAI was asked for and could not run                                     | <code>vm_capacity_fail: true</code> = red when a cluster or datastore is critical      |
| ESXi security settings                                                        | green; <b>red</b> when a host could not be set (the email says why)                               | -                                                                                      |
| NetApp ONTAP health report                                                    | <b>green</b> : the report and email say what is wrong                                             | <code>ontap_report_fail: true</code> = red on a critical finding or a cluster not read |
| Health checks, certificate reports, database health, POA&M, ServiceNow health | <b>red</b> when there are findings: AAP workflows start tickets (and ACT) from that               | <code>site_fail_on: []</code> = green, report only                                     |

So for a template whose job is to email a report to people, set `site_fail_on: []` and it stays green; the email is the alert. Keep the default on templates that feed a ticket workflow.

Red also means one of these, whatever the setting: a host could not be reached, or the job itself broke (a wrong setting, a missing credential). Those always need a look.

## What a report looks like

A title (also the subject) in a coloured band - green, amber, red or blue - then summary boxes with numbers, then sections: each a heading, a line of advice, and a table (or a list). In tables that have a severity, each row is tinted (light red, amber or blue) and the severity cell is filled in its colour, so the problems stand out at a glance. Every value from a host or vCenter is escaped, so a VM name cannot add markup to the email.

## Changing the look

The look is one file, `roles/site_email/templates/report.html.j2` (Jinja and HTML, styles inline because Outlook ignores most style sheets). To change it without touching the repository's file:

1. Copy it to `playbooks/files/email/report.html.j2` in your repository and edit the copy: colours, logo text, fonts, a banner line your site requires.
2. In `all.yml: report_email_html_template: "{{ playbook_dir }}/files/email/report.html.j2"`.
3. Add the line `playbooks/files/email/` to the file `.site-local` at the top of your repository (create it if it is not there). Without it, the next update from a release **deletes** your copy: the release does not have that file.

Check your copy after an update that changes the report layout (the CHANGELOG says so when that happens).

The **title** of a report (the coloured band, also the subject) is written by the playbook, not the look file: for the VMware reports, `_vm_report_subject` in `roles/vmware_vm/tasks/<report>.yml`. An edit there is undone by the next update, unless that file is in `.site-local` - and then you no longer get the release's fixes to it either.

## Emailing a report from another playbook

Describe the report as data, then include the role. For example, at the end of a play that runs on localhost:

YAML

```
- name: Email the report
  ansible.builtin.include_role:
    name: site_email
  vars:
    _report:
      title: "Certificate report: 3 expire within 30 days" # also the subject
      status: warning # ok | warning | critical | info: the colour of the band
      subtitle: "212 certificates on 48 servers" # optional
      summary: # optional: number boxes
        - {label: Certificates, value: 212}
        - {label: Expire within 30 days, value: 3, status: warning}
      sections:
        - title: Expire within 30 days
          status: warning
          text: "Renew them before the date shown." # optional
          columns: [Server, Subject, Expires]
          rows: "{{ my_rows }}" # a list of lists, one per row, in the order of columns
        - title: Notes
          lines: [one line, another line] # a bullet list instead of a table
      footer: "Source: playbooks/cert_report.yml" # optional
  when: report_email_to | default([]) | length > 0
```

A section has `columns` and `rows` (a table; an empty list shows "None."), or `lines`, or just `text`. To colour a table, add `row_status` (one status per row: `critical`, `warning`, `info`, `ok` or `'`; tints the row) and `cell_status` (per row, one status per cell; fills that cell), e.g. `row_status: [critical]` and `cell_status: [[critical, '', ']]` for a red row with a red first cell. In a play that runs on many hosts, build the rows from `hostvars` and add `run_once: true` to the include, so one email goes out, not one per host.

## Troubleshooting

The job's output says what happened to the email, in the task `Email | result` (or `No email`).

| You see                                                                                                 | It means                                                                        | Do                                                                                                                                                  |
|---------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------|
| No email: <code>report_email_to</code> is not set                                                       | no recipient for this template                                                  | add <code>report_email_to</code> to the template's Variables (or its survey)                                                                        |
| To email the report, set <code>report_email_smtp_host</code> ... and <code>report_email_from</code> ... | the relay is not set up                                                         | step 1 of "Turning it on", in <code>all.yml</code>                                                                                                  |
| the mail relay ... does not offer STARTTLS                                                              | the relay has no encryption on that port                                        | <code>report_email_security: none</code> (or <code>ssl</code> with port 465)                                                                        |
| TLS with the mail relay ... failed: ... certificate verify failed                                       | the execution environment does not trust the relay's certificate                | <code>report_email_ca_path</code> to the CA file (VMWARE.md, "Email the report"); list <code>playbooks/files/ca/</code> in <code>.site-local</code> |
| could not send the email through ...: ... timed out or Connection refused                               | the AAP node cannot reach the relay                                             | a firewall rule from the AAP node (and every execution node) to the relay's port                                                                    |
| the mail relay ... refused every recipient                                                              | the relay does not accept mail from the AAP node, or for those addresses        | ask the mail admins to add the AAP node's address to the relay's allowed list                                                                       |
| emailed "..." to ... but nothing arrives                                                                | the relay accepted it; something after it held it                               | look in junk / quarantine; ask the mail admins to trace it by the subject; check that <code>report_email_from</code> is an address they allow       |
| DRY RUN: would email ...                                                                                | the job ran as Check                                                            | nothing: a dry run never sends                                                                                                                      |
| The email shows plain text only                                                                         | <code>report_email_html: false</code> , or a mail gateway removed the HTML part | ask the mail admins whether HTML mail is stripped; the text copy has the same tables                                                                |
| The job is red, but the email came                                                                      | a check report found problems ( <code>site_fail_on</code> )                     | by design for ticket workflows; <code>site_fail_on: []</code> on email-only templates ("Green or red")                                              |

## 13. Secrets: where they live

### Recommendation

**Store every secret as an AAP credential, and inject it into jobs as environment variables or files through custom credential types.** Nothing secret goes into Git, inventory, job template variables, or surveys.

- AAP encrypts credential values at rest and never shows them again after you save them. People can be allowed to **use** a credential in a job without being able to **see** it.
- If your organization already runs **CyberArk, HashiCorp Vault, or Azure Key Vault**, keep the secrets there and link the AAP credential fields to it with AAP's credential lookup plugins (on the credential's form, the key icon next to a secret field). AAP fetches the value when a job starts; nothing else in this repository changes.
- Use **Ansible Vault** only for playbooks you must also run outside AAP. It is weaker: one vault password decrypts every value, and the encrypted values sit in Git forever (rotating the password does not un-publish old ciphertext).

## Where each secret lives

| Secret                                 | AAP credential type                                                      | Reaches the playbook as                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
|----------------------------------------|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| SSH key + sudo for the hosts           | built-in <b>Machine</b>                                                  | the SSH connection (never visible to tasks)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| WinRM login for Windows servers        | built-in <b>Machine</b> (user name and password; Kerberos or NTLM)       | the WinRM connection (never visible to tasks). Use a <b>separate</b> credential from the SSH one; see WINDOWS.md                                                                                                                                                                                                                                                                                                                                                                                                                     |
| Git access for the project             | built-in <b>Source Control</b>                                           | used only by the project sync                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| MySQL root / STIG Manager DB passwords | <b>STIG Manager database</b> (aap/credential_types/stigman_database.yml) | env STIGMAN_MYSQL_ROOT_PASSWORD, STIGMAN_DB_PASSWORD                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| TLS certificate, key, CA chain         | <b>TLS certificate</b> (tls_certificate.yml)                             | files; env TLS_CERT_FILE, TLS_KEY_FILE, TLS_CA_FILE hold the paths                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| Docker Hub / quay.io login for hosts   | <b>Container registry login (hosts)</b> (registry_login.yml)             | env REGISTRY_*, used with podman login --password-stdin                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| ACT model key                          | <b>ACT model key</b> (act_model_key.yml)                                 | env GENAI_KEY / ASKSAGE_API_KEY on the controller. <b>Note:</b> when ACT runs on a managed host, the job hands the key to that one ACT process on the task's stdin (never on a command line, so not in ps or the sudo log) and it is gone when the process ends. So the host's administrators and its process auditing can see it while ACT runs: use a key with a spending limit, and rotate it if a host is untrusted. <b>VM - alarms ACT analysis</b> runs ACT on the AAP node itself: there the key never reaches another server |
| ACT GenAI beta key                     | <b>ACT GenAI beta key</b> (act_genai_beta_key.yml)                       | env GENAI_BETA_KEY                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| Teams webhook URL                      | <b>Teams webhook</b> (teams_webhook.yml)                                 | env TEAMS_WEBHOOK_URL                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| ServiceNow API account                 | <b>ServiceNow API</b> (servicenow_api.yml)                               | env SN_HOST, SN_USERNAME, SN_PASSWORD (basic auth on the controller; never sent to hosts)                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| MariaDB monitor account                | <b>MariaDB monitor</b> (mariadb_monitor.yml)                             | env MARIADB_MONITOR_USER / _PASSWORD; the password goes to the client on stdin and reaches it only as MYSQL_PWD in its own environment (never a command line, never a fact)                                                                                                                                                                                                                                                                                                                                                          |

| Secret                                               | AAP credential type                                                                | Reaches the playbook as                                                                                                                                                                                                                     |
|------------------------------------------------------|------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| STIG Manager API client                              | <b>STIG Manager API</b> ( <code>stigman_api.yml</code> )                           | <code>env STIGMAN_TOKEN_URL, STIGMAN_CLIENT_ID, STIGMAN_CLIENT_SECRET</code> (token request is <code>no_log</code> )                                                                                                                        |
| Java keystore password                               | <b>Keystore password</b> ( <code>keystore_password.yml</code> )                    | <code>env KEYSTORE_PASSWORD</code> ; the inventory names the variable ( <code>password_env</code> ), never the value; the check reads it on stdin and <code>keytool</code> takes it with <code>-storepass:env</code> (never a command line) |
| vCenter account (VMware jobs)                        | built-in <b>VMware vCenter</b>                                                     | <code>env VMWARE_HOST, VMWARE_USER, VMWARE_PASSWORD</code> on the controller (the jobs talk only to vCenter)                                                                                                                                |
| NetApp ONTAP read-only account (ONTAP health report) | <b>NetApp ONTAP</b> ( <code>netapp_ontap.yml</code> )                              | <code>env ONTAP_USERNAME, ONTAP_PASSWORD</code> on the controller; sent only to the clusters' REST API over HTTPS (basic auth); the job sends GET requests only                                                                             |
| Mail relay login (emailed reports)                   | <b>SMTP relay</b> ( <code>smtp_relay.yml</code> ), only if the relay needs a login | <code>env SMTP_USERNAME / SMTP_PASSWORD</code> ; sent only over STARTTLS or SSL (refused with <code>report_email_security: none</code> )                                                                                                    |

Create one credential of each type **per environment** (dev / test / prod) so a test job can never use production secrets.

### Why environment variables and files, not extra variables

A credential type can inject into `extra_vars`, but then the secret is an ordinary play variable: any `debug` task, any failed task's error message, or `-vvv` output can print it. An environment variable or a file on the controller has to be read on purpose (`lookup('ansible.builtin.env', 'NAME')`), which keeps secrets out of the variable namespace unless a task asks for them. Files are the right choice for certificates and keys: AAP writes them to private temporary files that disappear when the job ends.

### Rules for every playbook in this repository

- no\_log**: `true` on every task that receives, compares, or writes a secret. On a loop, also set `loop_control: label`: to something harmless, or the item is printed.
- diff**: `false` on any task that writes secret material (copy/template of a key, a `.cnf` with a password). With AAP's *Show changes* enabled, a diff prints the file's contents.
- Never put a secret on a command line.** Command lines are visible in `ps`, in audit logs (`execve` records), and in verbose job output. Instead:
  - pass it on **stdin**: `podman secret create NAME -, podman login --password-stdin;`

- or in a **file** the program reads (`--defaults-extra-file` for MySQL clients), created with mode `0600` and removed afterwards;
  - never `mysql -pPassword` or `curl -u user:password`.
  - never through Ansible's `environment:` keyword on a Linux host with `become`: Ansible puts those variables on the `sudo /bin/sh -c '...'` command line, which `ps` shows and `sudo` writes to the host's journal / `/var/log/secure` - `no_log` does not prevent that. Send the secret on the task's `stdin:` to a small `sh` that exports it (see `roles/check_mariadb/tasks/health.yml`).
4. **On the hosts, secrets are podman secrets**, not environment lines in unit files:
- mounted as files where the image supports it (`MYSQL_ROOT_PASSWORD_FILE`, `MYSQL_PASSWORD_FILE`), with `mode=0400` and the container user's uid;
  - injected as an environment variable only when the application cannot read a file (`STIGMAN_DB_PASSWORD`): `Secret=NAME,type=env,target=VAR`.
- Podman keeps the value out of the unit file and out of `podman inspect` (checked on podman 5.8; verify on your RHEL's podman). It is still readable by root on the host and inside the container, so protect root and the container.
5. **Private keys on disk:** `0600`, owned by root; certificate and CA files `0644`.
6. **Don't register command output that can contain secrets** without `no_log`, and never `debug: var=` a registered result from such a task.
7. **Fail closed:** assert a secret is present (with `no_log`) instead of letting a task run with an empty password.
8. **Changing a secret in AAP is not the same as rotating it.** For a database password the value inside MySQL must change too (`ALTER USER`). The deploy role refuses to silently swap a deployed database password; rotation is a separate, deliberate run.

## The real trust boundary

AAP will not show a stored secret, but **any playbook it runs can print it**. So anyone who can:

- push to the branch or tag the AAP project uses,
- edit a job template (its playbook, credentials, or verbosity), or
- run a template with *Prompt on launch* for variables that change what the playbook does,

can extract every secret that template uses. Therefore:

- **protect the branch** (required reviews) and pin the project to a **tag**;
- give most people **Execute** on templates, not **Admin**; give only automation owners **Use** on credentials;
- keep *Prompt on launch* for variables off on templates that hold production credentials;
- review the **Activity Stream** (who changed which credential or template) regularly.

## Keeping secrets out of Git

- `.gitignore` in this repository blocks common key/cert/vault-password file names.
- `.pre-commit-config.yaml` runs **gitleaks** on each commit (`pre-commit install`).
- The inventory in `inventories/example/` holds placeholders only; the real inventory lives in your work Git or as an AAP inventory.
- If a secret is ever committed: rotate it first, then remove it from history. Removing it from history alone is not enough, because clones and caches keep it.

## Rotation checklist

| Secret                   | How                                                                                                                                                                          |
|--------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| TLS certificate          | update the <b>TLS certificate</b> credential, re-run <code>stigman_deploy.yml</code> (nginx restarts)                                                                        |
| Registry token           | update the credential; the next run logs in with it                                                                                                                          |
| STIG Manager DB password | password-rotation run: <b>ALTER USER</b> in MySQL, update the podman secret, restart the API; then update the AAP credential. The playbook for this comes with the day-2 set |
| SSH key                  | add the new public key to the hosts, update the <b>Machine</b> credential, remove the old key                                                                                |
| ACT / Teams              | update the credential; nothing else                                                                                                                                          |
| ServiceNow API account   | change the password in ServiceNow, update the <b>ServiceNow API</b> credential, then run <code>ServiceNow - test ticket</code> (SERVICENOW_SETUP.md) to prove it works       |
| MariaDB monitor account  | change it in the database ( <b>ALTER USER</b> ), update the <b>MariaDB monitor</b> credential, run a health check with <code>health_checks=mariadb</code> (MARIADB.md)       |
| WinRM login              | change it in the directory, update the Windows <b>Machine</b> credential, run <code>Windows connection test</code>                                                           |

## 14. Approved commands: what ACT may run by itself

ACT is a tool that runs **on each Linux server** when a job asks it to. This page explains, in plain words, which commands ACT runs by itself, which ones wait for a person, and how you add a command to the "run without asking" list.

**Short version.** There are three kinds of command:

| Kind                               | Examples                                                                    | What happens                                                                                                   |
|------------------------------------|-----------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------|
| <b>Reading</b>                     | <code>systemctl status chronyd, df -h, journalctl -u nginx, ss -tlnp</code> | ACT runs it. It never asks. <b>Nothing on this page limits reading.</b>                                        |
| <b>Changing</b> , not on your list | <code>systemctl restart nginx, dnf -y install httpd</code>                  | Refused. The job says <b>NEEDS APPROVAL</b> and shows the exact command. A person approves it in the workflow. |
| <b>Changing</b> , on your list     | the ones you wrote in <code>site_act_allow</code>                           | ACT runs it by itself (only at level <code>self-heal</code> ).                                                 |

The list is one setting: `site_act_allow`. You add a command by adding one line to it.

### Read this first: two safety rules that you cannot switch off

1. **Some commands are never run by a job**, whatever you write. ACT itself never pre-approves a *catastrophic* command (recursive delete, disk or volume destruction, power off, reboot). And the step that applies an approved fix (the *Apply approved ACT fix* job, and the service-watch fix step) has its own last check: reboot, shutdown, `kexec`, `mkfs`, `fdisk`, `usermod`, `passwd`, recursive deletes of system folders and a few more are refused **even after a person approved them** - also with a path in front (`/sbin/reboot`) or inside `bash -c '...' /$( )`. Patching and restarts have their own job: RUNBOOKS.md, section "Patch hosts".
2. **AAP itself is never fixed by a job**. Hosts in the groups `aap` and `aap_hosts` are diagnose-only: ACT reads and proposes, and the apply step prints the command for a person to run by hand. Nothing you set can remove those two groups (they are built into the code).

Everything below works inside those two rules.

### How a command becomes "approved": the three ways

#### Way 1: a person approves it (the workflow), nothing to configure

This is the default at level `diagnose`.

1. The health check finds a problem. ACT investigates by reading, then proposes the fix.
2. Because the fix changes something and is not on any list, ACT is **refused** and the job fails with **NEEDS APPROVAL - ... ACT proposes: `systemctl start chronyd`**.
3. The workflow stops at its **Approval** step. The approver reads the proposal and approves or denies.
4. *Apply approved ACT fix* runs **exactly the approved text**, no model involved, then the checks run again.

What the approver sees is what runs. To set this up, see `ADDING_ACT.md` and `WORKFLOWS_AND_SCHEDULES.md`.

Two details worth knowing:

- Only a **plain single command** can go through this path (see the character rule below). A proposal such as `cd /x && ./fix.sh` or `something | tee file` is listed in the report as a **manual fix**: a person runs it by hand. That is deliberate.
- The approval covers that command on the hosts that proposed it, once. It does not add the command to any list. Next week the same problem needs the same approval.

### Way 2: you put it on the list (`site_act_allow`), for fixes you trust

Use this for boring, safe, repeated fixes: start a stopped service, restart a hung one. It only takes effect at level `self-heal`. At `explain` and `diagnose`, the list is ignored.

YAML

```
site_act_allow:
- 'systemctl restart (chronyd|rsyslog|crond)'
- 'systemctl start (chronyd|rsyslog|crond)'
```

Where to put it, step by step, is in the next section.

### Way 3: `--allow` on the ACT command line (when you run ACT by hand)

If you run `act` on a server yourself, the same idea is `--allow`:

shell

```
act --allow 'systemctl restart (nginx|httpd)' "the web server is down; find out why and fix it"
```

`site_act_allow` is this, done for you: each list item becomes one `--allow` when the job calls ACT. (Full ACT reference: "Approved commands" in ACT's own README.)

At the prompt that ACT shows when you run it by hand (`run? [y]es / [n]o / [e]dit / [q]uit`) there is **no "always allow" answer**. This is on purpose: the list is the only way to make ACT stop asking, and the list is in Git, where a reviewer can see it.

## Add an approved command, step by step

You need: VS Code with this repository, and the command as you would type it on the server.

### Step 1: write the command exactly as ACT would run it

Try it by hand on a lab server first: `sudo systemctl restart chronyd`. Note the words in between: `systemctl restart chronyd` (drop `sudo`; ACT already runs as root on the host).

### Step 2: turn it into a pattern

A pattern is a *regular expression* that must match the **whole** command, from the first character to the last. `systemctl restart chronyd` on its own is already a valid pattern that matches exactly that command and nothing else. Start there. Widen it only when you have to:

| Pattern                                    | Matches                                | Does not match                                 |                                        |                        |
|--------------------------------------------|----------------------------------------|------------------------------------------------|----------------------------------------|------------------------|
| <code>systemctl restart chronyd</code>     | <code>systemctl restart chronyd</code> | <code>systemctl restart chronyd.service</code> |                                        |                        |
| <code>`systemctl restart (chronyd\`</code> | <code>rsyslog\`</code>                 | <code>crond)`</code>                           | <code>systemctl restart rsyslog</code> | <code>systemctl</code> |
| <code>`systemctl (start\`</code>           | <code>restart) (chronyd\`</code>       | <code>rsyslog)`</code>                         | <code>systemctl start rsyslog</code>   | <code>systemctl</code> |

| Pattern                                                       | Matches                                        | Does not match                                                    |
|---------------------------------------------------------------|------------------------------------------------|-------------------------------------------------------------------|
| <code>systemctl restart chronyd\.</code> <code>service</code> | <code>systemctl restart chronyd.service</code> | <code>systemctl restart chronydXservice</code>                    |
| <code>`podman restart (mysql\` nginx)`</code>                 | <code>nginx`</code>                            | <code>podman restart nginx</code> <code>podman rm -f nginx</code> |

(In the YAML file, write `|` normally. The `\|` above is only how a table shows a bar.)

Rules for the pattern:

- **The whole command must match.** `systemctl restart .*` looks like "restart anything", but it can only match a plain command. It is also far too wide: it would let ACT restart `sshd` and the network. Name the services.
- **Only a plain single command counts.** The command may contain only letters, digits and these characters: `space _ . / : = @ % + , -`. Anything else is never pre-approved, however you write the pattern: no `;` && `||` (chaining), no `|` (pipes), no `>` `<` (redirection), no `$` `( )` or backticks, no `*` or `?` (globs), no quotes, no `~`, no line breaks. Such a command still needs a person.
- **A dot in a pattern means "any character".** Write `\.` when the exact dot matters, as in `chronyd\.``service`.
- **No file edits.** A pattern covers `run` commands only; ACT changing a file (edit or write) always needs a person.
- **Dangerous commands are never approved by a pattern**, whatever you write: ACT's danger tier (package removal such as `dnf -y remove httpd, usermod -aG wheel alice, chmod 777, podman rm -f mysql, rm -f, ...`) and reboots, power-off and disk tools always need a person (tested with ACT 0.6.21). They come back as **NEEDS APPROVAL** / manual fixes.
- **Never put these on the list** anyway: `.*` on its own, `systemctl (stop|disable|mask) .*`, `kill .*`. Those are ordinary changes a pattern *can* approve, and far too wide: `systemctl stop .*` would let ACT stop `sshd`.

### Step 3: put it in the right settings file

`site_act_allow` is a setting like any other (VARIABLES.md explains where settings live). Pick the place that matches who should get it:

| Who gets the command                                 | Where to write it                                                                                                            |
|------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------|
| <b>Every host</b>                                    | <code>playbooks/group_vars/all.yml</code> : add the lines at the end of the file (copy the example below).                   |
| <b>One group</b> (for example <code>stigman</code> ) | <code>playbooks/group_vars/stigman.yml</code> (the file name is the AAP group name)                                          |
| <b>One host</b>                                      | <code>playbooks/host_vars/&lt;host name as in the inventory&gt;.yml</code> (create the folder and file if they do not exist) |
| <b>One template, every run</b>                       | the job template's <b>Variables</b> box (extra variables), YAML: <code>site_act_allow: ['systemctl restart chronyd']</code>  |

**A list replaces, it does not add.** If `all.yml` has two commands and `stigman.yml` has one command, the `stigman` hosts get **only that one**. Repeat the ones you keep. (The exact order of who wins is in VARIABLES.md.)

Example, `playbooks/group_vars/all.yml`:

```
YAML
site_act_allow:
  - 'systemctl (start|restart) (chronyd|rsyslog|crond)'
```

Use single quotes. Keep the indentation. Spaces, never tabs.

#### Step 4: save, commit, sync

In VS Code: save, **Commit**, **Sync Changes**. AAP picks the file up on the next run (the project syncs when *Update revision on launch* is ticked).

#### Step 5: use the level self-heal

Launch the check with `use_act = yes` and `site_act_level = self-heal` (the survey question on the Health check and Troubleshoot templates; `ADDING_ACT.md`).

## How to verify it worked

Do all three. They take ten minutes and catch the usual mistakes.

### 1. Check the pattern matches (and only what you meant). No server needed. In PowerShell:

```
$p = 'systemctl restart (chronyd|rsyslog|crond)'      # your pattern, without the quotes' worth of
YAML
foreach ($c in 'systemctl restart chronyd',          # should be True
               'systemctl restart sshd',             # should be False
               'systemctl restart chronyd; reboot') { # should be False
    '{0,-40} {1}' -f $c, ($c -cmatch ('^A(?:' + $p + ')\z'))
}
```

Expected output: `True, False, False`. (Verified. The same test in Python, on a machine that has `vendor/act/act`: load it with `importlib` and call `compile_preapproved/preapproved_pattern`; ACT applies exactly `\A(?:your pattern)\Z`.) Add a line for every command you want to *stay* refused and check it prints `False`.

**2. Check ACT itself agrees, on a lab server.** Log in, put your key in the environment the way `SECRETS.md` describes, and run ACT in the same mode a job uses. It will not ask anything; it either runs what you allowed or refuses:

```
shell
sudo systemctl stop chronyd          # break it on purpose (lab server only)
act --non-interactive --allow 'systemctl (start|restart) chronyd' \
  --result-file /tmp/r.json "chronyd is stopped; find out why and fix it"
echo "exit code: $?"                 # 0 = fixed, 4 = something was refused
python3 -m json.tool /tmp/r.json | less # see commands[].approval and denied[]
```

In `/tmp/r.json`: `commands[].approval` says `pre_approved` for what your list allowed (with the `pattern` that matched) and `auto` for reads. `denied[]` lists what was refused, each with `pre_approvable`: `true` means "a pattern could approve this" (add one if you want it); `false` means a person has to do it.

**3. Check the job.** On the same lab server, launch **Health check** with `use_act = yes`, `site_act_level = self-heal`. In the output:

- ACT | what ACT says shows `APPLIED (pre-approved): systemctl start chronyd;`

- the report afterwards says **cleared**: ....

Then break something that is **not** on your list and check the job says **NEEDS APPROVAL** and does not touch it. A dry run (job type *Check*) does not call ACT, so it cannot test this; see DRY\_RUNS.md.

### If it does not work

|                                                 | Why                                                                                                                                    | Fix                                                                                                         |                                                                                   |
|-------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------|
| <b>APPROVAL</b> for a command and you listed    | the pattern does not match the <b>whole</b> command; or a listed command contains a <code>;</code> , <code>\</code>                    | <b>, quote or glob; or the level is</b> explain/diagnose, <b>not</b> self-heal                              | run check 1 with the exact command from the job output; fix the pattern the level |
| command is in <b>manual</b> not <b>proposed</b> | it is not a plain single command                                                                                                       | ask ACT for a simpler fix, or apply by hand                                                                 |                                                                                   |
| something seems ignored                         | a more specific place wins (the group or host file, or the template's Variables box)                                                   | see the order in VARIABLES.md; the job prints a notice when a settings file overrides the AAP Variables box |                                                                                   |
| invalid pattern                                 | ACT does not run: the job shows <b>ACT error ... invalid --allow pattern '...'</b> and the reason (for instance an unbalanced bracket) | fix the YAML/regex                                                                                          |                                                                                   |
| command on an AAP host does not                 | by design: AAP is diagnose-only                                                                                                        | run it by hand from the printed line                                                                        |                                                                                   |
| approved command was not at apply time          | the apply-time guard (reboot, shutdown, <b>mkfs</b> , <b>usermod</b> , recursive delete of a system folder, ...)                       | do it by hand if you really mean it                                                                         |                                                                                   |

## Reference: what the guard rails are

| Where                                                                                | Rule                                                                                                                                                                | Can you change it?                                            |
|--------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------|
| ACT, every command                                                                   | proven read-only commands run without asking; everything else asks                                                                                                  | no (it is the design)                                         |
| ACT, <b>--allow</b> / <b>site_act_allow</b>                                          | anchored to the whole command; plain characters only; never catastrophic; never a file edit                                                                         | you choose the patterns                                       |
| ACT, <b>--auto</b>                                                                   | runs reads and ordinary changes without asking, still asks for the danger tier. Not used by these jobs                                                              | not used                                                      |
| <b>site_act_diagnose_only_groups</b>                                                 | hosts where ACT never fixes; the list <i>you</i> write replaces the shipped one ( <b>sat</b> , <b>idm</b> , ...), but <b>aap</b> , <b>aap_hosts</b> are always kept | you choose the groups; you cannot remove <b>aap/aap_hosts</b> |
| Apply-time guard ( <b>_apply_refuse</b> , <b>roles/site_findings/vars/main.yml</b> ) | last check before an approved command runs                                                                                                                          | edit the code to change it                                    |
| Host log                                                                             | <b>/var/log/site-automation-fix.log</b><br>on each host lists what the apply job ran, when, and for which AAP job and user                                          |                                                               |

The apply-time guard applies to the **approval** path (Way 1). Commands run by **self-heal** (Way 2) are run by ACT itself, so what protects you there is your list. Keep it short and boring.

Windows

`site_act_allow` is read by the Linux jobs. Windows servers use the approval path only (ACT proposes, a person approves); see WINDOWS.md. If you run ACT for Windows by hand, its `-Allow` option follows the same rules as Linux `--allow`:

|                                                                                                                | Linux ACT (--allow)                                                           | Windows ACT (-Allow)                                                         |
|----------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------|------------------------------------------------------------------------------|
| A pattern that names a danger-tier command (for example <code>dnf -y remove httpd, podman rm -f mysql</code> ) | <b>does not</b> approve it                                                    | <b>does not</b> approve it                                                   |
| A danger-tier command in an unattended run                                                                     | <b>always refused</b> ; only a person at an interactive prompt can approve it | the same                                                                     |
| Catastrophic commands                                                                                          | never                                                                         | never                                                                        |
| <code>git</code> , cmd built-ins, aliases such as <code>ls</code> and <code>dir</code> in an unattended run    | follow the rules above                                                        | <b>still ask</b> , because they were never on Windows ACT's unattended list  |
| Result file ( <code>--result-file / -ResultFile</code> )                                                       | written (since 0.6.16)                                                        | written (since 0.6.16; its <code>tokens</code> field is filled since 0.6.21) |

(ACT 0.6.20 for Linux, released for a few minutes on 2026-09-30, let a pattern approve a danger-tier command; 0.6.21 and later (this project carries 0.6.22) do not. The Windows rows are taken from ACT-Windows 0.6.22 and its self-tests and were not run by hand on a Windows server.)

# 15. The MariaDB and MySQL check, step by step

There are **two ways to run the database check**, and both check **MariaDB and MySQL** (MySQL 5.7, 8.0 and 8.4, MariaDB 10.x and 11.x, and Percona), installed on the host or **in podman containers** (one, several, found automatically, rootless):

| How                             | Job template                                                                                                                                                 | Use it when                                                                                                                |
|---------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------|
| <b>Its own template</b>         | <b>Database health - MariaDB / MySQL</b><br>( <code>playbooks/database_health.yml</code> ;<br><code>SETUP_AAP.md</code> )                                    | the database team wants its own job, schedule and tickets. Runs only this check, on every database server it is pointed at |
| <b>Part of the Health check</b> | <b>Health check</b> with <code>health_checks = mariadb, mysql</code> or <code>database</code> (all three are the same check; <code>daily</code> includes it) | you already run the Health check and want the databases in it                                                              |

This page shows how to turn the check on, how to give it a login, and how to see that it works. Nothing here needs you to know Ansible. It changes nothing on the database: the check only asks status questions and never reads your data.

Words used here:

- **container name** - the name podman gives a container. You read it with `podman ps -a`.
- **rootless** - the container belongs to an ordinary user (not root). Only that user can see it.
- **settings file** - a small file in `playbooks/group_vars/` that holds your choices for one group of hosts. See `VARIABLES.md`.

## The short version

1. On the database host, find the container name: `sudo podman ps -a` (rootless: `sudo -u THE_USER podman ps -a`).
2. Put the name in the settings file for your MariaDB group (section 3).
3. Ask the DBA for a read-only monitoring account and put it in an AAP credential (section 4).
4. Run the **Database health** job (or the *Health check* with `health_checks = mariadb`) and read "How to verify" below.

## 1. Which hosts the check runs on

The check runs on a host when **any** of these is true:

- the host is in one of the inventory groups `mariadb_hosts`, `mysql_hosts`, `database_hosts`, `mariadb` or `mysql` (any of the five works, for either engine), or
- the host has a container name set (section 2), or
- you set `check_mariadb_enabled: true` for the host or its group.

The **Database health** template does not need this: every server it runs on is checked. It runs on the groups `mariadb_hosts`, `mysql_hosts`, `database_hosts`, `mariadb` and `mysql` (since 0.6.0; before that only the first three). A **Limit** only narrows that list; it cannot add a server that is in none of them. For a server outside those groups, put `target: <group or host>` in the template's **Variables** box, or add the server to one of those groups. If `target` matches no host, the job stops with a message that lists the groups your inventory does have. On every other host, in the Health check, the job prints `skipped: the MariaDB/MySQL check is`

off for `<host>` and moves on. That is normal; it is not an error.

## 2. Tell it where MariaDB is

Pick **one** of the four ways. Put the lines in the settings file (section 3).

| Your situation                                                                                                                                                                                                                                                                        | Setting                                                                                                                                                                                                                         |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| MariaDB or MySQL is installed on the host itself (a service)                                                                                                                                                                                                                          | nothing to set. The check finds the service <code>mariadb</code> , <code>mysqld</code> or <code>mysql</code> and tells the engine and version apart by itself. To name the service: <code>check_mariadb_service: mariadb</code> |
| In a podman container, with <b>no</b> database server installed on the host, or started by a unit that runs podman ( <code>podman compose up -d</code> , <code>podman start</code> - such a unit shows <code>active (exited)</code> or <code>inactive</code> while the database runs) | usually nothing to set (0.7.0 and later): the check sees that and finds the container by its image. Naming it is still clearer: <code>check_mariadb_container: mariadb</code>                                                   |
| One container                                                                                                                                                                                                                                                                         | <code>check_mariadb_container: snow-mariadb</code>                                                                                                                                                                              |
| Several containers on the same host                                                                                                                                                                                                                                                   | <code>check_mariadb_containers: [snow-mariadb, other-db]</code>                                                                                                                                                                 |
| You do not want to list them: find them                                                                                                                                                                                                                                               | <code>check_mariadb_container_discover: true</code>                                                                                                                                                                             |

Details that matter:

- **Discovery** looks at every container on the host, running or not, and picks the ones whose **image name** matches `check_mariadb_container_image_regex` (default `(mariadb|mysql|percona)`). If your image has another name, change the pattern, for example `'(mariadb|snowdb)'`. Discovery that finds nothing is reported as a warning `mariadb:no-container`, so a wrong pattern is not silent.
- With **several** instances on one host, each is checked on its own. Their findings start with the container name, for example `[snow-mariadb] 143 of 151 connections in use`.
- The check runs the database's own client **inside** the container (`podman exec`), so you do not need a client on the host and you do not open any network port.
- The engine (MariaDB or MySQL) and its version are read from the server, so the same settings work for both.
- A container that is named but does not exist is reported as `mariadb:down` ("MariaDB/MySQL is not running (container ... is missing)"), not skipped.

### Rootless containers

When the containers belong to a user, name that user:

```
YAML
check_mariadb_container_user: appuser
```

The check then runs podman **as that user** (the job needs privilege escalation, which the templates in `SETUP_AAP.md` already have). If that user does not exist on the host you get the critical finding `mariadb:user`, with the hint `getent passwd <user>`. Leave the setting empty for containers that belong to root.

### 3. Where to put the settings

Open `playbooks/group_vars/mariadb.yml` (MySQL: `playbooks/group_vars/mysql.yml`) in VS Code. It already holds the lines below as comments. Remove the `#` in front of the ones you need and replace **CHANGE-ME**:

```
YAML

check_mariadb_enabled: true
check_mariadb_container: snow-mariadb
# check_mariadb_container_user: appuser      # only for rootless containers
```

A settings file applies to the group with the same name as the file. `mariadb.yml` is for the group `mariadb`, `mysql.yml` for `mysql`. If your hosts are in a group with another name (the example inventory uses `mariadb_hosts` and `mysql_hosts`), copy the file to `playbooks/group_vars/<your group name>.yml`. The check accepts all five group names, but the *settings* are only read from the file that matches the group. (The **Database health** template runs on the `_hosts` groups only: for a group called `mariadb` or `mysql`, give it a `target`, see section 1.) One host that differs from the rest: `playbooks/host_vars/<host name>.yml` (create it). Which place wins, and how to set the same thing in AAP's own Variables boxes, is in `VARIABLES.md`.

Then commit, push, and run **Sync** on the project in AAP.

### 4. The login: the "MariaDB monitor" credential

| Where the database runs        | What to do                                                                                                                                                                                               |
|--------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| MariaDB installed on the host  | Usually nothing. Run with privilege escalation and the check logs in as root through the local socket (RHEL's default).                                                                                  |
| MySQL installed on the host    | Usually the credential: MySQL's root account normally has a password.                                                                                                                                    |
| In a container (either engine) | Use the credential. The official images give root a password, and without it the job reports a finding that contains <b>Access denied</b> and ends with <b>attach the "MariaDB monitor" credential</b> . |

Steps:

1. Ask the DBA to run this **inside that database** (the account needs no data access). The privileges differ by engine:

```
SQL

CREATE USER 'aap_monitor'@'localhost' IDENTIFIED BY '<password>';

-- MySQL 5.7 / 8.0 / 8.4:
GRANT PROCESS, REPLICATION CLIENT ON *.* TO 'aap_monitor'@'localhost';
-- MySQL Group Replication only, also:
GRANT SELECT ON performance_schema.* TO 'aap_monitor'@'localhost';

-- MariaDB 10.5.9 and newer:
GRANT PROCESS, REPLICATION MONITOR ON *.* TO 'aap_monitor'@'localhost';
-- MariaDB before 10.5.9:
GRANT PROCESS, REPLICATION CLIENT ON *.* TO 'aap_monitor'@'localhost';
```

**On MariaDB 10.5.9 and newer, `REPLICATION CLIENT` is not enough:** it no longer covers the replication status, so the account needs **REPLICATION MONITOR**. **PROCESS** lets the check see other sessions (for long queries). Without the replication privilege the check still works and reports `mariadb:replication-query` (below) with

the grant to ask for. The host is `localhost` because the client runs inside the container.

2. In AAP create a credential of type **MariaDB monitor** (the type is in `aap/credential_types/mariadb_monitor.yml`; `SETUP_AAP.md` shows how to load it). Type the account name and password into the credential. **Never** put the password in a settings file, the inventory or Git: `SECRETS.md`.

3. Attach the credential to the **Database health**, **Health check** and **Troubleshoot** job templates.

Credentials attach to templates, not to groups; the check only uses it on hosts that have a database.

The password reaches the client only through its environment (`MYSQL_PWD`, and `podman exec -e MYSQL_PWD` for containers), never on a command line. If your client needs other options, for example a socket path or a TCP address, use `check_mariadb_client_args`: `['--socket=/var/lib/mysql/mysql.sock']`.

## 5. What it checks, and the limits you can change

Replication is read with the statement that fits the engine: `SHOW ALL SLAVES STATUS` (MariaDB), `SHOW REPLICA STATUS` (MySQL 8.0.22 and newer, 8.4) or `SHOW SLAVE STATUS` (older MySQL). The job's status line names the engine, for example `MySQL 8.4.11 - container snow-mysql (podman, image ...): up 312.4 h, 23/151 connections, ...`

Each limit is a setting. The defaults are in the table; change one the same way as in section 3.

| Finding                                                                     | Turns on when                                                                                       | Setting (default)                                                                                     |
|-----------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------|
| <code>mariadb:down</code>                                                   | the service or container is not running, or the named container does not exist                      | -                                                                                                     |
| <code>mariadb:connect</code>                                                | it runs but the status query failed (login, hang)                                                   | -                                                                                                     |
| <code>mariadb:restarted</code>                                              | restarted less than this many minutes ago                                                           | <code>check_mariadb_restart_warn_minutes</code> (60)                                                  |
| <code>mariadb:connections</code>                                            | connections in use, percent of <code>max_connections</code>                                         | <code>check_mariadb_conn_warn_pct</code> (80), <code>check_mariadb_conn_crit_pct</code> (95)          |
| <code>mariadb:refused</code>                                                | clients were refused because the limit was reached                                                  | -                                                                                                     |
| <code>mariadb:threads-running</code>                                        | queries executing at this moment                                                                    | <code>check_mariadb_threads_running_warn</code> (32)                                                  |
| <code>mariadb:long-queries</code>                                           | a query ran longer than this many seconds (never the query text)                                    | <code>check_mariadb_long_query_seconds</code> (300)                                                   |
| <code>mariadb:bufferpool</code>                                             | percent of reads served from memory (judged after 1 million reads)                                  | <code>check_mariadb_bufferpool_hit_warn</code> (95)                                                   |
| <code>mariadb:replica:&lt;name&gt;, mariadb:replica-lag:&lt;name&gt;</code> | replication stopped, or behind by seconds                                                           | <code>check_mariadb_replica_lag_warn</code> (300), <code>check_mariadb_replica_lag_crit</code> (1800) |
| <code>mariadb:replication-query</code>                                      | the replication status could not be read (usually a missing privilege; the finding names the grant) | section 4                                                                                             |

| Finding                                                           | Turns on when                                                                                                   | Setting (default)                                                                               |
|-------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------|
| <code>mariadb:galera</code> ,<br><code>mariadb:galera-size</code> | Galera / XtraDB Cluster node not healthy; cluster smaller than expected                                         | <code>check_mariadb_galera_size</code> (0 = not Galera; set 3 for three nodes)                  |
| <code>mariadb:group</code> ,<br><code>mariadb:group-size</code>   | MySQL Group Replication: a member is not ONLINE (or none is listed); fewer ONLINE members than expected         | <code>check_mariadb_group_size</code> (0 = do not check the size; set 3 for three members)      |
| <code>mariadb:group-query</code>                                  | the Group Replication members could not be read (needs <code>SELECT</code> on <code>performance_schema</code> ) | section 4                                                                                       |
| <code>mariadb:errorlog</code>                                     | <code>[ERROR]</code> lines in its log this recent (days)                                                        | <code>check_mariadb_error_since_days</code> (1)                                                 |
| <code>mariadb:disk</code>                                         | the disk holding the data directory is this full (percent)                                                      | <code>check_mariadb_disk_warn_pct</code> (85),<br><code>check_mariadb_disk_crit_pct</code> (95) |
| <code>mariadb:user</code>                                         | <code>check_mariadb_container_user</code> names a user that does not exist                                      | -                                                                                               |
| <code>mariadb:no-container</code>                                 | discovery found no container with a matching image                                                              | <code>check_mariadb_container_image_regex</code>                                                |

The finding ids keep the `mariadb:` prefix for MySQL too, so existing tickets and filters still match. What each finding means and the usual cause are in `RUNBOOKS.md`. Every setting with its default is in `VARIABLES_REFERENCE.md`. With several containers on one host, the finding id includes the container name, for example `mariadb:snow-mariadb:connections`.

**All settings** of this job, with their defaults: `check_mariadb`.

## How to verify

- The names are right.** On the database host: `sudo podman ps -a` (rootless: `sudo -u THE_USER podman ps -a`). The name you put in the settings file is in the `NAMES` column.
- The account works.** On the host, run the client inside the container as the monitoring account (it asks for the password; type it here, not in a file):

```
shell
sudo podman exec -it snow-mariadb mariadb --user=aap_monitor -p -e "SHOW GLOBAL STATUS LIKE 'Uptime'"
```

A number back means the login and the `PROCESS` rights work. (Some images have `mysql` instead of `mariadb`.)

- Run the job.** In AAP launch **Database health - MariaDB / MySQL** (first as *Check*, then as *Run*; `SETUP_AAP.md`), or the **Health check** with `health_checks = mysql` (or `mariadb`, or `database`) and a Limit of the database host. In the output, find the task `MariaDB | status`. It prints one line per instance, like `MariaDB 11.8.9 - container snow-mariadb (podman, image ...): up 312.4 h, 23/151 connections, 1 running, 0 long queries, replication connections 0`. That line proves the check reached the database. Each host then prints `HEALTHY` or its findings, and the job ends with `N healthy, M with findings`.

4. **No findings** means the database is healthy by these limits. To see a finding on purpose, set a limit that is sure to trip, for example `check_mariadb_conn_warn_pct: 1`, run again, look for `mariadb:connections`, then remove the line.

What was tested, and what was not

| Tested against real containers (throwaway rootless podman)                                                                                                                                                                                                                                                                                                | Tested only against a fake podman                                                                                                                                                                 |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| MySQL 8.0.46, MySQL 8.4.11, MariaDB 11.8.9: the engine and version are recognised; a replication source that never connects is reported as broken; the error log is read; a login with no credential gives <code>mariadb:connect</code> with the grant to ask for; the <code>REPLICA MONITOR</code> grant and revoke on MariaDB behave as described above | MySQL 5.7, and MySQL Group Replication ( <code>mariadb:group</code> , <code>group-size</code> , <code>group-query</code> ). They follow the documented statements but have not met a real server. |

The fake-podman test (`bash tests/mariadb/run_mariadb_test.sh`) covers one container, several, discovery, a stopped and a missing container, a rootless user that does not exist, MySQL 5.7, 8.0 and 8.4, replication, Group Replication, a denied grant, and that the password never appears on a command line. Nothing here has been run against **your** database: step 3 is the real proof.

## If it does not work

| You see                                                                                                                                        | Cause                                                                                                                        | Fix                                                                                                     |
|------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------|
| skipped: the MariaDB/MySQL check is off for <host> (Health check only)                                                                         | the host is in none of the five groups and no container is named                                                             | section 1, or use the <b>Database health</b> template                                                   |
| mariadb:replication-query                                                                                                                      | the account lacks the replication privilege (MariaDB 10.5.9+: <b>REPLICA MONITOR</b> , not <b>REPLICATION CLIENT</b> )       | section 4                                                                                               |
| mariadb:no-container                                                                                                                           | discovery matched nothing, or podman is run as the wrong user                                                                | podman ps -a as the right user; set check_mariadb_container_image_regex or check_mariadb_container_user |
| mariadb:user                                                                                                                                   | the rootless user name is wrong                                                                                              | getent passwd <user> on the host                                                                        |
| status query failed: ... Access denied ... and attach the "MariaDB monitor" credential                                                         | container root has a password                                                                                                | section 4                                                                                               |
| ... is not running (container ... is missing) but it is running                                                                                | the name is wrong, or the containers belong to another user                                                                  | compare with the podman ps -a output; check_mariadb_container_user                                      |
| The check runs on the wrong host, or not at all after a rename                                                                                 | a settings file only applies to the group named like the file                                                                | section 3                                                                                               |
| MariaDB is not running (service mariadb is inactive) - or a login error from the host's client - but podman ps shows the database container up | the host's mariadb.service only starts the container (podman compose up -d); before 0.7.0 the check took it for the database | 0.7.0 and later find the container by themselves; or name it: check_mariadb_container: mariadb          |
| Database health says skipping: no hosts matched (0.5.0), or target '...' matches no host                                                       | the server is in none of mariadb_hosts, mysql_hosts, database_hosts, mariadb, mysql (a Limit cannot add it)                  | section 1: target in the template's Variables, or add it to one of those groups                         |
| Your setting has no effect                                                                                                                     | another place wins                                                                                                           | VARIABLES.md                                                                                            |

## Appendix: credential types to paste into AAP

Each block below is one file in `aap/credential_types/`. In AAP: **Automation Execution > Infrastructure > Credential Types > Create credential type**. Paste the part under `inputs:` into **Input configuration** and the part under `injectors:` into **Injector configuration** (without those two top-level lines). Copy from the repository file where you can; PDF copy can change spaces.

### ServiceNow API (`servicenow_api.yml`)

Used by playbooks/`servicenow_tickets.yml`, `servicenow_health.yml` and `servicenow_test_ticket.yml` (REST Table API, basic auth only). The same variable names as the certified `servicenow.itsm` collection, so this credential also works if you move to that collection later.

YAML

```
inputs:
  fields:
    - id: instance
      type: string
      label: Instance URL
      help_text: e.g. https://servicenow.example.mil
    - id: username
      type: string
      label: API user name
    - id: password
      type: string
      label: API password
      secret: true
  required: [instance, username, password]
injectors:
  env:
    SN_HOST: "{{ instance }}"
    SN_USERNAME: "{{ username }}"
    SN_PASSWORD: "{{ password }}"
```

### MariaDB monitor (`mariadb_monitor.yml`)

For the mariadb health check. Needed when MariaDB runs in a container (the official images give root a password); optional on a host install, where the check can log in as root through the local socket. The password reaches the client only through its environment (`MYSQL_PWD`), never a command line. The account needs no data access at all (docs/MARIADB.md section 4 has the GRANT).

YAML

```
inputs:
  fields:
    - id: username
      type: string
      label: Monitor account
    - id: password
      type: string
      label: Password
      secret: true
  required: [username, password]
injectors:
  env:
    MARIADB_MONITOR_USER: "{{ username }}"
    MARIADB_MONITOR_PASSWORD: "{{ password }}"
```

## Keystore password (keystore\_password.yml)

Optional, for the certs check on a PKCS12 keystore (JKS keystores list without a password). In the inventory, name the variable, never the password: check\_certs\_keystores: [{path: /path/to/keystore, password\_env: KEYSTORE\_PASSWORD}]

```
YAML
inputs:
  fields:
    - id: password
      type: string
      label: Keystore password
      secret: true
      required: [password]
  injectors:
    env:
      KEYSTORE_PASSWORD: "{{ password }}"
```

## STIG Manager API (stigman\_api.yml)

Optional, for the STIG Manager cross-check in playbooks/poam\_status.yml. The client is created by your Keycloak realm owner (client-credentials grant, scope stig-manager:collection:read, and a grant in the collections to read) - this repository never changes the realm.

```
YAML
inputs:
  fields:
    - id: token_url
      type: string
      label: Token URL
      help_text: https://keycloak.example.mil/realms/<realm>/protocol/openid-connect/token
    - id: client_id
      type: string
      label: Client ID
    - id: client_secret
      type: string
      label: Client secret
      secret: true
      required: [token_url, client_id, client_secret]
  injectors:
    env:
      STIGMAN_TOKEN_URL: "{{ token_url }}"
      STIGMAN_CLIENT_ID: "{{ client_id }}"
      STIGMAN_CLIENT_SECRET: "{{ client_secret }}"
```

## ACT model key (act\_model\_key.yml)

YAML

```
inputs:
  fields:
    - id: genai_key
      type: string
      label: GenAI API key
      secret: true
    - id: asksage_key
      type: string
      label: AskSage API key
      secret: true
injectors:
  env:
    GENAI_KEY: "{{ genai_key }}"
    ASKSAGE_API_KEY: "{{ asksage_key }}"
```

## ACT GenAI beta key (act\_genai\_beta\_key.yml)

Only for ACT with site\_act\_provider=genai-beta (the GenAI.mil beta proxy). A separate type, so an "ACT model key" type that is already in use does not have to change (AAP refuses to change the fields of a credential type that credentials use). Attach it next to "ACT model key".

YAML

```
inputs:
  fields:
    - id: genai_beta_key
      type: string
      label: GenAI beta API key
      secret: true
      required: [genai_beta_key]
injectors:
  env:
    GENAI_BETA_KEY: "{{ genai_beta_key }}"
```

## STIG Manager database (stigman\_database.yml)

YAML

```
inputs:
  fields:
    - id: mysql_root_password
      type: string
      label: MySQL root password
      secret: true
      help_text: Used only when the database is first created (MYSQL_ROOT_PASSWORD_FILE).
    - id: stigman_db_password
      type: string
      label: STIG Manager database user password
      secret: true
      required: [mysql_root_password, stigman_db_password]
injectors:
  env:
    STIGMAN_MYSQL_ROOT_PASSWORD: "{{ mysql_root_password }}"
    STIGMAN_DB_PASSWORD: "{{ stigman_db_password }}"
```

## TLS certificate (tls\_certificate.yml)

AAP writes each value to a private temporary file inside the job's execution environment and puts the file's PATH in the environment variable; the playbook copies the files to the host.

```
YAML
inputs:
  fields:
    - id: certificate
      type: string
      label: Certificate (PEM; server certificate first, then intermediates)
      multiline: true
    - id: private_key
      type: string
      label: Private key (PEM)
      multiline: true
      secret: true
    - id: ca_chain
      type: string
      label: CA chain (PEM) - also used to trust your Keycloak/other internal services
      multiline: true
  required: [certificate, private_key]
injectors:
  file:
    template.cert: "{{ certificate }}"
    template.key: "{{ private_key }}"
    template.ca: "{{ ca_chain }}"
  env:
    TLS_CERT_FILE: "{{ tower.filename.cert }}"
    TLS_KEY_FILE: "{{ tower.filename.key }}"
    TLS_CA_FILE: "{{ tower.filename.ca }}"
```

## Container registry login (hosts) (registry\_login.yml)

For podman on the MANAGED HOSTS (the built-in "Container Registry" credential type is for pulling execution-environment images on AAP itself, not for the hosts).

```
YAML
inputs:
  fields:
    - id: registry
      type: string
      label: Registry (e.g. docker.io, quay.io)
    - id: username
      type: string
      label: Username
    - id: password
      type: string
      label: Password or access token
      secret: true
  required: [registry, username, password]
injectors:
  env:
    REGISTRY_HOST: "{{ registry }}"
    REGISTRY_USERNAME: "{{ username }}"
    REGISTRY_PASSWORD: "{{ password }}"
```

## Teams webhook (teams\_webhook.yml)

```
YAML
inputs:
  fields:
    - id: webhook_url
      type: string
      label: Teams Workflows webhook URL
      secret: true
      required: [webhook_url]
  injectors:
    env:
      TEAMS_WEBHOOK_URL: "{{ webhook_url }}"
```

## SMTP relay (smtp\_relay.yml)

Only when your mail relay needs a login to send. The jobs that email a report (report\_email\_to, roles/site\_email) read SMTP\_USERNAME / SMTP\_PASSWORD from it. A relay that accepts mail from the AAP server without a login needs no credential.

```
YAML
inputs:
  fields:
    - id: username
      type: string
      label: SMTP user name
    - id: password
      type: string
      label: SMTP password
      secret: true
      required: [username, password]
  injectors:
    env:
      SMTP_USERNAME: "{{ username }}"
      SMTP_PASSWORD: "{{ password }}"
```