

Settings Reference

Every setting of every job: its default, what it does, and where to set it

Site Automation - read it next to the setup guide | site-automation 0.12.1 | October 2026

What this is. Every setting the site-automation jobs read - the value it has when you set nothing, and what it does - with how to use one: where to put it (a template's Variables, `playbooks/group_vars/all.yml`, a group or a host), which place wins, and how to see what a job used.

How to read it

Keep it open next to the *Operations Runbooks Setup Guide*: each job's chapter there says which settings matter for it; this one lists them all, role by role.

Rendered from `docs/VARIABLES.md` and `docs/VARIABLES_REFERENCE.md` (generated from `roles/*/defaults/main.yml`) in the site-automation repository (version 0.12.1).

Contents

1. Settings: what they are, where they go, and who wins	4
The two-minute version	4
How to use a setting	4
1. What you can set	5
2. Where do I put a setting?	7
3. Who wins when the same setting is in two places?	8
Examples	10
How to verify a setting is in effect	10
If a setting does nothing	11
2. Every setting and its default	12
check_accounts	12
check_auditd	12
check_certs	13
check_containers	15
check_disk	16
check_fapolicyd	17
check_logging	17
check_mariadb	18
check_mounts	20
check_network	21
check_patching	22
check_performance	22
check_selinux	23
check_services	23
check_time	24
ontap_report	24
patch	26
poam	27
podman_discover	28
service_watch	29
servicenow	30
site_act	32

site_email	33
site_findings	34
stigman_stack	35
troubleshoot	36
vmware_vm	37
win_act	47
win_check_accounts	47
win_check_audit	48
win_check_certs	49
win_check_disk	49
win_check_eventlog	50
win_check_network	50
win_check_patching	51
win_check_performance	51
win_check_security	52
win_check_services	52
win_check_time	53
win_patch	53
win_troubleshoot	55

1. Settings: what they are, where they go, and who wins

A **variable** (also called a *setting*) is a named value the runbooks read: `check_disk_warn_pct: 85` means "warn when a disk is 85 % full". You change a variable to change how a check behaves, without touching any code. This page answers three questions:

1. **What can I set?** (find the name and the default)
2. **Where do I put it?** (for every host, for one group, for one host, for one run)
3. **What happens when the same setting is in two places?** (who wins)

You do not need to read it all. Do the two-minute version first.

The two-minute version

1. **Find the setting.** Open `VARIABLES_REFERENCE.md`, press Ctrl+F and search for a word (`disk`, `certificate`, `mariadb`). Every setting is there with its default.
2. **Put the line in one file.**
 - Every host: `playbooks/group_vars/all.yml`
 - Hosts in one AAP group: `playbooks/group_vars/<group name>.yml`
 - One host: `playbooks/host_vars/<host name>.yml`
3. **Commit and Sync Changes** in VS Code. Run the check.

That is all for a normal change. The rest of this page explains the details and the exceptions.

How to use a setting

Every job's chapter in these docs ends its settings part with an **All settings** link: the list of that job's settings in `VARIABLES_REFERENCE.md`, each with its default and what it does. Pick the place by **who the setting is for**:

The setting is about	Put it in	Example
hosts (thresholds, what a check looks at)	<code>playbooks/group_vars/all.yml</code> (every host), <code>playbooks/group_vars/<group>.yml</code> , <code>playbooks/host_vars/<host>.yml</code>	<code>check_disk_warn_pct: 90</code>
one job template (who gets its email, its rules)	the job template's Variables box in AAP	<code>report_email_to:</code> <code>ops@yoursite.mil</code> on the certificate report template
one run	a survey question (its variable name = the setting's name) or Prompt on launch	a survey question <code>health_checks</code> , multiple choice <code>daily / weekly</code>

Three examples:

```

1. playbooks/group_vars/all.yml - every host: disks warn at 90 %, a stricter audit log partition

check_disk_warn_pct: 90
check_disk_overrides:
  /var/log/audit: {warn: 70, crit: 85}

# 2. the "Certificate report" template's Variables - this job only: email it, stay green
report_email_to: [pki-team@yoursite.mil]
site_fail_on: []

# 3. playbooks/group_vars/win_web.yml - the Windows web servers only: IIS must run too
win_check_services_required_extra: [W3SVC, WAS]

```

- **A list replaces the default list** (it does not add to it): repeat the entries you keep. Details: "Three shapes of value" below.
- **Check what a job used:** in AAP, the job's **Details** tab shows the template's and the launch's variables; the job's output prints the decisions it made (thresholds in the finding texts, **REFUSED**: lines, the settings each report names in its email); its **artifacts** (Details > Artifacts) carry the results. A setting that seems ignored: If a setting does nothing.

1. What you can set

Every setting lives, with its default and an explanation, in the role that uses it:

`roles/<role>/defaults/main.yml`. Two easy ways to read them:

- **VARIABLES_REFERENCE.md**, the same information as one searchable page (generated from the role files, so it is never out of date). Start here.
- In VS Code: **Ctrl+Shift+F** (search in all files), type the setting or a word, and look at the hits under `roles\...\defaults\main.yml`.

Setting names start with the thing they belong to:

Starts with	Belongs to	Example
check_disk_ check_services_ check_mariadb_...	one health check (<code>roles/check_<name></code>)	<code>check_disk_warn_pct</code>
win_check_	one Windows health check	<code>win_check_disk_warn_pct</code>
site_act_	ACT (GenAI)	<code>site_act_level</code>
site_ (other)	reporting and settings	<code>site_fail_on</code>
patch_, win_patch_, automatic_restarts	patching	<code>patch_never_patch_groups</code>
servicenow_	ServiceNow tickets	<code>servicenow_assignment_group</code>
watch_	service watch	<code>watch_containers, watch_discover</code>
podman_discover_	finding the podman containers (the <code>containers</code> check and service watch share it)	<code>podman_discover_ignore</code>
poam_	POA&M	<code>poam_stigman_api</code>
stigman_	STIG Manager deployment	<code>stigman_scope</code>

Check names and what each finding means: RUNBOOKS.md.

Three shapes of value

YAML

```
check_disk_warn_pct: 85           # a number
check_fapolicyd_required: true    # true or false (lower case)
check_services_required: [sshd, crond] # a list: [a, b] or one "- item" per line
check_disk_overrides:             # a set of names with values (a "map")
  /var/log/audit: {warn: 70, crit: 85}
```

A list replaces the default list. It does not add to it. If the default is `[sshd, crond, rsyslog]` and you write `[chronyd]`, the hosts now require **only** `chronyd`. Copy the default list from the reference and add your entries to it. (Some settings have a separate `..._extra` version that *adds*, for example `check_services_required_extra`. The reference says which.)

Job variables (per run, not per host)

These choose what a run does. Put them in a job template's **Variables** box (*Extra variables*) or in a survey question, never in a settings file:

Variable	Meaning	Where it appears
<code>target</code>	which hosts or group to run on (default: all)	most playbooks
<code>health_checks</code>	which checks: <code>daily</code> , <code>weekly</code> , <code>all</code> , or a list	Health check
<code>ts_area</code> , <code>ts_service</code> , ...	what to troubleshoot	Troubleshoot
<code>use_act</code> , <code>site_act_level</code>	ask ACT, and what it may do	Health check, Troubleshoot
<code>site_fail_on</code>	which findings fail the job: <code>[critical]</code> or <code>[critical, warning]</code>	any check job

2. Where do I put a setting?

Place	Who gets it	How you change it	Good for
<code>roles/<role>/defaults/main.yml</code>	everyone	do not edit (updates replace it)	the built-in default
<code>playbooks/group_vars/all.yml</code>	every host	VS Code, commit, sync	site-wide values: the LDAP server, the SIEM, thresholds
<code>playbooks/group_vars/<group>.yml</code>	the hosts of that AAP group	VS Code	one kind of server: STIG Manager containers, MariaDB, AAP itself
<code>playbooks/host_vars/<host>.yml</code>	that one host	VS Code (create the folder and file)	one odd server: a bigger disk limit for <code>/var/log/audit</code>
AAP: inventory > Variables	every host	AAP web UI	values you prefer to keep in AAP
AAP: group > Variables	the hosts of that group	AAP web UI	same, per group
AAP: host > Variables	that host	AAP web UI	same, per host
Job template > Variables (<i>Extra variables</i>)	every run of that template	AAP web UI	"this template always uses <code>site_act_level: diagnose</code> "
Survey question / prompt on launch	this launch only	asked at launch	a choice made each time
AAP credential	the job	AAP web UI	secrets only (SECRETS.md)

File names matter. `playbooks/group_vars/<group>.yml` applies to the group with **exactly that name** in your AAP inventory, and `playbooks/host_vars/<host>.yml` to the host with exactly that name. A file for a group your inventory does not have is silently never used.

The files shipped in `playbooks/group_vars/` and the group each expects:

File	Group name it applies to	If your group has another name
<code>all.yml</code>	every host	(always works)
<code>aap.yml</code>	<code>aap</code>	copy it to <code>aap_hosts.yml</code> if your group is <code>aap_hosts</code> (the example inventory's name)
<code>mariadb.yml</code>	<code>mariadb</code>	the MariaDB check also runs for a group named <code>mariadb_hosts</code> , but the settings come from the file named after the group: copy it to <code>mariadb_hosts.yml</code>
<code>sat.yml</code> , <code>idm.yml</code> , <code>logstash.yml</code> , <code>stigman.yml</code> , <code>rhel8_all.yml</code>	the same name	copy or rename
<code>netapp.yml</code>	<code>netapp</code>	copy to <code>netapp_console_hosts.yml</code> if that is your group

(The built-in safety lists cover both spellings of the AAP and NetApp group names, so AAP stays protected whatever you call it.)

Updating the repository never overwrites your settings files. The update script (`scripts/update-from-release.ps1 / .sh`) protects `playbooks/group_vars/`, `playbooks/host_vars/`, `inventories/site/`, `poam/poam.csv` and anything in `.site-local`. If a new release adds a settings file you do not have, the script adds it **once** (the preview lists it as **YOURS**) and never changes it again. A setting that the release's `all.yml` mentions and yours does not is listed under **NEW SETTINGS**, so you can copy it over if you want it.

Adding "extras" to a host or a group in AAP

Where the boxes are in AAP 2.7 (older versions word a label differently):

- **Inventory: Automation Execution > Infrastructure > Inventories** > your inventory > **Edit inventory > Variables**.
- **Group:** the same inventory > **Groups** tab > the group > **Edit group > Variables**.
- **Host:** the same inventory > **Hosts** tab > the host > **Edit host > Variables**.
- **Job template: Automation Execution > Templates** > the template > **Edit template > Extra variables** (the template's *Variables* box).

The **Variables** box on a host or group takes the same YAML as a settings file. Typical entries:

a group's Variables box: hosts in this group need a bigger disk margin

```
check_disk_overrides:
  /var/lib/pgsql: {warn: 70, crit: 85}
```

a host's Variables box: its MariaDB containers belong to a rootless account

```
check_mariadb_container_user: svc_podman
```

a group's Variables box: never check or watch the test containers, or anything of one user

```
podman_discover_ignore: ['test-.*', 'devuser/.*']
```

Connection settings (`ansible_host`, `ansible_user`, `ansible_python_interpreter`, `ansible_connection: winrm`) belong here too. The *login* and *password* never do: they are credentials (SECRETS.md).

Prefer the project's settings files (this repository), because they are reviewed and versioned in Git and updates never overwrite them. Use the AAP boxes when AAP is where your team already manages the inventory. Either is fine; the rule is **one setting, one place**.

3. Who wins when the same setting is in two places?

Strongest first. This order was **proven by running a test playbook** on `ansible-core` 2.16 and 2.21 with the same value set in every place and removing the winner one place at a time (the same result on both versions).

#	Place	Notes
1	Extra variables: job template <i>Variables</i> , a survey answer, a prompt on launch, -e on the command line	beats everything below, including the built-in protection further down
2	<code>playbooks/host_vars/<host>.yml</code>	this project, one host

#	Place	Notes
3	the host's Variables box in AAP	
4	<code>playbooks/group_vars/<group>.yaml</code>	a child group beats its parent group
5	<code>playbooks/group_vars/all.yaml</code>	
6	the group's Variables box in AAP	a deeper group beats its parent
7	the inventory's Variables box in AAP	
8	<code>roles/<role>/defaults/main.yaml</code>	the built-in default

Things that surprise people:

- **A file in this project beats the AAP group and inventory boxes**, even when the box is on a more specific group. If you type a value in an AAP box and see no effect, look for the same setting in `playbooks/group_vars/`. Every job prints a notice, `NOTE - settings in this project's files win over the same setting in AAP's inventory Variables box`, listing which files define which settings. (Silence it with `site_settings_notice: false`.) Not shown for `playbooks/host_vars/`.
- **A host's own Variables box (3) beats a group file (4/5)**, but a `playbooks/host_vars` file (2) beats the host's box.
- **A survey answer or template extra variable beats everything**: it is the one place that can override the rest for a single run.
- Two sibling groups (neither is inside the other) that set the same variable for the same host: do not rely on the result. Put it in a host file, or in one group.
- An inventory **sourced from the project** (`inventories/site/`, SETUP_AAP.md step 5): AAP imports its `group_vars/all.yaml` into the inventory's Variables box and each `group_vars/<group>.yaml` into that group's box. So they rank as places 7 and 6, and the project's `playbooks/` files beat them. Keep a setting in one of the two, not both.
- **Updated from 0.4.1 or older?** Your `playbooks/group_vars/all.yaml` is yours, so the update left it as it was: it still sets `site_act_diagnose_only_groups`, `patch_never_patch_groups`, `patch_never_reboot_groups`, `site_act_provider` and `site_act_models` as real (uncommented) lines. They beat the same settings in AAP's inventory and group boxes. If you want those set in AAP (or the new built-in defaults), put `#` in front of those lines. See the upgrade notes in `CHANGELOG.md` (0.5.0).

Built-in protection: settings that cannot be weakened from a file

Some lists are protected in code, so a mistake in a settings file cannot expose the AAP host:

What	Built in	What your file can do
Never patched, never rebooted	<code>aap, aap_hosts</code> (<code>roles/patch/vars/main.yaml</code>)	your <code>patch_never_patch_groups</code> / <code>patch_never_reboot_groups</code> list replaces the shipped one (<code>sat, netapp, ...</code>): repeat the names you keep. <code>aap</code> and <code>aap_hosts</code> stay protected whatever you write
ACT never fixes by itself	<code>aap, aap_hosts</code> (<code>roles/site_act/vars/main.yaml</code>)	your <code>site_act_diagnose_only_groups</code> list replaces the shipped one (<code>sat, idm, netapp, ...</code>): repeat the names you keep. <code>aap</code> and <code>aap_hosts</code> stay protected whatever you write

Only an extra variable (place 1) can override these, and nothing in these docs asks you to do that.

Examples

Everyone: warn earlier on disks. `playbooks/group_vars/all.yml`:

```
YAML

check_disk_warn_pct: 80
check_disk_crit_pct: 92
```

One group: the STIG Manager hosts need these containers watched.

`playbooks/group_vars/stigman.yml` (shipped; already has `watch_containers`).

One host: a database server whose audit partition must stay emptier.

`playbooks/host_vars/db01.example.mil.yml` (the file name is the host's name in AAP):

```
YAML

check_disk_overrides:
  /var/log/audit: {warn: 70, crit: 85}
```

One template: the "Troubleshoot" template always uses Ask Sage. In the template's **Variables**:

```
YAML

site_act_provider: asksage
```

A list you want to extend. Where the role has a `..._extra` setting, use it: it *adds*.

```
YAML

check_services_required_extra: [tomcat]           # required on top of the built-in five
```

Where it has none, look up the default (`check_services_required` is `[sshd, crond, rsyslog, chronyd, auditd]`) and write the whole list with your addition:

```
YAML

check_services_required: [sshd, crond, rsyslog, chronyd, auditd, tomcat]
```

How to verify a setting is in effect

1. **Run the check on one host: Health check** with `target` = that host and `health_checks` = the check. The finding text quotes the limit it used (`/var is 96% full (critical at 92%)`), so you see the value that won.
2. **Break it on purpose** (or lower the limit) on a test host, and confirm the finding appears. A check that never fires proves nothing.
3. **Read the job's notice** (NOTE - `settings in this project's files ...`): it names every settings file that set something, so you can spot a file that overrides an AAP box.
4. **From the command line** (optional, in a terminal with Ansible and the repository): create a temporary file `playbooks/show_settings.yml`, **do not commit it**, and run it:

YAML

```
- name: Show the settings a host gets
hosts: "{{ target | default('all') }}"
gather_facts: false
tasks:
  - ansible.builtin.import_role: {name: check_disk}    # the role that owns the setting
    when: false                                         # loads its defaults, runs nothing
  - ansible.builtin.debug:
    msg: "{{ item }}" = "{{ lookup('vars', item, default='(not set)') }}"
    loop: [check_disk_warn_pct, check_disk_crit_pct]
```

shell

```
ansible-playbook -i inventories/site/hosts.yml playbooks/show_settings.yml -e
target=db01.example.mil
```

It shows the value from the role default plus the files in this project. (Values typed into AAP's Variables boxes exist only inside AAP; there, use steps 1-3.)

If a setting does nothing

Symptom	Likely cause
You set it in AAP and nothing changes	the same setting is in <code>playbooks/group_vars/</code> : that file wins (table above)
You set it in a file and nothing changes	the file name does not match the group or host name exactly; a spelling or capital letter
A list lost entries	a list replaces the default: repeat the entries you keep
YAML error at the start of the job	indentation (spaces, never tabs), a missing space after <code>:</code> , or a value with <code>:</code> inside it that needs quotes
A survey answer is ignored	the survey question's variable name is not exactly the setting's name

2. Every setting and its default

This page lists **every setting a role reads**, with the value it has when you set nothing, and what the role's own comment says. It is generated from `roles/<role>/defaults/main.yml`, so it is never out of date. Where to put a setting, and which place wins when it is in two, is in `VARIABLES.md`. Values are shortened to one line here: the role file has the full one.

A setting that is a list (`[a, b]`) **replaces** the default list when you set it; it does not add to it. Repeat the entries you want to keep.

Contents: `check_accounts`, `check_auditd`, `check_certs`, `check_containers`, `check_disk`, `check_fapolicyd`, `check_logging`, `check_mariadb`, `check_mounts`, `check_network`, `check_patching`, `check_performance`, `check_selinux`, `check_services`, `check_time`, `ontap_report`, `patch`, `poam`, `podman_discover`, `service_watch`, `servicenow`, `site_act`, `site_email`, `site_findings`, `stigman_stack`, `troubleshoot`, `vmware_vm`, `win_act`, `win_check_accounts`, `win_check_audit`, `win_check_certs`, `win_check_disk`, `win_check_eventlog`, `win_check_network`, `win_check_patching`, `win_check_performance`, `win_check_security`, `win_check_services`, `win_check_time`, `win_patch`, `win_troubleshoot`

check_accounts

File: `roles/check_accounts/defaults/main.yml`

Note

ACCOUNTS: the account problems auditors and attackers both look for, plus the one that quietly breaks automation - a service account password about to expire.

Setting	Default	What it does
<code>check_accounts_min_uid</code>	<code>1000</code>	people's accounts start here
<code>check_accounts_inactive_days</code>	<code>35</code>	the STIG wants unused accounts disabled after 35 days
<code>check_accounts_ignore_inactive</code>	<code>[]</code>	Accounts that may stay unused (break-glass, service accounts that never log in interactively)
<code>check_accounts_nopasswd_allowed</code>	<code>[]</code>	sudoers entries allowed to use NOPASSWD, e.g. the account AAP connects as: <code>[svc_aap, "%wheel"]</code>
<code>check_accounts_expiry_watch</code>	<code>[]</code>	Accounts whose password/account expiry matters - above all the account AAP logs in with. e.g. <code>[svc_aap]</code>
<code>check_accounts_expiry_warn_days</code>	<code>14</code>	a watched account's password or account expiring within this many days -> warning (3 days or less -> critical)

check_auditd

File: `roles/check_auditd/defaults/main.yml`

Note

AUDITD: the audit system records who did what (the STIG depends on it). Running, enabled in the kernel, rules loaded, no events lost, room for its log, and (optionally) forwarding. Also how MUCH it records: too much rotates the evidence away within hours, fills the disk (then auditd stops the host or puts it in single-user mode, as the STIG sets it up), drops records on the way to the SIEM, and uses up the SIEM licence.

Setting	Default	What it does
<code>check_auditd_min_rules</code>	5	fewer loaded rules than this = the STIG rules are missing
<code>check_auditd_log_warn_pct</code>	75	the audit log's filesystem
<code>check_auditd_log_crit_pct</code>	90	% used of that filesystem -> critical (warning from <code>check_auditd_log_warn_pct</code>)
<code>check_auditd_remote_required</code>	false	true = au-remote (audisp-remote) must be active
<code>check_auditd_exclude_accounts</code>	[]	Accounts whose audit events are expected in bulk and do NOT count toward the volume warning, e.g. the account a vulnerability scanner logs in with (a credentialed scan runs thousands of commands with sudo). Names or UIDs; a host that does not have the account skips it. The report still says how much they made, and the... (all of it: the role file)
<code>check_auditd_volume_hours</code>	24	look at the audit records of the last N hours
<code>check_auditd_volume_warn_mb_day</code>	200	warning above this many MB of audit records a day (excluded accounts not counted)
<code>check_auditd_retention_warn_hours</code>	24	With <code>max_log_file_action</code> ROTATE, auditd keeps <code>num_logs</code> files of <code>max_log_file</code> MB and deletes the oldest. Warning when what is left on disk reaches back fewer hours than this. The STIG wants a week (168) when the records are NOT sent off the host as they are written (RHEL-08-030660). 0 = off.
<code>check_auditd_volume_timeout</code>	120	The breakdown (which account, which rule, which program) reads the audit log files of the window at low CPU and disk priority, and only when the volume or the retention is over its limit. It stops after this many seconds on a host with very large logs (and says so).

check_certs

File: `roles/check_certs/defaults/main.yml`

Note
CERTS: certificates that expire soon - as files on the host, inside Java keystores, and as served by TLS endpoints (what a browser or client actually sees).

What it does

Certificate expiring within this many days -> warning (and amber in the certificate report)		
The certificate report colours these blue (an early heads-up; not a finding)		
Within this many days (or expired) -> critical		
Check certificate files (PEM) to check. Globs work. Every certificate in a file is checked (a chain file with an expiring intermediate is a problem too). e.g. [etc/stigman/*.crt, /etc/pki/tls/certs/server.crt]		
Also look for certificates in the usual places. Files with many certificates are CA bundles (hundreds of certs, noisy) and are skipped unless listed in check_certs_files.		
The folders searched (two levels deep) for *.crt, *.pem, *.cer		
certificates\	/localhost\.crt\$"	regular expression: found files whose p are skipped
Found file with more certificates than this is a CA bundle: skipped		
TLS endpoints, checked FROM this host, e.g. the host's own web server: check_certs_endpoints: - name: STIG Manager, host: 127.0.0.1, port: 443, servername: stigman.example.mil}		
Java keystores (Tomcat, the ServiceNow MID Server etc.). A JKS keystore lists without a password; a PKCS12 needs it - give the NAME of an environment variable that holds it (set by an AAP credential), or the password itself: check_certs_keystores: - path: /opt/servicenow/mid/agent/security/agent_keystore, passw... (all of it: the role file)		

What it does		
- Finding them by themselves (the certificate port turns both on; the health check does not) ---- - Listening TCP ports that answer TLS: each probed once, on the address it listens on, with <code>openssl client</code> (check_certs_probe_timeout seconds each, at most check_certs_discover_ports_max ports). Not these ports: pl... (all of it: the role file)		
Ports never probed (plain-text services: ssh, smtp, https, http, ldap, databases ...)		
Probe at most this many listening ports (the rest are listed as not probed)		
Seconds for each port probe		
Java keystores (*.jks, *.keystore, *.p12, *.pfx) in these folders, and the ones running Java programs (e.g. -Djavax.net.ssl.keyStore=...). A keystore that needs a password is listed as not checked: no password is ever guessed - give it in check_certs_keystores (password_env). The JDK's own trust stores (cacerts) are sk... (all of it: the role file)		
The folders searched for keystores		
How many folder levels deep the search goes		
search_cacerts\$\\	/jre/lib/security/	/lib/security/
Find at most this many found keystores per host		

check_containers

File: `roles/check_containers/defaults/main.yml`

Note
CONTAINERS: the podman containers on the host - root's (rootful) and every user's (rootless) - found by themselves (roles/podman_discover: which ones exist, which systemd unit runs each one, which should be running). Read-only. Which containers and users are looked at, and which are left out: the podman_discover_* settings.

Setting	Default	What it does
check_containers_required	[]	Containers that MUST exist and run on these hosts: a container name, or OWNER/NAME for one owner (root/web, alice/web). A name alone matches that container of any owner. Also matches a Quadlet file name without .container, e.g. [stigman-mysql, stigman-api, alice/web].
check_containers_recent_hours	24	A container that stopped with an error (exit code not 0) this many hours ago or less is reported (containers:crashed). Older crashes are history and are not reported.
check_containers_restart_warn	3	More restarts than this (podman's restart count, or its systemd unit's) = a restart loop.

check_disk

File: `roles/check_disk/defaults/main.yml`

Note

DISK: how full each LOCAL filesystem is (space and inodes), and - from the history of earlier runs - how many days until it is full. Network filesystems (NFS) are the mounts check's job: a hung NFS server can make `df` hang, so this check never touches them.

What it does				
% used -> warning				
% used -> critical				
inodes: too many small files fill a disk too				
% of inodes used -> critical (warning from check_disk_inode_warn_pct)				
Different limits for particular mount points, e.g. keep the audit log partition emptier: check_disk_overrides: /var/log/audit: {warn: 70, crit: 85}				
sys\	dev\	run)(\	\$)' - '^/etc/(hosts\	hostname\
filesystem types never checked (df -x)				

What it does

Days-until-full forecast. Each run adds a sample (date, bytes used) per mount to a small file on the host; once there is at least a day of history, the growth rate predicts when the filesystem fills. Check mode (a dry run) reads the history but does not add to it.

full within this many days at the current growth -> warning

full within this many days -> critical

where the forecast keeps its samples, on the host

samples older than this are dropped (the forecast uses what is left)

seconds for df

check_fapolicyd

File: `roles/check_fapolicyd/defaults/main.yml`

Note

FAPOLICYD (application allow-listing): running, enforcing, configuration and trust database valid, and recent denials - which is what you see when it blocks a program.

Setting	Default	What it does
<code>check_fapolicyd_required</code>	<code>false</code>	true = every host in the play must have fapolicyd. false = only check hosts that have it.
<code>check_fapolicyd_expect_enforcing</code>	<code>true</code>	permissive = 1 means it only logs, never blocks
<code>check_fapolicyd_check_trustdb</code>	<code>true</code>	Compare the trust database with the files on disk (can take a minute on big hosts; files changed outside RPM show up here, e.g. after a manual update).
<code>check_fapolicyd_denials_since</code>	<code>yesterday</code>	ausearch -ts value
<code>check_fapolicyd_denials_warn</code>	<code>1</code>	this many denials or more since check_fapolicyd_denials_since -> warning

check_logging

File: `roles/check_logging/defaults/main.yml`

Note

LOGGING: are logs kept, and do they reach the central log server (SIEM)?

Setting	Default	What it does
check_logging_journal_persistent	true	the journal must survive a reboot (/var/log/journal)
check_logging_remote	[]	The central log server(s) this host forwards to, e.g. check_logging_remote: - {name: SIEM, host: siem.example.mil, port: 6514}
check_logging_timeout	5	seconds to wait for each log server's TCP port
check_logging_queue_warn_mb	100	rsyslog keeps messages it cannot deliver in /var/lib/rsyslog. A big queue = the log server has been unreachable for a while.
check_logging_errors_since	" -24h"	Error-priority journal messages in the last day. Many = something is complaining a lot.
check_logging_errors_warn	500	this many error messages or more -> warning

check_mariadb

File: roles/check_mariadb/defaults/main.yml

Note

DATABASE health for MariaDB AND MySQL (read-only; role name kept for compatibility): running, accepting connections, connection headroom, queries that run too long, replication (MariaDB multi-source, MySQL REPLICA/SLAVE), Galera / Percona XtraDB Cluster, MySQL Group Replication, recent errors in its log, and room on its disk. It runs status queries only - no data is read and nothing is changed. The engine and version are detected; see docs/MARIADB.md.

	Default	What it does
oled	(worked out by the role; see its file)	Which hosts: members of any of the inventory groups mariadb_hosts, mysql_hosts, database_hosts, mariadb or mysql, or any host that names a container below (or turns discovery on). Set this true/false to override for a host or group.

	Default	What it does
container	(empty)	MariaDB or MySQL in podman containers (e.g. the ServiceNow database). The client inside each container is used (podman exec). Name them in one of these ways; all are checked, each as its own instance: check_mariadb_container: mariadb one container check_mariadb_containers: [mariadb, snow-db] several containers on the... (all of it: the role file)
containers	[]	several containers on one host, e.g. [mariadb, snow-db] (see above)
container_discover	auto	auto, true or false: find database containers by their image (see above)
container_paths	[/usr/sbin/mariabdd, /usr/libexec/mariabdd, /usr/sbin/mysqld, /usr/libexec/mysqld, /usr/b...	Where a database server installed on the host lives (any one present = a host install, for auto).
container_image_regex	``(mariadb\	mysql\percona)``
container_user	(empty)	ROOTLESS podman: the containers belong to a user (e.g. AAP's own MariaDB, or an application account), and only that user can see them. Name the user and the check runs podman AS that user. Empty = root's containers.
container_service	(empty)	The systemd unit (host installs). Empty = the first of mariadb, mysqld, mysql that exists.

	Default	What it does
ent_args	[]	How to log in. Without a credential: as root through the local socket - MariaDB's default for a host install on RHEL (no password needed when the job runs with sudo). A database in a CONTAINER usually gives root a password (the official images do), and MySQL host installs usually do too, so there attach the "MariaDB m... (all of it: the role file)
n_warn_pct	80	connections in use, % of max_connections
n_crit_pct	95	connections in use, % of max_connections -> critical
eads_running_warn	32	queries executing at this moment
g_query_seconds	300	a query running longer than this is "long"
lica_lag_warn	300	seconds a replica is behind its primary
lica_lag_crit	1800	seconds behind its primary -> critical
era_size	0	expected Galera / XtraDB Cluster size (0 = do not check the size)
up_size	0	expected MySQL Group Replication members (0 = do not check the size)
tart_warn_minutes	60	restarted less than this ago -> warning
ferpool_hit_warn	95	% of reads served from memory (only judged after 1M reads)
or_since_days	1	[ERROR] lines in its log this recent
k_warn_pct	85	the filesystem holding the data directory
k_crit_pct	95	% used of the data directory's filesystem -> critical

check_mounts

File: roles/check_mounts/defaults/main.yml

Note

MOUNTS: filesystems that should be mounted but are not, local filesystems that went read-only (the kernel does that after disk errors), and network mounts (NFS/CIFS) that hang. Diagnose only: this check never mounts, unmounts or remounts anything.

Setting	Default	What it does
<code>check_mounts_network_types</code>	<code>[nfs, nfs4, cifs, smb3]</code>	filesystem types counted as network mounts (checked for hangs and fullness)
<code>check_mounts_local_types</code>	<code>[xfs, ext4, ext3, btrfs, vfat]</code>	filesystem types checked for having gone read-only
<code>check_mounts_timeout</code>	5	seconds before a network mount counts as hung
<code>check_mounts_warn_pct</code>	85	Network mounts: how full (same idea as the disk check)
<code>check_mounts_crit_pct</code>	95	% used of a network mount -> critical (warning from <code>check_mounts_warn_pct</code>)
<code>check_mounts_ignore</code>	<code>[]</code>	fstab mount points that are allowed to be unmounted (regular expressions)

check_network

File: `roles/check_network/defaults/main.yml`

Note

NETWORK: can this host reach what it depends on? A default route, names that must resolve, and TCP ports that must answer (from THIS host, which is what matters when a service breaks).

Setting	Default	What it does
<code>check_network_require_default_route</code>	<code>true</code>	no default route -> finding (false for an isolated host)
<code>check_network_resolve_self</code>	<code>true</code>	this host's own full name (hostname -f) resolves
<code>check_network_dns_names</code>	<code>[]</code>	e.g. <code>[ldap.example.mil, servicenow.example.mil]</code>
<code>check_network_tcp</code>	<code>[]</code>	TCP services this host needs, e.g. <code>check_network_tcp: - {name: LDAP, host: ldap.example.mil, port: 636} - {name: SIEM (syslog over TLS), host: siem.example.mil, port: 6514}</code>
<code>check_network_timeout</code>	5	seconds for each name lookup and TCP port test

check_patching

File: `roles/check_patching/defaults/main.yml`

Note

PATCHING (readiness, read-only): pending security updates, a reboot still needed after the last updates, and how long ago the host was last updated. Installs nothing - patching itself is playbooks/patch_hosts.yml.

Setting	Default	What it does
<code>check_patching_max_age_days</code>	<code>35</code>	last package update older than this -> warning
<code>check_patching_security_warn</code>	<code>1</code>	this many pending security advisories -> warning
<code>check_patching_critical_is_critical</code>	<code>true</code>	a pending Critical-severity advisory -> critical
<code>check_patching_timeout</code>	<code>300</code>	seconds for dnf to talk to the repositories
<code>check_patching_reboot_check</code>	<code>true</code>	report a reboot still needed (needs-restarting -r, newer kernel installed than running)

check_performance

File: `roles/check_performance/defaults/main.yml`

Note

PERFORMANCE: is the host overloaded right now, or was something killed for lack of memory?

Setting	Default	What it does
<code>check_performance_load_warn</code>	<code>1.5</code>	15-minute load average PER CPU (1.0 = every CPU busy)
<code>check_performance_load_crit</code>	<code>3.0</code>	load per CPU -> critical
<code>check_performance_mem_avail_warn_pct</code>	<code>10</code>	memory still available for programs, % of total
<code>check_performance_mem_avail_crit_pct</code>	<code>5</code>	available memory below this % -> critical
<code>check_performance_swap_warn_pct</code>	<code>50</code>	swap in use, % of swap
<code>check_performance_swap_crit_pct</code>	<code>80</code>	swap in use above this % -> critical
<code>check_performance_iowait_warn</code>	<code>20</code>	% of CPU time waiting for disk (average over the sample)
<code>check_performance_iowait_crit</code>	<code>40</code>	I/O wait above this % -> critical

Setting	Default	What it does
<code>check_performance_steal_warn</code>	<code>10</code>	% of CPU taken by the hypervisor (VM starved)
<code>check_performance_sample_seconds</code>	<code>5</code>	how long vmstat samples
<code>check_performance_oom_since</code>	<code>"-24h"</code>	look for out-of-memory kills this far back
<code>check_performance_zombie_warn</code>	<code>20</code>	this many zombie processes or more -> warning

check_selinux

File: `roles/check_selinux/defaults/main.yml`

Note

SELINUX: enforcing (the STIG requires it), the config file agrees (so a reboot does not change it), no locally added permissive domains, and recent denials (AVCs) - which is what you see when SELinux blocks an application.

Setting	Default	What it does
<code>check_selinux_expected_mode</code>	<code>enforcing</code>	the mode SELinux must be in (and the config file must say); anything else -> critical
<code>check_selinux_avc_since</code>	<code>yesterday</code>	ausearch -ts: recent (10 min), today, yesterday, this-week
<code>check_selinux_avc_warn</code>	<code>1</code>	this many denials or more -> warning
<code>check_selinux_booleans</code>	<code>{}</code>	Booleans that must have a value, e.g. {httpd_can_network_connect: "on"}
<code>check_selinux_relabel_paths</code>	<code>[]</code>	Paths whose file labels are compared with the policy (a dry-run restorecon). Mislabeled files are a common cause of denials after files are copied or moved. Keep it to small trees. e.g. [/etc/nginx, /var/www, /etc/stigman]

check_services

File: `roles/check_services/defaults/main.yml`

Note

SERVICES: systemd units that failed, services that must be running (and start at boot), and services stuck restarting over and over.

Setting	Default	What it does
<code>check_services_required</code>	<code>[sshd, crond, rsyslog, chronyd, auditd]</code>	Must be running on every host. A host where one is not installed gets a warning (set <code>check_services_required</code> in the inventory to change the list for a group of hosts).
<code>check_services_required_extra</code>	<code>[]</code>	More for a group of hosts, e.g. in <code>group_vars/mariadb_hosts.yml: [mariadb]</code>
<code>check_services_ignore_failed</code>	<code>[]</code>	Failed units to ignore (regular expressions), e.g. a known-broken vendor unit.
<code>check_services_restart_warn</code>	<code>3</code>	More restarts than this since the unit started = a restart loop

check_time

File: `roles/check_time/defaults/main.yml`

Note

TIME: is the clock synchronized, and by how much is it off? A clock that is off by minutes breaks Kerberos/AD logins (5 minutes), TLS, log correlation and audit evidence.

Setting	Default	What it does
<code>check_time_offset_warn</code>	<code>0.5</code>	seconds
<code>check_time_offset_crit</code>	<code>5.0</code>	seconds off -> critical (warning from <code>check_time_offset_warn</code>)
<code>check_time_min_sources</code>	<code>1</code>	time servers chrony must be able to reach

ontap_report

File: `roles/ontap_report/defaults/main.yml`

Note

NetApp ONTAP health report (playbooks/ontap_health_report.yml). Read-only: GET requests to each cluster's REST API, with the account of a "NetApp ONTAP" credential (docs/NETAPP.md).

Setting	Default	What it does
<code>ontap_clusters</code>	<code>[]</code>	The clusters' management addresses (names or IPs), e.g. <code>[cluster1.yoursite.mil, cluster2.yoursite.mil]</code> . The credential's account must exist on each of them.
<code>ontap_validate_certs</code>	<code>true</code>	Check the clusters' HTTPS certificates. ONTAP's own certificate is often self-signed: then put its CA in <code>ontap_ca_path</code> (a file in the project, e.g. <code>"{{ playbook_dir }}/files/ca/ontap-ca.pem"</code>), or set false.

Setting	Default	What it does
ontap_ca_path	(empty)	the CA file (PEM) the clusters' certificates are checked with; empty = the execution environment's own trust
ontap_timeout	30	seconds for one request
ontap_volume_warn_pct	85	Thresholds (used %, hours, days, seconds): at or above "warn" = amber, at or above "crit" = red.
ontap_volume_crit_pct	90	volume used % -> red
ontap_aggr_warn_pct	85	aggregate used % -> amber
ontap_aggr_crit_pct	90	aggregate used % -> red
ontap_inode_warn_pct	85	inodes used % of a volume -> amber
ontap_inode_crit_pct	90	inodes used % -> red
ontap_snapmirror_lag_warn_hours	24	last good transfer longer ago than this
ontap_snapmirror_lag_crit_hours	48	last good transfer longer ago than this -> red
ontap_snapshot_reserve_warn_pct	100	snapshots using this much of their reserve (100 = they spill into the volume)
ontap_snapshot_reserve_crit_pct	150	snapshots using this much of their reserve -> red
ontap_cert_warn_days	60	certificates expiring within this many days
ontap_cert_crit_days	14	certificates expiring within this many days (or expired) -> red
ontap_time_drift_warn_seconds	60	a node's clock against the AAP server's
ontap_time_drift_crit_seconds	300	a node's clock this far off -> red
ontap_port_errors_warn	1	receive errors (CRC and other bad frames) since the node booted
ontap_port_errors_crit	10000	receive errors on a port since boot -> red
ontap_min_spares	1	fewer spare disks than this on a node = amber (0 = red)
ontap_ems_hours	24	the error events of the last this-many hours
ontap_ems_max	500	read at most this many events per cluster
ontap_report_max_rows	50	rows per table (problems first); the job's artifacts have all
ontap_snapshot_top	25	the snapshot space table: the volumes whose snapshots use the most

Setting	Default	What it does
ontap_report_fail	false	The job stays green when it finds problems: the report and email say what. true = mark it FAILED when there is a critical finding or a cluster could not be read, for workflows.

patch

File: roles/patch/defaults/main.yml

Note
PATCH: install updates with dnf, reboot if the updates need it, and prove the host came back: every service that was running before must be running after. One host at a time by default; the rollout stops at the first host that fails.

always

dnf-makecache\

packagekit\

fwupd\

poam

File: `roles/poam/defaults/main.yml`

Setting	Default	What it does
<code>poam_csv</code>	<code>"{{ playbook_dir }} / ../poam/poam.csv"</code>	POA&M STATUS: which Plan of Action & Milestones items are overdue or due soon, and (optional) which open CAT I / CAT II findings in STIG Manager have no POA&M item yet. The POA&M list is a CSV file in your work Git (export from eMASS or your spreadsheet, keep the columns below, save as CSV). POA&M data is sensitive: k... (all of it: the role file)
<code>poam_columns</code>	<code>id: POAM ID # unique item number weakness: Weakness # what is wrong severity: Severity #...</code>	Column names in YOUR file for each field this check needs (change the right-hand side).
<code>poam_closed_statuses</code>	<code>[Completed, Closed, Risk Accepted, Cancelled]</code>	items with one of these statuses are done: never reported
<code>poam_due_soon_days</code>	<code>30</code>	an open item due within this many days -> warning (past due -> critical)
<code>poam_date_formats</code>	<code>['%Y-%m-%d', '%m/%d/%Y', '%m/%d/%y', '%d-%b-%Y', '%d %b %Y', '%b %d, %Y']</code>	the date formats your file may use (Python strptime), tried in this order

Setting	Default	What it does
<code>poam_stigman_api</code>	(empty)	---- STIG Manager cross-check (optional) ----- Needs a Keycloak client that may use the client-credentials grant, with the scope stig-manager:collection:read, and a grant in the collections to read. Your realm owner creates it - this repository never changes the realm. T... (all of it: the role file)
<code>poam_stigman_collections</code>	<code>[]</code>	collection names; empty = every collection the client can read
<code>poam_stigman_severities</code>	<code>[high, medium]</code>	CAT I and CAT II findings must have a POA&M item
<code>poam_stigman_scope</code>	(empty)	Scope to ask for. Empty = the client's default scopes (they must include stig-manager:collection:read or stig-manager:collection).
<code>poam_validate_certs</code>	<code>true</code>	check STIG Manager's and Keycloak's HTTPS certificates

podman_discover

File: `roles/podman_discover/defaults/main.yml`

Note

PODMAN DISCOVERY (read-only): every podman container on the host, root's (rootful) and each user's (rootless), with how it is run (a systemd unit or plain podman) and whether it should be running. Used by the Health check `containers` and by Service watch. It changes nothing.

Setting	Default	What it does
<code>podman_discover_enabled</code>	<code>true</code>	false = do not look for containers at all (the containers check then says "skipped", and Service watch watches only its watch_containers list).
<code>podman_discover_rootless</code>	<code>true</code>	Also look at the containers of users (rootless podman). A user is looked at when a podman process (common) runs as them, when they have "linger" on (/var/lib/systemd/linger/NAME), or when they are in /etc/subuid AND have container storage (~/.local/share/containers/storage). Their containers are read AS that user (run... (all of it: the role file)
<code>podman_discover_users</code>	<code>[]</code>	More users to look at (names), e.g. an application account the rules above miss.

Setting	Default	What it does
<code>podman_discover_ignore</code>	<code>[]</code>	Containers to leave out, as regular expressions that must match the WHOLE container name or OWNER/NAME, e.g. <code>['test-.']</code> (every container whose name starts with <code>test-</code> , any owner), <code>['alice/.']</code> (all of alice's containers), <code>['root/web']</code> (root's web only).
<code>podman_discover_quadlet_dirs</code>	<code>[/etc/containers/systemd, /usr/share/containers/systemd]</code>	Where root's Quadlet files are (*.container). A user's are always read from <code>~/.config/containers/systemd</code> , <code>/etc/containers/systemd/users/UID</code> and <code>/etc/containers/systemd/users</code> .
<code>podman_discover_timeout</code>	<code>60</code>	Seconds one podman or systemctl call may take before it is given up (reported as a finding).

service_watch

File: `roles/service_watch/defaults/main.yml`

Setting	Default	What it does
<code>watch_discover</code>	<code>true</code>	---- What to watch ----- With <code>watch_discover</code> on (the default), service watch finds the containers itself: every podman container on the host that SHOULD be running - root's (rootful) and each user's (rootless) - as the Health check <code>containers</code> finds th... (all of it: the role file)
<code>watch_containers</code>	<code>[]</code>	Containers to watch in any case, in dependency order (a container's dependencies first). They are watched before the discovered ones, and one that is missing is reported as down. name: the container name, as <code>podman ps -a</code> shows it (or its Quadlet / service name) user: optional: the user that owns it (rootless podman)... (all of it: the role file)
<code>watch_url_ok_status</code>	<code>[200, 301, 302, 401, 403]</code>	HTTP codes that count as "answering"
<code>watch_find_units</code>	<code>true</code>	---- How each container is run (found by itself) ----- A container that a systemd unit runs (a Quadlet .container file, <code>podman generate systemd</code>, or a unit you wrote) must be started with <code>systemctl</code>, never with <code>podman start</code>: the unit owns it, and while the unit is stopped the co... (all of it: the role file)
<code>watch_quadlet_dirs</code>	<code>[/etc/containers/systemd, /usr/share/containers/systemd]</code>	where root's Quadlet files are looked for (to find the unit that runs a container)

Setting	Default	What it does
<code>watch_mode</code>	<code>approval</code>	---- What happens when one is down ----- approval: ACT only diagnoses. Its proposed fix waits for the AAP approval step, then <code>service_fix_approved.yml</code> runs exactly the approved command(s) - no model involved. When ACT proposes nothing (or cannot run), the standard fix... (all of it: the role file)
<code>watch_verify_retries</code>	<code>6</code>	re-checks after a fix (every <code>watch_verify_delay</code> seconds)
<code>watch_verify_delay</code>	<code>5</code>	seconds between those re-checks
<code>watch_record_file</code>	<code>/var/log/service-watch/incidents.jsonl</code>	---- Where incidents are recorded (on the host) ----- One JSON line per event, plus a syslog line (tag <code>service-watch</code>) that your log forwarding (Splunk, Elastic, rsyslog relay) picks up. AAP's job history keeps its own copy.
<code>watch_record_syslog</code>	<code>true</code>	also send each incident to syslog (tag <code>service-watch</code>)
<code>watch_act_env</code>	<code>{}</code>	---- ACT settings ----- Which provider and model: <code>site_act_provider</code> / <code>site_act_model</code> / <code>site_act_url</code> / <code>site_act_ca</code>, the same settings every runbook uses (roles/site_act/defaults/main.yml). The key comes from the "ACT model key" AAP credential, never f... (all of it: the role file)
<code>watch_act_timeout</code>	<code>1200</code>	seconds for one ACT run (the model's time per answer: <code>site_act_model_timeout</code>)
<code>watch_instructions</code>	(worked out by the role; see its file)	What ACT is told, on top of the <code>act_triage</code> role's standard instructions.

servicenow

File: `roles/servicenow/defaults/main.yml`

Note

SERVICENOW - open, update (and optionally resolve) incidents from check findings, and check that ServiceNow itself (the instance and its MID Servers) is healthy. Uses the REST Table API with `ansible.builtin.uri`: no extra collection to install. Connection: the "ServiceNow API" AAP credential sets `SN_HOST` (`https://<instance>`), `SN_USERNAME` and `SN_PASSWORD` in the job's environment. Without it, nothing is sent: the job prints what it WOULD have opened (a safe way to try it out).

	Default	What it does
<code>min_severity</code>	<code>warning</code>	Findings at or above this severity get a ticket: <code>warning</code> critical \
<code>table</code>	<code>incident</code>	the table tickets are opened in (incident)

	Default	What it does
<code>servicenow_assignment_group</code>	(empty)	Values for new incidents. Names work (the API is asked to accept display values), e.g. <code>servicenow_assignment_group: Linux Operations</code>
<code>servicenow_caller</code>	(empty)	user name; empty = the API account
<code>servicenow_category</code>	(empty)	e.g. Software
<code>servicenow_urgency</code>	<code>{critical: "1", warning: "2"}</code>	a new ticket's urgency, per finding severity (your instance's values)
<code>servicenow_impact</code>	<code>{critical: "2", warning: "3"}</code>	a new ticket's impact, per finding severity
<code>servicenow_set_ci</code>	<code>true</code>	set Configuration item = the host name (must exist in the CMDB)
<code>servicenow_note_cleared</code>	<code>true</code>	When a problem that had a ticket is gone, add a work note (and, if you want, resolve it).
<code>servicenow_resolve_cleared</code>	<code>false</code>	<code>true</code> = also resolve the ticket (<code>servicenow_resolved_state</code> , <code>servicenow_close_code</code>)
<code>servicenow_resolved_state</code>	<code>"6"</code>	Resolved
<code>servicenow_close_code</code>	<code>"Solution provided"</code>	must be a close code your instance has
<code>servicenow_note_repeats</code>	<code>true</code>	Add a work note to an open ticket each time its finding is seen again.
<code>servicenow_correlation_prefix</code>	<code>site-health</code>	Every ticket from these playbooks carries a correlation ID starting with this, so the same problem on the same host always updates the same ticket.
<code>servicenow_validate_certs</code>	<code>true</code>	check the instance's HTTPS certificate (see <code>servicenow_ca_path</code>)
<code>servicenow_ca_path</code>	(empty)	A CA bundle (PEM) for the instance's TLS certificate, when the execution environment does not trust its CA (e.g. a DoD root CA): a path on the controller, such as a file you add to the project: <code>servicenow_ca_path: "{{ playbook_dir }}/files/ca/servicenow-ca.pem"</code> . Empty = the execution environment's normal trust store.... (all of it: the role file)
<code>servicenow_timeout</code>	<code>30</code>	seconds for one API request
<code>servicenow_fail_when_findings</code>	<code>true</code>	Fail the job when it opened/updated tickets, so the workflow shows red while problems exist.

	Default	What it does
w_check_mid	true	servicenow_health.yml: also check the MID Servers (needs read access to the ecc_agent table)
w_slow_seconds	5	the API answering slower than this -> warning
w_test_resolve	true	servicenow_test_ticket.yml: resolve the test incident at the end (false = leave it open to look at)

site_act

File: roles/site_act/defaults/main.yml

What it does		
ACT (GenAI) on top of ANY check playbook. The check roles already did the cheap, certain part: they found WHAT is wrong (site_findings) and know which read-only command shows more (each finding's hint). This role hands that to ACT (the act_triage role) to find out WHY, and to propose - or, if you allow it, make - th... (all of it: the role file)		
self-heal only: regular expressions of fix commands ACT may run without a person. Each must match the WHOLE command. Keep them narrow, e.g. site_act_allow: ['systemctl restart (chronyd\	rsyslog\	crond)']
Hosts where ACT may only diagnose - never fix by itself, whatever site_act_level says: AAP itself (a fix could stop the job that is running it) and vendor-managed appliances. The apply step (act_fix_approved.yml) skips them too: a person applies the fix there by hand. AAP (groups aap, aap_hosts) is ALWAYS in this list... (all of it: the role file)		
---- Which model ACT talks to ----- site_act_provider genai GenAI.mil (the default) asksage Ask Sage (set site_act_url to your organization's instance) genai-beta the GenAI.mil beta proxy (preview models) site_act_models the model to use FOR EACH provider, e.... (all of it: the role file)		
the model for each provider (see above): the model follows the provider picked		
the API URL for each provider; a provider not listed uses ACT's default		
one model for this run, whatever the provider (e.g. a survey answer); empty = site_act_models		
one URL for this run; empty = site_act_urls		

What it does

a CA bundle file ON THE HOSTS for the provider's certificate;
empty = the host's trust store

Any other ACT environment settings (advanced; these win over the ones above), e.g. {HTTPS_PROXY: "http://proxy.yoursite.mil:8080"}. See the top of vendor/act/act for the list. Every ACT job streams the model's answer (ACT_STREAM=1), so a firewall or proxy idle timeout does not cut a long answer ("connection reset by p... (all of it: the role file)

Seconds the model may take for ONE answer, retries after a provider hiccup included (ACT's GENAI_TIMEOUT; ACT's own default, 90, is too short for a long analysis or a busy provider). Every ACT job uses it: the runbooks' use_act, the VMware alarms and capacity analyses, service watch.

seconds for one whole ACT run (several answers): keep it above site_act_model_timeout

How many hosts may run ACT at the same time in one job (0 = no limit: as many as the job's forks). Every host in a job uses the same key, and GenAI.mil's default quota is 60 requests and 200,000 tokens a minute per key: a fleet run with many hosts at once gets HTTP 429 (too many requests).

anything ACT should know about this site

extra evidence (troubleshoot.yml passes what it collected)

site_email

File: roles/site_email/defaults/main.yml

Note
EMAIL A REPORT through your mail relay (SMTP). Off until report_email_to has an address. Used by every report job: health checks, certificate reports, database health, POA&M, ServiceNow health, troubleshooting, the VMware jobs (docs/EMAIL_REPORTS.md). The relay settings belong in playbooks/group_vars/all.yml. If the relay needs a login, attach an "SMTP relay" credential (aap/credential_types/smtp_relay.yml) - never put a password in these settings.

Setting	Default	What it does
report_email_to	[]	Who gets it: a list, or text with commas / spaces / one per line (a survey answer). Empty = no email. Set it on each TEMPLATE (its Variables, or a survey): in all.yml every report job emails.
report_email_cc	[]	copied: a list, or text like report_email_to
report_email_from	(empty)	the sender, e.g. aap-noreply@yoursite.mil

Setting	Default	What it does
<code>report_email_smtp_host</code>	(empty)	your mail relay, e.g. smtp.yoursite.mil
<code>report_email_smtp_port</code>	25	25 or 587 with starttls, 465 with ssl
<code>report_email_security</code>	starttls	starttls = encrypt after connecting (the relay must offer it); ssl = encrypted from the start (port 465); none = not encrypted (only for a relay without encryption; then no login is sent).
<code>report_email_ca_path</code>	(empty)	a CA file that signed the relay's certificate (empty = the CAs the execution environment trusts)
<code>report_email_subject_prefix</code>	"[AAP]"	put before every subject, e.g. "[AAP] Certificate report: ..."
<code>report_email_html</code>	true	Send a formatted (HTML) version too, for reports that come with one (all VMware jobs). The email always carries a plain-text copy for mail clients that do not show HTML. false = plain text only.
<code>report_email_html_template</code>	report.html.j2	The look of the HTML version. Your own: copy roles/site_email/templates/report.html.j2 (e.g. to playbooks/files/email/report.html.j2), change it, and set report_email_html_template: "{{ playbook_dir }}/files/email/report.html.j2"

site_findings

File: `roles/site_findings/defaults/main.yml`

Note

Shared by every check playbook. A "finding" is one thing that is wrong on one host: check: the check that found it (disk, services, mariadb, ...) id: a stable name for this problem on this host, e.g. disk:/var. ServiceNow uses it to update the same ticket instead of opening a new one every run. severity: critical \ | warning summary: one line a person understands ("/var is 96% full (critical at 95%)") hint: a read-only command that shows more. A person can paste it into a shell; ACT (when you add it later) runs it to collect evidence. Check roles only ADD findings. This role starts the list, reports it, and decides pass/fail.

Setting	Default	What it does
<code>site_fail_on</code>	[critical, warning]	Which severities make the job FAIL. A failed job is what sends an AAP workflow down its "on failure" path (open tickets, ask ACT, wait for approval). [] = report only, never fail.
<code>site_syslog</code>	true	One syslog line per finding (tag below), so your log server / SIEM sees them too.
<code>site_syslog_tag</code>	site-health	the syslog tag those lines carry (to filter on in the SIEM)
<code>site_eventlog_source</code>	site-health	Windows: one Application event log entry per finding instead of syslog (source below).

Setting	Default	What it does
<code>site_os</code>	<code>linux</code>	--- Linux or Windows (set by the playbook; do not set) ----- Linux playbooks run the roles <code>check_<name></code> , the Windows ones (<code>win_*.yaml</code>) the roles <code>win_check_<name></code> .
<code>site_check_role_prefix</code>	<code>"{{ 'win_check_' if site_os == 'windows' else 'check_' }}"</code>	<code>check_</code> (Linux) or <code>win_check_</code> (Windows): set by the playbook from <code>site_os</code>
<code>site_windows_connections</code>	<code>[winrm, psrp, ansible.builtin.winrm, ansible.builtin.psrp]</code>	How a host is recognized as Windows: its connection, or its shell (Windows OpenSSH).
<code>site_settings_notice</code>	<code>true</code>	A notice in the job output when a project settings file (playbooks/group_vars/*.yaml) defines a setting: it wins over AAP's inventory Variables box. <code>false</code> = no notice.

stigman_stack

File: `roles/stigman_stack/defaults/main.yaml`

Note

STIG Manager stack on podman, managed by systemd through Quadlet unit files: `stigman-mysql` (MySQL 8.4) <- `stigman-api` (STIG Manager) <- `stigman-nginx` (TLS, port 443) Keycloak is EXTERNAL: the realm already exists and this role never touches it. Needs podman >= 4.4 (Quadlet): RHEL 9.2+ / RHEL 8.8+.

Setting	Default	What it does
<code>stigman_scope</code>	<code>system</code>	<code>system</code> = rootful (<code>/etc/containers/systemd, systemctl</code>) <- production user = rootless (<code>~/config/containers/systemd, systemctl --user</code>) <- tests / rootless hosts
<code>stigman_images</code>	<code>mysql: docker.io/library/mysql:8.4</code> <code>api: docker.io/nuwcdinvpt/stig-manager:1.6.17 # pin a...</code>	the container images (pin STIG Manager's version; change it deliberately)
<code>stigman_config_dir</code>	<code>/etc/stigman</code>	<code>nginx.conf</code> , TLS files, CA bundle
<code>stigman_https_port</code>	<code>443</code>	host port nginx publishes
<code>stigman_server_name</code>	<code>"{{ inventory_hostname }}"</code>	the name nginx answers to (and the TLS certificate is for)
<code>stigman_oidc_provider</code>	<code>"https://keycloak.example.mil/realms/stigman"</code>	Existing Keycloak realm (server-side URL used by the API; browser-side URL for the web client)

Setting	Default	What it does
<code>stigman_client_oidc_provider</code>	<code>"{{ stigman_oidc_provider }}"</code>	the Keycloak realm URL the browser uses (when it differs from the server-side one)
<code>stigman_client_id</code>	<code>stig-manager</code>	the Keycloak client the web app logs in with
<code>stigman_classification</code>	<code>U</code>	banner: U, CUI, C, S, TS, SCI, NONE
<code>stigman_db_schema</code>	<code>stigman</code>	the MySQL database STIG Manager uses
<code>stigman_db_user</code>	<code>stigman</code>	the MySQL account STIG Manager uses (its password: the credential)
<code>stigman_mysql_buffer_pool</code>	<code>1G</code>	STIG Manager docs: >= 8G for real deployments
<code>stigman_ready_retries</code>	<code>40</code>	How long to wait for the API to answer through nginx after (re)starting
<code>stigman_ready_delay</code>	<code>6</code>	seconds between those tries
<code>stigman_mysql_root_password</code>	<code>"{{ lookup('ansible.builtin.env', 'STIGMAN_MYSQL_ROOT_PASSWORD') }}"</code>	---- Secrets: NEVER set these in a vars file or inventory ----- They come from AAP credentials (aap/credential_types/*.yaml) as environment variables / files on the controller (execution environment). See docs/SECRETS.md.
<code>stigman_db_password</code>	<code>"{{ lookup('ansible.builtin.env', 'STIGMAN_DB_PASSWORD') }}"</code>	from the credential (STIGMAN_DB_PASSWORD), never set here
<code>stigman_tls_cert_src</code>	<code>"{{ lookup('ansible.builtin.env', 'TLS_CERT_FILE') }}"</code>	nginx's certificate file, from the TLS certificate credential (TLS_CERT_FILE)
<code>stigman_tls_key_src</code>	<code>"{{ lookup('ansible.builtin.env', 'TLS_KEY_FILE') }}"</code>	its private key file, from the credential (TLS_KEY_FILE)
<code>stigman_tls_ca_src</code>	<code>"{{ lookup('ansible.builtin.env', 'TLS_CA_FILE') }}"</code>	the CA bundle file, from the credential (TLS_CA_FILE)

troubleshoot

File: `roles/troubleshoot/defaults/main.yml`

What it does		
	TROUBLESHOOT: the commands an experienced admin runs first for a kind of problem, run for you (read-only), each with a line saying what it shows. The matching health checks run too, so you also get plain-words findings. Pick the area; for "service" and "network" say what. ts_area: overview \	disk \
	the systemd unit, for ts_area service (see above)	performance \
	host:port, for ts_area network	
	how far back logs are read (journalctl --since)	
	keep the last this-many lines of each command	
disk, mounts, time, - {why:...	What each area runs. why is printed above the command's output. Commands must be read-only.	

vmware_vm

File: roles/vmware_vm/defaults/main.yml

Note
VMWARE VMs: jobs that run against vCenter (not on the VMs), with the certified vmware.vmware collection. vCenter address and account come from a "VMware vCenter" credential on the job template (VMWARE_HOST, VMWARE_USER, VMWARE_PASSWORD). docs/VMWARE.md

Default	What it does
[]	The VMs, by their exact name in vCenter: a list, or (from a survey) one per line. No wildcards.
(empty)	only needed when the same VM name exists in two datacenters
cs true	Check vCenter's certificate. false only if the execution environment does not trust your CA yet.

	Default	What it does
	[]	Several vCenters (e.g. Enhanced Linked Mode), one report: the reports and the ESXi security settings read the credential's vCenter AND these (names, or name:port), with the same account. The same vCenter under two names is read once. docs/VMWARE.md, "Several vCenters". e.g. [vc01.yoursite.mil, vc02.yoursite.mil]
unread	false	true = the job shows failed when a vCenter could not be read (the report lists it in red either way)
groups	[aap, aap_hosts]	The AAP server is ALWAYS refused (hosts in the groups aap / aap_hosts, matched by VM name, guest host name and IP address), and so are the hosts of these groups and these VM names (* and ? allowed: DC*). List the AAP VM's name here when it differs from its host name, and vCenter's own VM. The snapshot cleanup never de... (all of it: the role file)
	[]	e.g. [AAP01, VCSA01, DC*] (* and ? allowed)
	guest	guest = ask the operating system (needs VMware Tools running; services stop cleanly) hard = reset / power off, like the button on a physical server (unsaved data is lost)
	600	shut down: seconds to wait for the VM to be off
	(empty)	The snapshot's name. Taking one: empty = aap-<date>-<time>. Deleting: the exact name to delete.
tion	(empty)	the AAP job number and user are added
	false	true = VMware Tools freezes the file systems first
all	false	delete: true = every snapshot of the VM (vm_snapshot_name empty)
	(empty)	The text to write.
	append	append (a new line, with date and user) \
		replace \
		clear

	Default	What it does	
	1	Which adapter: 1 = Network adapter 1, 2, ...	
	(empty)	the port group (standard or distributed) to move it to	
er	(empty)	Only the VMs in this vCenter folder ("" = every VM, or the ones in vm_names).	
	false	The job stays green (successful) when it finds VMs to fix: the report and email say what. true = mark the job FAILED when a VM is BIOS or has secure boot off, so a workflow's "on failure" path can react (open a ticket).	
t	85	Used space (% of capacity): a datastore at or above these is listed as a warning / critical.	
t	90	a datastore this % used or more is red (critical); amber from vm_datastore_warn_pct	
	false	The job stays green when datastores are full: the report and email say which. true = mark the job FAILED when one is at or above vm_datastore_warn_pct (or inaccessible), for workflows.	
days	3	Snapshots at least this many days old: listed apart by the report, deleted by the cleanup.	
regex	`(?i)keep\`	do.?not.?delete`	Never deleted by the cleanup (the report lists them, with the reason): - a snapshot whose name or description matches this
exclude_vms	[]	- a snapshot of these VMs: names, * and ? allowed (e.g. [DC*, VCSA01, AAP01]). The VMs in vmware_protected_vms and the AAP server are never cleaned up either.	
max_size_gb	1024	- a snapshot bigger than this many GB (1024 = 1 TB), or whose size cannot be read: merging a very big snapshot can take hours and slow the VM - do those by hand, at a quiet time.	

	Default	What it does
<code>_confirm</code>	<code>false</code>	The cleanup on its own (not after a workflow approval) deletes only when this is true - e.g. a survey question "Delete them?". <code>false</code> = it lists what it would delete, and deletes nothing.
<code>_max</code>	<code>50</code>	Refuse to delete more than this many snapshots in one run (a typo in the age would otherwise delete them all). Raise it on purpose for a big cleanup.
<code>timeout</code>	<code>1800</code>	seconds to wait for one deletion (vCenter merges it into the disk)
<code>email_only_</code>	<code>false</code>	The cleanup's email (report_email_to): true = only when it deleted something (or a deletion failed).
	<code>[vcenter, datacenter, cluster, host]</code>	Whose triggered alarms and configuration issues: vcenter, datacenter, cluster, host. Add vm, datastore or network for theirs too (VM alarms can be many).
	<code>24</code>	The events of the last this-many hours: every error and warning, failed logins, host connection events and VM shutdowns / power-offs / resets / HA restarts. 0 = alarms and host state only.
	<code>5000</code>	read at most this many events (the newest)
	<code>200</code>	rows per table in the report (the artifacts have them all)
<code>_cal</code>	<code>10</code>	This many failed logins of one user from one place, or more, is critical (fewer: a warning).
<code>_types</code>	<code>[BadUsernameSessionEvent, com.vmware.sso.LoginFailure, esx.audit.account.locked, esx.audi...</code>	The event types read whatever their category (vCenter's event type names). Add to these lists rather than replacing them; [] = none of that kind.
<code>_event_types</code>	<code>[HostConnectionLostEvent, HostNotRespondingEvent, HostDisconnectedEvent, HostReconnection...</code>	Events about a host losing (or failing to get back) its connection to vCenter: the Host connection table.

	Default	What it does	
types	[VmGuestShutdownEvent, VmGuestRebootEvent, VmPoweredOffEvent, VmResettingEvent, VmFailove...]	Events about a VM stopping, restarting or failing over (guest shutdown, power off, HA restart): the VM events table.	
	false	The job stays green when it finds alarms: the report and email say what. true = mark the job FAILED when there is a critical alarm or a host is not connected, for workflows.	
ems	40	ACT analysis (playbooks/vm_alarm_act_analysis.yml). ACT runs on the AAP side (the execution environment), never on the ESXi hosts; the model provider, model, URL, proxy and CA are the site_act_* settings (roles/site_act/defaults/main.yml), the key an "ACT model key" credential. analyze at most this many problems, the... (all of it: the role file)	
issues	true	configuration issues too (false = alarms and events only)	
if_none	false	true = email even when there is nothing to analyze	
url	true	first check that the AAP node can reach the model's URL	
timeout	"{{ site_act_model_timeout \	default(600) }}"	Seconds the model may take for its answer (ACT's GENAI_TIMEOUT, 90 by default, is too short for a long analysis): site_act_model_timeout (600) unless set here. A GENAI_TIMEOUT in site_act_env wins.
	[]	Which hosts: names (* and ? allowed; the short name matches too). [] = every host of vCenter (or of vmware_datacenter).	
ers	[]	only the hosts of these clusters (names, wildcards); [] = all	
le_hosts	[]	Never these hosts (names, wildcards). E.g. a host you opened SSH on for a support case: the nightly run would otherwise close it again.	

	Default	What it does
	<code>disabled</code>	The SSH service (TSM-SSH): disabled = stopped, and set to start manually; enabled = started, and set to start with the host; "" = leave it as it is.
	<code>disabled</code>	The ESXi Shell service (TSM), the host's local shell: the same choices.
ngs	<code>UserVars.ESXiShellTimeOut: 600</code> <code>UserVars.ESXiShellInteractiveTimeOut: 600</code>	Advanced settings, name: value; {} = none. The defaults (seconds): UserVars.ESXiShellTimeOut SSH and the ESXi Shell stop by themselves this long after someone starts them (0 = never) UserVars.ESXiShellInteractiveTimeOut an idle SSH or ESXi Shell session is logged out Others often set: UserVars.DcuiTimeOut (an idle DCU... (all of it: the role file)
own	<code>normal</code>	Lockdown mode: normal = the host is managed only through vCenter (the DCUI still works); disabled; "" = leave it as it is. strict is refused, and a host already strict is left alone.
own_exceptio	<code>[]</code>	Accounts that must keep logging in to the hosts directly under lockdown (a vulnerability scanner, a monitoring tool): host-local accounts, or Active Directory accounts with their own permission on the host (not vsphere.local, not an AD group). Added before lockdown is turned on; never removed.
only_if_cha	<code>true</code>	The email (report_email_to): true = only when a host was changed or failed; false = every run.
s	<code>[]</code>	---- vSphere capacity planning (playbooks/vm_capacity_report.yml) ---- only these clusters (names; * and ? allowed); [] = every cluster
_days	<code>90</code>	The history: vCenter's own daily statistics (cluster CPU and memory, datastore space), kept a year at statistics level 1. Read this many days; the growth is a straight line over the last vm_capacity_trend_days, with at least vm_capacity_min_samples daily values.

	Default	What it does
days	30	the growth (straight line) is fitted on this many most recent days
rules	7	fewer daily values than this in those days = 'not enough history' (no trend)
days	30	VMs added and removed (vCenter keeps events 30 days by default)
hosts	1	Plan for this many failed hosts per cluster (N+1); a larger HA admission control failover level wins.
memory_pct	80	memory used at most this %, with those hosts failed
cpu_pct	80	CPU used at most this %, with those hosts failed
vcpu_core_max	4	vCPUs per physical core (amber from here, red from 1.5 times it)
months	12	Hosts recommended per cluster (of its average host size) to keep memory and CPU at their targets with those hosts failed: now, and at the current growth within this many months.
memory_warn_pct	85	datastore used %: amber
memory_crit_pct	90	red, and the target of the storage runway
memory_warn_days	90	a runway shorter than this is amber
memory_crit_days	30	and red
local_datastores	false	a host's own (not shared) datastores count too
rows	100	rows per table (problems first); the artifacts have all
	false	The job stays green: the report says what is short. true = mark it FAILED when a cluster or datastore is critical (for workflows).

	Default	What it does
	false	GenAI: ACT reads the numbers and makes its own estimate, shown next to the job's (math vs GenAI). Needs an "ACT model key" credential on the template; the same site_act_* settings as the alarms analysis. ACT runs on the AAP side and runs no command.
items	30	items ACT estimates: everything overall, then clusters and datastores, problems first
model_timeout	"{{ site_act_model_timeout \ default(600) }}"	seconds the model may take for its answer
check_url	true	first check that the AAP node can reach the model's URL
table	(empty)	---- Deploy a VM from a ServiceNow ticket (playbooks/vm_deploy_from_ticket.yml) ---- The ticket (survey answer): INC..., TASK..., SCTASK..., RITM..., CHG... The table follows from the prefix; vm_deploy_ticket_table sets it for another prefix. The ticket says what to build in 'key: value' lines - in its description, or... (all of it: the role file)
table	(empty)	the ServiceNow table, for a ticket prefix the job does not know (e.g. u_vm_request)
	[]	What a ticket may ask for (names; * and ? allowed). Templates MUST be listed: nothing else is deployed. The others: [] = any. e.g. [rhel9-gold, win2022-gold]
	[]	clusters (and standalone hosts) a ticket may use; also the candidates of automatic placement
	[]	datastores (or datastore clusters) a ticket may use; also the candidates of automatic placement
	[]	VM folders a ticket may use (paths below the datacenter's VM folder, e.g. Linux/*)
	16	the most vCPUs a ticket may ask for
memory_gb	128	the most memory a ticket may ask for, in GB

	Default	What it does
vm_names	'^[A-Za-z0-9][A-Za-z0-9-]{0,62}\$'	The VM names allowed (a regular expression); e.g. '^[a-z0-9]{1,15}\$' for Windows-safe names.
ticket_group	(empty)	Only tickets assigned to this group (its name as ServiceNow shows it); "" = any group.
ticket_data	{}	What a ticket does not say: e.g. {cluster: PROD-CL01, datastore: DS-CLUSTER-01, folder: Linux/New}. Not set = the template's own cluster, datastore and folder; its CPU and memory.
ticket_sites	{}	Sites: a ticket's 'site: A' line (or its Location field, when it is one of these names) picks the placement - over vm_deploy_defaults; a cluster / datastore / folder the ticket names itself still wins. vm_deploy_sites: SiteA: {cluster: A-PROD-CL01, datastore: A-DSC01, folder: Servers/New} SiteB: {cluster: B-PROD-CL01,... (all of it: the role file)
require_site	true	with sites set up: a ticket must say its site (false = vm_deploy_defaults)
no_os_templates	[]	Templates with no operating system that boot from the network (PXE) into a deployment (e.g. an MECM task sequence): no guest customization, no waiting for an IP; the ticket gets the MAC address. e.g. [win2022-pxe]
max_disks	4	disks a ticket may ask for (disks_gb: 100, 200 = disk 1 grown to 100 GB, one 200 GB disk added)
max_disk_gb	2048	the biggest disk a ticket may ask for, in GB
total_disk_gb	4096	all its disks together, in GB

	Default	What it does
ct_pct	80	Automatic placement - when the ticket, its site and vm_deploy_defaults name no cluster / datastore: the job chooses among a site's 'clusters' / 'datastores' lists (e.g. SiteA: {clusters: [A-PROD-*]}), else among vm_deploy_clusters / vm_deploy_datastores; with neither, the template's own. a cluster fits when its memory... (all of it: the role file)
hosts	1	hosts to plan as failed when choosing a cluster
max_used_pc	85	a datastore fits when it stays under this % with the VM at its full size
	localdomain	the DNS domain Linux customization sets (a ticket's domain: wins)
	true	guest customization: host name = vm_name, every adapter on DHCP
	true	power the new VM on (false = leave it off, e.g. to check it first)
omization	900	seconds to wait for vCenter's customization result
	600	seconds to wait for an IPv4 address (DHCP, through VMware Tools)
ls	{}	Fields set on the ticket when the VM is up, e.g. to resolve an incident: {state: "6", close_code: "Solution provided", close_notes: "VM deployed"} (your instance's values)
	false	---- GenAI reads a request written in plain words (a proof of concept, off by default) ---- When the ticket's own 'key: value' lines do not give vm_name and template: ACT (GenAI, evidence mode, no commands) reads the ticket and proposes the VM - each value with the ticket's words it comes from, or a question to the re... (all of it: the role file)

	Default	What it does
catalog	<code>{}</code>	The templates GenAI may choose from, each with a plain description it can match the request to: {win2022-pxe: "Windows Server 2022 (installed over the network by MECM)", rhel9-gold: "Red Hat Enterprise Linux 9"}
approval	<code>true</code>	true: a GenAI proposal waits for a person - the workflow's approval (or a second run) - before it is built. false: one run reads, checks and builds, when every value is traced to the ticket and within the rules (questions or a refusal still stop it).
	(empty)	Which step this run is: "" = one job does it all; propose = the workflow's first step (read, ask GenAI if needed, check, write where it would go on the ticket - build nothing); build = the step after the approval (build what the ticket and its proposal say - never ask GenAI).
model_timeout	<code>"{{ site_act_model_timeout \</code>	<code>default(600) }}"</code> seconds the model may take for its answer

win_act

File: `roles/win_act/defaults/main.yml`

	Default	What it does
act_dir	<code>'C:\ProgramData\act'</code>	ACT FOR WINDOWS on the servers (vendor/act-windows/act.ps1, the PowerShell port of ACT). Installed into win_act_dir. That folder gets its own owner and permissions - Administrators and SYSTEM only, nothing inherited - because ordinary users may create folders and files in C:\ProgramData (and C:\temp): a non-admin coul... (all of it: the role file)
act_state	<code>present</code>	present \ absent (removes win_act_dir)
act_source	<code>"{{ role_path }}/../../../../vendor/act-windows/act.ps1"</code>	The ACT script this repository ships (vendor/act-windows, version in its VERSION file)

win_check_accounts

File: `roles/win_check_accounts/defaults/main.yml`

Note

WINDOWS ACCOUNTS (local accounts; domain controllers have none and are skipped) - the built-in Guest account is disabled -> critical if not - enabled local accounts: password never expires, password too old, no logon for a long time -> warning - members of the local Administrators group that you did not allow -> warning Names below allow wildcards (* and ?), e.g. 'YOURDOMAIN\Domain Admins', '*\svc_aap'.

Setting	Default	What it does
win_check_accounts_password_max_age_days	60	an enabled local account's password older than this -> warning (0 = off)
win_check_accounts_inactive_days	35	an enabled local account with no logon for this many days -> warning (0 = off)
win_check_accounts_password_never_expires_allowed	[]	accounts allowed "password never expires" (wildcards)
win_check_accounts_inactive_allowed	[]	accounts allowed to stay unused (wildcards)
win_check_accounts_admins_allowed	[]	Empty = list the members in the job output only. Set it to report anyone else, e.g. win_check_accounts_admins_allowed: ['*\Administrator', 'YOURDOMAIN\Domain Admins', 'YOURDOMAIN\svc_aap']

win_check_audit

File: roles/win_check_audit/defaults/main.yml

Note

WINDOWS AUDIT: the audit policy records what the STIG asks for, the event logs are big enough, and the log forwarding agent runs. One finding lists every missing audit setting.

	Default	What it does		
audit_subcategories	- {name: Credential Validation, guid: 0CCE923F-69AE-11D9-BED3-505054503030, need: Success...	Advanced audit policy subcategories (checked by GUID, so any Windows language works). need: Success \	Failure \	Success and
audit_security_log_min_k	196608	STIG		
audit_application_log_mi	32768	the Application log's maximum size must be at least this (KB)		
audit_system_log_min_kb	32768	the System log's maximum size must be at least this (KB)		

	Default	What it does
<code>audit_forwarder_services</code>	<code>[]</code>	Log forwarding agents that must run, e.g. [SplunkForwarder] or [nxlog]

win_check_certs

File: `roles/win_check_certs/defaults/main.yml`

Note
WINDOWS CERTIFICATES: every certificate in the machine stores below, and which of them are in use (https bindings of IIS / WSUS / ..., the WinRM HTTPS listener AAP connects through, Remote Desktop). Expiring soon -> warning / critical. Expired: critical when in use, or when it expired less than `win_check_certs_expired_report_days` ago (older leftovers are only listed in the certificate report).

Setting	Default	What it does
<code>win_check_certs_warn_days</code>	<code>30</code>	a certificate expiring within this many days -> warning (amber in the certificate report)
<code>win_check_certs_notice_days</code>	<code>60</code>	the certificate report colours these blue (an early heads-up; not a finding)
<code>win_check_certs_crit_days</code>	<code>7</code>	within this many days -> critical
<code>win_check_certs_expired_report_days</code>	<code>30</code>	an expired certificate NOT in use is a finding for this many days after it expired; older ones only in the report
<code>win_check_certs_stores</code>	<code>['Cert:\LocalMachine\My', 'Cert:\LocalMachine\WebHosting', 'Cert:\LocalMachine\RemoteDes...']</code>	the machine certificate stores read
<code>win_check_certs_ignore_subjects</code>	<code>[]</code>	wildcards, e.g. ['CN=localhost*']
<code>win_check_certs_endpoints</code>	<code>[]</code>	TLS services whose certificate a client actually gets, e.g. <code>win_check_certs_endpoints:</code> - {name: IIS, host: 127.0.0.1, port: 443, servername: web01.yoursite.mil}

win_check_disk

File: `roles/win_check_disk/defaults/main.yml`

Note

WINDOWS DISK: how full each local fixed volume (C:, D:, ...) is.

Setting	Default	What it does
<code>win_check_disk_warn_pct</code>	85	% used -> warning
<code>win_check_disk_crit_pct</code>	95	% used -> critical
<code>win_check_disk_min_free_gb</code>	0	also warn below this many GiB free (0 = off); big volumes
<code>win_check_disk_overrides</code>	{}	Different limits for particular volumes, e.g. a database volume kept full on purpose: win_check_disk_overrides: 'D:': {warn: 92, crit: 98}
<code>win_check_disk_ignore</code>	[]	volumes not to check, e.g. ['E:']

win_check_eventlog

File: `roles/win_check_eventlog/defaults/main.yml`

Note

WINDOWS EVENT LOG, the last win_check_eventlog_hours: - unexpected shutdown or power loss (41, 6008), a crash / blue screen (1001) -> critical - disk and file system errors (disk, Ntfs, storage drivers) -> critical - the Security log was cleared (1102) -> critical - many errors in System + Application (the finding names the top sources) -> warning - a System / Application / Security log that is full and not overwriting -> critical

Setting	Default	What it does
<code>win_check_eventlog_hours</code>	24	how many hours back the event logs are read
<code>win_check_eventlog_error_warn</code>	50	this many errors or more in System + Application -> warning

win_check_network

File: `roles/win_check_network/defaults/main.yml`

Note

WINDOWS NETWORK, from the server itself: - a default gateway and DNS servers are configured; - the names you list resolve, the TCP services you list answer; - domain members: the computer's trust with the domain works (nltest /sc_query). A broken trust is the classic "the trust relationship between this workstation and the primary domain failed" at logon.

Setting	Default	What it does
<code>win_check_network_dns_names</code>	<code>[]</code>	Names every Windows server must resolve, e.g. [dc01.yoursite.mil, ldap.yoursite.mil]
<code>win_check_network_tcp</code>	<code>[]</code>	Services every Windows server must reach, e.g. win_check_network_tcp: - {name: LDAPS, host: dc01.yoursite.mil, port: 636} - {name: SIEM, host: siem.yoursite.mil, port: 9997}
<code>win_check_network_secure_channel</code>	<code>true</code>	domain members: test the computer's trust with the domain (nltest /sc_query)
<code>win_check_network_timeout_ms</code>	<code>3000</code>	milliseconds for each TCP test

win_check_patching

File: `roles/win_check_patching/defaults/main.yml`

Note

WINDOWS PATCHING (read-only; installing updates is a later runbook) - a restart is pending for installed updates (Windows Update, component servicing, ConfigMgr) - days since the last update was installed (Get-HotFix) - optional: ask Windows Update which updates are missing (slower; needs full PowerShell)

Setting	Default	What it does
<code>win_check_patching_warn_days</code>	<code>35</code>	last update installed this many days ago or more -> warning
<code>win_check_patching_crit_days</code>	<code>60</code>	this many days or more -> critical
<code>win_check_patching_reboot_pending</code>	<code>true</code>	report a restart pending for installed updates
<code>win_check_patching_search_pending</code>	<code>false</code>	also ask Windows Update which updates are missing (slower)

win_check_performance

File: `roles/win_check_performance/defaults/main.yml`

Note

WINDOWS PERFORMANCE: CPU (average of three samples a second apart), memory, page file. The finding names the busiest processes, so you know where to look.

Setting	Default	What it does
<code>win_check_performance_cpu_warn_pct</code>	<code>90</code>	CPU busy % (average of three samples) -> warning
<code>win_check_performance_cpu_crit_pct</code>	<code>98</code>	CPU busy % -> critical

Setting	Default	What it does
<code>win_check_performance_mem_warn_pct</code>	90	memory in use % -> warning
<code>win_check_performance_mem_crit_pct</code>	97	memory in use % -> critical
<code>win_check_performance_pagefile_warn_pct</code>	80	page file in use % -> warning

win_check_security

File: `roles/win_check_security/defaults/main.yml`

Note

WINDOWS SECURITY BASELINE (a few high-value STIG items; not a full STIG scan: that is your SCAP / vulnerability scanner)

Setting	Default	What it does
<code>win_check_security_av</code>	auto	Antivirus: auto = check Microsoft Defender when it is the active antivirus (skipped when it is passive or off because another product protects the server); defender = always; none = never.
<code>win_check_security_av_signature_max_days</code>	7	Defender's signatures older than this many days -> finding
<code>win_check_security_av_services</code>	[]	Services of your endpoint protection that must run, e.g. Trellix / McAfee: [masvc, mfemms]
<code>win_check_security_firewall_profiles</code>	[Domain, Private, Public]	must be on
<code>win_check_security_smb1</code>	true	SMBv1 must be off
<code>win_check_security_rdp_nla</code>	true	Remote Desktop must require Network Level Authentication
<code>win_check_security_uac</code>	true	User Account Control must be on
<code>win_check_security_legacy_tls</code>	true	TLS 1.0 / 1.1 must be off for incoming connections

win_check_services

File: `roles/win_check_services/defaults/main.yml`

Note

WINDOWS SERVICES - every required service must be running (critical); - an automatic service that stopped WITH AN ERROR (exit code not 0) is a warning. Automatic services that stop cleanly by design (trigger-start, delayed) are not reported; - a service that crashed several times recently (System log events 7031 / 7034).

Setting	Default	What it does
<code>win_check_services_required</code>	[EventLog, RpcSs, Dnscache, W32Time, WinRM, mpssvc]	Must be running on every Windows server. Add your own per group with ..._required_extra, e.g. for IIS servers: win_check_services_required_extra: [W3SVC, WAS]
<code>win_check_services_required_extra</code>	[]	more services that must run, e.g. for a group of IIS servers (see above)
<code>win_check_services_ignore</code>	- gupdate - gupdatem - edgeupdate - edgeupdatem - MicrosoftEdgeElevationService - GoogleU...	Automatic services never to report (wildcards allowed: * and ?)
<code>win_check_services_crash_hours</code>	24	look back this far for unexpected stops
<code>win_check_services_crash_warn</code>	3	this many unexpected stops of one service -> warning

win_check_time

File: `roles/win_check_time/defaults/main.yml`

Note

WINDOWS TIME: the Windows Time service runs, syncs with a real source (not the local clock), synced recently, and how far off the clock is from its source. Kerberos (domain logons) fails when a clock is 5 minutes off.

Setting	Default	What it does
<code>win_check_time_offset_warn_sec</code>	1	seconds off its time source -> warning
<code>win_check_time_offset_crit_sec</code>	5	seconds off -> critical
<code>win_check_time_max_hours_since_sync</code>	24	last successful sync longer ago than this many hours -> finding

win_patch

File: `roles/win_patch/defaults/main.yml`

Note

WINDOWS PATCH through the ConfigMgr client: install what Software Center has for the server (the software updates deployed to it), restart if the updates need it, and prove the server came back: every automatic service that was running before runs again. One server at a time by default; the rollout stops at the first server that fails. Check mode lists and changes nothing.

Setting	Default	What it does	
<code>automatic_restarts</code>	<code>false</code>	NO AUTOMATIC RESTARTS unless you say so: this playbook restarts a server only when <code>automatic_restarts</code> is true (the job template's Variables, or the survey) and the updates need it. With <code>automatic_restarts</code> false (the default) it installs the updates and tells you which servers still need a restart; a person restarts them (all of it: the role file)	
<code>_patch_never_reboot_group</code>	<code>"{{ patch_never_reboot_groups \</code>	<code>default(['aap', 'aap_hosts', 'sat', 'netapp', 'netapp_con...'</code>	Servers in these groups are patched but NEVER restarted by this playbook (restart them by hand). Follows the Linux setting (<code>patch_never_reboot_groups</code>) when you set it, so ONE list covers both; the fallback is the Linux role's default (<code>roles/patch/defaults/main.yml</code>). <code>tests/test_protection.py</code> checks
<code>_patch_never_patch_groups</code>	<code>"{{ patch_never_patch_groups \</code>	<code>default(['no_patch', 'aap', 'aap_hosts', 'sat', 'netapp',...'</code>	Servers in these groups are skipped entirely (e.g. servers patched only by hand).
<code>_patch_scan_first</code>	<code>true</code>	ask the client to rescan and re-evaluate before installing	
<code>_patch_scan_wait_sec</code>	<code>120</code>	let that scan settle	
<code>_patch_timeout_min</code>	<code>120</code>	per round: download + install	
<code>_patch_poll_sec</code>	<code>30</code>	how often to look at the progress	
<code>_patch_max_rounds</code>	<code>2</code>	After a restart, look again: servicing-stack and cumulative updates often reveal more updates.	
<code>_patch_restart_delay_sec</code>	<code>30</code>	<code>shutdown.exe /t</code>	
<code>_patch_restart_timeout_mi</code>	<code>30</code>	how long the server may take to come back	
<code>_patch_restart_settle_sec</code>	<code>30</code>	after the port answers again, before trying to log on	
<code>_patch_reconnect_pause_se</code>	<code>20</code>	between logon attempts after the restart	
<code>_patch_services_grace_min</code>	<code>5</code>	After the restart, give delayed-start services this long before calling one missing.	

ing	Default	What it does	
_patch_services_ignore	['gupdate*', 'edgeupdate*', 'MicrosoftEdgeElevationService', 'GoogleUpdater*', 'MapsBroke...	Automatic services that come and go on their own (wildcards); not compared before/after.	
_patch_winrm_host	`"{{ ansible_host \	default(inventory_hostname) }}"`	Where AAP reaches WinRM, to the server go down and come back during the restart
_patch_winrm_port	`"{{ ansible_port \	default(5986) }}"`	its WinRM port (5986 = HTTPS)

win_troubleshoot

File: roles/win_troubleshoot/defaults/main.yml

isk \

performance \

service \

network \