SITE AUTOMATION

Settings Reference

Every setting of every job: its default, what it does, and where to set it

Site Automation - read it next to the setup guide | site-automation 1.32.0 | October 2026

What this is. Every setting the site-automation jobs read - the value it has when you set nothing, and what it does - with how to use one: where to put it (a template's Variables, `playbooks/group_vars/all.yml`, a group or a host), which place wins, and how to see what a job used.

How to read it

Keep it open next to the *Operations Runbooks Setup Guide*: each job's chapter there says which settings matter for it; this one lists them all, role by role.

Rendered from `docs/VARIABLES.md` and `docs/VARIABLES_REFERENCE.md` (generated from `roles/*/defaults/main.yml`) in the site-automation repository (version 1.32.0).

Host names and addresses in this guide are placeholders. Copy YAML from the repository files, not from the PDF.

Contents

1. Settings: what they are, where they go, and who wins 5

The two-minute version . . . 5
How to use a setting . . . 5
1. What you can set . . . 6
2. Where do I put a setting? . . . 8
3. Who wins when the same setting is in two places? . . . 10
Examples . . . 11
How to verify a setting is in effect . . . 12
If a setting does nothing . . . 12

2. Every setting and its default 13

aap_ee . . . 13
aap_email_question . . . 14
aap_health . . . 15
aap_template_audit . . . 22
ad_lockout . . . 23
ad_report . . . 24
api_availability . . . 27
cert_inventory . . . 28
check_accounts . . . 29
check_auditd . . . 29
check_certs . . . 30
check_containers . . . 37
check_disk . . . 37
check_fapolicyd . . . 39
check_logging . . . 40
check_mariadb . . . 40
check_mounts . . . 43
check_network . . . 44
check_patching . . . 44
check_performance . . . 45
check_selinux . . . 46
check_services . . . 46

check_sso . . . 47
check_stig . . . 47
check_time . . . 53
check_trust . . . 54
evaluate_stig_update . . . 58
fapolicyd . . . 59
fapolicyd_fix . . . 59
fapolicyd_rules . . . 60
horizon . . . 62
incident_docs . . . 71
log_analysis . . . 74
mail_relay . . . 78
net_certs . . . 81
net_config . . . 83
net_perf . . . 86
ontap_report . . . 88
patch . . . 90
pki_bundle . . . 93
poam . . . 95
podman_discover . . . 95
service_watch . . . 97
servicenow . . . 98
servicenow_insights . . . 108
servicenow_war . . . 117
site_act . . . 120
site_email . . . 122
site_findings . . . 129
stigman_ckl_import . . . 130
stigman_health . . . 132
stigman_stack . . . 134
stigman_update . . . 136
tenable_scan . . . 137
tmp_cleanup . . . 138
troubleshoot . . . 139

ucs_report . . . 140
vmware_ask_stack . . . 140
vmware_vm . . . 143
win_act . . . 175
win_check_accounts . . . 176
win_check_audit . . . 176
win_check_certs . . . 177
win_check_disk . . . 179
win_check_eventlog . . . 180
win_check_network . . . 180
win_check_patching . . . 180
win_check_performance . . . 181
win_check_security . . . 181
win_check_services . . . 182
win_check_stig . . . 183
win_check_time . . . 184
win_check_trust . . . 185
win_patch . . . 186
win_troubleshoot . . . 187
Settings of one host . . . 188

1. Settings: what they are, where they go, and who wins

A **variable** (also called a *setting*) is a named value the runbooks read: `check_disk_warn_pct: 85` means "warn when a disk is 85 % full". You change a variable to change how a check behaves, without touching any code. This page answers three questions:

1. **What can I set?** (find the name and the default)
2. **Where do I put it?** (for every host, for one group, for one host, for one run)
3. **What happens when the same setting is in two places?** (who wins)

You do not need to read it all. Do the two-minute version first.

The two-minute version

1. **Find the setting.** Open VARIABLES_REFERENCE.md, press Ctrl+F and search for a word (`disk`, `certificate`, `mariadb`). Every setting is there with its default.
2. **Put the line in one file.**
   - Every host: `playbooks/group_vars/all.yml`
   - Hosts in one AAP group: `playbooks/group_vars/<group name>.yml`
   - One host: `playbooks/host_vars/<host name>.yml`
3. **Commit and Sync Changes** in VS Code. Run the check.

That is all for a normal change. The rest of this page explains the details and the exceptions.

For one job template: `aap/extra_vars/<playbook>.yml` (the list) has every setting that template reads, switched off, with its default and what it does. Copy the file into the template's **Variables** box (YAML selected), delete the `#` in front of the settings you change, and change their values. Switch on only what you change: a template's Variables beat `all.yml`.

How to use a setting

Every job's chapter in these docs ends its settings part with an **All settings** link: the list of that job's settings in VARIABLES_REFERENCE.md, each with its default and what it does. Pick the place by **who the setting is for**:

The setting is about	Put it in	Example
hosts (thresholds, what a check looks at)	`playbooks/group_vars/all.yml` (every host), `playbooks/group_vars/<group>.yml`, `playbooks/host_vars/<host>.yml`	`check_disk_warn_pct: 90`
one job template (who gets its email, its rules)	the job template's **Variables** box in AAP	`report_email_to: ops@yoursite.mil` on the certificate report template
one run	a **survey** question (its variable name = the setting's name) or **Prompt on launch**	a survey question `health_checks`, multiple choice `daily` / `weekly`

Three examples:

1. playbooks/group_vars/all.yml - every host: disks warn at 90 %, a stricter audit log partition

```
check_disk_warn_pct: 90
check_disk_overrides:
  /var/log/audit: {warn: 70, crit: 85}

# 2. the "Certificate report" template's Variables - this job only: email it, stay green
report_email_to: [pki-team@yoursite.mil]
site_fail_on: []

# 3. playbooks/group_vars/win_web.yml - the Windows web servers only: IIS must run too
win_check_services_required_extra: [W3SVC, WAS]
```

- **A list replaces the default list** (it does not add to it): repeat the entries you keep. Details: "Three shapes of value" below.
- **Check what a job used:** in AAP, the job's **Details** tab shows the template's and the launch's variables; the job's output prints the decisions it made (thresholds in the finding texts, `REFUSED:` lines, the settings each report names in its email); its **artifacts** (Details > Artifacts) carry the results. A setting that seems ignored: If a setting does nothing.

1. What you can set

Every setting lives, with its default and an explanation, in the role that uses it: `roles/<role>/defaults/main.yml`. Two easy ways to read them:

- **VARIABLES_REFERENCE.md**, the same information as one searchable page (generated from the role files, so it is never out of date). Start here.
- In VS Code: **Ctrl+Shift+F** (search in all files), type the setting or a word, and look at the hits under `roles\...\defaults\main.yml`.

Setting names start with the thing they belong to:

Starts with	Belongs to	Example
`check_disk_`, `check_services_`, `check_mariadb_` ...	one health check (`roles/check_<name>`)	`check_disk_warn_pct`
`win_check_`	one Windows health check	`win_check_disk_warn_pct`
`site_act_`	ACT (GenAI)	`site_act_level`
`site_` (other)	reporting and settings	`site_fail_on`
`patch_`, `win_patch_`, `automatic_restarts`	patching	`patch_never_patch_groups`
`servicenow_`	ServiceNow tickets	`servicenow_assignment_group`
`watch_`	service watch	`watch_containers`, `watch_discover`
`podman_discover_`	finding the podman containers (the `containers` check and service watch share it)	`podman_discover_ignore`
`poam_`	POA&M	`poam_stigman_api`
`stigman_`	STIG Manager deployment	`stigman_scope`

Check names and what each finding means: RUNBOOKS.md.

Three shapes of value

YAML

```
check_disk_warn_pct: 85                    # a number
check_fapolicyd_required: true             # true or false (lower case)
check_services_required: [sshd, crond]     # a list: [a, b] or one "- item" per line
check_disk_overrides:                      # a set of names with values (a "map")
  /var/log/audit: {warn: 70, crit: 85}
```

A list replaces the default list. It does not add to it. If the default is `[sshd, crond, rsyslog]` and you write `[chronyd]`, the hosts now require **only** `chronyd`. Copy the default list from the reference and add your entries to it. (Some settings have a separate `..._extra` version that *adds*, for example `check_services_required_extra`. The reference says which.)

Job variables (per run, not per host)

These choose what a run does. Put them in a job template's **Variables** box (*Extra variables*) or in a survey question, never in a settings file:

Variable	Meaning	Where it appears
`target`	which hosts or group to run on (default: all)	most playbooks
`health_checks`	which checks: `daily`, `weekly`, `all`, or a list	Health check
`ts_area`, `ts_service`, ...	what to troubleshoot	Troubleshoot
`use_act`, `site_act_level`	ask ACT, and what it may do	Health check, Troubleshoot
`site_fail_on`	which findings fail the job: `[critical]` or `[critical, warning]`	any check job
`patch_serial`, `win_patch_serial`	how many hosts are patched at the same time (1). **Only** here: in `all.yml` or a group's or host's variables it does nothing	Patch hosts, Windows patch
`stigman_replace_secrets`	`true` for ONE launch: replace the database secrets when the database is created again (SECRETS.md)	STIG Manager - deploy

2. Where do I put a setting?

Place	Who gets it	How you change it	Good for
`roles/<role>/defaults/main.yml`	everyone	do **not** edit (updates replace it)	the built-in default
`playbooks/group_vars/all.yml`	every host	VS Code, commit, sync	site-wide values: the LDAP server, the SIEM, thresholds
`playbooks/group_vars/<group>.yml`	the hosts of that AAP group	VS Code	one kind of server: STIG Manager containers, MariaDB, AAP itself
`playbooks/host_vars/<host>.yml`	that one host	VS Code (create the folder and file)	one odd server: a bigger disk limit for `/var/log/audit`
AAP: inventory > Variables	every host	AAP web UI	values you prefer to keep in AAP
AAP: group > Variables	the hosts of that group	AAP web UI	same, per group
AAP: host > Variables	that host	AAP web UI	same, per host
Job template > Variables (*Extra variables*)	every run of that template	AAP web UI	"this template always uses `site_act_level: diagnose`"
Survey question / prompt on launch	this launch only	asked at launch	a choice made each time
AAP credential	the job	AAP web UI	**secrets only** (SECRETS.md)

File names matter. `playbooks/group_vars/<group>.yml` applies to the group with **exactly that name** in your AAP inventory, and `playbooks/host_vars/<host>.yml` to the host with exactly that name. A file for a group your inventory does not have is silently never used.

A few settings belong to one host - a network device's address, for example. They are at the end of the Settings Reference (settings of one host): set them in the host's Variables or `playbooks/host_vars/<host>.yml`, never in a template's Variables (there one value would be used for every host).

The files shipped in `playbooks/group_vars/` and the group each expects:

File	Group name it applies to	If your group has another name
`all.yml`	every host	(always works)
`aap.yml`	`aap`	copy it to `aap_hosts.yml` if your group is `aap_hosts` (the example inventory's name)
`mariadb.yml`	`mariadb`	the MariaDB check also runs for a group named `mariadb_hosts`, but the **settings** come from the file named after the group: copy it to `mariadb_hosts.yml`

File	Group name it applies to	If your group has another name
`sat.yml`, `idm.yml`, `logstash.yml`, `stigman.yml`, `rhel8_all.yml`	the same name	copy or rename
`netapp.yml`	`netapp`	copy to `netapp_console_hosts.yml` if that is your group

(The built-in safety lists cover both spellings of the AAP and NetApp group names, so AAP stays protected whatever you call it.)

Updating the repository never overwrites your settings files. The update script (`scripts/update-from-release.ps1` / `.sh`) protects `playbooks/group_vars/`, `playbooks/host_vars/`, `inventories/site/`, `poam/poam.csv` and anything in `.site-local`. If a new release adds a settings file you do not have, the script adds it **once** (the preview lists it as `YOURS`) and never changes it again. A setting that the release's `all.yml` mentions and yours does not is listed under `NEW SETTINGS`, so you can copy it over if you want it.

Adding "extras" to a host or a group in AAP

Where the boxes are in AAP 2.7 (older versions word a label differently):

- **Inventory**: **Automation Execution › Infrastructure › Inventories** › your inventory › **Edit inventory** › **Variables**.
- **Group**: the same inventory › **Groups** tab › the group › **Edit group** › **Variables**.
- **Host**: the same inventory › **Hosts** tab › the host › **Edit host** › **Variables**.
- **Job template**: **Automation Execution › Templates** › the template › **Edit template** › **Extra variables** (the template's *Variables* box).

The **Variables** box on a host or group takes the same YAML as a settings file. Typical entries:

a group's Variables box: hosts in this group need a bigger disk margin

```
check_disk_overrides:
  /var/lib/pgsql: {warn: 70, crit: 85}
```

a host's Variables box: its MariaDB containers belong to a rootless account

```
check_mariadb_container_user: svc_podman
```

a group's Variables box: never check or watch the test containers, or anything of one user

```
podman_discover_ignore: ['test-.*', 'devuser/.*']
```

Connection settings (`ansible_host`, `ansible_user`, `ansible_python_interpreter`, `ansible_connection: winrm`) belong here too. The *login and password* never do: they are credentials (SECRETS.md).

Prefer the project's settings files (this repository), because they are reviewed and versioned in Git and updates never overwrite them. Use the AAP boxes when AAP is where your team already manages the inventory. Either is fine; the rule is **one setting, one place**.

3. Who wins when the same setting is in two places?

Strongest first. This order was **proven by running a test playbook** on ansible-core 2.16 and 2.21 with the same value set in every place and removing the winner one place at a time (the same result on both versions).

#	Place	Notes
1	**Extra variables**: job template *Variables*, a survey answer, a prompt on launch, `-e` on the command line	beats everything below, including the built-in protection further down
2	`playbooks/host_vars/<host>.yml`	this project, one host
3	the **host's** Variables box in AAP	
4	`playbooks/group_vars/<group>.yml`	a child group beats its parent group
5	`playbooks/group_vars/all.yml`	
6	the **group's** Variables box in AAP	a deeper group beats its parent
7	the **inventory's** Variables box in AAP	
8	`roles/<role>/defaults/main.yml`	the built-in default

Things that surprise people:

- **A file in this project beats the AAP group and inventory boxes**, even when the box is on a more specific group. If you type a value in an AAP box and see no effect, look for the same setting in `playbooks/group_vars/`. Every job prints a notice, `NOTE - settings in this project's files win over the same setting in AAP's inventory Variables box`, listing which files define which settings. (Silence it with `site_settings_notice: false`.) Not shown for `playbooks/host_vars/`.
- **A host's own Variables box (3) beats a group file (4/5)**, but a `playbooks/host_vars` file (2) beats the host's box.
- **A survey answer or template extra variable beats everything**: it is the one place that can override the rest for a single run.
- Two sibling groups (neither is inside the other) that set the same variable for the same host: do not rely on the result. Put it in a host file, or in one group.
- An inventory **sourced from the project** (`inventories/site/`, SETUP_AAP.md step 5): AAP imports its `group_vars/all.yml` into the inventory's Variables box and each `group_vars/<group>.yml` into that group's box. So they rank as places 7 and 6, and the project's `playbooks/` files beat them. Keep a setting in one of the two, not both.
- **Updated from 0.4.1 or older?** Your `playbooks/group_vars/all.yml` is yours, so the update left it as it was: it still sets `site_act_diagnose_only_groups`, `patch_never_patch_groups`, `patch_never_reboot_groups`, `site_act_provider` and `site_act_models` as real (uncommented) lines. They beat the same settings in AAP's inventory and group boxes. If you want those set in AAP (or the new built-in defaults), put `#` in front of those lines. See the upgrade notes in `CHANGELOG.md` (0.5.0).

Built-in protection: settings that cannot be weakened from a file

Some lists are protected in code, so a mistake in a settings file cannot expose the AAP host:

What	Built in	What your file can do
Never patched, never rebooted	`aap`, `aap_hosts` (`roles/patch/vars/main.yml`)	your `patch_never_patch_groups` / `patch_never_reboot_groups` list **replaces** the shipped one (`sat`, `netapp`, ...): repeat the names you keep. `aap` and `aap_hosts` stay protected whatever you write
ACT never fixes by itself	`aap`, `aap_hosts` (`roles/site_act/vars/main.yml`)	your `site_act_diagnose_only_groups` list **replaces** the shipped one (`sat`, `idm`, `netapp`, ...): repeat the names you keep. `aap` and `aap_hosts` stay protected whatever you write

Only an extra variable (place 1) can override these, and nothing in these docs asks you to do that.

Examples

Everyone: warn earlier on disks. `playbooks/group_vars/all.yml`:

```yaml
check_disk_warn_pct: 80
check_disk_crit_pct: 92
```

One group: the STIG Manager hosts need these containers watched.
`playbooks/group_vars/stigman.yml` (shipped; already has `watch_containers`).

One host: a database server whose audit partition must stay emptier.
`playbooks/host_vars/db01.example.mil.yml` (the file name is the host's name in AAP):

```yaml
check_disk_overrides:
  /var/log/audit: {warn: 70, crit: 85}
```

One template: the "Troubleshoot" template always uses Ask Sage. In the template's **Variables**:

```yaml
site_act_provider: asksage
```

A list you want to extend. Where the role has a `..._extra` setting, use it: it *adds*.

```yaml
check_services_required_extra: [tomcat]        # required on top of the built-in five
```

Where it has none, look up the default (`check_services_required` is `[sshd, crond, rsyslog, chronyd, auditd]`) and write the whole list with your addition:

```yaml
check_services_required: [sshd, crond, rsyslog, chronyd, auditd, tomcat]
```

How to verify a setting is in effect

1. **Run the check on one host**: **Health check** with `target` = that host and `health_checks` = the check. The finding text quotes the limit it used (`/var is 96% full (critical at 92%)`), so you see the value that won.
2. **Break it on purpose** (or lower the limit) on a test host, and confirm the finding appears. A check that never fires proves nothing.
3. **Read the job's notice** (`NOTE - settings in this project's files ...`): it names every settings file that set something, so you can spot a file that overrides an AAP box.
4. **From the command line** (optional, in a terminal with Ansible and the repository): create a temporary file `playbooks/show_settings.yml`, **do not commit it**, and run it:

YAML

```yaml
- name: Show the settings a host gets
  hosts: "{{ target | default('all') }}"
  gather_facts: false
  tasks:
    - ansible.builtin.import_role: {name: check_disk}    # the role that owns the setting
      when: false                                        # loads its defaults, runs nothing
    - ansible.builtin.debug:
        msg: "{{ item }} = {{ lookup('vars', item, default='(not set)') }}"
      loop: [check_disk_warn_pct, check_disk_crit_pct]
```

shell

```
ansible-playbook -i inventories/site/hosts.yml playbooks/show_settings.yml -e
  target=db01.example.mil
```

It shows the value from the role default plus the files in this project. (Values typed into AAP's Variables boxes exist only inside AAP; there, use steps 1-3.)

If a setting does nothing

Symptom	Likely cause
You set it in AAP and nothing changes	the same setting is in `playbooks/group_vars/`: that file wins (table above)
You set it in a file and nothing changes	the file name does not match the group or host name exactly; a spelling or capital letter
A list lost entries	a list replaces the default: repeat the entries you keep
YAML error at the start of the job	indentation (spaces, never tabs), a missing space after `:`, or a value with `:` inside it that needs quotes
A survey answer is ignored	the survey question's **variable name** is not exactly the setting's name
`patch_serial` / `win_patch_serial` in a settings file changes nothing	Ansible picks how many hosts run at the same time before it reads any settings file: put it in the template's Variables, or a survey

2. Every setting and its default

This page lists **every setting a role reads**, with the value it has when you set nothing, and what the role's own comment says. It is generated from `roles/<role>/defaults/main.yml`, so it is never out of date. Where to put a setting, and which place wins when it is in two, is in VARIABLES.md. Values are shortened to one line here: the role file has the full one.

A setting that is a list (`[a, b]`) **replaces** the default list when you set it; it does not add to it. Repeat the entries you want to keep.

Contents: aap_ee, aap_email_question, aap_health, aap_template_audit, ad_lockout, ad_report, api_availability, cert_inventory, check_accounts, check_auditd, check_certs, check_containers, check_disk, check_fapolicyd, check_logging, check_mariadb, check_mounts, check_network, check_patching, check_performance, check_selinux, check_services, check_sso, check_stig, check_time, check_trust, evaluate_stig_update, fapolicyd, fapolicyd_fix, fapolicyd_rules, horizon, incident_docs, log_analysis, mail_relay, net_certs, net_config, net_perf, ontap_report, patch, pki_bundle, poam, podman_discover, service_watch, servicenow, servicenow_insights, servicenow_war, site_act, site_email, site_findings, stigman_ckl_import, stigman_health, stigman_stack, stigman_update, tenable_scan, tmp_cleanup, troubleshoot, ucs_report, vmware_ask_stack, vmware_vm, win_act, win_check_accounts, win_check_audit, win_check_certs, win_check_disk, win_check_eventlog, win_check_network, win_check_patching, win_check_performance, win_check_security, win_check_services, win_check_stig, win_check_time, win_check_trust, win_patch, win_troubleshoot, settings of one host

aap_ee

File: `roles/aap_ee/defaults/main.yml`

> **Note**
> SET THE EXECUTION ENVIRONMENT (playbooks/aap_ee.yml): puts one execution environment on many job templates at once, through AAP's API - after you add a new EE to AAP (docs/EXECUTION_ENVIRONMENT.md, chapter 6). Run it with Check first: it lists every template it would change, from which EE to which. A rebuild pushed with the next tag: aap_ee_image gives the EE that image (or change its Image in AAP by hand), and every template on it follows.

Setting	Default	What it does
`aap_ee_name`	(empty)	The execution environment to put on the templates: its name in AAP (Automation Execution > Infrastructure > Execution Environments), or its ID when two have that name.
`aap_ee_image`	(empty)	Its new Image (registry/name:tag): a rebuild's next tag, pushed - every job template on this EE then runs it; none need be picked below. The image it has now is refused (an execution node that has a tag keeps it). "" = as it is.
`aap_ee_project`	(empty)	Which job templates (at least one of the three, or only aap_ee_image; never every template in AAP): the templates of this AAP project (its name): your site-automation project

Setting	Default	What it does
aap_ee_from	[]	Only the templates on one of these execution environments now (names or IDs); none = those with no execution environment set, which use the project's or organization's default (left as they are otherwise). Empty = any.
aap_ee_templates	[]	more job templates, by name or ID - whatever they are on
aap_ee_exclude	[]	templates left out, by name or ID (jobs that stay on another EE)
aap_ee_ca_path	(empty)	a CA file for AAP's HTTPS certificate ("" = what the execution environment trusts)

aap_email_question

File: roles/aap_email_question/defaults/main.yml

Note
THE "ALSO EMAIL THE REPORT TO" QUESTION (playbooks/aap_email_question.yml): adds a text box to the survey of every report template and of the workflows that run them, through AAP's API, so whoever launches one can send that run's report to more people (report_email_also, added to the template's report_email_to). Run it once, and again after you add report templates. docs/EMAIL_REPORTS.md, "Also email the report to".

Setting	Default	What it does
aap_email_question_state	present	present = add the question; absent = take it out again
aap_email_question_playbooks	- aap_health_check.yml - aap_template_audit.yml - tenable_scan.yml - act_fix_approved.yml…	Which templates: the job templates whose playbook is one of these (the file name - your template names do not matter). Every playbook that can email a report is here; a test keeps the list complete.
aap_email_question_templates	[]	more job templates or workflows, by name or ID
aap_email_question_exclude	[]	templates left out, by name or ID
aap_email_question_workflows	true	also the workflows that run them (a job in a workflow is never asked its own survey)
aap_email_question_project	(empty)	only job templates of this AAP project (its name); empty = any project

Setting	Default	What it does
aap_email_question_ca_path	(empty)	a CA file for AAP's HTTPS certificate ("" = what the execution environment trusts)
aap_email_question_text	Also email the report to	The question, as whoever launches a template sees it. A text box of several lines (a pasted list keeps its line breaks), not required, empty by default.
aap_email_question_description	(worked out by the role; see its file)	The help text under the question.

aap_health

File: roles/aap_health/defaults/main.yml

Note
AAP HEALTH CHECK (playbooks/aap_health_check.yml): a read-only look at a containerized AAP 2.5+ host - the OS, the rootless containers and their systemd user units, the logs, the controller's database, the controller and gateway API, receptor and the certificates. Every check gives PASS, WARN or FAIL with a detail and a hint; an HTML and a JSON report are written on the host. Nothing is restarted, changed or installed. docs/AAP_HEALTH_CHECK.md

Setting	Default	What it does
aap_health_user	ansible	The account the AAP containers run as (rootless podman, systemd --user units).
aap_health_inventory_file	(empty)	The installer's inventory file (DB name and user, certificate paths are read from it). Empty = found by name (aap_health_inventory_name) in that account's home, at most aap_health_inventory_depth folders down.
aap_health_inventory_name	inventory-growth	the installer inventory file name looked for
aap_health_inventory_depth	2	how many folders down from the home it is looked for
aap_health_unit_regex	'^(automation-.*|postgresql.*|redis.*|receptor.*)\.service$'	The systemd user units that make up AAP (each runs the container of its name, or the --name in the unit).

Setting	Default	What it does
aap_health_cmd_timeout	60	seconds any one command may take (podman, psql, ausearch...) before it is given up
aap_url	https://localhost	Controller and gateway API. aap_username / aap_password: vault them, or attach a "Red Hat Ansible Automation Platform" credential (its CONTROLLER_USERNAME / CONTROLLER_PASSWORD are read). The calls are made from the AAP host.
aap_username	"{{ lookup('ansible.builtin.env', 'CONTROLLER_USERNAME') or lookup('ansible.builtin.env',…	the API login (an AAP administrator or auditor); vault it
aap_password	"{{ lookup('ansible.builtin.env', 'CONTROLLER_PASSWORD') or lookup('ansible.builtin.env',…	its password; vault it
aap_token	"{{ lookup('ansible.builtin.env', 'CONTROLLER_OAUTH_TOKEN') or lookup('ansible.builtin.en…	an API token instead of the username and password (sent as a Bearer token)
aap_health_validate_certs	false	Check the gateway's HTTPS certificate. Its own setting, not aap_validate_certs: an AAP credential on the template sets that one (its "Verify SSL"), and a certificate never names localhost - with aap_url https://localhost it can only fail.
aap_health_api_timeout	60	seconds for one API call
aap_health_api_slow_s	5	an inventory's hosts/?page_size=100 slower than this (or not 200): FAIL
aap_health_api_time_inventories	[]	inventory IDs whose hosts page is timed; [] = every inventory (up to aap_health_api_time_max)
aap_health_api_time_max	50	inventories timed at most
aap_health_focus_inventories	[]	inventory IDs always reported (db and api), even when fine, e.g. [8]

Setting	Default	What it does
aap_health_api_light_warn_s	1.0	the smallest requests (ping, me, config) slower than this: WARN
aap_health_api_light_fail_s	5.0	slower than this (or not 200): FAIL
aap_health_stuck_minutes	15	pending or waiting longer: FAIL
aap_health_running_hours	4	running longer: WARN
aap_health_cleanup_max_days	8	a cleanup system job (activity stream, job details) not run for longer: WARN
aap_health_capacity_warn_pct	10	an instance with less capacity left: WARN
aap_health_last_seen_minutes	10	an instance not seen for longer: WARN
aap_health_mem_warn_mb	2048	OS MemAvailable below: WARN
aap_health_mem_fail_mb	1024	below: FAIL
aap_health_swap_warn_pct	25	swap used at or above: WARN
aap_health_swap_fail_pct	75	at or above: FAIL
aap_health_load_warn	1.0	5-minute load per CPU
aap_health_load_fail	2.0	at or above: FAIL
aap_health_disk_warn_pct	80	space and inodes
aap_health_disk_fail_pct	90	at or above: FAIL
aap_health_disk_paths	[/, /home, /var, /tmp]	also the podman storage and the postgres data folder, found on the host
aap_health_time_warn_s	1.0	clock offset (chronyc tracking)
aap_health_time_fail_s	5.0	at or above (or not synchronised): FAIL
aap_health_selinux_users_ok	[unconfined_u, staff_u, sysadm_u]	SELinux users rootless podman works as
aap_health_dnf_recent_days	7	a dnf transaction this recent, and fapolicyd denials today: WARN
aap_health_fapolicyd_required	false	fapolicyd not running: INFO; true (a site that must run it) = FAIL

Setting	Default	What it does
aap_health_fqdn	(empty)	the host's name to resolve; empty = hostname -f
aap_health_container_mem_warn_mb	4096	Containers one container using more: WARN
aap_health_container_mem_fail_mb	8192	one container using more: FAIL
aap_health_min_uptime_minutes	10	a container up for less: WARN (restarted?)
aap_health_restart_warn	3	restarts (podman or systemd) at or above: WARN
aap_health_dangling_warn	10	dangling images or volumes at or above: WARN
aap_health_log_since	2h	Logs podman logs --since (2h, 30m, 1d)
aap_health_log_proxy_max_lines	50000	gateway proxy lines read at most (the last ones)
aap_health_log_match_max	2000	matching lines kept per container at most
aap_health_envoy_regex	(empty)	the proxy's access-log line, if not Envoy's default or JSON: named groups path, code, dur, up
aap_health_slow_ms	5000	a proxied request slower than this (upstream time): listed
aap_health_slow_max	15	slowest requests listed
aap_health_5xx_fail	20	5xx responses in the window at or above: FAIL (any: WARN)
aap_health_log_web_patterns	HARAKIRI: HARAKIRI worker killed: worker [0-9]+ killed|killed.*worker |SIGKILL|signal 9 Tr…	controller-web: name -> regex (case counts; awk-compatible)
aap_health_log_web_fail	[HARAKIRI, worker killed]	any of these: FAIL (others: WARN; a worker uWSGI only recycled: INFO)
aap_health_log_task_patterns	dispatcher error: "(ERROR|Exception).*dispatch|dispatch.*(ERROR|Exception)" callback rece…	controller-task: name -> regex (all: WARN)
aap_health_log_other_patterns	Traceback: Traceback ERROR: '\[ERROR\]| ERROR |^ERROR|level=error'	the gateway, hub and EDA containers: lines logged at level ERROR, Tracebacks

Setting	Default	What it does
aap_health_proxy_match	gateway-proxy	container names (regex) of the gateway proxy, controller web and task
aap_health_web_match	controller-web	the controller web container
aap_health_task_match	controller-task	the controller task container
aap_health_other_log_match	'^automation-(gateway|hub|eda)'	the containers whose Tracebacks and ERRORs are counted
aap_health_log_context_lines	5	lines kept after each matching line (an ERROR's message is often on the next lines)
aap_health_log_error_groups	5	the containers' ERRORs grouped by message: this many listed, with counts
aap_health_log_oom_journal	true	as root (become; quietly skipped without it): the kernel's OOM kills in the last 2 h
aap_health_log_oom_match	uwsgi	an OOM-killed process matching this (regex) makes controller-web FAIL
aap_health_pg_container	(empty)	Database (podman exec into the postgres container; peer login, else the installer inventory's controller_pg_password) empty = the container whose name matches aap_health_pg_match
aap_health_pg_match	postgres	the postgres container's name (regex)
aap_health_pg_data_dest	/var/lib/pgsql/data	the data folder inside the container (its host folder is found from the mount)
aap_health_db_statement_timeout	120	seconds per query
aap_health_db_vars_avg_warn_kb	10	host variables per inventory: average / largest host
aap_health_db_vars_avg_fail_kb	50	average above: FAIL
aap_health_db_vars_max_warn_kb	100	the largest host above: WARN

Setting	Default	What it does
aap_health_db_vars_max_fail_kb	500	above: FAIL
aap_health_db_host_jobs_warn	5000	job host summaries of one host above (the Hosts list reads each host's recent jobs): WARN
aap_health_db_host_groups_warn	100	groups one host is in, above: WARN (an inventory source's keyed groups can put a host in hundreds)
aap_health_db_dead_min	10000	a table with more dead rows than this (and 20% of its live ones): WARN (autovacuum behind)
aap_health_db_facts_avg_warn_kb	100	ansible_facts per host, average
aap_health_db_size_warn_gb	50	the database larger: WARN
aap_health_db_conn_warn_pct	80	of max_connections
aap_health_db_conn_fail_pct	95	at or above: FAIL
aap_health_db_long_s	30	a query running longer: WARN (aap_health_db_long_fail_s: FAIL)
aap_health_db_long_fail_s	300	a query running longer: FAIL
aap_health_db_idle_tx_s	60	idle in transaction longer: WARN
aap_health_db_jobevent_warn	20000000	job event rows (estimate) above: WARN
aap_health_db_job_keep_days	130	the oldest job older than this: WARN (Cleanup Job Details keeps 120 days by default)
aap_health_receptor_container	(empty)	Receptor empty = the container whose name matches aap_health_receptor_match
aap_health_receptor_match	receptor	the receptor container's name (regex)
aap_health_receptorctl_containers	[]	where receptorctl is run (the first that has it); [] = the receptor container, then controller task and web

Setting	Default	What it does
aap_health_receptor_socket	(empty)	empty = read from the container's receptor.conf (control-service filename)
aap_health_cert_dirs	[aap]	Certificates: the installer inventory's *_tls_cert files, the .cert/.crt/.pem files the containers mount, and these folders (in the AAP account's home) searched, plus aap_health_cert_files
aap_health_cert_files	[]	more certificate files on the AAP host (full paths)
aap_health_cert_warn_days	30	expires in fewer days: WARN
aap_health_cert_fail_days	7	in fewer days (or expired): FAIL
aap_health_trust_cas	[DoD Root CA 3, DoD Root CA 4, DoD Root CA 5, DoD Root CA 6]	labels that must be in the system trust store; [] = not checked
aap_health_report_dir	(empty)	Output on the AAP host; empty = <the AAP account's home>/aap-health
aap_health_top	10	top offenders listed per check
aap_health_email	true	Email the report like the other reports: to report_email_to, through the mail relay in all.yml (report_email_from, report_email_smtp_host ...; docs/EMAIL_REPORTS.md) - FAIL and WARN first, then every area, the sortable HTML attached. Nothing is sent while report_email_to is empty. false = never emailed (the launch box... (all of it: the role file)
aap_health_web	true	Put the report on the report website as well (report_web_dest in all.yml, like the other reports; docs/WEB_REPORTS.md): FAIL and WARN first, then every area. Nothing happens while report_web_dest is empty. false = never on the website.

aap_template_audit

File: `roles/aap_template_audit/defaults/main.yml`

> **Note**
> AAP - TEMPLATE VARIABLES AUDIT (playbooks/aap_template_audit.yml): reads every job template and workflow template through AAP's API and says which settings in their Variables box belong somewhere else - connection settings, addresses, site-wide settings, one set the same way in many templates, secrets (their names only, never their values) - and where. Read-only. docs/TEMPLATE_AUDIT.md

Setting	Default	What it does
`aap_template_audit_project`	(empty)	Which templates: every job template and workflow the AAP user can see, or these narrowed only the job templates of this AAP project (its name); workflows are all read
`aap_template_audit_workflows`	`true`	read the workflow templates too (false = the job templates only)
`aap_template_audit_exclude`	`[]`	templates left out, by name or ID
`aap_template_audit_repeat_min`	`3`	A setting with the same value in this many templates or more is flagged: one change would have to be made in each (0 = never)
`aap_template_audit_site_wide_roles`	`[site_email, site_act, site_findings]`	Site-wide: the settings of these roles are the ones "usually in all.yml" (as aap/extra_vars marks them), read from their defaults when the job runs - and these more
`aap_template_audit_site_wide`	`[]`	more site-wide settings, by name
`aap_template_audit_per_enclave`	`[]`	Settings that differ by enclave: each enclave has its own AAP, and they go in that AAP's all.yml (docs/ENCLAVES.md)
`aap_template_audit_per_template`	`[report_email_to, report_email_cc, report_email_also, site_fail_on, target]`	Settings that belong in a template (who gets its report, what it runs on): never flagged as site-wide or as an address - still flagged when they are a secret, or the same in many templates
`aap_template_audit_ignore`	`[]`	settings never flagged

Setting	Default	What it does
aap_template_audit_ca_path	(empty)	a CA file for AAP's HTTPS certificate ("" = what the execution environment trusts)
aap_template_audit_fail	false	true = the job fails when something is flagged (for a workflow's ticket step)

ad_lockout

File: roles/ad_lockout/defaults/main.yml

> **Note**
> Account lockout (playbooks/ad_lockout.yml): is the account locked out, and WHAT is locking it out - the computers its bad passwords came from, from the Security log of every domain controller. Read-only, unless ad_lockout_unlock. A PowerShell script on every DC (WinRM, the Windows Machine credential): the account needs to read the DCs' Security logs - see docs/ACTIVE_DIRECTORY.md, "Account lockout".

Setting	Default	What it does
ad_lockout_user	(empty)	The account to look at, as the service desk is given it: jdoe, CORP\jdoe or jdoe@corp.example.mil (a survey question, NOT required). Several: one per line, or separated by commas. "" = every account locked out now, each with the computer locking it out.
ad_lockout_hours	24	the Security log entries of the last this-many hours (raise it for a lockout of last week)
ad_lockout_max_users	25	at most this many accounts (asked for, or locked out now) in one run
ad_lockout_unlock	false	true = unlock the account (on the PDC emulator) after reading - only an account asked for by name, and only when it is locked out. A dry run (Check) says what it would unlock. It locks out again as soon as the source computer tries its old password again: the report says what to fix first.
ad_lockout_check_sources	false	true = also look on each source computer that is a host of this job's inventory (Windows, WinRM, the same credential - it needs administrator rights there): services and scheduled tasks running as the account, its sessions (a disconnected Remote Desktop session), processes and mapped drives. Read-only.
ad_lockout_check_max	5	look on at most this many source computers per account (the worst first)
ad_lockout_exchange_servers	[]	Your mail servers (Exchange / ActiveSync, by name or address, * and ? allowed, e.g. [EXCH*, 192.0.2.25]): bad passwords coming through one of them are a phone or a mail app with the old password.

Setting	Default	What it does
ad_lockout_max_events	2000	---- How much is read ---------- read at most this many Security log entries per account and DC (the newest)
ad_lockout_resolve_max	50	resolve at most this many addresses to names per DC (reverse DNS)
ad_lockout_max_rows	50	rows per table; the job's artifacts (ad_lockout) have all

ad_report

File: roles/ad_report/defaults/main.yml

Note
Active Directory health report (playbooks/ad_health_report.yml). Read-only: a PowerShell script on every domain controller (WinRM, the Windows Machine credential) and one more on one DC per domain - the account needs to read AD and run dcdiag / repadmin on the DCs; or, with no such account, reader mode from the AAP side as a plain domain user (ad_health_mode). See docs/ACTIVE_DIRECTORY.md.

Setting	Default	What it does
ad_health_mode	auto	---- Mode (1.4.0) ---------- admin: a PowerShell script on every DC over WinRM (an account that may log on to the DCs, Domain Admins). reader: no WinRM - read from the AAP side over LDAP, DNS and Kerberos as a plain domain user with nothing granted;... (all of it: the role file)
ad_reader_domain	(empty)	---- Reader mode ---------- the AD domain's DNS name ("" = learned: the DCs' names, their ansible_host, their rootDSE, the EE's DNS search domain)
ad_reader_user	(empty)	"" = the job's Machine credential (UPN svc@CORP.EXAMPLE.MIL, CORP\svc, or a DN)
ad_reader_password	(empty)	"" = the Machine credential's (a simple bind over LDAPS, else StartTLS on 389 - never in clear)
ad_reader_validate_certs	true	the LDAPS / StartTLS certificate must be trusted and name the DC before the password is sent
ad_reader_ca_file	(empty)	a CA file ("" = what the execution environment trusts: ee/*/ca)
ad_reader_kerberos	auto	auto: a TGT from each KDC with kinit (the Windows EE has it), else an AS-REQ; tgt; probe (AS-REQ only); off
ad_reader_dns_servers	[]	DNS servers asked for the SRV and A records ([] = the execution environment's)

Setting	Default	What it does
ad_reader_ports	{ldap: 389, ldaps: 636, kerberos: 88, dns: 53}	the ports the DCs answer on, and the DNS servers'
ad_reader_endpoints	{}	a DC reached at another address or port (NAT, a test): {dc01.corp.example.mil: {address: 192.0.2.11, ldaps: 636}}
ad_reader_timeout	5	seconds per connection
ad_reader_clock_warn_seconds	120	a DC's clock against the others' (rootDSE currentTime) -> amber
ad_reader_clock_crit_seconds	300	-> red (Kerberos fails at 5 minutes)
ad_reader_history	true	The privileged groups' members of the last run, to say who joined or left since (on top of the replication metadata, when the account may read it): this template's previous job's artifact in AAP - needs the "Red Hat Ansible Automation Platform" credential on the template - or, with ad_reader_history_file set, that fil... (all of it: the role file)
ad_reader_history_file	(empty)	e.g. /var/lib/site-health/ad-privileged-groups.json - "" = the AAP job artifact
ad_report_hours	24	---- Changes ---------- the changes of the last this-many hours (match it to the schedule: a daily job = 24)
ad_event_hours	24	the event log entries of the last this-many hours
ad_security_events	true	read each DC's Security log for WHO made a change (needs auditing of account and security group management)
ad_max_events	2000	read at most this many entries per event log and DC
ad_changes_max	500	look at most at this many changed objects per kind (created, changed accounts, GPO links)
ad_privileged_groups_extra	[DnsAdmins]	Privileged groups watched besides the built-in ones (Domain / Enterprise / Schema Admins, Administrators, Account / Server / Print / Backup Operators, Group Policy Creator Owners, Key Admins, Cert Publishers): names or SIDs of this domain's groups.
ad_groups_should_be_empty	[]	privileged groups that must have no member, e.g. [Schema Admins] - amber when they have one
ad_privileged_change_severity	warning	a member added to / removed from a privileged group: warning, critical or info
ad_schema_change_severity	warning	the schema version changed: warning, critical or info

Setting	Default	What it does
ad_dc_change_severity	info	a domain controller added or removed: warning, critical or info
ad_dcdiag_skip	[SystemLog]	---- What is read on each DC ---------- dcdiag tests to skip (SystemLog fails on any error in the System log: the Events table covers it)
ad_dcdiag_dns	true	also dcdiag's DNS tests that only read (basic, forwarders, delegation, record registration)
ad_trust_verify	true	verify outbound and two-way trusts (nltest /sc_verify)
ad_repl_warn_hours	6	---- Thresholds: at or above "warn" = amber, at or above "crit" = red ---------- a replication link with no success for this long (inter-site schedules: keep above their interval)
ad_repl_crit_hours	24	no success for this long (or a failing link) -> red
ad_repl_queue_warn	50	replication requests waiting in a DC's inbound queue
ad_time_warn_seconds	5	a DC's clock against the PDC emulator (the PDC: against its NTP source)
ad_time_crit_seconds	60	clock offset -> red (Kerberos fails at 300)
ad_time_sync_warn_hours	24	Windows Time last synchronized longer ago than this
ad_disk_warn_pct	15	free space of the database, log and SYSVOL volumes, % (at or below)
ad_disk_crit_pct	5	free % -> red
ad_cert_warn_days	30	the LDAPS / KDC certificate expires within this many days
ad_cert_crit_days	7	expires within this many days -> red
ad_backup_warn_days	7	the last system-state backup of a partition is older than this
ad_backup_crit_days	30	older than this -> red (always red past the tombstone lifetime)
ad_krbtgt_max_days	180	the krbtgt password is older than this -> amber (the STIG asks for 180)
ad_rid_warn_pct	10	RID pool left, % (at or below)
ad_rid_crit_pct	2	RID pool left -> red
ad_no_client_site_warn	1	this many client addresses in no AD subnet -> amber (fewer: blue)

Setting	Default	What it does
ad_report_max_rows	50	rows per table (problems first); the job's artifacts have all
ad_report_fail	false	The job stays green when it finds problems: the report and email say what. true = mark it FAILED when there is a critical finding or a DC could not be read, for workflows.

api_availability

File: roles/api_availability/defaults/main.yml

Note
AVAILABILITY CHECKS (playbooks/stigman_availability.yml, playbooks/servicenow_availability.yml): every 5 to 15 minutes, each address of the service asked a few times, read-only. A clean run is one line - no email. The job FAILS when the service is down (availability_fail_on), for a workflow's notification or ServiceNow ticket step; its own runs in AAP then give the availability over a day, a week and a month. docs/AVAILABILITY.md

Setting	Default	What it does
availability_attempts	3	times each address is asked in one run
availability_pause_seconds	2	seconds between two tries of one address
availability_timeout	10	seconds a try may take; no answer by then = a bad try
availability_slow_ms	3000	an answer slower than this (median of the tries), in milliseconds -> warning
availability_node_down	warning	One node (one server behind the service's address) that does not answer while the address still does: warning, or critical to fail the job for it too. Every node down with no service address set: critical.
availability_fail_on	[critical]	The job shows FAILED (a workflow's "On failure" path) when a run finds a condition of these severities.
availability_report_on	[critical]	The report is emailed (report_email_to) and published (report_web_dest) only when a run finds a condition of these severities - [critical, warning] also for slow, flaky or one node down; a clean run is one line in the job output.
availability_report_recovery	true	also a report when the service answers again after a failed run ("UP again")

Setting	Default	What it does
availability_history	true	The availability over these periods (hours), from this template's own runs in AAP: a run that ended well = up, a failed one = down. Needs a "Red Hat Ansible Automation Platform" credential on the template (Organization Auditor); without one the check still runs, without the percentages.
availability_history_hours	[24, 168, 720]	the periods: a day, a week, a month

cert_inventory

File: roles/cert_inventory/defaults/main.yml

Note
CERTIFICATE INVENTORY: every certificate report (Linux, Windows, VMware, network devices) can also write what it found - every certificate, with the same fields from each - for a dashboard (Power BI ...). Two ways, each on by itself: files (CSV + JSON) and a SQL Server table. A Check run writes nothing. The fields: id, source, host, location, subject, issuer, sans, not_after, days, status, tier (root / intermediate / ca-signed / self-signed), trust_store, in_use, public_key, sig_alg, serial, sha1, sha256, scanned_at, job (docs/CERT_INVENTORY.md).

Setting	Default	What it does
cert_inventory_files	false	---- files: <source>-certificates.csv and .json (source: linux, windows, vmware, net), replaced each run ---- true = write the files
cert_inventory_dir	/var/lib/site-automation/cert-inventory	where: on the AAP node the job runs on (a path AAP exposes to jobs: docs/CERT_INVENTORY.md), or on cert_inventory_host
cert_inventory_host	(empty)	an inventory host to write them on instead (a file server, or a Linux host with the share mounted); its credential must be on the job template
cert_inventory_history	false	true = also a dated copy (<source>-certificates-YYYY-MM-DD.csv) next to them, for trends
cert_inventory_sql	false	The login comes from the "SQL Server (certificate inventory)" credential (aap/credential_types/cert_inventory_sql.yml). The execution environment needs pymssql and community.general (the all-in-one EE has both). true = add / update the rows

Setting	Default	What it does
cert_inventory_sql_server	(empty)	the SQL Server (name or address)
cert_inventory_sql_port	1433	its port
cert_inventory_sql_database	(empty)	the database (SQL Server 2016 or later: the rows go in as JSON, OPENJSON)
cert_inventory_sql_table	dbo.cert_inventory	schema.table
cert_inventory_sql_create	false	true = create the table (and its view <table>_current) when it is not there - the account then needs CREATE TABLE

check_accounts

File: roles/check_accounts/defaults/main.yml

Note
ACCOUNTS: the account problems auditors and attackers both look for, plus the one that quietly breaks automation - a service account password about to expire.

Setting	Default	What it does
check_accounts_min_uid	1000	people's accounts start here
check_accounts_inactive_days	35	the STIG wants unused accounts disabled after 35 days
check_accounts_ignore_inactive	[]	Accounts that may stay unused (break-glass, service accounts that never log in interactively)
check_accounts_nopasswd_allowed	[]	sudoers entries allowed to use NOPASSWD, e.g. the account AAP connects as: [svc_aap, "%wheel"]
check_accounts_expiry_watch	[]	Accounts whose password/account expiry matters - above all the account AAP logs in with. e.g. [svc_aap]
check_accounts_expiry_warn_days	14	a watched account's password or account expiring within this many days -> warning (3 days or less -> critical)

check_auditd

File: roles/check_auditd/defaults/main.yml

Note
AUDITD: the audit system records who did what (the STIG depends on it). Running, enabled in the kernel, rules loaded, no events lost, room for its log, and (optionally) forwarding. Also how MUCH it records: too much rotates the evidence away within hours, fills the disk (then auditd stops the host or puts it in single-user mode, as the STIG sets it up), drops records on the way to the SIEM, and uses up the SIEM licence.

Setting	Default	What it does
`check_auditd_min_rules`	`5`	fewer loaded rules than this = the STIG rules are missing
`check_auditd_log_warn_pct`	`75`	the audit log's filesystem
`check_auditd_log_crit_pct`	`90`	% used of that filesystem -> critical (warning from check_auditd_log_warn_pct)
`check_auditd_remote_required`	`false`	true = au-remote (audisp-remote) must be active
`check_auditd_exclude_accounts`	`[]`	Accounts whose audit events are expected in bulk and do NOT count toward the volume warning, e.g. the account a vulnerability scanner logs in with (a credentialed scan runs thousands of commands with sudo). Names or UIDs; a host that does not have the account skips it. The report still says how much they made, and the... (all of it: the role file)
`check_auditd_volume_hours`	`24`	look at the audit records of the last N hours
`check_auditd_volume_warn_mb_day`	`200`	warning above this many MB of audit records a day (excluded accounts not counted)
`check_auditd_retention_warn_hours`	`24`	With max_log_file_action ROTATE, auditd keeps num_logs files of max_log_file MB and deletes the oldest. Warning when what is left on disk reaches back fewer hours than this. The STIG wants a week (168) when the records are NOT sent off the host as they are written (RHEL-08-030660). 0 = off.
`check_auditd_volume_timeout`	`120`	The breakdown (which account, which rule, which program) reads the audit log files of the window at low CPU and disk priority, and only when the volume or the retention is over its limit. It stops after this many seconds on a host with very large logs (and says so).

check_certs

File: `roles/check_certs/defaults/main.yml`

Note
CERTS: certificates that expire soon - as files on the host, inside Java keystores, and as served by TLS endpoints (what a browser or client actually sees).

Setting	Default	What it does
check_certs_warn_days	30	a certificate expiring within this many days -> warning (and amber in the certificate report)
check_certs_notice_days	60	the certificate report colours these blue (an early heads-up; not a finding)
check_certs_crit_days	14	within this many days (or expired) -> critical
check_certs_files	[]	Certificate files (PEM) to check. Globs work. Every certificate in a file is checked (a chain file with an expiring intermediate is a problem too). e.g. [/etc/stigman/*.crt, /etc/pki/tls/certs/server.crt]
check_certs_discover	true	Also look for certificates in the usual places. Files with many certificates are CA bundles (hundreds of roots, noisy) and are skipped unless listed in check_certs_files. Text (PEM) and binary (DER) certificates are both read; the names in any case (.crt, .CRT, .Pem ...).
check_certs_discover_dirs	[/etc/pki/tls/certs, /etc/pki/nginx, /etc/nginx, /etc/httpd/conf.d, /etc/stigman, /etc/st…	the folders searched (two levels deep) for *.crt, *.pem, *.cer, *.der (/etc/stigman/nginx/tls: STIG Manager's nginx certificate, three levels down)
check_certs_discover_extra_dirs	[]	more folders searched the same way, added to the ones above (no need to repeat them), e.g. [/home, /root, /opt/app/certs]
check_certs_discover_exclude	'ca-bundle|ca-certificates|/localhost\.crt$'	regular expression: found files whose path matches are skipped
check_certs_discover_max_per_file	10	a found file with more certificates than this is a CA bundle: listed as not checked (list it in check_certs_files to check each)

Setting	Default	What it does
check_certs_placeholder_subject	'(^|,)O=SomeOrganization(,|$)'	Regular expression on the subject of a package's placeholder certificate: postfix's /etc/pki/tls/certs/postfix.pem (O=SomeOrganization), self-signed for a year at install and never renewed. A found one is listed as not checked (no finding, no ticket) and none gets a CSR. A real certificate put in its place is checked;... (all of it: the role file)
check_certs_endpoints	[]	TLS endpoints, checked FROM this host, e.g. the site's own web server: check_certs_endpoints: - {name: STIG Manager, host: 127.0.0.1, port: 443, servername: stigman.example.mil}
check_certs_keystores	[]	Java keystores (Tomcat, the ServiceNow MID Server ...). A JKS keystore lists without a password; a PKCS12 one needs it - give the NAME of an environment variable that holds it (set by an AAP credential), never the password itself: check_certs_keystores: - {path: /opt/servicenow/mid/agent/security/agent_keystore, passw... (all of it: the role file)
check_certs_discover_ports	false	---- Finding them by themselves (the certificate report turns both on; the health check does not) ---- Listening TCP ports that answer TLS: each probed once, on the address it listens on, with `openssl s_client` (check_certs_probe_timeout seconds each, at most check_certs_discover_ports_max ports). Not these ports: pl... (all of it: the role file)
check_certs_discover_ports_skip	[22, 25, 53, 80, 110, 111, 123, 143, 323, 389, 631, 2049, 3306, 5355, 5432, 6379, 11211,…	ports never probed (plain-text services: ssh, smtp, dns, http, ldap, databases ...)
check_certs_discover_ports_max	25	probe at most this many listening ports (the rest are listed as not probed)
check_certs_probe_timeout	3	seconds for each port probe

Setting	Default	What it does
check_certs_discover_keystores	false	Java keystores (*.jks, *.keystore, *.p12, *.pfx) in these folders, and the ones running Java programs use (-Djavax.net.ssl.keyStore=...). A keystore that needs a password is listed as not checked: no password is ever guessed - give it in check_certs_keystores (password_env). The JDK's own trust stores (cacerts) are sk... (all of it: the role file)
check_certs_discover_keystore_dirs	[/etc, /opt, /srv, /usr/share/tomcat, /var/lib/tomcat, /var/lib/pki]	the folders searched for keystores
cert_report_discover	true	The certificate report and the CSR jobs (playbooks/cert_report.yml, cert_renew_csr.yml) switch on the port and keystore search above for their run, and the report the trust store too; the health check's certs check does not. false = those two jobs look only in the folders above and check_certs_files.
check_certs_discover_keystore_depth	5	how many folder levels deep the search goes
check_certs_discover_keystore_exclude	'/(cacerts|jssecacerts)$|/jre/lib/security/|/lib/security/'	regular expression: keystore paths skipped (the JDK's trust stores)
check_certs_discover_keystore_max	40	read at most this many found keystores per host
check_certs_discover_trust	false	The host's trust store: the CA certificates it trusts, put in by update-ca-trust (the DoD roots ...). Every file in these folders, whatever its name. The certificate report lists them with their dates - never a finding, a ticket or a CSR: a CA renews its own certificate (put its new one in the folder and run update-ca... (all of it: the role file)
check_certs_trust_dirs	[/etc/pki/ca-trust/source/anchors]	the folders read (not the folders below them)

Setting	Default	What it does
cert_skip_hosts	[]	---- Leaving hosts out (the certificate report and the renewal; not the health check) ---- Hosts whose certificates something else renews - Satellite and its Capsules (satellite-installer, never by hand), appliances - are left out of the certificate report and never get a CSR: listed as not checked, no ticket. Name th... (all of it: the role file)
cert_skip_groups	[sat]	AAP inventory groups, e.g. [sat] (the health check still watches the Satellite web UI: group_vars/sat.yml)
cert_renew_dir	/etc/pki/tls/renew	---- Renewal: a new key and CSR for each certificate that expires soon (playbooks/cert_renew_csr.yml) ---- The certificates the report flags (within check_certs_warn_days - 30 days or less - or expired) get a new key and a CSR with the same subject and names (subjectAltName), on the host; the CSR is attached to the ce... (all of it: the role file)
cert_renew_skip_hosts	[]	Hosts whose certificates are renewed by hand (IdM, AAP ...): in the certificate report and its tickets as usual - each ticket says it is renewed by hand - but the renewal makes no CSR for them (listed, with why). The same forms as cert_skip_hosts / cert_skip_groups (which leave a host out of the report too). e.g. [idm... (all of it: the role file)
cert_renew_skip_groups	[]	e.g. [idm]
cert_renew_extra_dirs	false	true = a CSR for certificates found in check_certs_discover_extra_dirs (/home, /root ...) too; false = they are in the report only (a file there also under check_certs_discover_dirs, or listed in check_certs_files, is renewed)
cert_renew_key_bits	2048	RSA key size, at least 2048 (a bigger old key keeps its size; an EC certificate gets a new key on the same curve)
cert_renew_digest	sha256	the CSR's signature: sha256, sha384 or sha512
cert_renew_max	10	at most this many NEW CSRs per host in one run (the rest come the next run)

Setting	Default	What it does
cert_renew_only	(empty)	regular expression: only certificates whose place matches - a file's path, a keystore's path, 'TLS port 8443, ...' (empty = every one), e.g. '^/etc/pki/tls/certs/'
cert_renew_issuer	(empty)	regular expression: only certificates whose issuer matches (empty = any CA), e.g. 'O=U\.S\. Government'
cert_renew_exclude	'/(katello|foreman|foreman-proxy|candlepin|pulp|puppet|puppetlabs|rhsm|qpid)/|/etc/pki/(c…	Certificates something else renews - never a CSR from here. The path (Satellite, Capsule, Puppet, subscription-manager on every registered host) and the issuer (Satellite's own CA, Puppet's CA). Check the report's Where column for your hosts' paths and add yours.
cert_renew_exclude_issuer	'O=Katello|Puppet CA'	regular expression on the issuer: Satellite's own CA (O=Katello), Puppet's CA - never a CSR from here
cert_renew_self_signed	false	Self-signed certificates (made on the host by an application or a person, not by a PKI) are reported and get a ticket, but no CSR - unless this is true: then the PKI is asked for a real certificate in their place (a STIG wants PKI certificates). Never a CA's own certificate (a root: it signs certificates), never one o... (all of it: the role file)
cert_renew_install_notes	/etc/stigman/: >- STIG Manager: put the issued certificate (with its chain) and the new k…	What to do with the issued certificate when something else deploys the file (a copy by hand would be overwritten): the start of the file's path -> the next step written on the ticket instead of "install it in place of the file". Add your own.
cert_renew_attach_ext	csr	The CSR's attachment on the ticket: <common name>-<8 characters of the CSR's SHA-256>.<ext> (a new CSR, a new name). An instance that refuses .csr attachments (its allowed extensions, or "Invalid file type: text/plain"): txt names every CSR .txt from the start. Refused, it is tried again as .txt (cert_renew_attach_ret... (all of it: the role file)
cert_renew_attach_type	text/plain	the attachment's Content-Type (a CSR is Base64 text)

Setting	Default	What it does
cert_renew_attach_retry_txt	true	a refused attachment is tried again as <name>.txt (text/plain); false: straight to the work note
cert_install_services	[]	---- Install: the issued certificate in place of the old one (playbooks/cert_install_check.yml -> approval -> playbooks/cert_install_approved.yml) ---- With cert_renew_servicenow: catalog, the PKI puts the issued certificate on the request that carries its CSR. The install check finds it there (the certificate whose p... (all of it: the role file)
cert_install_known_services	[]	The services whose configuration is searched for the certificate's file - [{dir, service, test}] (test: the command that checks the configuration before the reload). Empty = httpd (/etc/httpd, apachectl configtest), apache2, nginx (nginx -t), haproxy (haproxy -c), postfix (postfix check), dovecot (doveconf -n), slapd,... (all of it: the role file)
cert_install_never_reload	[]	services (globs) the job never reloads, e.g. ['slapd*']: their certificates are installed by a person
cert_install_key_files	{}	The old key's file, when it cannot be found - neither in the configuration that names the certificate, nor once in the usual folders (several files with the old key: held, to be named here): {certificate file: key file}, e.g. {/etc/pki/tls/certs/web01.crt: /etc/pki/tls/private/web01.key}
cert_install_key_dirs	[]	where else to look for the old key (empty = /etc/pki/tls/private, /etc/ssl/private, the services' folders ...)
cert_install_trust_check	true	the chain must build up to a root (the host's trust store, or the PKI's root sent with it); false = install without
cert_install_verify_wait	15	seconds the port has, after the reload, to serve the new certificate (then: put back)

Setting	Default	What it does
cert_restore_backup	(empty)	---- Restore: an old certificate back from its backup (playbooks/cert_restore.yml) ---- Each install copies the files it replaces to cert_renew_dir/backup/<date>-<certificate>/. The restore puts them back, reloads the service, checks the port serves the old certificate again, and the renewal waits again (the next inst... (all of it: the role file)
cert_restore_only	(empty)	or: the newest backup of each certificate whose name or subject matches (a regular expression), e.g. 'web01'

check_containers

File: roles/check_containers/defaults/main.yml

Note
CONTAINERS: the podman containers on the host - root's (rootful) and every user's (rootless) - found by themselves (roles/podman_discover: which ones exist, which systemd unit runs each one, which should be running). Read-only. Which containers and users are looked at, and which are left out: the podman_discover_* settings.

Setting	Default	What it does
check_containers_required	[]	Containers that MUST exist and run on these hosts: a container name, or OWNER/NAME for one owner (root/web, alice/web). A name alone matches that container of any owner. Also matches a Quadlet file name without .container, e.g. [stigman-mysql, stigman-api, alice/web].
check_containers_recent_hours	24	A container that stopped with an error (exit code not 0) this many hours ago or less is reported (containers:crashed). Older crashes are history and are not reported.
check_containers_restart_warn	3	More restarts than this (podman's restart count, or its systemd unit's) = a restart loop.

check_disk

File: roles/check_disk/defaults/main.yml

Note
DISK: how full each LOCAL filesystem is (space and inodes), and - from the history of earlier runs - how many days until it is full. Network filesystems (NFS) are the mounts check's job: a hung NFS server can make `df` hang, so this check never touches them.

Setting	Default	What it does
check_disk_warn_pct	85	% used -> warning
check_disk_crit_pct	95	% used -> critical
check_disk_inode_warn_pct	85	inodes: too many small files fill a disk too
check_disk_inode_crit_pct	95	% of inodes used -> critical (warning from check_disk_inode_warn_pct)
check_disk_overrides	{}	Different limits for particular mount points, e.g. keep the audit log partition emptier: check_disk_overrides: /var/log/audit: {warn: 70, crit: 85}
check_disk_ignore	- '^/(proc|sys|dev|run)(/|$)' - '^/etc/(hosts\|hostname\|resolv\.conf)$' # bind-mounted fil…	Mount points to skip (regular expressions).
check_disk_exclude_types	[tmpfs, devtmpfs, squashfs, overlay, iso9660, efivarfs, nsfs]	filesystem types never checked (df -x)
check_disk_forecast	true	Days-until-full forecast. Each run adds a sample (date, bytes used) per mount to a small file on the host; once there is at least a day of history, the growth rate predicts when the filesystem fills. Check mode (a dry run) reads the history but does not add to it.
check_disk_forecast_warn_days	14	full within this many days at the current growth -> warning
check_disk_forecast_crit_days	3	full within this many days -> critical
check_disk_history_file	/var/lib/site-health/disk-history.json	where the forecast keeps its samples, on the host
check_disk_history_keep_days	30	samples older than this are dropped (the forecast uses what is left)
check_disk_timeout	20	seconds for df

Setting	Default	What it does
check_disk_lvm	true	LVM: thin pools (data and metadata) and old-style snapshots, with the limits above (check_disk_warn_pct, check_disk_crit_pct); a snapshot that filled up (no longer valid) is critical. Volume groups' free space is listed in the disk space report, never a finding (a volume group with no free space is normal). Needs root... (all of it: the role file)
disk_report_warn_pct	85	The disk space report (playbooks/disk_report.yml) - not the health check: warning and critical at (% used, space or inodes; LVM thin pools and snapshots too). The report sets check_disk_warn_pct / check_disk_crit_pct from them. % used -> warning, in the disk space report
disk_report_crit_pct	90	% used -> critical, in the disk space report
disk_report_all	false	true = every filesystem and thin pool in the tables, not only the full ones

check_fapolicyd

File: roles/check_fapolicyd/defaults/main.yml

Note
FAPOLICYD (application allow-listing): running, enforcing, configuration and trust database valid, and recent denials - which is what you see when it blocks a program.

Setting	Default	What it does
check_fapolicyd_required	false	true = every host in the play must have fapolicyd. false = only check hosts that have it.
check_fapolicyd_expect_enforcing	true	permissive = 1 means it only logs, never blocks
check_fapolicyd_check_trustdb	true	Compare the trust database with the files on disk (can take a minute on big hosts; files changed outside RPM show up here, e.g. after a manual update).
check_fapolicyd_denials_since	yesterday	ausearch -ts value

Setting	Default	What it does
check_fapolicyd_denials_warn	1	this many denials or more since check_fapolicyd_denials_since -> warning

check_logging

File: roles/check_logging/defaults/main.yml

> **Note**
> LOGGING: are logs kept, and do they reach the central log server (SIEM)?

Setting	Default	What it does
check_logging_journal_persistent	true	the journal must survive a reboot (/var/log/journal)
check_logging_remote	[]	The central log server(s) this host forwards to, e.g. check_logging_remote: - {name: SIEM, host: siem.example.mil, port: 6514}
check_logging_timeout	5	seconds to wait for each log server's TCP port
check_logging_queue_warn_mb	100	rsyslog keeps messages it cannot deliver in /var/lib/rsyslog. A big queue = the log server has been unreachable for a while.
check_logging_errors_since	"-24h"	Error-priority journal messages in the last day. Many = something is complaining a lot.
check_logging_errors_warn	500	this many error messages or more -> warning

check_mariadb

File: roles/check_mariadb/defaults/main.yml

> **Note**
> DATABASE health for MariaDB AND MySQL (read-only; role name kept for compatibility): running, accepting connections, connection headroom, queries that run too long, replication (MariaDB multi-source, MySQL REPLICA/SLAVE), Galera / Percona XtraDB Cluster, MySQL Group Replication, recent errors in its log, and room on its disk. It runs status queries only - no data is read and nothing is changed. The engine and version are detected; see docs/MARIADB.md.

Setting	Default	What it does
check_mariadb_enabled	(worked out by the role; see its file)	Which hosts: members of any of the inventory groups mariadb_hosts, mysql_hosts, database_hosts, mariadb or mysql, or any host that names a container below (or turns discovery on). Set this true/false to override for a host or group.
check_mariadb_container	(empty)	MariaDB or MySQL in podman containers (e.g. the ServiceNow database). The client inside each container is used (podman exec). Name them in one of these ways; all are checked, each as its own instance: check_mariadb_container: mariadb one container check_mariadb_containers: [mariadb, snow-db] several containers on the... (all of it: the role file)
check_mariadb_containers	[]	several containers on one host, e.g. [mariadb, snow-db] (see above)
check_mariadb_container_discover	auto	auto, true or false: find database containers by their image (see above)
check_mariadb_server_paths	[/usr/sbin/mariadbd, /usr/libexec/mariadbd, /usr/sbin/mysqld, /usr/libexec/mysqld, /usr/b…	Where a database server installed on the host lives (any one present = a host install, for auto).
check_mariadb_container_image_regex	'(mariadb|mysql|percona)'	regular expression: a container whose image matches is a database (discovery)
check_mariadb_container_user	(empty)	ROOTLESS podman: the containers belong to a user (e.g. AAP's own MariaDB, or an application account), and only that user can see them. Name the user and the check runs podman AS that user. Empty = root's containers.
check_mariadb_service	(empty)	The systemd unit (host installs). Empty = the first of mariadb, mysqld, mysql that exists.
check_mariadb_client_args	[]	How to log in. Without a credential: as root through the local socket - MariaDB's default for a host install on RHEL (no password needed when the job runs with sudo). A database in a CONTAINER usually gives root a password (the official images do), and MySQL host installs usually do too, so there attach the "MariaDB m... (all of it: the role file)
check_mariadb_conn_warn_pct	80	connections in use, % of max_connections

Setting	Default	What it does
check_mariadb_conn_crit_pct	95	connections in use, % of max_connections -> critical
check_mariadb_threads_running_warn	32	queries executing at this moment
check_mariadb_long_query_seconds	300	a query running longer than this is "long"
check_mariadb_replica_lag_warn	300	seconds a replica is behind its primary
check_mariadb_replica_lag_crit	1800	seconds behind its primary -> critical
check_mariadb_galera_size	0	expected Galera / XtraDB Cluster size (0 = do not check the size)
check_mariadb_group_size	0	expected MySQL Group Replication members (0 = do not check the size)
check_mariadb_restart_warn_minutes	60	restarted less than this ago -> warning
check_mariadb_bufferpool_hit_warn	95	% of reads served from memory (only judged after 1M reads)
check_mariadb_error_since_days	1	[ERROR] lines in its log this recent
check_mariadb_disk_warn_pct	85	the filesystem holding the data directory
check_mariadb_disk_crit_pct	95	% used of the data directory's filesystem -> critical
check_mariadb_performance	true	---- PERFORMANCE (all warnings) ---------- ---------- false = only the health checks above. Counters that count since the server started (slow queries, temporary tables, joins, lock waits, deadlocks ...) are judged as an average since that start, and only once it has run chec... (all of it: the role file)
check_mariadb_perf_min_uptime_hours	2	judge the averages since the start only after this many hours
check_mariadb_slow_per_hour_warn	60	queries slower than long_query_time, per hour
check_mariadb_tmp_disk_pct_warn	25	% of temporary tables that went to disk (judged after 1000 of them)
check_mariadb_full_join_per_hour_warn	100	joins without a usable index (Select_full_join), per hour
check_mariadb_table_lock_wait_pct_warn	1	% of table locks (MyISAM / Aria) that had to wait (judged after 10000)
check_mariadb_row_lock_avg_ms_warn	500	average InnoDB row lock wait, ms (judged after 100 waits)

Setting	Default	What it does
check_mariadb_deadlocks_per_hour_warn	1	InnoDB deadlocks per hour (judged after 10 of them)
check_mariadb_innodb_waits_per_hour_warn	1	waits for a free buffer pool page / a full log buffer, per hour (after 10)
check_mariadb_aborted_connects_per_hour_warn	60	failed connection attempts per hour
check_mariadb_table_open_per_hour_warn	600	tables opened per hour while the table cache is full
check_mariadb_open_files_warn_pct	85	files open, % of open_files_limit
check_mariadb_lock_wait_seconds	30	a session waiting this long for a lock, right now
check_mariadb_long_trx_seconds	600	an InnoDB transaction open this long, right now
check_mariadb_history_length_warn	1000000	InnoDB history list length (old row versions not yet purged)
check_mariadb_sizes	false	---- SIZES (off by default) ---------- true = also read each schema's data and index bytes and the largest tables from information_schema (names and numbers only) and keep them as the fact check_mariadb_sizes_result for the job that asked (the ServiceNow healt... (all of it: the role file)
check_mariadb_sizes_tables	10	the largest tables listed

check_mounts

File: roles/check_mounts/defaults/main.yml

> **Note**
> MOUNTS: filesystems that should be mounted but are not, local filesystems that went read-only (the kernel does that after disk errors), and network mounts (NFS/CIFS) that hang. Diagnose only: this check never mounts, unmounts or remounts anything.

Setting	Default	What it does
check_mounts_network_types	[nfs, nfs4, cifs, smb3]	filesystem types counted as network mounts (checked for hangs and fullness)
check_mounts_local_types	[xfs, ext4, ext3, btrfs, vfat]	filesystem types checked for having gone read-only

Setting	Default	What it does
check_mounts_timeout	5	seconds before a network mount counts as hung
check_mounts_warn_pct	85	Network mounts: how full (same idea as the disk check)
check_mounts_crit_pct	95	% used of a network mount -> critical (warning from check_mounts_warn_pct)
check_mounts_ignore	[]	fstab mount points that are allowed to be unmounted (regular expressions)

check_network

File: roles/check_network/defaults/main.yml

Note
NETWORK: can this host reach what it depends on? A default route, names that must resolve, and TCP ports that must answer (from THIS host, which is what matters when a service breaks).

Setting	Default	What it does
check_network_require_default_route	true	no default route -> finding (false for an isolated host)
check_network_resolve_self	true	this host's own full name (hostname -f) resolves
check_network_dns_names	[]	e.g. [ldap.example.mil, servicenow.example.mil]
check_network_tcp	[]	TCP services this host needs, e.g. check_network_tcp: - {name: LDAP, host: ldap.example.mil, port: 636} - {name: SIEM (syslog over TLS), host: siem.example.mil, port: 6514}
check_network_timeout	5	seconds for each name lookup and TCP port test

check_patching

File: roles/check_patching/defaults/main.yml

Note
PATCHING (readiness, read-only): pending security updates, a reboot still needed after the last updates, and how long ago the host was last updated. Installs nothing - patching itself is playbooks/patch_hosts.yml.

Setting	Default	What it does
check_patching_max_age_days	35	last package update older than this -> warning
check_patching_security_warn	1	this many pending security advisories -> warning
check_patching_critical_is_critical	true	a pending Critical-severity advisory -> critical
check_patching_timeout	300	seconds for dnf to talk to the repositories
check_patching_reboot_check	true	report a reboot still needed (needs-restarting -r, newer kernel installed than running)

check_performance

File: roles/check_performance/defaults/main.yml

> **Note**
> PERFORMANCE: is the host overloaded right now, or was something killed for lack of memory?

Setting	Default	What it does
check_performance_load_warn	1.5	15-minute load average PER CPU (1.0 = every CPU busy)
check_performance_load_crit	3.0	load per CPU -> critical
check_performance_mem_avail_warn_pct	10	memory still available for programs, % of total
check_performance_mem_avail_crit_pct	5	available memory below this % -> critical
check_performance_swap_warn_pct	50	swap in use, % of swap
check_performance_swap_crit_pct	80	swap in use above this % -> critical
check_performance_iowait_warn	20	% of CPU time waiting for disk (average over the sample)
check_performance_iowait_crit	40	I/O wait above this % -> critical
check_performance_steal_warn	10	% of CPU taken by the hypervisor (VM starved)
check_performance_sample_seconds	5	how long vmstat samples
check_performance_oom_since	"-24h"	look for out-of-memory kills this far back
check_performance_zombie_warn	20	this many zombie processes or more -> warning

check_selinux

File: `roles/check_selinux/defaults/main.yml`

> **Note**
> SELINUX: enforcing (the STIG requires it), the config file agrees (so a reboot does not change it), no locally added permissive domains, and recent denials (AVCs) - which is what you see when SELinux blocks an application.

Setting	Default	What it does
`check_selinux_expected_mode`	`enforcing`	the mode SELinux must be in (and the config file must say); anything else -> critical
`check_selinux_avc_since`	`yesterday`	ausearch -ts: recent (10 min), today, yesterday, this-week
`check_selinux_avc_warn`	`1`	this many denials or more -> warning
`check_selinux_booleans`	`{}`	Booleans that must have a value, e.g. {httpd_can_network_connect: "on"}
`check_selinux_relabel_paths`	`[]`	Paths whose file labels are compared with the policy (a dry-run restorecon). Mislabeled files are a common cause of denials after files are copied or moved. Keep it to small trees. e.g. [/etc/nginx, /var/www, /etc/stigman]

check_services

File: `roles/check_services/defaults/main.yml`

> **Note**
> SERVICES: systemd units that failed, services that must be running (and start at boot), and services stuck restarting over and over.

Setting	Default	What it does
`check_services_required`	`[sshd, crond, rsyslog, chronyd, auditd]`	Must be running on every host. A host where one is not installed gets a warning (set check_services_required in the inventory to change the list for a group of hosts).
`check_services_required_extra`	`[]`	More for a group of hosts, e.g. in group_vars/mariadb_hosts.yml: [mariadb]
`check_services_ignore_failed`	`[]`	Failed units to ignore (regular expressions), e.g. a known-broken vendor unit.
`check_services_restart_warn`	`3`	More restarts than this since the unit started = a restart loop

check_sso

File: `roles/check_sso/defaults/main.yml`

> **Note**
> SSO (read-only): can domain users still log in? A Linux host joined to Active Directory or IdM (sssd, realmd, winbind) keeps working on cached logins for a while after it loses its domain - then logins fail. This check finds it before then: sssd (or winbind) running, each sssd domain online, the host's keytab still gets a ticket from a KDC (its machine account works), the AD join is valid (adcli testjoin), domain users resolve, and how often sssd went offline or could not reach a KDC recently. A host that is not joined to any domain is skipped.

Setting	Default	What it does
`check_sso_hours`	`24`	how far back the logs are read (sssd going offline, KDC unreachable)
`check_sso_offline_warn`	`1`	sssd went offline this many times in that window (and is online now) -> warning
`check_sso_kdc_warn`	`3`	"cannot contact any KDC" this many times -> warning
`check_sso_keytab`	`/etc/krb5.keytab`	the host's keytab
`check_sso_keytab_test`	`true`	get a ticket with it (kinit -k into memory: nothing is written) - proves the machine account works
`check_sso_testjoin`	`true`	adcli testjoin on hosts joined to Active Directory with adcli / realmd
`check_sso_test_users`	`[]`	domain users that must resolve (getent passwd), e.g. [svc_monitor@example.mil]; empty = none
`check_sso_sssd_dir`	`/etc/sssd`	sssd's configuration (sssd.conf, conf.d/): there = sssd is set up
`check_sso_smb_conf`	`/etc/samba/smb.conf`	Samba's: security = ads there = winbind is set up for Active Directory
`check_sso_sssd_log_dir`	`/var/log/sssd`	sssd's own logs, read with the journal
`check_sso_timeout`	`30`	seconds for each command that asks the domain (kinit, adcli, getent)

check_stig

File: `roles/check_stig/defaults/main.yml`

Note
STIG (drift): scans the Linux host against the DISA STIG with OpenSCAP - the SCAP Security Guide's STIG profile for its RHEL version (RHEL 9: V2R9, RHEL 10: V1R2 in scap-security-guide 0.1.82) - or with Evaluate-STIG (check_stig_scanner), keeps each scan's report on the host, and compares the result with the host's checklist in STIG Manager (read-only: this check never writes to STIG Manager) and with the host's earlier scans. A rule STIG Manager has as Not a Finding that fails now, and passed in an earlier scan, is drift; one that has failed in every scan is listed as "the scan and the checklist disagree"; a rule that passed in an earlier scan and fails now is new. playbooks/stig_drift_report.yml runs it; the health check runs it only when asked by name (health_checks: [stig] - in no shortcut: a scan takes minutes). docs/STIG_DRIFT.md.

Setting	Default	What it does
check_stig_scanner	openscap	---- The scan, on the host ---------- - openscap: oscap and the SCAP Security Guide (below); evaluate_stig: Evaluate-STIG, already on the host (the settings after these) - closer to checklists Evaluate-STIG filled. Changing it starts the host's... (all of it: the role file)
check_stig_install	false	true: install openscap-scanner and scap-security-guide when missing (dnf, from the host's own repositories); false: a host without them is listed as not scanned
check_stig_oscap	(empty)	the oscap program; empty = found in PATH (/usr/bin/oscap)
check_stig_profile	stig	the profile: stig, stig_gui (a host with a desktop), or a full profile id (xccdf_...)
check_stig_content_dir	/usr/share/xml/scap/ssg/content	where scap-security-guide puts its datastreams (ssg-rhel<major>-ds.xml)
check_stig_datastream	(empty)	another datastream on the host instead (e.g. a DISA SCAP benchmark, with check_stig_profile its profile id); empty = the SCAP Security Guide's for the host's RHEL version
check_stig_datastream_src	(empty)	a datastream on the AAP side (in your repository, from playbooks/: files/stig/ssg-rhel9-ds.xml) copied to the host first - newer content than the host's repositories have; empty = none
check_stig_tailoring	(empty)	a tailoring file on the host (your site's changes to the profile, made with SCAP Workbench); empty = none
check_stig_tailoring_src	(empty)	or a tailoring file on the AAP side (from playbooks/), copied to the host first; empty = none

Setting	Default	What it does
check_stig_timeout	1800	seconds the scan may take (OpenSCAP: one to five minutes; Evaluate-STIG: more)
check_stig_evaluate	(empty)	---- Evaluate-STIG (check_stig_scanner: evaluate_stig) ---------- Evaluate-STIG runs on the host as root with --Output CKLB into a folder of its own (deleted after), and the job reads the checklist of the host's RHEL STIG (check_stig_stigman_benchmarks: RHEL_9_STIG ..... (all of it: the role file)
check_stig_evaluate_select	{}	--SelectSTIG for each RHEL version, e.g. {"8": RHEL8, "9": RHEL9}: only the RHEL STIG is scanned (faster); its short name as your Evaluate-STIG lists it. {} = every STIG that applies to the host (the RHEL one is read)
check_stig_evaluate_args	[]	more arguments for Evaluate-STIG, e.g. ["--ScanType", "Classified", "--AnswerKey", "SITE"] (Evaluate-STIG.ps1: -ScanType ...); --Output and --OutputPath are the job's
check_stig_evaluate_answers	(empty)	Evaluate-STIG's answer files (docs/STIG_DRIFT.md, "Answer files from a Windows share"): copied to each host before it runs, into a folder of this run (removed after it) that Evaluate-STIG is given - nothing kept on the host. From a Windows file share, read over SMB from the AAP side (no mount; the "Windows file share"... (all of it: the role file)
check_stig_evaluate_answers_arg	--AFPath	Evaluate-STIG's option for their folder (the job checks the script has it; .ps1: -AFPath)
check_stig_evaluate_answers_patterns	["*.xml"]	the files taken (any case); sub-folders are not read
check_stig_evaluate_answers_user	(empty)	the share's account (user@YOURDOMAIN.MIL); empty = the "Windows file share" credential's
check_stig_evaluate_answers_password	(empty)	its password (vault it); empty = the credential's
check_stig_evaluate_answers_auth	negotiate	negotiate (Kerberos - FIPS too: user@YOURDOMAIN.MIL, the server's full name - else NTLM), ntlm, or kerberos (a ticket already in the job)
check_stig_evaluate_answers_encrypt	false	true = encrypt (SMB 3) even when the share does not require it

Setting	Default	What it does
check_stig_evaluate_answers_port	445	the share's port
check_stig_evaluate_answers_max_files	500	more files than this in the folder: not read (a wrong folder)
check_stig_evaluate_answers_max_mb	20	a file larger than this: not read
check_stig_evaluate_read	(empty)	read the newest checklist (.cklb or .ckl) Evaluate-STIG already wrote under this folder on the host - its own schedule - instead of running it, e.g. /opt/STIG_Compliance ; empty = run it
check_stig_evaluate_output	CKLB	Evaluate-STIG's --Output: CKLB, or CKL (the older XML checklist) if your version writes CKLB it then rejects itself
check_stig_evaluate_rejected	true	a checklist Evaluate-STIG's own schema check rejects (it moves it to its Bad_CKL folder and names it - a check of its that failed can do that): true = read it from there anyway, said in the report with the checks that failed; false = the host is not scanned
check_stig_pwsh	(empty)	PowerShell 7 for Evaluate-STIG.ps1; empty = pwsh found in PATH
check_stig_dir	/var/lib/site-automation/stig	---- Kept on the host ---------- each scan's HTML report, results (gzip) and summary (JSON), and latest-report.html / latest-summary.json; root-only. A dry run keeps nothing
check_stig_keep	10	how many scans are kept on the host (the oldest are deleted); the last one is what the next scan is compared with
check_stig_keep_report	true	keep oscap's HTML report of each scan (about 4 MB)
check_stig_keep_results	true	keep the XCCDF results of each scan (gzip, about 1 MB): what STIG Manager or eMASS can import

Setting	Default	What it does
check_stig_stigman_api	(empty)	---- STIG Manager (read-only) ---------- -- The job logs in as a Keycloak client with the client credentials grant - your autostig client, or another one - from the "STIG Manager API" credential (STIGMAN_TOKEN_URL, STIGMAN_CLIENT_ID, STIGMAN_CLIENT_SE... (all of it: the role file)
check_stig_stigman_token_url	(empty)	Keycloak's token endpoint (https://keycloak.example.mil/realms/stigman/protocol/openid-connect/token); empty = the credential's
check_stig_stigman_client_id	(empty)	the client, e.g. autostig; empty = the credential's
check_stig_stigman_client_secret	(empty)	its secret (vault it); empty = the credential's
check_stig_stigman_scope	(empty)	scopes to ask for, e.g. "stig-manager:collection:read"; empty = the client's default scopes
check_stig_stigman_validate_certs	true	check the HTTPS certificates of STIG Manager and Keycloak
check_stig_stigman_ca_path	(empty)	a CA bundle on the AAP side to check them with (the DoD roots); empty = the execution environment's
check_stig_stigman_timeout	60	seconds per request
check_stig_stigman_collection	(empty)	Finding each host's asset. Set the collection per group or per host (group_vars, host_vars, or the inventory's variables in AAP): the asset is looked for only there. Empty = every collection the client can read; a host whose name matches assets in two collections is listed and not compared until its collection is set.... (all of it: the role file)
check_stig_stigman_collections	[]	with no collection set: search only these collections ([] = every one the client can read)
check_stig_stigman_asset	(empty)	this host's asset name in STIG Manager; empty = the inventory name, then its short name (web01 for web01.example.mil); the asset's FQDN field matches too

Setting	Default	What it does
check_stig_stigman_benchmarks	{"8": RHEL_8_STIG, "9": RHEL_9_STIG, "10": RHEL_10_STIG}	the STIG Manager benchmark (STIG) for each RHEL version
check_stig_stigman_statuses	[saved, submitted, accepted]	the reviews that count as the checklist (a rejected review does not); [accepted] = only what was accepted
check_stig_stigman_missing_severity	warning	a host not found in STIG Manager (or its collection not readable): warning; critical; or none = listed in the report only
check_stig_severity	{high: critical, medium: warning, low: warning}	---- Findings ---------- the finding for a drifted or new fail, by the rule's CAT (high = CAT I)
check_stig_disagree_severity	none	a rule STIG Manager has as Not a Finding that has failed in every scan of the host (never passed, so nothing moved that the job saw): none = listed in the report only; warning or critical = a finding
check_stig_ignore_rules	[]	rules OpenSCAP is known to get wrong at your site: never a finding, listed in the report as left out. SV-257985, V-257985, RHEL-09-255045 or the SCAP rule's name (sshd_disable_root_login); per group or host too
check_stig_list_max	200	at most this many rows in the report's tables of rules left out and of fails STIG Manager has no review for
check_stig_recheck	[open, not_reviewed]	---- The STIG recheck (playbooks/stig_recheck.yml, win_stig_recheck.yml) ---------- what STIG Manager has for a rule that the recheck checks again: open (a review of Open that counts), not_reviewed (no review that counts, or Not Reviewed); the rest is left to the STIG drift report
check_stig_recheck_carry_forward	false	true = a rule DISA changed since the scan's STIG (held back otherwise) gets the earlier revision's Not a Finding / Not Applicable written to STIG Manager's revision when that review counts (check_stig_stigman_statuses) and the scan agrees; the comment says it was carried forward. For until Evaluate-STIG has STIG Manag... (all of it: the role file)

Setting	Default	What it does
check_stig_update_workflow	false	Workflow 37 - STIG Manager updated, with an approval (docs/STIG_DRIFT.md, "Updating STIG Manager, with an approval"): the recheck's check lists the reviews to write and hands them to the approval step; playbooks/stigman_update.yml writes them after it. This job itself never writes to STIG Manager. true on workflow 37'... (all of it: the role file)
check_stig_update	[fixed, now_na, answered]	what is written once approved: fixed (Open, passes now: Not a Finding), now_na (Open, Not Applicable now), answered (Not Reviewed: as the scan says - a fail is Open)
check_stig_update_status	submitted	the reviews' status as written: saved, submitted, or accepted (the collection must allow it, and the client's grant)
check_stig_update_skip_attachments	true	a review with attachments in STIG Manager (evidence a person attached) is left out of the plan and never written; false = written like the others
check_stig_update_detail_max	8000	at most this many characters of what the scan found, as each review's detail
check_stig_update_create_assets	false	A Linux host with no asset in STIG Manager (0.72.0): with check_stig_stigman_collection set, workflow 37's check scans it in full and plans its asset (named check_stig_stigman_asset, else the short name; the FQDN the inventory name) with its RHEL STIG, and the scan's answers as its reviews; created only after the appr... (all of it: the role file)

check_time

File: roles/check_time/defaults/main.yml

> **Note**
> TIME: is the clock synchronized, and by how much is it off? A clock that is off by minutes breaks Kerberos/AD logins (5 minutes), TLS, log correlation and audit evidence.

Setting	Default	What it does
check_time_offset_warn	0.5	seconds

Setting	Default	What it does
check_time_offset_crit	5.0	seconds off -> critical (warning from check_time_offset_warn)
check_time_min_sources	1	time servers chrony must be able to reach

check_trust

File: roles/check_trust/defaults/main.yml

Note
TRUST (read-only): which trust stores on the host still hold a certificate authority that was compromised or that the browsers dropped - the system bundles, p11-kit's live view and blocklist, Java cacerts, Python's certifi bundles, Node's built-in roots, the bundles inside running podman containers, certificate files an application carries - and the software with its own trust logic (Go programs, Node, NSS) below the version that fixed it. The list of certificates to look for ships with the check (roles/check_trust/vars/main.yml: October 2026, the eight eMudhra intermediates and their roots); add your own in check_trust_extra. The health check runs it in the weekly set; playbooks/trust_report.yml runs it on every host, vCenter and ESXi too.

Setting	Default	What it does
check_trust_extra	[]	More certificates to look for, besides the shipped list: name, sha256 (the certificate's fingerprint, with or without colons), spki_sha256 (its public key, optional), kind (root / intermediate), severity (critical / warning / notice), why. Get the fingerprints with: openssl x509 -in the.pem -noout -fingerprint -sha256... (all of it: the role file)
check_trust_severity	{}	another severity for a listed name: {"AAA Certificate Services": warning} once you decide to blocklist it; notice = listed only, no finding
check_trust_software_severity	warning	the finding for a program below the fixed version (a Go program, Node, NSS): warning or critical
check_trust_bundles	[/etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem, /etc/pki/ca-trust/extracted/openssl/c…	The system: the CA bundles every OpenSSL / Go / Python client reads (the ones that exist are read), p11-kit's live view (trust extract: what NSS and GnuTLS clients see, and the blocklist), the blocklist folders. the system CA bundles (RHEL, Debian / Ubuntu, Alpine, SUSE); the same file under two names is read once

Setting	Default	What it does
check_trust_blocklist_dirs	[/etc/pki/ca-trust/source/blocklist, /usr/share/pki/ca-trust-source/blocklist]	p11-kit blocklist folders: a certificate there is distrusted (listed as blocked, no finding)
check_trust_p11kit	true	read p11-kit's live anchors and blocklist too (trust extract), when the trust command is there
check_trust_nss_dirs	[/etc/pki/nssdb, /root/.pki/nssdb, /home/*/.pki/nssdb]	NSS databases read with certutil (the system's: curl and other NSS clients on older RHEL; root's and each user's: Chrome, Evolution)
check_trust_ca_certificates_fixed	(empty)	the ca-certificates package version that dropped the roots, e.g. '2026.2.82' (empty: not compared - the bundles themselves tell)
check_trust_nss_fixed	'3.131.1'	the NSS version that dropped emSign ECC Root CA - G3 (the package nss / libnss3 below it: a finding; empty: not compared)
check_trust_java	true	Java: every cacerts / jssecacerts under these folders (JDKs, and the keystores of running java programs), read directly (JKS) or with keytool (PKCS12), with the JDK's version.
check_trust_java_dirs	[/usr/lib/jvm, /usr/java, /etc/pki/java, /etc/pki/ca-trust/extracted/java, /opt, /usr/loc…	folders searched for cacerts (seven levels deep)
check_trust_python	true	Python: every certifi bundle (certifi/cacert.pem: requests, pip, urllib3 ... carry their own roots, not the system's) under these folders - system Python, venvs, applications.
check_trust_python_dirs	[/usr/lib, /usr/lib64, /usr/local, /opt, /srv, /var/lib, /root, /home]	folders searched for certifi/cacert.pem (twelve levels deep)
check_trust_node	true	Node: its built-in roots (not the system's: Node carries its own, unless started with --use-openssl-ca), and its version against the one that dropped the root, per major line.
check_trust_node_paths	[/usr/local/bin/node, /usr/local/node*/bin/node, /opt/*/bin/node, /opt/node*/bin/node, /r…	where node binaries are (globs); the one in PATH and running ones are found too

Setting	Default	What it does
check_trust_node_fixed	{"22": "22.23.4", "24": "24.22.1", "26": "26.12.1"}	Node's fixed version per major line (an older line than the oldest listed: unsupported, no fix)
check_trust_go	true	Go: a Go program carries its own certificate verification. Go 1.27.3 / 1.26.10 reject any chain with one of the compromised keys (CVE-2026-111943) - a program built with an older Go has to be rebuilt. The running programs, and the executables under these folders, with the Go version each was built with.
check_trust_go_dirs	[/usr/local/bin, /usr/local/sbin, /opt, /srv]	folders searched for Go programs (four levels deep, executables over 1 MB)
check_trust_go_fixed	{"1.26": "1.26.10", "1.27": "1.27.3"}	Go's fixed version per line (a program built with an older line: unsupported - rebuild it with a fixed Go)
check_trust_go_max	200	at most this many programs looked at per host
check_trust_containers	false	Containers: the CA bundle inside every running podman container (root's and each logged-in user's), read with podman exec - a container image carries its own ca-certificates (Red Hat: UBI images are affected too). Off by default (0.34.0): a shell started in each container that reads CA bundles and searches its whole f... (all of it: the role file)
check_trust_containers_max	30	at most this many containers per host
check_trust_find_dirs	[/etc, /opt, /srv, /usr/local, /var/lib, /root, /home]	Certificate files an application carries (its own bundle, a copied root): every *.pem, *.crt, *.cer, *.der, cacert*, ca-bundle*, ca-certificates*, and every Java keystore (*.jks, *.keystore, *truststore*; *.p12 / *.pfx with keytool and the default password) under these folders, on the host's own filesystems - only the... (all of it: the role file)
check_trust_find_max	2000	at most this many files read per host

Setting	Default	What it does
check_trust_find_exclude	'/(proc|sys|dev)/|/\.cache/|/containers/storage/|/docker/overlay2?/|/ca-trust/extracted/|…	regular expression on the path: matching files are skipped (container image layers, the copies update-ca-trust extracts, Debian's per-CA files - the bundles stand for them; this repository's own test certificates on the AAP controller)
check_trust_timeout	60	seconds for each search and each command (keytool, node, podman exec)
check_trust_remove	false	Removal (the Trust store report only, playbooks/trust_report.yml; the VMware trust store report has the same switch): take the listed certificates out of the stores that hold them. OFF by default - a read-only report. Check mode shows what would go and changes nothing. The system trust is blocklisted (RHEL / SUSE: a P… (all of it: the role file)
check_trust_remove_kinds	[system, java, python, file, nssdb]	which kinds of store to change: system (blocklist), java (every keystore), python (certifi), file (an application's bundle or certificate file), nssdb (certutil -D)
check_trust_remove_severities	[critical, warning]	only certificates at these severities go; a notice-level one (AAA Certificate Services) never, unless named here
check_trust_remove_backup_dir	/var/backups/site-automation/trust	every file changed is copied here first (a folder per run)
check_trust_remove_blocklist_dir	(empty)	the system blocklist folder; empty = the distribution's (/etc/pki/ca-trust/source/blocklist, SUSE /etc/pki/trust/blacklist)
check_trust_remove_update_cmd	(empty)	the command that rebuilds the system trust after blocklisting; empty = the distribution's (update-ca-trust, update-ca-certificates)
check_trust_keystore_password	changeit	The Java keystores' password (keytool's default). A real one belongs in the "Keystore password" credential (KEYSTORE_PASSWORD), attached to the template: it wins over this setting, and stays out of all.yml and the job's variables. keytool gets it in its own environment, never on the command line.

evaluate_stig_update

File: `roles/evaluate_stig_update/defaults/main.yml`

Setting	Default	What it does
`evaluate_stig_update_zip`	`files/evaluate-stig`	EVALUATE-STIG - UPDATE (playbooks/evaluate_stig_update.yml; docs/EVALUATE_STIG_UPDATE.md). The new Evaluate-STIG zip (downloaded by you - its site needs a CAC) is kept in YOUR repository, never the public one, and listed in .site-local. Each host's Evaluate-STIG is updated only when the zip is newer; Check only reads... (all of it: the role file)
`evaluate_stig_update_dest`	(empty)	Evaluate-STIG on the hosts: its folder (or Src/Evaluate-STIG, or the script); "" = check_stig_evaluate - the folder the STIG jobs run it from
`evaluate_stig_update_force`	`false`	true = put the zip's in place even when the host's is the same version or newer (or its version cannot be read)
`evaluate_stig_update_install_missing`	`false`	true = a host without Evaluate-STIG in that folder gets it; false = listed, left as it is
`evaluate_stig_update_keep`	`[AnswerFiles]`	folders of the old Evaluate-STIG whose files are carried into the new one (your own answer files); a file the zip also has: yours
`evaluate_stig_update_backups`	`1`	earlier folders kept beside it (<folder>.previous-<version>); 0 = none
`evaluate_stig_update_smoke`	`auto`	The quick run after the update: auto = Evaluate-STIG_Bash.sh --ListSupportedProducts (else pwsh Evaluate-STIG.ps1 -ListSupportedProducts); none; or your own command ({dir} = the script's folder). It fails: the old one is put back.
`evaluate_stig_update_smoke_timeout`	`600`	seconds the quick run may take
`evaluate_stig_update_tmp`	`/var/tmp`	where the zip is copied on the host before it is unpacked (deleted after)

Setting	Default	What it does
evaluate_stig_update_version_regex	(?:^[\t]*#[\t#]*Version:[\t]*|\$EvaluateStigVersion\s*=\s*["']])([0-9]+(?:\.[0-9]+)+)	Where the version is in Evaluate-STIG.ps1 (its first group): its header line "# # Version: 1.2604.0" (or a $EvaluateStigVersion = "..." line); the zip's file name is tried when it is not there.

fapolicyd

File: roles/fapolicyd/defaults/main.yml

> **Note**
> FAPOLICYD (playbooks/fapolicyd.yml, docs/RUNBOOKS.md "fapolicyd", workflow 36): fapolicyd in one pass on each host - what keeps it down (trust files it cannot read, SELinux labels, db_max_size, not enabled at boot, not running: it starts permissive), what it blocked or would block (a trust entry pinned to the file, a container's rule), its rules files cleaned up, back to enforcing on a quiet day. One line per change to approve; only approved lines are made.

Setting	Default	What it does
fapolicyd_approved	[]	The approved changes, each exactly as the report's "Changes to approve" gives it, e.g. - "trust file access: /etc/fapolicyd/trust.d/app root:root 0600 -> root:fapolicyd 0644" - "enable fapolicyd at boot (it is disabled)" - "permissive: 0 -> 1" (with the restart: a stopped fapolicyd starts blocking nothing) - "restart... (all of it: the role file)
fapolicyd_workflow	false	THE WORKFLOW (docs/WORKFLOWS_AND_SCHEDULES.md, 36): true on its first template, "fapolicyd - check". It reads, changes nothing, hands each host's changes to the approval step and the approved job (set_stats fapolicyd_workflow_plan) and fails on purpose - NEEDS APPROVAL. It holds back for a person: permissive: 0 -> 1 o... (all of it: the role file)
fapolicyd_since	today	blocks since (ausearch -ts): recent (10 minutes), today, yesterday, this-week, boot
fapolicyd_dir	/etc/fapolicyd	fapolicyd's folder: fapolicyd.conf, fapolicyd.trust, trust.d, rules.d
fapolicyd_wait	30	seconds fapolicyd may take to take a change (running, its new rules or trust)
fapolicyd_act	false	ACT (GenAI) explains what fapolicyd blocked, in the report - evidence only, ACT runs nothing on the host. Needs the "ACT model key" credential on the template and site_act's provider and model (docs/ADDING_ACT.md).

fapolicyd_fix

File: roles/fapolicyd_fix/defaults/main.yml

Setting	Default	What it does
fapolicyd_fix_approved	[]	FAPOLICYD - GETTING IT RUNNING (a part of playbooks/fapolicyd.yml, roles/fapolicyd; docs/RUNBOOKS.md "fapolicyd"): why fapolicyd is down, or near its limits, on each host - trust files it cannot read, their SELinux labels, db_max_size too small for the trust database, not enabled at boot, not running (it starts permis... (all of it: the role file)
fapolicyd_fix_dir	/etc/fapolicyd	set by roles/fapolicyd: fapolicyd_dir
fapolicyd_fix_binary	/usr/sbin/fapolicyd	read to see whether this fapolicyd sizes its database itself (db_max_size = auto)
fapolicyd_fix_owner	root	The trust files' owner, group and modes come from the fapolicyd package (rpm); these only when rpm gives none. the trust files' owner, when rpm gives none
fapolicyd_fix_group	fapolicyd	their group, when rpm gives none
fapolicyd_fix_file_mode	"0644"	a trust file's mode, when rpm gives none
fapolicyd_fix_dir_mode	"0750"	the folders' mode, when rpm gives none
fapolicyd_fix_large_mb	20	a trust file larger than this (MB) is listed for a person: every file in it is trusted
fapolicyd_fix_headroom	1.5	the proposed db_max_size: the estimate times this (where auto is not possible)
fapolicyd_fix_journal_lines	400	fapolicyd's journal lines read, this boot
fapolicyd_fix_wait	30	roles/fapolicyd sets it from fapolicyd_wait
fapolicyd_fix_timeout	120	seconds for each command the job reads with

fapolicyd_rules

File: roles/fapolicyd_rules/defaults/main.yml

Setting	Default	What it does
fapolicyd_rules_approved	[]	FAPOLICYD - BLOCKS AND RULES (a part of playbooks/fapolicyd.yml, roles/fapolicyd; docs/RUNBOOKS.md "fapolicyd"): what fapolicyd blocked on each host - or, permissive, would block - and the change that lets each through; the rules files' clean-up; permissive mode on or off. Each change is a line to approve; once approv... (all of it: the role file)
fapolicyd_rules_since	today	set by roles/fapolicyd: fapolicyd_since
fapolicyd_rules_dir	/etc/fapolicyd	set by roles/fapolicyd: fapolicyd_dir
fapolicyd_rules_own_rules	36-site-approved.rules	the rules file (in rules.d) this job writes its allow rules to
fapolicyd_rules_own_trust	site-approved	the trust file (in trust.d) this job writes its trust entries to
fapolicyd_rules_never_edit	[35-site-containers.rules]	rules files (in rules.d) reviewed, never changed (the old containers job's own)
fapolicyd_rules_when_enforcing	false	A clean-up that blocks more (a rule that opened too much, taken out) is made only while fapolicyd is permissive: a program it would break shows as a would-be block, not as an outage. true: made while it enforces too.
fapolicyd_rules_container_labels	'(^|:)(container_[a-z0-9_]*_t|spc_t)(:|$)'	A container's denial (as the containers job tells them apart): an SELinux label on the program or the file, a container program, or a path in container storage.
fapolicyd_rules_container_programs	[podman, crun, runc, conmon, buildah, skopeo, catatonit, containerd, containerd-shim-runc…	container programs

Setting	Default	What it does
fapolicyd_rules_container_paths	['^/var/lib/containers/', '/\.local/share/containers/', '^/run/containers/', '^/run/user/…	container storage
fapolicyd_rules_user_dirs	[/home/, /run/user/, /tmp/, /var/tmp/, /dev/shm/]	folders users can write: a file blocked there gets no trust entry (a person decides)
fapolicyd_rules_hash_max_mb	512	a file larger than this is not hashed: a person adds it (or large, below)
fapolicyd_rules_offer	[update, reload, reinstall, large, podman]	Lines for what only a person could do before - each still a line to approve: update a trusted file that changed since: its trust entry updated (its size and sha256 now) reload a file trusted as it is (or an RPM file that matches its package) that fapolicyd has not loaded: the trust database reloaded (fapolicyd-cli --u... (all of it: the role file)
fapolicyd_rules_large_mb	4096	with large: a file up to this many MB is hashed and gets its trust entry
fapolicyd_rules_max	200	at most this many block groups per host in the report
fapolicyd_rules_wait	30	roles/fapolicyd sets it from fapolicyd_wait
fapolicyd_rules_timeout	120	seconds for each command the job reads with
fapolicyd_rules_storage_conf	/etc/containers/storage.conf	podman's storage.conf: where root's and users' storage is

horizon

File: `roles/horizon/defaults/main.yml`

> **Note**
> VMware Horizon 8 (playbooks/horizon_*.yml): the Connection Server's REST API, from the AAP side, with the "Horizon API" credential (its address, domain, user name and password) - no collection needed. Every job asks each API path at the newest version first, down to v1, so it fits any Horizon 8 release. A dry run (Check) reads and says what it would do; it changes nothing. docs/HORIZON.md

Setting	Default	What it does
horizon_ca_path	(empty)	---- Connection ---------- a CA file for the Connection Server's HTTPS certificate ("" = what the execution environment trusts)
horizon_timeout	60	seconds per request
horizon_page_size	500	records per page when reading pools, desktops and sessions (Horizon's most: 1000)
horizon_api_newest	9	the newest API version to try (v9 down to v1): the first that answers is used
horizon_action_version	v1	the API version of the actions (reset, recover, log off, push image...)
horizon_protected_pools	[]	---- Pools no job changes ---------- Desktop pools no job fixes, logs off or pushes an image to: names as Horizon Console shows them, * and ? allowed, any case (e.g. [VIP-*, Kiosk]). The reports still show them.
horizon_cert_warn_days	30	---- Health report (playbooks/horizon_health_report.yml) ---------- a Connection Server's, gateway's or vCenter's certificate expires within this many days -> amber
horizon_cert_crit_days	7	within this many days (or expired, or not valid) -> red
horizon_license_warn_days	60	the license expires within this many days -> amber
horizon_report_max_rows	50	rows in the problems table (the job's artifacts have all)
horizon_backup_max_hours	26	a Connection Server's last configuration backup older than this -> amber
horizon_datastore_warn_pct	15	a datastore with less free than this, % -> amber (instant clones and new replicas need room)
horizon_datastore_crit_pct	5	less free than this, % -> red
horizon_health_fail	false	The health and capacity reports stay green when they find problems: the report says what. true = mark the job FAILED on a critical finding, for workflows and alerts.

Setting	Default	What it does
horizon_capacity_warn_pct	80	---- Capacity report (playbooks/horizon_capacity_report.yml) ---------- a pool's desktops in use (connected or disconnected) against its maximum, % -> amber
horizon_capacity_crit_pct	95	in use, % -> red (also red: no free desktop left)
horizon_broken_states	[AGENT_UNREACHABLE, AGENT_ERR_DISABLED, AGENT_ERR_INVALID_IP, AGENT_ERR_NEED_REBOOT, AGEN…	---- Broken desktops (the desktops check -> approval -> the fix job) ---------- A desktop in one of these states takes no user until it is fixed. Left out on purpose: states that pass by themselves (being made, customized, deleted, waiting for the agent at start-up, maintenance).
horizon_desktops_mode	report	report = the check lists them and stays green (a schedule). approval = when there is a desktop to fix, the check hands the list on and FAILS on purpose, so the workflow takes its "Run on fail" path to the approval step and then the fix job. Set approval on the workflow's check node.
horizon_fix_default_action	recover	what fixes a broken desktop: recover (instant clones: deleted and made again from the image), reset, restart, remove, rebuild
horizon_fix_actions	{AGENT_ERR_NEED_REBOOT: restart}	By state, when another action fits better (e.g. {AGENT_ERR_NEED_REBOOT: restart, PROVISIONING_ERROR: remove}).
horizon_fix_max	20	fix at most this many desktops per run: the rest are listed, the next run takes them
horizon_fix_exclude_machines	[]	desktops never fixed automatically: names, * and ? allowed (e.g. [W11-TEST-*])
horizon_stuck_hours	4	A desktop left in a passing state (being made, customized, deleted, maintenance, waiting for its agent) for this many hours is stuck: the check lists it. Counted from when it was made (or its last maintenance); a desktop whose time Horizon does not give is not counted.

Setting	Default	What it does
horizon_fix_stuck_actions	{}	What fixes a stuck desktop, by state (e.g. {CUSTOMIZING: recover, WAITING_FOR_AGENT: recover}). Not named = left for a person (the default: a stuck desktop often means vCenter, a datastore or the image).
horizon_fix_ghost_action	remove	A desktop Horizon lists but vCenter no longer has (missing in vCenter): remove its record. "" = left for a person.
horizon_fix_host_threshold	5	This many broken desktops on one ESXi host at once = the host is the problem: all of them are held back, and the check says to look at the host (Broadcom KB 431732). 0 = off.
horizon_fix_ad_gate	true	A Connection Server that cannot reach or bind to an AD domain holds back every recover and rebuild (a new clone could not join the domain): they wait until AD answers. false = off.
horizon_fix_hold_error_kinds	[domain, network, datastore]	A desktop whose clone error is one of these kinds is not recovered or rebuilt - the cause is outside the desktop, and a new clone would fail the same way: domain (join, trust), network (IP, DHCP), datastore (space). The others - timeout, activation, other - are recovered. [] = recover every kind.
horizon_clone_error_kinds	{}	More kinds of clone error, tried before the built-in ones: {kind: a regular expression on Horizon's message}, e.g. {security_agent: "carbon black|tanium"}. A kind named in horizon_fix_hold_error_kinds is held back.
horizon_events_hours	24	---- Events report (playbooks/horizon_events_report.yml) ---------- the events of the last this many hours
horizon_events_failed_warn	5	a user with this many failed logons or more in the window -> amber
horizon_events_failed_types	["BROKER_USER_AUTHFAILED*", "BROKER_USER_NOT_ENTITLED*", "BROKER_USER_NOT_AUTHORIZED*", "…	The event types of a failed logon (* = anything): bad password, locked or disabled account, not entitled...
horizon_events_error_severities	[ERROR, AUDIT_FAIL]	the events listed under Errors

Setting	Default	What it does
horizon_events_admin_modules	[Vlsi]	the event modules of what administrators changed (Horizon Console and its API)
horizon_settings_baseline	{}	---- Settings report (playbooks/horizon_settings_report.yml) ---------- What Horizon's settings should be: {area.setting: value}. area: general, security, pre_logon (Horizon's global settings) or connection_server (checked on each Connection Server). value: the value, or {min: N}, {max: N}, {... (all of it: the role file)
horizon_client_min_versions	{}	---- Sessions report (playbooks/horizon_sessions_report.yml) ---------- The oldest Horizon Client to accept, per client type as Horizon names it (WINDOWS, MAC, LINUX, IOS, ANDROID, HTML_ACCESS...) and '*' for any other, in the numbers Horizon shows for the client (e.g. 8.12.0, or 2312). Older... (all of it: the role file)
horizon_session_pending_minutes	10	a session pending this long: a launch that never finished
horizon_session_long_hours	72	a session started this many hours ago or more is listed (0 = off)
horizon_session_protocols	[BLAST]	the protocols you expect; a session on another is listed ([] = any)
horizon_logon_hours	24	---- Logon times report (playbooks/horizon_logon_report.yml) ---------- the desktop sessions started in the last this many hours are timed
horizon_logon_max	200	at most this many (the newest): one request each
horizon_logon_warn_s	60	a pool whose 95th percentile logon takes this many seconds or more -> amber
horizon_logon_crit_s	120	-> red
horizon_message	(empty)	---- Message users (playbooks/horizon_message_users.yml) ---------- the message (a survey question)
horizon_message_type	WARNING	INFO, WARNING or ERROR: how Horizon Client shows it

Setting	Default	What it does
horizon_message_pools	[]	to the sessions of these pools: names, * and ? allowed
horizon_message_host	(empty)	to the sessions on this ESXi host's desktops
horizon_message_users	[]	to these users' sessions: login names (jdoe or CORP\jdoe) or SIDs
horizon_message_all	false	true = to every session
horizon_maint_host	(empty)	---- Host maintenance (playbooks/horizon_host_maintenance_*.yml) ---------- the ESXi host, as vCenter names it (a survey question on the workflow)
horizon_maint_message	"Your desktop's server goes into maintenance in 15 minutes. Please save your work and log…	What the users on its desktops read (empty = no message). A survey question on the workflow.
horizon_maint_message_type	WARNING	INFO, WARNING or ERROR
horizon_maint_end_all	false	The end job brings back only the desktops the start job put in maintenance mode (handed on in the workflow). true = without that list, every desktop in maintenance mode on the host - also one put there for another reason (as before 1.22.0).
horizon_session_max_disconnected_hours	8	---- Disconnected sessions (playbooks/horizon_session_cleanup.yml) ---------- log off a session disconnected for this many hours or more
horizon_session_cleanup_max	100	log off at most this many sessions per run
horizon_session_cleanup_exclude_pools	[]	pools whose sessions are never logged off: names, * and ? allowed
horizon_session_cleanup_exclude_users	[]	users never logged off: login names (jdoe or CORP\jdoe) or SIDs
horizon_session_forced_logoff	true	log off a locked session too (a disconnected session is usually locked)
horizon_session_cleanup_report_only	false	true = list them, log off none (a dry run - Check - does the same)

Setting	Default	What it does
horizon_user	(empty)	---- User help (playbooks/horizon_user_help.yml) ---------- the user who cannot connect: jdoe, CORP\jdoe, jdoe@corp.example.mil or a SID (a survey question)
horizon_user_action	look	What to do: look (only show their sessions and desktops), logoff (all their sessions), disconnect (their connected sessions), reset or restart (their desktops), recover (their instant-clone desktops: made again from the image). Never a desktop another user has a session on, nor one in a protected pool.
horizon_push_pools	[]	---- Image push (the push plan -> approval -> the push job) ---------- the instant-clone pools to push to: names (a survey: one per line, or separated by commas)
horizon_push_base_vm	(empty)	the golden image VM, as vCenter names it (or its path)
horizon_push_snapshot	(empty)	its snapshot to push: the name, or the path when two have that name (e.g. /2026-10/2026-11)
horizon_push_start	(empty)	When the push starts: "" = as soon as the approved push job runs; or a time - "2026-11-01 22:00" (UTC) or with an offset, "2026-11-01T22:00-05:00". A time that has passed by the approval starts in a minute.
horizon_push_logoff_policy	WAIT_FOR_LOGOFF	WAIT_FOR_LOGOFF (each desktop is replaced when its user logs off) or FORCE_LOGOFF (users are logged off)
horizon_push_stop_on_first_error	true	stop the push at the first desktop that fails
horizon_test_user	(empty)	---- API test (playbooks/horizon_api_test.yml) ---------- also look up this user (e.g. jdoe), as the user-help job does

Setting	Default	What it does
horizon_fslogix_shares	[]	---- FSLogix profiles report (playbooks/horizon_fslogix_report.yml) ---------- The folders that hold the users' FSLogix folders - VHDLocations (and CCDLocations, or an ODFC share), one \\server\share\folder each (a survey: one per line). Read over SMB from the AAP side with the "Windows file share"... (all of it: the role file)
horizon_fslogix_size_mb	30000	the containers' cap: FSLogix's SizeInMBs (its default 30000)
horizon_fslogix_odfc_size_mb	0	the ODFC containers' SizeInMBs; 0 = the same as horizon_fslogix_size_mb
horizon_fslogix_warn_pct	85	a container this full against its cap, % -> amber
horizon_fslogix_crit_pct	95	this full, % -> red (the user gets errors when it fills)
horizon_fslogix_stale_days	90	a container not written for this many days -> stale (space to reclaim, once the user has left)
horizon_fslogix_share_warn_pct	15	a share with less free than this, % -> amber
horizon_fslogix_share_crit_pct	5	less free than this, % -> red
horizon_fslogix_in_use_minutes	30	a container (or its lock file) written this recently -> "likely in use" (a guess from the time)
horizon_fslogix_diff_left_hours	24	a differencing disk (RW.VHDX) not written for this many hours -> left behind, amber
horizon_fslogix_folder_pattern	"%sid%_%username%"	the users' folders: FSLogix's SIDDirNamePattern; "%username%_%sid%" with FlipFlopProfileDirectoryName
horizon_fslogix_profile_name	"Profile_%username%"	the profile container's name, no extension: FSLogix's VHDNamePattern
horizon_fslogix_odfc_name	"ODFC_%username%"	the ODFC container's name, no extension: its VHDNamePattern
horizon_fslogix_ignore	[]	file names to pass over (* and ?, any case), e.g. ["desktop.ini"]
horizon_fslogix_max_folders	20000	at most this many user folders per share (more: said in the report)

Setting	Default	What it does
horizon_fslogix_user	(empty)	the shares' account (user@YOURDOMAIN.MIL); empty = the "Windows file share" credential's
horizon_fslogix_password	(empty)	its password (vault it); empty = the credential's
horizon_fslogix_auth	negotiate	negotiate (Kerberos - FIPS too: user@YOURDOMAIN.MIL, the server's full name - else NTLM), ntlm, or kerberos (a ticket already in the job)
horizon_fslogix_encrypt	false	true = encrypt (SMB 3) even when the share does not require it
horizon_fslogix_port	445	the shares' port
horizon_internal_min_age_hours	24	---- Internal VMs report (playbooks/horizon_internal_vms_report.yml) ---------- An internal VM (cp-template, cp-replica, cp-parent) made less than this many hours ago is never called left over: a new pool or an image push makes them before Horizon shows the new image.
horizon_internal_other_pods	[]	Other Horizon pods that share these vCenters: their cloneprep.client.uuid (each VM's pod is in the job's artifacts). Their internal VMs are listed, never called left over.
horizon_internal_max_rows	100	rows per table in the report (the job's artifacts have all)
horizon_image_vm	(empty)	---- Golden image check (playbooks/horizon_image_check.yml) ---------- Before you snapshot and push a golden image: is the VM ready? Read-only. The vCenter side with the "VMware vCenter" credential; the Windows side over WinRM with the Machine credential, on the host the template's Limit nam... (all of it: the role file)
horizon_image_snapshot_warn	5	the current snapshot more than this many deep in its chain -> WARN
horizon_image_hw_min	14	a hardware version (VM compatibility) older than this -> WARN (a vTPM at the pool needs 14 or later); 0 = not checked
horizon_image_fslogix	true	FSLogix Apps should be in the image: missing -> WARN; false = not checked

Setting	Default	What it does
horizon_image_guest	true	false = the vCenter side only (no WinRM to the golden VM): the report says the guest was not checked
horizon_image_check_fail	false	The check stays green and says READY or NOT READY. true = mark the job FAILED on any FAIL, for workflows.

incident_docs

File: roles/incident_docs/defaults/main.yml

Note
OUTAGE AAR AND SOP - WORD DOCUMENTS (playbooks/incident_aar.yml, playbooks/incident_sop.yml; docs/INCIDENT_DOCS.md). What is typed in the job's survey becomes an after action report (AAR.docx) - and, when asked, an SOP - with the site's logo (report_logo_file) and marking banner, emailed to the addresses typed in the survey (report_email_also). Never on the report website. The survey itself: the job "AAP - incident document surveys" writes it (playbooks/incident_docs_survey.yml).

Setting	Default	What it does
incident_docs_banner	auto	The marking at the top and bottom of every page: auto = the reports' banner (report_banner_text, in its colour) when that banner is on - the classified enclave, or report_banner true (docs/ENCLAVES.md) - else CONTROLLED UNCLASSIFIED INFORMATION (CUI); your own text wins as it is; "" = none.
incident_docs_banner_color	(empty)	Its colour (bold text): "" = the marking's usual one - CUI purple (502B85), UNCLASSIFIED green (007A33), CONFIDENTIAL blue (0033A0), SECRET red (C8102E), TOP SECRET orange (FF8C00) - or, with auto, the reports' banner's; or your own, hex RRGGBB in quotes.
incident_docs_logo_file	(empty)	The logo at the top of the first page: report_logo_file, the reports' logo, unless this names another file (files/aar_logo.png = playbooks/files/aar_logo.png, in YOUR repository). "" = report_logo_file.
incident_docs_logo_height	64	pixels (its width follows)
incident_docs_font	Times New Roman	the documents' font
incident_docs_font_size	12	points
incident_docs_heading_color	"1F2D86"	the headings' colour (hex RRGGBB, in quotes) - dark blue

Setting	Default	What it does
incident_aar_instructions	- Save this file as AAR.docx. - Send it to your manager for review and approval before it…	The site's own instructions at the end of each document, a numbered list (your all.yml: your teams' names and steps). [] = none.
incident_sop_instructions	[]	the same at the end of the SOP; [] = none
incident_sop_status	Draft - review before use	the SOP's Status line
incident_docs_aar_file	AAR	The attachments' names: AAR.docx (AAR-<outage>.docx when the outage is named), SOP-<title>.docx. the AAR attachment's first part
incident_docs_sop_file	SOP	the SOP attachment's first part
incident_impact_levels	[MINOR, MODERATE, MAJOR, CRITICAL]	The impact levels the survey offers (the setup job writes them into the survey).
incident_docs_email_to	[]	Who gets the documents: the addresses typed in the survey (report_email_also - held to report_email_also_domains when that is set), plus these. The reports' usual recipients (report_email_to) are NOT used here: an AAR goes only where it is sent. Nor is report_email_cc.
incident_sop_draft_clear	true	The SOP job empties the draft the AAR job put in its form (incident_sop_draft) once it made the SOP from it - only when launched with that draft's outage. Needs the AAP credential on the SOP template (incident_docs_draft_credential).
incident_name	(empty)	What was typed: the jobs' surveys ask these (the setup job writes the questions); empty unless answered. The AAR's boxes and the SOP's steps take a line or several. the outage's name or number (in the file name and the email)
incident_start	(empty)	the AAR: its start date and time (required)
incident_end	(empty)	its end date and time (optional)
incident_impact_level	(empty)	one of incident_impact_levels
incident_impact	(empty)	why: "it is affecting users enterprise-wide"
incident_detection	(empty)	how the event was detected
incident_actions	(empty)	key response steps and the teams involved
incident_root_cause	(empty)	the reason for the outage (RFO)

Setting	Default	What it does
incident_recovery	(empty)	the fix action
incident_successes	(empty)	what went well
incident_improvements	(empty)	areas to improve
incident_tasks	(empty)	the improvements as tasks, each with who does it
incident_make_sop	"No"	Yes = the AAR job attaches an SOP too: a sop_ step box left empty = the AAR's Actions / Recovery
incident_sop_draft	"No"	Yes = the AAR job fills in the SOP job's survey from the AAR, to review and launch
sop_title	(empty)	the SOP's title
sop_purpose	(empty)	the symptoms that send someone to the SOP
sop_troubleshoot	(empty)	the steps to troubleshoot, one per line
sop_fix	(empty)	the steps to fix, one per line
sop_verify	(empty)	how to check it is fixed, one per line
sop_escalate	(empty)	who to call when the steps do not fix it
incident_act_tidy	false	ACT (GenAI, evidence mode: it runs no command, only reads what was typed). Each is a Yes/No in the survey; these are the defaults when the survey does not ask. Needs the "ACT model key" credential on the template when either is Yes. tidy the spelling and wording; facts, names and numbers stay as typed
incident_act_suggest	false	a "From ACT" section: what may be missing or better, to review
incident_act_sop	false	the AAR job's SOP: ACT writes the empty step boxes from the AAR's Actions and Recovery
incident_act_max_suggestions	8	suggestions from ACT at most
incident_act_extra_instructions	(empty)	the site's own words for ACT, e.g. "Our monitoring tool is called X."
incident_act_model_timeout	120	seconds GenAI may take to answer
incident_docs_survey_state	present	The survey setup job (playbooks/incident_docs_survey.yml): adds the questions to the job templates that run incident_aar.yml and incident_sop.yml (found by their playbook). A question already there is left as it is. present = add the questions; absent = take them out again

Setting	Default	What it does
incident_docs_survey_project	(empty)	only job templates of this AAP project (its name); empty = any project
incident_docs_survey_ca_path	(empty)	a CA file for AAP's HTTPS certificate ("" = what the execution environment trusts)
incident_docs_survey_max	8000	characters a text box takes at most (AAP's own default is 1024)
incident_docs_draft_credential	(empty)	The SOP draft (incident_sop_draft) is written through AAP's API by the AAR job: name a "Red Hat Ansible Automation Platform" credential here and the setup job attaches it to the AAR and SOP templates (its user: Admin on the SOP template). Empty = attach it by hand, or no draft.

log_analysis

File: roles/log_analysis/defaults/main.yml

Note
LOG ANALYSIS (playbooks/log_analysis.yml for a group of servers, playbooks/vm_log_analysis.yml for vCenter and its ESXi hosts through vCenter's API). Read-only: each host's warning, error and critical log lines of the window, grouped on the AAP side into patterns (the same line with other numbers, times and names is one pattern, counted across hosts), and ACT (GenAI) explains the top ones - likely cause, evidence, fix, confidence - in evidence mode: ACT runs no command anywhere. Every setting can be set for all hosts, per group or host in the inventory, on the template, or as a survey answer (a list may be typed as "a, b" or one per line). docs/LOG_ANALYSIS.md

Setting	Default	What it does
log_analysis_hours	24	---- the window ---------- The last this-many hours ...
log_analysis_since	(empty)	... or a fixed window, e.g. the outage: log_analysis_since "2026-10-23 08:00" and log_analysis_until "2026-10-23 10:30" (UTC, as the report shows times). Since alone = from then until now.
log_analysis_until	(empty)	the end of the fixed window (empty = now)

Setting	Default	What it does
log_analysis_source	auto	---- what to read ---------- What a host is: auto (Windows when its connection is WinRM / PSRP, else Linux), linux, windows, or vcsa (a vCenter appliance over SSH: Linux, with the vcsa profile on). Set it per group in the inventory.
log_analysis_profiles	[]	Profiles: one word switches on the logs of a product (roles/log_analysis/vars/main.yml): vcsa, mid_server, iis, sql_server, nginx, httpd, mariadb, podman, sssd, idm, auditd, aap. E.g. for a group of web servers: log_analysis_profiles: [nginx, sssd]
log_analysis_profiles_extra	{}	Your own profiles, or a shipped one changed (the same name replaces it), the shape of vars/main.yml: log_analysis_profiles_extra: {stigman: {files: [/opt/stigman/logs/*.log], units: [stigman-api]}}
log_analysis_priority	err	Linux (and the vCenter appliance): the journal's lines of this priority or worse (journalctl -p: emerg, alert, crit, err, warning, notice).
log_analysis_journal	true	false = no journal (text logs, units and containers only)
log_analysis_units	[]	Units and syslog identifiers whose journal is read at log_analysis_units_priority, lower than the rest: e.g. [sssd, nginx, stigman-api] (globs such as dirsrv@* work) and [sudo, sshd].
log_analysis_identifiers	[]	syslog identifiers read the same way, e.g. [sudo, sshd]
log_analysis_units_priority	warning	the priority those units and identifiers are read at
log_analysis_containers	[]	Podman containers whose output (podman logs) of the window is read, its lines that match log_analysis_match: a name for root's, user/name for a user's (rootless), e.g. [stigman-api, svc_app/web].

Setting	Default	What it does
log_analysis_files	[]	Text logs (globs), read when they changed in the window: their matching lines, the newest log_analysis_max_lines. Compressed rotations (.gz .xz .bz2 .zst) are left out. log_analysis_files is the base list; log_analysis_files_extra adds to it and to the profiles' files (never replaces them). e.g. [/var/log/messages, /v... (all of it: the role file)
log_analysis_files_extra	[]	e.g. [/opt/app/logs/*.log]
log_analysis_become	true	read as root (sudo), as the other checks do
log_analysis_win_logs	[System, Application]	Windows: the event logs read, their critical (1), error (2) and - log_analysis_win_level 3 - warning events.
log_analysis_win_logs_extra	[]	more channels, e.g. ["Microsoft-Windows-TaskScheduler/Operational", "Directory Service", "DNS Server"]
log_analysis_win_level	3	1 critical, 2 + error, 3 + warning
log_analysis_win_providers	[]	Event providers whose events are read at every level up to information, in the logs above (e.g. [Service Control Manager]); log_analysis_win_providers_only: true = only these providers' events.
log_analysis_win_providers_only	false	true = only those providers' events
log_analysis_win_event_ids	[]	Event IDs always read whatever their level: "Log:ID", or a number for every log read plus Security, e.g. [4625, 4740, "System:7031", "System:7034"]. Security needs an account that may read it.
log_analysis_win_files	[]	Text logs (wildcards), read like the Linux ones; _extra adds to the base list and the profiles', e.g. ['D:\Apps\\l*ogs*\.log'] (a MID Server's agent logs: the mid_server profile)
log_analysis_win_files_extra	[]	more Windows text logs, added to the list above
log_analysis_match	"error|fail|fatal|crit|panic|denied|refused|timed out|timeout|unable to|cannot|exception|…	The words that make a line of a text log or a container's output worth reading (grep -E, any case); _extra adds words, e.g. [SEVERE, 'HTTP/1.1" 5'].

Setting	Default	What it does
log_analysis_match_extra	[]	more words, added to log_analysis_match
log_analysis_vcenter_logs	[vpxd]	---- vCenter and ESXi through vCenter's API (playbooks/vm_log_analysis.yml) ---------- No SSH and no login to the hosts: vCenter hands over its own logs and each host's (the "VMware vCenter" credential; a read-only role plus Global > Diagnostics). Keys as vCenter names them. [] = not vCenter's own log
log_analysis_esxi_logs	[hostd, vmkernel, vpxa]	[] = no host logs; more: fdm, vobd, hostd-probe ...
log_analysis_esxi_hosts	[]	host names (* ? allowed); [] = every host
log_analysis_read_lines	20000	read back at most this many lines from the end of each log
log_analysis_max_lines	500	---- every source ---------- keep at most this many lines per log (the newest)
log_analysis_max_line_chars	400	a line longer than this is cut
log_analysis_ignore	- 'Event 10016:' # Windows: DCOM "application-specific permission settings" (harmless, by…	Lines left out everywhere before they are grouped (regular expressions, any case, matched against "<program> <line>"): noise every server logs. log_analysis_ignore is the shipped list; _extra adds yours, e.g. ['sshd.*Connection closed by authenticating user'].
log_analysis_ignore_extra	[]	more lines to leave out, added to log_analysis_ignore
log_analysis_min_count	1	a pattern seen fewer times than this is left out
log_analysis_act	true	---- ACT ---------- ACT runs on the AAP side (the execution environment) with the site_act_* settings (roles/site_act/defaults/main.yml); the key comes from an "ACT model key" credential. false = the patterns only, no ACT (no key needed).
log_analysis_question	(empty)	What you want to know, in your words - ACT looks at it first (a good survey question), e.g. "users report slow logins since 09:00".

Setting	Default	What it does
log_analysis_names_extra	[]	more names ACT replaces with placeholders before the model sees them (host names are always)
log_analysis_max_items	30	ACT analyzes at most this many patterns, the most severe first (the rest are listed)
log_analysis_act_batch	"{{ site_act_batch | default(5) }}"	patterns per question to ACT, each with only its evidence (0 = all in one question)
log_analysis_examples	3	example lines per pattern in the evidence, the report and the artifacts
log_analysis_max_chars	60000	the evidence ACT reads is cut at this many characters
log_analysis_check_url	true	first check that the AAP node can reach the model's URL
log_analysis_act_model_timeout	"{{ site_act_model_timeout | default(600) }}"	Seconds the model may take for its answer: site_act_model_timeout (600) unless set here.
log_analysis_email_if_clean	false	---- the report ---------- ---------- The report is printed, kept as the job artifact log_analysis, emailed (report_email_to, and the "Also email the report to" box) and published (report_web_dest). A run that finds nothing is not emailed unless log_analys... (all of it: the role file)
log_analysis_max_rows	200	rows per table in the report (the artifacts have them all)
log_analysis_fail	false	The job stays green when it finds patterns. It fails when ACT was wanted and could not run (after the report), and when no host could be read. true = also fail when a critical or error pattern was found.

mail_relay

File: roles/mail_relay/defaults/main.yml

Note
MAIL RELAY WHITELIST, with an approval: add IP addresses or subnets to the list of clients the mail relay lets relay - with a comment saying who or what they are for - or remove them. The plan job (playbooks/mail_relay_plan.yml) checks the request, reads the list on every relay server and emails what would change - "waiting for approval"; it changes nothing. After a person approves in AAP, the approved job (playbooks/mail_relay_approved.yml) writes exactly that (the file backed up first), checks the mail configuration and that Postfix reads the list as written, reloads Postfix and emails the result; a failed check puts the file back. The format job (playbooks/mail_relay_format.yml) puts each relay's list one address or subnet per line - subnets first, then the single addresses, in number order - each single address with a comment line above it: its name in DNS (looked up on the relay) and the folder the VM that has it is in, in vCenter ("# web01.example.mil Payroll [AAP]"); the list keeps exactly its addresses, a copy of main.cf is kept first. The unused job (playbooks/mail_relay_unused.yml) reads each relay's mail logs and plans the removal of the addresses that sent no mail in mail_relay_unused_days - the same approval and apply step as a removal. docs/MAIL_RELAY.md

Setting	Default	What it does
`mail_relay_action`	(empty)	---- The request (a survey on the workflow) ---------- add or remove: set in each workflow's Variables (one workflow per action)
`mail_relay_ip`	(empty)	one address (198.51.100.3), a subnet (198.51.100.0/24), or several separated by commas: "198.51.100.3, 198.51.100.4"
`mail_relay_comment`	(empty)	add: required - the name and/or service the addresses are for, e.g. "Payroll app (svc-payroll)". Written above the mynetworks line, one line per address: "# 198.51.100.3 Payroll app (svc-payroll) [added 2026-10-22 by jdoe]"
`mail_relay_requested_by`	(empty)	who asked (e.g. sent by Orchestrator); "" = the AAP user who launched it
`mail_relay_group`	`mail_relays`	---- The relay servers and their list ---------- the inventory group of the mail relay servers: every one is changed
`mail_relay_style`	`main_cf`	main_cf: the list is the mynetworks line of Postfix's main.cf ("mynetworks = 127.0.0.0/8, 192.0.2.10"); an address is added at the end of the line (", 198.51.100.3"). file: a list file with one address per line (mynetworks = /etc/postfix/mynetworks, or a Postfix access map), its comment on the line above it.
`mail_relay_file`	`/etc/postfix/main.cf`	the file with the list on each relay (a relay's own host_vars may differ)
`mail_relay_parameter`	`mynetworks`	main_cf: the parameter that holds the list

Setting	Default	What it does
mail_relay_postconf	postconf	main_cf: after a change, read the list back as Postfix does ("" = do not)
mail_relay_entry_suffix	(empty)	file: after the address on its line: " OK" for a Postfix access map
mail_relay_postmap	false	file: true = run postmap on the file after a change (a hash: or lmdb: map)
mail_relay_check_command	postfix check	run after a change; when it fails the file is put back ("" = no check)
mail_relay_reload_service	postfix	reloaded after a change ("" = no reload)
mail_relay_become	true	change the file as root (sudo)
mail_relay_allow_subnets	true	---- What a request may ask for ---------- a subnet (198.51.100.0/24) too; false = single addresses only
mail_relay_min_prefix	24	the widest subnet accepted (/24 = 256 addresses); 0.0.0.0/0 and the like never
mail_relay_max_ips	20	the most addresses in one request
mail_relay_comment_min	3	the comment's least characters
mail_relay_comment_max	200	the comment's most characters
mail_relay_format_dns	true	---- The format job (playbooks/mail_relay_format.yml) ---------- look up each single address's name on the relay (getent hosts: DNS, /etc/hosts)
mail_relay_format_dns_timeout	5	seconds for one lookup
mail_relay_format_vmware	true	find the VM that has the address (needs the "VMware vCenter" credential)
mail_relay_format_blank_lines	true	a blank line between the single addresses (Postfix skips it)
mail_relay_format_refresh	false	true = look up again the addresses this job named before (others kept)
mail_relay_format_reload	false	true = reload Postfix after (the list is the same: no need)
mail_relay_format_check_email	true	a dry run (Check) emails its report too, marked DRY RUN (false = only printed)

Setting	Default	What it does
mail_relay_unused_days	183	---- The unused addresses (playbooks/mail_relay_unused.yml: workflow "Mail relay - remove unused addresses") ------ an address that sent no mail through any relay in this many days is unused (183: six months)
mail_relay_unused_logs	["/var/log/maillog*"]	Postfix's log files on each relay (patterns; rotated and .gz copies too; a folder - /var/log/maillog_archive: every file in it and its folders)
mail_relay_unused_journal	auto	auto = read the journal when no log file is there; true = as well; false = never
mail_relay_unused_keep	[]	addresses or subnets never proposed, e.g. [198.51.100.7, 203.0.113.0/28] (127.0.0.0/8 never is)
mail_relay_unused_max	50	the most addresses one plan removes (the next run takes the rest)
mail_relay_unused_accept_short_logs	false	The logs must go back mail_relay_unused_days on every relay (RHEL keeps about 5 weeks of /var/log/maillog): else nothing is proposed. true = propose anyway, from what the logs cover (the email says how far back they go).
mail_relay_approve_url	(empty)	---- The approval email ---------- Both emails (waiting for approval, and the result) go to report_email_to (set it on the workflow). A link for the approvers in the first one: your AAP's Workflow Approvals page, copied from the browser ("" = no link).
mail_relay_web	false	The report website (report_web_dest): the plan, the result and the format's report are emailed only; true = also published there, as the other reports (a website that cannot be written then fails the job before the approval).

net_certs

File: roles/net_certs/defaults/main.yml

> **Note**
> NETWORK DEVICE CERTIFICATES (playbooks/net_cert_report.yml, net_cert_renew.yml): the TLS certificate each switch, firewall, printer, iLO / iDRAC, UPS, storage ... presents, read from the controller with a TLS hello - no login to the devices, nothing changed. The renewal: a ServiceNow incident per certificate due; on a Cisco IOS / IOS-XE device, optionally a CSR made on the device (net_cert_cisco_csr).

Setting	Default	What it does
net_cert_targets	inventory	---- which devices ---- inventory (the hosts of net_cert_group) / subnets (net_cert_subnets) / both
net_cert_group	net_cert_devices	the inventory group of the devices; a host's own net_cert_ports / net_cert_address / net_cert_servername override
net_cert_ports	[443]	the ports read (a host's own net_cert_ports wins), e.g. [443, 8443, 636]
net_cert_subnets	[]	CIDR ranges to sweep (net_cert_targets: subnets or both), e.g. ['192.0.2.0/24'] - every address, each port in net_cert_ports
net_cert_exclude	[]	never connected to: addresses, CIDR ranges, host name patterns (e.g. ['192.0.2.1', 'ups*'])
net_cert_max_addresses	4096	a sweep bigger than this (all subnets together) is refused: a /16 by mistake would take hours
net_cert_resolve_names	true	a swept address that answers: its DNS name (PTR) in the report, else the address
net_cert_timeout	5	seconds for each connection, and each handshake
net_cert_workers	32	connections at the same time
net_cert_only	[]	only these devices (names, wildcards), e.g. ['sw01*'] - to try the renewal on one
net_cert_warn_days	30	---- what is due ---- a certificate expiring within this many days is due: a warning incident (amber in the report)
net_cert_crit_days	7	within this many days (or expired): critical
net_cert_notice_days	60	the report colours these blue (a heads-up, no incident)
net_cert_unanswered_finding	false	true = an inventory device whose port gives no certificate gets an incident too (a warning)

Setting	Default	What it does
net_cert_cisco_csr	false	On: for a due certificate of an inventory host with ansible_network_os cisco.ios.ios (and its network_cli credential on the job template), the job makes a key on the device (crypto key generate rsa, not exportable), a trustpoint <prefix>-<expiry> (enrollment terminal) and its CSR (crypto pki enroll), and attaches the... (all of it: the role file)
net_cert_cisco_key_bits	2048	the new RSA key (2048, 3072, 4096)
net_cert_cisco_trustpoint_prefix	SA-RENEW	the trustpoint and key label: SA-RENEW-20270131 (the old certificate's expiry)

net_config

File: roles/net_config/defaults/main.yml

Note
NETWORK DEVICE CONFIGURATIONS: the backup (playbooks/net_config_backup.yml) keeps each device's running configuration - in a git repository (a commit when one changed) or on a file share (dated copies) - and can open a ServiceNow incident when one changed. The load (playbooks/net_config_load.yml) sends configuration lines, or restores a backup, after an approved change request - run it in Check first to see what it sends (the job does not check that you did: an approval step in a workflow does); a real load first keeps each device's running configuration (net_config_keep_before). Cisco IOS / IOS-XE (over its CLI: network_cli, a Machine credential) and Palo Alto Networks PAN-OS (over its XML API from the controller: the "Palo Alto PAN-OS API" credential). docs/NETWORK_CONFIG.md. A backup holds the device's secrets (password hashes, SNMP communities, keys): keep it in a private, internal repository or share only - never in this one.

Setting	Default	What it does
net_config_group	net_config_devices	The inventory group of the devices. Each one's platform: ansible_network_os (cisco.ios.ios, paloaltonetworks.panos), or net_config_platform (ios, panos) on the host.
net_config_only	[]	Only these devices (names, * and ? allowed), e.g. [rtr-lab*] - to try a job on one first. [] = every one.
net_config_store	git	git (a repository), share (a folder on a file server or the AAP node), or both.

Setting	Default	What it does
net_config_git_repo	(empty)	The repository (ssh://git@git.example.mil/netops/configs.git or https://...). Private and internal: the configurations hold secrets. Its login: the "Network config git repository" credential (an SSH key, or a username and token).
net_config_git_branch	main	the branch (made when the repository is new and empty)
net_config_git_path	configs	the folder in the repository: <path>/<device>.cfg (IOS), .xml (PAN-OS)
net_config_git_author_name	site-automation	the commits' author
net_config_git_author_email	site-automation@localhost	and its email
net_config_git_host_key_checking	accept-new	The git server's SSH host key: accept-new (the first connection learns it), or yes (only the keys in net_config_git_known_hosts - lines as in a known_hosts file). In an execution environment ssh's own known_hosts is empty every job, so accept-new learns the key each time: only net_config_git_known_hosts + yes really c... (all of it: the role file)
net_config_git_known_hosts	(empty)	the git server's host keys (known_hosts lines), for net_config_git_host_key_checking: yes
net_config_share_dir	/var/lib/site-automation/net-config	The folder for share: on the AAP node (a path AAP exposes to jobs) or on net_config_share_host - <dir>/<device>/<device>.cfg (the latest), and a dated copy beside it each time one changed.
net_config_share_host	(empty)	an inventory host to keep them on (Linux, or a Windows file server); its credential on the template
net_config_share_keep	30	dated copies kept per device (the oldest go)
net_config_ios_command	show running-config	---- what is read ---- IOS: the command whose output is kept
net_config_panos_port	443	PAN-OS: the XML API's HTTPS port
net_config_panos_scheme	https	PAN-OS: always https - http is for the tests' stand-in firewall only (it would send the API key in clear)

Setting	Default	What it does
net_config_panos_validate_certs	true	PAN-OS: check the firewall's certificate
net_config_panos_ca_path	(empty)	PAN-OS: a CA bundle (PEM) for that check ('' = the system's)
net_config_timeout	60	PAN-OS: seconds per request
net_config_diff_lines	60	lines of a diff shown in the report, the email and an incident (secrets hidden)
net_config_fail_on_failed	true	the backup job fails (red) when a device could not be backed up
net_config_tickets	false	true = an incident when a device's configuration changed (one per change, with the diff) and when one could not be backed up (noted as cleared at the next good backup). The "ServiceNow API" credential.
net_config_ticket_changes	true	an incident for a change
net_config_ticket_failures	true	an incident for a failed backup
net_config_load_mode	lines	lines (send configuration lines) or restore (put a device back to a backup).
net_config_lines	[]	IOS: the lines, as typed in configuration mode, e.g. ["ntp server 192.0.2.123", "banner motd ^Authorized use only^"]; under net_config_parents (e.g. ["interface GigabitEthernet1"]) when they belong to one.
net_config_parents	[]	IOS: the parent lines the lines go under (e.g. ["interface GigabitEthernet1"]); [] = global configuration
net_config_panos_entries	[]	PAN-OS: the changes - each {xpath, element, action}: action set (adds or replaces the element's children), edit (replaces the element at the xpath), delete (no element).
net_config_restore_ref	(empty)	restore: which backup - git: a commit, tag or branch ('' = the newest); share: a dated copy's file name ('' = the latest).
net_config_change	(empty)	The change request the load belongs to (CHG0012345). Its state must be one of net_config_change_states; the job writes on it what it did. A dry run without one only warns.

Setting	Default	What it does
net_config_require_change	true	false = a real load without a change request (not recommended)
net_config_change_states	[Implement]	the change states a load may run in (as ServiceNow shows them)
net_config_change_note	true	a work note on the change: each device, what was sent, the result
net_config_save	true	IOS: write memory after a change (false: only the running configuration - a reload undoes it)
net_config_ios_file	flash:site-automation-restore.cfg	IOS restore: where the backup is copied first (scp; ip scp server enable on the device)
net_config_panos_allow_pending	false	PAN-OS: change even with someone's uncommitted changes (they would be committed too)
net_config_max_devices	10	a load to more devices at once is refused (net_config_only, the template's Limit - or raise it)
net_config_keep_before	true	A real load first reads each device's running configuration and keeps it where the backup keeps them (net_config_store), named <time>-before-<change> (git: a tag; share: a dated copy) - restore it with net_config_load_mode: restore and net_config_restore_ref set to that name. A device whose copy could not be kept gets... (all of it: the role file)

net_perf

File: `roles/net_perf/defaults/main.yml`

Note
NETWORK PERFORMANCE (playbooks/net_perf_report.yml): from each Linux host - the latency, jitter and loss to the targets it depends on (ping, and a timed TCP connect), its OCSP responders answering a real request (from the host, where it matters), its interfaces' drops and errors since the last run, and how today compares with the host's own recent runs (kept on the host). iperf3 throughput only between the pairs you name: short, one at a time. Read-only, except the few files of the host's own history (net_perf_dir). docs/NETWORK_PERF.md.

Setting	Default	What it does
net_perf_targets	[]	---- Targets (from every host of the job) ---------- --- name, host; port (optional: a timed TCP connect too); thresholds of its own (optional: rtt_warn_ms, rtt_crit_ms, loss_warn_pct, loss_crit_pct). A target may also be set per group or host (net_perf_targets_extr... (all of it: the role file)
net_perf_targets_extra	[]	more targets for some hosts (group_vars / host_vars): added to net_perf_targets
net_perf_ping_count	20	pings per target
net_perf_ping_interval	0.2	seconds between them (0.2 is the least a non-root ping may use)
net_perf_tcp_tries	5	timed TCP connects per target with a port
net_perf_rtt_warn_ms	50	average round trip from this many ms -> warning
net_perf_rtt_crit_ms	200	... from this many -> critical
net_perf_loss_warn_pct	1	packet loss from this percent -> warning
net_perf_loss_crit_pct	10	... from this percent -> critical (100 % = the target does not answer at all)
net_perf_timeout	5	seconds a single ping or connect may wait
net_perf_ocsp	[]	---- OCSP responders (asked from every host) ---------- Simplest: site_ocsp_certs (below) - a certificate the responder knows, its OCSP address and issuer read from it. net_perf_ocsp, when set, is used instead: a url alone - an OCSP request about a made-up certific... (all of it: the role file)
site_ocsp_certs	[]	certificates the OCSP responders know (PEM files from playbooks/), for net_perf_ocsp and vm_sso_ocsp when those are empty, e.g. [files/ocsp/test-user.pem]
net_perf_ocsp_warn_ms	2000	an OCSP answer slower than this -> warning
net_perf_ocsp_severity	critical	an OCSP responder that fails: critical (smart card logins fail) or warning

Setting	Default	What it does
net_perf_nics	true	---- Interfaces (drops and errors since the last run) ---------- read the interfaces' counters (/sys/class/net/*/statistics)
net_perf_nic_exclude	"^(lo|veth|virbr|docker|podman|cni|flannel|vnet|tap|tun)"	interfaces left out (regular expression)
net_perf_sysfs	/sys/class/net	where the kernel shows the interfaces and their counters
net_perf_drops_warn_per_hour	100	dropped packets per hour (received + sent) from this many -> warning
net_perf_errors_warn_per_hour	10	errors per hour (CRC, frame, carrier ...) from this many -> warning
net_perf_dir	/var/lib/site-automation/netperf	---- History (on the host) and the trend ---- ---------- --- each run's summary and the last counters; root-only. A dry run keeps nothing
net_perf_keep	30	runs kept on the host
net_perf_slower_factor	2.0	a target's average round trip this many times its usual (the median of the kept runs) -> warning
net_perf_slower_min_ms	5	... and at least this many ms more than usual (2 ms against 1 ms is not news)
net_perf_iperf_pairs	[]	---- iperf3 between named pairs (optional, needs iperf3 on both hosts and port 5201 open) ---------- server and client are inventory names of this job's hosts. One pair at a time: the server runs iperf3 for that pair's test only (-1, as the login user, at most net_perf_iperf_seconds + 60 s), and is stopped a... (all of it: the role file)
net_perf_iperf_seconds	5	seconds per test (at most 30)
net_perf_iperf_port	5201	the TCP port iperf3 listens on, on the server of each pair

ontap_report

File: roles/ontap_report/defaults/main.yml

Note
NetApp ONTAP health report (playbooks/ontap_health_report.yml). Read-only: GET requests to each cluster's REST API, with the account of a "NetApp ONTAP" credential (docs/NETAPP.md).

Setting	Default	What it does
ontap_clusters	[]	The clusters' management addresses (names or IPs), e.g. [cluster1.yoursite.mil, cluster2.yoursite.mil]. The credential's account must exist on each of them.
ontap_validate_certs	true	Check the clusters' HTTPS certificates. ONTAP's own certificate is often self-signed: then put its CA in ontap_ca_path (a file in the project, e.g. "{{ playbook_dir }}/files/ca/ontap-ca.pem"), or set false.
ontap_ca_path	(empty)	the CA file (PEM) the clusters' certificates are checked with; empty = the execution environment's own trust
ontap_timeout	30	seconds for one request (no answer: asked once more)
ontap_tls_legacy	false	true = connect to older ONTAP whose TLS the execution environment refuses (UNSAFE_LEGACY_RENEGOTIATION_DISABLED, KEY_TOO_SMALL, SHA-1)
ontap_volume_warn_pct	85	Thresholds (used %, hours, days, seconds): at or above "warn" = amber, at or above "crit" = red.
ontap_volume_crit_pct	90	volume used % -> red
ontap_volume_min_gb	50	volumes smaller than this (GB) are left out of the Volumes table - full (root volumes and other small ones run full by design) or not online; 0 = every volume
ontap_aggr_warn_pct	85	aggregate used % -> amber
ontap_aggr_crit_pct	90	aggregate used % -> red
ontap_inode_warn_pct	85	inodes used % of a volume -> amber
ontap_inode_crit_pct	90	inodes used % -> red
ontap_snapmirror_lag_warn_hours	24	last good transfer longer ago than this
ontap_snapmirror_lag_crit_hours	48	last good transfer longer ago than this -> red
ontap_snapshot_reserve_warn_pct	100	snapshots using this much of their reserve (100 = they spill into the volume)
ontap_snapshot_reserve_crit_pct	150	snapshots using this much of their reserve -> red
ontap_cert_warn_days	60	certificates expiring within this many days
ontap_cert_crit_days	14	certificates expiring within this many days (or expired) -> red
ontap_time_drift_warn_seconds	60	a node's clock against the AAP server's
ontap_time_drift_crit_seconds	300	a node's clock this far off -> red
ontap_port_errors_warn	1	receive errors (CRC and other bad frames) since the node booted
ontap_port_errors_crit	10000	receive errors on a port since boot -> red

Setting	Default	What it does
ontap_min_spares	1	fewer spare disks than this on a node = amber (0 = red)
ontap_ems_hours	24	the error events of the last this-many hours
ontap_ems_max	500	read at most this many events per cluster
ontap_ems_ignore_ldap	auto	leave out events whose name has "ldap" in it: auto = on a cluster where no SVM has an LDAP client configured; true = always; false = never
ontap_ems_ignore	[]	event names to leave out (regular expressions, any case), e.g. ['^secd\.ldap\.', 'callhome\.battery\.low']
ontap_report_max_rows	50	rows per table (problems first); the job's artifacts have all
ontap_snapshot_top	25	the snapshot space table: the volumes whose snapshots use the most
ontap_report_fail	false	The job stays green when it finds problems: the report and email say what. true = mark it FAILED when there is a critical finding or a cluster could not be read, for workflows.
ontap_hardware_max_rows	2000	---- NetApp ONTAP hardware inventory (playbooks/ontap_hardware_report.yml) ---- controllers, shelves and disk models listed in the email and the page at most (the workbook always has all)
ontap_hardware_xlsx	true	attach every controller, shelf, shelf module, disk and port as an Excel workbook (.xlsx) to the email
ontap_hardware_fail	false	true = mark the job FAILED when a cluster could not be read (the report still goes out), for workflows

patch

File: roles/patch/defaults/main.yml

> **Note**
> PATCH: install updates with dnf, reboot if the updates need it, and prove the host came back: every service that was running before must be running after. One host at a time by default; the rollout stops at the first host that fails.

Setting	Default	What it does
patch_packages	['*']	what to update ('*' = everything that has an update)
patch_security_only	false	true = only updates that fix security advisories
patch_exclude	[]	never update these here, e.g. ['mariadb*', 'kernel*']

Setting	Default	What it does
patch_serial	1	How many hosts are patched at the same time (1 = one by one: the rollout stops at the first host that fails). ONLY in the job template's Variables, a survey or the launch: Ansible picks the batch size before it reads all.yml or any group or host variables, so set there it does nothing.
automatic_restarts	false	NO AUTOMATIC RESTARTS unless you say so: this playbook reboots a host only when automatic_restarts is true (set it in the job template's Variables, or answer the survey), and then only as patch_reboot says. With automatic_restarts false (the default) it installs the updates and tells you which hosts still need a reboo... (all of it: the role file)
patch_reboot	when_needed	Reboot (only with automatic_restarts: true): when_needed (needs-restarting says so) | never | always
patch_reboot_timeout	1800	seconds to wait for the host to come back
patch_never_reboot_groups	[aap, aap_hosts, sat, netapp, netapp_console_hosts]	Hosts in these groups are patched but NEVER rebooted by this playbook (reboot them by hand, in a maintenance window). Add your own group names; the built-in protection below is on always.
patch_never_patch_groups	[no_patch, aap, aap_hosts, sat, netapp, netapp_console_hosts]	Hosts in these groups are skipped entirely (Satellite is upgraded with satellite-maintain, not dnf).
patch_skip_unreachable	true	BUILT IN, whatever you set above (roles/patch/vars/main.yml): the AAP hosts (groups aap and aap_hosts) are never patched and never rebooted by a job, and neither is the machine the job runs from (a local connection). A package script can restart services under the running job. Set the two lists above only to ADD group... (all of it: the role file)
patch_skip_unsupported	true	RHEL 8 and later are patched with dnf, RHEL 7 and its clones with yum. A host with neither (not a RHEL-like Linux): true = it is listed as NO DNF OR YUM and left out - not patched, no snapshot - and the job is not failed for it; its last lines list every such host. false = it fails at the update.

Setting	Default	What it does
patch_skip_not_ready	true	A host that answers but where the job cannot become root (sudo refused, its password wrong or missing), or with less than patch_min_free_mb free in /var: true = it is listed as NOT READY and left out - nothing changed on it, nothing cleaned up - and the job is not failed for it; its last lines list every such host. fa... (all of it: the role file)
patch_min_free_mb	1024	free space needed in /var before starting
patch_update_check	true	Each host is asked first which updates it has (dnf check-update, yum on RHEL 7 - read-only, in Check mode too, with patch_security_only, patch_exclude and patch_packages): true = one with none is listed as UP TO DATE and left out - no snapshot, not patched, the job not failed for it. A check that does not answer (a re... (all of it: the role file)
patch_service_wait	120	seconds to wait for services to come back after a reboot
patch_ignore_services	'^(user@|user-runtime-dir@|session-|dnf-makecache|packagekit|fwupd|run-|dbus-:|setroubles…	Services that come and go on their own and are not compared before/after (regular expression). dbus-: = the helpers D-Bus starts on demand and that end on their own (dbus-:1.1-org.fedoraproject.SetroubleshootPrivileged@83.service)
patch_compare_services	enabled	Which services must run again after the update: enabled = the ones that start at boot (systemctl is-enabled: enabled, generated - an init script on RHEL 7); a service started on demand (static, indirect, transient: D-Bus, a socket, another unit) that ended on its own is listed, never failed. all = every service that w... (all of it: the role file)
patch_snapshot	false	true = just before a host is patched (in its batch: patch_serial), vCenter takes a snapshot of its VM, without memory, named patch_snapshot_name. Attach a "VMware vCenter" credential to the template as well as the Machine one. Every host's VM is found once, before any host is patched: the host variable vmware_vm_name... (all of it: the role file)
patch_snapshot_name	patching	the snapshot's name (the AAP job, user and date go in its description)

Setting	Default	What it does
patch_snapshot_quiesce	false	true = VMware Tools freezes the file systems first (slower; fails without Tools)
patch_snapshot_required	false	false = a host whose VM is in no vCenter is a physical server: patched without a snapshot, and the job says so. When a vCenter could not be read, such a host may be a VM there, and the job stops before anything is patched. true = a host without a VM stops the job too (every host must be a VM, each patched only after i... (all of it: the role file)
patch_snapshot_replace	false	A VM that still has a snapshot named patch_snapshot_name (the last patching's, not deleted yet) stops the job before anything is patched, and the job names them. true = delete that old snapshot, then take the new one.
patch_snapshot_reuse_hours	24	The job run again soon after (it stopped at a host, or was run twice): a snapshot named patch_snapshot_name younger than this many hours is that patching's - kept and reused, no second one, nothing stopped. Older: as above (stop, or patch_snapshot_replace). 0 = never reused (as before 1.1.0).
patch_protected_with_snapshot	false	---- Protected VMs ---- A host on the VMware protected lists is not patched: one of its groups is in vmware_protected_groups, or its name, short name, vmware_vm_name or VM in vCenter matches vmware_protected_vms or vm_snapshot_cleanup_exclude_vms (* and ? allowed, any case) - the VMs the snapshot cleanup never touches... (all of it: the role file)

pki_bundle

File: `roles/pki_bundle/defaults/main.yml`

Setting	Default	What it does
pki_bundle_url	https://dl.dod.cyber.mil/wp-content/uploads/pki-pke/zip/unclass-certificates_pkcs7_DoD.zip	PKI - DOD BUNDLE CHECK (playbooks/pki_bundle_check.yml; docs/CERT_INSTALL.md, "Keeping the CA bundle current"). On a schedule: the DoD PKI's PKCS7 bundle is downloaded, checked (DISA's signed list of its files, and every CA up to a root of yours by its signature) and compared with your copy (cert_install_ca_bundle). A... (all of it: the role file)
pki_bundle_proxy	(empty)	http://proxy.yoursite.mil:8080 when the execution environment reaches the internet through one
pki_bundle_validate_certs	true	check the download site's certificate
pki_bundle_ca_path	(empty)	a CA bundle for that check (the EE's own trust store when empty)
pki_bundle_timeout	120	seconds for the download
pki_bundle_warn_days	90	a CA in the bundle expiring within this many days is listed
pki_bundle_fail_when_newer	false	true = the job ends red when a newer bundle is there (an AAP notification on failure tells you)
pki_bundle_accept_unverified	false	A bundle that differs from yours while DISA's signature on it could not be checked (a FIPS-mode execution environment, a copy without the signed list): false = listed as NOT VERIFIED, not attached, not offered; true = offered as a newer one.
pki_bundle_email	true	Email the result like the other reports: to report_email_to, through the mail relay in all.yml (docs/EMAIL_REPORTS.md), the checked .p7b attached when it differs from yours. false = never emailed (the launch box "Also email the report to" still is).
pki_bundle_email_when	newer	newer = only when it differs from yours (or fails); always = every run
pki_bundle_web	false	true = also on the report website (report_web_dest)

poam

File: `roles/poam/defaults/main.yml`

Setting	Default	What it does
`poam_csv`	`"{{ playbook_dir }}/../poam/poam.csv"`	POA&M STATUS: which Plan of Action & Milestones items are overdue or due soon, and (optional) which open CAT I / CAT II findings in STIG Manager have no POA&M item yet. The POA&M list is a CSV file in your work Git (export from eMASS or your spreadsheet, keep the columns below, save as CSV). POA&M data is sensitive: k... (all of it: the role file)
`poam_columns`	`id: POAM ID # unique item number weakness: Weakness # what is wrong severity: Severity #…`	Column names in YOUR file for each field this check needs (change the right-hand side).
`poam_closed_statuses`	`[Completed, Closed, Risk Accepted, Cancelled]`	items with one of these statuses are done: never reported
`poam_due_soon_days`	`30`	an open item due within this many days -> warning (past due -> critical)
`poam_date_formats`	`['%Y-%m-%d', '%m/%d/%Y', '%m/%d/%y', '%d-%b-%Y', '%d %b %Y', '%b %d, %Y']`	the date formats your file may use (Python strptime), tried in this order
`poam_stigman_api`	(empty)	---- STIG Manager cross-check (optional) ---------- Needs a Keycloak client that may use the client-credentials grant, with the scope stig-manager:collection:read, and a grant in the collections to read. Your realm owner creates it - this repository never changes the realm. T... (all of it: the role file)
`poam_stigman_collections`	`[]`	collection names; empty = every collection the client can read
`poam_stigman_severities`	`[high, medium]`	CAT I and CAT II findings must have a POA&M item
`poam_stigman_scope`	(empty)	Scope to ask for. Empty = the client's default scopes (they must include stig-manager:collection:read or stig-manager:collection).
`poam_validate_certs`	`true`	check STIG Manager's and Keycloak's HTTPS certificates

podman_discover

File: `roles/podman_discover/defaults/main.yml`

Note
PODMAN DISCOVERY (read-only): every podman container on the host, root's (rootful) and each user's (rootless), with how it is run (a systemd unit or plain podman) and whether it should be running. Used by the Health check `containers` and by Service watch. It changes nothing.

Setting	Default	What it does
`podman_discover_enabled`	`true`	false = do not look for containers at all (the containers check then says "skipped", and Service watch watches only its watch_containers list).
`podman_discover_rootless`	`true`	Also look at the containers of users (rootless podman). A user is looked at when a podman process (conmon) runs as them, when they have "linger" on (/var/lib/systemd/linger/NAME), or when they are in /etc/subuid AND have container storage (~/.local/share/containers/storage). Their containers are read AS that user (run... (all of it: the role file)
`podman_discover_users`	`[]`	More users to look at (names), e.g. an application account the rules above miss.
`podman_discover_ignore`	`[]`	Containers to leave out, as regular expressions that must match the WHOLE container name or OWNER/NAME, e.g. ['test-.*'] *(every container whose name starts with test-, any owner), ['alice/.*']* (all of alice's containers), ['root/web'] (root's web only).
`podman_discover_quadlet_dirs`	`[/etc/containers/systemd, /usr/share/containers/systemd]`	Where root's Quadlet files are (*.container, and *.kube: a pod from Kubernetes YAML). A user's are always read from ~/.config/containers/systemd, /etc/containers/systemd/users/UID and /etc/containers/systemd/users.
`podman_discover_unit_dirs`	`[/etc/systemd/system, /usr/local/lib/systemd/system]`	Where root's own systemd units are looked for: a *.service that runs podman (`podman generate systemd`, `podman run`, `podman kube play`, `podman pod start`, podman-compose / docker-compose), and the podman-kube@ units enabled there (.wants links). A user's: ~/.config/systemd/user. Packaged units (/usr/lib/systemd/sys... (all of it: the role file)
`podman_discover_timeout`	`60`	Seconds one podman or systemctl call may take before it is given up (reported as a finding).

service_watch

File: `roles/service_watch/defaults/main.yml`

Setting	Default	What it does
`watch_discover`	`true`	---- What to watch ---------- With watch_discover on (the default), service watch finds the containers itself: every podman container on the host that SHOULD be running - root's (rootful) and each user's (rootless) - as the Health check `containers` finds th... (all of it: the role file)
`watch_containers`	`[]`	Containers to watch in any case, in dependency order (a container's dependencies first). They are watched before the discovered ones, and one that is missing is reported as down. name: the container name, as `podman ps -a` shows it (or its Quadlet / service name) user: optional: the user that owns it (rootless podman)... (all of it: the role file)
`watch_url_ok_status`	`[200, 301, 302, 401, 403]`	HTTP codes that count as "answering"
`watch_find_units`	`true`	---- How each container is run (found by itself) ---------- A container that a systemd unit runs (a Quadlet .container file, `podman generate systemd`, or a unit you wrote) must be started with systemctl, never with `podman start`: the unit owns it, and while the unit is stopped the co... (all of it: the role file)
`watch_quadlet_dirs`	`[/etc/containers/systemd, /usr/share/containers/systemd]`	where root's Quadlet files are looked for (to find the unit that runs a container)
`watch_mode`	`approval`	---- What happens when one is down ---------- approval: ACT only diagnoses. Its proposed fix waits for the AAP approval step, then service_fix_approved.yml runs exactly the approved command(s) - no model involved. When ACT proposes nothing (or cannot run), the standard fix... (all of it: the role file)
`watch_verify_retries`	`6`	re-checks after a fix (every watch_verify_delay seconds)
`watch_verify_delay`	`5`	seconds between those re-checks

Setting	Default	What it does
watch_record_file	/var/log/service-watch/incidents.jsonl	---- Where incidents are recorded (on the host) ---------- One JSON line per event, plus a syslog line (tag service-watch) that your log forwarding (Splunk, Elastic, rsyslog relay) picks up. AAP's job history keeps its own copy.
watch_record_syslog	true	also send each incident to syslog (tag service-watch)
watch_act_env	{}	---- ACT settings ---------- Which provider and model: site_act_provider / site_act_model / site_act_url / site_act_ca, the same settings every runbook uses (roles/site_act/defaults/main.yml). The key comes from the "ACT model key" AAP credential, never f... (all of it: the role file)
watch_act_timeout	1200	seconds for one ACT run (the model's time per answer: site_act_model_timeout)
watch_instructions	(worked out by the role; see its file)	What ACT is told, on top of the act_triage role's standard instructions.

servicenow

File: roles/servicenow/defaults/main.yml

Note
SERVICENOW - open, update (and optionally resolve) incidents from check findings, and check that ServiceNow itself (the instance and its MID Servers) is healthy. Uses the REST Table API with ansible.builtin.uri: no extra collection to install. Connection: the "ServiceNow API" AAP credential sets SN_HOST (https://<instance>), SN_USERNAME and SN_PASSWORD in the job's environment. Without it, nothing is sent: the job prints what it WOULD have opened (a safe way to try it out).

Setting	Default	What it does
servicenow_min_severity	warning	Findings at or above this severity get a ticket: critical | warning. A ticket already open stays open while its finding is still there, even below this (the certificate renewal opens one for every CSR)
servicenow_table	incident	the table tickets are opened in (incident)

Setting	Default	What it does
servicenow_assignment_group	(empty)	Values for new incidents. Names work (the API is asked to accept display values), e.g. servicenow_assignment_group: Linux Operations
servicenow_caller	(empty)	user name; empty = the API account
servicenow_category	(empty)	e.g. Software
servicenow_urgency	{critical: "1", warning: "2"}	a new ticket's urgency, per finding severity (your instance's values)
servicenow_impact	{critical: "2", warning: "3"}	a new ticket's impact, per finding severity
servicenow_set_ci	true	set Configuration item = the host name (must exist in the CMDB)
servicenow_fields	{}	More fields for every new ticket - for an instance that makes fields such as Location, Subcategory or Channel mandatory: {column name: value}. Names and choice labels work, e.g. servicenow_fields: {location: "Building 100", subcategory: "Operating System", contact_type: "Self-service"} A field's column name: right-cli... (all of it: the role file)
servicenow_note_cleared	true	When a problem that had a ticket is gone, add a work note (and, if you want, resolve it).
servicenow_resolve_cleared	false	true = also resolve the ticket (servicenow_resolved_state, servicenow_close_code)
servicenow_resolved_state	"6"	Resolved
servicenow_close_code	"Solution provided"	must be a close code your instance has
servicenow_resolve_fields	{}	More fields set when a ticket is resolved (servicenow_resolve_cleared, and the test ticket), for an instance that makes them mandatory: {column name: value}, e.g. {resolved_by: svc_aap_sn}
servicenow_note_repeats	true	Add a work note to an open ticket each time its finding is seen again.

Setting	Default	What it does
servicenow_correlation_prefix	site-health	Every ticket from these playbooks carries a correlation ID starting with this, so the same problem on the same host always updates the same ticket.
servicenow_validate_certs	true	check the instance's HTTPS certificate (see servicenow_ca_path)
servicenow_ca_path	(empty)	A CA bundle (PEM) for the instance's TLS certificate, when the execution environment does not trust its CA (e.g. a DoD root CA): a path on the controller, such as a file you add to the project: servicenow_ca_path: "{{ playbook_dir }}/files/ca/servicenow-ca.pem". Empty = the execution environment's normal trust store.... (all of it: the role file)
servicenow_timeout	30	seconds for one API request
servicenow_fail_when_findings	true	Fail the job when it opened/updated tickets, so the workflow shows red while problems exist.
servicenow_check_mid	true	servicenow_health.yml: also check the MID Servers (needs read access to the ecc_agent table)
servicenow_slow_seconds	5	the API answering slower than this -> warning
servicenow_availability_ui	true	---- servicenow_availability.yml: is the instance there (every 5 to 15 minutes; roles/api_availability) ---- Asked each run: the REST API (a read of servicenow_table with the "ServiceNow API" credential) and - with this - the login page (/login.do, no login: the instance's UI, or the single sign-on page it sends to).... (all of it: the role file)
servicenow_mid_names	[]	---- servicenow_health.yml: the MID Servers ---------- What the instance says about each MID Server (ecc_agent): Up and validated, its heartbeat (last_refreshed) no older than servicenow_mid_stale_minutes, the same version on every MID Server, the ones you expect all there; its... (all of it: the role file)

Setting	Default	What it does
servicenow_mid_expected	[]	these must exist (names): one missing -> critical
servicenow_mid_stale_minutes	15	Up, but its heartbeat older than this -> critical (it reports every few minutes)
servicenow_mid_version_check	true	MID Servers on different versions -> warning (one stuck in an upgrade)
servicenow_mid_queue_check	true	read ecc_queue (the account needs read on it, e.g. the mid_server role)
servicenow_mid_queue_minutes	10	output waiting longer than this for its MID Server counts as stuck
servicenow_mid_queue_warn	1	this many stuck messages -> warning
servicenow_mid_queue_critical	50	... this many -> critical
servicenow_mid_issues_check	true	read the open MID Server Issues (ecc_agent_issue)
servicenow_check_mid_hosts	true	The MID Server HOSTS (Windows, over WinRM), in the same job and report: the hosts of this group of the job's inventory - the Windows inventory, with its Machine (Windows) credential next to the ServiceNow API one. No such group (or false below): the instance's view only. Read-only. On each host: every MID Server servi... (all of it: the role file)
servicenow_mid_hosts_group	servicenow_mid_hosts	the inventory group of the MID Server hosts
servicenow_mid_service_match	"snc_mid*"	service names (* and ?); the display name "ServiceNow MID Server*" matches too
servicenow_mid_hours	24	how far back the logs are read for SEVERE lines and JVM restarts
servicenow_mid_log_stale_minutes	30	the agent log not written for this long while the service runs -> critical (hung)
servicenow_mid_severe_warn	20	this many SEVERE lines in servicenow_mid_hours -> warning
servicenow_mid_jvm_restarts_warn	1	this many JVM restarts (crash, hang, out of memory) -> warning
servicenow_mid_disk_warn_pct	10	free space of the agent's drive under this % -> warning

Setting	Default	What it does
servicenow_mid_disk_critical_pct	5	... under this % -> critical
servicenow_mid_url_check	true	config.xml's instance URL not this instance (a clone, a moved MID) -> critical
servicenow_health_part	all	---- servicenow_health.yml: a self-hosted instance ---------- The front-end (application) nodes, RHEL, over SSH with sudo (the Linux Machine credential): the hosts of this group. Read-only, on each: every node unit (snc_<env>.service) running and starting at boot; its stats page rea... (all of it: the role file)
servicenow_check_nodes	true	false = neither the nodes nor the database host is read
servicenow_nodes_group	servicenow_nodes	the inventory group of the front-end nodes
servicenow_db_group	servicenow_db_hosts	the inventory group of the database host(s) (a child of mariadb_hosts)
servicenow_check_db	auto	read the database host(s): auto = when the inventory has front-end nodes
servicenow_db_checks	[mariadb]	the checks on the database host (roles/check_<name>; e.g. add services)
servicenow_db_paths	[/database]	the database's folders: their filesystems reported, the data dir expected under one
servicenow_db_datadir	(empty)	the data dir on the host ("" = @@datadir as the server says; a container's: its path here)
servicenow_db_top_schemas	5	the largest schemas listed
servicenow_db_top_tables	10	the largest tables listed
servicenow_db_history	true	The database's size is kept on the database host, for its growth (change since the last run, per day, and when the disk fills): the one thing this job writes - a small JSON file, not in Check mode. false = no history, no growth.
servicenow_db_history_file	/var/lib/site-health/servicenow-db-history.json	on the database host, written by the job

Setting	Default	What it does
servicenow_db_history_keep_days	60	sizes older than this many days are dropped from the history
servicenow_db_fill_warn_days	30	at the database's growth, its filesystem full in fewer days -> warning
servicenow_db_fill_critical_days	7	... -> critical
servicenow_node_units	[snc_prod, snc_test, snc_dev]	The node units looked for: those installed are judged (a box runs one of them), those not installed are no finding; none of them on a host of the group -> critical. [] = every snc_*.service on the host. * and ? work.
servicenow_node_home	/glide	the install folder (nodes, logs, glide.properties)
servicenow_node_port	0	the stats page's port: 0 = glide.servlet.port in the node's glide.properties,
servicenow_node_stats	true	read each node's stats page
servicenow_node_stats_url	(empty)	Where: tried in turn until one answers with the numbers - https://<the node's own name>:<port>, http:// on it (its inventory name, then its FQDN), then 127.0.0.1 (http, https), then 127.0.0.1 with a name as the Host header (the node's, then the instance's); each /xmlstats.do then /stats.do. The instance's address (gli... (all of it: the role file)
servicenow_node_stats_auth	auto	A stats page that asks for a login (401 / 403, a redirect, a login page - or, with auto, the instance's normal web page in place of the numbers): auto = asked again with the ServiceNow API credential's login (basic auth) - sent only over HTTPS to the node's own name (its inventory name or FQDN), or to 127.0.0.1, never... (all of it: the role file)

Setting	Default	What it does
servicenow_node_labels	{}	A friendly name per node, shown beside its host in the report ("Node 1 (node01.corp.example.mil)"): the host variable servicenow_node_label on the node ("Node 1"), or here by host: {node01.corp.example.mil: Node 1, node03: Node 3} (the inventory name, or its short name). The host variable wins.
servicenow_node_stats_pages	[]	Nodes AAP does not log in to (off by default): each node's own stats page, read from the AAP side - no SSH to them. servicenow_node_stats_pages: - {name: "Node 1", url: "https://node01.corp.example.mil:16000/stats.do"} The same numbers, samples and thresholds as a node read over SSH; a redirect is never followed; the... (all of it: the role file)
servicenow_node_validate_certs	true	Check the stats page's HTTPS certificate and name, with the node's own trust store. A node whose certificate is self-signed or from a CA the node does not trust is not read over HTTPS: add that CA to the node's trust store, or set this false (the page is still read over plain HTTP on 127.0.0.1).
servicenow_node_stats_samples	3	reads of the stats page ...
servicenow_node_stats_interval	5	... this many seconds apart
servicenow_node_stats_timeout	10	seconds per read
servicenow_node_mem_warn_pct	85	JVM memory used of its max -> warning
servicenow_node_mem_critical_pct	95	... -> critical
servicenow_node_semaphore_queue_critical	10	transactions queued for a semaphore in every read: warning; this many -> critical
servicenow_node_scheduler_queue_warn	50	scheduled jobs waiting for a worker -> warning above this
servicenow_node_semaphore_wait_ms	1000	a slow transaction that waited this long for its semaphore is a semaphore problem

Setting	Default	What it does
servicenow_node_slow_tx_warn	200	slow transactions ("waited excessive time to run") in the log's hours -> warning
servicenow_node_log_hours	24	how far back the node's localhost log is read
servicenow_node_log_glob	"localhost_log*"	the log files in the node's logs folder
servicenow_node_log_groups	10	the most frequent log messages listed
servicenow_node_severe_warn	100	this many SEVERE lines in servicenow_node_log_hours -> warning
servicenow_node_disk_warn_pct	85	the install folder's or /'s filesystem used -> warning
servicenow_node_disk_critical_pct	95	... -> critical
servicenow_node_swap_warn_pct	50	swap used -> warning
servicenow_fapolicyd	true	fapolicyd on the nodes and the database host: active? (inactive: said, info - no finding), its denials since servicenow_fapolicyd_since by program and path (warning; a denial of the node's java or install folder, or of the database server or its data, critical)
servicenow_fapolicyd_since	today	ausearch -ts
servicenow_node_fapolicyd_critical	[java, "{{ servicenow_node_home }}"]	programs (names) or folders
servicenow_db_fapolicyd_critical	[mariadbd, mysqld, "{{ servicenow_db_paths | first | default('/database') }}"]	+ the data dir
servicenow_cluster_check	auto	sys_cluster_state: auto = when the inventory has front-end nodes; true; false
servicenow_cluster_stale_minutes	10	a node online but its last heartbeat older than this -> critical
servicenow_test_resolve	true	servicenow_test_ticket.yml: resolve the test incident at the end (false = leave it open to look at)

Setting	Default	What it does
cert_renew_servicenow	incident	How the CSRs the certificate renewals make (Linux, Windows, VMware, network devices) go to the PKI. incident: each one is attached to its certificate's incident, with a work note saying what to do next - a person submits it to the PKI. catalog: each one is ordered from your PKI's catalog item (cert_renew_catalog_item)... (all of it: the role file)
cert_renew_csr_memory_days	90	incident: a CSR on an incident of the certificate made in the last this-many days counts as made - also when that incident was resolved or closed since (the PKI may still be signing it): vCenter, ESXi and Cisco keep one pending key, and a second CSR would replace it. Older: a new CSR (the certificate is due again).
cert_renew_catalog_item	(empty)	The catalog item: its name exactly as the catalog shows it (e.g. "DoD PKI Certificate Request"), or its sys_id.
cert_renew_catalog_variables	{}	The item's questions - by their internal names, as the catalog test lists them - and the answers: {name: answer}. A reference (Requested For, Location ...): the record's name, or its sys_id. A choice: its value or its label. In an answer, {host}, {subject}, {names}, {expires}, {issuer} and {where} become the certifica... (all of it: the role file)
cert_renew_catalog_csr_variable	(empty)	The attachment question that gets the CSR (its internal name). Empty = the item's only attachment question.
cert_renew_catalog_names_variable	(empty)	A question that gets the certificate's names (subjectAltName), e.g. a "List any additional SANs" box. Empty = none.
cert_renew_catalog_reorder_states	["4", "7"]	A request closed in one of these states (a requested item's: 4 Closed Incomplete, 7 Closed Skipped - rejected or cancelled) is ordered again on the next run. In any other state - open, in progress, Closed Complete (the PKI issued it: install it) - the certificate is not ordered again.

Setting	Default	What it does
cert_renew_catalog_upload_table	ZZ_YYsc_cart_item	Where the CSR is uploaded before the order, for the attachment question (ServiceNow moves it to the request).
servicenow_catalog_test_order	false	servicenow_catalog_test.yml: true = also order the item once, with a test CSR - the PKI gets a real request (tell them first, or cancel it). false = only read the item and check the settings.
cert_install_attachment_exts	[cer, crt, pem, der, p7b, p7c, p7, cert, zip, txt]	---- Certificate install: the issued certificate, from its request (cert_install_check.yml, win_cert_install_check.yml) ---- The PKI puts the issued certificate on the request (RITM) that carries the CSR: attached to the RITM, to one of its attachment questions, to its REQ or to its catalog task (SCTASK) - one file or... (all of it: the role file)
cert_install_max_bytes	1048576	a bigger attachment is not read
cert_install_journal	true	also look for a certificate pasted into the request's work notes and comments
cert_install_min_days	30	the issued certificate must be valid at least this many more days (else held for a person)
cert_install_fewer_names	false	true = install a certificate that lacks names (subjectAltName) the old one had; false = held for a person
cert_install_ritm	(empty)	Linux and Windows: the request(s) to install from, named (the check templates' survey) - "RITM0012345", several with commas, or a list. Empty (the default) = each renewal's own catalog request, found by its CSR (cert_renew_servicenow: catalog). Named: only these requests are read, however they were opened - by hand, w... (all of it: the role file)

Setting	Default	What it does
cert_install_ca_bundle	(empty)	A CA bundle in YOUR repository that completes the chain when the PKI sends back only the .cer - e.g. the DoD PKI's PKCS7 bundle from cyber.mil: files/dod_pki_bundle.p7b (playbooks/files/; the .p7b, DER or PEM, a .pem, or the zip as downloaded). Only the CAs that signed the certificate are taken from it - its issuing C... (all of it: the role file)

servicenow_insights

File: roles/servicenow_insights/defaults/main.yml

Note
TICKET INSIGHTS (playbooks/servicenow_insights.yml, docs/SERVICENOW_INSIGHTS.md): months of ServiceNow tickets in one report for management - the most common incident issues and their fix, the most requested services, keeping up (opened, resolved, the backlog and its age), P1/P2s, problem candidates, tickets passed between groups, changes that failed and what followed, close notes that say nothing, the All-stars of each team, knowledge article drafts. A web page for the whole site and one per team (report_web_dest), an email (report_email_to) with the workbook. ACT (GenAI) groups the incidents into issues and writes the fixes; names, phone numbers, email addresses, host names and IP addresses are replaced before anything reaches the model and put back after. From the AAP side with the "ServiceNow API" credential (the account needs read on the tables: itil). Read-only, in Check too.

Setting	Default	What it does
servicenow_insights_months	6	THE WINDOW. The last servicenow_insights_months months up to now (servicenow_insights_until), compared with the same length before it. servicenow_insights_days > 0 wins over the months (e.g. 30, 60). Dates and weeks follow report_time_zone (the email settings; "" = UTC).
servicenow_insights_days	0	> 0: the last N days instead of the months
servicenow_insights_until	(empty)	the last day (through its end, 2026-10-31), "" = now
servicenow_insights_compare	true	read the window before as well: the changes against it (the tiles, rank)
servicenow_insights_groups	[]	The assignment groups reported on, exactly as ServiceNow names them: empty = every group (the whole site). A list, or text with one name per line or separated by ; or ,.

Setting	Default	What it does
servicenow_insights_team_pages	true	Each team's own page (every section of the report for that team) on the report website, and the Teams side by side table on the site's page. The teams are the groups named above, else the groups with at least servicenow_insights_team_min tickets in the window, the busiest first, at most servicenow_insights_teams_max.
servicenow_insights_teams_max	60	at most this many teams with their own page
servicenow_insights_team_min	10	a group with at least this many tickets in the window is a team (no groups named)
servicenow_insights_hierarchy_source	parents	SITES, DIRECTORATES AND TEAMS (0.97.0): a page for each site and directorate with tickets, under All; teams as above. parents (the default): ServiceNow's group tree (sys_user_group, read only) - the site is the root group (its name's first " - " part), the directorate the group one below the root, the team the group (... (all of it: the role file)
servicenow_insights_site_roots	[]	parents: the sites' root groups (exact names); [] = parentless groups with children
servicenow_insights_hierarchy_map	{}	map: {"<exact group name>": [Site, Level1, ...]}; parents / names: {"<group>": {site: X, directorate: Y, team: Z}}, wins over either
servicenow_insights_unmapped_bucket	Unknown	map: where the groups with tickets that are not in the map go (last)
servicenow_insights_custom_view	true	map: the All page's Custom view (tick units and groups: their figures)
servicenow_insights_custom_view_max_rows	30000	its tickets carried in the page at most, about 55 bytes each. More: no Custom view
servicenow_insights_custom_view_presets	{}	map: the Custom view's own quick choices, besides the split's two ("<label>", "non-<label>"). Either {name: {levels: [SPLIT]}} (the groups whose map path has every level listed) or {name: {pattern: "- *Desk*"}}.

Setting	Default	What it does
servicenow_insights_periods	[all, this_month, 1m, 2m, 2w, 1w]	map: the Custom view's periods (1.10.0), calendar ones, each against the one before. all = the window; this_month = the report end's month so far; 1m, 2m = the last whole calendar month(s); 1w, 2w = the last whole weeks. A from / to too.
servicenow_insights_week_start	mon	the day a week starts on, for 1w and 2w (mon ... sun)
servicenow_insights_suffix_directorates	[]	names: [NOVA] - a group whose name ends in "- NOVA" is in NOVA
servicenow_insights_directorates	[]	names: second parts that name a directorate (e.g. [CODEX])
servicenow_insights_directorate_tags	{}	names: {COMPONENTX: N6} - a whole word in a group's name puts it there
servicenow_insights_directorate_alias	{}	names: {J602: J60} - rename or merge directorates
servicenow_insights_use_group_parents	true	read the groups' parents in ServiceNow (sys_user_group); refused: by the names
servicenow_insights_known_groups_only	false	true: a group not in the map goes in Unknown (to find and map later)
servicenow_insights_exclude_groups	[]	no team page for these (names or patterns, e.g. ["* - Legacy*"]); tickets still count
servicenow_insights_inactive_teams	false	a group inactive in ServiceNow gets no team page; its tickets still count
servicenow_insights_skip_group_types	[approver, permission, survey, system, idea, legacy]	no team page for these group types
servicenow_insights_split_label	(empty)	THE SPLIT: with both set, every page has All | LABEL | non-LABEL - its figures for the groups whose name, or a parent team's below the root, matches the pattern (* and ?, any case), and for the rest. E.g. label NOVA, pattern "*- NOVA". the switch's label (e.g. NOVA); "" = no switch
servicenow_insights_split_pattern	(empty)	the groups it picks: e.g. "- *NOVA" (* and ?, any case)
servicenow_insights_split_by	both	map: pattern (the name), level (the label is one of the group's map levels) or both
servicenow_insights_split_teams	true	false: the switch on All, the sites and directorates only (fewer files)

Setting	Default	What it does
servicenow_insights_sla_source	auto	YOUR SLOs (1.7.0; docs/SERVICENOW_INSIGHTS.md, "Your SLOs, priorities and categories"). Where met or missed comes from: auto ServiceNow's own SLAs (task_sla) when the account reads them (itil does), else the targets below task_sla ServiceNow's SLAs: each SLA definition a Response or a Resolution by its name (the two l... (all of it: the role file)
servicenow_insights_sla_response_match	[response]	an SLA definition whose name has one of these is a Response
servicenow_insights_sla_resolution_match	[resolution, resolve]	... one of these a Resolution; any other: Other, shown by name
servicenow_insights_business_hours	{}	Business hours for the targets: {} = calendar hours. E.g. {days: [mon, tue, wed, thu, fri], start: "07:30", end: "16:30", holidays: ["2026-12-25"]} - in report_time_zone; an end before the start is a shift past midnight.
servicenow_insights_sla_targets	incident: {"1": 8, "2": 16, "3": 144, "4": 240, "5": 480} # P1 8 h, P2 16 h, P3 6 days, P…	SLA TARGETS, in hours from when a ticket was opened to when it was resolved (incidents) or closed (requested items) - business hours when set: a ticket met its SLA when it was done within its target. By table, then by priority ("1" to "5", or "P1"; "*" = every other priority). A number is the resolution; {response: 0.... (all of it: the role file)
servicenow_insights_sla_items	{}	A catalog item's own target in hours, beating its priority's: {"New account / onboarding": 48, "Software install": 120}
servicenow_insights_team_email_to	{}	Also email a team's own report to these people (the team's page, the same look): {"Ops - Windows": [lead@yoursite.mil]}.
servicenow_insights_tables	[incident, sc_req_item, sc_task, change_request, change_task, problem]	The tables read: incident, sc_req_item (requested items: the most requested services), sc_task (catalog tasks: the All-stars count their closures), change_request, change_task (1.2.0: the changes' tasks - each team's, and which team a change waited on last), problem. Take one out to leave its sections out. A table the... (all of it: the role file)
servicenow_insights_priority_names	{"1": Critical, "2": High, "3": Moderate, "4": Low, "5": Exempt}	PRIORITIES AND CATEGORIES (1.7.0). The priorities' names on the report, by number:

Setting	Default	What it does
servicenow_insights_priority_matrix	"1": {"1": "1", "2": "2", "3": "3"} "2": {"1": "2", "2": "3", "3": "4"} "3": {"1": "3", "…	The priority matrix, impact -> urgency -> priority (the standard 3 x 3): an incident whose priority is not the one its impact and urgency give was set by hand - counted, per group. [] = not checked.
servicenow_insights_categories	[]	Your official list of incident categories, e.g. [Email, Network Connectivity, Workstations, Other]: the report counts the incidents in a category not on it (and names those categories). [] = no list: every category as it is.
servicenow_insights_core_hours	{}	THE LIFE CYCLE AND THE FLOW (1.9.0) The core hours, for After hours (incidents opened outside them): {} = servicenow_insights_business_hours, else Mon-Fri 07:30-16:30; the same form, e.g. {days: [mon, tue, wed, thu, fri], start: "07:30", end: "16:30"}.
servicenow_insights_short_desc_chars	15	Ticket quality: a short description under this many characters
servicenow_insights_hold_many	2	On hold: incidents put on hold more than this many times are counted
servicenow_insights_repeat_ci_days	7	Hot spots: another incident on the same CI within this many days is a repeat
servicenow_insights_person_metrics	false	Workload: the median time to resolve per person, only when true (open and resolved per person are always shown).
servicenow_insights_sla_risk_pct	75	WHAT TO ACT ON FIRST (1.7.0) SLA at risk now: an open incident whose running SLA (task_sla) has used this much of its time (business time where ServiceNow keeps it), or is breached and still open - listed worst first.
servicenow_insights_hop_limit	3	The 3-hop rule: an incident transferred more than this many times after its first assignment (reassignment_count over hop_limit + 1) is listed while open, and counted per group.

Setting	Default	What it does
servicenow_insights_awaiting_caller_days	6	On hold: Awaiting Caller longer than this many business days (servicenow_insights_business_hours' days, else Mon-Fri) is listed - the process resolves it with No Resolution Provided by then. On hold since: this metric_instance definition (the incidents' state durations), else when their SLA was paused (task_sla); "" =... (all of it: the role file)
servicenow_insights_state_metric	Incident State Duration	the incidents' state durations (metric_instance); "" = not read
servicenow_insights_scrutiny_codes	[Duplicate, Known Error, No Resolution Provided]	The resolution codes counted per group, and the one each of whose tickets is listed for leadership's review.
servicenow_insights_review_code	No Resolution Provided	each incident resolved with it is listed for review
servicenow_insights_close_after_days	7	Resolved tickets close by themselves after this many days: a resolved incident older than that and not closed is listed.
servicenow_insights_outcome_map	Workaround: Workaround Workaround provided: Workaround Duplicate: Duplicate Known Error:…	How incidents ended: a close code (as shown, any case) -> its outcome; any other code is Fixed. Unable to resolve = these outcomes, of those resolved.
servicenow_insights_unable_outcomes	[Known Error, No Resolution Provided]	the outcomes counted as unable to resolve
servicenow_insights_service_desk_groups	[]	The service desk's assignment groups (exact names), for its own first-contact resolution; [] = the site's only.
servicenow_insights_survey	true	The incident surveys (asmt_metric_result, read only, best-effort): the average answer per group and month. The metrics read by name ([] = every metric whose answer is a number); false = not read.
servicenow_insights_survey_metrics	[]	the survey metrics read, by name; [] = every metric answered with a number
servicenow_insights_max_records	20000	at most this many records per table (the rest are said: shorten the window)

Setting	Default	What it does
servicenow_insights_audit	true	The assignment history of incidents passed around: the paths of Passed between groups. From sys_audit, else (0.97.0) the out-of-box incident metric "Assignment Group" (metric_instance) - both read only; the account needs read on one of them; without either the counts are still there and the report says which it could... (all of it: the role file)
servicenow_insights_handoff_source	auto	auto (sys_audit, else the metric), audit, metric (sys_audit never asked), none
servicenow_insights_members	true	The groups' members (sys_user_grmember): each team's leaderboard lists its members, and a ticket resolved by someone outside the team counts for the member it was assigned to. The account needs read on sys_user_grmember (itil has it on most instances); without it the leaderboard lists everyone the team's tickets were... (all of it: the role file)
servicenow_insights_top_issues	10	WHAT THE REPORT SHOWS incident issues listed (the workbook has every one)
servicenow_insights_top_requests	10	catalog items listed
servicenow_insights_old_days	30	open longer than this: flagged (Over 30 days)
servicenow_insights_handoff_min	3	an incident reassigned this many times or more is "passed between groups"
servicenow_insights_problem_min	3	a problem candidate: this many incidents of one issue on one CI ...
servicenow_insights_problem_days	30	... within this many days
servicenow_insights_repeat_days	14	reopened, or came back (a new incident of the same caller and issue), within this many days
servicenow_insights_reopen_min_percent	0.2	Reopen is "not used" when fewer than this share (%) of the incidents were ever reopened: the report then counts the incidents that came back instead (the same person, the same issue, a new ticket).

Setting	Default	What it does
servicenow_insights_change_follow_hours	72	Changes: an incident on the change's CI within this many hours of its end is shown as following it ("Same CI"; a Caused by link is always shown). A failed change: its close code is one of these (any case).
servicenow_insights_change_failed_codes	[unsuccessful]	a change closed with one of these codes failed
servicenow_insights_change_backout_words	[back out, backed out, backout, rolled back, rollback, roll back]	in its close code or notes: backed out
servicenow_insights_bare_chars	20	CLOSE NOTES. Bare: shorter than servicenow_insights_bare_chars, or only one of servicenow_insights_bare_phrases (any case, the punctuation ignored). Thin: longer, but no action word. Good: says what was done.
servicenow_insights_bare_phrases	[fixed, done, resolved, closed, complete, completed, resolved per user, closed per user,…	The close notes that say nothing, whole (any case, the punctuation ignored).
servicenow_insights_bare_weight	0.5	ALL-STARS (each team): a ticket with a bare close note counts this much (0.5 = half); awards need this many resolved.
servicenow_insights_star_min	25	awards need this many tickets resolved
servicenow_insights_act	false	ACT (GenAI): groups the incidents into issues (a list carried from batch to batch), writes each top issue's fix from its close notes (or suggests one, marked, where the notes say too little), drafts knowledge articles and writes the summaries. Attach an "ACT model key" credential (site_act_provider, site_act_models);... (all of it: the role file)
servicenow_insights_act_batch	100	incident descriptions per question
servicenow_insights_act_max_descriptions	2000	the most frequent descriptions grouped; the rest go by their category
servicenow_insights_act_max_issues	40	at most this many issues
servicenow_insights_act_fix_issues	15	the top issues ACT writes a fix for ...
servicenow_insights_act_fix_batch	3	... this many per question

Setting	Default	What it does
servicenow_insights_act_notes	40	close notes per issue it reads
servicenow_insights_act_work_notes	true	... and the work notes of those issues' incidents (0.96.0): what was tried while they were open, read from sys_journal_field (GET only; the account needs read on it - refused, the page says so and the fixes come from the close notes), masked like everything else before ACT reads them. The newest ..._max per issue, eac... (all of it: the role file)
servicenow_insights_act_work_notes_max	20	the newest work notes per issue ACT reads
servicenow_insights_act_work_note_chars	1500	each cut to this many characters
servicenow_insights_kb_min	50	a knowledge article draft for an issue with this many incidents and a fix that agrees
servicenow_insights_act_model_timeout	600	seconds the model may take for an answer
servicenow_insights_act_check_url	true	first check that this AAP node reaches the model's URL
servicenow_insights_mask_names	[]	more words to replace before ACT (a project name, a building)
servicenow_insights_mask_keep	[]	words that look like host names but are not (e.g. [win11, rhel9]): sent as they are
servicenow_insights_db	false	THE HISTORY (1.13.0, docs/REPORT_DATABASE.md): true = each run writes its figures to the report database (MySQL 8.4, the "Report database" credential) - the window's, and each calendar month in it (the month so far marked partial) - and the pages show the months as Trends, for each site, unit and team. Figures only: c... (all of it: the role file)
servicenow_insights_db_required	false	true: the database not reached (or no credential) fails the job
servicenow_insights_db_create	true	make the tables and views on the first run (false: made by hand, from the doc)
servicenow_insights_db_months	12	the months the trend shows
servicenow_insights_db_timeout	10	seconds to connect

Setting	Default	What it does
servicenow_insights_backfill_months	0	THE BACKFILL: > 0 = this launch reads that many past calendar months from ServiceNow once, writes each month's figures to the report database and stops - no ACT, no pages, no email. Run it once after setup (12), as a separate launch. Safe to run twice: each month replaces what was there.
servicenow_insights_pages_per_call	40	THE PAGES (1.14.0): each site's, unit's and team's page is written this many at a time, in one task after the All page, when site_report_web writes the website itself (a folder, an SMB share, an upload). 0 = the site_email role once per page, as before (report_web_copy and report_web_host always do that).
servicenow_insights_xlsx	true	The workbook on the email: every incident with its issue and close-note grade, the issues, requests, changes, people.
servicenow_insights_fail	true	The job fails when ServiceNow could not be read at all, or a group named is not there. false = a report only.

servicenow_war

File: roles/servicenow_war/defaults/main.yml

Note
WEEKLY ACTIVITY REPORT (playbooks/servicenow_war.yml, docs/SERVICENOW_WAR.md): everything a ServiceNow assignment group had on its plate in a week - incidents, catalog tasks, changes, change tasks, problem tasks - what was opened, what was completed and when (the date of closure), what is still open and for how long, who had what; the lines to paste into the WAR; the same as an Excel workbook; emailed (on the report website only with servicenow_war_web). From the AAP side with the "ServiceNow API" credential (the account needs read on the tables: itil). Read-only, in Check too.

Setting	Default	What it does
servicenow_war_groups	[]	The assignment group(s), exactly as ServiceNow names them: a list, or text - one name, or several on separate lines or separated by ; or , (a survey or a schedule's answer; never cut at spaces or dashes). Every task of the week that is assigned to one of them now. e.g. [Ops - Example group]. One template serves every... (all of it: the role file)

Setting	Default	What it does
servicenow_war_users	[]	Also the tasks assigned to these people, whatever their group (user names, or full names). Empty = the groups' tasks.
servicenow_war_web	false	The report website: false (the default) = this report is not published there, whatever report_web_dest says - only emailed (and in the job output and artifact); true = published like the other reports.
servicenow_war_web_per_team	true	On the report website (servicenow_war_web: true) each team's runs go in their own folder ({template}/<the group, lower case with dashes>/...): its own latest page and its own line on the index, when one template runs for several teams. false: one folder for all.
servicenow_war_tables	[incident, sc_task, change_request, change_task, problem_task]	The tables read: incident (tickets), sc_task (catalog tasks, SCTASK), change_request (changes), change_task (CTASK), problem_task (PTASK); also problem, sc_req_item (RITM), sc_request (REQ), or any task table of your instance. "task" = every kind of task in one read (the parent table: everything, whatever its class).... (all of it: the role file)
servicenow_war_fields_extra	[]	more columns of your own read from every table, e.g. [u_customer, u_site]: on the workbook
servicenow_war_period	days	THE WEEK. days = the last servicenow_war_days days up to now (a Friday afternoon run: the week so far); last-week = the previous calendar week, from servicenow_war_week_start to the day before it (a Monday morning run: last week, whole); this-week = this calendar week so far; dates = from servicenow_war_since to servi... (all of it: the role file)
servicenow_war_days	7	period days: how far back
servicenow_war_week_start	monday	the day a calendar week starts on (last-week, this-week): monday ... sunday
servicenow_war_since	(empty)	period dates: the first day (2026-10-20, or 2026-10-20 08:00); given, it wins over the period
servicenow_war_until	(empty)	period dates: the last day (through its end), or a date and time; "" = now

Setting	Default	What it does
servicenow_war_old_days	30	What the report flags: a task open at the end of the week for more than this many days is amber (Older than N days).
servicenow_war_open_scope	worked	Which open tasks are listed (Open at the end of the week, the WAR's in-progress lines, ACT's evidence). worked (the default) = the ones new this week or updated in it - the week's work; the rest, the backlog nobody touched (a task open 90 days and not looked at), is counted in the tiles (Open at the end, Older than N... (all of it: the role file)
servicenow_war_cancelled_states	[Canceled, Cancelled, Closed Skipped, Closed Incomplete]	Cancelled is not completed. A task closed in the week in one of these states (as ServiceNow shows them, any case) - or with a close code holding one of these words - is not counted as completed, not in the WAR lines and not among ACT's accomplishments: it is listed in its own section at the bottom, "Cancelled this wee... (all of it: the role file)
servicenow_war_max_rows	300	rows per table in the report (the workbook has every row)
servicenow_war_notes_chars	200	how much of each close note the Completed table shows (0 = the column stays empty; the workbook has it whole)
servicenow_war_max_records	5000	at most this many records read per table (the rest are said)
servicenow_war_sla	true	The tasks' SLAs (task_sla): those breached, and those still running with servicenow_war_sla_warn_percent of their time or more used (close to breaching) - tiles, a section, the WAR lines, the workbook, ACT's evidence. The API account needs read on task_sla (the itil role has it); when it cannot read it the report says... (all of it: the role file)
servicenow_war_sla_warn_percent	75	an SLA at this share of its time or more is close to breaching (business time where it has a schedule)
servicenow_war_xlsx	true	Attach the Excel workbook to the email: every task with every column on one sheet, then Completed, Opened, Open at the end, By person, By type - for sorting, filtering and pivoting (whatever report_email_xlsx says). false = only with report_email_xlsx.

Setting	Default	What it does
servicenow_war_act	false	ACT (GenAI) writes the narrative - accomplishments, in progress, what deserves attention - from the week's tables, ready to paste, at the bottom after the one-line-per-task version (both stay, to choose from): attach an "ACT model key" credential (site_act_provider, site_act_models). It runs on the AAP side: this node... (all of it: the role file)
servicenow_war_act_tone	positive	The voice of ACT's narrative: positive (the default) = confident and forward-looking, the team's achievements and their impact first, open work as progress under way, issues as next steps - still only what the evidence says; neutral = factual, like an operations log; or your own words, e.g. "upbeat and brief, for the... (all of it: the role file)
servicenow_war_act_max_items	150	at most this many tasks per set go to the model (the biggest weeks)
servicenow_war_act_model_timeout	600	seconds the model may take for its answer
servicenow_war_act_check_url	true	first check that this AAP node reaches the model's URL (a clear message when not)
servicenow_war_fail	true	The job fails when ServiceNow could not be read at all, or a group named is not there. false = a report only (a table the account cannot read is always only said).

site_act

File: roles/site_act/defaults/main.yml

Setting	Default	What it does
site_act_level	explain	ACT (GenAI) on top of ANY check playbook. The check roles already did the cheap, certain part: they found WHAT is wrong (site_findings) and know which read-only command shows more (each finding's hint). This role hands that to ACT (the act_triage role) to find out WHY, and to propose - or, if you allow it, make - th... (all of it: the role file)

Setting	Default	What it does
use_act	false	The runbooks (Health check, Database health, Troubleshoot): true = ACT looks at what the job found. A job variable, per template or launch - not all.yml (every check job would then call the model).
site_act_enabled	true	ACT at all (1.3.0): false = no job calls ACT - no request, no model endpoint or key needed, no error; each report says "ACT is off (site_act_enabled false)" in one line and the job stays green (an enclave without a model). Usually in all.yml (docs/ENCLAVES.md).
site_act_allow	[]	self-heal only: regular expressions of fix commands ACT may run without a person. Each must match the WHOLE command. Keep them narrow, e.g. site_act_allow: ['systemctl restart (chronyd|rsyslog|crond)']
site_act_diagnose_only_groups	[aap, aap_hosts, sat, idm, netapp, netapp_console_hosts]	Hosts where ACT may only diagnose - never fix by itself, whatever site_act_level says: AAP itself (a fix could stop the job that is running it) and vendor-managed appliances. The apply step (act_fix_approved.yml) skips them too: a person applies the fix there by hand. AAP (groups aap, aap_hosts) is ALWAYS in this list... (all of it: the role file)
site_act_provider	genai	---- Which model ACT talks to ---------- site_act_provider genai GenAI.mil (the default) asksage Ask Sage - RETIRED FOR NOW: do not pick it (kept for when it returns) genai-beta the GenAI.mil beta proxy (preview models) site_act_models the model to use FOR EACH pr... (all of it: the role file)
site_act_models	{genai: gemini-4-argon-withtools, asksage: gpt-5.6-sol-gov}	the model for each provider (see above): the model follows the provider picked
site_act_urls	{}	the API URL for each provider; a provider not listed uses ACT's default
site_act_model	(empty)	one model for this run, whatever the provider (e.g. a survey answer); empty = site_act_models
site_act_url	(empty)	one URL for this run; empty = site_act_urls
site_act_ca	(empty)	a CA bundle file ON THE HOSTS for the provider's certificate; empty = the host's trust store

Setting	Default	What it does
site_act_env	{}	Any other ACT environment settings (advanced; these win over the ones above), e.g. {HTTPS_PROXY: "http://proxy.yoursite.mil:8080"}. See the top of vendor/act/act for the list. Every ACT job streams the model's answer (ACT_STREAM=1), so a firewall or proxy idle timeout does not cut a long answer ("connection reset by p... (all of it: the role file)
site_act_model_timeout	600	Seconds the model may take for ONE answer, retries after a provider hiccup included (ACT's GENAI_TIMEOUT; ACT's own default, 90, is too short for a long analysis or a busy provider). Every ACT job uses it: the runbooks' use_act, the VMware alarms and capacity analyses, service watch.
site_act_timeout	1200	seconds for one whole ACT run (several answers): keep it above site_act_model_timeout
site_act_concurrency	3	How many hosts may run ACT at the same time in one job (0 = no limit: as many as the job's forks). Every host in a job uses the same key, and GenAI.mil's default quota is 60 requests and 200,000 tokens a minute per key: a fleet run with many hosts at once gets HTTP 429 (too many requests).
site_act_batch	5	The jobs where ACT reads evidence (the VMware alarms analysis and capacity planning, the log analysis) ask it about this many items per question: a long answer about many items is what a gateway or proxy cuts off, or what runs past site_act_model_timeout. 15 items are 3 questions of 5, in one report. 0 = all in one qu... (all of it: the role file)
site_act_extra_instructions	(empty)	anything ACT should know about this site
site_act_evidence_text	(empty)	extra evidence (troubleshoot.yml passes what it collected)

site_email

File: `roles/site_email/defaults/main.yml`

Note
EMAIL A REPORT through your mail relay (SMTP). Off until report_email_to has an address. Used by every report job: health checks, certificate reports, database health, POA&M, ServiceNow health, troubleshooting, the VMware jobs (docs/EMAIL_REPORTS.md). The relay settings belong in playbooks/group_vars/all.yml. If the relay needs a login, attach an "SMTP relay" credential (aap/credential_types/smtp_relay.yml) - never put a password in these settings.

Setting	Default	What it does
report_email_to	[]	Who gets it: a list, or text with commas / spaces / one per line (a survey answer). Empty = no email. Set it on each TEMPLATE (its Variables, or a survey): in all.yml every report job emails.
report_email_cc	[]	copied: a list, or text like report_email_to
report_email_also	[]	"Also email the report to": more people for one run, asked when the job is launched - a survey question (playbooks/aap_email_question.yml adds it to the report templates). Added to report_email_to, never in its place; empty = only report_email_to. A list, or text: addresses separated by commas, semicolons, spaces or o... (all of it: the role file)
report_email_also_domains	[]	The domains an address in report_email_also may be in (a subdomain too), e.g. [yoursite.mil]; any other is listed in the job and gets no email. Empty = any domain (your mail relay may still refuse some).
report_email_from	(empty)	the sender, e.g. aap-noreply@yoursite.mil
report_email_smtp_host	(empty)	your mail relay, e.g. smtp.yoursite.mil
report_email_smtp_port	25	25 or 587 with starttls, 465 with ssl
report_email_security	starttls	starttls = encrypt after connecting (the relay must offer it); ssl = encrypted from the start (port 465); none = not encrypted (only for a relay without encryption; then no login is sent).
report_email_ca_path	(empty)	a CA file that signed the relay's certificate (empty = the CAs the execution environment trusts)
report_email_subject_prefix	"[AAP]"	put before every subject, e.g. "[AAP] Certificate report: ..."
site_enclave	(empty)	Which enclave this AAP serves: unclassified or classified; empty = found from the job's hosts - classified when a target host's name or address (inventory_hostname, ansible_host) ends in one of site_enclave_classified_domains, else unclassified. site_enclave, when set, always wins.

Setting	Default	What it does
site_enclave_classified_domains	[]	DNS suffixes of the classified enclave, e.g. [classified.example.mil]; [] = never found classified from the hosts.
report_banner	auto	The banner at the top and bottom of every report: each email look and the plain text (its first and last line), every report website page and the index, the ticket insights pages, print, and the first row of each workbook sheet. auto = on in the classified enclave (site_enclave), off in the unclassified; true / false... (all of it: the role file)
report_banner_text	SECRET	the banner's text, e.g. SECRET, SECRET//NOFORN, UNCLASSIFIED, CUI
report_banner_color	(empty)	The banner's colour (#rrggbb); empty = the standard colour of the text: UNCLASSIFIED green, CUI purple, CONFIDENTIAL blue, SECRET red, TOP SECRET orange; the text black or white, whichever reads 4.5:1 or better.
report_banner_subject	true	with the banner on: "(SECRET) " before the email subject (and the prefix)
report_email_html	true	Send a formatted (HTML) version too, for reports that come with one (all VMware jobs). The email always carries a plain-text copy for mail clients that do not show HTML. false = plain text only.
report_email_xlsx	false	Attach the report's tables to the email as an Excel workbook (.xlsx): a sheet per table. Some reports attach their own whatever this says (VM - inventory report: every VM on one sheet). The website gets no workbook.
report_email_theme	standard	The look of the HTML version (docs/EMAIL_REPORTS.md, "The look"): standard = a header band in the report's colour (green, amber, red, blue) with number boxes under it; command = a dark header with your logo and report_email_header_text, the job's last runs as a strip of bars (report_history), the numbers in one row; p... (all of it: the role file)
report_email_header_text	Site automation report	command, portal, glass: the line with the logo

Setting	Default	What it does
report_logo_file	(empty)	Your logo in the header (every look but standard): a PNG, JPEG or GIF in your repository - a path from playbooks/ (files/report_logo.png = playbooks/files/report_logo.png) or a full path. Attached to each email, so keep it small: about 160 px tall, under 100 KB. Updating with scripts/update-from-release: list the file... (all of it: the role file)
report_logo_height	84	pixels in the email (its width follows); portal: at most 56, glass: at most 80
report_history	true	command, portal, glass: the job's last runs, read from AAP - attach a "Red Hat Ansible Automation Platform" credential to the template, whose user can see the template's jobs (Organization Auditor on its organization, docs/EMAIL_REPORTS.md); without it the email goes out without the strip.
report_history_runs	14	bars in the strip, this run included
report_history_aap_ca_path	(empty)	a CA file for AAP's own HTTPS certificate ("" = what the execution environment trusts)
report_email_html_template	report.html.j2	Your own look instead: copy roles/site_email/templates/report.html.j2 (or report_command / _portal / _glass) to e.g. playbooks/files/email/report.html.j2, change it, and set report_email_html_template: "{{ playbook_dir }}/files/email/report.html.j2" report.html.j2 (the default) = the look report_email_theme picks.
report_time_zone	(empty)	The time zone of the times in a report, besides UTC: e.g. America/New_York (a name as `timedatectl list-timezones` prints them) shows "2026-10-25 10:03:12 EDT (14:03 UTC)". vCenter, ESXi and their log files keep UTC; the vSphere Client shows the browser's time zone. "" = UTC only. Used by the VMware health and alarms... (all of it: the role file)

Setting	Default	What it does
report_web_dest	(empty)	THE REPORT WEBSITE (docs/WEB_REPORTS.md): every report job also writes its report as a web page - off until report_web_dest is set. Works without an email (report_email_to empty). Each run: its page, a copy as the template's latest page, the report's data as JSON beside each (for your own site), and an index of every... (all of it: the role file)
report_web_url	(empty)	the address people open that folder at, e.g. https://reports.yoursite.mil ("" = dest when it is https://)
report_web_theme	atlas	The website's look (docs/WEB_REPORTS.md, "The look"): atlas = made for a browser, never an email - a wide page whose tables use all of it (columns sized by what they hold: numbers right-aligned, short values whole, paths broken at / and \), a bar at the top that stays (the sections, Compact / Comfortable rows, Light /... (all of it: the role file)
report_web_page	"{template}/{date}_{time}_job-{job}.html"	The names inside the folder, with {names} filled in: {template} (the AAP template's name, lower case with dashes; the play's name outside AAP), {template_id}, {job} (the AAP job; local outside AAP), {status} (ok, warning, critical, info), {date} (2026-10-23), {time} (0605), {year}, {month}, {day} - the time the job wr... (all of it: the role file)
report_web_latest	"{template}/latest.html"	always the template's newest run ("" = none)
report_web_index	index.html	every report's last runs, a tab per area ("" = none)
report_web_history	5	the runs each report keeps: listed under its line of the index, the pages of older runs deleted
report_web_delete_old	true	false = the older runs' pages stay in the folder (only the index forgets them)

Setting	Default	What it does
report_web_index_prune	true	A line of the index whose latest page is no longer there (its folder deleted by hand) is taken off the index by the next report that writes it. false = lines stay until a report writes them again. Not with report_web_host (that server's folder is not read from the AAP side): there, delete the line from index.json by h... (all of it: the role file)
report_web_tabs	- {tab: Certificates, match: ["*certificate*", "*trust store*", "PKI *"]} - {tab: VMware,…	The index page's tabs: a report is on the first tab one of whose names its AAP template's name matches (any case, * = anything); else on the tab of its name up to " - " (Backup - nightly: Backup); else on Other. A tab with no report is not shown. Set it in all.yml, for every template alike: the index is one page.
report_web_names	{}	your own {names}, e.g. {site: "{{ site_code }}"} for "{site}/{template}/..."
report_web_json	true	the report's data as .json beside each page and index.json beside the index
report_web_smb_auth	negotiate	A share (report_web_dest \\server\share\folder): how the job logs in, and the connection. negotiate (Kerberos - FIPS too: user@YOURDOMAIN.MIL, the server's full name - else NTLM), ntlm, or kerberos (a ticket already in the job)
report_web_smb_encrypt	false	true = encrypt (SMB 3) even when the share does not require it; one that requires it is always
report_web_smb_port	445	the share's port
report_web_ca_path	(empty)	a CA file for an https:// report_web_dest, or report_web_host's WinRM certificate ("" = what the EE trusts)
report_web_retries	3	A share, an upload address or a folder that fails for a reason that passes (Kerberos "Message stream modified" while a file server's machine password changes, no KDC reached, a connection reset, a session gone, HTTP 503): tried again from a new login. A real error (password, access, no such share or folder) is not. Tr... (all of it: the role file)
report_web_retry_delay	30	seconds between two tries

Setting	Default	What it does
report_web_required	false	false = still not published after the retries: a warning, the email says so at the top, the job goes on; true = the job fails
report_web_host	(empty)	Or on a Windows server (a file server, or the IIS server itself): report_web_host is that server, report_web_dest the folder on it (D:\reports, there already). The pages are made on the AAP side and written there over WinRM, in one PowerShell script per run (Constrained Language Mode too): nothing to mount, no WebDAV.... (all of it: the role file)
report_web_host_connection	winrm	winrm or psrp
report_web_host_port	5986	WinRM: 5986 (https), 5985 (http)
report_web_host_scheme	https	https, or http (Kerberos and NTLM encrypt each message then)
report_web_host_transport	ntlm	ntlm, kerberos (Kerberos in the EE), credssp or basic
report_web_host_cert_validation	validate	validate (the CA: report_web_ca_path) or ignore
report_web_copy	false	THE COPY JOB instead (docs/WEB_REPORTS.md, "The copy job"): for report templates whose execution environment cannot reach the web server (no WinRM or Kerberos in it). true = the report jobs write nothing there - each keeps its report in AAP (the job artifact site_report_page) and the one job template "Report website -... (all of it: the role file)
report_web_copy_max_kb	256	the kept report's greatest size, packed (a bigger one: its longest tables cut, said on the page)
report_web_copy_jobs	[]	The copy job itself (playbooks/report_web_copy.yml): it reaches report_web_host as a host of its own inventory when it is one (that host's connection settings and the template's Machine credential - your Windows inventory), else with the report_web_host_* settings above. Its "Red Hat Ansible Automation Platform" crede... (all of it: the role file)

Setting	Default	What it does
report_web_copy_hours	24	its first run (no copy before it): the report jobs of the last 24 hours
report_web_copy_max_pages	50	pages a run writes at most; the rest are the next run's

site_findings

File: roles/site_findings/defaults/main.yml

Note
Shared by every check playbook. A "finding" is one thing that is wrong on one host: check: the check that found it (disk, services, mariadb, ...) id: a stable name for this problem on this host, e.g. disk:/var. ServiceNow uses it to update the same ticket instead of opening a new one every run. severity: critical \| warning summary: one line a person understands ("/var is 96% full (critical at 95%)") hint: a read-only command that shows more. A person can paste it into a shell; ACT (when you add it later) runs it to collect evidence. Check roles only ADD findings. This role starts the list, reports it, and decides pass/fail.

Setting	Default	What it does
site_fail_on	[critical, warning]	Which severities make the job FAIL. A failed job is what sends an AAP workflow down its "on failure" path (open tickets, ask ACT, wait for approval). [] = report only, never fail.
site_unreachable_fails	false	A host that cannot be reached (switched off, no network, ssh or WinRM refused, its login refused) in the health checks, the certificate report and renewal, database health, the service watch and troubleshoot: listed as NOT REACHED and left out, and the job is not failed for it. It still counts as checked-for, so the S... (all of it: the role file)
site_syslog	true	One syslog line per finding (tag below), so your log server / SIEM sees them too.
site_syslog_tag	site-health	the syslog tag those lines carry (to filter on in the SIEM)
site_eventlog_source	site-health	Windows: one Application event log entry per finding instead of syslog (source below).
site_os	linux	---- Linux or Windows (set by the playbook; do not set) ---------- Linux playbooks run the roles check_<name>, the Windows ones (win_*.yml) the roles win_check_<name>.
site_check_role_prefix	"{{ 'win_check_' if site_os == 'windows' else 'check_' }}"	check_ (Linux) or win_check_ (Windows): set by the playbook from site_os

Setting	Default	What it does
site_windows_connections	[winrm, psrp, ansible.builtin.winrm, ansible.builtin.psrp]	How a host is recognized as Windows: its connection, or its shell (Windows OpenSSH).
site_settings_notice	true	A notice in the job output when a project settings file (playbooks/group_vars/*.yml) defines a setting: it wins over AAP's inventory Variables box. false = no notice.

stigman_ckl_import

File: roles/stigman_ckl_import/defaults/main.yml

Setting	Default	What it does
stigman_ckl_import_src	(empty)	STIG MANAGER - IMPORT CHECKLISTS (playbooks/stigman_ckl_import.yml): workflow 40's check. It reads the checklists (CKL from STIG Viewer 2, CKLB from STIG Viewer 3 - VMware's InSpec checklists, Evaluate-STIG's, a person's) from your share, finds each one's asset by its host name (HOST_NAME) and its STIG among those ins... (all of it: the role file)
stigman_ckl_import_patterns	["*.ckl", "*.cklb"]	the files taken (any case)
stigman_ckl_import_recurse	true	read the sub-folders too (one per host, one per month ...)
stigman_ckl_import_user	(empty)	the share's account (user@YOURDOMAIN.MIL); empty = the "Windows file share" credential's
stigman_ckl_import_password	(empty)	its password (vault it); empty = the credential's
stigman_ckl_import_auth	negotiate	negotiate (Kerberos - FIPS too: user@YOURDOMAIN.MIL, the server's full name - else NTLM), ntlm, or kerberos (a ticket already in the job)
stigman_ckl_import_encrypt	false	true = encrypt (SMB 3) even when the share does not require it
stigman_ckl_import_port	445	the share's port
stigman_ckl_import_max_files	500	at most this many files (more: nothing read - name a narrower folder)
stigman_ckl_import_max_mb	50	a file larger than this is left out (and said)
stigman_ckl_import_collection	(empty)	the STIG Manager collection (its name) the checklists' assets are in - or are created in (the staging collection)

Setting	Default	What it does
stigman_ckl_import_collections	[]	Where else a host's asset is looked for (0.76.0): collection names or ids, * and ? ('*' = every collection the client reads). Its reviews go to each of them whose asset has the STIG matched; the staging collection above is searched last, and written to only when none of these has the asset; a host in none of them is c... (all of it: the role file)
stigman_ckl_import_create_assets	true	A host name with no asset in the collection: the asset is planned too - created once approved, with every STIG its checklists matched (the update's client needs the Manage grant). false = listed, not written: create it yourself.
stigman_ckl_import_routes	[]	Where each host goes (1.26.0): an ordered list of rules, the first that fits a host wins. A rule: collection (its name or id) and any of stigs (the installed STIG its checklist matched, its id or title: "*Windows_Server*", "RHEL_*", "VMW_*"), names (the host name or FQDN, or its short name: "esx*", "vc*"), roles and t... (all of it: the role file)
stigman_ckl_import_route_move	staging	An asset only in the staging collection whose rule names another collection: staging = moved there (after the approval, or at once with stigman_ckl_import_auto: all). A move keeps its reviews and their history, its STIGs, metadata and labels; a Restricted grant in the staging collection that reached it would lose it -... (all of it: the role file)
stigman_ckl_import_assign_stigs	false	An asset that is there without the STIG its checklist matched: false = listed (assign it yourself); true = the assignment is planned, listed and written with the approval (or at once with auto: all), then its reviews. The client needs Manage there.
stigman_ckl_import_benchmarks	{}	A checklist's STIG id -> the installed STIG to write to, instead of matching by rules (fnmatch, any case), e.g. {"VMware Aria Operations 8.x Apache Profile": "VMW_Aria_Operations_8-x_Apache_STIG"}
stigman_ckl_import_min_match	0.8	the share of a checklist STIG's rules the installed STIG must have (0.8 = 80%)
stigman_ckl_import_candidates	25	at most this many installed STIGs read to match one (the likeliest by name first)
stigman_ckl_import_search_all	false	true = no match among those: read every installed STIG (slow with hundreds)
stigman_ckl_import_other_revision	false	true = write a rule the checklist gives another revision's id of (SV-...r123 vs r456); false = listed, held back
stigman_ckl_import_overwrite_accepted	false	true = write over an Accepted review too; false = listed, left for the reviewers

Setting	Default	What it does
stigman_ckl_import_status	submitted	the reviews' status: saved, submitted or accepted
stigman_ckl_import_detail_max	8000	a longer finding detail is cut
stigman_ckl_import_list_max	200	at most this many rows in the report's tables
stigman_ckl_import_workflow	false	true on workflow 40's check only: the plan handed on (set_stats stigman_update_plan), the job fails on purpose (NEEDS APPROVAL)
stigman_ckl_import_auto	none	Written at once, no approval (0.76.0): none (everything waits for workflow 40's approval), existing (the reviews of assets that are there are written by this job; a new asset and its reviews wait for the approval) or all (the new assets too). The same rules as the approval's write: never over Accepted, never with atta... (all of it: the role file)
stigman_ckl_import_move	false	Move each file done with - all of it in STIG Manager: written, the same already, or left for good (Accepted, not in the STIG, another revision) - to <stigman_ckl_import_move_to>/<date>/ under stigman_ckl_import_src, if not written again since it was read (the share's account then needs Modify). One that waits for the... (all of it: the role file)
stigman_ckl_import_move_to	imported	the folder under stigman_ckl_import_src (never read)

stigman_health

File: roles/stigman_health/defaults/main.yml

> **Note**
> STIG MANAGER - HEALTH and AVAILABILITY (playbooks/stigman_health.yml, playbooks/stigman_availability.yml), read-only, through STIG Manager's API from the AAP side. docs/STIGMAN_HEALTH.md

Setting	Default	What it does
stigman_health_api	(empty)	The address people use, e.g. https://stigman.example.mil/api (the load balancer's, or the one name the nodes share); empty = the nodes only.

Setting	Default	What it does
stigman_health_nodes	[]	Each node (an API server of STIG Manager) on its own: [{name: stig01, url: "https://stig01.example.mil/api"}]. [] = each host of the inventory group stigman_health_nodes_group, at stigman_health_node_url ({host} = its inventory name; a host's own stigman_health_node_url in host_vars wins).
stigman_health_nodes_group	stigman	the inventory group of the nodes
stigman_health_node_url	"https://{host}/api"	each node's API ({host} = its inventory name)
stigman_health_db_group	stigman_db	The database host(s): the inventory group whose hosts the health check also reads over SSH (the MariaDB container, its disk) - with the nodes' hosts (their containers, their disk). stigman_health_hosts: false = the API only. the inventory group of the database host(s)
stigman_health_hosts	true	false = the API only, no SSH to the hosts
stigman_health_token_url	(empty)	The health check logs in as a Keycloak client (client credentials) - from the "STIG Manager API" credential (STIGMAN_TOKEN_URL, STIGMAN_CLIENT_ID, STIGMAN_CLIENT_SECRET), or these; no client = the checks without a login only. Keycloak's token endpoint; empty = the credential's
stigman_health_client_id	(empty)	the client; empty = the credential's
stigman_health_client_secret	(empty)	vault it
stigman_health_scope	(empty)	e.g. "stig-manager:user:read stig-manager:collection:read stig-manager:op:read"; empty = the client's default scopes
stigman_health_appinfo	true	The deployment's own numbers (/op/appinfo?elevate=true: the database's connections and uptime, each node's memory and uptime): the client needs the scope stig-manager:op:read and admin in STIG Manager. false = not asked.
stigman_health_validate_certs	true	check the HTTPS certificates of STIG Manager and Keycloak
stigman_health_ca_path	(empty)	a CA bundle on the AAP side to check them with (the DoD roots); empty = the execution environment's
stigman_health_timeout	20	seconds per request

Setting	Default	What it does
stigman_health_slow_ms	2000	an answer slower than this, in milliseconds -> warning
stigman_health_node_down	critical	a node that does not answer, or is not "available" (warning = only the address counts)
stigman_health_conn_warn_pct	80	the database's connections in use, % of max_connections -> warning
stigman_health_conn_crit_pct	95	-> critical
stigman_health_restart_minutes	60	a node or the database restarted less than this ago -> warning
stigman_health_version	(empty)	the STIG Manager version every node must run, e.g. 1.6.18; empty = only that they are alike
stigman_availability_nodes	true	The availability check (every 5 to 15 minutes) asks /op/state of the address and of each node: the API, its database and Keycloak all up. false = the address only (the nodes are then not asked).

stigman_stack

File: roles/stigman_stack/defaults/main.yml

Note
STIG Manager stack on podman, managed by systemd through Quadlet unit files: stigman-mysql (MySQL 8.4) <- stigman-api (STIG Manager) <- stigman-nginx (TLS, port 443) Keycloak is EXTERNAL: the realm already exists and this role never touches it. Needs podman >= 4.4 (Quadlet): RHEL 9.2+ / RHEL 8.8+.

Setting	Default	What it does
stigman_scope	system	system = rootful (/etc/containers/systemd, systemctl) <- production user = rootless (~/.config/containers/systemd, systemctl --user) <- tests / rootless hosts
stigman_images	mysql: docker.io/library/mysql:8.4 api: docker.io/nuwcdivnpt/stig-manager:1.6.17 # pin a…	the container images (pin STIG Manager's version; change it deliberately)
stigman_config_dir	/etc/stigman	nginx.conf, TLS files, CA bundle
stigman_https_port	443	host port nginx publishes
stigman_server_name	"{{ inventory_hostname }}"	the name nginx answers to (and the TLS certificate is for)

Setting	Default	What it does
stigman_oidc_provider	"https://keycloak.example.mil/realms/stigman"	Existing Keycloak realm (server-side URL used by the API; browser-side URL for the web client)
stigman_client_oidc_provider	"{{ stigman_oidc_provider }}"	the Keycloak realm URL the browser uses (when it differs from the server-side one)
stigman_client_id	stig-manager	the Keycloak client the web app logs in with
stigman_classification	U	banner: U, CUI, C, S, TS, SCI, NONE
stigman_db_schema	stigman	the MySQL database STIG Manager uses
stigman_db_user	stigman	the MySQL account STIG Manager uses (its password: the credential)
stigman_mysql_buffer_pool	1G	STIG Manager docs: >= 8G for real deployments
stigman_ready_retries	40	How long to wait after (re)starting: for MySQL to be healthy, then for the API to answer through nginx (each)
stigman_ready_delay	6	seconds between those tries
stigman_mysql_root_password	"{{ lookup('ansible.builtin.env', 'STIGMAN_MYSQL_ROOT_PASSWORD') }}"	---- Secrets: NEVER set these in a vars file or inventory ---------- They come from AAP credentials (aap/credential_types/*.yml) as environment variables / files on the controller (execution environment). See docs/SECRETS.md.
stigman_db_password	"{{ lookup('ansible.builtin.env', 'STIGMAN_DB_PASSWORD') }}"	from the credential (STIGMAN_DB_PASSWORD), never set here
stigman_tls_cert_src	"{{ lookup('ansible.builtin.env', 'TLS_CERT_FILE') }}"	nginx's certificate file, from the TLS certificate credential (TLS_CERT_FILE)
stigman_tls_key_src	"{{ lookup('ansible.builtin.env', 'TLS_KEY_FILE') }}"	its private key file, from the credential (TLS_KEY_FILE)
stigman_tls_ca_src	"{{ lookup('ansible.builtin.env', 'TLS_CA_FILE') }}"	the CA bundle file, from the credential (TLS_CA_FILE)

Setting	Default	What it does
stigman_replace_secrets	false	A deploy refuses to change a database password secret already deployed: in MySQL the password stays the old one, and STIG Manager would lose its database. true = replace the secrets anyway - ONE launch only, when the database is being created again (Prompt on launch, Variables); never save it in the template.

stigman_update

File: roles/stigman_update/defaults/main.yml

Setting	Default	What it does
stigman_update_api	"{{ check_stig_stigman_api | default('') }}"	STIG MANAGER - UPDATE (playbooks/stigman_update.yml): workflow 37's (and 40's) step after the approval. It writes the reviews the STIG recheck's (or the checklist import's) check listed (stigman_update_plan, handed on by the workflow) - each only if STIG Manager still has it as the check read it - and the new assets i... (all of it: the role file)
stigman_update_token_url	"{{ check_stig_stigman_token_url | default('') }}"	empty = the credential's
stigman_update_client_id	"{{ check_stig_stigman_client_id | default('') }}"	empty = the credential's
stigman_update_client_secret	"{{ check_stig_stigman_client_secret | default('') }}"	empty = the credential's (vault it)
stigman_update_scope	(empty)	scopes to ask for, e.g. "stig-manager:collection"; empty = the client's default scopes
stigman_update_validate_certs	"{{ check_stig_stigman_validate_certs | default(true) }}"	check the HTTPS certificates
stigman_update_ca_path	"{{ check_stig_stigman_ca_path | default('') }}"	a CA bundle on the AAP side to check them with
stigman_update_timeout	"{{ check_stig_stigman_timeout | default(60) }}"	seconds per request

Setting	Default	What it does
stigman_update_skip_attachments	"{{ check_stig_update_skip_attachments | default(true) }}"	a review with attachments in STIG Manager (evidence a person attached) is never written; false = write it like the others

tenable_scan

File: roles/tenable_scan/defaults/main.yml

Note
ACAS SCAN (playbooks/tenable_scan.yml, docs/TENABLE.md): kick off a Tenable.sc (ACAS) scan of the hosts of an inventory group, from the scanner of each host's own datacenter - Tenable.sc's scan zones say which scanners cover which IP ranges - one scan per zone. From the AAP side, with the "Tenable.sc API" credential. Nothing is launched without tenable_scan_policy and tenable_scan_repository: a first run lists the zones, scanners, policies and repositories the account may use. A dry run (Check) chooses the zones and says what it would launch.

Setting	Default	What it does
tenable_scan_target	(empty)	The hosts: a group or pattern of this inventory (the job runs on localhost: not the Limit), and/or names or addresses not in the inventory. "" and [] = nothing to scan: the job lists what Tenable.sc offers. a group or pattern of the inventory (web_servers, dc1_*)
tenable_scan_targets	[]	names or addresses not in the inventory, e.g. [198.51.100.5, nas01.example.mil]
tenable_scan_policy	(empty)	the scan policy (name or id): the first run lists them
tenable_scan_repository	(empty)	the repository the results go to (name or id)
tenable_scan_credentials	[]	scan credentials (names or ids) the scans log in to the hosts with
tenable_scan_name	"AAP {target} - {zone}"	The scans' name: {target} = tenable_scan_target (or the number of hosts), {zone} = the zone. A scan of that name from an earlier run is updated (its hosts) and launched again, not made twice.
tenable_scan_description	"Launched by site-automation (AAP), playbooks/tenable_scan.yml"	the scans' description in Tenable.sc

Setting	Default	What it does
tenable_scan_sites	{}	The zone (the scanners) of each host - the first rule that applies: 1. the host's own tenable_zone (an inventory variable: a zone's name) 2. its site, through tenable_scan_sites: a group of the host, or the site code tenable_scan_site_regex takes from its name, mapped to a zone's name - e.g. {dc_east: "Datacenter East... (all of it: the role file)
tenable_scan_site_regex	(empty)	e.g. '^([a-z]{2}[0-9])': the first group is the site code in the host's name (xy2vmp02 -> xy2)
tenable_scan_zone_default	(empty)	rule 4: a zone's name, auto, or "" (a host no rule gives a zone is not scanned)
tenable_scan_wait	false	true = wait for the scans and report the findings by severity and by host
tenable_scan_wait_minutes	120	at most this long
tenable_scan_poll_seconds	60	seconds between two looks at a running scan
tenable_scan_max_targets	256	at most this many hosts per run (narrow tenable_scan_target, or raise it)
tenable_scan_max_rows	200	rows per table in the report
tenable_scan_fail	true	The job fails when Tenable.sc could not be read, a policy or repository named is not there, or a scan could not be launched. false = a report only. A host no rule gives a zone is said (amber), never a failure.
tenable_ca_path	(empty)	a CA file for Tenable.sc's HTTPS certificate ("" = what the execution environment trusts)
tenable_timeout	60	seconds per request

tmp_cleanup

File: roles/tmp_cleanup/defaults/main.yml

Setting	Default	What it does
tmp_cleanup_dirs	- {path: /tmp, days: 10} - {path: /var/tmp, days: 30}	TMP CLEANUP: the old files in the tmp folders - measured by the disk space report on a host whose filesystem is full, deleted by "Linux - tmp cleanup" after the approval (workflow 42). Old = accessed, changed and modified more than days ago (as systemd-tmpfiles ages them; 10 days for /tmp and 30 for /var/tmp are its o... (all of it: the role file)
tmp_cleanup_keep	- 'systemd-private-*' - '.X11-unix' - '.ICE-unix' - '.XIM-unix' - '.font-unix' - '.Test-u…	Never touched (names, * and ? allowed), wherever they are under a tmp folder: what systemd, X, ssh, Kerberos and the databases keep there while they run.
tmp_cleanup_confirm	false	On its own (not in workflow 42): it deletes only with tmp_cleanup_confirm: true; without it, it lists what it would delete. In the workflow it cleans the hosts the disk space report planned, after the approval.
tmp_cleanup_timeout	600	seconds per folder

troubleshoot

File: roles/troubleshoot/defaults/main.yml

Setting	Default	What it does
ts_area	overview	TROUBLESHOOT: the commands an experienced admin runs first for a kind of problem, run for you (read-only), each with a line saying what it shows. The matching health checks run too, so you also get plain-words findings. Pick the area; for "service" and "network" say what. ts_area: overview | disk | performance | servi... (all of it: the role file)
ts_service	(empty)	the systemd unit, for ts_area service (see above)
ts_target	(empty)	host:port, for ts_area network
ts_since	"-2h"	how far back logs are read (journalctl --since)
ts_output_lines	60	keep the last this-many lines of each command
ts_areas	overview: checks: [disk, mounts, services, performance, time, logging] commands: - {why:…	What each area runs. why is printed above the command's output. Commands must be read-only.

ucs_report

File: `roles/ucs_report/defaults/main.yml`

> **Note**
> Cisco UCS inventory (playbooks/ucs_inventory_report.yml). Read-only: XML API queries to each UCS Manager domain over HTTPS (port 443), with the account of a "Cisco UCS Manager" credential - its read-only role is enough (docs/UCS.md).

Setting	Default	What it does
`ucs_domains`	`[]`	Each domain's UCS Manager address: the fabric interconnects' cluster (virtual) address, by name or IP, e.g. [ucs01.yoursite.mil, ucs02.yoursite.mil]. The credential's account must be able to log in to each of them.
`ucs_validate_certs`	`true`	Check UCS Manager's HTTPS certificate. Its own certificate is self-signed unless a key ring with your CA's certificate was made: then put that CA in ucs_ca_path (a file in the project, e.g. "{{ playbook_dir }}/files/ca/ucs-ca.pem"), or set false.
`ucs_ca_path`	(empty)	the CA file (PEM) UCS Manager's certificate is checked with; empty = the execution environment's own trust
`ucs_timeout`	`60`	seconds for one request (no answer: asked once more)
`ucs_tls_legacy`	`false`	true = connect to an older UCS Manager whose TLS the execution environment refuses (UNSAFE_LEGACY_RENEGOTIATION_DISABLED, KEY_TOO_SMALL, SHA-1)
`ucs_inventory_max_rows`	`2000`	servers, chassis and faults listed in the email and the page at most (the workbook always has all)
`ucs_inventory_xlsx`	`true`	attach every fabric interconnect, chassis, IO module, server, adapter, CPU, DIMM, storage controller, power supply, firmware, service profile and fault as an Excel workbook (.xlsx) to the email
`ucs_inventory_fault_severities`	`[critical, major]`	UCS Manager's faults listed in the report (all of them are in the workbook); [] = no fault table
`ucs_inventory_fail`	`false`	true = mark the job FAILED when a domain could not be read (the report still goes out), for workflows
`ucs_inventory_raw`	`false`	true = also print the first objects of each class as UCS Manager gave them (to see what a field is called on your UCS Manager)

vmware_ask_stack

File: `roles/vmware_ask_stack/defaults/main.yml`

Note
VMware ask - the web page of "VM - ask a question" - one container on podman, managed by systemd through a Quadlet unit file: vmware-ask-app (the service: questions, GenAI through ACT, the vCenter reads). It sits behind the reverse proxy you have - STIG Manager's nginx, in its podman container - at /ask/: that nginx does TLS and the client certificates (mTLS, the CAC), and passes who it verified. The role joins the service to that nginx's podman network and writes the nginx lines to add (vmware_ask_config_dir/nginx-ask.conf); it never changes the nginx, its files, its network or any other container (docs/VMWARE_ASK.md). Needs podman >= 4.4 (Quadlet): RHEL 9.2+ / RHEL 8.8+.

Setting	Default	What it does
`vmware_ask_scope`	`system`	system = rootful (/etc/containers/systemd, systemctl) <- production user = rootless (~/.config/containers/systemd, systemctl --user) <- tests / rootless hosts
`vmware_ask_image`	`aap.example.mil/site-ee-all:1`	the service's image: the site's execution environment (ee/all-in-one) - Python, ansible-core and pyVmomi are in it already, so nothing is installed on the server or downloaded at start. Pull it from your automation hub.
`vmware_ask_config_dir`	`/etc/vmware-ask`	the service's files, and nginx-ask.conf (the nginx lines to add)
`vmware_ask_network`	(empty)	the podman network of the reverse proxy's nginx container (required): `sudo podman ps --format "{{.Names}} {{.Networks}}"` shows it. It must exist with DNS on (a network you made, not podman's default "podman"); the service joins it, the role never creates or changes it
`vmware_ask_path`	`/ask/`	the page's path on the reverse proxy (nginx-ask.conf's location)
`vmware_ask_public_url`	`"https://{{ inventory_hostname }}{{ vmware_ask_path }}"`	the page's address, as people open it (said at the end, and logged)
`vmware_ask_title`	`Ask VMware`	the page's title
`vmware_ask_classification`	`U`	banner: U, CUI, C, S, TS, SCI, NONE
`vmware_ask_vcenters`	`"{{ vmware_vcenters | default([]) }}"`	vCenter: the "VMware vCenter" credential's account (read-only is enough) and vCenter, and these vCenters with it more vCenters, the same account (as vmware_vcenters)
`vmware_ask_validate_certs`	`"{{ vmware_validate_certs | default(true) }}"`	check vCenter's certificate (the CAs: vmware_ask_trust_host_cas)
`vmware_ask_datacenter`	(empty)	only this datacenter; "" = all

Setting	Default	What it does
vmware_ask_refresh_minutes	10	read vCenter every this many minutes (the answers are as of the last read)
vmware_ask_refresh_timeout	1800	seconds one read may take
vmware_ask_event_days	"{{ vm_ask_event_days | default(30) }}"	a VM's last power off, restart and vMotion: from vCenter's events of this many days
vmware_ask_max_events	"{{ vm_ask_max_events | default(50000) }}"	at most this many of those events per vCenter
vmware_ask_genai	true	GenAI: ACT's settings (site_act_provider, site_act_models, site_act_urls ...: roles/site_act/defaults/main.yml) and the key of the "ACT model key" credential. No key: the job's own reader answers the simplest questions. false = only the job's own reader
vmware_ask_local_fallback	"{{ vm_ask_local_fallback | default(true) }}"	GenAI did not run or gave no query: the job's own reader tries it (said)
vmware_ask_genai_retry	"{{ vm_ask_genai_retry | default(true) }}"	GenAI's answer held no query: asked once more, told what it answered
vmware_ask_model_timeout	"{{ vm_ask_model_timeout | default(120) }}"	seconds GenAI may take to answer
vmware_ask_extra_instructions	"{{ vm_ask_extra_instructions | default('') }}"	more words for GenAI, e.g. your naming
vmware_ask_max_rows	"{{ vm_ask_max_rows | default(500) }}"	rows shown on the page (the CSV has every row)
vmware_ask_keep_days	365	keep who asked what this many days (0 = for ever)
vmware_ask_parallel	4	questions with GenAI at once (the rest wait)
vmware_ask_per_user	2	questions one person may have running at once
vmware_ask_trust_host_cas	true	the service trusts what the server trusts (/etc/pki/ca-trust): vCenter, GenAI
vmware_ask_ready_retries	30	How long to wait after (re)starting for the service to answer
vmware_ask_ready_delay	4	seconds between those tries

Setting	Default	What it does
vmware_ask_vcenter_host	"{{ lookup('ansible.builtin.env', 'VMWARE_HOST') }}"	---- Secrets: NEVER set these in a vars file or inventory ---------- They come from AAP credentials as environment variables / files on the controller (execution environment). the "VMware vCenter" credential
vmware_ask_vcenter_user	"{{ lookup('ansible.builtin.env', 'VMWARE_USER') }}"	its account
vmware_ask_vcenter_password	"{{ lookup('ansible.builtin.env', 'VMWARE_PASSWORD') }}"	its password: a podman secret on the server

vmware_vm

File: roles/vmware_vm/defaults/main.yml

Note
VMWARE VMs: jobs that run against vCenter (not on the VMs), with the certified vmware.vmware collection. vCenter address and account come from a "VMware vCenter" credential on the job template (VMWARE_HOST, VMWARE_USER, VMWARE_PASSWORD). docs/VMWARE.md

Setting	Default	What it does
vm_names	[]	The VMs, by their exact name in vCenter: a list, or (from a survey) one per line. No wildcards.
vmware_datacenter	(empty)	only needed when the same VM name exists in two datacenters
vmware_validate_certs	true	Check vCenter's certificate. false only if the execution environment does not trust your CA yet.
vmware_vcenters	[]	Several vCenters (e.g. Enhanced Linked Mode), one report: the reports and the ESXi security settings read the credential's vCenter AND these (names, or name:port), with the same account. The same vCenter under two names is read once. docs/VMWARE.md, "Several vCenters". e.g. [vc01.yoursite.mil, vc02.yoursite.mil]
vmware_vcenter_fail_unread	false	true = the job shows failed when a vCenter could not be read (the report lists it in red either way)

Setting	Default	What it does
vmware_protected_groups	[aap, aap_hosts]	The AAP server is ALWAYS refused (hosts in the groups aap / aap_hosts, matched by VM name, guest host name and IP address), and so are the hosts of these groups and these VM names (* and ? allowed: DC*). List the AAP VM's name here when it differs from its host name, and vCenter's own VM. The snapshot cleanup and the... (all of it: the role file)
vmware_protected_vms	[]	e.g. [AAP01, VCSA01, DC*] (* and ? allowed)
vm_power_mode	guest	guest = ask the operating system (needs VMware Tools running; services stop cleanly) hard = reset / power off, like the button on a physical server (unsaved data is lost)
vm_power_timeout	600	shut down: seconds to wait for the VM to be off
vm_snapshot_name	(empty)	The snapshot's name. Taking one: empty = aap-<date>-<time>. Deleting: the exact name to delete.
vm_snapshot_description	(empty)	the AAP job number and user are added
vm_snapshot_quiesce	false	true = VMware Tools freezes the file systems first
vm_snapshot_delete_all	false	delete: true = every snapshot of the VM (vm_snapshot_name empty)
vm_notes	(empty)	The text to write.
vm_notes_mode	append	append (a new line, with date and user) | replace | clear
vm_nic	1	Which adapter: 1 = Network adapter 1, 2, ...
vm_portgroup	(empty)	the port group (standard or distributed) to move it to
vm_secure_boot_folder	(empty)	Only the VMs in this vCenter folder ("" = every VM, or the ones in vm_names).

Setting	Default	What it does
vm_secure_boot_fail	false	The job stays green (successful) when it finds VMs to fix: the report and email say what. true = mark the job FAILED when a VM is BIOS or has secure boot off, so a workflow's "on failure" path can react (open a ticket).
vm_datastore_warn_pct	85	Used space (% of capacity): a datastore at or above these is listed as a warning / critical.
vm_datastore_crit_pct	90	a datastore this % used or more is red (critical); amber from vm_datastore_warn_pct
vm_datastore_fail	false	The job stays green when datastores are full: the report and email say which. true = mark the job FAILED when one is at or above vm_datastore_warn_pct (or inaccessible), for workflows.
vm_hostnet_services	[vmotion]	Read-only: vmkping between the hosts of each cluster on these VMkernel services' adapters (full size, don't fragment), the adapters' MTUs, the storage latency. The pings run on the hosts through vCenter (esxcli): the vCenter account needs Host > CIM > CIM interaction on them (docs/VMWARE.md, "The ESXi network check").... (all of it: the role file)
vm_hostnet_targets	[]	pinged from every host, e.g. [{name: NFS array, host: 192.0.2.50, interface: vmk2}]
vm_hostnet_hosts	[]	only these hosts (names, * and ?); [] = every host
vm_hostnet_clusters	[]	only these clusters
vm_hostnet_exclude	[]	never these hosts
vm_hostnet_ping	true	false = the adapters and the storage latency only, no pings
vm_hostnet_ping_count	3	pings per pair
vm_hostnet_max_pairs	400	pings at most per vCenter (16 hosts in one cluster: 240)
vm_hostnet_workers	4	pings at the same time

Setting	Default	What it does
vm_hostnet_rtt_warn_ms	5	a round trip this slow or slower: warning (vMotion within a site is well under 1 ms)
vm_hostnet_rtt_crit_ms	150	... critical (vMotion's own limit)
vm_hostnet_latency	true	read the storage latency (the last hour of the hosts' real-time statistics)
vm_hostnet_latency_minutes	60	minutes of the hosts' real-time statistics read (at most 60)
vm_hostnet_latency_warn_ms	20	a datastore's average read or write latency from this: warning
vm_hostnet_latency_crit_ms	50	... critical
vm_hostnet_raw	false	true = the first ping and the first failed one whole, in the job's output: the host's parameters of the command, what was sent, the request, the answer or the fault (the first run, or a host that refuses a parameter)
vm_hostnet_fail	false	true = the job shows failed when the report has a warning or worse (for workflows)
vm_inventory_templates	false	---- VM inventory report (playbooks/vm_inventory_report.yml) ---- true = templates counted and listed as VMs (false: left out, their number in the subtitle)
vm_inventory_by_folder	true	a table per vCenter and folder; false = one table of every VM (vCenter and folder columns)
vm_inventory_max_rows	2000	VMs listed in the email and the page at most (2000: about 2.5 MB of email; the workbook always has every VM)
vm_inventory_xlsx	true	attach every VM as an Excel workbook (.xlsx) to the email
vm_inventory_security	true	read each VM's encryption, Secure Boot and vTPM: the section "VMs without encryption, EFI or vTPM", its counts, the website's columns and the workbook (false: not read)
vm_inventory_security_skip_templates	true	templates left out of that section and its counts (when vm_inventory_templates lists them)

Setting	Default	What it does
vm_inventory_security_skip_off	false	true = powered-off VMs left out of that section and its counts
vm_inventory_hidden_columns	[SB, Tools ver.]	the website's VM tables: these of the extra columns (Enc, EFI, SB, vTPM, HW, Tools, Tools ver. - or their full names) start hidden; the Columns menu shows them, and keeps a viewer's choice
vm_tools_templates	false	---- VMware Tools inventory (playbooks/vm_tools_report.yml) ---- true = templates counted and listed as VMs (false: left out, their number in the subtitle)
vm_tools_by_folder	true	a table per vCenter and folder; false = one table of every VM (vCenter and folder columns)
vm_tools_max_rows	2000	VMs listed in the email and the page at most (the workbook always has every VM)
vm_tools_xlsx	true	attach every VM, the Tools versions and the hardware versions as an Excel workbook (.xlsx) to the email
vm_ask_question	(empty)	---- VM - ask a question (playbooks/vm_ask.yml, docs/VM_ASK.md) ---- the question, in plain words (the survey's Textarea), e.g. "all VMs with dbv in the name", "VMs on datastores under 10% free"
vm_ask_query	{}	or the query itself, in the job's own form (docs/VM_ASK.md, "The query"): GenAI is not asked
vm_ask_genai	true	GenAI (ACT, evidence mode: it runs no command) turns the question into a query; false = only the job's own reader (simple questions)
vm_ask_local_fallback	true	GenAI did not run (no key, not reachable) or its answer held no query: the job's own reader tries it (simple questions only; the report says so)
vm_ask_genai_retry	true	GenAI's answer held no query (a plan or a step of its own instead, or prose): it is asked once more, told what it answered; false = not
vm_ask_model_timeout	120	seconds GenAI may take to answer

Setting	Default	What it does
vm_ask_event_days	30	a VM's last power off and vMotion (and a restart inside the guest) come from vCenter's events of this many days: older ones are not found (vCenter keeps events for a limited time, often 30 days)
vm_ask_max_events	50000	read at most this many of those events per vCenter (the newest)
vm_ask_extra_instructions	(empty)	more words for GenAI, e.g. your naming: "a name with DBV is a database server; ACAS is the scanner's folder"
vm_ask_max_rows	500	rows listed in the email and on the page at most (the workbook has every row)
vm_ask_print_rows	50	rows printed in the job's output
vm_ask_artifact_rows	200	rows in the job artifact vm_ask (for the next step of a workflow)
vm_ask_xlsx	true	attach every row of the answer as an Excel workbook (.xlsx) to the email
vm_versions_max_rows	2000	---- VMware version report (playbooks/vmware_version_report.yml) ---- ESXi hosts listed in the email and the page at most (the workbook always has every host)
vm_versions_xlsx	true	attach every vCenter and ESXi host as an Excel workbook (.xlsx) to the email
vm_hardware_vibs	true	---- VMware hardware inventory (playbooks/vmware_hardware_report.yml) ---- also read every installed VIB of each host (the workbook's VIBs sheet; one call per host)
vm_hardware_max_rows	2000	ESXi hosts listed in the email and the page at most (the workbook always has every host)
vm_hardware_xlsx	true	attach every host, NIC, storage adapter, software component and VIB as an Excel workbook

Setting	Default	What it does
vm_rightsize_days	30	---- VM right-sizing report (playbooks/vm_rightsizing_report.yml) ---- vCenter's own statistics only (statistics level 1, vCenter's default, is enough). Suggestions only: nothing is changed. days of history (at most 30: what vCenter keeps of 2-hour averages)
vm_rightsize_year	true	also each VM's busiest day of the last year (daily averages): month-end and yearly peaks
vm_rightsize_min_days	7	fewer days of statistics than this = 'Not enough history' (no suggestion)
vm_rightsize_clusters	[]	only the VMs of these clusters (names; * and ? allowed); [] = every VM
vm_rightsize_vms	[]	only these VMs (names; * and ? allowed)
vm_rightsize_exclude	[]	never these VMs (names; * and ? allowed)
vm_rightsize_cpu_target_pct	70	vCPUs suggested so that the 95th percentile of CPU use would be about this
vm_rightsize_cpu_high_pct	85	CPU at this or more 95% of the time = too small
vm_rightsize_ready_warn_pct	5	CPU ready (waiting for a physical CPU, per vCPU) at this or more is said
vm_rightsize_min_vcpus	2	never fewer vCPUs than this
vm_rightsize_min_vcpus_windows	2	... for Windows
vm_rightsize_cpu_max_cut_pct	50	never suggest more than this % fewer vCPUs at once
vm_rightsize_mem_headroom_pct	100	memory suggested = active memory plus this % (active memory understates real need)
vm_rightsize_mem_high_pct	90	active memory at this % or more of what the VM has = too small
vm_rightsize_min_mem_gb	2	never less memory than this (GB)
vm_rightsize_min_mem_gb_windows	4	... for Windows
vm_rightsize_mem_max_cut_pct	50	never suggest more than this % less memory at once

Setting	Default	What it does
vm_rightsize_min_change_pct	25	a smaller memory size is suggested only when it is this % (and 2 GB) smaller or more
vm_rightsize_taken_mb	100	memory ballooned or swapped by the host at this many MB or more (its peak) is said
vm_rightsize_disk_warn_pct	90	a guest disk this % used or more (VMware Tools' numbers) is listed
vm_rightsize_max_rows	500	rows per table in the email and the page (the workbook always has every VM)
vm_rightsize_xlsx	true	attach every VM with every number as an Excel workbook (.xlsx) to the email
vm_rightsize_batch	10	VMs per statistics query (vCenter refuses large ones: vpxd.stats.maxQueryMetrics)
vm_rightsize_workers	4	statistics queries at a time
vm_rightsize_act	false	GenAI (ACT) reviews the suggestions: needs an "ACT model key" credential (docs/VMWARE.md, "ACT analysis")
vm_rightsize_act_max_items	40	suggestions ACT reviews: too small first, then too big (the most freed first)
vm_rightsize_act_batch	"{{ site_act_batch | default(5) }}"	suggestions per question to ACT (0 = all in one)
vm_rightsize_act_model_timeout	"{{ site_act_model_timeout | default(600) }}"	seconds the model may take for its answer
vm_rightsize_act_check_url	true	first check that the AAP node can reach the model's URL
vm_snapshot_max_age_days	5	Snapshots at least this many days old: listed apart by the report, deleted by the cleanup.
vm_snapshot_keep_regex	'(?i)keep|do.?not.?delete'	Never deleted by the cleanup (the report lists them, with the reason). Names given as text (a survey): separated by commas or line breaks, not spaces. - a snapshot whose name or description matches this

Setting	Default	What it does
vm_snapshot_cleanup_exclude_vms	[]	- a snapshot of these VMs: names, * and ? allowed (e.g. [DC*, VCSA01, AAP01]). The VMs in vmware_protected_vms and the AAP server are never cleaned up either. Patch hosts does not patch these VMs. (docs/RUNBOOKS.md, Patch hosts)
vm_snapshot_cleanup_exclude_datacenters	[]	- a snapshot of a VM in these vCenter datacenters: names, * and ? allowed, any case (e.g. [DRSITE])
vm_snapshot_cleanup_exclude_clusters	[]	- a snapshot of a VM running in these clusters: names, * and ? allowed, any case (e.g. [SQL*, MGMT*]). A VM on a host that is in no cluster is not matched.
vm_snapshot_cleanup_name_words	[]	- with words here, a snapshot whose name contains none of them: only a name with one is deleted (any case; Patch also matches Patching and Patches - and dispatch). E.g. [Patching, Patches, Patch]: only the patching snapshots. [] = any name. Set it in all.yml, so the report lists the same.
vm_snapshot_cleanup_max_size_gb	1024	- a snapshot bigger than this many GB (1024 = 1 TB), or whose size cannot be read: merging a very big snapshot can take hours and slow the VM - do those by hand, at a quiet time.
vm_snapshot_cleanup_confirm	false	The cleanup on its own (not after a workflow approval) deletes only when this is true - e.g. a survey question "Delete them?". false = it lists what it would delete, and deletes nothing.
vm_snapshot_cleanup_max	50	Refuse to delete more than this many snapshots in one run (a typo in the age would otherwise delete them all). Raise it on purpose for a big cleanup.
vm_snapshot_delete_timeout	1800	seconds to wait for one deletion (vCenter merges it into the disk)

Setting	Default	What it does
vm_snapshot_cleanup_email_only_if_deleted	false	The cleanup's email (report_email_to): true = only when it deleted something (or a deletion failed).
vm_pick_snapshot_templates	[]	vm_pick_refresh.yml (on a schedule) writes what may be done right now as the choices of one multiple-choice question in the survey of these templates: job templates or workflow job templates, by name (or {template: name or ID, kind: job or workflow} when a job template and a workflow share the name). In a workflow onl... (all of it: the role file)
vm_pick_snapshot_question	vm_snapshot_picks	the question's Answer variable name
vm_pick_reboot_templates	[]	e.g. ["VM - restart picked VMs (workflow)"]
vm_pick_reboot_question	vm_reboot_picks	the question's Answer variable name
vm_pick_max_choices	200	a list never has more choices than this (the rest is counted in the job's output)
vm_pick_snapshot_min_age_days	0	The snapshots offered: only those the snapshot cleanup could delete (the same holds - a name that says keep, the excluded VMs, datacenters and clusters, the protected VMs and the AAP server, bigger than vm_snapshot_cleanup_max_size_gb; not the name words), at least this many days old. 0 = any age.
vm_pick_reboot_vms	[]	The VMs offered for a restart: powered on, never the AAP server or a protected VM (vmware_protected_*), among these (names, * and ?, any case; [] = every VM; survey text: separated by commas or line breaks) ...
vm_pick_reboot_clusters	[]	e.g. [PROD-*]
vm_pick_reboot_folders	[]	folder paths as vCenter shows them, e.g. ["/DC1/vm/Linux*"]
vm_pick_reboot_exclude_vms	[]	... and never these (e.g. [SQL*, DC*])

Setting	Default	What it does
vm_pick_reboot_only_hung	false	true = only the VMs that look hung: VMware Tools installed, but no heartbeat (red) or Tools not running, powered on for vm_pick_hung_min_minutes at least (a VM still starting has neither yet). Restart those with vm_power_mode: hard - a guest restart needs VMware Tools running.
vm_pick_hung_min_minutes	15	minutes powered on before a VM without a heartbeat counts as hung
vm_pick_aap_ca_path	(empty)	a CA file for AAP's own HTTPS certificate ("" = what the execution environment trusts)
vm_snapshot_picks	[]	The survey's answers (the picked choices: a list, or text one per line). the snapshots to delete
vm_reboot_picks	[]	the VMs to restart
vm_pick_confirm	false	vm_snapshot_delete_picked.yml / vm_reboot_picked.yml on their own (not a workflow): true = do it now. false = list exactly what would be done (the plan), and do nothing. Inside a workflow it is never used: the first step plans, the step after the approval does exactly the plan.
vm_pick_max	20	refuse to restart, or delete, more than this many in one run
vm_alarm_types	[vcenter, datacenter, cluster, host]	Whose triggered alarms and configuration issues: vcenter, datacenter, cluster, host. Add vm, datastore or network for theirs too (VM alarms can be many).
vm_alarm_hours	24	The events of the last this-many hours: every error and warning, failed logins, host connection events and VM shutdowns / power-offs / resets / HA restarts. 0 = alarms and host state only.
vm_alarm_max_events	5000	read at most this many events (the newest)
vm_alarm_max_alarm_changes	2000	read at most this many alarm status changes (the newest; 0 = no "Alarms that fired")

Setting	Default	What it does
vm_alarm_max_rows	200	rows per table in the report (the artifacts have them all)
vm_alarm_login_critical	10	This many failed logins of one user from one place, or more, is critical (fewer: a warning).
vm_alarm_login_event_types	[BadUsernameSessionEvent, com.vmware.sso.LoginFailure, esx.audit.account.locked, esx.audi…	The event types read whatever their category (vCenter's event type names). Add to these lists rather than replacing them; [] = none of that kind.
vm_alarm_connection_event_types	[HostConnectionLostEvent, HostNotRespondingEvent, HostDisconnectedEvent, HostReconnection…	Events about a host losing (or failing to get back) its connection to vCenter: the Host connection table.
vm_alarm_vm_event_types	[VmGuestShutdownEvent, VmGuestRebootEvent, VmPoweredOffEvent, VmResettingEvent, VmFailove…	Events about a VM stopping, restarting or failing over (guest shutdown, power off, HA restart): the VM events table.
vm_alarm_fail	false	The job stays green when it finds alarms: the report and email say what. true = mark the job FAILED when there is a critical alarm or a host is not connected, for workflows.
vm_alarm_act_max_items	40	ACT analysis (playbooks/vm_alarm_act_analysis.yml). ACT runs on the AAP side (the execution environment), never on the ESXi hosts; the model provider, model, URL, proxy and CA are the site_act_* settings (roles/site_act/defaults/main.yml), the key an "ACT model key" credential. analyze at most this many problems, the… (all of it: the role file)
vm_alarm_act_batch	"{{ site_act_batch | default(5) }}"	problems per question to ACT: a long answer about many problems is what a gateway or proxy cuts off (0 = all in one question)
vm_alarm_act_config_issues	true	configuration issues too (false = alarms and events only)
vm_alarm_act_email_if_none	false	true = email even when there is nothing to analyze
vm_alarm_act_check_url	true	first check that the AAP node can reach the model's URL

Setting	Default	What it does
vm_alarm_act_model_timeout	"{{ site_act_model_timeout | default(600) }}"	Seconds the model may take for its answer (ACT's GENAI_TIMEOUT, 90 by default, is too short for a long analysis): site_act_model_timeout (600) unless set here. A GENAI_TIMEOUT in site_act_env wins.
vm_health_require_ha	true	A cluster of 2 or more hosts without vSphere HA is a warning (false = HA off is on purpose).
vm_health_require_drs	true	A cluster of 2 or more hosts without DRS is a warning (false where DRS is off on purpose or not licensed, e.g. vSphere Standard).
vm_health_cpu_warn_pct	85	CPU and memory used, in percent of a host (or of a cluster's hosts running VMs): CPU used: a warning from this percent
vm_health_cpu_crit_pct	95	CPU used: critical from this percent
vm_health_mem_warn_pct	90	memory used: a warning from this percent
vm_health_mem_crit_pct	95	memory used: critical from this percent
vm_health_sensors	true	read the hosts' hardware sensors (false = skip them)
vm_health_clock	true	Each connected host's clock against vCenter's, and its NTP setting: a host writes its own logs (hostd, vmkernel) by its clock, vCenter its events by its own - off clocks are why a disconnection's time in a host's log does not match vCenter's. No NTP server set, or the NTP service stopped, is a warning. vCenter's clock... (all of it: the role file)
vm_health_clock_warn_seconds	60	a host's clock this many seconds off vCenter's (ahead or behind) -> warning
vm_health_clock_crit_seconds	300	-> critical (Active Directory logins fail from 5 minutes off)
vm_health_max_rows	500	rows per table in the report (the artifacts have them all)

Setting	Default	What it does
vm_health_fail	false	The job stays green when it finds problems: the report and email say what. true = mark the job FAILED when there is a critical problem (a host not connected, an HA agent isolated ...), for workflows.
vm_net_drops	true	Network drops and errors in the health report: packets dropped and network errors of every ESXi host (by physical NIC) and powered-on VM (by virtual NIC), from vCenter's statistics. The past day needs vCenter's 5-minute statistics at level 2 (the default is 1); at level 1 the last hour is read, and the report says so.... (all of it: the role file)
vm_net_drops_vms	true	the powered-on VMs too (false = the hosts only: quicker where there are thousands of VMs)
vm_net_drops_hours	24	hours to read when vCenter keeps the past day (at most 24: what the 5-minute interval keeps)
vm_net_drops_warn_per_hour	100	dropped packets an hour (received + sent) of one host or VM from this many -> warning
vm_net_errors_warn_per_hour	10	network errors an hour (received + sent) of one host or VM from this many -> warning
vm_watch_window_minutes	15	The HA and storage events of the last this many minutes (set it to the schedule's interval).
vm_watch_min_up_minutes	15	A VM powered on for less than this is not judged (a VM still starting has no heartbeat yet).
vm_watch_confirm_checks	3	A VM is hung only when it has no guest heartbeat AND does not answer on the network, in every check, AND (vm_watch_traffic) vCenter saw it send no network packets: a suspect is checked this many times in one run (1 = once), vm_watch_confirm_seconds apart.
vm_watch_confirm_seconds	120	seconds between those checks

Setting	Default	What it does
vm_watch_probe	true	The network probe, from the AAP node, of the suspects only (a VM with a heartbeat is never probed): ping, and a TCP connect. A connection made or refused counts as an answer (the guest is alive); a timeout - a firewall between AAP and the guest, often - is no answer, and never makes a VM hung by itself: only with no h... (all of it: the role file)
vm_watch_probe_port	auto	TCP port also tried: auto = 22, or 3389 for a Windows guest; a number (e.g. 5985); 0 = ping only
vm_watch_probe_timeout	3	seconds per address
vm_watch_max_suspects	50	probe and confirm at most this many suspects in one run
vm_watch_traffic	true	The third signal, from vCenter itself (no network path needed): a silent suspect that sent network packets since the first check - vCenter's real-time statistics, 20-second samples - is alive, not hung (Tools died, or a firewall stops the probe). A hung guest sends none. false = the heartbeat and the probe only
vm_watch_traffic_packets	1	packets sent from which a VM counts as alive
vm_watch_one_signal_severity	warning	A VM with no heartbeat that answers on the network (or cannot be probed): VMware Tools died, not the guest. warning = in the tickets (servicenow_min_severity warning); info = the job's output only.
vm_watch_skip_vms	[]	VMs never judged: names, * and ? allowed (e.g. [LAB-*])
vm_watch_skip_folders	[]	VM folders never judged, as vCenter shows them (e.g. ["/DC1/vm/Lab*"])
vm_watch_protect	[]	Phase 2 (not built yet) would restart the hung VMs - never these (names, * and ?), never the AAP server or the VMs of vmware_protected_vms / vmware_protected_groups: their hung state is a ticket only. e.g. [DC*, SQL*, VCSA*]

Setting	Default	What it does
vm_watch_max_actions	5	more hung VMs at once than this = something bigger: none would be restarted past it
vm_watch_fail_on	[critical]	The job shows FAILED (a workflow's "On failure" path) when it finds a condition of these severities.
vm_watch_report_on	[critical]	The report is emailed (report_email_to) and published (report_web_dest) only when it finds a condition of these severities; a clean run is one line in the job output.
vm_watch_sso	true	the watch also asks: is single sign-on down now (vm_sso_*, below)? false = not asked
vm_sso	true	---- Single sign-on and OCSP (vCenter 8.0): the health report (playbooks/vm_health_report.yml) and the watch ---- vCenter's SSO services, the STS signing certificate (health report only), the identity providers, a login by a test domain account (the "vCenter SSO test account" credential), and the OCSP responders smart... (all of it: the role file)
vm_sso_services	[sts, lookupsvc, trustmanagement, vmdird, vmafdd]	SSO services that must be started (as vCenter's and the appliance's service lists name them; one not listed is skipped)
vm_sso_sts_warn_days	60	the STS signing certificate expires within this many days -> warning
vm_sso_sts_crit_days	14	... within this many -> critical (expired: critical)
vm_sso_test_user	(empty)	the test domain account; empty = the "vCenter SSO test account" credential's (VMWARE_SSO_TEST_USER); none = no login test

Setting	Default	What it does
vm_sso_ocsp	[]	The OCSP responders. Simplest: site_ocsp_certs (below) - a certificate the responder knows, its OCSP address and its issuer read from the certificate itself. vm_sso_ocsp, when set, is used instead: [{url: http://ocsp.example.mil}] asks only that it answers; with cert (a PEM file from playbooks/, e.g. files/ocsp/test-u... (all of it: the role file)
site_ocsp_certs	[]	certificates the OCSP responders know (PEM files from playbooks/), for vm_sso_ocsp and net_perf_ocsp when those are empty, e.g. [files/ocsp/test-user.pem]
vm_sso_ocsp_severity	critical	an OCSP responder that fails: critical (smart card logins fail) or warning
vm_sso_crl	true	The CRL of each certificate above (its CRL Distribution Point, or crl_url in a vm_sso_ocsp entry; crl: false skips one), in the health report only - from the AAP node: downloaded, issued and signed by the CA, not past its next update, not listing the test certificate. vCenter falls back to the CRL when OCSP does not a... (all of it: the role file)
vm_sso_crl_severity	warning	a CRL that fails: warning (OCSP still answers smart card logins) or critical
vm_sso_crl_timeout	120	seconds for one CRL download (DoD CRLs are tens of MB)
vm_sso_crl_max_mb	200	a CRL larger than this is not read
vm_sso_timeout	30	seconds per request
vm_created_days	7	The window: the VMs created in the last this many days, up to the moment the job runs (7 = the past week; run weekly, each VM is in exactly one report). VMs created and deleted since are listed too, as long as vCenter keeps their events (often 30 days).
esxi_security_hosts	[]	Which hosts: names (* and ? allowed; the short name matches too). [] = every host of vCenter (or of vmware_datacenter).

Setting	Default	What it does
esxi_security_clusters	[]	only the hosts of these clusters (names, wildcards); [] = all
esxi_security_exclude_hosts	[]	Never these hosts (names, wildcards). E.g. a host you opened SSH on for a support case: the nightly run would otherwise close it again.
esxi_security_ssh	disabled	The SSH service (TSM-SSH): disabled = stopped, and set to start manually; enabled = started, and set to start with the host; "" = leave it as it is.
esxi_security_shell	disabled	The ESXi Shell service (TSM), the host's local shell: the same choices.
esxi_security_settings	UserVars.ESXiShellTimeOut: 600 UserVars.ESXiShellInteractiveTimeOut: 600	Advanced settings, name: value; {} = none. The defaults (seconds): UserVars.ESXiShellTimeOut SSH and the ESXi Shell stop by themselves this long after someone starts them (0 = never) UserVars.ESXiShellInteractiveTimeOut an idle SSH or ESXi Shell session is logged out Others often set: UserVars.DcuiTimeOut (an idle DCU... (all of it: the role file)
esxi_security_lockdown	normal	Lockdown mode: normal = the host is managed only through vCenter (the DCUI still works); disabled; "" = leave it as it is. strict is refused, and a host already strict is left alone.
esxi_security_lockdown_exception_users	[]	Accounts that must keep logging in to the hosts directly under lockdown (a vulnerability scanner, a monitoring tool): host-local accounts, or Active Directory accounts with their own permission on the host (not vsphere.local, not an AD group). Added before lockdown is turned on; never removed.
esxi_security_email_only_if_changed	true	The email (report_email_to): true = only when a host was changed or failed; false = every run.

Setting	Default	What it does
esxi_security_vcsa_ssh	(empty)	The vCenter appliance's own access - its VAMI's Access page (https://vcenter:5480 > Access) - on every vCenter of the job: enabled, disabled, or "" = leave it (the default: nothing changes until you set one). Set through vCenter's API with the same "VMware vCenter" credential: its account must be in vCenter's single s... (all of it: the role file)
esxi_security_vcsa_dcui	(empty)	DCUI - the appliance's console screen.
esxi_security_vcsa_consolecli	(empty)	Console CLI - the appliance shell on its console.
esxi_security_vcsa_shell	(empty)	Bash Shell.
esxi_security_vcsa_shell_timeout	3600	Seconds Bash Shell stays on when this job turns it on (1 to 86400).
vmware_cert_warn_days	30	A certificate of vCenter (its Machine SSL certificate, the one its web pages show) or of an ESXi host that expires within this many days (or has expired) is due: a ServiceNow incident, and a CSR made where its key stays - vCenter, or the host - attached to the incident with what to do next. To try the job on certifica... (all of it: the role file)
vmware_cert_crit_days	7	within this many days (or expired) the incident is critical, else warning
vmware_cert_esxi	true	the ESXi hosts' certificates too; false = vCenter's only
vmware_cert_only	[]	Only these vCenters and hosts: names, any case, * and ? allowed, the short name matches too - e.g. one host to try the job on. [] = every vCenter of the job and every host connected to it.
vmware_cert_make_csr	true	false = the incidents only, no CSR. A CSR is made once per certificate: the incident is its record, and a second CSR would replace the key that waits on vCenter or the host for the issued certificate.

Setting	Default	What it does
vmware_cert_vmca_csr	false	A certificate signed by vCenter's own CA (VMCA - how vCenter gives ESXi hosts their certificates unless told otherwise) is renewed in vCenter (its Renew button): its incident says so, and no CSR is made. true = a CSR for those too, to replace them with certificates of your PKI.
vmware_cert_vmca_issuers	[]	Issuers taken as VMCA besides VMCA's own root (CN=CA with DC=..., or OU=VMware Engineering): any case, * and ? allowed, e.g. ['*O=vc01.yoursite.mil*'].
vmware_cert_csr_fields	{}	vCenter's CSR: its subject and names are the current certificate's. vCenter's API wants an organization, organization unit, locality, state and country - one the certificate does not have goes empty. Set any of them here (the PKI sets the subject by its own rules anyway), e.g. vmware_cert_csr_fields: {organization: U.... (all of it: the role file)
vmware_cert_csr_without_servicenow	false	Without the ServiceNow API credential no CSR is made (nothing would record it: the next run would make another one). true = make them anyway, and print them in the job output.
vmware_cert_install_trust_check	true	The certificate the PKI put on the catalog request of a vCenter's or a host's CSR, installed where its key waits. A host's certificate: vCenter must trust its issuer, or it refuses the host after the replacement - the plan adds the CA chain (from the request) to vCenter's trusted roots when it is not there. false = do... (all of it: the role file)
vmware_cert_install_ritm	(empty)	The request to install from: its number (RITM0012345) - the check reads that request alone, for the vCenter or host its CSR was made for (from the CSR's name), and hands it to the approval. The workflow's survey asks it.

Setting	Default	What it does
vmware_cert_install_search	false	true = without a request named, look up the request of every certificate of vCenter and its hosts (each by its CSR's name and expiry) - for a schedule on the workflow, once you trust it. false = a request must be named.
vmware_cert_request_for	[]	VMware - certificate request (playbooks/vmware_cert_request.yml): the vCenters and ESXi hosts to order a certificate for now - exact names, no * or ? (each one gets a CSR and a request); required. e.g. [esx01.yoursite.mil]
vmware_cert_install_host_wait	300	Seconds a host may take, after the install, until vCenter shows its new certificate (reconnected if vCenter lost it).
vmware_cert_install_vcenter_wait	1800	Seconds vCenter may take to serve its new certificate on its port: before vCenter 8.0 Update 2 its services restart.
vmware_trust_esxi	true	---- VMware trust store report (playbooks/vmware_trust_report.yml) ---- vCenter's trusted roots and each ESXi host's trusted CAs; false = vCenter only (the ESXi hosts left out)
vm_capacity_clusters	[]	---- vSphere capacity planning (playbooks/vm_capacity_report.yml) ---- only these clusters (names; * and ? allowed); [] = every cluster
vm_capacity_history_days	120	The history: vCenter's own daily statistics (cluster CPU and memory, datastore space), kept a year at statistics level 1. Read this many days; the growth is a straight line over the last vm_capacity_trend_days, with at least vm_capacity_min_samples daily values.
vm_capacity_trend_days	90	The growth (straight line) is fitted on this many most recent days. Growth comes in jumps (VMs are deployed in batches): over 30 days one build-out becomes the rate and the runway reads days when it is months; 90 days averages the batches out. 30 reacts faster to a real change of rate.

Setting	Default	What it does
vm_capacity_min_samples	7	fewer daily values than this in those days = 'not enough history' (no trend)
vm_capacity_events_days	30	VMs added and removed (vCenter keeps events 30 days by default)
vm_capacity_failover_hosts	1	Plan for this many failed hosts per cluster (N+1); a larger HA admission control failover level wins.
vm_capacity_mem_target_pct	80	memory used at most this %, with those hosts failed
vm_capacity_cpu_target_pct	80	CPU used at most this %, with those hosts failed
vm_capacity_vcpu_per_core_max	4	vCPUs per physical core (amber from here, red from 1.5 times it)
vm_capacity_horizon_months	12	Hosts recommended per cluster (of its average host size) to keep memory and CPU at their targets with those hosts failed: now, and at the current growth within this many months.
vm_capacity_storage_warn_pct	85	datastore used %: amber
vm_capacity_storage_crit_pct	90	red, and the target of the storage runway
vm_capacity_runway_warn_days	90	a runway shorter than this is amber
vm_capacity_runway_crit_days	30	and red
vm_capacity_include_local_datastores	false	a host's own (not shared) datastores count too
vm_capacity_max_rows	100	rows per table (problems first); the artifacts have all
vm_capacity_fail	false	The job stays green: the report says what is short. true = mark it FAILED when a cluster or datastore is critical (for workflows).
vm_capacity_act	false	GenAI: ACT reads the numbers and makes its own estimate, shown next to the job's (math vs GenAI). Needs an "ACT model key" credential on the template; the same site_act_* settings as the alarms analysis. ACT runs on the AAP side and runs no command.

Setting	Default	What it does
vm_capacity_act_max_items	30	items ACT estimates: everything overall, then clusters and datastores, problems first
vm_capacity_act_batch	"{{ site_act_batch | default(5) }}"	items per question to ACT (each question has the whole evidence; 0 = all in one)
vm_capacity_act_model_timeout	"{{ site_act_model_timeout | default(600) }}"	seconds the model may take for its answer
vm_capacity_act_check_url	true	first check that the AAP node can reach the model's URL
vm_deploy_ticket	(empty)	---- Deploy a VM from a ServiceNow ticket (playbooks/vm_deploy_from_ticket.yml) ---- The ticket (survey answer): INC..., TASK..., SCTASK..., RITM..., CHG... - or your instance's own prefixes (INCN..., CHGN...): the job asks ServiceNow which table a ticket is in. The ticket says what to build in 'key: value' lines - in... (all of it: the role file)
vm_deploy_ticket_table	(empty)	leave it empty (the job finds each ticket's table); a table of your own, e.g. u_vm_request: for numbers whose prefix the job does not know (CTASK stays change_task)
vm_deploy_ticket_prefixes	{}	Your instance's own number prefixes and their tables, e.g. {INCN: incident, CHGN: change_request}. Rarely needed: for a prefix it does not know, the job asks ServiceNow's task table (and takes INCN as INC, CHGN as CHG when it may not read that table). A table as people say it is understood: change / change request ->... (all of it: the role file)
vm_deploy_parent_tables	[change_request]	The request may also be on the ticket's parent - e.g. a task (SCTASK, CTASK) whose change holds the requirements: the parent's 'key: value' lines (and, with vm_deploy_plain_words, its plain words) are read too, under the ticket's own (what the ticket says wins). A parent that cannot be read, is closed, or is not ready... (all of it: the role file)

Setting	Default	What it does
vm_deploy_parent_fields	[change_request, parent]	The ticket's fields that point to its parent - the first one set is used: a change task's change_request, any ticket's parent (or your instance's own field, e.g. u_change):
vm_deploy_parent_text_fields	[description]	The parent's fields with the request, in this order (a later one wins), e.g. [description, implementation_plan]; comments and work_notes may be listed too (oldest first). GenAI reads them as well.
vm_deploy_parent_wins	false	What wins, from the least to the most: the parent (the change), the ticket's short description and description, its work notes, its comments (each oldest first) - a 'key: value' line and plain words alike: the later one wins, however each is written. So the change task wins over its change, and a comment wins over a w... (all of it: the role file)
vm_deploy_notes_also	[]	Every note and comment this job writes on the ticket goes on these too (closing never does: only the ticket itself is closed): parent = the change a task belongs to (when it was read); tasks = when the job is launched with a change, its open change tasks assigned to vm_deploy_assignment_group (and matching vm_deploy_n... (all of it: the role file)
vm_deploy_notes_tasks_query	(empty)	An extra ServiceNow query for the change's tasks that get the notes, e.g. "short_descriptionLIKEto deploy" ("" = the group's)
vm_deploy_notes_tasks_limit	50	the change's tasks that get the notes (and are closed), at most - more is a warning
vm_deploy_close_tasks	false	Launched with a change: also close its tasks found that way (with vm_deploy_done_fields_by_table's change_task entry, when every VM is built) - the change itself is not closed by this. They get the notes too.

Setting	Default	What it does
vm_deploy_ready_by_table	change_request: {approval: Approved, state: Implement}	Before anything is built, a ticket of these tables - the ticket itself, or the parent it belongs to (a change task's change) - must have these fields: one value or a list (any of them), written as the ticket's form shows them (upper / lower case ignored). Checked by the dry run, the propose step AND right before the b... (all of it: the role file)
vm_deploy_templates	[]	What a ticket may ask for (names; * and ? allowed). Templates MUST be listed: nothing else is deployed. The others: [] = any. e.g. [rhel9-gold, win2022-gold]
vm_deploy_clusters	[]	clusters (and standalone hosts) a ticket may use; also the candidates of automatic placement
vm_deploy_datastores	[]	datastores (or datastore clusters) a ticket may use; also the candidates of automatic placement
vm_deploy_folders	[]	VM folders a ticket may use (paths below the datacenter's VM folder, e.g. Linux/*)
vm_deploy_template_folders	[]	The folders the templates are taken from (paths below the datacenter's VM folder; * and ?), e.g. [Staging]: a template named in a ticket is used only from there - a copy of the same name elsewhere is not (only elsewhere: refused, nothing built). With vm_deploy_folders: [Staging] and no folder in vm_deploy_defaults, th... (all of it: the role file)
vm_deploy_max_cpu	16	the most vCPUs a ticket may ask for
vm_deploy_max_memory_gb	128	the most memory a ticket may ask for, in GB
vm_deploy_max_vms	10	A ticket may ask for several VMs ("2 VMs named web01 and web02", "vm_name: web01, web02", or "Hostnames:" and one name per line): one VM each, the same size - or each its own, from a line of its own - built one after the other, the ticket closed when all are built. More names than this: refused, nothing built (the lim... (all of it: the role file)

Setting	Default	What it does
vm_deploy_name_regex	'^[A-Za-z0-9][A-Za-z0-9-]{0,62}$'	The VM names allowed (a regular expression); e.g. '^[a-z0-9-]{1,15}$' for Windows-safe names.
vm_deploy_assignment_group	(empty)	Only tickets assigned to this group (its name as ServiceNow shows it), or to one of a list of groups (["Group A", "Group B"]); "" = any group (the pickup needs one).
vm_deploy_assignment_group_by_table	{}	The same for the ticket the job is launched with, by its table - in place of vm_deploy_assignment_group for that table; "" = any group. e.g. {change_request: ""}: launched with a change, the change may be assigned to any group (it must still be approved: vm_deploy_ready_by_table); its tasks that get the notes (vm_depl... (all of it: the role file)
vm_deploy_defaults	{}	What a ticket does not say: e.g. {cluster: PROD-CL01, datastore: DS-CLUSTER-01, folder: Linux/New}. Not set = the template's own cluster, datastore and folder; its CPU and memory.
vm_deploy_sites	{}	Sites: a ticket's 'site: A' line (or its Location field, or with plain words a site's name or its words in vm_deploy_plain_sites - each VM by its own line) picks the placement - over vm_deploy_defaults and the name rules; a cluster / datastore / folder the ticket names itself still wins. A site in another vCenter: its... (all of it: the role file)
vm_deploy_site_required	true	with sites set up: a ticket must say its site (false = vm_deploy_defaults); a name rule counts too
vm_deploy_name_rules	[]	Name rules: the VM's name picks where it goes - the first rule whose 'match' fits the name (* and ?, upper / lower case alike) fills in what the ticket does not say: a cluster (and datastores, folder ...) or a site. The ticket's own lines, and a site the ticket names, still win. A name used at several sites: 'sites' -... (all of it: the role file)

Setting	Default	What it does
vm_deploy_pxe_templates	[]	Templates with no operating system yet, that install it themselves: they boot from the network (PXE, e.g. an MECM task sequence) or from their CD / ISO (e.g. a kickstart). No guest customization, no waiting for an IP address; the MAC address and the computer name go on the ticket. e.g. [win2022-pxe]
vm_deploy_guestinfo	true	Give a VM from one of those empty templates the ticket's name and number as VM settings its installer can read: guestinfo.hostname and guestinfo.ticket (a kickstart with VMware Tools in the installer: vmtoolsd --cmd "info-get guestinfo.hostname"). Needs the vCenter privilege Change configuration > Advanced configurati... (all of it: the role file)
vm_deploy_template_aliases	{}	Short names for long template names - for the ticket's 'template:' line and GenAI's catalog (whose keys may be these short names): {short name: the template's name in vCenter}, e.g. {win2022: "Windows 2022 Standard Template - Site 1", rhel9: "RHEL 9 Server Template"} A ticket may also write a template's own name in ot... (all of it: the role file)
vm_deploy_template_specs	{}	The customization specification (vCenter > Policies and Profiles) for each template, when the ticket, its site and the name rules name none: {template name or pattern: spec}. A Windows template that is not network-boot needs one (Sysprep: administrator password, domain or workgroup), e.g. {"Windows 2022*": Win2022-Sys... (all of it: the role file)
vm_deploy_max_disks	5	disks a ticket may ask for, disk 1 included (disks_gb: 100, 200 = disk 1 grown to 100 GB, one 200 GB disk added)
vm_deploy_max_disk_gb	2048	the biggest disk a ticket may ask for, in GB
vm_deploy_max_total_disk_gb	4096	all its disks together, in GB

Setting	Default	What it does
vm_deploy_mem_target_pct	80	Automatic placement - when the ticket, its site and vm_deploy_defaults name no cluster / datastore: the job chooses among a site's 'clusters' / 'datastores' lists (e.g. SiteA: {clusters: [A-PROD-*]}), else among vm_deploy_clusters / vm_deploy_datastores; with neither, the template's own. a cluster fits when its memory... (all of it: the role file)
vm_deploy_failover_hosts	1	hosts to plan as failed when choosing a cluster
vm_deploy_datastore_max_used_pct	85	a datastore fits when it stays under this % with the VM at its full size
vm_deploy_datastore_rules	[]	Datastores by the VM's name: the first rule whose match fits the name (* and ?, any case) gives the datastores it goes on - automatic placement chooses among them on its cluster, the least used first, over a site's or vm_deploy_defaults' datastore (one the ticket names must be one of them). A VM no rule matches never... (all of it: the role file)
vm_deploy_datastore_balance	false	true = a VM whose ticket, site, name rule and vm_deploy_defaults name no datastore (and no vm_deploy_datastores) goes on the least used shared datastore of its cluster - the location's - the VMs of one ticket spread out (each counts the ones before it); false = the template's own datastore.
vm_deploy_resume	true	Run again for the same ticket (a job that stopped: failed, cancelled, timed out), pick up where it stopped: each VM carries its ticket in its notes from the clone on - one it built is left as it is, one it started and did not finish (cloned, not yet powered on) is finished where it is; the work notes say which. false... (all of it: the role file)
vm_deploy_domain	localdomain	the DNS domain Linux customization sets (a ticket's domain: wins)
vm_deploy_customize	true	guest customization: host name = vm_name, every adapter on DHCP

Setting	Default	What it does
vm_deploy_power_on	true	power the new VM on (false = leave it off, e.g. to check it first)
vm_deploy_wait_customization	900	seconds to wait for vCenter's customization result
vm_deploy_wait_ip	600	seconds to wait for an IPv4 address (DHCP, through VMware Tools)
vm_deploy_done_fields	{}	Fields set on the ticket when the VM is up, e.g. to resolve an incident: {state: "6", close_code: "Solution provided", close_notes: "VM deployed"} (your instance's values)
vm_deploy_done_fields_by_table	{}	The same for one table (the ticket's table: incident, sc_task, change_task ...), in place of vm_deploy_done_fields for it - e.g. close the task, never the change it belongs to: {sc_task: {state: "3", close_notes: "VM deployed by AAP"}} (3 = Closed Complete; your instance's values) Sent on its own, after the result not... (all of it: the role file)
vm_deploy_done_fields_issues_by_table	{}	Over those fields when every VM is built but something did not finish: no IPv4 address in time, no customization result, the guest's host name not applied, no network at power on, guestinfo not set - e.g. {change_task: {close_code: successful_issues}} (your instance's values; {} = closed the same way). What happened i... (all of it: the role file)
vm_deploy_done_fields_pxe	true	Also for a VM from a network-boot template (vm_deploy_pxe_templates), which is powered on and still installing its OS from MECM when the job ends. false = that ticket stays open for whoever checks the install.

Setting	Default	What it does
vm_deploy_service	(empty)	---- The server's service, picked at launch: words in the VM's notes, an AAP inventory and group ---- What the new server is for, picked by whoever launches the deploy: a survey question (multiple choice, NOT required - the scheduled pickup starts the workflow without it), e.g. sql or dc. Every VM of the ticket gets i... (all of it: the role file)
vm_deploy_services	{}	The services the survey offers: each one's AAP inventory and group, and the words for the VM's notes (note: left out = "Service: <its name>."). {} = any service is only written in the notes. vm_deploy_services: sql: {inventory: Windows servers, group: win_sql, note: SQL Server} dc: {inventory: Windows servers, group:... (all of it: the role file)
vm_deploy_inventory_name	"{name}"	The name a VM gets in the inventory: write it the way your inventory already names servers, or the same server shows up twice in AAP. {name} = the VM's name, e.g. "{name}.yoursite.mil". A service's own name: wins (e.g. {inventory: Windows servers, name: "{name}.corp.yoursite.mil"}).
vm_deploy_inventory_lowercase	true	the inventory name in lower case (APP01 -> app01); false = as the ticket writes it
vm_deploy_aap_ca_path	(empty)	a CA file for AAP's own HTTPS certificate ("" = vm_deploy_poll_aap_ca_path, else what the execution environment trusts)
vm_deploy_plain_words	false	---- A request in plain words, without GenAI (off by default) ---- true = also read the ticket's own words, by the rules below, so a requester can write "a Red Hat 9 server named ab01web05, medium" instead of 'key: value' lines. Never a guess: two different answers in the text are a question to the requester (a commen... (all of it: the role file)

Setting	Default	What it does
vm_deploy_plain_templates	{}	The words for each template (any case; "windows 2022" also matches "Windows Server 2022", "server 2025" also "a new 2025 VM"). An allowed template's own name and its short names (vm_deploy_template_aliases) count too. An OS no template is for (Server 2025, only a 2022 template) is a question - never another version; t... (all of it: the role file)
vm_deploy_plain_sizes	{}	Size words and what they mean, e.g. {small: {cpu: 2, memory_gb: 8}, medium: {cpu: 4, memory_gb: 16}, large: {cpu: 8, memory_gb: 32}}. "4 vCPU" and "16 GB RAM" in the ticket are read too, and win in the same text.
vm_deploy_plain_sites	{}	The words for each site (vm_deploy_sites), as people write the place - any case, the spaces optional ("Site Two" is also "SiteTwo"). A site's own name counts too. A change that lists its servers one per line ('new VM / ab01web05 / ... Location: Site Two') puts each VM at the site its own line names; the CPU, memory an... (all of it: the role file)
vm_deploy_genai	false	---- GenAI reads a request written in plain words (a proof of concept, off by default) ---- When the ticket's own 'key: value' lines do not give vm_name and template: ACT (GenAI, evidence mode, no commands) copies out what the ticket says - the names, the operating system as written, CPU, memory, disks - each with the... (all of it: the role file)
vm_deploy_genai_in_check	false	A dry run (Check) asks GenAI too (true): it shows GenAI's proposal and where the VM would go, and writes nothing on the ticket - no proposal note, no questions to the requester. false = Check only says GenAI would be asked.

Setting	Default	What it does
vm_deploy_template_catalog	{}	Optional: templates GenAI may name by key, each with a plain description (the key: the template's exact name). Not needed when the templates' names say their OS: the job finds the template from the OS the ticket names. A key for another version than the ticket asks for is dropped. e.g. {win2022-pxe: "Windows Server 20... (all of it: the role file)
vm_deploy_genai_approval	true	true: a GenAI proposal waits for a person - the workflow's approval (or a second run) - before it is built. false: one run reads, checks and builds, when every value is traced to the ticket and within the rules (questions or a refusal still stop it).
vm_deploy_phase	(empty)	Which step this run is: "" = one job does it all; propose = the workflow's first step (read, ask GenAI if needed, check, write where it would go on the ticket - build nothing); build = the step after the approval (build what the ticket and its proposal say - never ask GenAI).
vm_deploy_genai_model_timeout	"{{ site_act_model_timeout | default(600) }}"	seconds the model may take for its answer
vm_deploy_poll_enabled	true	---- Pick up VM request tickets on a schedule (playbooks/vm_deploy_poll.yml) ---- Every run (an AAP schedule, e.g. every 15 minutes) asks ServiceNow for NEW VM requests - active, assigned to vm_deploy_assignment_group (required), marked by vm_deploy_poll_prefix and/or vm_deploy_poll_query - and starts the deploy workf... (all of it: the role file)
vm_deploy_poll_tables	[incident]	the ServiceNow tables to look in: incident, sc_task, task ...
vm_deploy_poll_prefix	"VM request"	a VM request's short description starts with this ("" = mark them by vm_deploy_poll_query)

Setting	Default	What it does
vm_deploy_poll_query	(empty)	An extra ServiceNow encoded query, ANDed with the rest, e.g. "category=vm_request" or "cmdb_ci.name=VM Request" ("" = none). ^NQ (a second query) is refused.
vm_deploy_poll_states	["1"]	the ticket states to pick up (state values: "1" = New)
vm_deploy_poll_max	3	the most tickets started per run; the rest wait for the next run
vm_deploy_poll_scan	100	the most tickets read per table and run, oldest first (tickets AAP already has count too)
vm_deploy_poll_workflow	"VM - deploy from ServiceNow ticket (workflow)"	the workflow job template started per ticket
vm_deploy_poll_dry_run	false	true = only list what it would pick up: no note, no workflow (Check does the same)
vm_deploy_poll_claim_fields	{}	Fields set on the ticket together with the "Picked up" note, e.g. {state: "2"} (In Progress) - so the ticket leaves the states picked up as well ({} = only the note).
vm_deploy_poll_aap_ca_path	(empty)	a CA file for AAP's own HTTPS certificate ("" = what the execution environment trusts)

win_act

File: roles/win_act/defaults/main.yml

Setting	Default	What it does
win_act_dir	'C:\ProgramData\act'	ACT FOR WINDOWS on the servers (vendor/act-windows/act.ps1, the PowerShell port of ACT). Installed into win_act_dir. That folder gets its own owner and permissions - Administrators and SYSTEM only, nothing inherited - because ordinary users may create folders and files in C:\ProgramData (and C:\temp): a non-admin coul... (all of it: the role file)
win_act_state	present	present | absent (removes win_act_dir)

Setting	Default	What it does
win_act_source	"{{ role_path }}/../../vendor/act-windows/act.ps1"	The ACT script this repository ships (vendor/act-windows, version in its VERSION file)

win_check_accounts

File: roles/win_check_accounts/defaults/main.yml

Note
WINDOWS ACCOUNTS (local accounts; domain controllers have none and are skipped) - the built-in Guest account is disabled -> critical if not - enabled local accounts: password never expires, password too old, no logon for a long time -> warning - members of the local Administrators group that you did not allow -> warning Names below allow wildcards (* and ?), e.g. 'YOURDOMAIN\Domain Admins', '*\svc_aap'.

Setting	Default	What it does
win_check_accounts_password_max_age_days	60	an enabled local account's password older than this -> warning (0 = off)
win_check_accounts_inactive_days	35	an enabled local account with no logon for this many days -> warning (0 = off)
win_check_accounts_password_never_expires_allowed	[]	accounts allowed "password never expires" (wildcards)
win_check_accounts_inactive_allowed	[]	accounts allowed to stay unused (wildcards)
win_check_accounts_admins_allowed	[]	Empty = list the members in the job output only. Set it to report anyone else, e.g. win_check_accounts_admins_allowed: ['*\Administrator', 'YOURDOMAIN\Domain Admins', 'YOURDOMAIN\svc_aap']

win_check_audit

File: roles/win_check_audit/defaults/main.yml

Note
WINDOWS AUDIT: the audit policy records what the STIG asks for, the event logs are big enough, and the log forwarding agent runs. One finding lists every missing audit setting.

Setting	Default	What it does
win_check_audit_subcategories	- {name: Credential Validation, guid: 0CCE923F-69AE-11D9-BED3-505054503030, need: Success…	Advanced audit policy subcategories (checked by GUID, so any Windows language works). need: Success | Failure | Success and Failure

Setting	Default	What it does
win_check_audit_security_log_min_kb	196608	STIG
win_check_audit_application_log_min_kb	32768	the Application log's maximum size must be at least this (KB)
win_check_audit_system_log_min_kb	32768	the System log's maximum size must be at least this (KB)
win_check_audit_forwarder_services	[]	Log forwarding agents that must run, e.g. [SplunkForwarder] or [nxlog]

win_check_certs

File: roles/win_check_certs/defaults/main.yml

Note
WINDOWS CERTIFICATES: every certificate in the machine stores below, and which of them are in use (https bindings of IIS / WSUS / ..., the WinRM HTTPS listener AAP connects through, Remote Desktop). Expiring soon -> warning / critical. Expired: critical when in use, or when it expired less than win_check_certs_expired_report_days ago (older leftovers are only listed in the certificate report).

Setting	Default	What it does
win_check_certs_warn_days	30	a certificate expiring within this many days -> warning (amber in the certificate report)
win_check_certs_notice_days	60	the certificate report colours these blue (an early heads-up; not a finding)
win_check_certs_crit_days	7	within this many days -> critical
win_check_certs_expired_report_days	30	an expired certificate NOT in use is a finding for this many days after it expired; older ones only in the report
win_check_certs_stores	['Cert:\LocalMachine\My', 'Cert:\LocalMachine\WebHosting', 'Cert:\LocalMachine\Remote Des…	the machine certificate stores read
win_check_certs_trust	false	also every other machine store: Trusted Root, Intermediate CAs, Other People, ... (the certificate report turns it on); with its private key checked as above, without it only listed (trust store)

Setting	Default	What it does
win_check_certs_trust_skip_stores	['Disallowed', 'AuthRoot', 'REQUEST', 'SMS']	stores never read: the untrusted certificates, the roots Windows Update keeps, pending requests (a CSR's key), the MECM client's own
win_check_certs_remote_desktop	false	false = Remote Desktop's certificates are left out of the report and the renewal (Windows or the GPO renews them): the Remote Desktop store, the one RDP uses, those for Remote Desktop only; true = checked like the others
win_check_certs_ignore_subjects	[]	wildcards, e.g. ['CN=localhost*']
win_check_certs_endpoints	[]	TLS services whose certificate a client actually gets, e.g. win_check_certs_endpoints: - {name: IIS, host: 127.0.0.1, port: 443, servername: web01.yoursite.mil}
win_cert_renew_dir	'C:\ProgramData\site-automation\renew'	---- Renewal: a new key and CSR for each certificate that expires soon (playbooks/win_cert_renew_csr.yml) ---- The certificates the Windows certificate report flags (within win_check_certs_warn_days, or expired) that are the server's own - not a trust store's - get a new key and a CSR with the same subject, names and... (all of it: the role file)
win_cert_renew_key_bits	2048	RSA key size, at least 2048 (a bigger old key keeps its size; an EC certificate gets a new key on the same curve)
win_cert_renew_digest	sha256	the CSR's signature: sha256, sha384 or sha512
win_cert_renew_max	10	at most this many NEW CSRs per server in one run (the rest come the next run)
win_cert_renew_self_signed	false	true = a self-signed certificate gets a CSR too, for the PKI to issue a real one in its place (never one only for localhost)
win_cert_renew_issuer	(empty)	regular expression: only certificates whose issuer matches (empty = any CA), e.g. 'O=U\.S\. Government'
win_cert_renew_exclude_issuer	(empty)	regular expression: certificates whose issuer matches are never renewed here (renewed by their own tool)

Setting	Default	What it does
win_cert_renew_only	(empty)	regular expression on the certificate's place, e.g. 'used by https' (empty = every one)
win_cert_renew_exportable	false	true = the new key can be exported (to move it to another server); false keeps it on this one
win_cert_install_rebind	[https, rdp, winrm]	---- Install: the issued certificate (playbooks/win_cert_install_check.yml -> approval -> playbooks/win_cert_install_approved.yml) ---- The PKI's certificate, from the request carrying the CSR (cert_renew_servicenow: catalog), is put together with its pending key by certreq -accept -machine (in Personal); its chain's... (all of it: the role file)
win_cert_install_verify_wait	15	seconds an https port has, after the rebinding, to serve the new certificate (full PowerShell only)
win_cert_restore_backup	(empty)	---- Restore: the bindings back on the old certificate (playbooks/win_cert_restore.yml) ---- a backup's folder name, as the restore's Check lists it
win_cert_restore_only	(empty)	or: the newest backup of each certificate whose thumbprint or subject matches (a regular expression)

win_check_disk

File: roles/win_check_disk/defaults/main.yml

Note
WINDOWS DISK: how full each local fixed volume (C:, D:, ...) is.

Setting	Default	What it does
win_check_disk_warn_pct	85	% used -> warning
win_check_disk_crit_pct	95	% used -> critical
win_check_disk_min_free_gb	0	also warn below this many GiB free (0 = off): big volumes
win_check_disk_overrides	{}	Different limits for particular volumes, e.g. a database volume kept full on purpose: win_check_disk_overrides: 'D:': {warn: 92, crit: 98}
win_check_disk_ignore	[]	volumes not to check, e.g. ['E:']

win_check_eventlog

File: `roles/win_check_eventlog/defaults/main.yml`

> **Note**
> WINDOWS EVENT LOG, the last win_check_eventlog_hours: - unexpected shutdown or power loss (41, 6008), a crash / blue screen (1001) -> critical - disk and file system errors (disk, Ntfs, storage drivers) -> critical - the Security log was cleared (1102) -> critical - many errors in System + Application (the finding names the top sources) -> warning - a System / Application / Security log that is full and not overwriting -> critical

Setting	Default	What it does
`win_check_eventlog_hours`	`24`	how many hours back the event logs are read
`win_check_eventlog_error_warn`	`50`	this many errors or more in System + Application -> warning

win_check_network

File: `roles/win_check_network/defaults/main.yml`

> **Note**
> WINDOWS NETWORK, from the server itself: - a default gateway and DNS servers are configured; - the names you list resolve, the TCP services you list answer; - domain members: the computer's trust with the domain works (nltest /sc_query). A broken trust is the classic "the trust relationship between this workstation and the primary domain failed" at logon.

Setting	Default	What it does
`win_check_network_dns_names`	`[]`	Names every Windows server must resolve, e.g. [dc01.yoursite.mil, ldap.yoursite.mil]
`win_check_network_tcp`	`[]`	Services every Windows server must reach, e.g. win_check_network_tcp: - {name: LDAPS, host: dc01.yoursite.mil, port: 636} - {name: SIEM, host: siem.yoursite.mil, port: 9997}
`win_check_network_secure_channel`	`true`	domain members: test the computer's trust with the domain (nltest /sc_query)
`win_check_network_timeout_ms`	`3000`	milliseconds for each TCP test

win_check_patching

File: `roles/win_check_patching/defaults/main.yml`

Note
WINDOWS PATCHING (read-only; installing updates is a later runbook) - a restart is pending for installed updates (Windows Update, component servicing, ConfigMgr) - days since the last update was installed (Get-HotFix) - optional: ask Windows Update which updates are missing (slower; needs full PowerShell)

Setting	Default	What it does
`win_check_patching_warn_days`	`35`	last update installed this many days ago or more -> warning
`win_check_patching_crit_days`	`60`	this many days or more -> critical
`win_check_patching_reboot_pending`	`true`	report a restart pending for installed updates
`win_check_patching_search_pending`	`false`	also ask Windows Update which updates are missing (slower)

win_check_performance

File: `roles/win_check_performance/defaults/main.yml`

Note
WINDOWS PERFORMANCE: CPU (average of three samples a second apart), memory, page file. The finding names the busiest processes, so you know where to look.

Setting	Default	What it does
`win_check_performance_cpu_warn_pct`	`90`	CPU busy % (average of three samples) -> warning
`win_check_performance_cpu_crit_pct`	`98`	CPU busy % -> critical
`win_check_performance_mem_warn_pct`	`90`	memory in use % -> warning
`win_check_performance_mem_crit_pct`	`97`	memory in use % -> critical
`win_check_performance_pagefile_warn_pct`	`80`	page file in use % -> warning

win_check_security

File: `roles/win_check_security/defaults/main.yml`

Note
WINDOWS SECURITY BASELINE (a few high-value STIG items; not a full STIG scan: that is your SCAP / vulnerability scanner)

Setting	Default	What it does
win_check_security_av	auto	Antivirus: auto = check Microsoft Defender when it is the active antivirus (skipped when it is passive or off because another product protects the server); defender = always; none = never.
win_check_security_av_signature_max_days	7	Defender's signatures older than this many days -> finding
win_check_security_av_services	[]	Services of your endpoint protection that must run, e.g. Trellix / McAfee: [masvc, mfemms]
win_check_security_firewall_profiles	[Domain, Private, Public]	must be on
win_check_security_smb1	true	SMBv1 must be off
win_check_security_rdp_nla	true	Remote Desktop must require Network Level Authentication
win_check_security_uac	true	User Account Control must be on
win_check_security_legacy_tls	true	TLS 1.0 / 1.1 must be off for incoming connections

win_check_services

File: roles/win_check_services/defaults/main.yml

Note
WINDOWS SERVICES - every required service must be running (critical); - an automatic service that stopped WITH AN ERROR (exit code not 0) is a warning. Automatic services that stop cleanly by design (trigger-start, delayed) are not reported; - a service that crashed several times recently (System log events 7031 / 7034).

Setting	Default	What it does
win_check_services_required	[EventLog, RpcSs, Dnscache, W32Time, WinRM, mpssvc]	Must be running on every Windows server. Add your own per group with ..._required_extra, e.g. for IIS servers: win_check_services_required_extra: [W3SVC, WAS]
win_check_services_required_extra	[]	more services that must run, e.g. for a group of IIS servers (see above)
win_check_services_ignore	- gupdate - gupdatem - edgeupdate - edgeupdatem - MicrosoftEdgeElevationService - GoogleU…	Automatic services never to report (wildcards allowed: * and ?)
win_check_services_crash_hours	24	look back this far for unexpected stops

Setting	Default	What it does
win_check_services_crash_warn	3	this many unexpected stops of one service -> warning

win_check_stig

File: roles/win_check_stig/defaults/main.yml

> **Note**
> WINDOWS STIG DRIFT (playbooks/win_stig_drift_report.yml): Evaluate-STIG on the server - every STIG it finds that applies (Windows Server 2019 / 2022 / 2025, Defender, Firewall, .NET, IIS, Edge ...) - each checklist compared with STIG Manager (read-only) and with the server's earlier scans, as on Linux. STIG Manager, the findings and the rules left out are the check_stig_* settings (roles/check_stig/defaults/main.yml), shared with Linux. The job never installs or updates Evaluate-STIG. docs/STIG_DRIFT.md.

Setting	Default	What it does
win_check_stig_evaluate	(empty)	---- Evaluate-STIG on the server ---------- -- its folder on the server (or Evaluate-STIG.ps1), e.g. 'C:\Evaluate-STIG'; empty = found: a scheduled task that runs Evaluate-STIG.ps1, else a folder named Evaluate-STIG* on the local drives (win_check_stig_s... (all of it: the role file)
win_check_stig_evaluate_read	(empty)	read the checklists Evaluate-STIG already wrote instead of running it: auto = the folder its scheduled task writes to (-OutputPath), else its default C:\Users\Public\Documents\STIG_Compliance; or a folder - this server's newest checklist of each STIG (not its Previous folders, not another server's); empty = run it
win_check_stig_evaluate_select	[]	-SelectSTIG: only these STIGs, by the short names your Evaluate-STIG lists (Evaluate-STIG.ps1 -ListSupportedProducts), e.g. [WinServer2022, IIS10Server, IIS10Site]; [] = every STIG that applies to the server
win_check_stig_evaluate_exclude	[]	-ExcludeSTIG: leave these out (short names, as above); [] = none
win_check_stig_evaluate_args	[]	more arguments for Evaluate-STIG, e.g. ["-ScanType", "Classified", "-AnswerKey", "SITE"]; -Output and -OutputPath are the job's
win_check_stig_evaluate_output	CKLB	-Output: CKLB, or CKL (the older XML checklist)

Setting	Default	What it does
win_check_stig_evaluate_rejected	true	a checklist Evaluate-STIG's own schema check rejects (moved to its Bad_CKL folder): true = read it from there anyway, said in the report; false = the server is not scanned
win_check_stig_run_as	task	task = as SYSTEM in a one-time scheduled task (deleted after; outside the WinRM session's memory and time limits - Evaluate-STIG wants SYSTEM or an administrator); direct = in the job's own WinRM session, as its user
win_check_stig_powershell	(empty)	the PowerShell that runs Evaluate-STIG.ps1; empty = Windows PowerShell 5.1 (C:\Windows\System32\WindowsPowerShell\v1.0\powershell.exe)
win_check_stig_timeout	3600	seconds Evaluate-STIG may take (every STIG of a server: five to thirty minutes); then it is stopped and the server not scanned
win_check_stig_search	[]	folders to search for Evaluate-STIG when win_check_stig_evaluate is empty and no scheduled task runs it; [] = the local drives
win_check_stig_search_depth	4	how many folder levels down the search goes (C:\Users\admin\Desktop\Evaluate-STIG is 4)
win_check_stig_dir	'C:\ProgramData\site-automation\stig'	---- Kept on the server ---------- one folder per checklist (a STIG, or a STIG of one web site or database instance): each run's summary (JSON) and checklist (STIG Manager imports it), latest-summary.json; only SYSTEM and Administrators can o... (all of it: the role file)
win_check_stig_keep	10	runs kept per checklist (the oldest deleted); when each rule last passed is kept in every summary, so it outlives them

win_check_time

File: roles/win_check_time/defaults/main.yml

Note
WINDOWS TIME: the Windows Time service runs, syncs with a real source (not the local clock), synced recently, and how far off the clock is from its source. Kerberos (domain logons) fails when a clock is 5 minutes off.

Setting	Default	What it does
`win_check_time_offset_warn_sec`	`1`	seconds off its time source -> warning
`win_check_time_offset_crit_sec`	`5`	seconds off -> critical
`win_check_time_max_hours_since_sync`	`24`	last successful sync longer ago than this many hours -> finding

win_check_trust

File: `roles/win_check_trust/defaults/main.yml`

Note
WINDOWS TRUST STORES: which still hold a certificate authority that was compromised or that the browsers dropped - the machine certificate stores, the running account's, every Java cacerts and Python certifi bundle found - matched against the trust check's list (roles/check_trust/vars/main.yml, plus check_trust_extra; a severity changed in check_trust_severity). A certificate in the Disallowed store is blocked: listed, no finding. Removal (check_trust_remove: true, on the Windows trust store report): the certificate goes into the Disallowed store and out of the store it was in, Java entries are deleted, certifi bundles cut - each backed up to win_check_trust_backup_dir first. Check mode shows what would go. docs/TRUST_STORES.md

Setting	Default	What it does
`win_check_trust_stores`	`['Cert:\LocalMachine\Root', 'Cert:\LocalMachine\CA', 'Cert:\LocalMachine\AuthRoot', 'Cert…`	the certificate stores read (CurrentUser: the account AAP connects with)
`win_check_trust_block_stores`	`['Cert:\LocalMachine\Disallowed', 'Cert:\CurrentUser\Disallowed']`	the Disallowed stores: a certificate there is distrusted (blocked)
`win_check_trust_java`	`true`	every Java cacerts under the folders below (JKS natively; PKCS12 with the JDK's keytool)
`win_check_trust_java_dirs`	`['C:\Program Files\Java', 'C:\Program Files (x86)\Java', 'C:\Program Files\Eclipse Adopti…`	folders searched for cacerts (seven levels deep)
`win_check_trust_python`	`true`	every certifi cacert.pem under the folders below
`win_check_trust_python_dirs`	`['C:\Program Files\Python*', 'C:\Program Files (x86)\Python*', 'C:\Python*', 'C:\ProgramD…`	folders searched for certifi (seven levels deep)

Setting	Default	What it does
win_check_trust_backup_dir	'C:\ProgramData\site-automation\trust-backup'	the removal's backups: every certificate and file changed, a folder per run
win_check_trust_timeout	60	seconds for keytool

win_patch

File: roles/win_patch/defaults/main.yml

Note
WINDOWS PATCH through the ConfigMgr client: install what Software Center has for the server (the software updates deployed to it), restart if the updates need it, and prove the server came back: every automatic service that was running before runs again. One server at a time by default; the rollout stops at the first server that fails. Check mode lists and changes nothing.

Setting	Default	What it does
automatic_restarts	false	NO AUTOMATIC RESTARTS unless you say so: this playbook restarts a server only when automatic_restarts is true (the job template's Variables, or the survey) and the updates need it. With automatic_restarts false (the default) it installs the updates and tells you which servers still need a restart; a person restarts th... (all of it: the role file)
win_patch_serial	1	How many servers are patched at the same time (1 = one by one: the rollout stops at the first server that fails). ONLY in the job template's Variables, a survey or the launch: Ansible picks the batch size before it reads all.yml or any group or host variables, so set there it does nothing.
win_patch_never_reboot_groups	"{{ patch_never_reboot_groups | default(['aap', 'aap_hosts', 'sat', 'netapp', 'netapp_con…	Servers in these groups are patched but NEVER restarted by this playbook (restart them by hand). Follows the Linux setting (patch_never_reboot_groups) when you set it, so ONE list covers both; the fallback is the Linux role's default (roles/patch/defaults/main.yml; tests/test_protection.py checks it).
win_patch_never_patch_groups	"{{ patch_never_patch_groups | default(['no_patch', 'aap', 'aap_hosts', 'sat', 'netapp',…	Servers in these groups are skipped entirely (e.g. servers patched only by hand).

Setting	Default	What it does
win_patch_scan_first	true	ask the client to rescan and re-evaluate before installing
win_patch_scan_wait_sec	120	let that scan settle
win_patch_timeout_min	120	per round: download + install
win_patch_poll_sec	30	how often to look at the progress
win_patch_max_rounds	2	After a restart, look again: servicing-stack and cumulative updates often reveal more updates.
win_patch_restart_delay_sec	30	shutdown.exe /t
win_patch_restart_timeout_min	30	how long the server may take to come back
win_patch_restart_settle_sec	30	after the port answers again, before trying to log on
win_patch_reconnect_pause_sec	20	between logon attempts after the restart
win_patch_services_grace_min	5	After the restart, give delayed-start services this long before calling one missing.
win_patch_services_ignore	['gupdate*', 'edgeupdate*', 'MicrosoftEdgeElevationService', 'GoogleUpdater*', 'MapsBroke…	Automatic services that come and go on their own (wildcards); not compared before/after.
win_patch_winrm_host	"{{ ansible_host | default(inventory_hostname) }}"	Where AAP reaches WinRM, to see the server go down and come back during the restart
win_patch_winrm_port	"{{ ansible_port | default(5986) }}"	its WinRM port (5986 = HTTPS)

win_troubleshoot

File: roles/win_troubleshoot/defaults/main.yml

Setting	Default	What it does
ts_area	overview	WINDOWS TROUBLESHOOT: the PowerShell commands an experienced Windows admin runs first for a kind of problem, run for you (read-only), each with a line saying what it shows. The matching health checks run too, so you also get plain-words findings. ts_area: overview | disk | performance | service | network | login | tim... (all of it: the role file)

Setting	Default	What it does
ts_service	(empty)	the service name, for ts_area service (see above)
ts_target	(empty)	host:port, for ts_area network
ts_hours	2	how many hours back the event logs are read
ts_output_lines	60	keep the last this-many lines of each command
win_ts_areas	overview: checks: [disk, services, performance, time, eventlog] commands: - {why: "Window…	What each area runs. why is printed above the command's output. Commands must be read-only.

Settings of one host

Set these on the host itself: its host variables in the inventory, or the host's **Variables** box in AAP. Never in a job template's Variables - there one value would be used for every host. When a host has none, the role works the value out from the host, so no role file has a default for them.

Setting	Role	When not set	What it does
net_config_address	net_config	the host's ansible_host, else its name	The address the backup and load jobs connect to, when it is not ansible_host (PAN-OS)
net_config_platform	net_config	from ansible_network_os	ios or panos: the platform of a host that has no ansible_network_os (cisco.ios.ios / paloaltonetworks.panos)
net_cert_address	net_certs	the host's ansible_host, else its name	The address the certificate report connects to (a host of net_cert_group)
net_cert_servername	net_certs	the host's name	The name sent in the TLS hello (SNI): the certificate a device with several names shows for this one
servicenow_node_label	servicenow	from servicenow_node_labels, else the host's short name	The friendly name of one self-hosted node in the health report ("Node 1"), set on that node's host