Settings Reference

Every setting of every job: its default, what it does, and where to set it

Site Automation - read it next to the setup guide | site-automation 0.25.0 | October 2026

What this is. Every setting the site-automation jobs read - the value it has when you set nothing, and what it does - with how to use one: where to put it (a template's Variables, `playbooks/group_vars/all.yml`, a group or a host), which place wins, and how to see what a job used.

> **How to read it**
> Keep it open next to the *Operations Runbooks Setup Guide*: each job's chapter there says which settings matter for it; this one lists them all, role by role.
> Rendered from `docs/VARIABLES.md` and `docs/VARIABLES_REFERENCE.md` (generated from `roles/*/defaults/main.yml`) in the site-automation repository (version 0.25.0).

Host names and addresses in this guide are placeholders. Copy YAML from the repository files, not from the PDF.

Contents

1. Settings: what they are, where they go, and who wins 4

The two-minute version . . . 4
How to use a setting . . . 4
1. What you can set . . . 5
2. Where do I put a setting? . . . 7
3. Who wins when the same setting is in two places? . . . 8
Examples . . . 10
How to verify a setting is in effect . . . 10
If a setting does nothing . . . 11

2. Every setting and its default 12

aap_email_question . . . 12
ad_lockout . . . 13
ad_report . . . 14
cert_inventory . . . 16
check_accounts . . . 17
check_auditd . . . 17
check_certs . . . 18
check_containers . . . 24
check_disk . . . 25
check_fapolicyd . . . 26
check_logging . . . 27
check_mariadb . . . 27
check_mounts . . . 30
check_network . . . 31
check_patching . . . 31
check_performance . . . 32
check_selinux . . . 32
check_services . . . 33
check_time . . . 33
check_trust . . . 34
horizon . . . 37
net_certs . . . 40

net_config . . . 42
ontap_report . . . 45
patch . . . 46
poam . . . 48
podman_discover . . . 49
service_watch . . . 50
servicenow . . . 51
site_act . . . 55
site_email . . . 57
site_findings . . . 60
stigman_stack . . . 61
troubleshoot . . . 63
vmware_vm . . . 63
win_act . . . 84
win_check_accounts . . . 85
win_check_audit . . . 85
win_check_certs . . . 86
win_check_disk . . . 88
win_check_eventlog . . . 89
win_check_network . . . 89
win_check_patching . . . 89
win_check_performance . . . 90
win_check_security . . . 90
win_check_services . . . 91
win_check_time . . . 92
win_check_trust . . . 92
win_patch . . . 93
win_troubleshoot . . . 94

1. Settings: what they are, where they go, and who wins

A **variable** (also called a *setting*) is a named value the runbooks read: `check_disk_warn_pct: 85` means "warn when a disk is 85 % full". You change a variable to change how a check behaves, without touching any code. This page answers three questions:

1. **What can I set?** (find the name and the default)
2. **Where do I put it?** (for every host, for one group, for one host, for one run)
3. **What happens when the same setting is in two places?** (who wins)

You do not need to read it all. Do the two-minute version first.

The two-minute version

1. **Find the setting.** Open VARIABLES_REFERENCE.md, press Ctrl+F and search for a word (`disk`, `certificate`, `mariadb`). Every setting is there with its default.
2. **Put the line in one file.**
   - Every host: `playbooks/group_vars/all.yml`
   - Hosts in one AAP group: `playbooks/group_vars/<group name>.yml`
   - One host: `playbooks/host_vars/<host name>.yml`
3. **Commit and Sync Changes** in VS Code. Run the check.

That is all for a normal change. The rest of this page explains the details and the exceptions.

For one job template: `aap/extra_vars/<playbook>.yml` (the list) has every setting that template reads, switched off, with its default and what it does. Copy the file into the template's **Variables** box (YAML selected), delete the `#` in front of the settings you change, and change their values. Switch on only what you change: a template's Variables beat `all.yml`.

How to use a setting

Every job's chapter in these docs ends its settings part with an **All settings** link: the list of that job's settings in VARIABLES_REFERENCE.md, each with its default and what it does. Pick the place by **who the setting is for**:

The setting is about	Put it in	Example
hosts (thresholds, what a check looks at)	`playbooks/group_vars/all.yml` (every host), `playbooks/group_vars/<group>.yml`, `playbooks/host_vars/<host>.yml`	`check_disk_warn_pct: 90`
one job template (who gets its email, its rules)	the job template's **Variables** box in AAP	`report_email_to: ops@yoursite.mil` on the certificate report template
one run	a **survey** question (its variable name = the setting's name) or **Prompt on launch**	a survey question `health_checks`, multiple choice `daily` / `weekly`

Three examples:

1. playbooks/group_vars/all.yml - every host: disks warn at 90 %, a stricter audit log partition

```
check_disk_warn_pct: 90
check_disk_overrides:
  /var/log/audit: {warn: 70, crit: 85}

# 2. the "Certificate report" template's Variables - this job only: email it, stay green
report_email_to: [pki-team@yoursite.mil]
site_fail_on: []

# 3. playbooks/group_vars/win_web.yml - the Windows web servers only: IIS must run too
win_check_services_required_extra: [W3SVC, WAS]
```

- **A list replaces the default list** (it does not add to it): repeat the entries you keep. Details: "Three shapes of value" below.
- **Check what a job used:** in AAP, the job's **Details** tab shows the template's and the launch's variables; the job's output prints the decisions it made (thresholds in the finding texts, `REFUSED:` lines, the settings each report names in its email); its **artifacts** (Details > Artifacts) carry the results. A setting that seems ignored: If a setting does nothing.

1. What you can set

Every setting lives, with its default and an explanation, in the role that uses it: `roles/<role>/defaults/main.yml`. Two easy ways to read them:

- **VARIABLES_REFERENCE.md**, the same information as one searchable page (generated from the role files, so it is never out of date). Start here.
- In VS Code: **Ctrl+Shift+F** (search in all files), type the setting or a word, and look at the hits under `roles\...\defaults\main.yml`.

Setting names start with the thing they belong to:

Starts with	Belongs to	Example
`check_disk_`, `check_services_`, `check_mariadb_` ...	one health check (`roles/check_<name>`)	`check_disk_warn_pct`
`win_check_`	one Windows health check	`win_check_disk_warn_pct`
`site_act_`	ACT (GenAI)	`site_act_level`
`site_` (other)	reporting and settings	`site_fail_on`
`patch_`, `win_patch_`, `automatic_restarts`	patching	`patch_never_patch_groups`
`servicenow_`	ServiceNow tickets	`servicenow_assignment_group`
`watch_`	service watch	`watch_containers`, `watch_discover`
`podman_discover_`	finding the podman containers (the `containers` check and service watch share it)	`podman_discover_ignore`
`poam_`	POA&M	`poam_stigman_api`
`stigman_`	STIG Manager deployment	`stigman_scope`

Check names and what each finding means: RUNBOOKS.md.

Three shapes of value

```yaml
check_disk_warn_pct: 85                    # a number
check_fapolicyd_required: true             # true or false (lower case)
check_services_required: [sshd, crond]     # a list: [a, b] or one "- item" per line
check_disk_overrides:                      # a set of names with values (a "map")
  /var/log/audit: {warn: 70, crit: 85}
```

A list replaces the default list. It does not add to it. If the default is `[sshd, crond, rsyslog]` and you write `[chronyd]`, the hosts now require **only** `chronyd`. Copy the default list from the reference and add your entries to it. (Some settings have a separate `..._extra` version that *adds*, for example `check_services_required_extra`. The reference says which.)

Job variables (per run, not per host)

These choose what a run does. Put them in a job template's **Variables** box (*Extra variables*) or in a survey question, never in a settings file:

Variable	Meaning	Where it appears
`target`	which hosts or group to run on (default: all)	most playbooks
`health_checks`	which checks: `daily`, `weekly`, `all`, or a list	Health check
`ts_area`, `ts_service`, ...	what to troubleshoot	Troubleshoot
`use_act`, `site_act_level`	ask ACT, and what it may do	Health check, Troubleshoot
`site_fail_on`	which findings fail the job: `[critical]` or `[critical, warning]`	any check job

2. Where do I put a setting?

Place	Who gets it	How you change it	Good for
`roles/<role>/defaults/main.yml`	everyone	do **not** edit (updates replace it)	the built-in default
`playbooks/group_vars/all.yml`	every host	VS Code, commit, sync	site-wide values: the LDAP server, the SIEM, thresholds
`playbooks/group_vars/<group>.yml`	the hosts of that AAP group	VS Code	one kind of server: STIG Manager containers, MariaDB, AAP itself
`playbooks/host_vars/<host>.yml`	that one host	VS Code (create the folder and file)	one odd server: a bigger disk limit for `/var/log/audit`
AAP: inventory > Variables	every host	AAP web UI	values you prefer to keep in AAP
AAP: group > Variables	the hosts of that group	AAP web UI	same, per group
AAP: host > Variables	that host	AAP web UI	same, per host
Job template > Variables (*Extra variables*)	every run of that template	AAP web UI	"this template always uses `site_act_level: diagnose`"
Survey question / prompt on launch	this launch only	asked at launch	a choice made each time
AAP credential	the job	AAP web UI	**secrets only** (SECRETS.md)

File names matter. `playbooks/group_vars/<group>.yml` applies to the group with **exactly that name** in your AAP inventory, and `playbooks/host_vars/<host>.yml` to the host with exactly that name. A file for a group your inventory does not have is silently never used.

The files shipped in `playbooks/group_vars/` and the group each expects:

File	Group name it applies to	If your group has another name
`all.yml`	every host	(always works)
`aap.yml`	`aap`	copy it to `aap_hosts.yml` if your group is `aap_hosts` (the example inventory's name)
`mariadb.yml`	`mariadb`	the MariaDB check also runs for a group named `mariadb_hosts`, but the **settings** come from the file named after the group: copy it to `mariadb_hosts.yml`
`sat.yml`, `idm.yml`, `logstash.yml`, `stigman.yml`, `rhel8_all.yml`	the same name	copy or rename
`netapp.yml`	`netapp`	copy to `netapp_console_hosts.yml` if that is your group

(The built-in safety lists cover both spellings of the AAP and NetApp group names, so AAP stays protected whatever you call it.)

Updating the repository never overwrites your settings files. The update script (`scripts/update-from-release.ps1` / `.sh`) protects `playbooks/group_vars/`, `playbooks/host_vars/`, `inventories/site/`, `poam/poam.csv` and anything in `.site-local`. If a new release adds a settings file you do not have, the script adds it **once** (the preview lists it as `YOURS`) and never changes it again. A setting that the release's `all.yml` mentions and yours does not is listed under `NEW SETTINGS`, so you can copy it over if you want it.

Adding "extras" to a host or a group in AAP

Where the boxes are in AAP 2.7 (older versions word a label differently):

- **Inventory**: **Automation Execution › Infrastructure › Inventories** › your inventory › **Edit inventory** › **Variables**.
- **Group**: the same inventory › **Groups** tab › the group › **Edit group** › **Variables**.
- **Host**: the same inventory › **Hosts** tab › the host › **Edit host** › **Variables**.
- **Job template**: **Automation Execution › Templates** › the template › **Edit template** › **Extra variables** (the template's *Variables* box).

The **Variables** box on a host or group takes the same YAML as a settings file. Typical entries:

a group's Variables box: hosts in this group need a bigger disk margin

```
check_disk_overrides:
  /var/lib/pgsql: {warn: 70, crit: 85}
```

a host's Variables box: its MariaDB containers belong to a rootless account

```
check_mariadb_container_user: svc_podman
```

a group's Variables box: never check or watch the test containers, or anything of one user

```
podman_discover_ignore: ['test-.*', 'devuser/.*']
```

Connection settings (`ansible_host`, `ansible_user`, `ansible_python_interpreter`, `ansible_connection: winrm`) belong here too. The *login and password* never do: they are credentials (SECRETS.md).

Prefer the project's settings files (this repository), because they are reviewed and versioned in Git and updates never overwrite them. Use the AAP boxes when AAP is where your team already manages the inventory. Either is fine; the rule is **one setting, one place**.

3. Who wins when the same setting is in two places?

Strongest first. This order was **proven by running a test playbook** on ansible-core 2.16 and 2.21 with the same value set in every place and removing the winner one place at a time (the same result on both versions).

#	Place	Notes
1	**Extra variables**: job template *Variables*, a survey answer, a prompt on launch, `-e` on the command line	beats everything below, including the built-in protection further down
2	`playbooks/host_vars/<host>.yml`	this project, one host

#	Place	Notes
3	the **host's** Variables box in AAP	
4	`playbooks/group_vars/<group>.yml`	a child group beats its parent group
5	`playbooks/group_vars/all.yml`	
6	the **group's** Variables box in AAP	a deeper group beats its parent
7	the **inventory's** Variables box in AAP	
8	`roles/<role>/defaults/main.yml`	the built-in default

Things that surprise people:

- **A file in this project beats the AAP group and inventory boxes**, even when the box is on a more specific group. If you type a value in an AAP box and see no effect, look for the same setting in `playbooks/group_vars/`. Every job prints a notice, `NOTE - settings in this project's files win over the same setting in AAP's inventory Variables box`, listing which files define which settings. (Silence it with `site_settings_notice: false`.) Not shown for `playbooks/host_vars/`.
- **A host's own Variables box (3) beats a group file (4/5)**, but a `playbooks/host_vars` file (2) beats the host's box.
- **A survey answer or template extra variable beats everything**: it is the one place that can override the rest for a single run.
- Two sibling groups (neither is inside the other) that set the same variable for the same host: do not rely on the result. Put it in a host file, or in one group.
- An inventory **sourced from the project** (`inventories/site/`, SETUP_AAP.md step 5): AAP imports its `group_vars/all.yml` into the inventory's Variables box and each `group_vars/<group>.yml` into that group's box. So they rank as places 7 and 6, and the project's `playbooks/` files beat them. Keep a setting in one of the two, not both.
- **Updated from 0.4.1 or older?** Your `playbooks/group_vars/all.yml` is yours, so the update left it as it was: it still sets `site_act_diagnose_only_groups`, `patch_never_patch_groups`, `patch_never_reboot_groups`, `site_act_provider` and `site_act_models` as real (uncommented) lines. They beat the same settings in AAP's inventory and group boxes. If you want those set in AAP (or the new built-in defaults), put `#` in front of those lines. See the upgrade notes in `CHANGELOG.md` (0.5.0).

Built-in protection: settings that cannot be weakened from a file

Some lists are protected in code, so a mistake in a settings file cannot expose the AAP host:

What	Built in	What your file can do
Never patched, never rebooted	`aap`, `aap_hosts` (`roles/patch/vars/main.yml`)	your `patch_never_patch_groups` / `patch_never_reboot_groups` list **replaces** the shipped one (`sat`, `netapp`, ...): repeat the names you keep. `aap` and `aap_hosts` stay protected whatever you write
ACT never fixes by itself	`aap`, `aap_hosts` (`roles/site_act/vars/main.yml`)	your `site_act_diagnose_only_groups` list **replaces** the shipped one (`sat`, `idm`, `netapp`, ...): repeat the names you keep. `aap` and `aap_hosts` stay protected whatever you write

Only an extra variable (place 1) can override these, and nothing in these docs asks you to do that.

Examples

Everyone: warn earlier on disks. `playbooks/group_vars/all.yml`:

```yaml
check_disk_warn_pct: 80
check_disk_crit_pct: 92
```

One group: the STIG Manager hosts need these containers watched.
`playbooks/group_vars/stigman.yml` (shipped; already has `watch_containers`).

One host: a database server whose audit partition must stay emptier.
`playbooks/host_vars/db01.example.mil.yml` (the file name is the host's name in AAP):

```yaml
check_disk_overrides:
  /var/log/audit: {warn: 70, crit: 85}
```

One template: the "Troubleshoot" template always uses Ask Sage. In the template's **Variables**:

```yaml
site_act_provider: asksage
```

A list you want to extend. Where the role has a `..._extra` setting, use it: it *adds*.

```yaml
check_services_required_extra: [tomcat]        # required on top of the built-in five
```

Where it has none, look up the default (`check_services_required` is `[sshd, crond, rsyslog, chronyd, auditd]`) and write the whole list with your addition:

```yaml
check_services_required: [sshd, crond, rsyslog, chronyd, auditd, tomcat]
```

How to verify a setting is in effect

1. **Run the check on one host**: **Health check** with `target` = that host and `health_checks` = the check. The finding text quotes the limit it used (`/var is 96% full (critical at 92%)`), so you see the value that won.
2. **Break it on purpose** (or lower the limit) on a test host, and confirm the finding appears. A check that never fires proves nothing.
3. **Read the job's notice** (`NOTE - settings in this project's files ...`): it names every settings file that set something, so you can spot a file that overrides an AAP box.
4. **From the command line** (optional, in a terminal with Ansible and the repository): create a temporary file `playbooks/show_settings.yml`, **do not commit it**, and run it:

YAML

```yaml
- name: Show the settings a host gets
  hosts: "{{ target | default('all') }}"
  gather_facts: false
  tasks:
    - ansible.builtin.import_role: {name: check_disk}    # the role that owns the setting
      when: false                                        # loads its defaults, runs nothing
    - ansible.builtin.debug:
        msg: "{{ item }} = {{ lookup('vars', item, default='(not set)') }}"
      loop: [check_disk_warn_pct, check_disk_crit_pct]
```

shell

```
ansible-playbook -i inventories/site/hosts.yml playbooks/show_settings.yml -e
  target=db01.example.mil
```

It shows the value from the role default plus the files in this project. (Values typed into AAP's Variables boxes exist only inside AAP; there, use steps 1-3.)

If a setting does nothing

Symptom	Likely cause
You set it in AAP and nothing changes	the same setting is in `playbooks/group_vars/`: that file wins (table above)
You set it in a file and nothing changes	the file name does not match the group or host name exactly; a spelling or capital letter
A list lost entries	a list replaces the default: repeat the entries you keep
YAML error at the start of the job	indentation (spaces, never tabs), a missing space after `:`, or a value with `:` inside it that needs quotes
A survey answer is ignored	the survey question's **variable name** is not exactly the setting's name

2. Every setting and its default

This page lists **every setting a role reads**, with the value it has when you set nothing, and what the role's own comment says. It is generated from `roles/<role>/defaults/main.yml`, so it is never out of date. Where to put a setting, and which place wins when it is in two, is in VARIABLES.md. Values are shortened to one line here: the role file has the full one.

A setting that is a list (`[a, b]`) **replaces** the default list when you set it; it does not add to it. Repeat the entries you want to keep.

Contents: aap_email_question, ad_lockout, ad_report, cert_inventory, check_accounts, check_auditd, check_certs, check_containers, check_disk, check_fapolicyd, check_logging, check_mariadb, check_mounts, check_network, check_patching, check_performance, check_selinux, check_services, check_time, check_trust, horizon, net_certs, net_config, ontap_report, patch, poam, podman_discover, service_watch, servicenow, site_act, site_email, site_findings, stigman_stack, troubleshoot, vmware_vm, win_act, win_check_accounts, win_check_audit, win_check_certs, win_check_disk, win_check_eventlog, win_check_network, win_check_patching, win_check_performance, win_check_security, win_check_services, win_check_time, win_check_trust, win_patch, win_troubleshoot

aap_email_question

File: `roles/aap_email_question/defaults/main.yml`

> **Note**
> THE "ALSO EMAIL THE REPORT TO" QUESTION (playbooks/aap_email_question.yml): adds a text box to the survey of every report template and of the workflows that run them, through AAP's API, so whoever launches one can send that run's report to more people (report_email_also, added to the template's report_email_to). Run it once, and again after you add report templates. docs/EMAIL_REPORTS.md, "Also email the report to".

Setting	Default	What it does
`aap_email_question_state`	`present`	present = add the question; absent = take it out again
`aap_email_question_playbooks`	`- act_fix_approved.yml - ad_health_report.yml - ad_lockout.yml - cert_report.yml - databa…`	Which templates: the job templates whose playbook is one of these (the file name - your template names do not matter). Every playbook that can email a report is here; a test keeps the list complete.
`aap_email_question_templates`	`[]`	more job templates or workflows, by name or ID
`aap_email_question_exclude`	`[]`	templates left out, by name or ID
`aap_email_question_workflows`	`true`	also the workflows that run them (a job in a workflow is never asked its own survey)

Setting	Default	What it does
`aap_email_question_project`	(empty)	only job templates of this AAP project (its name); empty = any project
`aap_email_question_ca_path`	(empty)	a CA file for AAP's HTTPS certificate ("" = what the execution environment trusts)
`aap_email_question_text`	`Also email the report to`	The question, as whoever launches a template sees it. A text box of several lines (a pasted list keeps its line breaks), not required, empty by default.
`aap_email_question_description`	(worked out by the role; see its file)	The help text under the question.

ad_lockout

File: `roles/ad_lockout/defaults/main.yml`

> **Note**
> Account lockout (playbooks/ad_lockout.yml): is the account locked out, and WHAT is locking it out - the computers its bad passwords came from, from the Security log of every domain controller. Read-only, unless ad_lockout_unlock. A PowerShell script on every DC (WinRM, the Windows Machine credential): the account needs to read the DCs' Security logs - see docs/ACTIVE_DIRECTORY.md, "Account lockout".

Setting	Default	What it does
`ad_lockout_user`	(empty)	The account to look at, as the service desk is given it: jdoe, CORP\jdoe or jdoe@corp.example.mil (a survey question, NOT required). Several: one per line, or separated by commas. "" = every account locked out now, each with the computer locking it out.
`ad_lockout_hours`	`24`	the Security log entries of the last this-many hours (raise it for a lockout of last week)
`ad_lockout_max_users`	`25`	at most this many accounts (asked for, or locked out now) in one run
`ad_lockout_unlock`	`false`	true = unlock the account (on the PDC emulator) after reading - only an account asked for by name, and only when it is locked out. A dry run (Check) says what it would unlock. It locks out again as soon as the source computer tries its old password again: the report says what to fix first.
`ad_lockout_check_sources`	`false`	true = also look on each source computer that is a host of this job's inventory (Windows, WinRM, the same credential - it needs administrator rights there): services and scheduled tasks running as the account, its sessions (a disconnected Remote Desktop session), processes and mapped drives. Read-only.

Setting	Default	What it does
ad_lockout_check_max	5	look on at most this many source computers per account (the worst first)
ad_lockout_exchange_servers	[]	Your mail servers (Exchange / ActiveSync, by name or address, * and ? allowed, e.g. [EXCH*, 192.0.2.25]): bad passwords coming through one of them are a phone or a mail app with the old password.
ad_lockout_max_events	2000	---- How much is read ---------- read at most this many Security log entries per account and DC (the newest)
ad_lockout_resolve_max	50	resolve at most this many addresses to names per DC (reverse DNS)
ad_lockout_max_rows	50	rows per table; the job's artifacts (ad_lockout) have all

ad_report

File: roles/ad_report/defaults/main.yml

> **Note**
> Active Directory health report (playbooks/ad_health_report.yml). Read-only: a PowerShell script on every domain controller (WinRM, the Windows Machine credential) and one more on one DC per domain. The account needs to read AD and run dcdiag / repadmin on the DCs: see docs/ACTIVE_DIRECTORY.md.

Setting	Default	What it does
ad_report_hours	24	---- Changes ---------- the changes of the last this-many hours (match it to the schedule: a daily job = 24)
ad_event_hours	24	the event log entries of the last this-many hours
ad_security_events	true	read each DC's Security log for WHO made a change (needs auditing of account and security group management)
ad_max_events	2000	read at most this many entries per event log and DC
ad_changes_max	500	look at most at this many changed objects per kind (created, changed accounts, GPO links)
ad_privileged_groups_extra	[DnsAdmins]	Privileged groups watched besides the built-in ones (Domain / Enterprise / Schema Admins, Administrators, Account / Server / Print / Backup Operators, Group Policy Creator Owners, Key Admins, Cert Publishers): names or SIDs of this domain's groups.
ad_groups_should_be_empty	[]	privileged groups that must have no member, e.g. [Schema Admins] - amber when they have one

Setting	Default	What it does
ad_privileged_change_severity	warning	a member added to / removed from a privileged group: warning, critical or info
ad_schema_change_severity	warning	the schema version changed: warning, critical or info
ad_dc_change_severity	info	a domain controller added or removed: warning, critical or info
ad_dcdiag_skip	[SystemLog]	---- What is read on each DC ---------- dcdiag tests to skip (SystemLog fails on any error in the System log: the Events table covers it)
ad_dcdiag_dns	true	also dcdiag's DNS tests that only read (basic, forwarders, delegation, record registration)
ad_trust_verify	true	verify outbound and two-way trusts (nltest /sc_verify)
ad_repl_warn_hours	6	---- Thresholds: at or above "warn" = amber, at or above "crit" = red ---------- a replication link with no success for this long (inter-site schedules: keep above their interval)
ad_repl_crit_hours	24	no success for this long (or a failing link) -> red
ad_repl_queue_warn	50	replication requests waiting in a DC's inbound queue
ad_time_warn_seconds	5	a DC's clock against the PDC emulator (the PDC: against its NTP source)
ad_time_crit_seconds	60	clock offset -> red (Kerberos fails at 300)
ad_time_sync_warn_hours	24	Windows Time last synchronized longer ago than this
ad_disk_warn_pct	15	free space of the database, log and SYSVOL volumes, % (at or below)
ad_disk_crit_pct	5	free % -> red
ad_cert_warn_days	30	the LDAPS / KDC certificate expires within this many days
ad_cert_crit_days	7	expires within this many days -> red
ad_backup_warn_days	7	the last system-state backup of a partition is older than this
ad_backup_crit_days	30	older than this -> red (always red past the tombstone lifetime)
ad_krbtgt_max_days	180	the krbtgt password is older than this -> amber (the STIG asks for 180)
ad_rid_warn_pct	10	RID pool left, % (at or below)
ad_rid_crit_pct	2	RID pool left -> red
ad_no_client_site_warn	1	this many client addresses in no AD subnet -> amber (fewer: blue)

Setting	Default	What it does
ad_report_max_rows	50	rows per table (problems first); the job's artifacts have all
ad_report_fail	false	The job stays green when it finds problems: the report and email say what. true = mark it FAILED when there is a critical finding or a DC could not be read, for workflows.

cert_inventory

File: roles/cert_inventory/defaults/main.yml

> **Note**
> CERTIFICATE INVENTORY: every certificate report (Linux, Windows, VMware, network devices) can also write what it found - every certificate, with the same fields from each - for a dashboard (Power BI ...). Two ways, each on by itself: files (CSV + JSON) and a SQL Server table. A Check run writes nothing. The fields: id, source, host, location, subject, issuer, sans, not_after, days, status, tier (root / intermediate / ca-signed / self-signed), trust_store, in_use, public_key, sig_alg, serial, sha1, sha256, scanned_at, job (docs/CERT_INVENTORY.md).

Setting	Default	What it does
cert_inventory_files	false	---- files: <source>-certificates.csv and .json (source: linux, windows, vmware, net), replaced each run ---- true = write the files
cert_inventory_dir	/var/lib/site-automation/cert-inventory	where: on the AAP node the job runs on (a path AAP exposes to jobs: docs/CERT_INVENTORY.md), or on cert_inventory_host
cert_inventory_host	(empty)	an inventory host to write them on instead (a file server, or a Linux host with the share mounted); its credential must be on the job template
cert_inventory_history	false	true = also a dated copy (<source>-certificates-YYYY-MM-DD.csv) next to them, for trends
cert_inventory_sql	false	The login comes from the "SQL Server (certificate inventory)" credential (aap/credential_types/cert_inventory_sql.yml). The execution environment needs pymssql and community.general (the all-in-one EE has both). true = add / update the rows
cert_inventory_sql_server	(empty)	the SQL Server (name or address)
cert_inventory_sql_port	1433	its port
cert_inventory_sql_database	(empty)	the database (SQL Server 2016 or later: the rows go in as JSON, OPENJSON)

Setting	Default	What it does
cert_inventory_sql_table	dbo.cert_inventory	schema.table
cert_inventory_sql_create	false	true = create the table (and its view <table>_current) when it is not there - the account then needs CREATE TABLE

check_accounts

File: roles/check_accounts/defaults/main.yml

> **Note**
> ACCOUNTS: the account problems auditors and attackers both look for, plus the one that quietly breaks automation - a service account password about to expire.

Setting	Default	What it does
check_accounts_min_uid	1000	people's accounts start here
check_accounts_inactive_days	35	the STIG wants unused accounts disabled after 35 days
check_accounts_ignore_inactive	[]	Accounts that may stay unused (break-glass, service accounts that never log in interactively)
check_accounts_nopasswd_allowed	[]	sudoers entries allowed to use NOPASSWD, e.g. the account AAP connects as: [svc_aap, "%wheel"]
check_accounts_expiry_watch	[]	Accounts whose password/account expiry matters - above all the account AAP logs in with. e.g. [svc_aap]
check_accounts_expiry_warn_days	14	a watched account's password or account expiring within this many days -> warning (3 days or less -> critical)

check_auditd

File: roles/check_auditd/defaults/main.yml

> **Note**
> AUDITD: the audit system records who did what (the STIG depends on it). Running, enabled in the kernel, rules loaded, no events lost, room for its log, and (optionally) forwarding. Also how MUCH it records: too much rotates the evidence away within hours, fills the disk (then auditd stops the host or puts it in single-user mode, as the STIG sets it up), drops records on the way to the SIEM, and uses up the SIEM licence.

Setting	Default	What it does
check_auditd_min_rules	5	fewer loaded rules than this = the STIG rules are missing
check_auditd_log_warn_pct	75	the audit log's filesystem
check_auditd_log_crit_pct	90	% used of that filesystem -> critical (warning from check_auditd_log_warn_pct)
check_auditd_remote_required	false	true = au-remote (audisp-remote) must be active
check_auditd_exclude_accounts	[]	Accounts whose audit events are expected in bulk and do NOT count toward the volume warning, e.g. the account a vulnerability scanner logs in with (a credentialed scan runs thousands of commands with sudo). Names or UIDs; a host that does not have the account skips it. The report still says how much they made, and the... (all of it: the role file)
check_auditd_volume_hours	24	look at the audit records of the last N hours
check_auditd_volume_warn_mb_day	200	warning above this many MB of audit records a day (excluded accounts not counted)
check_auditd_retention_warn_hours	24	With max_log_file_action ROTATE, auditd keeps num_logs files of max_log_file MB and deletes the oldest. Warning when what is left on disk reaches back fewer hours than this. The STIG wants a week (168) when the records are NOT sent off the host as they are written (RHEL-08-030660). 0 = off.
check_auditd_volume_timeout	120	The breakdown (which account, which rule, which program) reads the audit log files of the window at low CPU and disk priority, and only when the volume or the retention is over its limit. It stops after this many seconds on a host with very large logs (and says so).

check_certs

File: `roles/check_certs/defaults/main.yml`

> **Note**
> CERTS: certificates that expire soon - as files on the host, inside Java keystores, and as served by TLS endpoints (what a browser or client actually sees).

Setting	Default	What it does
check_certs_warn_days	30	a certificate expiring within this many days -> warning (and amber in the certificate report)
check_certs_notice_days	60	the certificate report colours these blue (an early heads-up; not a finding)
check_certs_crit_days	14	within this many days (or expired) -> critical
check_certs_files	[]	Certificate files (PEM) to check. Globs work. Every certificate in a file is checked (a chain file with an expiring intermediate is a problem too). e.g. [/etc/stigman/*.crt, /etc/pki/tls/certs/server.crt]
check_certs_discover	true	Also look for certificates in the usual places. Files with many certificates are CA bundles (hundreds of roots, noisy) and are skipped unless listed in check_certs_files. Text (PEM) and binary (DER) certificates are both read; the names in any case (.crt, .CRT, .Pem ...).
check_certs_discover_dirs	[/etc/pki/tls/certs, /etc/pki/nginx, /etc/nginx, /etc/httpd/conf.d, /etc/stigman, /etc/st…	the folders searched (two levels deep) for *.crt, *.pem, *.cer, *.der (/etc/stigman/nginx/tls: STIG Manager's nginx certificate, three levels down)
check_certs_discover_extra_dirs	[]	more folders searched the same way, added to the ones above (no need to repeat them), e.g. [/home, /root, /opt/app/certs]
check_certs_discover_exclude	'ca-bundle|ca-certificates|/localhost\.crt$'	regular expression: found files whose path matches are skipped
check_certs_discover_max_per_file	10	a found file with more certificates than this is a CA bundle: listed as not checked (list it in check_certs_files to check each)
check_certs_placeholder_subject	'(^|,)O=SomeOrganization(,|$)'	Regular expression on the subject of a package's placeholder certificate: postfix's /etc/pki/tls/certs/postfix.pem (O=SomeOrganization), self-signed for a year at install and never renewed. A found one is listed as not checked (no finding, no ticket) and none gets a CSR. A real certificate put in its place is checked;... (all of it: the role file)

Setting	Default	What it does
check_certs_endpoints	[]	TLS endpoints, checked FROM this host, e.g. the site's own web server: check_certs_endpoints: - {name: STIG Manager, host: 127.0.0.1, port: 443, servername: stigman.example.mil}
check_certs_keystores	[]	Java keystores (Tomcat, the ServiceNow MID Server ...). A JKS keystore lists without a password; a PKCS12 one needs it - give the NAME of an environment variable that holds it (set by an AAP credential), never the password itself: check_certs_keystores: - {path: /opt/servicenow/mid/agent/security/agent_keystore, passw... (all of it: the role file)
check_certs_discover_ports	false	---- Finding them by themselves (the certificate report turns both on; the health check does not) ---- Listening TCP ports that answer TLS: each probed once, on the address it listens on, with `openssl s_client` (check_certs_probe_timeout seconds each, at most check_certs_discover_ports_max ports). Not these ports: pl... (all of it: the role file)
check_certs_discover_ports_skip	[22, 25, 53, 80, 110, 111, 123, 143, 323, 389, 631, 2049, 3306, 5355, 5432, 6379, 11211,…	ports never probed (plain-text services: ssh, smtp, dns, http, ldap, databases ...)
check_certs_discover_ports_max	25	probe at most this many listening ports (the rest are listed as not probed)
check_certs_probe_timeout	3	seconds for each port probe
check_certs_discover_keystores	false	Java keystores (*.jks, *.keystore, *.p12, *.pfx) in these folders, and the ones running Java programs use (-Djavax.net.ssl.keyStore=...). A keystore that needs a password is listed as not checked: no password is ever guessed - give it in check_certs_keystores (password_env). The JDK's own trust stores (cacerts) are sk... (all of it: the role file)
check_certs_discover_keystore_dirs	[/etc, /opt, /srv, /usr/share/tomcat, /var/lib/tomcat, /var/lib/pki]	the folders searched for keystores

Setting	Default	What it does
check_certs_discover_keystore_depth	5	how many folder levels deep the search goes
check_certs_discover_keystore_exclude	'/(cacerts|jssecacerts)$|/jre/lib/security/|/lib/security/'	regular expression: keystore paths skipped (the JDK's trust stores)
check_certs_discover_keystore_max	40	read at most this many found keystores per host
check_certs_discover_trust	false	The host's trust store: the CA certificates it trusts, put in by update-ca-trust (the DoD roots ...). Every file in these folders, whatever its name. The certificate report lists them with their dates - never a finding, a ticket or a CSR: a CA renews its own certificate (put its new one in the folder and run update-ca... (all of it: the role file)
check_certs_trust_dirs	[/etc/pki/ca-trust/source/anchors]	the folders read (not the folders below them)
cert_skip_hosts	[]	---- Leaving hosts out (the certificate report and the renewal; not the health check) ---- Hosts whose certificates something else renews - Satellite and its Capsules (satellite-installer, never by hand), appliances - are left out of the certificate report and never get a CSR: listed as not checked, no ticket. Name th... (all of it: the role file)
cert_skip_groups	[sat]	AAP inventory groups, e.g. [sat] (the health check still watches the Satellite web UI: group_vars/sat.yml)
cert_renew_dir	/etc/pki/tls/renew	---- Renewal: a new key and CSR for each certificate that expires soon (playbooks/cert_renew_csr.yml) ---- The certificates the report flags (within check_certs_warn_days - 30 days or less - or expired) get a new key and a CSR with the same subject and names (subjectAltName), on the host; the CSR is attached to the ce... (all of it: the role file)

Setting	Default	What it does
cert_renew_skip_hosts	[]	Hosts whose certificates are renewed by hand (IdM, AAP ...): in the certificate report and its tickets as usual - each ticket says it is renewed by hand - but the renewal makes no CSR for them (listed, with why). The same forms as cert_skip_hosts / cert_skip_groups (which leave a host out of the report too). e.g. [idm... (all of it: the role file)
cert_renew_skip_groups	[]	e.g. [idm]
cert_renew_extra_dirs	false	true = a CSR for certificates found in check_certs_discover_extra_dirs (/home, /root ...) too; false = they are in the report only (a file there also under check_certs_discover_dirs, or listed in check_certs_files, is renewed)
cert_renew_key_bits	2048	RSA key size, at least 2048 (a bigger old key keeps its size; an EC certificate gets a new key on the same curve)
cert_renew_digest	sha256	the CSR's signature: sha256, sha384 or sha512
cert_renew_max	10	at most this many NEW CSRs per host in one run (the rest come the next run)
cert_renew_only	(empty)	regular expression: only certificates whose place matches - a file's path, a keystore's path, 'TLS port 8443, ...' (empty = every one), e.g. '^/etc/pki/tls/certs/'
cert_renew_issuer	(empty)	regular expression: only certificates whose issuer matches (empty = any CA), e.g. 'O=U\.S\. Government'
cert_renew_exclude	'/(katello|foreman|foreman-proxy|candlepin|pulp|puppet|puppetlabs|rhsm|qpid)/|/etc/pki/(c…	Certificates something else renews - never a CSR from here. The path (Satellite, Capsule, Puppet, subscription-manager on every registered host) and the issuer (Satellite's own CA, Puppet's CA). Check the report's Where column for your hosts' paths and add yours.
cert_renew_exclude_issuer	'O=Katello|Puppet CA'	regular expression on the issuer: Satellite's own CA (O=Katello), Puppet's CA - never a CSR from here

Setting	Default	What it does
cert_renew_self_signed	false	Self-signed certificates (made on the host by an application or a person, not by a PKI) are reported and get a ticket, but no CSR - unless this is true: then the PKI is asked for a real certificate in their place (a STIG wants PKI certificates). Never a CA's own certificate (a root: it signs certificates), never one o… (all of it: the role file)
cert_renew_install_notes	/etc/stigman/: >- STIG Manager: put the issued certificate (with its chain) and the new k…	What to do with the issued certificate when something else deploys the file (a copy by hand would be overwritten): the start of the file's path -> the next step written on the ticket instead of "install it in place of the file". Add your own.
cert_renew_attach_ext	csr	The CSR's attachment on the ticket: <common name>-<8 characters of the CSR's SHA-256>.<ext> (a new CSR, a new name). An instance that refuses .csr attachments (its allowed extensions, or "Invalid file type: text/plain"): txt names every CSR .txt from the start. Refused, it is tried again as .txt (cert_renew_attach_ret… (all of it: the role file)
cert_renew_attach_type	text/plain	the attachment's Content-Type (a CSR is Base64 text)
cert_renew_attach_retry_txt	true	a refused attachment is tried again as <name>.txt (text/plain); false: straight to the work note
cert_install_services	[]	---- Install: the issued certificate in place of the old one (playbooks/cert_install_check.yml -> approval -> playbooks/cert_install_approved.yml) ---- With cert_renew_servicenow: catalog, the PKI puts the issued certificate on the request that carries its CSR. The install check finds it there (the certificate whose p… (all of it: the role file)

Setting	Default	What it does
cert_install_known_services	[]	The services whose configuration is searched for the certificate's file - [{dir, service, test}] (test: the command that checks the configuration before the reload). Empty = httpd (/etc/httpd, apachectl configtest), apache2, nginx (nginx -t), haproxy (haproxy -c), postfix (postfix check), dovecot (doveconf -n), slapd,... (all of it: the role file)
cert_install_never_reload	[]	services (globs) the job never reloads, e.g. ['slapd*']: their certificates are installed by a person
cert_install_key_files	{}	The old key's file, when it cannot be found - neither in the configuration that names the certificate, nor once in the usual folders (several files with the old key: held, to be named here): {certificate file: key file}, e.g. {/etc/pki/tls/certs/web01.crt: /etc/pki/tls/private/web01.key}
cert_install_key_dirs	[]	where else to look for the old key (empty = /etc/pki/tls/private, /etc/ssl/private, the services' folders ...)
cert_install_trust_check	true	the chain must build up to a root (the host's trust store, or the PKI's root sent with it); false = install without
cert_install_verify_wait	15	seconds the port has, after the reload, to serve the new certificate (then: put back)
cert_restore_backup	(empty)	---- Restore: an old certificate back from its backup (playbooks/cert_restore.yml) ---- Each install copies the files it replaces to cert_renew_dir/backup/<date>-<certificate>/. The restore puts them back, reloads the service, checks the port serves the old certificate again, and the renewal waits again (the next inst... (all of it: the role file)
cert_restore_only	(empty)	or: the newest backup of each certificate whose name or subject matches (a regular expression), e.g. 'web01'

check_containers

File: roles/check_containers/defaults/main.yml

Note
CONTAINERS: the podman containers on the host - root's (rootful) and every user's (rootless) - found by themselves (roles/podman_discover: which ones exist, which systemd unit runs each one, which should be running). Read-only. Which containers and users are looked at, and which are left out: the podman_discover_* settings.

Setting	Default	What it does
`check_containers_required`	`[]`	Containers that MUST exist and run on these hosts: a container name, or OWNER/NAME for one owner (root/web, alice/web). A name alone matches that container of any owner. Also matches a Quadlet file name without .container, e.g. [stigman-mysql, stigman-api, alice/web].
`check_containers_recent_hours`	`24`	A container that stopped with an error (exit code not 0) this many hours ago or less is reported (containers:crashed). Older crashes are history and are not reported.
`check_containers_restart_warn`	`3`	More restarts than this (podman's restart count, or its systemd unit's) = a restart loop.

check_disk

File: `roles/check_disk/defaults/main.yml`

Note
DISK: how full each LOCAL filesystem is (space and inodes), and - from the history of earlier runs - how many days until it is full. Network filesystems (NFS) are the mounts check's job: a hung NFS server can make `df` hang, so this check never touches them.

Setting	Default	What it does
`check_disk_warn_pct`	`85`	% used -> warning
`check_disk_crit_pct`	`95`	% used -> critical
`check_disk_inode_warn_pct`	`85`	inodes: too many small files fill a disk too
`check_disk_inode_crit_pct`	`95`	% of inodes used -> critical (warning from check_disk_inode_warn_pct)
`check_disk_overrides`	`{}`	Different limits for particular mount points, e.g. keep the audit log partition emptier: check_disk_overrides: /var/log/audit: {warn: 70, crit: 85}

Setting	Default	What it does
check_disk_ignore	- '^/(proc|sys|dev|run)(/|$)' - '^/etc/(hosts\|hostname\|resolv\.conf)$' # bind-mounted fil…	Mount points to skip (regular expressions).
check_disk_exclude_types	[tmpfs, devtmpfs, squashfs, overlay, iso9660, efivarfs, nsfs]	filesystem types never checked (df -x)
check_disk_forecast	true	Days-until-full forecast. Each run adds a sample (date, bytes used) per mount to a small file on the host; once there is at least a day of history, the growth rate predicts when the filesystem fills. Check mode (a dry run) reads the history but does not add to it.
check_disk_forecast_warn_days	14	full within this many days at the current growth -> warning
check_disk_forecast_crit_days	3	full within this many days -> critical
check_disk_history_file	/var/lib/site-health/disk-history.json	where the forecast keeps its samples, on the host
check_disk_history_keep_days	30	samples older than this are dropped (the forecast uses what is left)
check_disk_timeout	20	seconds for df

check_fapolicyd

File: roles/check_fapolicyd/defaults/main.yml

> **Note**
> FAPOLICYD (application allow-listing): running, enforcing, configuration and trust database valid, and recent denials - which is what you see when it blocks a program.

Setting	Default	What it does
check_fapolicyd_required	false	true = every host in the play must have fapolicyd. false = only check hosts that have it.
check_fapolicyd_expect_enforcing	true	permissive = 1 means it only logs, never blocks
check_fapolicyd_check_trustdb	true	Compare the trust database with the files on disk (can take a minute on big hosts; files changed outside RPM show up here, e.g. after a manual update).
check_fapolicyd_denials_since	yesterday	ausearch -ts value

Setting	Default	What it does
check_fapolicyd_denials_warn	1	this many denials or more since check_fapolicyd_denials_since -> warning

check_logging

File: roles/check_logging/defaults/main.yml

> **Note**
> LOGGING: are logs kept, and do they reach the central log server (SIEM)?

Setting	Default	What it does
check_logging_journal_persistent	true	the journal must survive a reboot (/var/log/journal)
check_logging_remote	[]	The central log server(s) this host forwards to, e.g. check_logging_remote: - {name: SIEM, host: siem.example.mil, port: 6514}
check_logging_timeout	5	seconds to wait for each log server's TCP port
check_logging_queue_warn_mb	100	rsyslog keeps messages it cannot deliver in /var/lib/rsyslog. A big queue = the log server has been unreachable for a while.
check_logging_errors_since	"-24h"	Error-priority journal messages in the last day. Many = something is complaining a lot.
check_logging_errors_warn	500	this many error messages or more -> warning

check_mariadb

File: roles/check_mariadb/defaults/main.yml

> **Note**
> DATABASE health for MariaDB AND MySQL (read-only; role name kept for compatibility): running, accepting connections, connection headroom, queries that run too long, replication (MariaDB multi-source, MySQL REPLICA/SLAVE), Galera / Percona XtraDB Cluster, MySQL Group Replication, recent errors in its log, and room on its disk. It runs status queries only - no data is read and nothing is changed. The engine and version are detected; see docs/MARIADB.md.

Setting	Default	What it does
check_mariadb_enabled	(worked out by the role; see its file)	Which hosts: members of any of the inventory groups mariadb_hosts, mysql_hosts, database_hosts, mariadb or mysql, or any host that names a container below (or turns discovery on). Set this true/false to override for a host or group.
check_mariadb_container	(empty)	MariaDB or MySQL in podman containers (e.g. the ServiceNow database). The client inside each container is used (podman exec). Name them in one of these ways; all are checked, each as its own instance: check_mariadb_container: mariadb one container check_mariadb_containers: [mariadb, snow-db] several containers on the... (all of it: the role file)
check_mariadb_containers	[]	several containers on one host, e.g. [mariadb, snow-db] (see above)
check_mariadb_container_discover	auto	auto, true or false: find database containers by their image (see above)
check_mariadb_server_paths	[/usr/sbin/mariadbd, /usr/libexec/mariadbd, /usr/sbin/mysqld, /usr/libexec/mysqld, /usr/b…	Where a database server installed on the host lives (any one present = a host install, for auto).
check_mariadb_container_image_regex	'(mariadb|mysql|percona)'	regular expression: a container whose image matches is a database (discovery)
check_mariadb_container_user	(empty)	ROOTLESS podman: the containers belong to a user (e.g. AAP's own MariaDB, or an application account), and only that user can see them. Name the user and the check runs podman AS that user. Empty = root's containers.
check_mariadb_service	(empty)	The systemd unit (host installs). Empty = the first of mariadb, mysqld, mysql that exists.
check_mariadb_client_args	[]	How to log in. Without a credential: as root through the local socket - MariaDB's default for a host install on RHEL (no password needed when the job runs with sudo). A database in a CONTAINER usually gives root a password (the official images do), and MySQL host installs usually do too, so there attach the "MariaDB m... (all of it: the role file)
check_mariadb_conn_warn_pct	80	connections in use, % of max_connections

Setting	Default	What it does
check_mariadb_conn_crit_pct	95	connections in use, % of max_connections -> critical
check_mariadb_threads_running_warn	32	queries executing at this moment
check_mariadb_long_query_seconds	300	a query running longer than this is "long"
check_mariadb_replica_lag_warn	300	seconds a replica is behind its primary
check_mariadb_replica_lag_crit	1800	seconds behind its primary -> critical
check_mariadb_galera_size	0	expected Galera / XtraDB Cluster size (0 = do not check the size)
check_mariadb_group_size	0	expected MySQL Group Replication members (0 = do not check the size)
check_mariadb_restart_warn_minutes	60	restarted less than this ago -> warning
check_mariadb_bufferpool_hit_warn	95	% of reads served from memory (only judged after 1M reads)
check_mariadb_error_since_days	1	[ERROR] lines in its log this recent
check_mariadb_disk_warn_pct	85	the filesystem holding the data directory
check_mariadb_disk_crit_pct	95	% used of the data directory's filesystem -> critical
check_mariadb_performance	true	---- PERFORMANCE (all warnings) ---------- false = only the health checks above. Counters that count since the server started (slow queries, temporary tables, joins, lock waits, deadlocks ...) are judged as an average since that start, and only once it has run chec... (all of it: the role file)
check_mariadb_perf_min_uptime_hours	2	judge the averages since the start only after this many hours
check_mariadb_slow_per_hour_warn	60	queries slower than long_query_time, per hour
check_mariadb_tmp_disk_pct_warn	25	% of temporary tables that went to disk (judged after 1000 of them)
check_mariadb_full_join_per_hour_warn	100	joins without a usable index (Select_full_join), per hour
check_mariadb_table_lock_wait_pct_warn	1	% of table locks (MyISAM / Aria) that had to wait (judged after 10000)
check_mariadb_row_lock_avg_ms_warn	500	average InnoDB row lock wait, ms (judged after 100 waits)

Setting	Default	What it does
check_mariadb_deadlocks_per_hour_warn	1	InnoDB deadlocks per hour (judged after 10 of them)
check_mariadb_innodb_waits_per_hour_warn	1	waits for a free buffer pool page / a full log buffer, per hour (after 10)
check_mariadb_aborted_connects_per_hour_warn	60	failed connection attempts per hour
check_mariadb_table_open_per_hour_warn	600	tables opened per hour while the table cache is full
check_mariadb_open_files_warn_pct	85	files open, % of open_files_limit
check_mariadb_lock_wait_seconds	30	a session waiting this long for a lock, right now
check_mariadb_long_trx_seconds	600	an InnoDB transaction open this long, right now
check_mariadb_history_length_warn	1000000	InnoDB history list length (old row versions not yet purged)

check_mounts

File: roles/check_mounts/defaults/main.yml

> **Note**
> MOUNTS: filesystems that should be mounted but are not, local filesystems that went read-only (the kernel does that after disk errors), and network mounts (NFS/CIFS) that hang. Diagnose only: this check never mounts, unmounts or remounts anything.

Setting	Default	What it does
check_mounts_network_types	[nfs, nfs4, cifs, smb3]	filesystem types counted as network mounts (checked for hangs and fullness)
check_mounts_local_types	[xfs, ext4, ext3, btrfs, vfat]	filesystem types checked for having gone read-only
check_mounts_timeout	5	seconds before a network mount counts as hung
check_mounts_warn_pct	85	Network mounts: how full (same idea as the disk check)
check_mounts_crit_pct	95	% used of a network mount -> critical (warning from check_mounts_warn_pct)
check_mounts_ignore	[]	fstab mount points that are allowed to be unmounted (regular expressions)

check_network

File: `roles/check_network/defaults/main.yml`

> **Note**
> NETWORK: can this host reach what it depends on? A default route, names that must resolve, and TCP ports that must answer (from THIS host, which is what matters when a service breaks).

Setting	Default	What it does
`check_network_require_default_route`	`true`	no default route -> finding (false for an isolated host)
`check_network_resolve_self`	`true`	this host's own full name (hostname -f) resolves
`check_network_dns_names`	`[]`	e.g. [ldap.example.mil, servicenow.example.mil]
`check_network_tcp`	`[]`	TCP services this host needs, e.g. check_network_tcp: - {name: LDAP, host: ldap.example.mil, port: 636} - {name: SIEM (syslog over TLS), host: siem.example.mil, port: 6514}
`check_network_timeout`	`5`	seconds for each name lookup and TCP port test

check_patching

File: `roles/check_patching/defaults/main.yml`

> **Note**
> PATCHING (readiness, read-only): pending security updates, a reboot still needed after the last updates, and how long ago the host was last updated. Installs nothing - patching itself is playbooks/patch_hosts.yml.

Setting	Default	What it does
`check_patching_max_age_days`	`35`	last package update older than this -> warning
`check_patching_security_warn`	`1`	this many pending security advisories -> warning
`check_patching_critical_is_critical`	`true`	a pending Critical-severity advisory -> critical
`check_patching_timeout`	`300`	seconds for dnf to talk to the repositories
`check_patching_reboot_check`	`true`	report a reboot still needed (needs-restarting -r, newer kernel installed than running)

check_performance

File: `roles/check_performance/defaults/main.yml`

> **Note**
> PERFORMANCE: is the host overloaded right now, or was something killed for lack of memory?

Setting	Default	What it does
`check_performance_load_warn`	`1.5`	15-minute load average PER CPU (1.0 = every CPU busy)
`check_performance_load_crit`	`3.0`	load per CPU -> critical
`check_performance_mem_avail_warn_pct`	`10`	memory still available for programs, % of total
`check_performance_mem_avail_crit_pct`	`5`	available memory below this % -> critical
`check_performance_swap_warn_pct`	`50`	swap in use, % of swap
`check_performance_swap_crit_pct`	`80`	swap in use above this % -> critical
`check_performance_iowait_warn`	`20`	% of CPU time waiting for disk (average over the sample)
`check_performance_iowait_crit`	`40`	I/O wait above this % -> critical
`check_performance_steal_warn`	`10`	% of CPU taken by the hypervisor (VM starved)
`check_performance_sample_seconds`	`5`	how long vmstat samples
`check_performance_oom_since`	`"-24h"`	look for out-of-memory kills this far back
`check_performance_zombie_warn`	`20`	this many zombie processes or more -> warning

check_selinux

File: `roles/check_selinux/defaults/main.yml`

> **Note**
> SELINUX: enforcing (the STIG requires it), the config file agrees (so a reboot does not change it), no locally added permissive domains, and recent denials (AVCs) - which is what you see when SELinux blocks an application.

Setting	Default	What it does
`check_selinux_expected_mode`	`enforcing`	the mode SELinux must be in (and the config file must say); anything else -> critical

Setting	Default	What it does
check_selinux_avc_since	yesterday	ausearch -ts: recent (10 min), today, yesterday, this-week
check_selinux_avc_warn	1	this many denials or more -> warning
check_selinux_booleans	{}	Booleans that must have a value, e.g. {httpd_can_network_connect: "on"}
check_selinux_relabel_paths	[]	Paths whose file labels are compared with the policy (a dry-run restorecon). Mislabeled files are a common cause of denials after files are copied or moved. Keep it to small trees. e.g. [/etc/nginx, /var/www, /etc/stigman]

check_services

File: roles/check_services/defaults/main.yml

> **Note**
> SERVICES: systemd units that failed, services that must be running (and start at boot), and services stuck restarting over and over.

Setting	Default	What it does
check_services_required	[sshd, crond, rsyslog, chronyd, auditd]	Must be running on every host. A host where one is not installed gets a warning (set check_services_required in the inventory to change the list for a group of hosts).
check_services_required_extra	[]	More for a group of hosts, e.g. in group_vars/mariadb_hosts.yml: [mariadb]
check_services_ignore_failed	[]	Failed units to ignore (regular expressions), e.g. a known-broken vendor unit.
check_services_restart_warn	3	More restarts than this since the unit started = a restart loop

check_time

File: roles/check_time/defaults/main.yml

> **Note**
> TIME: is the clock synchronized, and by how much is it off? A clock that is off by minutes breaks Kerberos/AD logins (5 minutes), TLS, log correlation and audit evidence.

Setting	Default	What it does
check_time_offset_warn	0.5	seconds

Setting	Default	What it does
check_time_offset_crit	5.0	seconds off -> critical (warning from check_time_offset_warn)
check_time_min_sources	1	time servers chrony must be able to reach

check_trust

File: roles/check_trust/defaults/main.yml

> **Note**
> TRUST (read-only): which trust stores on the host still hold a certificate authority that was compromised or that the browsers dropped - the system bundles, p11-kit's live view and blocklist, Java cacerts, Python's certifi bundles, Node's built-in roots, the bundles inside running podman containers, certificate files an application carries - and the software with its own trust logic (Go programs, Node, NSS) below the version that fixed it. The list of certificates to look for ships with the check (roles/check_trust/vars/main.yml: October 2026, the eight eMudhra intermediates and their roots); add your own in check_trust_extra. The health check runs it in the weekly set; playbooks/trust_report.yml runs it on every host, vCenter and ESXi too.

Setting	Default	What it does
check_trust_extra	[]	More certificates to look for, besides the shipped list: name, sha256 (the certificate's fingerprint, with or without colons), spki_sha256 (its public key, optional), kind (root / intermediate), severity (critical / warning / notice), why. Get the fingerprints with: openssl x509 -in the.pem -noout -fingerprint -sha256... (all of it: the role file)
check_trust_severity	{}	another severity for a listed name: {"AAA Certificate Services": warning} once you decide to blocklist it; notice = listed only, no finding
check_trust_software_severity	warning	the finding for a program below the fixed version (a Go program, Node, NSS): warning or critical
check_trust_bundles	[/etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem, /etc/pki/ca-trust/extracted/openssl/c…	The system: the CA bundles every OpenSSL / Go / Python client reads (the ones that exist are read), p11-kit's live view (trust extract: what NSS and GnuTLS clients see, and the blocklist), the blocklist folders. the system CA bundles (RHEL, Debian / Ubuntu, Alpine, SUSE); the same file under two names is read once

Setting	Default	What it does
check_trust_blocklist_dirs	[/etc/pki/ca-trust/source/blocklist, /usr/share/pki/ca-trust-source/blocklist]	p11-kit blocklist folders: a certificate there is distrusted (listed as blocked, no finding)
check_trust_p11kit	true	read p11-kit's live anchors and blocklist too (trust extract), when the trust command is there
check_trust_nss_dirs	[/etc/pki/nssdb, /root/.pki/nssdb, /home/*/.pki/nssdb]	NSS databases read with certutil (the system's: curl and other NSS clients on older RHEL; root's and each user's: Chrome, Evolution)
check_trust_ca_certificates_fixed	(empty)	the ca-certificates package version that dropped the roots, e.g. '2026.2.82' (empty: not compared - the bundles themselves tell)
check_trust_nss_fixed	'3.131.1'	the NSS version that dropped emSign ECC Root CA - G3 (the package nss / libnss3 below it: a finding; empty: not compared)
check_trust_java	true	Java: every cacerts / jssecacerts under these folders (JDKs, and the keystores of running java programs), read directly (JKS) or with keytool (PKCS12), with the JDK's version.
check_trust_java_dirs	[/usr/lib/jvm, /usr/java, /etc/pki/java, /etc/pki/ca-trust/extracted/java, /opt, /usr/loc…	folders searched for cacerts (seven levels deep)
check_trust_python	true	Python: every certifi bundle (certifi/cacert.pem: requests, pip, urllib3 ... carry their own roots, not the system's) under these folders - system Python, venvs, applications.
check_trust_python_dirs	[/usr/lib, /usr/lib64, /usr/local, /opt, /srv, /var/lib, /root, /home]	folders searched for certifi/cacert.pem (twelve levels deep)
check_trust_node	true	Node: its built-in roots (not the system's: Node carries its own, unless started with --use-openssl-ca), and its version against the one that dropped the root, per major line.
check_trust_node_paths	[/usr/local/bin/node, /usr/local/node*/bin/node, /opt/*/bin/node, /opt/node*/bin/node, /r…	where node binaries are (globs); the one in PATH and running ones are found too

Setting	Default	What it does
check_trust_node_fixed	{"22": "22.23.4", "24": "24.22.1", "26": "26.12.1"}	Node's fixed version per major line (an older line than the oldest listed: unsupported, no fix)
check_trust_go	true	Go: a Go program carries its own certificate verification. Go 1.27.3 / 1.26.10 reject any chain with one of the compromised keys (CVE-2026-111943) - a program built with an older Go has to be rebuilt. The running programs, and the executables under these folders, with the Go version each was built with.
check_trust_go_dirs	[/usr/local/bin, /usr/local/sbin, /opt, /srv]	folders searched for Go programs (four levels deep, executables over 1 MB)
check_trust_go_fixed	{"1.26": "1.26.10", "1.27": "1.27.3"}	Go's fixed version per line (a program built with an older line: unsupported - rebuild it with a fixed Go)
check_trust_go_max	200	at most this many programs looked at per host
check_trust_containers	true	Containers: the CA bundle inside every running podman container (root's and each logged-in user's), read with podman exec - a container image carries its own ca-certificates (Red Hat: UBI images are affected too).
check_trust_containers_max	30	at most this many containers per host
check_trust_find_dirs	[/etc, /opt, /srv, /usr/local, /var/lib, /root, /home]	Certificate files an application carries (its own bundle, a copied root): every *.pem, *.crt, *.cer, *.der, cacert*, ca-bundle*, ca-certificates*, and every Java keystore (*.jks, *.keystore, *truststore*; *.p12 / *.pfx with keytool and the default password) under these folders, on the host's own filesystems - only the... (all of it: the role file)
check_trust_find_max	2000	at most this many files read per host

Setting	Default	What it does
check_trust_find_exclude	'/(proc|sys|dev)/|/\.cache/|/containers/storage/|/docker/overlay2?/|/ca-trust/extracted/|…	regular expression on the path: matching files are skipped (container image layers, the copies update-ca-trust extracts, Debian's per-CA files - the bundles stand for them; this repository's own test certificates on the AAP controller)
check_trust_timeout	60	seconds for each search and each command (keytool, node, podman exec)
check_trust_remove	false	Removal (the Trust store report only, playbooks/trust_report.yml; the VMware trust store report has the same switch): take the listed certificates out of the stores that hold them. OFF by default - a read-only report. Check mode shows what would go and changes nothing. The system trust is blocklisted (RHEL / SUSE: a P… (all of it: the role file)
check_trust_remove_kinds	[system, java, python, file, nssdb]	which kinds of store to change: system (blocklist), java (every keystore), python (certifi), file (an application's bundle or certificate file), nssdb (certutil -D)
check_trust_remove_severities	[critical, warning]	only certificates at these severities go; a notice-level one (AAA Certificate Services) never, unless named here
check_trust_remove_backup_dir	/var/backups/site-automation/trust	every file changed is copied here first (a folder per run)
check_trust_remove_blocklist_dir	(empty)	the system blocklist folder; empty = the distribution's (/etc/pki/ca-trust/source/blocklist, SUSE /etc/pki/trust/blacklist)
check_trust_remove_update_cmd	(empty)	the command that rebuilds the system trust after blocklisting; empty = the distribution's (update-ca-trust, update-ca-certificates)
check_trust_keystore_password	changeit	the Java keystores' password (keytool's default)

horizon

File: roles/horizon/defaults/main.yml

Note
VMware Horizon 8 (playbooks/horizon_*.yml): the Connection Server's REST API, from the AAP side, with the "Horizon API" credential (its address, domain, user name and password) - no collection needed. Every job asks each API path at the newest version first, down to v1, so it fits any Horizon 8 release. A dry run (Check) reads and says what it would do; it changes nothing. docs/HORIZON.md

Setting	Default	What it does
`horizon_ca_path`	(empty)	---- Connection ---------- a CA file for the Connection Server's HTTPS certificate ("" = what the execution environment trusts)
`horizon_timeout`	`60`	seconds per request
`horizon_page_size`	`500`	records per page when reading pools, desktops and sessions (Horizon's most: 1000)
`horizon_api_newest`	`9`	the newest API version to try (v9 down to v1): the first that answers is used
`horizon_action_version`	`v1`	the API version of the actions (reset, recover, log off, push image...)
`horizon_protected_pools`	`[]`	---- Pools no job changes ---------- Desktop pools no job fixes, logs off or pushes an image to: names as Horizon Console shows them, * and ? allowed, any case (e.g. [VIP-*, Kiosk]). The reports still show them.
`horizon_cert_warn_days`	`30`	---- Health report (playbooks/horizon_health_report.yml) ---------- a Connection Server's, gateway's or vCenter's certificate expires within this many days -> amber
`horizon_cert_crit_days`	`7`	within this many days (or expired, or not valid) -> red
`horizon_license_warn_days`	`60`	the license expires within this many days -> amber
`horizon_report_max_rows`	`50`	rows in the problems table (the job's artifacts have all)
`horizon_health_fail`	`false`	The health and capacity reports stay green when they find problems: the report says what. true = mark the job FAILED on a critical finding, for workflows and alerts.
`horizon_capacity_warn_pct`	`80`	---- Capacity report (playbooks/horizon_capacity_report.yml) ---------- a pool's desktops in use (connected or disconnected) against its maximum, % -> amber

Setting	Default	What it does
horizon_capacity_crit_pct	95	in use, % -> red (also red: no free desktop left)
horizon_broken_states	[AGENT_UNREACHABLE, AGENT_ERR_DISABLED, AGENT_ERR_INVALID_IP, AGENT_ERR_NEED_REBOOT, AGEN…	---- Broken desktops (the desktops check -> approval -> the fix job) ---------- A desktop in one of these states takes no user until it is fixed. Left out on purpose: states that pass by themselves (being made, customized, deleted, waiting for the agent at start-up, maintenance).
horizon_desktops_mode	report	report = the check lists them and stays green (a schedule). approval = when there is a desktop to fix, the check hands the list on and FAILS on purpose, so the workflow takes its "Run on fail" path to the approval step and then the fix job. Set approval on the workflow's check node.
horizon_fix_default_action	recover	what fixes a broken desktop: recover (instant clones: deleted and made again from the image), reset, restart, remove, rebuild
horizon_fix_actions	{AGENT_ERR_NEED_REBOOT: restart}	By state, when another action fits better (e.g. {AGENT_ERR_NEED_REBOOT: restart, PROVISIONING_ERROR: remove}).
horizon_fix_max	20	fix at most this many desktops per run: the rest are listed, the next run takes them
horizon_fix_exclude_machines	[]	desktops never fixed automatically: names, * and ? allowed (e.g. [W11-TEST-*])
horizon_session_max_disconnected_hours	8	---- Disconnected sessions (playbooks/horizon_session_cleanup.yml) ---------- log off a session disconnected for this many hours or more
horizon_session_cleanup_max	100	log off at most this many sessions per run
horizon_session_cleanup_exclude_pools	[]	pools whose sessions are never logged off: names, * and ? allowed
horizon_session_cleanup_exclude_users	[]	users never logged off: login names (jdoe or CORP\jdoe) or SIDs
horizon_session_forced_logoff	true	log off a locked session too (a disconnected session is usually locked)
horizon_session_cleanup_report_only	false	true = list them, log off none (a dry run - Check - does the same)

Setting	Default	What it does
horizon_user	(empty)	---- User help (playbooks/horizon_user_help.yml) ---------- the user who cannot connect: jdoe, CORP\jdoe, jdoe@corp.example.mil or a SID (a survey question)
horizon_user_action	look	What to do: look (only show their sessions and desktops), logoff (all their sessions), disconnect (their connected sessions), reset or restart (their desktops), recover (their instant-clone desktops: made again from the image). Never a desktop another user has a session on, nor one in a protected pool.
horizon_push_pools	[]	---- Image push (the push plan -> approval -> the push job) ---------- the instant-clone pools to push to: names (a survey: one per line, or separated by commas)
horizon_push_base_vm	(empty)	the golden image VM, as vCenter names it (or its path)
horizon_push_snapshot	(empty)	its snapshot to push: the name, or the path when two have that name (e.g. /2026-10/2026-11)
horizon_push_start	(empty)	When the push starts: "" = as soon as the approved push job runs; or a time - "2026-11-01 22:00" (UTC) or with an offset, "2026-11-01T22:00-05:00". A time that has passed by the approval starts in a minute.
horizon_push_logoff_policy	WAIT_FOR_LOGOFF	WAIT_FOR_LOGOFF (each desktop is replaced when its user logs off) or FORCE_LOGOFF (users are logged off)
horizon_push_stop_on_first_error	true	stop the push at the first desktop that fails
horizon_test_user	(empty)	---- API test (playbooks/horizon_api_test.yml) ---------- also look up this user (e.g. jdoe), as the user-help job does

net_certs

File: roles/net_certs/defaults/main.yml

Note
NETWORK DEVICE CERTIFICATES (playbooks/net_cert_report.yml, net_cert_renew.yml): the TLS certificate each switch, firewall, printer, iLO / iDRAC, UPS, storage … presents, read from the controller with a TLS hello - no login to the devices, nothing changed. The renewal: a ServiceNow incident per certificate due; on a Cisco IOS / IOS-XE device, optionally a CSR made on the device (net_cert_cisco_csr).

Setting	Default	What it does
net_cert_targets	inventory	---- which devices ---- inventory (the hosts of net_cert_group) / subnets (net_cert_subnets) / both
net_cert_group	net_cert_devices	the inventory group of the devices; a host's own net_cert_ports / net_cert_address / net_cert_servername override
net_cert_ports	[443]	the ports read (a host's own net_cert_ports wins), e.g. [443, 8443, 636]
net_cert_subnets	[]	CIDR ranges to sweep (net_cert_targets: subnets or both), e.g. ['192.0.2.0/24'] - every address, each port in net_cert_ports
net_cert_exclude	[]	never connected to: addresses, CIDR ranges, host name patterns (e.g. ['192.0.2.1', 'ups*'])
net_cert_max_addresses	4096	a sweep bigger than this (all subnets together) is refused: a /16 by mistake would take hours
net_cert_resolve_names	true	a swept address that answers: its DNS name (PTR) in the report, else the address
net_cert_timeout	5	seconds for each connection, and each handshake
net_cert_workers	32	connections at the same time
net_cert_only	[]	only these devices (names, wildcards), e.g. ['sw01*'] - to try the renewal on one
net_cert_warn_days	30	---- what is due ---- a certificate expiring within this many days is due: a warning incident (amber in the report)
net_cert_crit_days	7	within this many days (or expired): critical

Setting	Default	What it does
net_cert_notice_days	60	the report colours these blue (a heads-up, no incident)
net_cert_unanswered_finding	false	true = an inventory device whose port gives no certificate gets an incident too (a warning)
net_cert_cisco_csr	false	On: for a due certificate of an inventory host with ansible_network_os cisco.ios.ios (and its network_cli credential on the job template), the job makes a key on the device (crypto key generate rsa, not exportable), a trustpoint <prefix>-<expiry> (enrollment terminal) and its CSR (crypto pki enroll), and attaches the... (all of it: the role file)
net_cert_cisco_key_bits	2048	the new RSA key (2048, 3072, 4096)
net_cert_cisco_trustpoint_prefix	SA-RENEW	the trustpoint and key label: SA-RENEW-20270131 (the old certificate's expiry)

net_config

File: roles/net_config/defaults/main.yml

> **Note**
> NETWORK DEVICE CONFIGURATIONS: the backup (playbooks/net_config_backup.yml) keeps each device's running configuration - in a git repository (a commit when one changed) or on a file share (dated copies) - and can open a ServiceNow incident when one changed. The load (playbooks/net_config_load.yml) sends configuration lines, or restores a backup, after an approved change request - always a dry run (Check) first. Cisco IOS / IOS-XE (over its CLI: network_cli, a Machine credential) and Palo Alto Networks PAN-OS (over its XML API from the controller: the "Palo Alto PAN-OS API" credential). docs/NETWORK_CONFIG.md. A backup holds the device's secrets (password hashes, SNMP communities, keys): keep it in a private, internal repository or share only - never in this one.

Setting	Default	What it does
net_config_group	net_config_devices	The inventory group of the devices. Each one's platform: ansible_network_os (cisco.ios.ios, paloaltonetworks.panos), or net_config_platform (ios, panos) on the host.
net_config_only	[]	Only these devices (names, * and ? allowed), e.g. [rtr-lab*] - to try a job on one first. [] = every one.

Setting	Default	What it does
net_config_store	git	git (a repository), share (a folder on a file server or the AAP node), or both.
net_config_git_repo	(empty)	The repository (ssh://git@git.example.mil/netops/configs.git or https://...). Private and internal: the configurations hold secrets. Its login: the "Network config git repository" credential (an SSH key, or a username and token).
net_config_git_branch	main	the branch (made when the repository is new and empty)
net_config_git_path	configs	the folder in the repository: <path>/<device>.cfg (IOS), .xml (PAN-OS)
net_config_git_author_name	site-automation	the commits' author
net_config_git_author_email	site-automation@localhost	and its email
net_config_git_host_key_checking	accept-new	The git server's SSH host key: accept-new (the first connection learns it), or yes (only the keys in net_config_git_known_hosts - lines as in a known_hosts file). In an execution environment ssh's own known_hosts is empty every job, so accept-new learns the key each time: only net_config_git_known_hosts + yes really c... (all of it: the role file)
net_config_git_known_hosts	(empty)	the git server's host keys (known_hosts lines), for net_config_git_host_key_checking: yes
net_config_share_dir	/var/lib/site-automation/net-config	The folder for share: on the AAP node (a path AAP exposes to jobs) or on net_config_share_host - <dir>/<device>/<device>.cfg (the latest), and a dated copy beside it each time one changed.
net_config_share_host	(empty)	an inventory host to keep them on (Linux, or a Windows file server); its credential on the template
net_config_share_keep	30	dated copies kept per device (the oldest go)
net_config_ios_command	show running-config	---- what is read ---- IOS: the command whose output is kept
net_config_panos_port	443	PAN-OS: the XML API's HTTPS port

Setting	Default	What it does
net_config_panos_validate_certs	true	PAN-OS: check the firewall's certificate
net_config_panos_ca_path	(empty)	PAN-OS: a CA bundle (PEM) for that check ('' = the system's)
net_config_timeout	60	PAN-OS: seconds per request
net_config_diff_lines	60	lines of a diff shown in the report, the email and an incident (secrets hidden)
net_config_fail_on_failed	true	the backup job fails (red) when a device could not be backed up
net_config_tickets	false	true = an incident when a device's configuration changed (one per change, with the diff) and when one could not be backed up (noted as cleared at the next good backup). The "ServiceNow API" credential.
net_config_ticket_changes	true	an incident for a change
net_config_ticket_failures	true	an incident for a failed backup
net_config_load_mode	lines	lines (send configuration lines) or restore (put a device back to a backup).
net_config_lines	[]	IOS: the lines, as typed in configuration mode, e.g. ["ntp server 192.0.2.123", "banner motd ^Authorized use only^"]; under net_config_parents (e.g. ["interface GigabitEthernet1"]) when they belong to one.
net_config_parents	[]	IOS: the parent lines the lines go under (e.g. ["interface GigabitEthernet1"]); [] = global configuration
net_config_panos_entries	[]	PAN-OS: the changes - each {xpath, element, action}: action set (adds or replaces the element's children), edit (replaces the element at the xpath), delete (no element).
net_config_restore_ref	(empty)	restore: which backup - git: a commit, tag or branch ('' = the newest); share: a dated copy's file name ('' = the latest).
net_config_change	(empty)	The change request the load belongs to (CHG0012345). Its state must be one of net_config_change_states; the job writes on it what it did. A dry run without one only warns.

Setting	Default	What it does
net_config_require_change	true	false = a real load without a change request (not recommended)
net_config_change_states	[Implement]	the change states a load may run in (as ServiceNow shows them)
net_config_change_note	true	a work note on the change: each device, what was sent, the result
net_config_save	true	IOS: write memory after a change (false: only the running configuration - a reload undoes it)
net_config_ios_file	flash:site-automation-restore.cfg	IOS restore: where the backup is copied first (scp; ip scp server enable on the device)
net_config_panos_allow_pending	false	PAN-OS: change even with someone's uncommitted changes (they would be committed too)
net_config_max_devices	10	a load to more devices at once is refused (net_config_only, the template's Limit - or raise it)

ontap_report

File: roles/ontap_report/defaults/main.yml

> **Note**
> NetApp ONTAP health report (playbooks/ontap_health_report.yml). Read-only: GET requests to each cluster's REST API, with the account of a "NetApp ONTAP" credential (docs/NETAPP.md).

Setting	Default	What it does
ontap_clusters	[]	The clusters' management addresses (names or IPs), e.g. [cluster1.yoursite.mil, cluster2.yoursite.mil]. The credential's account must exist on each of them.
ontap_validate_certs	true	Check the clusters' HTTPS certificates. ONTAP's own certificate is often self-signed: then put its CA in ontap_ca_path (a file in the project, e.g. "{{ playbook_dir }}/files/ca/ontap-ca.pem"), or set false.
ontap_ca_path	(empty)	the CA file (PEM) the clusters' certificates are checked with; empty = the execution environment's own trust
ontap_timeout	30	seconds for one request (no answer: asked once more)
ontap_tls_legacy	false	true = connect to older ONTAP whose TLS the execution environment refuses (UNSAFE_LEGACY_RENEGOTIATION_DISABLED, KEY_TOO_SMALL, SHA-1)

Setting	Default	What it does
ontap_volume_warn_pct	85	Thresholds (used %, hours, days, seconds): at or above "warn" = amber, at or above "crit" = red.
ontap_volume_crit_pct	90	volume used % -> red
ontap_aggr_warn_pct	85	aggregate used % -> amber
ontap_aggr_crit_pct	90	aggregate used % -> red
ontap_inode_warn_pct	85	inodes used % of a volume -> amber
ontap_inode_crit_pct	90	inodes used % -> red
ontap_snapmirror_lag_warn_hours	24	last good transfer longer ago than this
ontap_snapmirror_lag_crit_hours	48	last good transfer longer ago than this -> red
ontap_snapshot_reserve_warn_pct	100	snapshots using this much of their reserve (100 = they spill into the volume)
ontap_snapshot_reserve_crit_pct	150	snapshots using this much of their reserve -> red
ontap_cert_warn_days	60	certificates expiring within this many days
ontap_cert_crit_days	14	certificates expiring within this many days (or expired) -> red
ontap_time_drift_warn_seconds	60	a node's clock against the AAP server's
ontap_time_drift_crit_seconds	300	a node's clock this far off -> red
ontap_port_errors_warn	1	receive errors (CRC and other bad frames) since the node booted
ontap_port_errors_crit	10000	receive errors on a port since boot -> red
ontap_min_spares	1	fewer spare disks than this on a node = amber (0 = red)
ontap_ems_hours	24	the error events of the last this-many hours
ontap_ems_max	500	read at most this many events per cluster
ontap_report_max_rows	50	rows per table (problems first); the job's artifacts have all
ontap_snapshot_top	25	the snapshot space table: the volumes whose snapshots use the most
ontap_report_fail	false	The job stays green when it finds problems: the report and email say what. true = mark it FAILED when there is a critical finding or a cluster could not be read, for workflows.

patch

File: roles/patch/defaults/main.yml

Note
PATCH: install updates with dnf, reboot if the updates need it, and prove the host came back: every service that was running before must be running after. One host at a time by default; the rollout stops at the first host that fails.

Setting	Default	What it does
patch_packages	['*']	what to update ('*' = everything that has an update)
patch_security_only	false	true = only updates that fix security advisories
patch_exclude	[]	never update these here, e.g. ['mariadb*', 'kernel*']
automatic_restarts	false	NO AUTOMATIC RESTARTS unless you say so: this playbook reboots a host only when automatic_restarts is true (set it in the job template's Variables, or answer the survey), and then only as patch_reboot says. With automatic_restarts false (the default) it installs the updates and tells you which hosts still need a reboo... (all of it: the role file)
patch_reboot	when_needed	Reboot (only with automatic_restarts: true): when_needed (needs-restarting says so) | never | always
patch_reboot_timeout	1800	seconds to wait for the host to come back
patch_never_reboot_groups	[aap, aap_hosts, sat, netapp, netapp_console_hosts]	Hosts in these groups are patched but NEVER rebooted by this playbook (reboot them by hand, in a maintenance window). Add your own group names; the built-in protection below is on always.
patch_never_patch_groups	[no_patch, aap, aap_hosts, sat, netapp, netapp_console_hosts]	Hosts in these groups are skipped entirely (Satellite is upgraded with satellite-maintain, not dnf).

Setting	Default	What it does
patch_min_free_mb	1024	BUILT IN, whatever you set above (roles/patch/vars/main.yml): the AAP hosts (groups aap and aap_hosts) are never patched and never rebooted by a job, and neither is the machine the job runs from (a local connection). A package script can restart services under the running job. Set the two lists above only to ADD group... (all of it: the role file)
patch_service_wait	120	seconds to wait for services to come back after a reboot
patch_ignore_services	'^(user@|user-runtime-dir@|session-|dnf-makecache|packagekit|fwupd|run-|systemd-(hostname…	Services that come and go on their own and are not compared before/after (regular expression)

poam

File: roles/poam/defaults/main.yml

Setting	Default	What it does
poam_csv	"{{ playbook_dir }}/../poam/poam.csv"	POA&M STATUS: which Plan of Action & Milestones items are overdue or due soon, and (optional) which open CAT I / CAT II findings in STIG Manager have no POA&M item yet. The POA&M list is a CSV file in your work Git (export from eMASS or your spreadsheet, keep the columns below, save as CSV). POA&M data is sensitive: k... (all of it: the role file)
poam_columns	id: POAM ID # unique item number weakness: Weakness # what is wrong severity: Severity #…	Column names in YOUR file for each field this check needs (change the right-hand side).
poam_closed_statuses	[Completed, Closed, Risk Accepted, Cancelled]	items with one of these statuses are done: never reported
poam_due_soon_days	30	an open item due within this many days -> warning (past due -> critical)
poam_date_formats	['%Y-%m-%d', '%m/%d/%Y', '%m/%d/%y', '%d-%b-%Y', '%d %b %Y', '%b %d, %Y']	the date formats your file may use (Python strptime), tried in this order

Setting	Default	What it does
poam_stigman_api	(empty)	---- STIG Manager cross-check (optional) ---------- Needs a Keycloak client that may use the client-credentials grant, with the scope stig-manager:collection:read, and a grant in the collections to read. Your realm owner creates it - this repository never changes the realm. T... (all of it: the role file)
poam_stigman_collections	[]	collection names; empty = every collection the client can read
poam_stigman_severities	[high, medium]	CAT I and CAT II findings must have a POA&M item
poam_stigman_scope	(empty)	Scope to ask for. Empty = the client's default scopes (they must include stig-manager:collection:read or stig-manager:collection).
poam_validate_certs	true	check STIG Manager's and Keycloak's HTTPS certificates

podman_discover

File: roles/podman_discover/defaults/main.yml

> **Note**
> PODMAN DISCOVERY (read-only): every podman container on the host, root's (rootful) and each user's (rootless), with how it is run (a systemd unit or plain podman) and whether it should be running. Used by the Health check containers and by Service watch. It changes nothing.

Setting	Default	What it does
podman_discover_enabled	true	false = do not look for containers at all (the containers check then says "skipped", and Service watch watches only its watch_containers list).
podman_discover_rootless	true	Also look at the containers of users (rootless podman). A user is looked at when a podman process (conmon) runs as them, when they have "linger" on (/var/lib/systemd/linger/NAME), or when they are in /etc/subuid AND have container storage (~/.local/share/containers/storage). Their containers are read AS that user (run... (all of it: the role file)
podman_discover_users	[]	More users to look at (names), e.g. an application account the rules above miss.

Setting	Default	What it does
podman_discover_ignore	[]	Containers to leave out, as regular expressions that must match the WHOLE container name or OWNER/NAME, e.g. ['test-.*'] (every container whose name starts with test-, any owner), ['alice/.*'] (all of alice's containers), ['root/web'] (root's web only).
podman_discover_quadlet_dirs	[/etc/containers/systemd, /usr/share/containers/systemd]	Where root's Quadlet files are (*.container, and *.kube: a pod from Kubernetes YAML). A user's are always read from ~/.config/containers/systemd, /etc/containers/systemd/users/UID and /etc/containers/systemd/users.
podman_discover_unit_dirs	[/etc/systemd/system, /usr/local/lib/systemd/system]	Where root's own systemd units are looked for: a *.service that runs podman (podman generate systemd, podman run, podman kube play, podman pod start, podman-compose / docker-compose), and the podman-kube@ units enabled there (.wants links). A user's: ~/.config/systemd/user. Packaged units (/usr/lib/systemd/sys... (all of it: the role file)
podman_discover_timeout	60	Seconds one podman or systemctl call may take before it is given up (reported as a finding).

service_watch

File: roles/service_watch/defaults/main.yml

Setting	Default	What it does
watch_discover	true	---- What to watch ---------- With watch_discover on (the default), service watch finds the containers itself: every podman container on the host that SHOULD be running - root's (rootful) and each user's (rootless) - as the Health check containers finds th... (all of it: the role file)
watch_containers	[]	Containers to watch in any case, in dependency order (a container's dependencies first). They are watched before the discovered ones, and one that is missing is reported as down. name: the container name, as podman ps -a shows it (or its Quadlet / service name) user: optional: the user that owns it (rootless podman)... (all of it: the role file)
watch_url_ok_status	[200, 301, 302, 401, 403]	HTTP codes that count as "answering"

Setting	Default	What it does
watch_find_units	true	---- How each container is run (found by itself) ---------- A container that a systemd unit runs (a Quadlet .container file, podman generate systemd, or a unit you wrote) must be started with systemctl, never with podman start: the unit owns it, and while the unit is stopped the co... (all of it: the role file)
watch_quadlet_dirs	[/etc/containers/systemd, /usr/share/containers/systemd]	where root's Quadlet files are looked for (to find the unit that runs a container)
watch_mode	approval	---- What happens when one is down ---------- approval: ACT only diagnoses. Its proposed fix waits for the AAP approval step, then service_fix_approved.yml runs exactly the approved command(s) - no model involved. When ACT proposes nothing (or cannot run), the standard fix... (all of it: the role file)
watch_verify_retries	6	re-checks after a fix (every watch_verify_delay seconds)
watch_verify_delay	5	seconds between those re-checks
watch_record_file	/var/log/service-watch/incidents.jsonl	---- Where incidents are recorded (on the host) ---------- One JSON line per event, plus a syslog line (tag service-watch) that your log forwarding (Splunk, Elastic, rsyslog relay) picks up. AAP's job history keeps its own copy.
watch_record_syslog	true	also send each incident to syslog (tag service-watch)
watch_act_env	{}	---- ACT settings ---------- Which provider and model: site_act_provider / site_act_model / site_act_url / site_act_ca, the same settings every runbook uses (roles/site_act/defaults/main.yml). The key comes from the "ACT model key" AAP credential, never f... (all of it: the role file)
watch_act_timeout	1200	seconds for one ACT run (the model's time per answer: site_act_model_timeout)
watch_instructions	(worked out by the role; see its file)	What ACT is told, on top of the act_triage role's standard instructions.

servicenow

File: roles/servicenow/defaults/main.yml

Note
SERVICENOW - open, update (and optionally resolve) incidents from check findings, and check that ServiceNow itself (the instance and its MID Servers) is healthy. Uses the REST Table API with ansible.builtin.uri: no extra collection to install. Connection: the "ServiceNow API" AAP credential sets SN_HOST (https://<instance>), SN_USERNAME and SN_PASSWORD in the job's environment. Without it, nothing is sent: the job prints what it WOULD have opened (a safe way to try it out).

Setting	Default	What it does
`servicenow_min_severity`	`warning`	Findings at or above this severity get a ticket: critical | warning. A ticket already open stays open while its finding is still there, even below this (the certificate renewal opens one for every CSR)
`servicenow_table`	`incident`	the table tickets are opened in (incident)
`servicenow_assignment_group`	(empty)	Values for new incidents. Names work (the API is asked to accept display values), e.g. servicenow_assignment_group: Linux Operations
`servicenow_caller`	(empty)	user name; empty = the API account
`servicenow_category`	(empty)	e.g. Software
`servicenow_urgency`	`{critical: "1", warning: "2"}`	a new ticket's urgency, per finding severity (your instance's values)
`servicenow_impact`	`{critical: "2", warning: "3"}`	a new ticket's impact, per finding severity
`servicenow_set_ci`	`true`	set Configuration item = the host name (must exist in the CMDB)
`servicenow_fields`	`{}`	More fields for every new ticket - for an instance that makes fields such as Location, Subcategory or Channel mandatory: {column name: value}. Names and choice labels work, e.g. servicenow_fields: {location: "Building 100", subcategory: "Operating System", contact_type: "Self-service"} A field's column name: right-cli... (all of it: the role file)
`servicenow_note_cleared`	`true`	When a problem that had a ticket is gone, add a work note (and, if you want, resolve it).
`servicenow_resolve_cleared`	`false`	true = also resolve the ticket (servicenow_resolved_state, servicenow_close_code)
`servicenow_resolved_state`	`"6"`	Resolved

Setting	Default	What it does
servicenow_close_code	"Solution provided"	must be a close code your instance has
servicenow_resolve_fields	{}	More fields set when a ticket is resolved (servicenow_resolve_cleared, and the test ticket), for an instance that makes them mandatory: {column name: value}, e.g. {resolved_by: svc_aap_sn}
servicenow_note_repeats	true	Add a work note to an open ticket each time its finding is seen again.
servicenow_correlation_prefix	site-health	Every ticket from these playbooks carries a correlation ID starting with this, so the same problem on the same host always updates the same ticket.
servicenow_validate_certs	true	check the instance's HTTPS certificate (see servicenow_ca_path)
servicenow_ca_path	(empty)	A CA bundle (PEM) for the instance's TLS certificate, when the execution environment does not trust its CA (e.g. a DoD root CA): a path on the controller, such as a file you add to the project: servicenow_ca_path: "{{ playbook_dir }}/files/ca/servicenow-ca.pem". Empty = the execution environment's normal trust store.... (all of it: the role file)
servicenow_timeout	30	seconds for one API request
servicenow_fail_when_findings	true	Fail the job when it opened/updated tickets, so the workflow shows red while problems exist.
servicenow_check_mid	true	servicenow_health.yml: also check the MID Servers (needs read access to the ecc_agent table)
servicenow_slow_seconds	5	the API answering slower than this -> warning
servicenow_test_resolve	true	servicenow_test_ticket.yml: resolve the test incident at the end (false = leave it open to look at)
cert_renew_servicenow	incident	How the CSRs the certificate renewals make (Linux, Windows, VMware, network devices) go to the PKI. incident: each one is attached to its certificate's incident, with a work note saying what to do next - a person submits it to the PKI. catalog: each one is ordered from your PKI's catalog item (cert_renew_catalog_item)... (all of it: the role file)

Setting	Default	What it does
`cert_renew_csr_memory_days`	`90`	incident: a CSR on an incident of the certificate made in the last this-many days counts as made - also when that incident was resolved or closed since (the PKI may still be signing it): vCenter, ESXi and Cisco keep one pending key, and a second CSR would replace it. Older: a new CSR (the certificate is due again).
`cert_renew_catalog_item`	(empty)	The catalog item: its name exactly as the catalog shows it (e.g. "DoD PKI Certificate Request"), or its sys_id.
`cert_renew_catalog_variables`	`{}`	The item's questions - by their internal names, as the catalog test lists them - and the answers: {name: answer}. A reference (Requested For, Location ...): the record's name, or its sys_id. A choice: its value or its label. In an answer, {host}, {subject}, {names}, {expires}, {issuer} and {where} become the certifica... (all of it: the role file)
`cert_renew_catalog_csr_variable`	(empty)	The attachment question that gets the CSR (its internal name). Empty = the item's only attachment question.
`cert_renew_catalog_names_variable`	(empty)	A question that gets the certificate's names (subjectAltName), e.g. a "List any additional SANs" box. Empty = none.
`cert_renew_catalog_reorder_states`	`["4", "7"]`	A request closed in one of these states (a requested item's: 4 Closed Incomplete, 7 Closed Skipped - rejected or cancelled) is ordered again on the next run. In any other state - open, in progress, Closed Complete (the PKI issued it: install it) - the certificate is not ordered again.
`cert_renew_catalog_upload_table`	`ZZ_YYsc_cart_item`	Where the CSR is uploaded before the order, for the attachment question (ServiceNow moves it to the request).
`servicenow_catalog_test_order`	`false`	servicenow_catalog_test.yml: true = also order the item once, with a test CSR - the PKI gets a real request (tell them first, or cancel it). false = only read the item and check the settings.

Setting	Default	What it does
cert_install_attachment_exts	[cer, crt, pem, der, p7b, p7c, p7, cert, zip, txt]	---- Certificate install: the issued certificate, from its request (cert_install_check.yml, win_cert_install_check.yml) ---- The PKI puts the issued certificate on the request (RITM) that carries the CSR: attached to the RITM, to one of its attachment questions, to its REQ or to its catalog task (SCTASK) - one file or... (all of it: the role file)
cert_install_max_bytes	1048576	a bigger attachment is not read
cert_install_journal	true	also look for a certificate pasted into the request's work notes and comments
cert_install_min_days	30	the issued certificate must be valid at least this many more days (else held for a person)
cert_install_fewer_names	false	true = install a certificate that lacks names (subjectAltName) the old one had; false = held for a person

site_act

File: roles/site_act/defaults/main.yml

Setting	Default	What it does
site_act_level	explain	ACT (GenAI) on top of ANY check playbook. The check roles already did the cheap, certain part: they found WHAT is wrong (site_findings) and know which read-only command shows more (each finding's hint). This role hands that to ACT (the act_triage role) to find out WHY, and to propose - or, if you allow it, make - th... (all of it: the role file)
site_act_allow	[]	self-heal only: regular expressions of fix commands ACT may run without a person. Each must match the WHOLE command. Keep them narrow, e.g. site_act_allow: ['systemctl restart (chronyd|rsyslog|crond)']
site_act_diagnose_only_groups	[aap, aap_hosts, sat, idm, netapp, netapp_console_hosts]	Hosts where ACT may only diagnose - never fix by itself, whatever site_act_level says: AAP itself (a fix could stop the job that is running it) and vendor-managed appliances. The apply step (act_fix_approved.yml) skips them too: a person applies the fix there by hand. AAP (groups aap, aap_hosts) is ALWAYS in this list... (all of it: the role file)

Setting	Default	What it does
site_act_provider	genai	---- Which model ACT talks to ---------- site_act_provider genai GenAI.mil (the default) asksage Ask Sage (set site_act_url to your organization's instance) genai-beta the GenAI.mil beta proxy (preview models) site_act_models the model to use FOR EACH provider, e.... (all of it: the role file)
site_act_models	{genai: gemini-3.8-flash, asksage: gpt-5.6-sol-gov}	the model for each provider (see above): the model follows the provider picked
site_act_urls	{}	the API URL for each provider; a provider not listed uses ACT's default
site_act_model	(empty)	one model for this run, whatever the provider (e.g. a survey answer); empty = site_act_models
site_act_url	(empty)	one URL for this run; empty = site_act_urls
site_act_ca	(empty)	a CA bundle file ON THE HOSTS for the provider's certificate; empty = the host's trust store
site_act_env	{}	Any other ACT environment settings (advanced; these win over the ones above), e.g. {HTTPS_PROXY: "http://proxy.yoursite.mil:8080"}. See the top of vendor/act/act for the list. Every ACT job streams the model's answer (ACT_STREAM=1), so a firewall or proxy idle timeout does not cut a long answer ("connection reset by p... (all of it: the role file)
site_act_model_timeout	600	Seconds the model may take for ONE answer, retries after a provider hiccup included (ACT's GENAI_TIMEOUT; ACT's own default, 90, is too short for a long analysis or a busy provider). Every ACT job uses it: the runbooks' use_act, the VMware alarms and capacity analyses, service watch.
site_act_timeout	1200	seconds for one whole ACT run (several answers): keep it above site_act_model_timeout
site_act_concurrency	3	How many hosts may run ACT at the same time in one job (0 = no limit: as many as the job's forks). Every host in a job uses the same key, and GenAI.mil's default quota is 60 requests and 200,000 tokens a minute per key: a fleet run with many hosts at once gets HTTP 429 (too many requests).
site_act_extra_instructions	(empty)	anything ACT should know about this site

Setting	Default	What it does
site_act_evidence_text	(empty)	extra evidence (troubleshoot.yml passes what it collected)

site_email

File: roles/site_email/defaults/main.yml

> **Note**
> EMAIL A REPORT through your mail relay (SMTP). Off until report_email_to has an address. Used by every report job: health checks, certificate reports, database health, POA&M, ServiceNow health, troubleshooting, the VMware jobs (docs/EMAIL_REPORTS.md). The relay settings belong in playbooks/group_vars/all.yml. If the relay needs a login, attach an "SMTP relay" credential (aap/credential_types/smtp_relay.yml) - never put a password in these settings.

Setting	Default	What it does
report_email_to	[]	Who gets it: a list, or text with commas / spaces / one per line (a survey answer). Empty = no email. Set it on each TEMPLATE (its Variables, or a survey): in all.yml every report job emails.
report_email_cc	[]	copied: a list, or text like report_email_to
report_email_also	[]	"Also email the report to": more people for one run, asked when the job is launched - a survey question (playbooks/aap_email_question.yml adds it to the report templates). Added to report_email_to, never in its place; empty = only report_email_to. A list, or text: addresses separated by commas, semicolons, spaces or o... (all of it: the role file)
report_email_also_domains	[]	The domains an address in report_email_also may be in (a subdomain too), e.g. [yoursite.mil]; any other is listed in the job and gets no email. Empty = any domain (your mail relay may still refuse some).
report_email_from	(empty)	the sender, e.g. aap-noreply@yoursite.mil
report_email_smtp_host	(empty)	your mail relay, e.g. smtp.yoursite.mil
report_email_smtp_port	25	25 or 587 with starttls, 465 with ssl

Setting	Default	What it does
report_email_security	starttls	starttls = encrypt after connecting (the relay must offer it); ssl = encrypted from the start (port 465); none = not encrypted (only for a relay without encryption; then no login is sent).
report_email_ca_path	(empty)	a CA file that signed the relay's certificate (empty = the CAs the execution environment trusts)
report_email_subject_prefix	"[AAP]"	put before every subject, e.g. "[AAP] Certificate report: ..."
report_email_html	true	Send a formatted (HTML) version too, for reports that come with one (all VMware jobs). The email always carries a plain-text copy for mail clients that do not show HTML. false = plain text only.
report_email_theme	standard	The look of the HTML version (docs/EMAIL_REPORTS.md, "The look"): standard = a header band in the report's colour (green, amber, red, blue) with number boxes under it; command = a dark header with your logo and report_email_header_text, the job's last runs as a strip of bars (report_history), the numbers in one row; p... (all of it: the role file)
report_email_header_text	Site automation report	command, portal, glass: the line with the logo
report_logo_file	(empty)	Your logo in the header (every look but standard): a PNG, JPEG or GIF in your repository - a path from playbooks/ (files/report_logo.png = playbooks/files/report_logo.png) or a full path. Attached to each email, so keep it small: about 160 px tall, under 100 KB. Updating with scripts/update-from-release: list the file... (all of it: the role file)
report_logo_height	84	pixels in the email (its width follows); portal: at most 56, glass: at most 80
report_history	true	command, portal, glass: the job's last runs, read from AAP - attach a "Red Hat Ansible Automation Platform" credential to the template, whose user can see the template's jobs (Organization Auditor on its organization, docs/EMAIL_REPORTS.md); without it the email goes out without the strip.

Setting	Default	What it does
report_history_runs	14	bars in the strip, this run included
report_history_aap_ca_path	(empty)	a CA file for AAP's own HTTPS certificate ("" = what the execution environment trusts)
report_email_html_template	report.html.j2	Your own look instead: copy roles/site_email/templates/report.html.j2 (or report_command / _portal / _glass) to e.g. playbooks/files/email/report.html.j2, change it, and set report_email_html_template: "{{ playbook_dir }}/files/email/report.html.j2" report.html.j2 (the default) = the look report_email_theme picks.
report_web_dest	(empty)	THE REPORT WEBSITE (docs/WEB_REPORTS.md): every report job also writes its report as a web page - off until report_web_dest is set. Works without an email (report_email_to empty). Each run: its page, a copy as the template's latest page, the report's data as JSON beside each (for your own site), and an index of every... (all of it: the role file)
report_web_url	(empty)	the address people open that folder at, e.g. https://reports.yoursite.mil ("" = dest when it is https://)
report_web_theme	portal	the page's look: standard, command, portal or glass (as report_email_theme)
report_web_page	"{template}/{date}_{time}_job-{job}.html"	The names inside the folder, with {names} filled in: {template} (the AAP template's name, lower case with dashes; the play's name outside AAP), {template_id}, {job} (the AAP job; local outside AAP), {status} (ok, warning, critical, info), {date} (2026-10-23), {time} (0605), {year}, {month}, {day} - the time the job wr... (all of it: the role file)
report_web_latest	"{template}/latest.html"	always the template's newest run ("" = none)
report_web_index	index.html	every report's last run ("" = none)
report_web_names	{}	your own {names}, e.g. {site: "{{ site_code }}"} for "{site}/{template}/..."
report_web_json	true	the report's data as .json beside each page and index.json beside the index

Setting	Default	What it does
report_web_ca_path	(empty)	a CA file for an https:// report_web_dest, or report_web_host's WinRM certificate ("" = what the EE trusts)
report_web_host	(empty)	Or on a Windows server (a file server, or the IIS server itself): report_web_host is that server, report_web_dest the folder on it (D:\reports, there already). The pages are made on the AAP side and written there over WinRM, in one PowerShell script per run (Constrained Language Mode too): nothing to mount, no WebDAV.... (all of it: the role file)
report_web_host_connection	winrm	winrm or psrp
report_web_host_port	5986	WinRM: 5986 (https), 5985 (http)
report_web_host_scheme	https	https, or http (Kerberos and NTLM encrypt each message then)
report_web_host_transport	ntlm	ntlm, kerberos (Kerberos in the EE), credssp or basic
report_web_host_cert_validation	validate	validate (the CA: report_web_ca_path) or ignore

site_findings

File: roles/site_findings/defaults/main.yml

> **Note**
> Shared by every check playbook. A "finding" is one thing that is wrong on one host: check: the check that found it (disk, services, mariadb, ...) id: a stable name for this problem on this host, e.g. disk:/var. ServiceNow uses it to update the same ticket instead of opening a new one every run. severity: critical \| warning summary: one line a person understands ("/var is 96% full (critical at 95%)") hint: a read-only command that shows more. A person can paste it into a shell; ACT (when you add it later) runs it to collect evidence. Check roles only ADD findings. This role starts the list, reports it, and decides pass/fail.

Setting	Default	What it does
site_fail_on	[critical, warning]	Which severities make the job FAIL. A failed job is what sends an AAP workflow down its "on failure" path (open tickets, ask ACT, wait for approval). [] = report only, never fail.

Setting	Default	What it does
site_unreachable_fails	false	A host that cannot be reached (switched off, no network, ssh or WinRM refused, its login refused) in the health checks, the certificate report and renewal, database health, the service watch and troubleshoot: listed as NOT REACHED and left out, and the job is not failed for it. It still counts as checked-for, so the S... (all of it: the role file)
site_syslog	true	One syslog line per finding (tag below), so your log server / SIEM sees them too.
site_syslog_tag	site-health	the syslog tag those lines carry (to filter on in the SIEM)
site_eventlog_source	site-health	Windows: one Application event log entry per finding instead of syslog (source below).
site_os	linux	---- Linux or Windows (set by the playbook; do not set) ---------- Linux playbooks run the roles check_<name>, the Windows ones (win_*.yml) the roles win_check_<name>.
site_check_role_prefix	"{{ 'win_check_' if site_os == 'windows' else 'check_' }}"	check_ (Linux) or win_check_ (Windows): set by the playbook from site_os
site_windows_connections	[winrm, psrp, ansible.builtin.winrm, ansible.builtin.psrp]	How a host is recognized as Windows: its connection, or its shell (Windows OpenSSH).
site_settings_notice	true	A notice in the job output when a project settings file (playbooks/group_vars/*.yml) defines a setting: it wins over AAP's inventory Variables box. false = no notice.

stigman_stack

File: roles/stigman_stack/defaults/main.yml

> **Note**
> STIG Manager stack on podman, managed by systemd through Quadlet unit files: stigman-mysql (MySQL 8.4) <- stigman-api (STIG Manager) <- stigman-nginx (TLS, port 443) Keycloak is EXTERNAL: the realm already exists and this role never touches it. Needs podman >= 4.4 (Quadlet): RHEL 9.2+ / RHEL 8.8+.

Setting	Default	What it does
stigman_scope	system	system = rootful (/etc/containers/systemd, systemctl) <- production user = rootless (~/.config/containers/systemd, systemctl --user) <- tests / rootless hosts

Setting	Default	What it does
stigman_images	mysql: docker.io/library/mysql:8.4 api: docker.io/nuwcdivnpt/stig-manager:1.6.17 # pin a…	the container images (pin STIG Manager's version; change it deliberately)
stigman_config_dir	/etc/stigman	nginx.conf, TLS files, CA bundle
stigman_https_port	443	host port nginx publishes
stigman_server_name	"{{ inventory_hostname }}"	the name nginx answers to (and the TLS certificate is for)
stigman_oidc_provider	"https://keycloak.example.mil/realms/stigman"	Existing Keycloak realm (server-side URL used by the API; browser-side URL for the web client)
stigman_client_oidc_provider	"{{ stigman_oidc_provider }}"	the Keycloak realm URL the browser uses (when it differs from the server-side one)
stigman_client_id	stig-manager	the Keycloak client the web app logs in with
stigman_classification	U	banner: U, CUI, C, S, TS, SCI, NONE
stigman_db_schema	stigman	the MySQL database STIG Manager uses
stigman_db_user	stigman	the MySQL account STIG Manager uses (its password: the credential)
stigman_mysql_buffer_pool	1G	STIG Manager docs: >= 8G for real deployments
stigman_ready_retries	40	How long to wait after (re)starting: for MySQL to be healthy, then for the API to answer through nginx (each)
stigman_ready_delay	6	seconds between those tries
stigman_mysql_root_password	"{{ lookup('ansible.builtin.env', 'STIGMAN_MYSQL_ROOT_PASSWORD') }}"	---- Secrets: NEVER set these in a vars file or inventory ---------- They come from AAP credentials (aap/credential_types/*.yml) as environment variables / files on the controller (execution environment). See docs/SECRETS.md.
stigman_db_password	"{{ lookup('ansible.builtin.env', 'STIGMAN_DB_PASSWORD') }}"	from the credential (STIGMAN_DB_PASSWORD), never set here
stigman_tls_cert_src	"{{ lookup('ansible.builtin.env', 'TLS_CERT_FILE') }}"	nginx's certificate file, from the TLS certificate credential (TLS_CERT_FILE)
stigman_tls_key_src	"{{ lookup('ansible.builtin.env', 'TLS_KEY_FILE') }}"	its private key file, from the credential (TLS_KEY_FILE)

Setting	Default	What it does
stigman_tls_ca_src	"{{ lookup('ansible.builtin.env', 'TLS_CA_FILE') }}"	the CA bundle file, from the credential (TLS_CA_FILE)

troubleshoot

File: roles/troubleshoot/defaults/main.yml

Setting	Default	What it does
ts_area	overview	TROUBLESHOOT: the commands an experienced admin runs first for a kind of problem, run for you (read-only), each with a line saying what it shows. The matching health checks run too, so you also get plain-words findings. Pick the area; for "service" and "network" say what. ts_area: overview | disk | performance | servi... (all of it: the role file)
ts_service	(empty)	the systemd unit, for ts_area service (see above)
ts_target	(empty)	host:port, for ts_area network
ts_since	"-2h"	how far back logs are read (journalctl --since)
ts_output_lines	60	keep the last this-many lines of each command
ts_areas	overview: checks: [disk, mounts, services, performance, time, logging] commands: - {why:…	What each area runs. why is printed above the command's output. Commands must be read-only.

vmware_vm

File: roles/vmware_vm/defaults/main.yml

> **Note**
> VMWARE VMs: jobs that run against vCenter (not on the VMs), with the certified vmware.vmware collection. vCenter address and account come from a "VMware vCenter" credential on the job template (VMWARE_HOST, VMWARE_USER, VMWARE_PASSWORD). docs/VMWARE.md

Setting	Default	What it does
vm_names	[]	The VMs, by their exact name in vCenter: a list, or (from a survey) one per line. No wildcards.
vmware_datacenter	(empty)	only needed when the same VM name exists in two datacenters

Setting	Default	What it does
vmware_validate_certs	true	Check vCenter's certificate. false only if the execution environment does not trust your CA yet.
vmware_vcenters	[]	Several vCenters (e.g. Enhanced Linked Mode), one report: the reports and the ESXi security settings read the credential's vCenter AND these (names, or name:port), with the same account. The same vCenter under two names is read once. docs/VMWARE.md, "Several vCenters". e.g. [vc01.yoursite.mil, vc02.yoursite.mil]
vmware_vcenter_fail_unread	false	true = the job shows failed when a vCenter could not be read (the report lists it in red either way)
vmware_protected_groups	[aap, aap_hosts]	The AAP server is ALWAYS refused (hosts in the groups aap / aap_hosts, matched by VM name, guest host name and IP address), and so are the hosts of these groups and these VM names (* and ? allowed: DC*). List the AAP VM's name here when it differs from its host name, and vCenter's own VM. The snapshot cleanup never de... (all of it: the role file)
vmware_protected_vms	[]	e.g. [AAP01, VCSA01, DC*] (* and ? allowed)
vm_power_mode	guest	guest = ask the operating system (needs VMware Tools running; services stop cleanly) hard = reset / power off, like the button on a physical server (unsaved data is lost)
vm_power_timeout	600	shut down: seconds to wait for the VM to be off
vm_snapshot_name	(empty)	The snapshot's name. Taking one: empty = aap-<date>-<time>. Deleting: the exact name to delete.
vm_snapshot_description	(empty)	the AAP job number and user are added
vm_snapshot_quiesce	false	true = VMware Tools freezes the file systems first
vm_snapshot_delete_all	false	delete: true = every snapshot of the VM (vm_snapshot_name empty)
vm_notes	(empty)	The text to write.

Setting	Default	What it does
vm_notes_mode	append	append (a new line, with date and user) | replace | clear
vm_nic	1	Which adapter: 1 = Network adapter 1, 2, ...
vm_portgroup	(empty)	the port group (standard or distributed) to move it to
vm_secure_boot_folder	(empty)	Only the VMs in this vCenter folder ("" = every VM, or the ones in vm_names).
vm_secure_boot_fail	false	The job stays green (successful) when it finds VMs to fix: the report and email say what. true = mark the job FAILED when a VM is BIOS or has secure boot off, so a workflow's "on failure" path can react (open a ticket).
vm_datastore_warn_pct	85	Used space (% of capacity): a datastore at or above these is listed as a warning / critical.
vm_datastore_crit_pct	90	a datastore this % used or more is red (critical); amber from vm_datastore_warn_pct
vm_datastore_fail	false	The job stays green when datastores are full: the report and email say which. true = mark the job FAILED when one is at or above vm_datastore_warn_pct (or inaccessible), for workflows.
vm_snapshot_max_age_days	5	Snapshots at least this many days old: listed apart by the report, deleted by the cleanup.
vm_snapshot_keep_regex	'(?i)keep|do.?not.?delete'	Never deleted by the cleanup (the report lists them, with the reason): - a snapshot whose name or description matches this
vm_snapshot_cleanup_exclude_vms	[]	- a snapshot of these VMs: names, * and ? allowed (e.g. [DC*, VCSA01, AAP01]). The VMs in vmware_protected_vms and the AAP server are never cleaned up either.
vm_snapshot_cleanup_exclude_datacenters	[]	- a snapshot of a VM in these vCenter datacenters: names, * and ? allowed, any case (e.g. [DRSITE])

Setting	Default	What it does
vm_snapshot_cleanup_exclude_clusters	[]	- a snapshot of a VM running in these clusters: names, * and ? allowed, any case (e.g. [SQL*, MGMT*]). A VM on a host that is in no cluster is not matched.
vm_snapshot_cleanup_name_words	[]	- with words here, a snapshot whose name contains none of them: only a name with one is deleted (any case; Patch also matches Patching and Patches - and dispatch). E.g. [Patching, Patches, Patch]: only the patching snapshots. [] = any name. Set it in all.yml, so the report lists the same.
vm_snapshot_cleanup_max_size_gb	1024	- a snapshot bigger than this many GB (1024 = 1 TB), or whose size cannot be read: merging a very big snapshot can take hours and slow the VM - do those by hand, at a quiet time.
vm_snapshot_cleanup_confirm	false	The cleanup on its own (not after a workflow approval) deletes only when this is true - e.g. a survey question "Delete them?". false = it lists what it would delete, and deletes nothing.
vm_snapshot_cleanup_max	50	Refuse to delete more than this many snapshots in one run (a typo in the age would otherwise delete them all). Raise it on purpose for a big cleanup.
vm_snapshot_delete_timeout	1800	seconds to wait for one deletion (vCenter merges it into the disk)
vm_snapshot_cleanup_email_only_if_deleted	false	The cleanup's email (report_email_to): true = only when it deleted something (or a deletion failed).
vm_pick_snapshot_templates	[]	vm_pick_refresh.yml (on a schedule) writes what may be done right now as the choices of one multiple-choice question in the survey of these templates: job templates or workflow job templates, by name (or {template: name or ID, kind: job or workflow} when a job template and a workflow share the name). In a workflow onl... (all of it: the role file)
vm_pick_snapshot_question	vm_snapshot_picks	the question's Answer variable name

Setting	Default	What it does
vm_pick_reboot_templates	[]	e.g. ["VM - restart picked VMs (workflow)"]
vm_pick_reboot_question	vm_reboot_picks	the question's Answer variable name
vm_pick_max_choices	200	a list never has more choices than this (the rest is counted in the job's output)
vm_pick_snapshot_min_age_days	0	The snapshots offered: only those the snapshot cleanup could delete (the same holds - a name that says keep, the excluded VMs, datacenters and clusters, the protected VMs and the AAP server, bigger than vm_snapshot_cleanup_max_size_gb; not the name words), at least this many days old. 0 = any age.
vm_pick_reboot_vms	[]	The VMs offered for a restart: powered on, never the AAP server or a protected VM (vmware_protected_*), among these (names, * and ?, any case; [] = every VM) ...
vm_pick_reboot_clusters	[]	e.g. [PROD-*]
vm_pick_reboot_folders	[]	folder paths as vCenter shows them, e.g. ["/DC1/vm/Linux*"]
vm_pick_reboot_exclude_vms	[]	... and never these (e.g. [SQL*, DC*])
vm_pick_reboot_only_hung	false	true = only the VMs that look hung: VMware Tools installed, but no heartbeat (red) or Tools not running, powered on for vm_pick_hung_min_minutes at least (a VM still starting has neither yet). Restart those with vm_power_mode: hard - a guest restart needs VMware Tools running.
vm_pick_hung_min_minutes	15	minutes powered on before a VM without a heartbeat counts as hung
vm_pick_aap_ca_path	(empty)	a CA file for AAP's own HTTPS certificate ("" = what the execution environment trusts)
vm_snapshot_picks	[]	The survey's answers (the picked choices: a list, or text one per line). the snapshots to delete
vm_reboot_picks	[]	the VMs to restart

Setting	Default	What it does
vm_pick_confirm	false	vm_snapshot_delete_picked.yml / vm_reboot_picked.yml on their own (not a workflow): true = do it now. false = list exactly what would be done (the plan), and do nothing. Inside a workflow it is never used: the first step plans, the step after the approval does exactly the plan.
vm_pick_max	20	refuse to restart, or delete, more than this many in one run
vm_alarm_types	[vcenter, datacenter, cluster, host]	Whose triggered alarms and configuration issues: vcenter, datacenter, cluster, host. Add vm, datastore or network for theirs too (VM alarms can be many).
vm_alarm_hours	24	The events of the last this-many hours: every error and warning, failed logins, host connection events and VM shutdowns / power-offs / resets / HA restarts. 0 = alarms and host state only.
vm_alarm_max_events	5000	read at most this many events (the newest)
vm_alarm_max_alarm_changes	2000	read at most this many alarm status changes (the newest; 0 = no "Alarms that fired")
vm_alarm_max_rows	200	rows per table in the report (the artifacts have them all)
vm_alarm_login_critical	10	This many failed logins of one user from one place, or more, is critical (fewer: a warning).
vm_alarm_login_event_types	[BadUsernameSessionEvent, com.vmware.sso.LoginFailure, esx.audit.account.locked, esx.audi…	The event types read whatever their category (vCenter's event type names). Add to these lists rather than replacing them; [] = none of that kind.
vm_alarm_connection_event_types	[HostConnectionLostEvent, HostNotRespondingEvent, HostDisconnectedEvent, HostReconnection…	Events about a host losing (or failing to get back) its connection to vCenter: the Host connection table.
vm_alarm_vm_event_types	[VmGuestShutdownEvent, VmGuestRebootEvent, VmPoweredOffEvent, VmResettingEvent, VmFailove…	Events about a VM stopping, restarting or failing over (guest shutdown, power off, HA restart): the VM events table.

Setting	Default	What it does
vm_alarm_fail	false	The job stays green when it finds alarms: the report and email say what. true = mark the job FAILED when there is a critical alarm or a host is not connected, for workflows.
vm_alarm_act_max_items	40	ACT analysis (playbooks/vm_alarm_act_analysis.yml). ACT runs on the AAP side (the execution environment), never on the ESXi hosts; the model provider, model, URL, proxy and CA are the site_act_* settings (roles/site_act/defaults/main.yml), the key an "ACT model key" credential. analyze at most this many problems, the... (all of it: the role file)
vm_alarm_act_config_issues	true	configuration issues too (false = alarms and events only)
vm_alarm_act_email_if_none	false	true = email even when there is nothing to analyze
vm_alarm_act_check_url	true	first check that the AAP node can reach the model's URL
vm_alarm_act_model_timeout	"{{ site_act_model_timeout | default(600) }}"	Seconds the model may take for its answer (ACT's GENAI_TIMEOUT, 90 by default, is too short for a long analysis): site_act_model_timeout (600) unless set here. A GENAI_TIMEOUT in site_act_env wins.
vm_health_require_ha	true	A cluster of 2 or more hosts without vSphere HA is a warning (false = HA off is on purpose).
vm_health_require_drs	true	A cluster of 2 or more hosts without DRS is a warning (false where DRS is off on purpose or not licensed, e.g. vSphere Standard).
vm_health_cpu_warn_pct	85	CPU and memory used, in percent of a host (or of a cluster's hosts running VMs): CPU used: a warning from this percent
vm_health_cpu_crit_pct	95	CPU used: critical from this percent
vm_health_mem_warn_pct	90	memory used: a warning from this percent
vm_health_mem_crit_pct	95	memory used: critical from this percent

Setting	Default	What it does
vm_health_sensors	true	read the hosts' hardware sensors (false = skip them)
vm_health_max_rows	500	rows per table in the report (the artifacts have them all)
vm_health_fail	false	The job stays green when it finds problems: the report and email say what. true = mark the job FAILED when there is a critical problem (a host not connected, an HA agent isolated ...), for workflows.
vm_created_days	7	The window: the VMs created in the last this many days, up to the moment the job runs (7 = the past week; run weekly, each VM is in exactly one report). VMs created and deleted since are listed too, as long as vCenter keeps their events (often 30 days).
esxi_security_hosts	[]	Which hosts: names (* and ? allowed; the short name matches too). [] = every host of vCenter (or of vmware_datacenter).
esxi_security_clusters	[]	only the hosts of these clusters (names, wildcards); [] = all
esxi_security_exclude_hosts	[]	Never these hosts (names, wildcards). E.g. a host you opened SSH on for a support case: the nightly run would otherwise close it again.
esxi_security_ssh	disabled	The SSH service (TSM-SSH): disabled = stopped, and set to start manually; enabled = started, and set to start with the host; "" = leave it as it is.
esxi_security_shell	disabled	The ESXi Shell service (TSM), the host's local shell: the same choices.
esxi_security_settings	UserVars.ESXiShellTimeOut: 600 UserVars.ESXiShellInteractiveTimeOut: 600	Advanced settings, name: value; {} = none. The defaults (seconds): UserVars.ESXiShellTimeOut SSH and the ESXi Shell stop by themselves this long after someone starts them (0 = never) UserVars.ESXiShellInteractiveTimeOut an idle SSH or ESXi Shell session is logged out Others often set: UserVars.DcuiTimeOut (an idle DCU... (all of it: the role file)

Setting	Default	What it does
esxi_security_lockdown	normal	Lockdown mode: normal = the host is managed only through vCenter (the DCUI still works); disabled; "" = leave it as it is. strict is refused, and a host already strict is left alone.
esxi_security_lockdown_exception_users	[]	Accounts that must keep logging in to the hosts directly under lockdown (a vulnerability scanner, a monitoring tool): host-local accounts, or Active Directory accounts with their own permission on the host (not vsphere.local, not an AD group). Added before lockdown is turned on; never removed.
esxi_security_email_only_if_changed	true	The email (report_email_to): true = only when a host was changed or failed; false = every run.
vmware_cert_warn_days	30	A certificate of vCenter (its Machine SSL certificate, the one its web pages show) or of an ESXi host that expires within this many days (or has expired) is due: a ServiceNow incident, and a CSR made where its key stays - vCenter, or the host - attached to the incident with what to do next. To try the job on certifica... (all of it: the role file)
vmware_cert_crit_days	7	within this many days (or expired) the incident is critical, else warning
vmware_cert_esxi	true	the ESXi hosts' certificates too; false = vCenter's only
vmware_cert_only	[]	Only these vCenters and hosts: names, any case, * and ? allowed, the short name matches too - e.g. one host to try the job on. [] = every vCenter of the job and every host connected to it.
vmware_cert_make_csr	true	false = the incidents only, no CSR. A CSR is made once per certificate: the incident is its record, and a second CSR would replace the key that waits on vCenter or the host for the issued certificate.

Setting	Default	What it does
vmware_cert_vmca_csr	false	A certificate signed by vCenter's own CA (VMCA - how vCenter gives ESXi hosts their certificates unless told otherwise) is renewed in vCenter (its Renew button): its incident says so, and no CSR is made. true = a CSR for those too, to replace them with certificates of your PKI.
vmware_cert_vmca_issuers	[]	Issuers taken as VMCA besides VMCA's own root (CN=CA with DC=..., or OU=VMware Engineering): any case, * and ? allowed, e.g. ['*O=vc01.yoursite.mil*'].
vmware_cert_csr_fields	{}	vCenter's CSR: its subject and names are the current certificate's. vCenter's API wants an organization, organization unit, locality, state and country - one the certificate does not have goes empty. Set any of them here (the PKI sets the subject by its own rules anyway), e.g. vmware_cert_csr_fields: {organization: U.... (all of it: the role file)
vmware_cert_csr_without_servicenow	false	Without the ServiceNow API credential no CSR is made (nothing would record it: the next run would make another one). true = make them anyway, and print them in the job output.
vm_capacity_clusters	[]	---- vSphere capacity planning (playbooks/vm_capacity_report.yml) ---- only these clusters (names; * and ? allowed); [] = every cluster
vm_capacity_history_days	90	The history: vCenter's own daily statistics (cluster CPU and memory, datastore space), kept a year at statistics level 1. Read this many days; the growth is a straight line over the last vm_capacity_trend_days, with at least vm_capacity_min_samples daily values.
vm_capacity_trend_days	30	the growth (straight line) is fitted on this many most recent days
vm_capacity_min_samples	7	fewer daily values than this in those days = 'not enough history' (no trend)

Setting	Default	What it does
vm_capacity_events_days	30	VMs added and removed (vCenter keeps events 30 days by default)
vm_capacity_failover_hosts	1	Plan for this many failed hosts per cluster (N+1); a larger HA admission control failover level wins.
vm_capacity_mem_target_pct	80	memory used at most this %, with those hosts failed
vm_capacity_cpu_target_pct	80	CPU used at most this %, with those hosts failed
vm_capacity_vcpu_per_core_max	4	vCPUs per physical core (amber from here, red from 1.5 times it)
vm_capacity_horizon_months	12	Hosts recommended per cluster (of its average host size) to keep memory and CPU at their targets with those hosts failed: now, and at the current growth within this many months.
vm_capacity_storage_warn_pct	85	datastore used %: amber
vm_capacity_storage_crit_pct	90	red, and the target of the storage runway
vm_capacity_runway_warn_days	90	a runway shorter than this is amber
vm_capacity_runway_crit_days	30	and red
vm_capacity_include_local_datastores	false	a host's own (not shared) datastores count too
vm_capacity_max_rows	100	rows per table (problems first); the artifacts have all
vm_capacity_fail	false	The job stays green: the report says what is short. true = mark it FAILED when a cluster or datastore is critical (for workflows).
vm_capacity_act	false	GenAI: ACT reads the numbers and makes its own estimate, shown next to the job's (math vs GenAI). Needs an "ACT model key" credential on the template; the same site_act_* settings as the alarms analysis. ACT runs on the AAP side and runs no command.
vm_capacity_act_max_items	30	items ACT estimates: everything overall, then clusters and datastores, problems first

Setting	Default	What it does
vm_capacity_act_model_timeout	"{{ site_act_model_timeout | default(600) }}"	seconds the model may take for its answer
vm_capacity_act_check_url	true	first check that the AAP node can reach the model's URL
vm_deploy_ticket	(empty)	---- Deploy a VM from a ServiceNow ticket (playbooks/vm_deploy_from_ticket.yml) ---- The ticket (survey answer): INC..., TASK..., SCTASK..., RITM..., CHG... - or your instance's own prefixes (INCN..., CHGN...): the job asks ServiceNow which table a ticket is in. The ticket says what to build in 'key: value' lines - in... (all of it: the role file)
vm_deploy_ticket_table	(empty)	leave it empty (the job finds each ticket's table); a table of your own, e.g. u_vm_request: for numbers whose prefix the job does not know (CTASK stays change_task)
vm_deploy_ticket_prefixes	{}	Your instance's own number prefixes and their tables, e.g. {INCN: incident, CHGN: change_request}. Rarely needed: for a prefix it does not know, the job asks ServiceNow's task table (and takes INCN as INC, CHGN as CHG when it may not read that table). A table as people say it is understood: change / change request ->... (all of it: the role file)
vm_deploy_parent_tables	[change_request]	The request may also be on the ticket's parent - e.g. a task (SCTASK, CTASK) whose change holds the requirements: the parent's 'key: value' lines (and, with vm_deploy_plain_words, its plain words) are read too, under the ticket's own (what the ticket says wins). A parent that cannot be read, is closed, or is not ready... (all of it: the role file)
vm_deploy_parent_fields	[change_request, parent]	The ticket's fields that point to its parent - the first one set is used: a change task's change_request, any ticket's parent (or your instance's own field, e.g. u_change):

Setting	Default	What it does
vm_deploy_parent_text_fields	[description]	The parent's fields with the request, in this order (a later one wins), e.g. [description, implementation_plan]; comments and work_notes may be listed too (oldest first). GenAI reads them as well.
vm_deploy_parent_wins	false	What wins, from the least to the most: the parent (the change), the ticket's short description and description, its work notes, its comments (each oldest first) - a 'key: value' line and plain words alike: the later one wins, however each is written. So the change task wins over its change, and a comment wins over a w... (all of it: the role file)
vm_deploy_notes_also	[]	Every note and comment this job writes on the ticket goes on these too (closing never does: only the ticket itself is closed): parent = the change a task belongs to (when it was read); tasks = when the job is launched with a change, its open change tasks assigned to vm_deploy_assignment_group (and matching vm_deploy_n... (all of it: the role file)
vm_deploy_notes_tasks_query	(empty)	An extra ServiceNow query for the change's tasks that get the notes, e.g. "short_descriptionLIKEto deploy" ("" = the group's)
vm_deploy_notes_tasks_limit	50	the change's tasks that get the notes (and are closed), at most - more is a warning
vm_deploy_close_tasks	false	Launched with a change: also close its tasks found that way (with vm_deploy_done_fields_by_table's change_task entry, when every VM is built) - the change itself is not closed by this. They get the notes too.

Setting	Default	What it does
vm_deploy_ready_by_table	change_request: {approval: Approved, state: Implement}	Before anything is built, a ticket of these tables - the ticket itself, or the parent it belongs to (a change task's change) - must have these fields: one value or a list (any of them), written as the ticket's form shows them (upper / lower case ignored). Checked by the dry run, the propose step AND right before the b... (all of it: the role file)
vm_deploy_templates	[]	What a ticket may ask for (names; * and ? allowed). Templates MUST be listed: nothing else is deployed. The others: [] = any. e.g. [rhel9-gold, win2022-gold]
vm_deploy_clusters	[]	clusters (and standalone hosts) a ticket may use; also the candidates of automatic placement
vm_deploy_datastores	[]	datastores (or datastore clusters) a ticket may use; also the candidates of automatic placement
vm_deploy_folders	[]	VM folders a ticket may use (paths below the datacenter's VM folder, e.g. Linux/*)
vm_deploy_max_cpu	16	the most vCPUs a ticket may ask for
vm_deploy_max_memory_gb	128	the most memory a ticket may ask for, in GB
vm_deploy_max_vms	10	A ticket may ask for several VMs ("2 VMs named web01 and web02", "vm_name: web01, web02", or "Hostnames:" and one name per line): one VM each, the same size - or each its own, from a line of its own - built one after the other, the ticket closed when all are built. More names than this: refused, nothing built (the lim... (all of it: the role file)
vm_deploy_name_regex	'^[A-Za-z0-9][A-Za-z0-9-]{0,62}$'	The VM names allowed (a regular expression); e.g. '^[a-z0-9-]{1,15}$' for Windows-safe names.
vm_deploy_assignment_group	(empty)	Only tickets assigned to this group (its name as ServiceNow shows it), or to one of a list of groups (["Group A", "Group B"]); "" = any group (the pickup needs one).

Setting	Default	What it does
vm_deploy_assignment_group_by_table	{}	The same for the ticket the job is launched with, by its table - in place of vm_deploy_assignment_group for that table; "" = any group. e.g. {change_request: ""}: launched with a change, the change may be assigned to any group (it must still be approved: vm_deploy_ready_by_table); its tasks that get the notes (vm_depl... (all of it: the role file)
vm_deploy_defaults	{}	What a ticket does not say: e.g. {cluster: PROD-CL01, datastore: DS-CLUSTER-01, folder: Linux/New}. Not set = the template's own cluster, datastore and folder; its CPU and memory.
vm_deploy_sites	{}	Sites: a ticket's 'site: A' line (or its Location field, or with plain words a site's name or its words in vm_deploy_plain_sites - each VM by its own line) picks the placement - over vm_deploy_defaults and the name rules; a cluster / datastore / folder the ticket names itself still wins. A site in another vCenter: its... (all of it: the role file)
vm_deploy_site_required	true	with sites set up: a ticket must say its site (false = vm_deploy_defaults); a name rule counts too
vm_deploy_name_rules	[]	Name rules: the VM's name picks where it goes - the first rule whose 'match' fits the name (* and ?, upper / lower case alike) fills in what the ticket does not say: a cluster (and datastores, folder ...) or a site. The ticket's own lines, and a site the ticket names, still win. A name used at several sites: 'sites' -... (all of it: the role file)
vm_deploy_pxe_templates	[]	Templates with no operating system yet, that install it themselves: they boot from the network (PXE, e.g. an MECM task sequence) or from their CD / ISO (e.g. a kickstart). No guest customization, no waiting for an IP address; the MAC address and the computer name go on the ticket. e.g. [win2022-pxe]

Setting	Default	What it does
vm_deploy_guestinfo	true	Give a VM from one of those empty templates the ticket's name and number as VM settings its installer can read: guestinfo.hostname and guestinfo.ticket (a kickstart with VMware Tools in the installer: vmtoolsd --cmd "info-get guestinfo.hostname"). Needs the vCenter privilege Change configuration > Advanced configurati... (all of it: the role file)
vm_deploy_template_aliases	{}	Short names for long template names - for the ticket's 'template:' line and GenAI's catalog (whose keys may be these short names): {short name: the template's name in vCenter}, e.g. {win2022: "Windows 2022 Standard Template - Site 1", rhel9: "RHEL 9 Server Template"} A ticket may also write a template's own name in ot... (all of it: the role file)
vm_deploy_template_specs	{}	The customization specification (vCenter > Policies and Profiles) for each template, when the ticket, its site and the name rules name none: {template name or pattern: spec}. A Windows template that is not network-boot needs one (Sysprep: administrator password, domain or workgroup), e.g. {"Windows 2022*": Win2022-Sys... (all of it: the role file)
vm_deploy_max_disks	5	disks a ticket may ask for, disk 1 included (disks_gb: 100, 200 = disk 1 grown to 100 GB, one 200 GB disk added)
vm_deploy_max_disk_gb	2048	the biggest disk a ticket may ask for, in GB
vm_deploy_max_total_disk_gb	4096	all its disks together, in GB

Setting	Default	What it does
vm_deploy_mem_target_pct	80	Automatic placement - when the ticket, its site and vm_deploy_defaults name no cluster / datastore: the job chooses among a site's 'clusters' / 'datastores' lists (e.g. SiteA: {clusters: [A-PROD-*]}), else among vm_deploy_clusters / vm_deploy_datastores; with neither, the template's own. a cluster fits when its memory... (all of it: the role file)
vm_deploy_failover_hosts	1	hosts to plan as failed when choosing a cluster
vm_deploy_datastore_max_used_pct	85	a datastore fits when it stays under this % with the VM at its full size
vm_deploy_datastore_rules	[]	Datastores by the VM's name: the first rule whose match fits the name (* and ?, any case) gives the datastores it goes on - automatic placement chooses among them on its cluster, the least used first, over a site's or vm_deploy_defaults' datastore (one the ticket names must be one of them). A VM no rule matches never... (all of it: the role file)
vm_deploy_datastore_balance	false	true = a VM whose ticket, site, name rule and vm_deploy_defaults name no datastore (and no vm_deploy_datastores) goes on the least used shared datastore of its cluster - the location's - the VMs of one ticket spread out (each counts the ones before it); false = the template's own datastore.
vm_deploy_resume	true	Run again for the same ticket (a job that stopped: failed, cancelled, timed out), pick up where it stopped: each VM carries its ticket in its notes from the clone on - one it built is left as it is, one it started and did not finish (cloned, not yet powered on) is finished where it is; the work notes say which. false... (all of it: the role file)
vm_deploy_domain	localdomain	the DNS domain Linux customization sets (a ticket's domain: wins)
vm_deploy_customize	true	guest customization: host name = vm_name, every adapter on DHCP

Setting	Default	What it does
vm_deploy_power_on	true	power the new VM on (false = leave it off, e.g. to check it first)
vm_deploy_wait_customization	900	seconds to wait for vCenter's customization result
vm_deploy_wait_ip	600	seconds to wait for an IPv4 address (DHCP, through VMware Tools)
vm_deploy_done_fields	{}	Fields set on the ticket when the VM is up, e.g. to resolve an incident: {state: "6", close_code: "Solution provided", close_notes: "VM deployed"} (your instance's values)
vm_deploy_done_fields_by_table	{}	The same for one table (the ticket's table: incident, sc_task, change_task ...), in place of vm_deploy_done_fields for it - e.g. close the task, never the change it belongs to: {sc_task: {state: "3", close_notes: "VM deployed by AAP"}} (3 = Closed Complete; your instance's values) Sent on its own, after the result not... (all of it: the role file)
vm_deploy_done_fields_issues_by_table	{}	Over those fields when every VM is built but something did not finish: no IPv4 address in time, no customization result, the guest's host name not applied, no network at power on, guestinfo not set - e.g. {change_task: {close_code: successful_issues}} (your instance's values; {} = closed the same way). What happened i... (all of it: the role file)
vm_deploy_done_fields_pxe	true	Also for a VM from a network-boot template (vm_deploy_pxe_templates), which is powered on and still installing its OS from MECM when the job ends. false = that ticket stays open for whoever checks the install.

Setting	Default	What it does
vm_deploy_service	(empty)	---- The server's service, picked at launch: words in the VM's notes, an AAP inventory and group ---- What the new server is for, picked by whoever launches the deploy: a survey question (multiple choice, NOT required - the scheduled pickup starts the workflow without it), e.g. sql or dc. Every VM of the ticket gets i... (all of it: the role file)
vm_deploy_services	{}	The services the survey offers: each one's AAP inventory and group, and the words for the VM's notes (note: left out = "Service: <its name>."). {} = any service is only written in the notes. vm_deploy_services: sql: {inventory: Windows servers, group: win_sql, note: SQL Server} dc: {inventory: Windows servers, group:... (all of it: the role file)
vm_deploy_inventory_name	"{name}"	The name a VM gets in the inventory: write it the way your inventory already names servers, or the same server shows up twice in AAP. {name} = the VM's name, e.g. "{name}.yoursite.mil". A service's own name: wins (e.g. {inventory: Windows servers, name: "{name}.corp.yoursite.mil"}).
vm_deploy_inventory_lowercase	true	the inventory name in lower case (APP01 -> app01); false = as the ticket writes it
vm_deploy_aap_ca_path	(empty)	a CA file for AAP's own HTTPS certificate ("" = vm_deploy_poll_aap_ca_path, else what the execution environment trusts)
vm_deploy_plain_words	false	---- A request in plain words, without GenAI (off by default) ---- true = also read the ticket's own words, by the rules below, so a requester can write "a Red Hat 9 server named ab01web05, medium" instead of 'key: value' lines. Never a guess: two different answers in the text are a question to the requester (a commen... (all of it: the role file)

Setting	Default	What it does
vm_deploy_plain_templates	{}	The words for each template (any case; "windows 2022" also matches "Windows Server 2022", "server 2025" also "a new 2025 VM"). An allowed template's own name and its short names (vm_deploy_template_aliases) count too. An OS no template is for (Server 2025, only a 2022 template) is a question - never another version; t... (all of it: the role file)
vm_deploy_plain_sizes	{}	Size words and what they mean, e.g. {small: {cpu: 2, memory_gb: 8}, medium: {cpu: 4, memory_gb: 16}, large: {cpu: 8, memory_gb: 32}}. "4 vCPU" and "16 GB RAM" in the ticket are read too, and win in the same text.
vm_deploy_plain_sites	{}	The words for each site (vm_deploy_sites), as people write the place - any case, the spaces optional ("Site Two" is also "SiteTwo"). A site's own name counts too. A change that lists its servers one per line ('new VM / ab01web05 / ... Location: Site Two') puts each VM at the site its own line names; the CPU, memory an... (all of it: the role file)
vm_deploy_genai	false	---- GenAI reads a request written in plain words (a proof of concept, off by default) ---- When the ticket's own 'key: value' lines do not give vm_name and template: ACT (GenAI, evidence mode, no commands) copies out what the ticket says - the names, the operating system as written, CPU, memory, disks - each with the... (all of it: the role file)
vm_deploy_genai_in_check	false	A dry run (Check) asks GenAI too (true): it shows GenAI's proposal and where the VM would go, and writes nothing on the ticket - no proposal note, no questions to the requester. false = Check only says GenAI would be asked.

Setting	Default	What it does
vm_deploy_template_catalog	{}	Optional: templates GenAI may name by key, each with a plain description (the key: the template's exact name). Not needed when the templates' names say their OS: the job finds the template from the OS the ticket names. A key for another version than the ticket asks for is dropped. e.g. {win2022-pxe: "Windows Server 20... (all of it: the role file)
vm_deploy_genai_approval	true	true: a GenAI proposal waits for a person - the workflow's approval (or a second run) - before it is built. false: one run reads, checks and builds, when every value is traced to the ticket and within the rules (questions or a refusal still stop it).
vm_deploy_phase	(empty)	Which step this run is: "" = one job does it all; propose = the workflow's first step (read, ask GenAI if needed, check, write where it would go on the ticket - build nothing); build = the step after the approval (build what the ticket and its proposal say - never ask GenAI).
vm_deploy_genai_model_timeout	"{{ site_act_model_timeout | default(600) }}"	seconds the model may take for its answer
vm_deploy_poll_enabled	true	---- Pick up VM request tickets on a schedule (playbooks/vm_deploy_poll.yml) ---- Every run (an AAP schedule, e.g. every 15 minutes) asks ServiceNow for NEW VM requests - active, assigned to vm_deploy_assignment_group (required), marked by vm_deploy_poll_prefix and/or vm_deploy_poll_query - and starts the deploy workf... (all of it: the role file)
vm_deploy_poll_tables	[incident]	the ServiceNow tables to look in: incident, sc_task, task ...
vm_deploy_poll_prefix	"VM request"	a VM request's short description starts with this ("" = mark them by vm_deploy_poll_query)

Setting	Default	What it does
vm_deploy_poll_query	(empty)	An extra ServiceNow encoded query, ANDed with the rest, e.g. "category=vm_request" or "cmdb_ci.name=VM Request" ("" = none). ^NQ (a second query) is refused.
vm_deploy_poll_states	["1"]	the ticket states to pick up (state values: "1" = New)
vm_deploy_poll_max	3	the most tickets started per run; the rest wait for the next run
vm_deploy_poll_scan	100	the most tickets read per table and run, oldest first (tickets AAP already has count too)
vm_deploy_poll_workflow	"VM - deploy from ServiceNow ticket (workflow)"	the workflow job template started per ticket
vm_deploy_poll_dry_run	false	true = only list what it would pick up: no note, no workflow (Check does the same)
vm_deploy_poll_claim_fields	{}	Fields set on the ticket together with the "Picked up" note, e.g. {state: "2"} (In Progress) - so the ticket leaves the states picked up as well ({} = only the note).
vm_deploy_poll_aap_ca_path	(empty)	a CA file for AAP's own HTTPS certificate ("" = what the execution environment trusts)

win_act

File: roles/win_act/defaults/main.yml

Setting	Default	What it does
win_act_dir	'C:\ProgramData\act'	ACT FOR WINDOWS on the servers (vendor/act-windows/act.ps1, the PowerShell port of ACT). Installed into win_act_dir. That folder gets its own owner and permissions - Administrators and SYSTEM only, nothing inherited - because ordinary users may create folders and files in C:\ProgramData (and C:\temp): a non-admin coul... (all of it: the role file)
win_act_state	present	present | absent (removes win_act_dir)

Setting	Default	What it does
win_act_source	"{{ role_path }}/../../vendor/act-windows/act.ps1"	The ACT script this repository ships (vendor/act-windows, version in its VERSION file)

win_check_accounts

File: roles/win_check_accounts/defaults/main.yml

> **Note**
> WINDOWS ACCOUNTS (local accounts; domain controllers have none and are skipped) - the built-in Guest account is disabled -> critical if not - enabled local accounts: password never expires, password too old, no logon for a long time -> warning - members of the local Administrators group that you did not allow -> warning Names below allow wildcards (* and ?), e.g. 'YOURDOMAIN\Domain Admins', '*\svc_aap'.

Setting	Default	What it does
win_check_accounts_password_max_age_days	60	an enabled local account's password older than this -> warning (0 = off)
win_check_accounts_inactive_days	35	an enabled local account with no logon for this many days -> warning (0 = off)
win_check_accounts_password_never_expires_allowed	[]	accounts allowed "password never expires" (wildcards)
win_check_accounts_inactive_allowed	[]	accounts allowed to stay unused (wildcards)
win_check_accounts_admins_allowed	[]	Empty = list the members in the job output only. Set it to report anyone else, e.g. win_check_accounts_admins_allowed: ['*\Administrator', 'YOURDOMAIN\Domain Admins', 'YOURDOMAIN\svc_aap']

win_check_audit

File: roles/win_check_audit/defaults/main.yml

> **Note**
> WINDOWS AUDIT: the audit policy records what the STIG asks for, the event logs are big enough, and the log forwarding agent runs. One finding lists every missing audit setting.

Setting	Default	What it does
win_check_audit_subcategories	- {name: Credential Validation, guid: 0CCE923F-69AE-11D9-BED3-505054503030, need: Success…	Advanced audit policy subcategories (checked by GUID, so any Windows language works). need: Success | Failure | Success and Failure

Setting	Default	What it does
win_check_audit_security_log_min_kb	196608	STIG
win_check_audit_application_log_min_kb	32768	the Application log's maximum size must be at least this (KB)
win_check_audit_system_log_min_kb	32768	the System log's maximum size must be at least this (KB)
win_check_audit_forwarder_services	[]	Log forwarding agents that must run, e.g. [SplunkForwarder] or [nxlog]

win_check_certs

File: roles/win_check_certs/defaults/main.yml

> **Note**
> WINDOWS CERTIFICATES: every certificate in the machine stores below, and which of them are in use (https bindings of IIS / WSUS / ..., the WinRM HTTPS listener AAP connects through, Remote Desktop). Expiring soon -> warning / critical. Expired: critical when in use, or when it expired less than win_check_certs_expired_report_days ago (older leftovers are only listed in the certificate report).

Setting	Default	What it does
win_check_certs_warn_days	30	a certificate expiring within this many days -> warning (amber in the certificate report)
win_check_certs_notice_days	60	the certificate report colours these blue (an early heads-up; not a finding)
win_check_certs_crit_days	7	within this many days -> critical
win_check_certs_expired_report_days	30	an expired certificate NOT in use is a finding for this many days after it expired; older ones only in the report
win_check_certs_stores	['Cert:\LocalMachine\My', 'Cert:\LocalMachine\WebHosting', 'Cert:\LocalMachine\Remote Des…	the machine certificate stores read
win_check_certs_trust	false	also every other machine store: Trusted Root, Intermediate CAs, Other People, ... (the certificate report turns it on); with its private key checked as above, without it only listed (trust store)

Setting	Default	What it does
win_check_certs_trust_skip_stores	['Disallowed', 'AuthRoot', 'REQUEST', 'SMS']	stores never read: the untrusted certificates, the roots Windows Update keeps, pending requests (a CSR's key), the MECM client's own
win_check_certs_remote_desktop	false	false = Remote Desktop's certificates are left out of the report and the renewal (Windows or the GPO renews them): the Remote Desktop store, the one RDP uses, those for Remote Desktop only; true = checked like the others
win_check_certs_ignore_subjects	[]	wildcards, e.g. ['CN=localhost*']
win_check_certs_endpoints	[]	TLS services whose certificate a client actually gets, e.g. win_check_certs_endpoints: - {name: IIS, host: 127.0.0.1, port: 443, servername: web01.yoursite.mil}
win_cert_renew_dir	'C:\ProgramData\site-automation\renew'	---- Renewal: a new key and CSR for each certificate that expires soon (playbooks/win_cert_renew_csr.yml) ---- The certificates the Windows certificate report flags (within win_check_certs_warn_days, or expired) that are the server's own - not a trust store's - get a new key and a CSR with the same subject, names and... (all of it: the role file)
win_cert_renew_key_bits	2048	RSA key size, at least 2048 (a bigger old key keeps its size; an EC certificate gets a new key on the same curve)
win_cert_renew_digest	sha256	the CSR's signature: sha256, sha384 or sha512
win_cert_renew_max	10	at most this many NEW CSRs per server in one run (the rest come the next run)
win_cert_renew_self_signed	false	true = a self-signed certificate gets a CSR too, for the PKI to issue a real one in its place (never one only for localhost)
win_cert_renew_issuer	(empty)	regular expression: only certificates whose issuer matches (empty = any CA), e.g. 'O=U\.S\. Government'
win_cert_renew_exclude_issuer	(empty)	regular expression: certificates whose issuer matches are never renewed here (renewed by their own tool)

Setting	Default	What it does
win_cert_renew_only	(empty)	regular expression on the certificate's place, e.g. 'used by https' (empty = every one)
win_cert_renew_exportable	false	true = the new key can be exported (to move it to another server); false keeps it on this one
win_cert_install_rebind	[https, rdp, winrm]	---- Install: the issued certificate (playbooks/win_cert_install_check.yml -> approval -> playbooks/win_cert_install_approved.yml) ---- The PKI's certificate, from the request carrying the CSR (cert_renew_servicenow: catalog), is put together with its pending key by certreq -accept -machine (in Personal); its chain's... (all of it: the role file)
win_cert_install_verify_wait	15	seconds an https port has, after the rebinding, to serve the new certificate (full PowerShell only)
win_cert_restore_backup	(empty)	---- Restore: the bindings back on the old certificate (playbooks/win_cert_restore.yml) ---- a backup's folder name, as the restore's Check lists it
win_cert_restore_only	(empty)	or: the newest backup of each certificate whose thumbprint or subject matches (a regular expression)

win_check_disk

File: roles/win_check_disk/defaults/main.yml

> **Note**
> WINDOWS DISK: how full each local fixed volume (C:, D:, ...) is.

Setting	Default	What it does
win_check_disk_warn_pct	85	% used -> warning
win_check_disk_crit_pct	95	% used -> critical
win_check_disk_min_free_gb	0	also warn below this many GiB free (0 = off): big volumes
win_check_disk_overrides	{}	Different limits for particular volumes, e.g. a database volume kept full on purpose: win_check_disk_overrides: 'D:': {warn: 92, crit: 98}
win_check_disk_ignore	[]	volumes not to check, e.g. ['E:']

win_check_eventlog

File: `roles/win_check_eventlog/defaults/main.yml`

> **Note**
> WINDOWS EVENT LOG, the last win_check_eventlog_hours: - unexpected shutdown or power loss (41, 6008), a crash / blue screen (1001) -> critical - disk and file system errors (disk, Ntfs, storage drivers) -> critical - the Security log was cleared (1102) -> critical - many errors in System + Application (the finding names the top sources) -> warning - a System / Application / Security log that is full and not overwriting -> critical

Setting	Default	What it does
`win_check_eventlog_hours`	`24`	how many hours back the event logs are read
`win_check_eventlog_error_warn`	`50`	this many errors or more in System + Application -> warning

win_check_network

File: `roles/win_check_network/defaults/main.yml`

> **Note**
> WINDOWS NETWORK, from the server itself: - a default gateway and DNS servers are configured; - the names you list resolve, the TCP services you list answer; - domain members: the computer's trust with the domain works (nltest /sc_query). A broken trust is the classic "the trust relationship between this workstation and the primary domain failed" at logon.

Setting	Default	What it does
`win_check_network_dns_names`	`[]`	Names every Windows server must resolve, e.g. [dc01.yoursite.mil, ldap.yoursite.mil]
`win_check_network_tcp`	`[]`	Services every Windows server must reach, e.g. win_check_network_tcp: - {name: LDAPS, host: dc01.yoursite.mil, port: 636} - {name: SIEM, host: siem.yoursite.mil, port: 9997}
`win_check_network_secure_channel`	`true`	domain members: test the computer's trust with the domain (nltest /sc_query)
`win_check_network_timeout_ms`	`3000`	milliseconds for each TCP test

win_check_patching

File: `roles/win_check_patching/defaults/main.yml`

Note
WINDOWS PATCHING (read-only; installing updates is a later runbook) - a restart is pending for installed updates (Windows Update, component servicing, ConfigMgr) - days since the last update was installed (Get-HotFix) - optional: ask Windows Update which updates are missing (slower; needs full PowerShell)

Setting	Default	What it does
`win_check_patching_warn_days`	`35`	last update installed this many days ago or more -> warning
`win_check_patching_crit_days`	`60`	this many days or more -> critical
`win_check_patching_reboot_pending`	`true`	report a restart pending for installed updates
`win_check_patching_search_pending`	`false`	also ask Windows Update which updates are missing (slower)

win_check_performance

File: `roles/win_check_performance/defaults/main.yml`

Note
WINDOWS PERFORMANCE: CPU (average of three samples a second apart), memory, page file. The finding names the busiest processes, so you know where to look.

Setting	Default	What it does
`win_check_performance_cpu_warn_pct`	`90`	CPU busy % (average of three samples) -> warning
`win_check_performance_cpu_crit_pct`	`98`	CPU busy % -> critical
`win_check_performance_mem_warn_pct`	`90`	memory in use % -> warning
`win_check_performance_mem_crit_pct`	`97`	memory in use % -> critical
`win_check_performance_pagefile_warn_pct`	`80`	page file in use % -> warning

win_check_security

File: `roles/win_check_security/defaults/main.yml`

Note
WINDOWS SECURITY BASELINE (a few high-value STIG items; not a full STIG scan: that is your SCAP / vulnerability scanner)

Setting	Default	What it does
win_check_security_av	auto	Antivirus: auto = check Microsoft Defender when it is the active antivirus (skipped when it is passive or off because another product protects the server); defender = always; none = never.
win_check_security_av_signature_max_days	7	Defender's signatures older than this many days -> finding
win_check_security_av_services	[]	Services of your endpoint protection that must run, e.g. Trellix / McAfee: [masvc, mfemms]
win_check_security_firewall_profiles	[Domain, Private, Public]	must be on
win_check_security_smb1	true	SMBv1 must be off
win_check_security_rdp_nla	true	Remote Desktop must require Network Level Authentication
win_check_security_uac	true	User Account Control must be on
win_check_security_legacy_tls	true	TLS 1.0 / 1.1 must be off for incoming connections

win_check_services

File: roles/win_check_services/defaults/main.yml

> **Note**
> WINDOWS SERVICES - every required service must be running (critical); - an automatic service that stopped WITH AN ERROR (exit code not 0) is a warning. Automatic services that stop cleanly by design (trigger-start, delayed) are not reported; - a service that crashed several times recently (System log events 7031 / 7034).

Setting	Default	What it does
win_check_services_required	[EventLog, RpcSs, Dnscache, W32Time, WinRM, mpssvc]	Must be running on every Windows server. Add your own per group with ..._required_extra, e.g. for IIS servers: win_check_services_required_extra: [W3SVC, WAS]
win_check_services_required_extra	[]	more services that must run, e.g. for a group of IIS servers (see above)
win_check_services_ignore	- gupdate - gupdatem - edgeupdate - edgeupdatem - MicrosoftEdgeElevationService - GoogleU…	Automatic services never to report (wildcards allowed: * and ?)
win_check_services_crash_hours	24	look back this far for unexpected stops

Setting	Default	What it does
win_check_services_crash_warn	3	this many unexpected stops of one service -> warning

win_check_time

File: roles/win_check_time/defaults/main.yml

> **Note**
> WINDOWS TIME: the Windows Time service runs, syncs with a real source (not the local clock), synced recently, and how far off the clock is from its source. Kerberos (domain logons) fails when a clock is 5 minutes off.

Setting	Default	What it does
win_check_time_offset_warn_sec	1	seconds off its time source -> warning
win_check_time_offset_crit_sec	5	seconds off -> critical
win_check_time_max_hours_since_sync	24	last successful sync longer ago than this many hours -> finding

win_check_trust

File: roles/win_check_trust/defaults/main.yml

> **Note**
> WINDOWS TRUST STORES: which still hold a certificate authority that was compromised or that the browsers dropped - the machine certificate stores, the running account's, every Java cacerts and Python certifi bundle found - matched against the trust check's list (roles/check_trust/vars/main.yml, plus check_trust_extra; a severity changed in check_trust_severity). A certificate in the Disallowed store is blocked: listed, no finding. Removal (check_trust_remove: true, on the Windows trust store report): the certificate goes into the Disallowed store and out of the store it was in, Java entries are deleted, certifi bundles cut - each backed up to win_check_trust_backup_dir first. Check mode shows what would go. docs/TRUST_STORES.md

Setting	Default	What it does
win_check_trust_stores	['Cert:\LocalMachine\Root', 'Cert:\LocalMachine\CA', 'Cert:\LocalMachine\AuthRoot', 'Cert…	the certificate stores read (CurrentUser: the account AAP connects with)
win_check_trust_block_stores	['Cert:\LocalMachine\Disallowed', 'Cert:\CurrentUser\Disallowed']	the Disallowed stores: a certificate there is distrusted (blocked)
win_check_trust_java	true	every Java cacerts under the folders below (JKS natively; PKCS12 with the JDK's keytool)

Setting	Default	What it does
win_check_trust_java_dirs	['C:\Program Files\Java', 'C:\Program Files (x86)\Java', 'C:\Program Files\Eclipse Adopti…	folders searched for cacerts (seven levels deep)
win_check_trust_python	true	every certifi cacert.pem under the folders below
win_check_trust_python_dirs	['C:\Program Files\Python*', 'C:\Program Files (x86)\Python*', 'C:\Python*', 'C:\ProgramD…	folders searched for certifi (seven levels deep)
win_check_trust_backup_dir	'C:\ProgramData\site-automation\trust-backup'	the removal's backups: every certificate and file changed, a folder per run
win_check_trust_timeout	60	seconds for keytool

win_patch

File: roles/win_patch/defaults/main.yml

> **Note**
> WINDOWS PATCH through the ConfigMgr client: install what Software Center has for the server (the software updates deployed to it), restart if the updates need it, and prove the server came back: every automatic service that was running before runs again. One server at a time by default; the rollout stops at the first server that fails. Check mode lists and changes nothing.

Setting	Default	What it does
automatic_restarts	false	NO AUTOMATIC RESTARTS unless you say so: this playbook restarts a server only when automatic_restarts is true (the job template's Variables, or the survey) and the updates need it. With automatic_restarts false (the default) it installs the updates and tells you which servers still need a restart; a person restarts th… (all of it: the role file)
win_patch_never_reboot_groups	"{{ patch_never_reboot_groups | default(['aap', 'aap_hosts', 'sat', 'netapp', 'netapp_con…	Servers in these groups are patched but NEVER restarted by this playbook (restart them by hand). Follows the Linux setting (patch_never_reboot_groups) when you set it, so ONE list covers both; the fallback is the Linux role's default (roles/patch/defaults/main.yml; tests/test_protection.py checks it).
win_patch_never_patch_groups	"{{ patch_never_patch_groups | default(['no_patch', 'aap', 'aap_hosts', 'sat', 'netapp',…	Servers in these groups are skipped entirely (e.g. servers patched only by hand).

Setting	Default	What it does
win_patch_scan_first	true	ask the client to rescan and re-evaluate before installing
win_patch_scan_wait_sec	120	let that scan settle
win_patch_timeout_min	120	per round: download + install
win_patch_poll_sec	30	how often to look at the progress
win_patch_max_rounds	2	After a restart, look again: servicing-stack and cumulative updates often reveal more updates.
win_patch_restart_delay_sec	30	shutdown.exe /t
win_patch_restart_timeout_min	30	how long the server may take to come back
win_patch_restart_settle_sec	30	after the port answers again, before trying to log on
win_patch_reconnect_pause_sec	20	between logon attempts after the restart
win_patch_services_grace_min	5	After the restart, give delayed-start services this long before calling one missing.
win_patch_services_ignore	['gupdate*', 'edgeupdate*', 'MicrosoftEdgeElevationService', 'GoogleUpdater*', 'MapsBroke…	Automatic services that come and go on their own (wildcards); not compared before/after.
win_patch_winrm_host	"{{ ansible_host | default(inventory_hostname) }}"	Where AAP reaches WinRM, to see the server go down and come back during the restart
win_patch_winrm_port	"{{ ansible_port | default(5986) }}"	its WinRM port (5986 = HTTPS)

win_troubleshoot

File: roles/win_troubleshoot/defaults/main.yml

Setting	Default	What it does
ts_area	overview	WINDOWS TROUBLESHOOT: the PowerShell commands an experienced Windows admin runs first for a kind of problem, run for you (read-only), each with a line saying what it shows. The matching health checks run too, so you also get plain-words findings. ts_area: overview | disk | performance | service | network | login | tim... (all of it: the role file)

Setting	Default	What it does
ts_service	(empty)	the service name, for ts_area service (see above)
ts_target	(empty)	host:port, for ts_area network
ts_hours	2	how many hours back the event logs are read
ts_output_lines	60	keep the last this-many lines of each command
win_ts_areas	overview: checks: [disk, services, performance, time, eventlog] commands: - {why: "Window…	What each area runs. why is printed above the command's output. Commands must be read-only.